

HP-UX System Administrator's Guide: Logical Volume Management

HP-UX 11i Version 3

Abstract

This document describes how to configure, administer, and troubleshoot the Logical Volume Manager (LVM) product on the HP-UX Version 3 platform.

The *HP-UX System Administrator's Guide* is written for administrators of HP-UX systems of all skill levels needing to administer HP-UX systems beginning with HP-UX Release 11i Version 3.

While many topics in this set apply to previous releases, much has changed in HP-UX 11i Version 3. Therefore, for information about prior releases, see *Managing Systems and Workgroups: A Guide for System Administrators*.



© Copyright 2011 Hewlett-Packard Development Company, L.P.

Legal Notices

Confidential computer software. Valid license from HP required for possession, use or copying. Consistent with FAR 12.211 and 12.212, Commercial Computer Software, Computer Software Documentation, and Technical Data for Commercial Items are licensed to the U.S. Government under vendor's standard commercial license. The information contained herein is subject to change without notice. The only warranties for HP products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. HP shall not be liable for technical or editorial errors or omissions contained herein. UNIX is a registered trademark of The Open Group.

Trademark Notices

VERITAS™ is a registered trademark of Symantec Corporation.

Contents

1 Introduction.....	8
LVM Features.....	8
LVM Architecture.....	9
Physical versus Logical Extents.....	10
LVM Volume Group Versions.....	11
LVM Device File Usage.....	12
Legacy Device Files versus Persistent Device Files.....	12
Cluster-wide Device Special Files.....	13
Naming Conventions for LVM.....	13
Device Number Format.....	15
Version 1.0 Device Number Format.....	16
Version 2.x Device Number Format.....	16
LVM Disk Layout.....	16
Boot Data Reserved Area.....	16
Logical Interface Format Area.....	17
Physical Volume Reserved Area.....	17
Volume Group Reserved Area.....	17
User Data Area.....	17
LVM Limitations.....	18
Shared LVM.....	18
2 Configuring LVM.....	20
Planning Your LVM Configuration.....	20
Setting Up Different Types of Logical Volumes.....	20
Setting Up Logical Volumes for Raw Data Storage.....	20
Setting Up Logical Volumes for File Systems.....	21
File System Logical Volume Guidelines.....	22
Setting Up Logical Volumes for Swap.....	22
Swap Logical Volume Guidelines.....	22
Setting Up Logical Volumes for Dump.....	23
Dump Logical Volume Guidelines.....	23
Setting Up Snapshot Logical Volumes for Backup.....	23
Planning for Availability.....	24
Increasing Data Availability Through Mirroring.....	24
Mirror Write Behavior Control.....	24
Allocation Policy.....	24
Scheduling Policy.....	25
Synchronization Policy.....	25
Synchronizing a Mirrored Logical Volume.....	26
Increasing Disk Redundancy Through Disk Sparing.....	27
Increasing Hardware Path Redundancy Through Multipathing.....	28
Setting Up Multipathing to a Physical Volume (LVM PVlinks).....	28
Planning for Performance.....	29
General Performance Factors.....	29
Memory Usage.....	29
CPU Usage.....	29
Disk Space Usage.....	29
Internal Performance Factors.....	29
Scheduling Policy.....	29
Mirror Write Consistency Cache.....	29
Disk Spanning.....	30
Number of Volume Groups.....	30

Physical Volume Groups.....	30
Snapshots and Performance.....	30
Increasing Performance Through Disk Striping.....	31
Determining Optimum Stripe Size.....	32
Interactions Between Mirroring and Striping.....	33
Increasing Performance Through I/O Channel Separation.....	33
Configuring LVM I/O Timeout Parameters.....	33
Planning for Recovery.....	34
Preparing for LVM System Recovery.....	35
Example Script for LVM Configuration Recording.....	37
3 Administering LVM.....	38
Administration Tools.....	38
Displaying LVM Information.....	40
Information on Volume Groups.....	41
Information on Physical Volumes.....	41
Information on Logical Volumes.....	42
Common LVM Tasks.....	43
Initializing a Disk for LVM Use.....	43
Creating a Volume Group.....	44
Creating the Volume Group Device File.....	44
Creating a Version 1.0 Volume Group.....	44
Creating a Version 2.x Volume Group.....	45
Migrating a Volume Group to a Different Version: vgversion.....	46
Determining the Version of a Volume Group.....	46
Command Syntax.....	47
Review Mode: vgversion -r.....	47
Example of Migrating a Volume Group to a Different Version.....	48
Migration Recovery.....	49
Example Recovery.....	50
Adding a Disk to a Volume Group.....	51
Removing a Disk from a Volume Group.....	51
Creating a Logical Volume.....	52
Creating a Striped Logical Volume.....	52
Creating a Mirrored Logical Volume.....	53
Extending a Logical Volume.....	53
Extending a Logical Volume to a Specific Disk.....	54
Reducing a Logical Volume.....	55
Adding a Mirror to a Logical Volume.....	55
Removing a Mirror from a Logical Volume.....	56
Renaming a Logical Volume.....	56
Removing a Logical Volume.....	57
Exporting a Volume Group.....	57
Importing a Volume Group.....	58
Modifying Volume Group Parameters.....	58
vgmodify for a Version 1.0 Volume Group.....	59
Example: vgmodify for a Version 1.0 Volume Group.....	59
vgmodify for a Version 2.x Volume Group.....	62
Example: vgmodify for a Version 2.x Volume Group.....	63
Quiescing and Resuming a Volume Group.....	65
Renaming a Volume Group.....	66
Splitting a Volume Group.....	67
Removing a Volume Group.....	68
Backing Up a Mirrored Logical Volume.....	68
Backing Up and Restoring Volume Group Configuration.....	69

Moving and Reconfiguring Your Disks.....	70
Moving Disks Within a System.....	71
Moving Disks Between Systems.....	72
Moving Data to a Different Physical Volume	73
Moving a Root Disk to Another Disk.....	73
pvmove Command Syntax.....	74
Moving Data for Disk Space Balancing: Auto Re-balancing.....	75
Creating a Spare Disk.....	76
Reinstating a Spare Disk.....	77
Modifying Physical Volume Characteristics.....	77
Handling Size Increases.....	78
Handling Size Decreases.....	84
Supported LUN Changes per Release.....	84
Changing Physical Volume Boot Types.....	84
Disabling a Path to a Physical Volume.....	87
Creating an Alternate Boot Disk.....	88
Mirroring the Boot Disk.....	90
Mirroring the Boot Disk on HP 9000 Servers.....	90
Mirroring the Boot Disk on HP Integrity Servers.....	92
Migrating a Volume Group to New Disks: vgmove.....	95
Example of Migrating a Volume Group to New Disks.....	95
Migrating a Logical Volume to New Disks: lvmove.....	96
Administering File System Logical Volumes.....	97
Creating a File System.....	97
Extending a File System.....	98
Reducing the Size of a File System.....	99
Backing Up a VxFS Snapshot File System.....	101
Administering Swap Logical Volumes.....	101
Creating a Swap Logical Volume.....	102
Extending a Swap Device.....	102
Reducing the Size of a Swap Device.....	102
Administering Dump Logical Volumes.....	102
Creating a Dump Logical Volume.....	102
Removing a Dump Logical Volume.....	103
Creating and Administering Snapshot Logical Volumes.....	103
Types of Snapshots.....	103
Creating a Snapshot Logical Volume.....	104
Removing Snapshot Logical Volumes.....	104
Displaying Snapshot Information.....	104
Managing the lvmpud Daemon.....	105
Hardware Issues.....	106
Integrating Cloned LUNs.....	106
4 Troubleshooting LVM.....	107
Troubleshooting Overview.....	107
Information Collection.....	107
Consistency Checks.....	108
Maintenance Mode Boot.....	108
I/O Errors.....	109
Recoverable Errors.....	109
Temporarily Unavailable Device.....	110
Permanently Unavailable Device.....	110
Nonrecoverable Errors.....	110
Media Errors.....	111
Missing Device When the Volume Group Was Activated.....	111

Volume Group Activation Failures.....	111
Quorum Problems with a Nonroot Volume Group.....	111
Quorum Problems with a Root Volume Group.....	112
Version 2.x Volume Group Activation Failures.....	112
Root Volume Group Scanning.....	114
LVM Boot Failures.....	115
Insufficient Quorum.....	115
Corrupted LVM Data Structures on Disk.....	115
Corrupted LVM Configuration File.....	115
Problems After Reducing the Size of a Logical Volume.....	115
Disk Troubleshooting and Recovery Procedures.....	116
Step 1: Preparing for Disk Recovery.....	116
Defining a Recovery Strategy.....	116
Using LVM Online Disk Replacement (LVM OLR).....	116
Mirroring Critical Information, Especially the Root Volume Group.....	117
Step 2: Recognizing a Failing Disk.....	118
I/O Errors in the System Log.....	118
Disk Failure Notification Messages from Diagnostics.....	118
LVM Command Errors.....	118
Step 3: Confirming Disk Failure.....	120
Step 4: Determining Action for Disk Removal or Replacement.....	123
Step 5: Removing a Bad Disk.....	126
Removing a Mirror Copy from a Disk.....	126
Removing Mirror Copy from Ghost Disk.....	126
Moving the Physical Extents to Another Disk.....	127
Removing the Disk from the Volume Group	127
Step 6: Replacing a Bad Disk (Persistent DSFs).....	129
Replacing a Mirrored Nonboot Disk.....	129
Replacing an Unmirrored Nonboot Disk.....	131
Replacing a Mirrored Boot Disk.....	134
Replacing an Unmirrored Boot Disk.....	137
Step 7: Replacing a Bad Disk (Legacy DSFs).....	138
Reporting Problems.....	139
5 Support and Other Resources.....	141
New and Changed Information in This Edition.....	141
About this Series.....	141
Typographic Conventions.....	141
Examples and Shells.....	142
Related Information.....	143
Finding HP-UX Information.....	143
HP-UX 11i Release Names and Operating System Version Identifiers.....	144
Determining Your System Version.....	144
Publication History.....	144
HP Insight Remote Support Software.....	145
HP Encourages Your Comments.....	146
A LVM Specifications and Limitations.....	147
Determining LVM's Maximum Limits on a System.....	150
B LVM Command Summary.....	152
C Volume Group Provisioning Tips.....	155
Choosing an Optimal Extent Size for a Version 1.0 Volume Group.....	155
Choosing an Optimal Extent Size for a Version 2.x Volume Group.....	156

D Striped and Mirrored Logical Volumes.....	157
Summary of Hardware Raid Configuration.....	157
LVM Implementation of RAID Levels in HP-UX.....	157
LVM Striped and Mirrored Logical Volume Configuration.....	158
Examples.....	158
Compatibility Note.....	160
E LVM I/O Timeout Parameters.....	161
Logical Volume Timeout (LV timeout).....	161
Physical Volume Timeout (PV timeout).....	161
Timeout Differences: 11i v2 and 11i v3.....	162
F Warning and Error Messages.....	163
Matching Error Messages to Physical Disks and Volume Groups.....	163
Messages For All LVM Commands.....	164
lvchange(1M).....	164
lvextend(1M).....	164
lvlnboot(1 M).....	165
pvchange(1M).....	166
vgcfgbackup(1M).....	166
vgcfgrestore(1M).....	166
vgchange(1M).....	167
vgcreate(1M).....	169
vgdisplay(1M).....	170
vgextend(1M).....	171
vgimport(1M).....	171
vgmodify(1M).....	172
vgversion(1M).....	172
Log Files and Trace Files: /var/adm/syslog/syslog.log.....	173
Glossary.....	177
Index.....	179

1 Introduction

This chapter addresses the following topics:

- “[LVM Features](#)” (page 8)
- “[LVM Architecture](#)” (page 9)
- “[Physical versus Logical Extents](#)” (page 10)
- “[LVM Volume Group Versions](#)” (page 11)
- “[LVM Device File Usage](#)” (page 12)
- “[LVM Disk Layout](#)” (page 16)
- “[LVM Limitations](#)” (page 18)
- “[Shared LVM](#)” (page 18)

LVM Features

Logical Volume Manager (LVM) is a storage management system that lets you allocate and manage disk space for file systems or raw data. Historically, you treated your disks individually and in terms of fixed-sized partitions; each disk or partition held a file system, swap space, boot area, or raw data. With LVM, you do not need to assign a disk or fixed-sized partition to a single purpose. Instead, consider the disks as a pool (or volume) of data storage, consisting of equal-sized extents. Extents are allocated into virtual storage devices known as [logical volumes](#), which can be treated as disks.

LVM provides the following capabilities:

- A logical volume size can be dynamically reduced or expanded to meet changing data needs. For example, a logical volume can be as small or large as the file system mounted to it requires. The file system can be extended without rebuilding it or the logical volume; reducing a file system is more complex, and may require recreating the file system.
- Small chunks of unused space from several disks can be combined to create a usable volume.
- A logical volume can exceed the size of a single physical disk. This feature is called [disk spanning](#), because a single file system (and individual files) can span disks.
- Up to six copies of identical data can be stored and updated simultaneously using LVM. This feature is called [mirroring](#) a logical volume, and requires an optional product, HP MirrorDisk/UX. See “[Increasing Data Availability Through Mirroring](#)” (page 24).
- Mirrored data can be configured to automatically create a new mirror to a separate disk when one of the mirror copies fails. This feature is called [sparing](#), and requires an optional product, HP MirrorDisk/UX. See “[Increasing Disk Redundancy Through Disk Sparing](#)” (page 27).
- A logical volume can be created so that logically contiguous data blocks (for example, chunks of the same file) are distributed across multiple disks, which speeds I/O throughput for large files when they are read and written sequentially. This feature is called [striping](#). Striping can be used in conjunction with mirroring. See “[Increasing Performance Through Disk Striping](#)” (page 31).
- A point-in-time image of a logical volume can be created, called a [snapshot](#) logical volume. See “[Creating and Administering Snapshot Logical Volumes](#)” (page 103).
- Devices accessed through multiple links can be configured to improve availability. If the primary link to a device fails, LVM can switch automatically to an alternate link. This feature is called [multipathing](#). See “[Increasing Hardware Path Redundancy Through Multipathing](#)” (page 28).

LVM Architecture

An LVM system starts by initializing disks for LVM usage. An LVM disk is known as a [physical volume](#) (PV). A disk is marked as an LVM physical volume using either the HP System Management Homepage (HP SMH) or the `pvcreate` command. Physical volumes use the same device special files as traditional HP-UX disk devices.

LVM divides each physical volume into addressable units called [physical extents](#) (PEs). Starting after the LVM metadata at the beginning of the disk, extents are allocated sequentially, with an index starting at zero and incrementing by one for each unit. The physical extent size is configurable at the time you form a volume group and applies to all disks in the volume group. You can select a size from 1 MB to 256 MB.

Physical volumes are organized into [volume groups](#) (VGs). A volume group can consist of one or more physical volumes, and there can be more than one volume group in the system. Once created, the volume group, not the disk, is the entity that represents data storage. Thus, whereas earlier you moved disks from one system to another, with LVM, you move a volume group from one system to another. Therefore, it is often convenient to have multiple volume groups on a system.

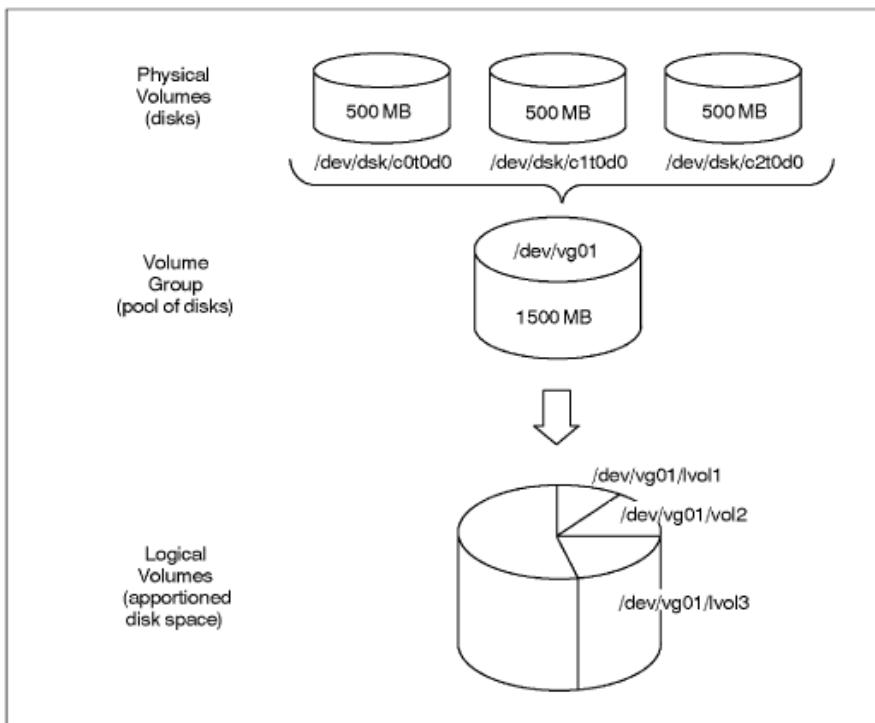
The pool of disk space that is represented by a volume group can be divided into [logical volumes](#) (LVs) of various sizes. Once created, logical volumes can be treated just like disk partitions. They are accessible through device special files. A logical volume can span a number of physical volumes in a volume group or represent only part of one physical volume.

The basic allocation units for a logical volume are called [logical extents](#) (LEs). A logical extent is mapped to a physical extent. Thus, if the physical extent size is 4 MB, the logical extent size is also 4 MB. The size of a logical volume is determined by the number of logical extents configured.

Starting with volume group Version 2.2, LVM introduced a new type of logical volume: [snapshot logical volume](#). Snapshots are point-in-time image of a logical volume. Snapshots allow you to create another copy of the logical volume (which can be used for a backup) without taking up as much of the physical space as the size of the logical volume. See "[Creating and Administering Snapshot Logical Volumes](#)" (page 103) for more information about snapshots.

You assign file systems, swap, dump, or raw data to logical volumes. For example, in [Figure 1](#), logical volume `/dev/vg01/lvol1` might contain a file system, logical volume `/dev/vg01/lvol2` might contain swap space, and logical volume `/dev/vg01/lvol3` might contain raw data. You can use HP SMH to create a file system in a logical volume of a specified size, then mount the file system. Alternately, you can use LVM commands to create, then extend a logical volume to allocate sufficient space for file systems or raw data. You then create and mount new file systems or install your application in the logical volume.

Figure 1 Disk Space Partitioned Into Logical Volumes



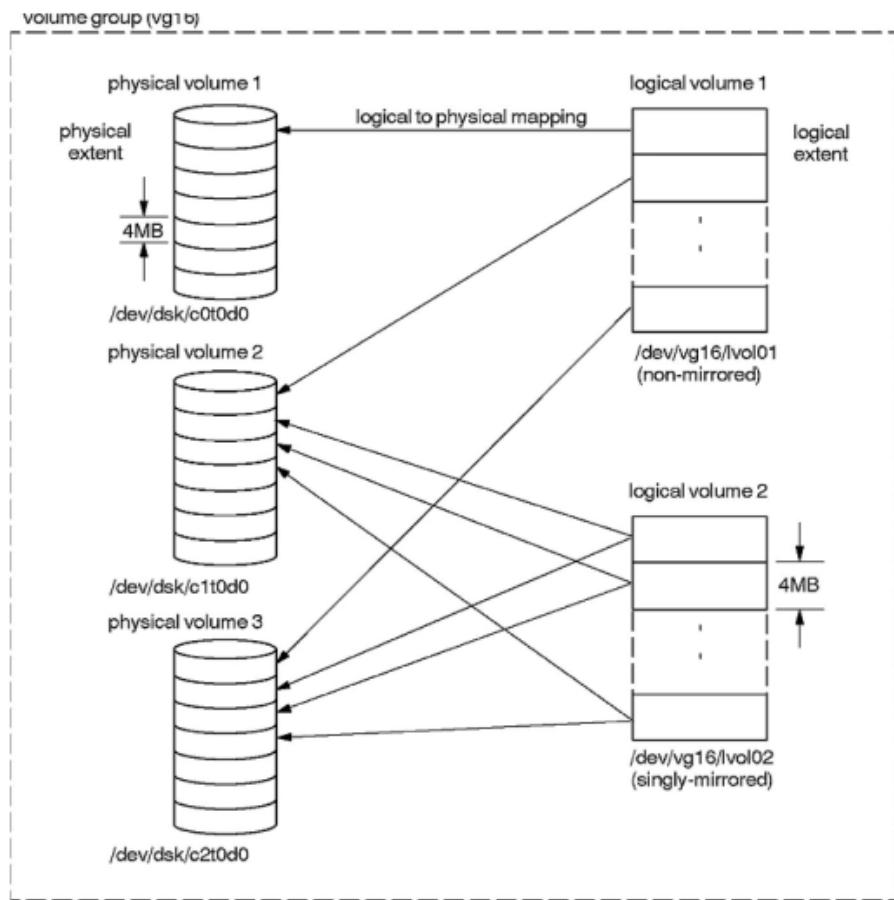
Physical versus Logical Extents

When LVM allocates disk space to a logical volume, it automatically creates a mapping of the logical extents to physical extents. This mapping depends on the policy chosen when creating the logical volume. Logical extents are allocated sequentially, starting at zero, for each logical volume. LVM uses this mapping to access the data, regardless of where it physically resides. Commands are provided for you to examine this mapping; see `pvdisplay(1M)` and `lvdisplay(1M)`.

Except for mirrored, striped, or striped-mirrored logical volumes, each logical extent is mapped to one physical extent. For mirrored logical volumes, each logical extent is mapped to multiple physical extents, depending on the number of mirror copies. For example, if one mirror copy exists, then each logical extent maps to two physical extents, one extent for the original and one for the mirror copy. For more information on mirroring, see “[Increasing Data Availability Through Mirroring](#)” (page 24). For information on striped logical volumes, see “[Increasing Performance Through Disk Striping](#)” (page 31). Also refer to the book *Disk and File Management Tasks on HP-UX*.

Figure 2 shows an example of several types of mapping available between physical extents and logical extents within a volume group.

Figure 2 Physical Extents and Logical Extents



As shown in [Figure 2](#), the contents of the first logical volume are contained on all three physical volumes in the volume group. Because the second logical volume is mirrored, each logical extent is mapped to more than one physical extent. In this case, there are two physical extents containing the data, each on both the second and third disks within the volume group.

By default, LVM assigns physical extents to logical volumes by selecting available physical extents from disks in the order in which they appear in the LVM configuration files, `/etc/lvmtab` and `/etc/lvmtab_p`. As a system administrator, you can bypass this default assignment and control which disks are used by a logical volume (see ["Extending a Logical Volume to a Specific Disk" \(page 54\)](#)).

If a logical volume is to be used for root, boot, primary swap, or dump, the physical extents must be **contiguous**, which means that the physical extents must be allocated in increasing order with no gaps on a single physical volume. For logical volumes that are not being used for root, boot, primary swap or dump, physical extents that correspond to contiguous logical extents within a logical volume can be **noncontiguous** on a physical volume or reside on entirely different disks. As a result, a file system created within one logical volume can reside on more than one disk.

LVM Volume Group Versions

As of the March 2010 release of HP-UX 11i Version 3, LVM supports four versions of volume groups. All information and tasks in this document apply to all volume group versions except where noted.

Version 1.0 is the version supported on all current and previous versions of HP-UX 11i. The procedures and command syntax for managing Version 1.0 volume groups are unchanged from previous releases. When creating a new volume group, `vgcreate` defaults to Version 1.0.

Version 2.0, 2.1, and 2.2 enable the configuration of larger volume groups, logical volumes, physical volumes, and other parameters. Version 2.1 is identical to Version 2.0, but allows a greater number of volume groups, physical volumes, and logical volumes. Version 2.x volume groups are managed exactly like Version 1.0 volume groups, with the following exceptions:

- Version 2.x volume groups have simpler options to the `vgcreate` command. When creating a Version 2.x volume group, you specify only the extent size and the maximum size to which the volume group can grow. This gives LVM greater flexibility in managing space; you can use the same parameters for a volume group with many small PVs and for a volume group with a few large PVs. For more information on volume group creation, see “[Creating a Volume Group](#)” (page 44)
- Version 2.0 volume groups are not recognized on previous releases of HP-UX, including versions of HP-UX 11i Version 3 before March 2008. Version 2.1 volume groups are not recognized on previous releases of HP-UX, including versions of HP-UX 11i Version 3 before September 2008. Version 2.2 volume groups are only supported on systems operating with the HP-UX 11i Version 3 March 2010 Update and forward.
- Version 2.0 and 2.1 volume groups do not support root, boot, swap, or dump logical volumes. The `1vlnboot` and `1vrmbboot` commands display an error message if run on a Version 2.0 or 2.1 volume group. Note that Version 2.2 and higher do support root, boot, swap, or dump logical volumes.
- Version 2.0 and 2.1 volume groups do not support bootable physical volumes. You cannot add a physical volume created with `pvcreate -B` to a Version 2.0 or 2.1 volume group. Note that Version 2.2 and higher do support bootable physical volumes.
- Version 2.2 or higher volume groups support creation of snapshots of logical volumes.
- Version 2.x volume groups do not support disk sparing. Using the `-z` option to the `vgextend` or `pvchange` command displays an error message.
- The `pvck` command is not supported on Version 2.x volume groups.
- Some HP-UX products do not support Version 2.x volume groups. For more information, see the *HP-UX Logical Volume Manager and MirrorDisk/UX Release Notes* for your release.

To compare Version 1.0 and Version 2.x volume groups, see the tables that list maximum limits for different volume group versions in “[LVM Specifications and Limitations](#)” (page 147). You can also display volume group limits with the `1vmadm` command; an example is shown in “[LVM Specifications and Limitations](#)” (page 147) . The `1vmadm` is fully covered in the `lvmadm(1M)` manual page.

See “[Creating and Administering Snapshot Logical Volumes](#)” (page 103) for information about snapshot logical volumes.

LVM Device File Usage

All LVM components are represented by device special files located in the `/dev` directory. Device special files act as agents for managing the interactions with the disk space. The LVM device files are created by both HP SMH and HP-UX commands. This section describes the device special files used by LVM, and naming conventions for LVM objects.

HP-UX 11i v3 introduces a new representation of mass storage devices called the **agile view**. In this representation, the device special file (DSF) name for each disk no longer contains path (or link) information. The multipathed disk has a single **persistent DSF** regardless of the number of physical paths to it. The **legacy view**, represented by the **legacy DSF**, continue to exist.

Legacy Device Files versus Persistent Device Files

As of HP-UX 11i Version 3, disk devices can be represented by two different types of device files in the `/dev` directory, **legacy** and **persistent**.

Legacy device files were the only type of mass storage device files in releases prior to HP-UX 11i Version 3. They have hardware path information such as SCSI bus, target, and LUN encoded in the device file name and minor number. For example, the legacy device file `/dev/dsk/c3t2d0` represents the disk at card instance 3, target address 2, and lun address 0.

Persistent device files are not tied to the physical hardware path to a disk, but instead map to the disk's unique worldwide identifier (WWID). Thus, the device file is unchanged if the disk is moved from one interface to another, moved from one switch or hub port to another, or presented from a different target port to the host. The name of a persistent device file follows a simpler naming convention: `/dev/disk/diskn`, where *n* is the instance number assigned to the disk. Neither the device file name nor the minor number contain any hardware path information.

In addition, if the disk has multiple hardware paths, it is represented by a single persistent device file. Persistent device files transparently handle multipathed disks and supersede LVM's multipathing functionality described in “[Increasing Hardware Path Redundancy Through Multipathing](#)” (page 28). If a disk has multiple hardware paths, which LVM refers to as **pvlinks**, the persistent device special file acts as a single access point for all the links. I/O requests are distributed across all available links by the mass storage stack, with a choice of load balancing algorithms. If a link fails, the mass storage stack automatically disables the failed link and I/O continues on all remaining links. Any failed or nonresponsive links are monitored, so that when a failed link recovers, it is automatically and transparently reincorporated into any load balancing. New disks and links are also automatically discovered and added to load balancing. If the disk's connectivity changes—addition, removal, or modification of a link—applications using the persistent device file are not affected, provided at least one link is still active. New disks are automatically discovered.

You can use either persistent or legacy device files for LVM disks. LVM recommends the use of persistent device special files, because they support a greater variety of load balancing options.

To facilitate migration from legacy DSF to persistent DSF, HP provides the `/usr/contrib/bin/vgdsf` script, which works for both root and non-root volume groups. For more detail on the HP-UX 11i v3 mass storage stack, its new features (such as the persistent DSFs), and their benefits, see *The Next Generation Mass Storage, HP-UX 11i v3 White Paper* on <http://www.hp.com/go/hpux-core-docs>.

NOTE: To use LVM's alternate link functionality, you must use legacy device files, and disable multipathing through those legacy device files, as described in “[Increasing Hardware Path Redundancy Through Multipathing](#)” (page 28).

Cluster-wide Device Special Files

HP-UX 11i Version 3, September 2010 Release introduces the support of cluster-wide device special files (cDSF) for nodes running with Serviceguard, A.11.20.

A cluster device special files provides a consistent set of device special files across a set of cluster nodes. The cluster device special files of a LUN will be the same on any node in a specified set of nodes that share the LUN. The set of cluster nodes across which cluster device special files need to be created can be specified using the `cmsetdsfgroup` command. The cluster device special files can be displayed using the `io_cdsf_config` command. See the *Managing Serviceguard Eighteenth Edition* manual for more detail on cluster device special files

Naming Conventions for LVM

You must refer to LVM devices or volume groups by name when using them within HP SMH or with HP-UX commands. By default, the LVM device files created by both HP SMH and HP-UX commands follow a standard naming convention. However, you can choose customized names for volume groups and logical volumes.

Physical Volume Names

Physical volumes are identified by their device file names, as follows:

Table 1 Physical Volume Naming Conventions

Device File Name	Type of Device
/dev/disk/diskn	Persistent block device file
/dev/disk/diskn_p2	Persistent block device file, partition 2
/dev/rdisk/diskn	Persistent character device file
/dev/rdisk/diskn_p2	Persistent character device file, partition 2
/dev/dsk/cntndn	Legacy block device file
/dev/dsk/cntndns2	Legacy block device file, partition 2
/dev/rdsk/cntndn	Legacy character device file
/dev/rdsk/cntndns2	Legacy character device file, partition 2
/dev/cdisk/diskn	Cluster block device file
/dev/rcdisk/diskn	Cluster character device file

Each disk has a block device file and a character or raw device file, the latter identified by the `r`. Which name you use depends on what task you are doing with the disk.

For the boot disk on HP Integrity servers, make sure to use the device files with the `_p2` suffix or `s2` suffix, because they represent the HP-UX partition on the boot disk. On HP 9000 servers, use the device file without a partition number.

Use a physical volume's raw device file for the following tasks only:

- Preparing a physical volume for LVM using the `pvcreate` command. Here you use the device file for the disk; for example, `/dev/rdisk/disk14`. (The absence of a partition suffix indicates that you are referring to the entire disk.)
- Removing LVM information from a physical volume using the `pvremove` command.
- Restoring your volume group configuration using the `vgcfgrestore` command.
- Performing a consistency check on a physical volume using the `pvck` command.
- Modifying the volume group identifier on a physical volume using the `vgchgid` command.
- Changing the disk type of a physical volume using the `vgmodify` command.

For all other tasks, use the block device file. For example, when you add a physical volume to a volume group using the `vgextend` command, you use the disk's block device file for the disk, such as `/dev/disk/disk14`.

All disk device files are created automatically when a new disk is discovered. For more information, see `insf(1M)`.

Volume Group Names

Each volume group must have a unique name, up to 255 characters. For example, typical volume group names are `vg01`, `vgroot`, or `vg_sales`. Although the name does not need to start with `vg`, HP recommends using this prefix. By default, HP SMH uses the names of the form `/dev/vgnn`. The number `nn` starts at 00 and is incremented in the order that volume groups are created. By default, the root volume group is `vg00`.

Logical Volume Names

Logical volumes are identified by their device file names, which can either be assigned by you or assigned by default when you create a logical volume using the `lvcreate` command.

When assigned by you, you can choose any name up to 255 characters.

When assigned by default, these names take the form `/dev/vg n /lvol N` (the block device file form) and `/dev/vg n /rlvol N` (the character device file form). The number N starts at 1 and increments in the order that logical volumes are created within each volume group.

When LVM creates a logical volume, it creates both block and character device files. LVM then places the device files for a logical volume in the appropriate volume group directory. For example, the default block name for the first logical volume created in volume group `vg01` has the following full path name:

```
/dev/vg01/lvol1
```

If you create a logical volume to contain a sales database, you can name it explicitly as follows:

```
/dev/vg01/sales_db_lv
```

After the logical volume in the previous example is created, it has two device files: `/dev/vg01/sales_db_lv` for the block device file and `/dev/vg01/rsales_db_lv` for the character, or raw, device file.

Physical Volume Group Names

Physical volume groups are useful for mirroring and are discussed in “[Increasing Performance Through I/O Channel Separation](#)” (page 33). The only naming restriction is that within a volume group, each physical volume group must have its own unique name. For example, the volume group `/dev/vg02` might have two physical volume groups named `pvg1` and `pvg2`.

Snapshot Logical Volume Names

The `lvcreate` command is used to create a snapshot logical volume, and its `-n` option can specify the name of the snapshot logical volume. If this option is not used, then the snapshot logical volume name defaults to the string comprising of the original logical volume name, the tag S , and the snapshot volume minor number N , where N is the decimal equivalent of the least significant 12 bits of the snapshot volume's minor number. The range for N is 1 to 2047.

Additionally, you can use the `lvcreate -t` option to append the timestamp to the snapshot volume name. The timestamp is of the form: `YYYYMMDD_HH.MM.SS`. The default is to not append a timestamp to the name of the snapshot. The `lvcreate(1M)` manpage provides full details on creating snapshot logical volumes. See “[Creating and Administering Snapshot Logical Volumes](#)” (page 103) for more information about snapshot logical volumes.

Device Number Format

The device files associated with LVM reside in the `/dev` directory. For each volume group, a directory under `/dev` is named after the volume group. In that directory is a single “group” device file and separate block and character device files for each logical volume.

The following is a sample listing:

```
# ls -1 /dev/vg01
total 0
crw-r--r-- 1 root      root      64 0x010000 Mar 28 2004 group
brw-r---- 1 root      root      64 0x010001 Jul 29 16:53 lvol1
brw-r---- 1 root      root      64 0x010002 Jul 29 16:53 lvol2
crw-r---- 1 root      root      64 0x010001 Mar 28 2004 rlvol1
crw-r---- 1 root      root      64 0x010002 Mar 28 2004 rlvol2
```

By default, volume group numbering begins with zero (`vg00`), while logical volumes begin with one (`lvol1`). This is because the logical volume number corresponds to the minor number and the volume group's group file is assigned minor number 0.

Physical volumes use the device files associated with their disk. LVM does not create device files for physical volumes.

Version 1.0 Device Number Format

Table 2 lists the format of the device file number for Version 1.0 volume groups.

Table 2 Version 1.0 Device Number Format

Major Number	Volume Group Number	Reserved	Logical Volume Number
64	0-0xff	0	0-0xff 0=group file

For Version 1.0 volume groups, the major number for LVM device files is 64. The volume group number is encoded into the top eight bits of the minor number, and the logical volume number is encoded into the low eight bits. Logical volume number 0 is reserved for the group file.

Version 2.x Device Number Format

Table 3 lists the format of the device file number for Version 2.x volume groups.

Table 3 Version 2.x Device Number Format

Major Number	Volume Group Number	Logical Volume Number
128	0-0x7ff	0-0x7ff 0=group file

For Version 2.x volume groups, the major number for LVM device files is 128. The volume group number is encoded into the top twelve bits of the minor number, and the logical volume number is encoded into the low twelve bits. Logical volume number 0 is reserved for the group file.

NOTE: The most significant bit of the volume group number and logical volume number fields are reserved and must be zero.

The device number format is subject to change.

LVM Disk Layout

NOTE: This information applies only to disks belonging to Version 1.0 and Version 2.2 (or higher) volume groups.

There are two kinds of LVM disk layouts, one for boot disks and another for all other LVM disks. These differ in their data structures. Nonbootable disks have two reserved areas: the physical volume reserved area (PVRA) and the volume group reserved area (VGRA). Bootable disks have a PVRA and VGRA, and additional sectors reserved for the boot data reserved area (BDRA) and boot LIF.

Boot Data Reserved Area

The BDRA contains the information needed to configure the root, primary swap, and dump logical volumes, and to mount the root file system.

Information about the LVM disk data structures in the BDRA is maintained with the `lvolnboot` and `lvrmboot` commands. The following is a sample output:

```
# lvolnboot -v
Boot Definitions for Volume Group /dev/vg00:
Physical Volumes belonging in Root Volume Group:
        /dev/dsk/c3t0d0 -- Boot Disk
        /dev/dsk/c4t0d0 -- Boot Disk
        /dev/dsk/c5t0d0
        /dev/dsk/c12t0d0 -- Boot Disk
Root: lvol1      on: /dev/dsk/c3t0d0
```

```

Swap: lvol2      /dev/dsk/c4t0d0
          on: /dev/dsk/c3t0d0
                  /dev/dsk/c4t0d0
Dump: lvol2      on: /dev/dsk/c3t0d0

```

The physical volumes designated "Boot Disk" are bootable, having been initialized with `mkboot` and `pvcreate -B`. Multiple lines for `lvol1` and `lvol2` indicate that the root and swap logical volumes are being mirrored.

Logical Interface Format Area

LVM boot disks contain a Logical Interface Format (LIF) area, in which is stored a `LABEL` file. On HP 9000 servers, the LIF area contains boot utilities such as the initial system loader (ISL), the kernel boot loader (HPUX), the autoboot file (AUTO), and offline diagnostics.

The `LABEL` file is created and maintained by `lvolboot` and `lvrmboot`. It contains information about the starting point and size of boot-relevant logical volumes, including the boot file system (`/stand`). Utilities can use the `LABEL` file to access the root, primary swap, and dump logical volumes without actually using LVM.

Physical Volume Reserved Area

The physical volume reserved area (PVRA) contains information describing the physical volume, such as its unique identifier, physical extent information, and pointers to other LVM structures on the disk.

Volume Group Reserved Area

The volume group reserved area (VGRA) describes the volume group to which the disk belongs. The information is replicated on all of the physical volumes and updated whenever a configuration change is made. Among other data, it contains the following information:

- A list of physical volumes in the volume group, including physical volume status and size, and a map of physical extents to logical volumes.
- A list of logical volumes in the volume group (including the status and capabilities of each logical volume), its scheduling and allocation policies, and the number of mirror copies.
- A volume group header containing the VGID and three configurable parameters:
 - the number of physical volumes allowed in the volume group
 - the maximum number of logical volumes allowed in the volume group
 - the maximum number of physical extents allowed per physical volume

Since each physical extent is recorded in the VGRA, the extent size has a direct bearing on the size of the VGRA. In most cases, the default extent size is sufficient. However, if you encounter problems, consider that the VGRA is a fixed size and a high-capacity physical volume might exceed the total number of physical extents allowed. As a result, you might need to use a larger-than-default extent size on high-capacity LVM disks. Conversely, if all LVM disks in a volume group are small, the default number of extents might make the VGRA too large, wasting disk and memory space. A smaller-than-default extent size or number of physical extents might be preferable. A high-capacity physical volume might be unusable in a volume group whose extent size is small or set with a small number of physical extents per disk.

User Data Area

The user data area is the region of the LVM disk used to store all user data, including file systems, virtual memory system (swap), or user applications.

LVM Limitations

LVM is a sophisticated subsystem. It requires time to learn, it requires maintenance, and in rare cases, things can go wrong.

HP recommends using logical volumes as the preferred method for managing disks. Use LVM on file and application servers. On servers that have only a single disk and are used only to store the operating system and for swap, a “whole-disk” approach is simpler and easier to manage. LVM is not necessary on such systems.

By default, LVM configurations are automatically backed up each time you change them, in the default directory, /etc/lvmconf. Mirroring provides insurance against data loss that is not available under the whole-disk method.

NOTE: This default directory /etc/lvmconf can be changed for volume groups Version 2.x by configuring a new path in the `LVMP_CONF_PATH_NON_BOOT` variable defined in the /etc/lvmrc file. See `vgcfgbackup(1M)` for details.

Additional limitations to LVM include the following:

- Both LVM disks and non-LVM disks can exist simultaneously on your system, but a given disk or partition must be managed entirely by either LVM or non-LVM methods. That is, you cannot combine these techniques for use with a single disk or partition.
- On an HP Integrity server, LVM supports partitioning of the root disk and its mirrors only, and supports only one HP-UX partition on any disk.
- Floppy disks, optical disks, and CD-ROMs do not support logical volumes.
- You must use an LVM or VERITAS™ Volume Manager (VxVM) disk for your root disk.
- To use LVM, a disk must be first initialized into a physical volume.
- To be allocatable for storage, a physical volume must be assigned to a volume group.
- A physical volume can belong to only one volume group.
- The extent size of a volume group is fixed when the volume group is created. It cannot be changed without recreating the volume group.

Shared LVM

Shared LVM (SLVM) allows multiple systems in a Serviceguard cluster to share (read/write) disk resources in the form of volume groups. SLVM is designed to be used by specialized distributed applications that use raw access to disks, rather than going through a file system.

Shared mode is configured with the `vgchange` command. The `vgdisplay` command will show the current activation mode for the volume group.

The following commands have restrictions when operated on volume groups that are activated in shared mode:

- For volume group Version 1.0 or 2.0, `lvchange`, `lvcreate`, `lvextend`, `lvmerge`, `lvmmove`, `lvreduce`, `lvremove`, `lvsplit`, `vgextend`, `vgmodify`, `vgmove`, and `vgreduce` cannot be used in shared mode.
- The `pvmove` command cannot be used in shared mode for volume group Version 1.0.
- For `pvchange`, only the `-a` option (for attaching or detaching a path to the specified physical volume) is allowed in shared mode. All other `pvchange` options are not supported in shared mode.

To replace disks that are part of a volume group shared by a Serviceguard cluster, the physical volume must be detached and attached (using `pvchange`) independently on each of the systems in the cluster.

- The `lvolnboot` and `lvrmbboot` commands cannot be used in shared mode.

Volume group Version 2.2 and higher with snapshot logical volumes configured cannot be activated in shared mode. Also, snapshots cannot be created off logical volumes belonging to shared volume groups.

Synchronization of mirrored volume groups will be slower on shared volume groups.

Beginning with the September 2009 Update, LVM provides new options in LVM commands to let system administrators reconfigure shared volume groups, logical volumes, and physical volumes on multiple nodes in a cluster, while user applications and data reads/writes remain available and active. `lvmppud` is the daemon that handles LVM online shared volume group reconfiguration. The online reconfiguration of shared volume groups is only available on volume groups version 2.1 and higher. See the *SLVM Online Reconfiguration* white paper for details.

For Version 1.0 and 2.0 volume groups, you can use the `vgchange -x` option (the SLVM SNOR feature) to change the configuration of a shared volume group, while keeping it active only on a single node. Using this procedure, applications on at least one node remain available during the volume group reconfiguration.

The shared LVM features, options, and restrictions are fully covered in the LVM command manpages. See “[LVM Command Summary](#)” (page 152) for supported LVM commands. Each LVM command has a manpage that is installed with the LVM product.

2 Configuring LVM

By default, the LVM commands are already installed on your system. This chapter discusses issues to consider when setting up your logical volumes. It addresses the following topics:

- “[Planning Your LVM Configuration](#)” (page 20)
- “[Setting Up Different Types of Logical Volumes](#)” (page 20)
- “[Planning for Availability](#)” (page 24)
- “[Planning for Recovery](#)” (page 34)
- “[Planning for Performance](#)” (page 29)

Planning Your LVM Configuration

Using logical volumes requires some planning. Some of the issues to consider for planning purposes are discussed in this chapter. Consider these issues before setting up or modifying logical volumes on your system.

- For what purpose will you use a logical volume? For raw data or a file system? As a swap area or dump area? See “[Setting Up Different Types of Logical Volumes](#)” (page 20).
- How big will you make a logical volume?
- Does your data require high availability? If so, consider mirroring your logical volume across multiple disks, as described in “[Increasing Data Availability Through Mirroring](#)” (page 24). Also consider setting up spare disks to handle mirror failure, as described in “[Increasing Disk Redundancy Through Disk Sparing](#)” (page 27).
- Is I/O performance very important to you? If so, consider striping your logical volumes across multiple disks, as described in “[Increasing Performance Through Disk Striping](#)” (page 31), and separating I/O channels, as described in “[Increasing Performance Through I/O Channel Separation](#)” (page 33). For additional recommendations for performance, read “[Planning for Performance](#)” (page 29).
- Do you need to balance high availability and I/O performance? If so, consider striping and mirroring your logical volume, as described in “[Increasing Data Availability Through Mirroring](#)” (page 24) and “[Increasing Performance Through Disk Striping](#)” (page 31).
- Is it important to recover from disk failures quickly? If so, see “[Planning for Recovery](#)” (page 34).

Setting Up Different Types of Logical Volumes

This section contains information on setting up special logical volumes.

Setting Up Logical Volumes for Raw Data Storage

You can optimize raw I/O performance by planning your logical volumes specifically for raw data storage. To create a raw data logical volume (such as for a database), consider how large the logical volume must be and how such a logical volume is distributed over your disks.

Typically, you specify the size of a logical volume in megabytes. However, a logical volume must be a multiple of the extent size used in the volume group. For example, if a database partition requires 2002 MB and the logical extent size is 4 MB, LVM creates a logical volume that is 2004 MB (or 501 logical extents).

If you plan to use logical volumes heavily for raw data storage (such as for setting up database partitions), consider how the logical volumes are distributed over your disks.

By default, LVM assigns disk space for a logical volume from one physical volume, uses the space on this physical volume entirely, then assigns space from each successive physical volume in the same manner. LVM uses the physical volumes in the order in which they appear in /etc/lvmtab

and /etc/lvmtab_p, which means that data of a logical volume might not be evenly distributed over all the physical volumes within your volume group.

As a result, when I/O access to the logical volumes occurs, one or more disks within the volume group might be heavily used, while the others might be lightly used, or not used at all. This arrangement does not provide optimum I/O performance.

As a better alternative, you can set up your logical volume on specific disks in an interleaved manner, thus balancing the I/O access and optimizing performance (see “[Extending a Logical Volume](#)” (page 53)).

Because there are no HP-UX commands that identify that the contents of a logical volume are being used for raw data, use recognizable names for the logical volumes you create for raw data. In this way, you can recognize the contents of such a logical volume.

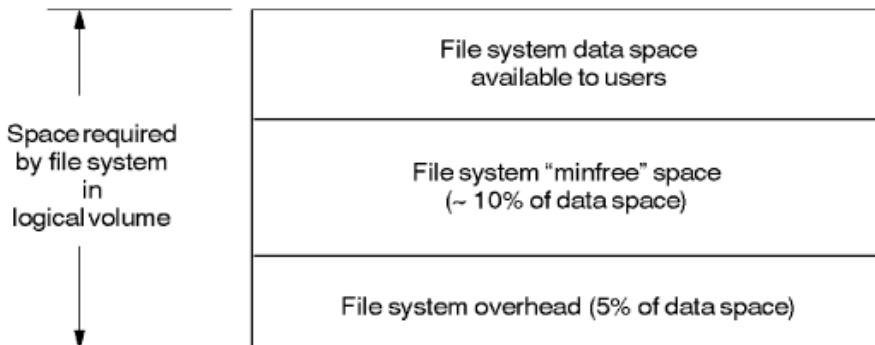
Setting Up Logical Volumes for File Systems

File systems reside in a logical volume just as they do within disk partitions or nonpartitioned disks. Two types of file systems can be used in a logical volume: Hierarchical File Systems (HFS) and Journaled File Systems (JFS) (VxFS).

Choosing the Initial Size of File System Logical Volumes

When determining the required space for a file system, consider the three major components shown in [Figure 3](#).

Figure 3 File System Space Components



To estimate how big to make a logical volume that will contain your file system, follow these steps:

1. Estimate how much disk space users will need for their data in the future. Allow for any anticipated changes, which usually include additional growth. (Use the du command to see how much disk space is currently used.)
2. Add 10% to the above amount for a “minfree” area; this area is reserved to maintain performance.
3. Add another 5% for file system overhead; this includes all data structures required to maintain the file system.
4. Round up to the next integer multiple of the logical extent size used in this logical volume to find the size in logical extents. (This step is performed automatically when you create a logical volume.)

For example, if a group of users will require 60 MB space for file system data, the following estimate allows for expected growth. Add 6 MB for the minfree space. Then add 3 MB for file system overhead for a total estimate of 69 MB required by the file system, and for the logical volume that contains the file system. If you are creating the logical volume in a volume group that has an extent size of 4 MB, round 69 to 72 to make it divisible by 4 MB.

Although these estimates are not precise, they suffice for planning a file system size. Create your file system large enough to be useful for some time before increasing its size.



TIP: Because increasing the size of a file system is usually easier than reducing its size, be conservative in estimating how large to create a file system.

An exception is the root file system. As a contiguous logical volume, the root file system is difficult to extend.

Resizing File System Logical Volumes

If your users have outgrown the space originally allocated for the file system, you can increase its size by first enlarging the logical volume it resides in using the `lvextend` command, then using the `extdeffs` command to enlarge the file system contained in the logical volume.

Decreasing the size of a file system can be difficult. Based on the type of file system, you might not be able to decrease its size. However, you can create a new smaller file system to take its place.

For more information on resizing file system logical volumes, see [“Administering File System Logical Volumes” \(page 97\)](#).

File System Logical Volume Guidelines

Use the following guidelines when configuring file system logical volumes:

- If you create a file system that spans LVM disks, be sure that the logical volume in which the file system resides spans identical disk types for best system performance.
- By default, LVM creates logical volumes on available disks, not necessarily with regard for best performance. A file system can span two disks with different characteristics, in which case the file system performance can possibly be impaired.

You can control which physical volumes contain the physical extents of a logical volume by following these steps:

1. Create a logical volume without specifying a size using the `lvcreate` command or HP SMH. When you do not specify a size, by default, no physical extents are allocated for the logical volume.
 2. Extend the logical volume (that is, allocate space) to the specific physical volumes you want to contain the file system using the `lvextend` command.
- The root or boot logical volume is limited to either 2 GB or 4 GB, depending on your processor.

Setting Up Logical Volumes for Swap

NOTE: Version 2.0 and 2.1 volume groups do not support swap logical volumes.

This section explains what to consider when using logical volumes as swap devices. For information on managing your system swap space, including determining how much and what type of swap space the system needs, see *HP-UX System Administrator’s Guide: Configuration Management*.

When configured as swap, logical volumes are treated as **device swap space**. Device swap space occupies a logical volume or partition, which is typically reserved expressly for swapping purposes. This space can also be configured as a dump area (See [“Dump Logical Volume Guidelines” \(page 23\)](#)).

Swap Logical Volume Guidelines

Use the following guidelines when configuring swap logical volumes:

- Interleave device swap areas for better performance.

Two swap areas on different disks perform better than one swap area with the equivalent amount of space. This configuration allows *interleaved swapping*, which means the swap areas are written to concurrently, thus enhancing performance.

When using LVM, set up secondary swap areas within logical volumes that are on different disks using `lvextend`.

If you have only one disk and must increase swap space, try to move the primary swap area to a larger region.

- Similar-sized device swap areas work best.
Device swap areas must have similar sizes for best performance. Otherwise, when all space in the smaller device swap area is used, only the larger swap area is available, making interleaving impossible.
- By default, primary swap is located on the same disk as the root file system. By default, the `/stand/system` system kernel configuration file contains the configuration information for primary swap.
- If you are using logical volumes as secondary swap, allocate the secondary swap to reside on a disk other than the root disk for better performance.

Setting Up Logical Volumes for Dump

NOTE: Version 2.0 and 2.1 volume groups do not support dump logical volumes.

This section explains what to consider when using logical volumes as dump devices. A dump area is disk space used to write an image of the core memory after a system crash. The analysis of a core dump can be useful in troubleshooting and restoring the system to working order.

By default, the primary swap device also serves as a dump area when no dump area is specifically designated. Although you are not required to retain primary swap as your dump area, doing so conserves disk space. You can configure a different or multiple dump devices on your system. To do this, create a logical volume as a dump device. This device can also be used for swap.

For information on adding, removing, or modifying dump devices, and configuring the dump algorithms, see the *HP-UX System Administrator's Guide: Configuration Management*.

Dump Logical Volume Guidelines

Use the following guidelines when configuring dump logical volumes:

- HP recommends using logical volumes for dump area rather than disk partitions.
- A dump logical volume can exist only within the root volume group; that is, the volume group that contains the root logical volume.
- You can use any secondary swap logical volume as a dump area too, provided the swap area is in the root volume group.

Setting Up Snapshot Logical Volumes for Backup

Beginning with HP-UX 11i v3 March 2010 Update, LVM introduces snapshot logical volumes, which can be used to capture point-in-time images of logical volumes quickly and without requiring as much space as the size of the logical volume. To back up a logical volume while it is in use, a snapshot of the logical volume can be taken and used for backup while the original logical volume remains online and is being modified. The snapshot can be deleted after the backup. The approach, unlike the `lvsplit` approach, does not require the user to reduce a mirror copy from the original logical volume.

Refer to the *LVM Snapshot Logical Volumes* white paper for more details on backups using snapshot logical volumes.

Planning for Availability

This section describes LVM features that can improve the availability and redundancy of your data. It addresses the following topics:

- “[Increasing Data Availability Through Mirroring](#)” (page 24)
- “[Increasing Disk Redundancy Through Disk Sparing](#)” (page 27)
- “[Increasing Hardware Path Redundancy Through Multipathing](#)” (page 28)

Increasing Data Availability Through Mirroring

NOTE: Mirroring requires an optional product, HP MirrorDisk/UX.

Mirroring means storing identical copies of data in logical volumes, preferably on separate disks. This redundancy has several advantages:

- If you mirror the root file system and swap, the operating system can tolerate a root disk failure because critical data is available on more than one LVM disk.
- If you mirror the logical volumes used by a particular application, the application continues to run even if a disk fails.
- If an I/O channel fails, LVM can recover the data from the duplicate source.
- Mirroring speeds read-intensive applications by enabling the hardware to read data from the most convenient LVM disk.
- You can back up one copy of the data while another copy continues to run.

Mirroring maps one logical extent to two or more sets of physical extents. The number of logical extents remains constant, but the number of used physical extents (and therefore, occupied disk space) changes depending on the number of mirrored copies. Mirroring increases data protection and system availability, but consumes twice as much disk space (or as many times more as there are mirror copies), so use disk mirroring for volatile, mission-critical data only.

Mirrored logical volumes must belong to the same volume group; you cannot mirror across volume groups.

This section contains the following information:

- “[Mirror Write Behavior Control](#)” (page 24)
- “[Synchronizing a Mirrored Logical Volume](#)” (page 26)

To learn more about basic mirroring tasks, see *Disk and File Management Tasks on HP-UX* published by Prentice Hall PTR, 1997.

Mirror Write Behavior Control

Three policies govern how mirrored logical extents are written to physical extents: the **allocation policy**, the **scheduling policy** for disk writes, and the **synchronization policy** for crash recovery. These policies can be set using HP SMH, the `lvcreate` command, or the `lvchange` command.

Allocation Policy

Mirrored extents can be allocated on physical volumes by **strict** or **nonstrict**, **contiguous** or **noncontiguous** policies. By default, the allocation policy of mirrored logical volumes is set to strict, noncontiguous.

Strict and Nonstrict Allocation

Strict allocation requires logical extents to be mirrored to physical extents on different physical volumes. Nonstrict allocation allows logical extents to be mirrored to physical extents that may be

on the same physical volume. The `-s y` and `-s n` options to the `lvcreate` or `lvchange` commands set strict or nonstrict allocation.

⚠ **CAUTION:** Using nonstrict allocation can reduce the redundancy created by LVM mirroring because a logical extent can be mirrored to different physical extents on the same disk. Therefore, the failure of this one disk makes both copies of the data unavailable.

Contiguous and Noncontiguous Allocation

Contiguous allocation has three characteristics: the physical extents are allocated in ascending order, no gap exists between physical extents within a mirror copy, and all physical extents of a mirror copy reside on a single physical volume. Noncontiguous allocation allows logical extents to be mapped to nonconsecutive physical extents. The `-c y` and `-c n` options to the `lvcreate` or `lvchange` commands set contiguous or noncontiguous allocation.

NOTE: When the logical volumes allocated from the root volume group are mirrored, each must be set up with contiguous allocation.

Scheduling Policy

The LVM scheduler converts logical I/O requests into one or more physical I/O requests, then schedules them for processing at the hardware level. Scheduling occurs for both mirrored and nonmirrored data.

Two I/O scheduling policies are available: **parallel** and **sequential**.

Parallel Scheduling

The parallel scheduling policy is used by default with mirroring for maximum I/O performance. Parallel scheduling causes mirror operations to write simultaneously to all copies. LVM optimizes reads by reading from the physical volume with the fewest outstanding I/O operations. The `-d p` option to the `lvcreate` or `lvchange` command sets the scheduling policy to parallel for a logical volume.

Sequential Scheduling

The sequential scheduling policy causes mirror write operations to proceed sequentially; that is, LVM waits for one mirror write to complete before it begins the next mirror write. Likewise, LVM mirrors are read in a predefined order. On a practical level, sequential policy is used only for extreme caution in maintaining consistency of mirrors. The `-d s` option to the `lvcreate` or `lvchange` command sets the scheduling policy to sequential for a logical volume.

Synchronization Policy

You can maintain consistency of mirrored data by enabling or disabling two features of your logical volume: the Mirror Write Cache and Mirror Consistency Recovery.

Synchronization Using Mirror Write Cache

The **Mirror Write Cache (MWC)** provides a fast resynchronization of data following a system crash or failure, but at a potential performance cost for routine system use.

The MWC keeps track of where I/O writes are occurring on the volume group, and periodically records this activity in an ondisk data structure. An extra disk write is required for every mirrored write not already recorded on the physical volume. This slows down runtime I/O write processing and degrades performance when you access the disk at random; when writing to an area of the disk that is already recorded, the performance is not impaired. Upon system reboot after crash, the operating system uses the MWC to resynchronize inconsistent data blocks quickly.

The frequency of extra disk writes is small for sequentially accessed logical volumes (such as database logs), but increases when access is more random. Therefore, logical volumes containing

database data or file systems with few or infrequently written large files (greater than 256K) must not use the MWC when runtime performance is more important than crash recovery time.

The `-M` option to the `lvcreate` or `lvchange` command controls the MWC.

Synchronization Using Mirror Consistency Recovery

When the **Mirror Consistency Recovery** is enabled, LVM does not impact runtime I/O performance. However, following a system crash, for any logical volumes using Mirror Consistency Recovery, the entire data space resynchronizes when you activate the volume group. Synchronization can be performed in the background without interfering with reboot or access; however, during this time I/O performance and redundancy are degraded.

Synchronization with No Mirror Consistency Mechanism

When Mirror Consistency Recovery is disabled, the operating system's runtime behavior is identical to that of the previous approach. However, following a crash, LVM does *not* perform any resynchronization of data. This approach is useful for swap volumes and for volumes used by an application (such as a database) with its own means of maintaining or recovering consistent data, such as transaction log files. However, database log files themselves can be configured as a mirrored logical volume to use the MWC.

The `-c` option to the `lvcreate` or `lvchange` command controls the use of the Mirror Consistency Recovery.

Synchronizing a Mirrored Logical Volume

The data in a mirrored copy or copies of a logical volume can become out of sync, or "stale." For example, mirrored data becomes stale if LVM cannot access a disk as a result of disk power failure. Under such circumstances, for each mirrored copy to re-establish identical data, synchronization must occur. Usually, synchronization occurs automatically, but sometimes it must be done manually.

Automatic Synchronization

If you activate a volume group that is *not currently active*, either automatically at boot time or later with the `vgchange` command, then LVM automatically synchronizes the mirrored copies of all logical volumes with the Mirror Consistency Recovery policy enabled. It replaces data in physical extents marked as stale with data from nonstale extents. Otherwise, no automatic synchronization occurs and manual synchronization is necessary.

LVM also automatically synchronizes mirrored data in the following cases:

- When you increase the number of mirror copies of a logical volume using the `-m` option of `lvmmerge`, the newly added physical extents are synchronized.
- When a disk comes back online after experiencing a power failure.

Manual Synchronization

If you look at the status of a logical volume using `lvdisplay -v`, you can verify if the logical volume contains any stale data. You can then identify which disk contains the stale physical extents. Manually synchronize the data in one or more logical volumes using either the `lvsync` command or all logical volumes in one or more volume groups using the `vgsync` command. For more information, see `lvdisplay(1M)`, `vgsync(1M)`, and `lvsync(1M)`.

Parallel Synchronization

By default, the `lvsync` command synchronizes logical volumes serially. In other words, it acts on the logical volumes specified on the command line one at a time, waiting until a volume finishes synchronization before starting the next. Starting with the September 2007 release of HP-UX 11i Version 3, you can use the `-T` option to synchronize logical volumes in parallel. With the `-T` option, `lvsync` spawns multiple threads to simultaneously synchronize all logical volumes belonging to the same volume group, often reducing the total synchronization time.

TIP: The `vgchange`, `lvmerge`, and `lvextend` commands support the `-s` option to suppress the automatic synchronization of stale extents. If you are performing multiple mirror-related tasks, you can suppress the extent synchronization until you have finished all the tasks, then run `lvsync` with `-T` to synchronize all the mirrored volumes in parallel. For example, you can use `vgchange -s` with `lvsync -T` to reduce the activation time for volume groups with mirrored logical volumes. For another example, see “[Mirroring the Boot Disk](#)” (page 90).

Increasing Disk Redundancy Through Disk Sparing

NOTE: Version 2.x volume groups do not support disk sparing.

Disk sparing requires the optional product HP MirrorDisk/UX.

MirrorDisk/UX is not available for shared LVM environments within a high availability cluster across more than two nodes. You cannot configure sparing within these environments. In such cases, HP recommends that you use hardware mirroring through RAID devices, which can support their own form of sparing.

If a disk containing mirrored data fails, replace the disk as soon as possible, as described in “[Disk Troubleshooting and Recovery Procedures](#)” (page 116). Before you replace the disk, the data in your logical volume does not have an extra mirrored copy unless you have set up more than one mirror copy. Even with multi-way mirroring, your level of security is reduced because of the loss of one mirror copy.

To prevent this possibility, you can use one or more spare disks within each of your volume groups to serve as substitute devices in the event of disk failure. With this configuration, LVM automatically reconfigures the volume group so that the spare physical volume takes the place of a failed device without any intervention required. That is, a copy of the data from all the logical volumes currently on the failed disk is created on the substitute physical volume. This process is referred to as automatic sparing, or **sparing**. Sparing occurs while the logical volume remains available to users. You can then schedule the replacement of the failed disk at a time of minimal inconvenience to you and your users. At that time, you copy the data from the spare disk back to the original disk or its replacement and return the spare disk to its role as a standby empty disk.

For sparing to occur, the following conditions must be met:

- All logical volumes in the volume group must have been configured with strict mirroring so that mirrored copies are maintained on separate disks because LVM copies the data onto the spare from an undamaged disk rather than from the defective disk itself.
- At least one physical volume must be available as a standby spare; if your last spare is already in use as a result of a prior disk failure, it cannot serve as a currently available spare.
- The available spare must be at least as large as the failed disk.

The spare physical volume disk space is not available for extent allocation for any other purpose than in the event of serving as a substitute disk in the event of disk failure. Therefore, its physical extents are not included in the counts shown under `total PE` or `free PE` in the output of the `pvdisplay` and `vgdisplay` commands.

NOTE: If it is important to maintain comparable performance in the event of disk failure, configure a spare physical volume to each bus. However, if more than one disk on the same bus fails, even with this strategy, performance will be impacted.

The `pvdisplay` and `vgdisplay` commands provide information on whether a given physical volume is an empty standby spare or currently holding data as a spare in use, along with information on any physical volume that is currently unavailable but had data spared.

Increasing Hardware Path Redundancy Through Multipathing

Your hardware might provide the capability for dual cabling (dual controllers) to the same physical volume. If so, LVM can be configured with multiple paths to the same physical volume. If the primary link fails, an automatic switch to an alternate link occurs. Using this type of multipathing (also called pvlinks) increases availability.

NOTE:

As of HP-UX 11i Version 3, the mass storage stack supports native multipathing without using LVM pvlinks. Native multipathing provides more load balancing algorithms and path management options than LVM. HP recommends using native multipathing to manage multipathed devices instead of using LVM's alternate links.

HP recommends converting volume groups with alternate links to use native multipathing by the use of persistent DSFs. The `/usr/contrib/bin/vgdsf` script, `vgscan -N` command, or `vgimport -s -N` commands perform this conversion.

For backward compatibility, you can use existing pvlinks. However, you must use legacy device special files for physical volumes and disable native multipathing for those legacy device special files using the `scsimgr` command. For more information, see the white paper *LVM Migration from Legacy to Agile Naming Model*, available at <http://www.hp.com/go/hpux-core-docs>.

Setting Up Multipathing to a Physical Volume (LVM PVlinks)

To use an alternate link, you can create a volume group with `vgcreate`, specifying both the primary link and the alternate link device file names. Both links must represent paths to the same physical volume. (Do not run `pvcreate` on the alternate link; it must already be the same physical volume as the primary link.) When you indicate two device file names, both referring to the same disk using `vgcreate`, LVM configures the first one as the primary link and the second one as the alternate link.

For example, if a disk has two cables and you want to make one the primary link and the other an alternate link, enter the following command:

```
# vgcreate /dev/vg01 /dev/dsk/c3t0d0 /dev/dsk/c5t0d0
```

To add an alternate link to a physical volume that is already part of a volume group, use `vgextend` to indicate the new link to the physical volume. For example, if `/dev/dsk/c2t0d0` is already part of your volume group but you want to add another connection to the physical volume, enter the following command:

```
# vgextend /dev/vg02 /dev/dsk/c4t0d0
```

If the primary link fails, LVM automatically switches from the primary controller to the alternate controller. However, you can also tell LVM to switch to a different controller at any time using the `pvchange` command. For example:

```
# pvchange -s /dev/dsk/c2t1d0
```

After the primary link has recovered, LVM automatically switches back from the alternate controller to the original controller unless you previously instructed it not to by using `pvchange` as follows:

```
# pvchange -s n /dev/dsk/c2t2d0
```

NOTE: You can also disable the automatic switchback by using the `-p` option to `pvchange` to disable proactive polling. For more information, see `pvchange(1M)`.

View the current links to a physical volume using `vgdisplay` with the `-v` option.

Planning for Performance

This section describes strategies to obtain the best possible performance using LVM. It addresses the following topics:

- “General Performance Factors” (page 29)
- “Internal Performance Factors” (page 29)
- “Increasing Performance Through Disk Striping” (page 31)
- “Increasing Performance Through I/O Channel Separation” (page 33)

General Performance Factors

The following factors affect overall system performance, but not necessarily the performance of LVM.

Memory Usage

The amount of memory used by LVM is based on the values used at volume group creation time and on the number of open logical volumes. The largest portion of LVM memory is used for extent maps. The memory used is proportional to the maximum number of physical volumes multiplied by the maximum number of physical extents per physical volume for each volume group.

The other factors to be concerned with regarding memory parameters are expected system growth and number of logical volumes required. You can set the volume group maximum parameters to exactly what is required on the system today. However, if you want to extend the volume group by another disk (or perhaps replace one disk with a larger disk), you must use the `vgmodify` command.

CPU Usage

Compared to the non-LVM case, no significant impact to system CPU usage (by observing idle time) has been observed.

With LVM, extra CPU cycles are required to perform mirror write consistency cache operations, which is the only configurable option that impacts CPU usage.

Disk Space Usage

LVM reserves some disk space on each physical volume for its own metadata. The amount of space used is proportional to the maximum values used at volume group creation time.

Internal Performance Factors

The following factors directly affect the performance of I/O through LVM.

Scheduling Policy

The scheduling policy is significant only with mirroring. When mirroring, the sequential scheduling policy requires more time to perform writes proportional to the number of mirrors. For instance, a logical volume with three copies of data requires three times as long to perform a write using the sequential scheduling policy, as compared to the parallel policy. Read requests are always directed to only one device. Under the parallel scheduling policy, LVM directs each read request to the least busy device. Under the sequential scheduling policy, LVM directs all read requests to the device shown on the left hand side of an `lvdisplay -v` output.

Mirror Write Consistency Cache

The purpose of the Mirror Write Consistency cache (MWC) is to provide a list of mirrored areas that might be out of sync. When a volume group is activated, LVM copies all areas with an entry in the MWC from one of the good copies to all the other copies. This process ensures that the mirrors are consistent but does not guarantee the quality of the data.

On each write request to a mirrored logical volume that uses MWC, LVM potentially introduces one extra serial disk write to maintain the MWC. Whether this condition occurs depends on the degree to which accesses are random.

The more random the accesses, the higher the probability of missing the MWC. Getting an MWC entry can involve waiting for one to be available. If all the MWC entries are currently being used by I/O in progress, a given request might have to wait in a queue of requests until an entry becomes available.

Another performance consideration for mirrored logical volumes is the method of reconciling inconsistencies between mirror copies after a system crash. Two methods of resynchronization are available: Mirror Consistency Recovery (MCR) and none. Whether you use the MWC depends on which aspect of system performance is more important to your environment, run time or recovery time.

For example, a customer using mirroring on a database system might choose "none" for the database logical volume because the database logging mechanism already provides consistency recovery. The logical volume used for the log uses the MWC if quick recovery time was an issue, or MCR if higher runtime performance is required. A database log is typically used by one process and is sequentially accessed, which means it suffers little performance degradation using MWC because the cache is hit most of the time.

Disk Spanning

For disk areas that see the most intensive use by multiple processes, HP recommends spreading the data space for this disk area across as many physical volumes as possible.

Number of Volume Groups

The number of volume groups is directly related to the MWC issues. Because there is only one MWC per volume group, disk space that is used for many small random write requests must be kept in distinct volume groups if possible when the MWC is being used. This is the only performance consideration that affects the decision regarding the number of volume groups.

Physical Volume Groups

This factor can be used to enforce the separation of different mirror copies across I/O channels. You must define the physical volume groups. This factor increases the availability by decreasing the single points of failure and provides faster I/O throughput because of less contention at the hardware level.

For example, in a system with several disk devices on each card and several cards on each bus converter, create physical volume groups so that all disks off of one bus converter are in one group and all the disks on the other are in another group. This configuration ensures that all mirrors are created with devices accessed through different I/O paths.

Snapshots and Performance

Logical volumes that have snapshots associated with them could experience an increase in latencies associated with writes. This is also applicable for writes on (writable) snapshots themselves. The increased latencies are seen only for writes that fall on regions (unshare units) that are shared with a successor/predecessor. The reason for this increased latency is the data unsharing that has to precede the actual writes on the original logical volume or a snapshot. Note that the increase in latencies is reduced as data gets unshared between the original logical volume and its snapshots.

Performance is also dependent on the unshare unit size configured for the snapshots of a logical volume. The supported values of the snapshot unshare unit are 512K, 1024K, 2048K, and 4096K. While choosing the unshare unit size one has to consider the kind of application that will be writing to the snapshots or the original logical volume, and also consider the tradeoff between performance requirements and metadata space that will have to be provisioned on disk.

The larger the unit of unshare, the more will the latency be when data has to be unshared and the lesser the metadata space needed on disk to track the sharing relationship between snapshots and the original logical volume. If the application performs large and occasional writes, then it is recommended that a larger unshare unit is used. If the writes are small and occasional, then a smaller unshare unit is recommended. If the write pattern is going to be such that they are small and sequential, then again a large unshare unit is recommended because after the first unshare, the following consecutive I/Os that fall within the same unshare unit will not incur the overhead of a unshare operation.

Refer to the *LVM Snapshot Logical Volumes* white paper for more information about snapshots and performance.

Increasing Performance Through Disk Striping

Disk striping distributes logically contiguous data blocks (for example, chunks of the same file) across multiple disks, which speeds I/O throughput for large files when they are read and written sequentially (but not necessarily when access is random).

The disadvantage of disk striping is that the loss of a single disk can result in damage to many files because files are purposely spread across two or more disks.

Consider using disk striping on file systems where large files are stored, if those files are normally read and written sequentially and I/O performance is important.

When you use disk striping, you create a logical volume that spans multiple disks, allowing successive blocks of data to go to logical extents on different disks. For example, a three-way striped logical volume has data allocated on three disks, with each disk storing every third block of data. The size of each of these blocks is called the **stripe size** of the logical volume. The stripe size (in K) must be a power of two in the range 4 to 32768 for a Version 1.0 volume group, and a power of two in the range 4 to 262144 for a Version 2.x volume group.

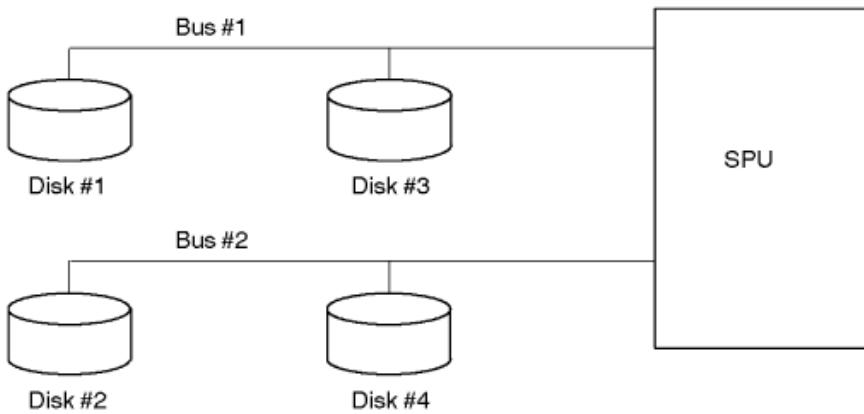
Disk striping can increase the performance of applications that read and write *large, sequentially accessed files*. Data access is performed over the multiple disks simultaneously, resulting in a decreased amount of required time as compared to the same operation on a single disk. If all of the striped disks have their own controllers, each can process data simultaneously.

You can use standard commands to manage your striped logical volumes. For example, the `lvcreate`, `diskinfo`, `newfs`, `fsck`, and `mount` commands all work with striped logical volumes.

The following guidelines, most of which apply to LVM disk usage, apply to striped logical volumes for performance reasons:

- Best performance results from a striped logical volume that spans similar disks. The more closely you match the disks by speed, capacity, and interface type, the better the performance you can expect. When striping across several disks of varying speeds, performance is no faster than that of the *slowest* disk.
- If you have more than one interface card or bus to which you can connect disks, distribute the disks as evenly as possible among them. That is, each interface card or bus must have roughly the same number of disks attached to it. You can achieve the best I/O performance when you use more than one bus and interleave the stripes of the logical volume. For example, if you have two buses with two disks on each bus, order the disks so that disk 1 is on bus 1, disk 2 is on bus 2, disk 3 is on bus 1, and disk 4 is on bus 2, as shown in [Figure 4](#).

Figure 4 Interleaving Disks Among Buses



- Increasing the number of disks might not improve performance because the maximum efficiency that can be achieved by combining disks in a striped logical volume is limited by the maximum throughput of the file system itself and by the buses to which the disks are attached.
- Disk striping is highly beneficial for applications with few users and large, sequential transfers. However, applications that exhibit small, concurrent, random I/O (such as databases) often see no performance gain through disk striping. Consider four disks with a stripe size of 512 bytes. Each 2K request is sent to all disks. One 2K request completes in about the same time when the disk has been striped. However, several 2K requests are all serialized because all the disks must seek for each request. On the nonstriped system, performance might actually be better because each disk might service separate requests in parallel.

Determining Optimum Stripe Size

The logical volume stripe size identifies the size of each of the blocks of data that make up the stripe. You can set the stripe size to a power of two in the range 4 to 32768 for a Version 1.0 volume group, or a power of two in the range 4 to 262144 for a Version 2.x volume group. The default is 8192.

NOTE: The stripe size of a logical volume is not related to the physical sector size of a disk, which is typically 512 bytes.

How you intend to use the striped logical volume determines what stripe size you assign to it.

For best results follow these guidelines:

- If you plan to use the striped logical volume for an HFS file system, select the stripe size that most closely reflects the block size of the file system. The `newfs` command enables you to specify a block size when you build the file system and provides a default block size of 8K for HFS.
- If you plan to use the striped logical volume for a JFS (VxFS) file system, use the largest available size, 64K. For I/O purposes, VxFS combines blocks into extents, which are variable in size and may be very large. The configured block size, 1K by default, is not significant in this context.
- If you plan to use the striped logical volume as swap space, set the stripe size to 16K for best performance. For more information on configuring swap, see “[Administering Swap Logical Volumes](#)” (page 101).
- If you plan to use the striped logical volume as a raw data partition (for example, for a database application that uses the device directly), the stripe size must be as large or larger than the I/O size for the application.

You might need to experiment to determine the optimum stripe size for your particular situation. To change the stripe size, re-create the logical volume.

Interactions Between Mirroring and Striping

Mirroring a striped logical volume improves the read I/O performance in a same way that it does for a nonstriped logical volume. Simultaneous read I/O requests targeting a single logical extent are served by two or three different physical volumes instead of one. A striped and mirrored logical volume follows a strict allocation policy; that is, the data is always mirrored on different physical volumes.

The appendix “[Striped and Mirrored Logical Volumes](#)” (page 157) provides information about logical volumes that are *both* striped and mirrored.

Increasing Performance Through I/O Channel Separation

[I/O channel separation](#) is an approach to LVM configuration requiring that mirrored copies of data reside on LVM disks accessed using separate host bus adapters (HBAs) and cables. I/O channel separation achieves higher availability and better performance by reducing the number of single points of possible hardware failure. If you mirror data on two separate disks, but through one card, your system can fail if the card fails.

You can separate I/O channels on a system with multiple HBAs and a single bus, by mirroring disks across different HBAs. You can further ensure channel separation by establishing a policy called **PVG-strict** allocation, which requires logical extents to be mirrored in separate physical volume groups. **Physical volume groups** are subgroups of physical volumes within a volume group.

An ASCII file, `/etc/lvmpvg`, contains all the mapping information for the physical volume group, but the mapping is not recorded on disk. Physical volume groups have no fixed naming convention; you can name them PVG0, PVG1, and so on. The `/etc/lvmpvg` file is created and updated using the `vgcreate`, `vgextend`, and `vgreduce` commands, but you can edit the file with a text editor.

I/O channel separation is useful for databases, because it heightens availability (LVM has more flexibility in reading data on the most accessible logical extent), resulting in better performance. If you define your physical volume groups to span I/O devices, you ensure against data loss even if one HBA fails.

When using physical volume groups, consider using a PVG-strict allocation policy for logical volumes.

Configuring LVM I/O Timeout Parameters

You can configure the following timeout parameters for LVM I/Os.

- **Logical volume timeout (LV timeout).**

This controls how long LVM retries a logical I/O after a recoverable physical I/O error. LV timeout can be configured for a specific logical volume using `lvchange`.

- **Physical volume timeout (PV timeout).**

This is the time budget set by LVM for each physical I/O originating from LVM. The mass storage stack underlying LVM manages completion of the physical I/O subject to this timeout, resulting in success or failure for the request. PV timeout can be configured for a specific physical volume using `pvchange`.

- **Mass storage stack timeout parameters**

There are mass storage stack tunables that affect LVM I/O timeout behavior: `path_fail_secs` and `transient_secs`.

For details on these timeout parameters, refer to Appendix “[LVM I/O Timeout Parameters](#)” (page 161).

Planning for Recovery

Flexibility in configuration, one of the major benefits of LVM, can also be a source of problems in recovery. The following are guidelines to help create a configuration that minimizes recovery time:

- Keep the number of disks in the root volume group to a minimum; HP recommends using three disks, even if the root volume group is mirrored.

Root volume groups with many disks make reinstallation difficult because of the complexity of recovering LVM configurations of accessory disks within the root volume group.

A small root volume group is quickly recovered. In some cases, you can reinstall a minimal system, restore a backup, and be back online within three hours of diagnosis and replacement of hardware. Another benefit is that exact match to the previous root disk layout is not required.

Three disks in the root volume group are better than two, because of quorum restrictions. With a two-disk root volume group, the loss of one disk can require you to override quorum to activate the volume group; if you must reboot to replace the disk, overriding quorum requires you to interrupt the boot process. If you have three disks in the volume group and they are isolated from each other such that a hardware failure only affects one of them, then failure of only one disk enables the system to maintain quorum.

There are two reasons to expand the root volume group beyond a minimal size.

- A very small root disk.

In this case, HP recommends migrating or installing to a larger disk.

- Providing for dump-to-swap for large memory systems.

Swap volumes targeted for dump must be in the root volume group. A better solution is to configure an extra dedicated disk for dump up front.

- Do not use network-based backup programs such as Omniback or Networker for basic root volume group backups. The complexity associated with these utilities can significantly delay the resumption of processing.

HP recommends using Ignite-UX and the `make_net_recovery` command to back up and recover your root volume group.

- Create adequate documentation.

Output from `ioscan -kf`, `vgcfgrestore -lv` for all groups and `vgscan -pv`, and `lvolboot -v` are very minimal requirements. Recovery from almost any problem is possible if these and output from `vgdisplay -v` for all groups and `lvdisplay -v` for all volumes are available. Extent mappings are critical to recovery of LVM volumes with corrupted headers. Additionally, output from `pvdisplay -v` for all physical volumes, although not as important, provides complete volume group information. Hard copy is not required or even necessarily practical, but accessibility during recovery is important and should be planned for.

- Many small groups are better than fewer large groups when configuring auxiliary volume groups.

Reloading dozens of gigabytes of data because one disk is missing out of a group after a root disk rebuild can be difficult. Also the scope of a catastrophic single disk failure within a group is minimized with many small groups.

Size implies complexity. The larger the number of disks in a single group, the more chances an administrator has to create an error that affects the entire group. It is simpler to identify and import smaller groups if necessary. It is also simpler to conceptualize and map smaller groups when required.

Preparing for LVM System Recovery

To ensure that the system data and configuration are recoverable in the event of a system failure, follow these steps:

1. Load any patches for LVM.
2. Use Ignite-UX to create a recovery image of your root volume group. Although Ignite-UX is not intended to be used to back up all system data, you can use it with other data recovery applications to create a method of total system recovery.
3. Perform regular backups of the other important data on your system.
Without a valid backup, you risk losing some or all of your data.
4. Regularly print out your system configuration.

The configuration details stored on the system might not be accessible during a recovery. A printed copy is an invaluable reference. HP recommends printing the configuration details once a week *and* every time a change is made. Some of the commands create large amounts of output. An alternative to printing them is to output the information to a file and then storing the file on tape which enables quick recovery of the information when needed. You can include this configuration file with the backup in step 3.

The easiest way to save the configuration is to set up a cron job to run regularly, so that the system backs it up automatically.

Use the following commands to get a useful output:

```
/usr/sbin/ioscan -fk  
/usr/sbin/vgdisplay -v  
/usr/sbin/lvlnboot -v  
/usr/sbin/lvdisplay -v /dev/vgXX/lvYY (for every logical volume)  
/usr/sbin/pvdisplay -v /dev/dsk/c#t#d0 (for every LVM disk)  
lp /etc/fstab
```

As an alternative, you can write an intelligent script that detects any changes in the configuration and only print out those changes. An example script is included at the end of this section.

5. Back up the LVM configuration after every configuration change.

The `vgcfgbackup` command copies the LVM headers from the system area of the disk to a disk file, which resides in the `/etc/lvmconf` directory. This default directory can be changed for volume group Version 2.x by configuring a new path in the `LVMP_CONF_PATH_NON_BOOT` variable in the `/etc/lvmrc` file (see `vgcfgbackup(1M)` for more information). After this information is stored in a disk file, it can be backed up to tape during file system backups.

The information in this file enables you to replace the LVM headers on the disk in the event of the disk being replaced or your LVM configuration becoming corrupted.

It is important that these configuration backups are taken whenever you make a change to any part of the LVM configuration. By default all the commands perform a backup so a manual `vfcfgbackup` after each command is not required.

Do this task on a regular basis, whether you have made changes or not. Enter the following command:

```
# /usr/sbin/vgcfgbackup /dev/vgXX (for every volume group)
```

NOTE: For volume group Version 2.2 and higher: volume groups that have snapshots on which data unsharing is occurring, the LVM configuration backup file might not always be in sync with the LVM metadata on disk. LVM ensures that the configuration for volume group Version 2.2 and higher is the latest by automatically backing it up during deactivation, unless backup has been disabled by the `-A n` option. If the configuration backup during deactivation will be bypassed, then to ensure that the volume group's configuration is the latest, it is recommended that you close all the logical volumes of the volume group and back up the configuration using `vgcfgbackup` prior to deactivating the volume group.

6. Update the boot structures after every change to the root volume group.

This task is only required if you are using LVM on your boot disk. Whenever you make changes to the root volume group, which is usually named `/dev/vg00`, the BDRA on the boot disk must be updated. This is normally performed automatically by the LVM commands. To update the BDRA manually, enter the following command:

```
# /usr/sbin/lvlnboot -R
```

See [Chapter 4: Troubleshooting LVM](#) for detailed information about recovery.

Example Script for LVM Configuration Recording

The following example script captures the current LVM and I/O configurations. If they differ from the previously captured configuration, the script prints the updated configuration files and notifies the system administrator.

```
#!/usr/bin/ksh
WORKDIR=/lvmbbackup # directory is regularly backed up
LOG=$WORKDIR/log
SYSADM=root
if [ -f "$LOG" ]
then
    rm -f "$LOG"
fi
if [ ! -d "$WORKDIR" ]
then
    echo "missing directory $WORKDIR" exit 1
fi
cd $WORKDIR
/usr/sbin/vgdisplay -v -F > vgdisplay.new
LVMVGS=`grep vg_name vgdisplay.new | cut -d: -f1 | cut -d= -f2` 
LVMPVOLS=`grep pv_name vgdisplay.new | cut -d: -f1 | cut -d= -f2 | cut -d, -f1` 
LVMLVOLS=`grep lv_name vgdisplay.new | cut -d: -f1 | cut -d= -f2` 
/usr/sbin/pvdisplay -v $LVMPVOLS > pvdisplay.new
/usr/sbin/lvdisplay -v $LVMLVOLS > lvdisplay.new
/usr/sbin/lvlnboot -v > lvlnboot.new 2> /dev/null
/usr/sbin/iосan -fk > iосan.new
cp /etc/fstab fstab.new
for CURRENT in *new.
do
    ORIG=${CURRENT%.new}
    if diff $CURRENT $ORIG > /dev/null then
        # files are the same....do nothing
        rm $CURRENT
    else
        # files differ...make the new file the current file, move old
        # one to file.old.
        echo `date` "The config for $ORIG has changed." >> $LOG
        echo "Copy of the new $ORIG config has been printed" >> $LOG
        lp $CURRENT
        mv $ORIG ${ORIG}old.
        mv $CURRENT $ORIG
    fi
done
if [ -s "$LOG" ]
then
    mailx -s "LVM configs have changed" $SYSADM < $LOG
fi
exit 0
```

3 Administering LVM

This section contains information on the day-to-day operation of LVM. It addresses the following topics:

- “Administration Tools” (page 38)
- “Displaying LVM Information” (page 40)
- “Common LVM Tasks” (page 43)
- “Moving and Reconfiguring Your Disks” (page 70)
- “Migrating a Volume Group to New Disks: `vgmove`” (page 95)
- “Migrating a Logical Volume to New Disks: `lvmove`” (page 96)
- “Administering File System Logical Volumes” (page 97)
- “Administering Swap Logical Volumes” (page 101)
- “Administering Dump Logical Volumes” (page 102)
- “Creating and Administering Snapshot Logical Volumes” (page 103)
- “Managing the `lvmpud` Daemon” Page 105
- “Hardware Issues” (page 106)

Administration Tools

HP-UX provides two tools to manage your LVM configuration:

- **HP System Management Homepage (HP SMH):** An HP-UX tool that provides an easy-to-use GUI for performing most LVM tasks. HP SMH minimizes or eliminates the need for detailed knowledge of administration commands, thus saving time. Use HP SMH to manage your LVM configuration whenever possible, especially when performing a new task. HP SMH can help you with the following tasks:
 - Creating or removing volume groups
 - Adding or removing disks within volume groups
 - Activating and deactivating volume groups
 - Exporting and importing volume groups
 - Creating, removing, or modifying logical volumes
 - Increasing the size of logical volumes
 - Creating or increasing the size of a file system in a logical volume
 - Creating and modifying mirrored logical volumes
 - Creating striped logical volumes
 - Splitting a mirrored logical volume and merging a mirror copy
 - Adding and removing mirror copies within mirrored logical volumes
 - Creating and modifying physical volume groups

To use HP SMH, enter the command `/usr/sbin/smh`.

For help using HP SMH, see the HP SMH online help.

- **LVM command-line interface:** LVM has a number of low-level user commands to perform LVM tasks, described in “Physical Volume Management Commands” (page 39), “Volume Group

[Management Commands](#) (page 39), and [Logical Volume Management Commands](#) (page 40).

The following tables provide an overview of which commands perform a given task. For more information, see the LVM individual manpages.

Table 4 Physical Volume Management Commands

Task	Command
Changing the characteristics of a physical volume in a volume group	<code>pvchange</code>
Creating a physical volume for use in a volume group	<code>pvcreate</code>
Displaying information about physical volumes in a volume group	<code>pvdisplay</code>
Moving data from one physical volume to another	<code>pvmove</code>
Removing a physical volume from LVM control	<code>pvremove</code>
Checking or repairing a physical volume	<code>pvck</code> ¹
Checking if a disk volume is under LVM control	<code>lvmchk</code>

1 The `pvck` command is not supported on Version 2.x volume groups.

Table 5 Volume Group Management Commands

Task	Command
Creating a volume group	<code>vgcreate</code>
Removing a volume group	<code>vgremove</code>
Activating, deactivating, or changing the characteristics of a volume group	<code>vgchange</code>
Modifying the configuration parameters of a volume group; handling physical volume size changes	<code>vgmodify</code>
Backing up volume group configuration information	<code>vgcfgbackup</code>
Restoring volume group configuration from a configuration file	<code>vgcfgrestore</code>
Displaying information about volume groups	<code>vgdisplay</code>
Exporting a volume group and its associated logical volumes	<code>vgexport</code>
Importing a volume group onto the system; adding an existing volume group back into the LVM configuration files	<code>vgimport</code>
Scanning all physical volumes for logical volumes and volume groups; recovering the LVM configuration files	<code>vgscan</code>
Adding a disk to a volume group	<code>vgextend</code>
Removing a disk from a volume group	<code>vgreduce</code>
Synchronizing mirrored logical volumes in a volume group	<code>vgsync</code>
Modifying the volume group ID on a physical volume	<code>vgchgid</code>
Migrating a volume group from legacy to persistent device files	<code>vgdsv</code>
Converting persistent device special files in a volume group to cluster device special files. (See note after table.)	<code>vgcdsf</code> ¹
Displaying limits associated with a volume group version	<code>lvmadm</code>
Handling online shared LVM reconfiguration, and pre-allocation of extents for space-efficient snapshots	<code>lvpmpud</code>

Table 5 Volume Group Management Commands (continued)

Task	Command
Migrating a volume group to a different volume group version	vgversion
Migrating a volume group to new disks	vgmove

1 To convert the cDSFs of a volume group in a particular node back to their corresponding persistent DSFs, use the vgscan -f command. For example:

```
# vgscan -f vgtest  
*** LVMTAB has been updated successfully.
```

Repeat the above command on all the nodes in the cluster where this shared volume group is present (i.e., imported). Use the lvmadm -l command to verify the conversion of cDSFs to persistent DSF.

Table 6 Logical Volume Management Commands

Task	Command
Creating a logical volume	lvcreate
Modifying a logical volume	lvchange
Displaying information about logical volumes	lvdisplay
Increasing the size of a logical volume by allocating disk space	lvextend
Migrating a logical volume to new disks	lvmmove
Decreasing the size of a logical volume	lvreduce
Removing the allocation of disk space for one or more logical volumes within a volume group	lvremove
Preparing a logical volume to be a root, primary swap, or dump volume; updating the boot information on the boot physical volume	lvlnboot ¹
Removing the link that makes a logical volume a root, primary swap, or dump volume	lvrmboot ¹
Splitting a mirrored logical volume into two logical volumes	lvsplit
Merging two logical volumes into one mirrored logical volume	lvmerge
Synchronizing mirror copies in a mirrored logical volume	lvsync

1 The lvlnboot and lvrmboot commands are not supported on Version 2.0 and 2.1 volume groups.

The command-line interface is more powerful—and thus more dangerous—than HP SMH and offers options that are not available using HP SMH. For example, the following tasks cannot currently be done by HP SMH. For these tasks, use the HP-UX commands:

- “[Creating a Spare Disk](#)” (page 76)
- “[Reinstating a Spare Disk](#)” (page 77)
- “[Synchronizing a Mirrored Logical Volume](#)” (page 26)
- “[Moving Data to a Different Physical Volume](#)” (page 73)
- “[Disk Troubleshooting and Recovery Procedures](#)” (page 116)

The remainder of this chapter explains how to do LVM tasks using HP-UX commands. However, HP SMH is the tool of choice for most administration work.

Displaying LVM Information

To display information about volume groups, logical volumes, or physical volumes, use one of three commands. Each command supports the -v option to display detailed output and the -F option to help with scripting.

NOTE: For volume group Version 2.2 or higher, when snapshots are involved, additional fields are displayed by these commands. See the individual command manpages for full description of the fields displayed. The “[Creating and Administering Snapshot Logical Volumes](#)” (page 103) section provides a summary of additional data displayed for snapshots.

Information on Volume Groups

Use the `vgdisplay` command to show information about volume groups. For example:

```
# vgdisplay -v vg01
-- Volume groups --
VG Name          /dev/vg01
VG Write Access  read/write
VG Status        available
Max LV          255
Cur LV          1
Open LV          1
Max PV          16
Cur PV          1
Act PV          1
Max PE per PV   1016
VGDA            2
PE Size (Mbytes) 4
Total PE        508
Alloc PE        508
Free PE         0
Total PVG       0
Total Spare PVs 0
Total Spare PVs in use 0
VG Version      1.0
VG Max Size    1082g
VG Max Extents 69248

-- Logical volumes --
LV Name          /dev/vg01/lvol1
LV Status        available/syncd
LV Size (Mbytes) 2032
Current LE       125
Allocated PE     508
Used PV          1

-- Physical volumes --
PV Name          /dev/disk/disk42
PV Status        available
Total PE        508
Free PE         0
Autoswitch      On
Proactive Polling On
```

Use the `vgdisplay` command to verify whether the LVM configuration in memory has problems. If the LVM configuration is working properly, there are no error messages, and the display shows the following:

- The status is available (or available/exclusive for Serviceguard volume groups).
- All the physical volumes are active; that is, the current number of physical volumes (`Cur PV`) is equal to number of active physical volumes (`Act PV`).
- All the logical volumes are open; that is, `Cur LV` is equal to `Open LV`.

Information on Physical Volumes

Use the `pvdisplay` command to show information about physical volumes. For example:

```

# pvdisplay -v /dev/disk/disk47
-- Physical volumes --
PV Name                  /dev/disk/disk47
VG Name                  /dev/vg00
PV Status                available
Allocatable              yes
VGDA                     2
Cur LV                   9
PE Size (Mbytes)        4
Total PE                 1023
Free PE                  494
Allocated PE             529
Stale PE                 0
IO Timeout (Seconds)    default
Autoswitch               On
Proactive Polling        On

-- Distribution of physical volume --
LV Name      LE of LV  PE for LV
/dev/vg00/lvol1 25      25
/dev/vg00/lvol2 25      25
/dev/vg00/lvol3 50      50

--- Physical extents ---
PE  Status   LV          LE
0000 current  /dev/vg00/lvol1 0000
0001 current  /dev/vg00/lvol1 0001
0002 current  /dev/vg00/lvol1 0002
...
1021 free     0000
1022 free     0000

```

If the physical volume is functioning properly, Stale PE is 0.

Information on Logical Volumes

Use the lvdisplay command to show information about logical volumes. For example:

```

# lvdisplay -v /dev/vg00/lvol1
-- Logical volumes --
LV Name                  /dev/vg00/lvol1
VG Name                  /dev/vg00
LV Permission            read/write
LV Status                available/syncd
Mirror copies            0
Consistency Recovery    MWC
Schedule                 parallel
LV Size (Mbytes)         100
Current LE                25
Allocated PE              25
Stripes                  0
Stripe Size (Kbytes)    0
Bad block                 off
Allocation               strict/contiguous
IO Timeout (Seconds)    default

-- Distribution of logical volume --
PV Name      LE on PV  PE on PV
/dev/disk/disk42 25      25

-- Logical extents --
LE  PV1          PE1  Status 1
0000 /dev/disk/disk42 0000 current
0001 /dev/disk/disk42 0001 current
0002 /dev/disk/disk42 0002 current

```

Common LVM Tasks

The section addresses the following topics:

- “[Initializing a Disk for LVM Use](#)” (page 43)
- “[Creating a Volume Group](#)” (page 44)
- “[Migrating a Volume Group to a Different Version: vgversion](#)” (page 46)
- “[Adding a Disk to a Volume Group](#)” (page 51)
- “[Removing a Disk from a Volume Group](#)” (page 51)
- “[Creating a Logical Volume](#)” (page 52)
- “[Extending a Logical Volume](#)” (page 53)
- “[Reducing a Logical Volume](#)” (page 55)
- “[Adding a Mirror to a Logical Volume](#)” (page 55)
- “[Removing a Mirror from a Logical Volume](#)” (page 56)
- “[Renaming a Logical Volume](#)” (page 56)
- “[Removing a Logical Volume](#)” (page 57)
- “[Exporting a Volume Group](#)” (page 57)
- “[Importing a Volume Group](#)” (page 58)
- “[Modifying Volume Group Parameters](#)” (page 58)
- “[Quiescing and Resuming a Volume Group](#)” (page 65)
- “[Renaming a Volume Group](#)” (page 66)
- “[Splitting a Volume Group](#)” (page 67)
- “[Removing a Volume Group](#)” (page 68)
- “[Backing Up a Mirrored Logical Volume](#)” (page 68)
- “[Backing Up and Restoring Volume Group Configuration](#)” (page 69)

Initializing a Disk for LVM Use

△ CAUTION: Initializing a disk using `pvcreate` results in the loss of any existing data currently on the disk.

NOTE: If your disk is already connected to the system, skip the first four steps of this procedure.

To initialize a disk for use as a physical volume, follow these steps:

1. Shut down and power off the system.
2. Connect the disk to the system and power supply. For detailed information and instructions on adding a particular type of disk, see your device documentation.
3. Power on the disk.
4. Boot the system.
5. Determine the disk's associated device file. To show the disks attached to the system and their device file names, enter the `ioscan` command with the `-f`, `-N`, and `-n` options. For example:
`# ioscan -f -n -N -C disk`
For more information, see [`ioscan\(1M\)`](#).

6. Initialize the disk as a physical volume using the `pvcreate` command. For example:

```
# pvcreate /dev/rdisk/disk3
```

Use the character device file for the disk.

If you are initializing a disk for use as a boot device, add the `-B` option to `pvcreate` to reserve an area on the disk for a LIF volume and boot utilities. If you are creating a boot disk on an HP Integrity server, make sure the device file specifies the HP-UX partition number (2). For example:

```
# pvcreate -B /dev/rdisk/disk3_p2
```

NOTE: Version 2.0 and 2.1 volume groups do not support bootable physical volumes. Do not use the `-B` option if the disk will be used in a Version 2.0 or 2.1 volume group.

After a disk is initialized, it is called a physical volume.

Creating a Volume Group

To create a volume group, use the `vgcreate` command. The options differ depending on whether you are creating a Version 1.0 volume group or a Version 2.x volume group.

Creating the Volume Group Device File

As of the March 2008 release of HP-UX 11i Version 3, the `vgcreate` command automatically creates the device file `/dev/vgname/group` to manage the volume group, regardless of the volume group version. If you are using an HP-UX release before March 2008, or if you want to specify the minor number of the group file, you must create `/dev/vgname/group` before running the `vgcreate` command.

If the group file does not exist and `vgcreate` can not create it, `vgcreate` displays the following message:

```
vgcreate: "/dev/vgname/group": not a character device.
```

To create the volume group device file, follow these steps:

1. Create a directory for the volume group. For example:

```
# mkdir /dev/vgname
```

By convention, `vgname` is `vgnn`, where `nn` is a unique number across all volume groups. However, you can choose any unique name up to 255 characters.

2. Create a device file named `group` in the volume group directory with the `mknod` command. For example:

```
# mknod /dev/vgname/group c major 0xminor
```

The `c` following the device file name specifies that `group` is a character device file.

`major` is the major number for the `group` device file. For a Version 1.0 volume group, it is 64. For a Version 2.x volume group, it is 128.

`minor` is the minor number for the `group` file in hexadecimal. For a Version 1.0 volume group, `minor` has the form `0xnn0000`, where `nn` is a unique number across all Version 1.0 volume groups. For a Version 2.x volume group, `minor` has the form `0xnnn000`, where `nnn` is a unique number across all Version 2.x volume groups.

For more information on `mknod`, see [mknod\(1M\)](#); for more information on major numbers and minor numbers, see ["Device Number Format" \(page 15\)](#).

Creating a Version 1.0 Volume Group

To create a Version 1.0 volume group, use the `vgcreate` command, specifying each physical volume to be included. For example:

```
# vgcreate /dev/vgname /dev/disk/disk3
```

Use the block device file to include each disk in your volume group. You can assign all the physical volumes to the volume group with one command, or create the volume group with a single physical volume. No physical volume can already be part of an existing volume group.

You can set volume group attributes using the following options:

- V 1.0 Version 1.0 volume group (default)
- s *pe_size* Size of a physical extent in MB (default 4)
- e *max_pe* Maximum number of physical extents per physical volume (default 1016)
- l *max_lv* Maximum number of logical volumes (default 255)
- p *max_pv* Maximum number of physical volumes (default 255)

The size of a physical volume is limited by *pe_size* times *max_pe*. If you plan to assign a disk larger than approximately 4 GB (1016 * 4 MB) to this volume group, use a larger value of *pe_size* or *max_pe*.

The size of the LVM metadata on each disk depends on *max_lv*, *max_pv*, and *max_pe*. If the vgcreate options would cause the metadata to exceed its available space, vgcreate does not create the volume group. You must select new values of *max_lv*, *max_pv*, and *max_pe*. For example, if you plan to use disks larger than 100 GB, consider reducing *max_pv*. For recommendations on choosing an optimal extent size, see [Appendix C \(page 155\)](#).

Creating a Version 2.x Volume Group

For Version 2.x volume groups, the vgcreate command does not require maximum values for the number of physical volumes (-p), number of logical volumes (-l), or extents per physical volume (-e). Instead you must specify only the extent size (-s) and the maximum size to which the volume group can grow (-S). For example:

```
# vgcreate -V 2.0 -s pe_size -S vg_size /dev/disk1 /dev/disk2
```

Use the block device file to include each disk in your volume group. You can assign all the physical volumes to the volume group with one command, or create the volume group with a single physical volume. No physical volume can already be part of an existing volume group.

You must use the following options:

- V 2.x Version 2.x (such as 2.0, 2.1, or 2.2) volume group
- s *pe_size* Size of a physical extent in MB
- S *vg_size* Maximum future size of the volume group

The size of a volume group is the sum of the user data space on all physical volumes assigned to the volume group. *vg_size* is not the size of the volume group at creation; it is the size to which the volume group can grow in the future. This value can be specified in megabytes, gigabytes, terabytes, or petabytes, by adding the character m, g, t, or p, respectively. For example, to specify a maximum size of two terabytes, use -S 2t.

In a Version 2.x volume group, the number of physical extents in a volume group has an architectural limit, so your choice of physical extent size affects the maximum size of the volume group. To display the maximum volume group size for a given physical extent size, use the -E option to vgcreate with the -s option. For example:

```
# vgcreate -V 2.0 -E -s 256
Max_VG_size=2p:extent_size=256m
```

Conversely, to display the minimum physical extent size for a given volume group size, use the -E option to vgcreate with -S. For example:

```
# vgcreate -V 2.0 -E -S 2t
Max_VG_size=2t:extent_size=1m
```

For volume group Version 2.2 or higher, a new vgcreate-U option is introduced to configure the size of the unshare unit for snapshots in the volume group. Once the unshare unit size is

specified, it cannot be changed. The default unshare unit size is 1024 KB if the **-U** option is not used. Below is an example of **vgcreate** with the **-U** option:

```
# vgcreate -V 2.2 -S 4t -s 8 -U 2048 /dev/vg01 /dev/disk/disk20
```

Migrating a Volume Group to a Different Version: **vgversion**

Beginning with the HP-UX 11i v3 March 2009 Update, LVM offers the new command **vgversion**, which enables you to migrate the current version of an existing volume group to any other version, except for migrating to Version 1.0. Therefore, the following migrations are supported:

Table 7 Supported Volume Group Version Migrations

Original Version	Supported Target Version
1.0	2.0 or 2.1 (migration does not support boot, swap, dump PVs)
1.0	2.2 (migration does not support boot PVs)
2.x	2.y

NOTE: Because Version 2.0 and 2.1 volume groups do not support bootable physical volumes, and because at least one physical volume in the root volume group must be a bootable physical volume, the root volume group, usually named **vg00**, cannot be migrated to Version 2.0 or 2.1.

Version 2.2 (or higher) volume groups do support bootable PVs; however, the migration of Version 1.0 bootable PVs to Version 2.2 (or higher) is currently not supported for the HP-UX 11i v3 March 2010 Update.

For detailed information on the versions of volume groups and the benefits of migrating your volume groups to Version 2.x, see ["LVM Volume Group Versions" \(page 11\)](#).

For detailed information on the **vgversion** command, see the **vgversion(1M)** manpage.

Determining the Version of a Volume Group

To determine the current version of an existing volume group, use the **vgdisplay** command:

```
# vgdisplay vg01
--- Volume groups ---
VG Name          /dev/vg01
VG Write Access  read/write
VG Status        available
Max LV          255
Cur LV          4
Open LV          4
Max PV          16
Cur PV          1
Act PV          1
Max PE per PV   1279
VGDA            2
PE Size (Mbytes) 4
Total PE        1279
Alloc PE        100
Free PE         1179
Total PVG       0
Total Spare PVs 0
Total Spare PVs in use 0
VG Version      1.0
VG Max Size     81856m
VG Max Extents  20464
```

The **VG Version** field indicates the target volume group is Version 1.0

Command Syntax

The vgversion syntax is:

```
vgversion [-r] [-v] -V [-U unshare_unit] vg_version_new vg_name
```

where

-r	is Review mode. This allows you to review the operation before performing the actual volume group version migration.
-v	is Verbose mode.
-U <i>unshare_unit</i>	sets the unit at which data will be unshared between a logical volume and its snapshots, in the new volume group. This is only applicable for migration to volume group Version 2.2 or higher. For more information about snapshots see " Creating and Administering Snapshot Logical Volumes " (page 103).
-V <i>vg_version_new</i>	is the new version to which you are migrating (2..x).
<i>vg_name</i>	is the target volume group (for example, vg01)



TIP:

- For a successful migration, the target volume group must meet certain conditions at execution time; for example, the target volume group must be inactive. However, these conditions do *not* need to be met when you are using Review mode (the -r option). For more information on using the -r option, see "[Review Mode: vgversion -r](#)" (page 47).
- Like other commands that run at the HP-UX shell level, the exit value of vgversion for a failed migration or failed migration review will be non-zero:

```
# echo $?
```

```
1
```

Review Mode: vgversion -r

vgversion allows you to review (-r) its stdout and stderr messages to determine whether an attempted migration will succeed or fail without performing the actual migration.

Migration Success

In Review mode, when the migration will succeed, you should get messages that conclude with text similar to the following:

```
...
Volume Group version can be successfully changed to ...
Review complete. Volume group not modified.
```

Migration Failure

A migration of a volume group can fail for various reasons, including:

- the target volume group must not be active
- the target volume group must be cluster-unaware
- no physical volumes of the target volume group should be used as the cluster lock disk
- the physical volumes of the volume group must have enough space for the new metadata
- the existing volume group has any features that are not supported in the target volume group version

(For a list of all conditions, see the latest *vgversion(1M)* manpage.)

When the migration will fail, the vgversion -r output messages should indicate the problem and possible solutions. For example, if you are migrating a volume group from Version 1.0 to 2.1,

Version 2.1 requires more metadata. Thus, it is possible that there is not enough space in the LUN for the increase in metadata. In this example, vgversion -r should display the following:

```
#vgversion -V 2.1 -r -v vg01
Performing "vgchange -a r -l -p -s vg01" to collect data
Activated volume group
Volume group "vg01" has been successfully activated.
The space required for Volume Group version 2.1 metadata on Physical Volume
/dev/disk/disk12 is 8448 KB, but available free space is 1024 KB.
0 free user extents from the end of the Physical Volume /dev/disk/disk12
will be utilized to accommodate the Volume Group version 2.1 metadata.
vgversion: More space for metadata is required. Take one of the following actions:
1. Free 8 extents more from the end of the Physical Volume
2. Increase the disk size by 7424 KB to avoid using the free user extents
3. Increase the disk size by 7424 KB and use the free user extents
```

Warning: Volume Group version 2.1 does not support bad block relocation. The bad block relocation policy of all logical volumes will be set to NONE.

```
Deactivating Volume Group "vg01"
Volume group "vg01" has been successfully deactivated.
Review complete. Volume group not modified
```

Using one of the solutions displayed or using a solution of your own, you must resolve the problem in order for the migration to be successful. In this case, we could increase the LUN size to allow for the additional metadata. After you believe you have resolved the problem(s) displayed, you can try vgversion -r again to determine whether the migration will be successful. Once the vgversion -r output shows a successful migration, you can execute the actual migration using vgversion without the -r option.



TIP: Use the -v (Verbose) option along with the -r (Review mode) option to display full details of the problem. In the above example, if you omit the -v option, the above output would not include the following details:

The space required for Volume Group version 2.1 metadata on Physical Volume /dev/disk/disk12 is 8448 KB, but available free space is 1024 KB.

Example of Migrating a Volume Group to a Different Version

In this example:

- The target volume is vg01.
- We will be migrating it from version 1.0 to 2.1.
- Assume that the target volume group is currently cluster-unaware but activated.

The process for performing a migration are:

1. Verify the current version of the target volume group:

```
# vgdisplay vg01 | grep -i version
VG Version          1.0
```

This shows the current version is 1.0.

2. Review the migration:

```
# vgversion -r -v -V 2.1 vg01
The space required for Volume Group version 2.1 metadata on
Physical Volume /dev/disk/disk12 is 8448 KB, but available free space
is 1024 KB.
8 free user extents from the end of the Physical Volume
/dev/disk/disk12 will be utilized to accommodate the Volume Group
version 2.1 metadata.
Warning: Volume Group version 2.1 does not support bad block
```

relocation. The bad block relocation policy of all logical volumes will be set to NONE.

Volume Group version can be successfully changed to 2.1
Review complete. Volume group not modified

3. After messages from the review indicate a successful migration, you can begin the actual migration:
 - a. *Unlike in Review mode, the target volume group must meet certain conditions at execution time, including being de-activated. Since the target volume group is active, we must de-activate it before performing the actual migration:*

```
# vgchange -a n vg01
```
 - b. Now, you can perform the actual migration:

```
# vgversion -v -V 2.1 vg01
Performing "vgchange -a y -l -p -s vg01" to collect data
Activated volume group
Volume group "vg01" has been successfully activated.
The space required for Volume Group version 2.1 metadata on
Physical Volume /dev/disk/disk12 is 8448 KB, but available free
space is 1024 KB.
8 free user extents from the end of the Physical Volume /dev/disk/disk12
will be utilized to accommodate the Volume Group version 2.1 metadata.
Warning: Volume Group version 2.1 does not support bad block
relocation. The bad block relocation policy of all logical volumes
will be set to NONE.
```

```
Old Volume Group configuration for "/dev/vg01" has been saved in
"/etc/lvmconf/vgversion_vg01/vg01_1.0.conf"
Deactivating Volume Group "vg01"
Volume group "vg01" has been successfully deactivated.
```

```
New Volume Group configuration for "/dev/vg01" has been saved in
"/etc/lvmconf/vgversion_vg01/vg01_2.1.conf"
Removing the Volume Group /dev/vg01 from /etc/lvmtab
```

```
Applying the configuration to all Physical Volumes from
"/etc/lvmconf/vgversion_vg01/vg01_2.1.conf"
Volume Group configuration has been restored to /dev/rdisk/disk12
Creating the Volume Group of version 2.1 with minor number 0x5000.
Adding the Volume Group /dev/vg01 to /etc/lvmtab_p
Original Volume Group Version was 1.0
New Volume Group Version is 2.1
Volume Group version has been successfully changed to 2.1"
```

The output shows that the migration was successful.

4. To verify the migration:
 - a. You can use vgdisplay to verify the migration. vgdisplay requires its target volume group to be activated. Since the target volume group could not be active during the migration, we must activate it:

```
# vgchange -a y vg01
```
 - b. Once activated, we can use vgdisplay to verify the migration:

```
# vgdisplay vg01 | grep -i version
VG Version
          2.1
```

The version shows the correct final version of 2.1.

Migration Recovery

When running vgversion, recovery configuration files and a recovery script are created so that you can restore the target volume group to its original version in case there are problems during the actual migration.

CAUTION: The recovery script should be run only in cases where the migration unexpectedly fails, such as an interruption during migration execution.

- The recovery script should not be used to “undo” a successful migration. For a successful vgversion migration, you should use only a subsequent vgversion execution (and not the recovery script) to reach the newly desired volume group version. Note that a migration to 1.0 is not supported, so no return path is available once a migration from Version 1.0 to Version 2.x is successful.
- Further, when needed, the recovery script should be run only immediately after the attempted migration fails. Because the recovery script contains only configuration information corresponding to the time of migration, any volume group configuration changes subsequent to the migration would not be included in the recovery script, and therefore will be lost when the recovery script is run.

File Location

The backup files, including the recovery script, exist in

`/etc/lvmconf/vgversion_vg_name`

where

`vg_name` is the name of the target volume group at the time of the actual conversion.

Command Syntax

The syntax for the recovery script is:

`vgversion_vg_name_restore vg_name_vg_version_old.conf`

where

`vg_name` is the name of the target volume group

`vg_version_old` is the original version number of the volume group (1.0 or 2.x)

NOTE: The target volume group should be de-activated or all physical volumes of the target volume group should be detached prior to running the recovery script.

Example Recovery

Assume there was an unexpected interruption, such as a power failure, during the above migration of vg01 and that we need to return the volume group to its pre-migration state, where the volume group was of version 1.0:

1. If needed, de-activate the target volume group:

```
# vgchange -a n vg01
```

2. cd to the directory where the configuration backup files and recovery script exist. Using the above file location naming convention:

```
# cd /etc/lvmconf/vgversion_vg01
```

3. Run the recovery script. Using the above syntax:

```
# vgversion_vg01_restore vg01_1.0.conf
```

4. Once the recovery has successfully completed, you can now re-attempt the migration using vgversion:

```
# vgversion -v -V 2.1 vg01
```

For more information on recovery, see the `vgversion(1M)` manpage.

NOTE: Once the recovery is complete using the restore script, an immediate vgversion operation in review mode will fail. You need to de-active the volume group and activate it again before running vgversion in review mode. This reset of volume group is not needed for a vgversion operation not in review mode.

Adding a Disk to a Volume Group

Often, as new disks are added to a system, they must be added to an existing volume group rather than creating a whole new volume group. If new disks are being added for user data, such as file systems or databases, do not add them to the root volume group. Instead, leave the root volume group as only the disks containing the root file system and system file systems such as /usr, /tmp, and so on.

To add a disk to a volume group, follow these steps:

1. Initialize the disk as a physical volume by using the pvcreate command, as described in “[Initializing a Disk for LVM Use](#)” (page 43).
2. Add the physical volume to the volume group using the vgextend command and the block device file for the disk. For example:

```
# vgextend /dev/vgname /dev/disk/disk3
```

Removing a Disk from a Volume Group

To remove a disk from a volume group, follow these steps:

1. Make sure the disk has no assigned physical extents, using the pvdisplay command. For example:

```
# pvdisplay /dev/disk/disk3
-- Physical volumes --
PV Name          /dev/disk/disk3
VG Name          /dev/vg00
PV Status        available
Allocatable      yes
VGDA             2
Cur LV           9
PE Size (Mbytes) 4
Total PE         1023
Free PE          494
Allocated PE     529
Stale PE         0
IO Timeout (Seconds) default
Autoswitch       On
Proactive Polling On

-- Distribution of physical volume --
LV Name          LE of LV  PE for LV
/dev/vg00/lvol1   25      25
/dev/vg00/lvol2   25      25
/dev/vg00/lvol3   50      50

--- Physical extents ---
PE   Status    LV          LE
0000 current   /dev/vg00/lvol1 0000
0001 current   /dev/vg00/lvol1 0001
0002 current   /dev/vg00/lvol1 0002
1021 free      0000
1022 free      0000
```

Check that the number of free physical extents (Free PE) matches the total number of physical extents (Total PE). If they are not the same, do one of the following tasks:

- Move the physical extents onto another physical volume in the volume group. See “[Moving Data to a Different Physical Volume](#)” (page 73).
 - Remove the logical volumes from the disk, as described in “[Removing a Logical Volume](#)” (page 57). The logical volumes with physical extents on the disk are shown at the end of the `pvdisplay` listing.
2. After the disk no longer holds any physical extents, use the `vgreduce` command to remove it from the volume group. For example:

```
# vgreduce /dev/vg01 /dev/disk/disk3
```

-
- ① **IMPORTANT:** If you are using LVM pvlinks, as described in “[Increasing Hardware Path Redundancy Through Multipathing](#)” (page 28), you must run the `vgreduce` command for each link to the disk.
-

Creating a Logical Volume

To create a logical volume, follow these steps:

1. Decide how much disk space the logical volume needs.

For example, you can add 200 MB of device swap space, or you might have a new project that you expect to grow to 10 GB.

2. Find a volume group that has adequate free space.

To determine if there is sufficient disk space available for the logical volume within a volume group, use the `vgdisplay` command to calculate this information. `vgdisplay` outputs data on one or more volume groups, including the physical extent size (under PE Size (MBytes)) and the number of available physical extents (under Free PE). By multiplying these two figures, you get the number of megabytes available within the volume group. For more information, see `vgdisplay(1M)`.

3. Create the logical volume using `lvcreate`. For example:

```
# lvcreate -L size_in_MB /dev/vg01
```

This command creates the logical volume `/dev/vg01/lvol1` with LVM automatically assigning the `n` in `lvoln`.

When LVM creates the logical volume, it creates block and character device files for that logical volume and places them in the `/dev/vg01` directory.

Creating a Striped Logical Volume

To create a striped logical volume, use `lvcreate` with the `-i` and the `-I` options to specify the number of disks and stripe width, respectively. For example, suppose you want to stripe across three disks with a stripe size of 32K. Your logical volume size is 240 MB. To create the striped logical volume, enter the following command:

```
# lvcreate -i 3 -I 32 -L 240 -n lvol1 /dev/vg01
```

The `lvcreate` command automatically rounds up the size of the logical volume to a multiple of the number of disks times the extent size. For example, if you have three disks you want to stripe across and the extent size is 4 MB, even though you indicate a logical volume size of 200 MB (`-L 200`), `lvcreate` creates a 204 MB logical volume because 200 is not a multiple of 12.

NOTE: When you stripe across multiple disks, the striped volume size cannot exceed the capacity of the smallest disk multiplied by the number of disks used in the striping.

Creating a Mirrored Logical Volume

To create a mirrored logical volume, use `lvcreate` with the `-m` option to select the number of mirror copies. To control how the mirror copies are managed, choose from the following options:

Strict, Nonstrict, or PVG-strict Extent Allocation

<code>-s y</code>	Strict allocation (default)
<code>-s n</code>	Nonstrict allocation
<code>-s g</code>	PVG-strict allocation

Contiguous or Noncontiguous Extent Allocation

<code>-C y</code>	Contiguous allocation
<code>-C n</code>	Noncontiguous allocation (default)

Mirror Scheduling Policy

<code>-d p</code>	Parallel scheduling (default)
<code>-d s</code>	Sequential scheduling

Mirror Consistency Policy

<code>-M y</code>	MWC enable (default, optimal mirror resynchronization during crash recovery)
<code>-M n -c y</code>	MCR enable (full mirror resynchronization during crash recovery)
<code>-M n -c n</code>	MCR disable (no mirror resynchronization during crash recovery)

For example, to create a 240 MB mirrored logical volume with one mirror copy, nonstrict allocation, parallel scheduling, and no mirror resynchronization, enter the following command:

```
# lvcreate -m 1 -s n -d p -M n -c n -L 240 -n lvol1 /dev/vg01
```

TIP: To change the characteristics of an existing mirrored logical volume, use the `lvchange` command. It supports the `-C`, `-c`, `-d`, `-M`, and `-s` options. For more information, see [lvchange\(1M\)](#).

Extending a Logical Volume

NOTE: Adding space to a logical volume does not automatically assign that space to the entity using that logical volume. For example, if you want to add space to a file system contained in a logical volume, you must run `extendfs` after extending the logical volume. See [“Administering File System Logical Volumes” \(page 97\)](#) and [“Administering Swap Logical Volumes” \(page 101\)](#) for more information.

1. Decide how much more disk space the logical volume needs.

For example, you can add 200 MB of swap space, or an existing project might need an additional 1 GB.

2. Find out if any space is available using the `vgdisplay` command. For example:

```
# vgdisplay vg00
--- Volume groups ---
VG Name                  /dev/vg00
VG Write Access          read/write
VG Status                available
Max LV                  255
Cur LV                  8
Open LV                 8
Max PV                  16
Cur PV                  1
```

Act PV	1
Max PE per PV	2000
VGDA	2
PE Size (Mbytes)	4
Total PE	249
Alloc PE	170
Free PE	79
Total PVG	0
Total Spare PVs	0
Total Spare PVs in use	0
VG Version	1.0
VG Max Size	1082g
VG Max Extents	69248

The Free PE entry indicates the number of 4 MB extents available, in this case, 79 (316 MB).

- Extend the logical volume. For example:

```
# lvextend -L 332 /dev/vg00/lvol7
```

This increases the size of this volume to 332 MB.

NOTE: On the HP-UX 11i v3 March 2010 Update, the size of a logical volume cannot be extended if it has snapshots associated with it. With the HP-UX 11i v3 September 2010 Update, this limitation is removed; and logical volumes with snapshots can be extended.. For information about snapshot logical volumes, see “[Creating and Administering Snapshot Logical Volumes](#)” ([page 103](#)).

Extending a Logical Volume to a Specific Disk

For performance reasons, you can force a logical volume to span disks. For example, if you want to create a 30 GB logical volume and put 10 GB on the first disk, another 10 GB on the second disk, and 10 GB on the third disk, then assuming that the extent size is 4 MB, the logical volume requires a total of 7680 extents. To extend the logical volume, follow these steps:

- After making the disks physical volumes and creating your volume group, create a logical volume named lvol1 of size 0. For example:
- ```
lvcreate -n lvol1 /dev/vg01
```
- Allocate a third of the extents to the logical volume on the first physical volume. For example:
- ```
# lvextend -l 2560 /dev/vg01/lvol1 /dev/disk/disk7
```
- Increase the total number of physical extents allocated to the logical volume for the remaining physical volumes by 2560. In each case, the additional 2560 extents are allocated to the disk specified. For example:

```
# lvextend -l 5120 /dev/vg01/lvol1 /dev/disk/disk8
# lvextend -l 7680 /dev/vg01/lvol1 /dev/disk/disk9
```

When you use the -l option with lvextend, you specify space in logical extents.

For another example, you have two disks in a volume group, both identical models. You currently have a 24 GB logical volume that resides on only one of the disks. You want to extend the logical volume size to 40 GB and ensure that the 16 GB increase is allocated to the other disk.

Extend the logical volume to a specific disk as follows:

```
# lvextend -L 40960 /dev/vg01/lvol2 /dev/disk/disk3
```

Here, when you use the -L option (uppercase), you specify space in megabytes, not logical extents.

For complete information on command options, see [lvextend\(1M\)](#).

Reducing a Logical Volume

⚠ **CAUTION:** Before you reduce a logical volume, you must notify the users of that logical volume. For example, before reducing a logical volume that contains a file system, *back up the file system*. Even if the file system currently occupies less space than the new (reduced) size of the logical volume, you will almost certainly lose data when you reduce the logical volume. See “[Administering File System Logical Volumes](#)” (page 97) and “[Administering Swap Logical Volumes](#)” (page 101) for the appropriate procedures for file systems and swap devices.

To reduce a logical volume, follow these steps:

1. To find out what applications are using the logical volume, use the `fuser` command. For example:
`# fuser -cu /dev/vg01/lvol5`
If the logical volume is in use, ensure the underlying applications can handle the size reduction. You might have to stop the applications.
2. Decide on the new size of the logical volume.
For example, if the logical volume is mounted to a file system, the new size must be greater than the space the data in the file system currently occupies. The `bdf` command shows the size of all mounted volumes. The first column shows the space allocated to the volume; the second shows how much is actually being used. The new size of the logical volume must be larger than the size shown in the second column of the `bdf` output.
3. Reduce the size of the logical volume as follows:
`# lvreduce -L 500 /dev/vg01/lvol5`
This command reduces the logical volume `/dev/vg01/lvol5` to 500 MB.

NOTE: On the HP-UX 11i v3 March 2010 Update, the size of a logical volume cannot be reduced if it has any type of snapshots associated with it. With the HP-UX 11i v3 September 2010 Update, only logical volumes with associated space efficient snapshots cannot be reduced. For information about snapshots, see “[Creating and Administering Snapshot Logical Volumes](#)” (page 103).

Adding a Mirror to a Logical Volume

NOTE: Mirroring requires the optional product HP MirrorDisk/UX.

💡 **TIP:** This task is easier to perform with HP SMH. HP SMH confirms that enough disk space is available for the mirror copy and that the available space meets any allocation policy.

To add a mirror to a logical volume, follow these steps:

1. Decide how many mirror copies to create.
For this example, you will create one mirror; that is, you will keep two copies of the data online, the original and one mirror copy.
2. Make sure that there is enough free space in the volume group that contains the logical volume you want to mirror.

The volume group needs at least as much free space as the logical volume you want to mirror currently has allocated to it—that is, you will double the amount of physical space this volume requires.

To use strict mirroring (which HP recommends because it keeps the mirror copy on a separate disk) this free space must be on a disk or disks not currently used by the volume you want to mirror.

3. Use the `lvextend` command with the `-m` option to add the number of additional copies you want. For example:

```
# lvextend -m 1 /dev/vg00/lvol1
```

This adds a single mirror copy of the given logical volume.

To force the mirror copy onto a specific physical volume, add it at end of the command line. For example:

```
# lvextend -m 1 /dev/vg00/lvol1 /dev/disk/disk4
```

NOTE: On the HP-UX 11i v3 March 2010 Update, the number of mirror copies of a logical volume cannot be changed if the logical volume has snapshots associated with it. With the HP-UX 11i v3 September 2010 Update, this limitation has been removed; and logical volumes with associated snapshots can have their mirror count changed. For information about snapshots see "[Creating and Administering Snapshot Logical Volumes](#)" (page 103).

Removing a Mirror from a Logical Volume

To remove a mirror copy, use the `lvreduce` command, specifying the number of mirror copies you want to leave. For example, to remove all mirrors of a logical volume, enter the following command:

```
# lvreduce -m 0 /dev/vg00/lvol1
```

This reduces the number of mirror copies to 0, so only the original copy is left.

To remove the mirror copy from a specific disk, use `lvreduce` and specify the disk from which to remove the mirror copy. For example:

```
# lvreduce -m 0 /dev/vg00/lvol1 /dev/disk/disk4
```

NOTE: On the HP-UX 11i v3 March 2010 Update, the number of mirror copies of a logical volume cannot be changed if the logical volume has snapshots associated with it. With the HP-UX 11i v3 September 2010 Update, this limitation has been removed; and logical volumes with associated snapshots can have their mirror count changed. For information about snapshots see "[Creating and Administering Snapshot Logical Volumes](#)" (page 103).

Renaming a Logical Volume

To change the name of a logical volume, follow these steps:

1. Make sure that the logical volume has two existing device files, a block device file and a character or raw device file. They must have the same name, except that the character device file name has a leading `r`. For example, to rename a logical volume in volume group `vg00` from `lvol1` to `database`, list the contents of the `/dev/vg00` directory. For example:

```
# cd /dev/vg00
# ls -l
total 0
crw-r---- 1 root      sys  64 0x000000 Nov 16 02:49 group
brw-r---- 1 root      sys  64 0x000001 Nov 16 02:49 lvol1
brw-r---- 1 root      sys  64 0x000002 Nov 16 02:49 lvol2
brw-r---- 1 root      sys  64 0x000003 Nov 16 02:49 lvol3
brw-r---- 1 root      sys  64 0x000004 Nov 16 02:49 lvol4
crw-r---- 1 root      sys  64 0x000001 Nov 16 02:49 rlvvol1
crw-r---- 1 root      sys  64 0x000002 Nov 16 02:49 rlvvol2
crw-r---- 1 root      sys  64 0x000003 Nov 16 02:49 rlvvol3
crw-r---- 1 root      sys  64 0x000004 Nov 16 02:49 rlvvol4
```

2. Use the `mv` command to rename both files. For example:

```
# mv /dev/vg00/lvol1 /dev/vg00/database
# mv /dev/vg00/rpvvol1 /dev/vg00/rdatabase
```

3. Update all references to the old name in any other files on the system. These include /etc/fstab for mounted file systems or swap devices and existing mapfiles from a vgexport command.

Removing a Logical Volume

⚠ **CAUTION:** Removing a logical volume makes its contents unavailable and likely to be overwritten. In particular, any file system contained in the logical volume is destroyed.

To remove a logical volume, follow these steps:

1. Make sure that the logical volume is not in use either as a file system or as raw disk space for an application. Use the fuser command as follows:

```
# fuser -cu /dev/vg01/lvol5
```

If the logical volume is in use, confirm that the underlying applications no longer need it. You might need to stop the applications.
 2. Use the lvremove command to remove the logical volume. For example:

```
# lvremove /dev/vg01/lvol5
```

You can now use this space to extend an existing logical volume or build a new logical volume.
 3. For volume group Version 2.2 and higher, a snapshot and all its predecessors can be removed using a single lvremove command, with the new -F option. See *lvremove(1M)* for more information.
-

NOTE: A logical volume with associated snapshots cannot be removed. First, all of its snapshots have to be deleted, then the original logical volume can be deleted.

Exporting a Volume Group

Exporting a volume group removes all data concerning the volume group from the system, while leaving the data on the disks intact. The disks of an exported volume can be physically moved or connected to another system, and the volume group can be imported there.

Exporting a volume group removes information about the volume group and its associated physical volumes from /etc/lvmtab and /etc/lvmtab_p, and removes the volume group's directory with device files in the /dev directory.

1. Make sure that none of the logical volumes in the volume group are in use. You might need to stop applications using any logical volumes in the volume group, and unmount file systems contained in the volume group.

Use the fuser command on each logical volume. For example:

```
# fuser -cu /dev/vgnn/lvoln
```

2. Deactivate the volume group. For example:

```
# vgchange -a n vgnn
```

3. Use the vgexport command to export the volume group. For example:

```
# vgexport -v -m /tmp/vgnn.map vgnn
```

If you are planning to move the volume group to another system, use the -m option to vgexport to create a mapfile. This ASCII file contains the logical volume names because they are not stored on the disks. You must create a mapfile if you do not use the default names /dev/vgnn/lvoln for the logical volumes in the volume group.

If there are several disks in the volume group, use the -s option with vgexport; this option adds the volume group identifier (VGID) to the mapfile. When the volume group is imported, you can avoid specifying all the disks by name. See “[Importing a Volume Group \(page 58\)](#)”.

When `vgexport` completes, all information about the volume group has been removed from the system. The disks can now be moved to a different system, and the volume group can be imported there.

Importing a Volume Group

To import a volume group, follow these steps:

1. Connect the disks to the system.
2. If you are using an HP-UX release before March 2008, or if you want to specify the minor number of the volume group device file, create it using the procedure in “[Creating the Volume Group Device File](#)” (page 44).
3. Use the `vgimport` command to import the volume group:

```
# vgimport -v -N -m /tmp/vgnn.map /dev/vgnn list_of_disks
```

If there are several disks in the volume group and the VGID was saved in the mapfile (that is, the `vgexport` command was run with the `-s` and `-m` options), you can avoid specifying all of them in the `vgimport` command line by using the `-s` option. This causes `vgimport` to scan all the disks on the system. Any physical volumes with a VGID matching the one in the mapfile are included automatically into the volume group.

4. Activate the volume group as follows:

```
# vgchange -a y vgnn
```

NOTE: If the volume group contains any multipathed disks, HP recommends using HP-UX's native multipathing that is a superset of LVM's alternate links. See “[Increasing Hardware Path Redundancy Through Multipathing](#)” (page 28) for more information.

If you want to use LVM's alternate link features, importing the volume group has several implications:

- You must omit the `-N` option to the `vgimport` command.
 - The `vgimport` sets the first link found as the primary link for all physical volumes. If the links are not in the desired order after the import, use `vgreduce` and `vgextend` on the primary link for each physical volume for which you want to change the primary.
 - The tunable `maxfiles` must be more than double the number of disks free.
-

Modifying Volume Group Parameters

When you create a volume group, you set certain characteristics of the volume group. Using the `vgmodify` command you can adjust some of the volume group parameters without removing and re-creating the volume group or having to move your data. The `vgmodify` command enables you to alter different parameters for Version 1.0 volume groups and Version 2.x volume groups.

For version 1.0 volume groups, you can use `vgmodify` to do the following:

- Detect and handle physical volume size changes.
- Modify the maximum number of physical extents that can be allocated per physical volume.
- Modify the maximum number of physical volumes that the volume group can contain.
- Modify the maximum number of logical volumes that the volume group can contain.
- Change a disk type from boot to non-boot or vice versa.
- Handle physical volume LUN expansion. See “[Modifying Physical Volume Characteristics](#)” (page 77) for more details
- Prepare a physical volume for a LUN contraction. See “[Modifying Physical Volume Characteristics](#)” (page 77) for more details.

For version 2.x volume groups, you can use `vgmodify` to do the following:

- Detect and handle physical volume size changes
- Modify the maximum volume group size.
- Handle physical volume LUN expansion. See “[Modifying Physical Volume Characteristics](#)” ([page 77](#)) for more details
- Prepare a physical volume for a LUN contraction. See “[Modifying Physical Volume Characteristics](#)” ([page 77](#)) for more details.

Beginning with the March 2009 Update, the `vgmodify` command can be run online (volume group activated and applications running) for Version 1.0 and 2.x volume groups. Additionally, for Version 1.0 volume groups, `vgmodify` can be run with the volume group offline (deactivated).

vgmodify for a Version 1.0 Volume Group

For Version 1.0 volume groups, use the following procedure to adjust the maximum number of PE per PV, the maximum number of PVs in the VG, and the maximum number of LVs in the VG:

1. Run `vgmodify` to collect information about the volume group.

Save the output from these three commands:

```
# vgmodify -o -r vgnn  
# vgmodify -v -t vgnn  
# vgmodify -v -n -t vgnn
```

The `-o` option attempts to optimize the values by making full use of the existing LVM metadata space. The `-t` option reports the optimized range of settings without renumbering physical extents; the `-n` option enables renumbering of physical extents.

2. Based on the information collected in the previous step, choose new values for the volume group parameters.
3. The new values may increase the size of the volume group reserved area (VGRA) on each physical volume. The VGRA resides in the LVM header, so increasing its size may require moving the first physical extent of any user data on physical volume. Use the `pvmove` command to move the first physical extent to another location.
4. Review the values by running `vgmodify` with the new settings and the `-r` option.
5. Deactivate the volume group, if needed.



TIP: Beginning with the March 2009 Update of HP-UX, the volume group can be active when running `vgmodify`.

6. Commit the new values by running `vgmodify` without the `-r` option.
7. If needed, activate the volume group. Run the `vgdisplay` command to verify the settings have been applied.

Example: vgmodify for a Version 1.0 Volume Group

As an example, you expect to add larger disks to the volume group `vg32`. You want to increase the maximum number of physical extents per physical volume (`max_pe`) and the maximum number of physical volumes (`max_pv`). Here are the steps involved:

1. Run `vgmodify` to collect information about the volume group.

Save the output from these three commands:

```
# vgmodify -o -r vg32  
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

```
New configuration requires "max_pes" are increased from 1016 to 6652
The current and new Volume Group parameters differ.
An update to the Volume Group IS required
```

```
New Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	6652
PE Size (Mbytes)	32
VGRA Size (Kbytes)	896

```
Review complete. Volume group not modified
```

```
# vgmodify -v -t vg32
```

```
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

```
VGRA space (Kbytes) on Physical Volumes with extents in use:
```

PV	current	-n
/dev/rdisk/disk6	896	32768
/dev/rdisk/disk5	896	32768
Summary	896	32768

```
Volume Group optimized settings (no PEs renumbered) :
```

max_pv(-p) max_pe(-e) Disk size (Mb)		
2	53756	1720193
3	35836	1146753
...		
213	296	9473
255	252	8065

```
# vgmodify -v -n -t vg32
```

```
Volume Group configuration for /dev/vg32 has been saved in
/etc/lvmconf/vg32.conf
```

```
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

```
VGRA space (Kbytes) on Physical Volumes with extents in use:
```

PV	current	-n
/dev/rdisk/disk6	896	32768
/dev/rdisk/disk5	896	32768
Summary	896	32768

```
Physical Extent zero is not free on all PVs. You will not achieve these
values until the first extent is made free (see pvmove(1M)) on all the
following disks:
```

```
/dev/rdisk/disk6
/dev/rdisk/disk5
```

```
Volume Group optimized settings (PEs renumbered lower) :
```

max_pv(-p) max_pe(-e) Disk size (Mb)		
61	65535	2097152
62	65532	2097056
...		
252	16048	513568
255	15868	507808

2. Based on the output of vgmodify -n -t, choose 255 for max_pv and 15868 for max_pe.
3. Since the new values require physical extent 0 to be free, use pvmove to move it to another location:

```
# pvmove /dev/disk/disk5:0 /dev/disk/disk5
```

```
Transferring logical extents of logical volume "/dev/vg32/lvol2"...
Physical volume "/dev/disk/disk5" has been successfully moved.
```

```
Volume Group configuration for /dev/vg32 has been saved in  
/etc/lvmconf/vg32.conf
```

```
# pvmove /dev/disk/disk6:0 /dev/disk/disk6  
Transferring logical extents of logical volume "/dev/vg32/lvol1"....  
Physical volume "/dev/disk/disk6" has been successfully moved.  
Volume Group configuration for /dev/vg32 has been saved in  
/etc/lvmconf/vg32.conf
```

4. Preview the changes by using the -r option to vgmodify:

```
# vgmodify -p 255 -e 15868 -r -n vg32  
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

```
The current and new Volume Group parameters differ.  
An update to the Volume Group IS required
```

```
New Volume Group settings:
```

Max LV	255
Max PV	255
Max PE per PV	15868
PE Size (Mbytes)	32
VGRA Size (Kbytes)	32640

```
Review complete. Volume group not modified
```

5. Deactivate the volume group:

```
# vgchange -a n vg32  
Volume group "vg32" has been successfully changed.
```

6. Commit the new values:

```
# vgmodify -p 255 -e 15868 -n vg32  
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

```
The current and new Volume Group parameters differ.  
An update to the Volume Group IS required
```

```
New Volume Group settings:
```

Max LV	255
Max PV	255
Max PE per PV	15868
PE Size (Mbytes)	32
VGRA Size (Kbytes)	32640

```
New Volume Group configuration for "vg32" has been saved in  
"/etc/lvmconf/vg32.conf"
```

```
Old Volume Group configuration for "vg32" has been saved in  
"/etc/lvmconf/vg32.conf.old"
```

```
Starting the modification by writing to all Physical Volumes  
Applying the configuration to all Physical Volumes from  
"/etc/lvmconf/vg32.conf"
```

```
Completed the modification process.
```

```
New Volume Group configuration for "vg32" has been saved in  
"/etc/lvmconf/vg32.conf.old"
```

```
Volume group "vg32" has been successfully changed.
```

7. Activate the volume group and verify the changes:

```
# vgchange -a y vg32  
Activated volume group  
Volume group "vg32" has been successfully changed.
```

```
# vgdisplay vg32
--- Volume groups ---
VG Name                  /dev/vg32
VG Write Access          read/write
VG Status                available
Max LV                  255
Cur LV                  0
Open LV                 0
Max PV                  255
Cur PV                  2
Act PV                  2
Max PE per PV           15868
VGDA                    4
PE Size (Mbytes)        32
Total PE                1084
Alloc PE                0
Free PE                 1084
Total PVG               0
Total Spare PVs         0
Total Spare PVs in use 0
VG Version              1.0
```

vgmodify for a Version 2.x Volume Group

If the maximum volume group size (chosen when the Version 2.x volume group was created by the `vgcreate` command) needs to be increased, then the `vgmodify -S` option can be used to increase the maximum volume group size.

For instance, the maximum volume group size might need to be increased to allow more extents to be added to the volume group through the `vgextend` command or through LUN expansion (see the `vgmodify -E` option described in “[Modifying Physical Volume Characteristics](#)” (page 77)).

The maximum volume group size can also be reduced through the `vgmodify -S` option to make more extents available on the physical volumes.

Increasing the maximum volume group size might require that additional space be added to every physical volume through LUN expansion. In this case, you can invoke the `-E` and `-S` options together to handle the LUN expansions and increase the maximum VG size in a single step.

LVM might also use free extents at the ends of the physical volumes to increase the maximum volume group size. If insufficient free extents are available at the end of some physical volumes, and those physical volume cannot be expanded further via LUN expansion, it might be necessary to free the extents at the ends of those physical volumes by using the `pvmove` command to relocate the portions of the logical volumes that are using those extents.

When using the `vgmodify -S` option to increase the maximum volume group size, note that the maximum volume group size is not increased until all physical volumes in the volume group are reconfigured.

When using the `vgmodify -S` option to decrease the maximum volume group size, note that the maximum volume group size is decreased prior to reconfiguring the physical volumes.

To increase the maximum size of a Version 2.x volume group, follow these steps:

1. Use the `vgmodify` review mode to verify that there is enough free space on all of the physical volumes to reconfigure them to the desired larger maximum VG size. The `vgmodify` command checks for free extents at the ends of the PVs that it can utilize, and also checks for free spaces resulting from prior LUN expansions. You must provide the `-E` option and the `-S` option for review mode to check for LUN expansions. It might be necessary to use the `pvmove` command to relocate the portions of logical volumes residing at the ends of the PVs.

```
# vgmodify -r -a -E -S 64t vg1
```

2. If review mode indicates that the maximum VG size can be increased, perform the actual reprovisioning reconfiguration. This operation reconfigures every PV in the VG to the new (larger) maximum VG size. This operation also automatically adds new extents to PVs where the new extents were previously not added because the volume group was at its maximum number of extents.

```
# vgmodify -a -E -S 64t vg1
```
3. Verify the result of the reprovisioning reconfiguration by running the `vgdisplay -v` command to check the new maximum VG size, the new maximum number of extents for the VG, and additional extents that might have been added to some of the PVs..
4. If any `vgextend` were previously unable to add PVs to the volume group, rerun those commands now. The `vgextend` commands must succeed in adding more PVs to the VG. Run `vgdisplay -v` to verify that more free extents are now available.
5. Expand logical volumes using `lvextend -l` as necessary to allow them to accommodate more user data. You can also create additional logical volumes using `lvcreate`. You can also add additional mirrors using `lvextend -m`.

To decrease the maximum Volume Group size for a Version 2.x volume group, follow these steps:

1. Use the review mode of `vgmodify` to verify that the volume group maximum size can be decreased to the desired lower value.

```
# vgmodify -r -a -S 32t vg1
```
2. If review mode indicates that the maximum VG size can be decreased, perform the actual reprovisioning reconfiguration. The `vgmodify` command reconfigures every PV in the VG to reduce the amount of space being used for LVM configuration data. The unused space is made available as new extents for user data

```
# vgmodify -a -S 32t vg1
```
3. Verify the result of the reprovisioning reconfiguration by running the `vgdisplay -v` command to check the new maximum VG size, the new maximum number of extents for the VG, the new total number of extents for the VG, and the new number of free extents for the VG.

Example: `vgmodify` for a Version 2.x Volume Group

Here is an example of online reprovisioning when a volume group has reached its maximum VG size limit. In this example, the `vgmodify -S` option is used to increase the maximum VG size for a 2.x volume group that has reached its maximum VG size limit. Here is the scenario:

The volume group `/dev/vg1` has two physical volumes, `/dev/disk/disk46` and `/dev/disk/disk47`.

The physical volume `/dev/disk/disk46` currently has 38398 extents and some extents at the end of the physical volume are free. The physical volume `/dev/disk/disk47` currently has 25602 extents. The volume group currently has 64000 extents ($38398 + 25602 = 64000$).

The maximum VG size as shown by `vgdisplay` is 500 GB. The maximum number of extents in the volume group is 64000.

The `diskinfo` command shows that both physical volumes are roughly 300 GB in size. However, since the maximum VG size is only 500 GB, roughly 100 GB is currently unavailable for use on `/dev/disk/disk47`. In addition, any attempt to add more physical volumes through the `vgextend` command will fail since the volume group is already at its maximum VG size limit.

To make the full 300 GB on `/dev/disk/disk47` available for use, and to allow additional physical volumes to be added to the volume group through the `vgextend` command, the `vgmodify -S` option can be used to increase the maximum VG size. In this example, it is increased from 500 GB to 8 TB to allow room for future growth. The steps involved are:

1. Run `vgdisplay`, which shows that the maximum VG size is 500 GB and the maximum number of extents in the volume group (determined from the maximum VG size) is 64000.

```
# vgdisplay -v vg1
--- Volume groups ---
VG Name                  /dev/vg1
...
VG Version               2.1
VG Max Size              500 GB
VG Max Entents           64000
...
```

- Run `diskinfo` on the physical volumes to display their size.

```
# diskinfo /dev/rdisk/disk46
# diskinfo /dev/rdisk/disk47
```

- Run online `vgmodify` in review mode to verify that all physical volumes can be reconfigured to the new (larger) maximum VG size of 8TB.

```
# vgmodify -r -a -E -S 8t vg1
Physical volume "/dev/disk/disk46" requires reconfiguration to be
provisioned to the requested maximum volume group size 8388608 MB.
Current number of extents: 38398
Number of extents after reconfiguration: 38396
...
Physical volume "/dev/disk/disk46" was not changed.

Physical volume "/dev/disk/disk4" requires reconfiguration to be
provisioned to the requested maximum volume group size 8388608 MB.
Current number of extents: 25602
Number of extents after reconfiguration: 25602
...
Physical volume "/dev/disk/disk47" was not changed.
```

In this example, all physical volumes in the volume group can be reconfigured, so the maximum Volume Group size of `/dev/vg1` can be increased from 500 GB to 8 TB.

Note that two free extents will be used for the reconfiguration on `/dev/disk/disk46`, reducing the extents on that physical volume from 38398 to 39396.

- Note that additional extents will be added to `/dev/disk/disk47` when the `vgmodify -S` option is run in change mode, as shown here.

```
# vgmodify -a -E -S 8t vg1
Reconfiguration of physical volume "/dev/disk/disk46" for the
requested maximum volume group size 8388608 MB succeeded.
Previous number of extents: 38398
Number of extents after reconfiguration: 38396
Physical volume "/dev/disk/disk46" was changed.
```

Volume Group configuration for `/dev/vg1` has been saved in
`/etc/lvmconf/vg1.conf`

```
Reconfiguration of physical volume "/dev/disk/disk47" for the
requested maximum volume group size 8388608 MB succeeded.
Previous number of extents: 25604
Number of extents after reconfiguration: 25604
Physical volume "/dev/disk/disk47" was changed.
```

The maximum volume group size for `" /dev/vg1 "` has been
increased from 512000 MB to 8388608 MB.

Volume Group configuration for `/dev/vg1` has been saved in
`/etc/lvmconf/vg1.conf`

As noted for review mode, two extents were taken from `/dev/disk/disk46` for the reconfiguration (a decrease from 38398 to 38396). This allowed two extents to be added to `/dev/disk/disk47` (an increase from 25602 to 25604) before the maximum volume group size was increased

5. Note that the number of extents for /dev/disk/disk47 is increased from 25604 to 38396 after the maximum VG size is increased, as shown by vgdisplay -v here. This command confirms that the maximum VG size was increased from 500 GB to 8 TB, and that the number of extents for /dev/disk/disk47 was increased from 25602 extents to 38396 extents, as shown by the vgdisplay command below.

```
# vgdisplay -v vg1
--- Volume groups ---
VG Name                               /dev/vg1
...
PE Size (Mbytes)                      8
Total PE                             76792
Alloc PE                            51196
Free PE                             25596
...
VG Version                           2.1
VG Max Size                          8t
VG Max Extents                      1048576

--- Logical volumes ---
LV Name                               /dev/vg1/lvol1
LV Status available/syncd
LV Size (Mbytes)                      204784
Current LE                           25598
Allocated PE                          51196
Used PV                             2

--- Physical volumes ---
PV Name                               /dev/disk/disk46
PV Status                            available
Total PE                            38396
Free PE                             12798
...
PV Name                               /dev/disk/disk47
PV Status                            available
Total PE                            38396
Free PE                             12798
```

Quiescing and Resuming a Volume Group

If you plan to use a disk management utility to create a backup image or “snapshot” of all the disks in a volume group, you must make sure that LVM is not writing to any of the disks when the snapshot is being taken; otherwise, some disks can contain partially written or inconsistent LVM metadata. To keep the volume group disk image in a consistent state, you must either deactivate the volume group or quiesce it.

Deactivating the volume group requires you to close all the logical volumes in the volume group, which can be disruptive. For example, you must unmount any file system using a logical volume in the volume group. However, temporarily quiescing the volume group enables you to keep the volume group activated and the logical volumes open during the snapshot operation, minimizing the impact to your system.

You can quiesce both read and write operations to the volume group, or just write operations. While a volume group is quiesced, the vgdisplay command reports the volume group access mode as quiesced. The indicated I/O operations queue until the volume group is resumed, and commands that modify the volume group configuration fail immediately.

NOTE: Individual physical volumes or logical volumes cannot be quiesced using this feature. To temporarily quiesce a physical volume to disable or replace it, see “[Disabling a Path to a Physical Volume](#)” (page 87). To quiesce a logical volume, quiesce or deactivate the volume group. To provide a stable image of a logical volume without deactivating the volume group, mirror the logical volume, then split off one of the mirrors, as described in “[Backing Up a Mirrored Logical Volume](#)” (page 68).

Quiescing a volume group is not persistent across reboots.

To quiesce a volume group, use the `vgchange` command with the `-Q` option as follows:

```
# vgchange -Q mode vgnn
```

The `mode` parameter can be either `rw`, which blocks both read and write operations, or `w`, which permits read operations but blocks write operations.

By default, the volume group remains quiesced until it is explicitly resumed. You can specify a maximum quiesce time in seconds using the `-t` option. If the quiesce time expires, the volume group is resumed automatically. For example, to quiesce volume group `vg08` for a maximum of ten minutes (600 seconds) but permitting read operations, enter the following command:

```
# vgchange -Q w -t 600 vg08
```

To resume a quiesced volume group, use the `vgchange` command with the `-R` option as follows:

```
# vgchange -R vgnn
```

Renaming a Volume Group

To change the name of a volume group, export it, then import it using the new name. For more detailed information on how to export and import a volume group, see “[Exporting a Volume Group](#)” (page 57) and “[Importing a Volume Group](#)” (page 58).

To rename the volume group `vg01` to `vgdb`, follow these steps:

1. Deactivate the volume group as follows:

```
# vgchange -a n vg01
```

2. If you want to retain the same minor number for the volume group, examine the volume group's `group` file as follows:

```
# ls -l /dev/vg01/group
crw-r--r-- 1 root sys 64 0x010000 Mar 28 2004 /dev/vg01/group
```

For this example, the volume group major number is 64, and the minor number is `0x010000`.

3. Export the volume group as follows:

```
# vgexport -m vg01.map vg01
```

4. If you are using an HP-UX release before March 2008, or if you want to specify the minor number of the volume group device file, create it for the volume group's new name, using the procedure in “[Creating the Volume Group Device File](#)” (page 44).

Since the `group` file in this example has a major number of 64 and a minor number of `0x010000`, enter the following commands:

```
# mkdir /dev/vgdb
# mknod /dev/vgdb/group c 64 0x010000
```

5. Import the volume group under its new name as follows:

```
# vgimport -m vg01.map /dev/vgdb
```

6. Back up the volume group configuration information as follows:

```
# vgcfgbackup /dev/vgdb
```

7. Activate the volume group as follows:

```
# vgchange -a y /dev/vgdb
```

8. Remove saved configuration information based on the old volume group name as follows:

```
# rm /etc/lvmconf/vg01.conf
```

Note: if your volume group is Version 2.x and does not have any bootable physical volumes, and if you have configured a new path for the configuration file using the `LVMP_CONF_PATH_NON_BOOT` variable in the `/etc/lvmrc` file, you need to remove the configuration file from the new path.

9. Update all references to the old name in any other files on the system. These include `/etc/fstab` for mounted file systems or swap devices, and existing mapfiles from a `vgexport` command.

Splitting a Volume Group

You can use the `vgchgid` to split an existing volume group into two or more volume groups, provided that the physical volumes to be split are self-contained; that is, any logical volumes on the physical volumes must be wholly contained on those physical volumes. For example, a splittable volume group can have logical volumes 1, 2, and 3 on physical volumes 0 and 1, and logical volumes 4, 5, and 6 on physical volumes 2, 3, 4, and 5.

In this example, volume group `vgold` contains physical volumes `/dev/disk/disk0` through `/dev/disk/disk5`. Logical volumes `lvol1`, `lvol2`, and `lvol3` are on physical volumes `/dev/disk/disk0` and `/dev/disk/disk1`, and logical volumes `lvol4`, `lvol5`, and `lvol6` are on the remaining physical volumes.

To keep `/dev/disk/disk0` and `/dev/disk/disk1` in `vgold` and split the remaining physical volumes into a new volume group named `vgnew`, follow these steps:

1. Deactivate the volume group as follows:

```
# vgchange -a n vgold
```

2. Export the volume group as follows:

```
# vgexport vgold
```

3. Change the VGID on the physical volumes to be assigned to the new volume group as follows:

```
# vgchgid -f /dev/rdisk/disk2 /dev/rdisk/disk3 \
    /dev/rdisk/disk4 /dev/rdisk/disk5
```

4. If you are using an HP-UX release before March 2008, or if you want to specify the minor number of the `vgold` group file, create it using the procedure in “[Creating the Volume Group Device File](#)” (page 44).

5. If you are using an HP-UX release before March 2008, or if you want to specify the minor number of the `vgnew` group file, create it using the procedure in “[Creating the Volume Group Device File](#)” (page 44).

6. Import the physical volumes in the old volume group as follows:

```
# vgimport /dev/vgold /dev/rdisk/disk0 /dev/rdisk/disk1
```

7. Import the physical volumes in the new volume group as follows:

```
# vgimport /dev/vgnew /dev/rdisk/disk2 /dev/rdisk/disk3 \
    /dev/rdisk/disk4 /dev/rdisk/disk5
```

8. Activate the volume groups. Disable quorum checks for the old volume group (it is missing over half its disks) as follows:

```
# vgchange -a y -q n /dev/vgold
# vgchange -a y /dev/vgnew
```

9. The logical volumes are currently defined in both volume groups. Remove the duplicate logical volumes from the volume group that no longer contains them as follows:

```
# lvremove -f vgold/lvol4 vgold/lvol5 vgold/lvol6
# lvremove -f vgnew/lvol1 vgnew/lvol2 vgnew/lvol3
```

10. The physical volumes are currently defined in both volume groups. Remove the missing physical volumes from both volume groups as follows:

```
# vgreduce -f vgold
# vgreduce -f vgnew
```

11. Enable quorum checks for the old volume group as follows:

```
# vgchange -a y -q y /dev/vgold
```

On completion, the original volume group contains three logical volumes (lvol1, lvol2, and lvol3) with physical volumes /dev/disk/disk0 and /dev/disk/disk1. The new volume group vgnew contains three logical volumes (lvol4, lvol5, and lvol6) across physical volumes /dev/disk/disk2, /dev/disk/disk3, /dev/disk/disk4, and /dev/disk/disk5.

Removing a Volume Group



TIP: It is easier to export a volume group than to remove it, because removing the volume group requires that you remove all the logical volumes and physical volumes from the volume group before running vgremove. In addition, exporting the volume group leaves the LVM information on the disks untouched, which is an advantage if you want to reimport the volume group later. For the procedure to export a volume group, see “[Exporting a Volume Group](#)” (page 57).

To remove a volume group, follow these steps:

1. Back up all user data.
2. Find the names of all logical volumes and physical volumes in the volume group. Enter the following command:

```
# vgdisplay -v /dev/vgnn
```

3. Make sure that none of those logical volumes are in use. This may require stopping applications using any logical volumes in the volume group, and unmounting file systems contained in the volume group.

Use the fuser command on each logical volume:

```
# fuser -cu /dev/vgnn/lvoln
```

4. Remove each of the logical volumes as follows:

```
# lvremove /dev/vgnn/lvoln
```

For more information, see “[Removing a Logical Volume](#)” (page 57).

5. Remove all but one of the physical volumes as follows:

```
# vgreduce /dev/vgnn /dev/disk/diskn
```

For more information, see “[Removing a Disk from a Volume Group](#)” (page 51).

6. Use the vgremove command to remove the volume group as follows:

```
# vgremove vgnn
```

Backing Up a Mirrored Logical Volume

NOTE: Mirroring requires the optional product HP MirrorDisk/UX.

You can split a mirrored logical volume into two logical volumes to perform a backup on an offline copy while the other copy stays online. When you complete the backup of the offline copy, you can merge the two logical volumes back into one. To bring the two copies back in synchronization, LVM updates the physical extents in the offline copy based on changes made to the copy that remained in use.

You can use HP SMH to split and merge logical volumes, or use the lvsplit and lvmerge commands.

To back up a mirrored logical volume containing a file system, using lvsplit and lvmerge, follow these steps:

1. Split the logical volume `/dev/vg00/lvol1` into two separate logical volumes as follows:

```
# lvsplit /dev/vg00/lvol1
```

This creates the new logical volume `/dev/vg00/lvol1b`. The original logical volume `/dev/vg00/lvol1` remains online.
2. Perform a file system consistency check on the logical volume to be backed up as follows:

```
# fsck /dev/vg00/lvol1b
```
3. Mount the file system as follows:

```
# mkdir /backup_dir
# mount /dev/vg00/lvol1b /backup_dir
```
4. Perform the backup using the utility of your choice.
5. Unmount the file system as follows:

```
# umount /backup_dir
```
6. Merge the split logical volume back with the original logical volume as follows:

```
# lvmerge /dev/vg00/lvol1b /dev/vg00/lvol1
```

NOTE: `lvsplit` is not supported on snapshot logical volumes. `lvmerge` is not supported for snapshots as well as logical volumes having snapshots.

Backing Up and Restoring Volume Group Configuration

It is important that volume group configuration information be saved whenever you make *any* change to the configuration such as:

- Adding or removing disks to a volume group
- Changing the disks in a root volume group
- Creating or removing logical volumes
- Extending or reducing logical volumes

Unlike fixed disk partitions or nonpartitioned disks that begin and end at known locations on a given disk, each volume group configuration is unique, changes, and uses space on several disks.

If you back up your volume group configuration, you can restore a corrupted or lost LVM configuration in the event of a disk failure or corruption of your LVM configuration information.

The `vgcfgbackup` command creates or updates a backup file containing the volume group configuration; it *does not back up the data within your logical volumes*. To simplify the backup process, `vgcfgbackup` is invoked automatically whenever you make a configuration change as a result of using any of the following commands:

<code>lvchange</code>	<code>lvcreate</code>	<code>lvextend</code>	<code>lvlnboot</code>	<code>lvmerge</code>
<code>lvreduce</code>	<code>lvremove</code>	<code>lvrmboot</code>	<code>lvsplit</code>	<code>vgcreate</code>
<code>pvchange</code>	<code>pvmove</code>	<code>vgchange</code> ¹	<code>vgextend</code>	<code>vgmodify</code>
<code>vgreduce</code>				

1 For volume group Version 2.2 and higher, the `vgchange` command automatically takes a backup of the volume group configuration during deactivation, provided they are not activated in read-only mode. In all earlier versions, there is no automatic backup when `vgchange` is used.

When snapshots are involved, the volume group configuration changes with respect to the snapshot data unsharing while writes are occurring on the snapshot tree. So, it is recommended that the automatic backup of the volume group configuration not be overridden by the `-A n` option during volume group deactivation.

You can display LVM configuration information previously backed up with `vgcfgbackup` or restore it using `vgcfgrestore`.

By default, `vgcfgbackup` saves the configuration of a volume group to the file `volume_group_name.conf` in the default directory `/etc/lvmconf/`. You can override this default directory setting for volume group Version 2.x by configuring a new path in the `LVMP_CONF_PATH_NON_BOOT` variable in the `/etc/lvmrc` file. For more information see `vgcfgbackup(1M)`

NOTE: After a cold installation of the HP-UX 11i v3 March 2010 release, the `/etc/lvmrc` file will contain the `LVMP_CONF_PATH_NON_BOOT` variable definition as below:

`LVMP_CONF_PATH_NON_BOOT=""`

But if you have upgraded your system from earlier version to HP-UX 11i v3 March 2010, this default definition will not exist. In such case, you can manually define the variable in `/etc/lvmrc` to configure a new path.

You can also run `vgcfgbackup` at the command line, saving the backup file in any directory you indicate. If you do, first run `vgdisplay` with the `-v` option to ensure that the all logical volumes in the volume group are shown as available/syncd. If so, then run the following command:

```
# vgcfgbackup -f pathname/filename volume_group_name
```

If you use the `-f` option with `vgcfgbackup` to save the configuration file in a non-default location, take note of and retain its location. For information on command options, see `vgcfgbackup(1M)`. Make sure backups of the root volume group are on the root file system, in case these are required during recovery.

To run `vgcfgrestore`, the physical volume must be detached. If all the data on the physical volume is mirrored and the mirror copies are current and available, you can temporarily detach the physical volume using `pvchange`, perform the `vgcfgrestore`, and reattach the physical volume. For example, to restore volume group configuration data for `/dev/disk/disk5`, a disk in the volume group `/dev/vgsales`, enter the following commands:

```
# pvchange -a n /dev/disk/disk5
# vgcfgrestore -n /dev/vgsales /dev/rdisk/disk5
# pvchange -a y /dev/disk/disk5
```

If the physical volume is not mirrored or the mirror copies are not current and available, you must deactivate the volume group with `vgchange`, perform the `vgcfgrestore`, and activate the volume group as follows:

```
# vgchange -a n /dev/vgsales
# vgcfgrestore -n /dev/vgsales /dev/rdisk/disk5
# vgchange -a y /dev/vgsales
```

These examples restore the LVM configuration to the disk from the default backup location in `/etc/lvmconf/vgsales.conf`.

For information on command options, see `vgcfgrestore(1M)`.

Moving and Reconfiguring Your Disks

This section addresses the following topics:

- “Moving Disks Within a System” (page 71)
- “Moving Disks Between Systems” (page 72)
- “Moving Data to a Different Physical Volume” (page 73)
- “Modifying Physical Volume Characteristics” (page 77)
- “Disabling a Path to a Physical Volume” (page 87)
- “Creating an Alternate Boot Disk” (page 88)
- “Mirroring the Boot Disk” (page 90)

- “[Mirroring the Boot Disk on HP 9000 Servers](#)” (page 90)
- “[Mirroring the Boot Disk on HP Integrity Servers](#)” (page 92)

You might need to do the following tasks:

- Move the disks in a volume group to different hardware locations on a system.
- Move entire volume groups of disks from one system to another.

△ CAUTION: Moving a disk that is part of your root volume group is not recommended. For more information, see [Configuring HP-UX for Peripherals](#).

The `/etc/lvmtab` and `/etc/lvmtab_p` files contain information about the mapping of LVM disks on a system to volume groups; that is, volume group names and lists of the physical volumes included in volume groups. When you do either of the previous tasks, these configuration files must be changed to reflect the new hardware locations and device files for the disks. However, you cannot edit these files directly because they are not text files. Instead, you must use `vgexport` and `vgimport` to reconfigure the volume groups, which records configuration changes in the LVM configuration files.

Moving Disks Within a System

There are two procedures for moving the disks in a volume group to different hardware locations on a system. Choose a procedure depending on whether you use persistent or legacy device files for your physical volumes; the types of device files are described in “[Legacy Device Files versus Persistent Device Files](#)” (page 12).

LVM Configuration with Persistent Device Files

If your LVM configuration uses persistent device files, follow these steps:

1. Be sure that you have an up-to-date backup for both the data within the volume group and the volume group configuration.
2. Deactivate the volume group by entering the following command:
`# vgchange -a n /dev/vgnn`
3. Physically move your disks to their desired new locations.
4. Activate the volume group as follows:
`# vgchange -a y /dev/vgnn`

LVM Configuration with Legacy Device Files

The names of legacy device files change when the hardware paths to their physical devices change. Therefore, you must update the LVM configuration by exporting and importing the volume group to use the new legacy device files. Follow these steps:

1. Be sure that you have an up-to-date backup for both the data within the volume group and the volume group configuration.
2. Deactivate the volume group as follows:
`# vgchange -a n /dev/vgnn`
3. If you want to retain the same minor number for the volume group, examine the volume group's group file as follows:
`# ls -l /dev/vgnn/group`
crw-r--r-- 1 root sys 64 0x010000 Mar 28 2004 /dev/vgnn/group
- For this example, the volume group major number is 64, and the minor number is 0x010000.
4. Remove the volume group device files and its entry from the LVM configuration files by entering the following command:
`# vgexport -v -s -m /tmp/vgnn.map /dev/vgnn`

5. Physically move your disks to their desired new locations.
6. To view the new locations, enter the following command:

```
# vgscan -v
```
7. If you are using an HPUX release before March 2008, or if you want to retain the minor number of the volume group device file, create it using the procedure in "[Creating the Volume Group Device File](#)" (page 44).

Since the group file in this example has a major number of 64 and a minor number of 0x01000000, enter the following commands:

```
# mkdir /dev/vgnn
# mknod /dev/vgnn/group c 64 0x010000
```

8. Add the volume group entry back to the LVM configuration files using the `vgimport` command as follows:

```
# vgimport -v -s -m /tmp/vgnn.map /dev/vgnn
```
9. Activate the newly imported volume group as follows:

```
# vgchange -a y /dev/vgnn
```
10. Back up the volume group configuration as follows:

```
# vcfgbackup /dev/vgnn
```

Moving Disks Between Systems

To move the disks in a volume group to different hardware locations on a different system, export the volume group from one system, physically move the disks to the other system, and import the volume group there. The procedures for exporting and importing a volume are described in "[Exporting a Volume Group](#)" (page 57) and "[Importing a Volume Group](#)" (page 58). They are illustrated in the following example.

NOTE: If the volume group contains any multipathed disks, see the note under "[Importing a Volume Group](#)" (page 58).

To move the three disks in the volume group `/dev/vg_planning` to another system, follow these steps:

1. If any of the logical volumes contain a file system, unmount the file system. If any of the logical volumes are used as secondary swap, disable swap and reboot the system; for information on secondary swap, see *HP-UX System Administrator's Guide: Configuration Management*.
2. Make the volume group and its associated logical volumes unavailable to users as follows:

```
# vgchange -a n /dev/vg_planning
```
3. Preview the removal of the volume group information from the LVM configuration files using the following command:

```
# vgexport -p -v -s -m /tmp/vg_planning.map /dev/vg_planning
```

With the `-m` option, you can specify the name of a map file that will hold the information that is removed from the LVM configuration files. The map file contains the names of all logical volumes in the volume group. You use this map file when you set up the volume group on the new system.
4. If the preview is satisfactory, remove the volume group information as follows:

```
# vgexport -v -s -m /tmp/vg_planning.map /dev/vg_planning
```

The `vgexport` command removes the volume group from the system and creates the `/tmp/vg_planning.map` file.
5. Connect the disks to the new system and copy the `/tmp/vg_planning.map` file to the new system.

6. If you are using an HP-UX release before March 2008, create the volume group device file using the procedure in “[Creating the Volume Group Device File](#)” (page 44).
7. To get device file information about the disks, run the `ioscan` command:
`# ioscan -funN -C disk`
8. To preview the import operation, run the `vgimport` command with the `-p` option:
`# vgimport -p -N -v -s -m /tmp/vg_planning.map /dev/vg_planning`
9. To import the volume group, run `vgimport` without the `-p` option as follows:
`# vgimport -N -v -s -m /tmp/vg_planning.map /dev/vg_planning`
10. Activate the newly imported volume group as follows:
`# vgchange -a y /dev/vg_planning`

Moving Data to a Different Physical Volume

You can use the `pvmove` command to move data contained in logical volumes from one disk to another disk or to move data between disks within a volume group.

For example, you can move the data from a specific logical volume from one disk to another, to use the vacated space on the first disk for another purpose. To move the data in logical volume `/dev/vg01/markets` from the disk `/dev/disk/disk4` to the disk `/dev/disk/disk7`, enter the following:

```
# pvmove -n /dev/vg01/markets /dev/disk/disk4 /dev/disk/disk7
```

On the other hand, you can move all the data contained on one disk, regardless of which logical volume it is associated with, to another disk within the same volume group. For example, do this to remove a disk from a volume group. You can use `pvmove` to move the data to other specified disks or let LVM move the data to appropriate available space within the volume group, subject to any mirroring allocation policies.

To move all data off disk `/dev/disk/disk3` and relocate it at the destination disk `/dev/disk/disk5`, enter the following command:

```
# pvmove /dev/disk/disk3 /dev/disk/disk5
```

To move all data off disk `/dev/disk/disk3` and let LVM transfer the data to available space within the volume group, enter the following command:

```
# pvmove /dev/disk/disk3
```

In each of the previous instances, if space does not exist on the destination disk, the `pvmove` command fails. The `pvmove` also provides a `-p` option for you to preview the operation before actually performing the move.

Moving a Root Disk to Another Disk

If the data to be moved is the root disk, additional steps are required. Use the following steps on an HP 9000 system as an example, to move root disk `/dev/disk/disk1` (source disk) to disk `/dev/disk/disk4` (destination disk) staying within the same volume group:

1. To make the destination disk a bootable LVM disk, enter:
`# pvcreate -B /dev/rdisk/disk4`
2. To make the disk bootable. Enter:
`# mkboot /dev/rdisk/disk4`
`# mkboot -a "hpx -a (;0)/stand/vmunix" /dev/rdisk/disk4`
3. To extend your root volume group with the destination disk, enter:
`# vgextend /dev/vg00 /dev/disk/disk4`
4. To move all physical extents from the source disk to the destination disk, enter:
`# pvmove /dev/disk/disk1 /dev/disk/disk4`

5. To reduce the source disk from the volume group, enter:

```
# vgreduce /dev/vg00 /dev/disk/disk1
```
6. To shut down and reboot from the new root disk in maintenance mode, enter:

```
ISL> hpx -lm (;0)/stand/vmunix
```
7. In maintenance mode, to update the BDRA and the LABEL file, enter:

```
# vgchange -a y /dev/vg00
# lvnboot -b /dev/vg00/lvol1
# lvnboot -s /dev/vg00/lvol2
# lvnboot -r /dev/vg00/lvol3
# lvnboot -Rv
# vgchange -a n /dev/vg00
```
8. Reboot the system normally.

pvmove Command Syntax

Beginning with the September 2008 Update, pvmove provides these options:

-p	Provides a preview of the move but does not actually perform the move.
-e <i>no_of_extents</i>	Moves the last number of physical extents, specified by <i>no_of_extents</i> .
<i>de</i>	Specifies the starting location of the destination physical extents within a destination physical volume.
<i>se1 [-se2]</i>	Defines the source physical extent range, provided along with source physical volume.

Beginning with the September 2009 Update, pvmove additionally provides these options:

-a	Moves data to achieve auto-rebalance of disk space usage within a volume group. Supported for Version 2.x volume groups only. See " Moving Data for Disk Space Balancing: Auto Re-balancing " (page 75).
-f <i>pv_path</i>	Moves data from the specified physical path, <i>pv_path</i> , to remaining physical volumes to achieve balance. Supported only with the -a option.
-s	Provides a summary preview report of the data move, but does not actually move data. Used only with the -a and -p options for a preview of data auto-rebalance.

See the *pvmove(1M)* manpage for details on all its options.

NOTE: The pvmmove command is not an atomic operation; it moves data extent by extent. The following might happen upon abnormal pvmmove termination by a system crash or kill -9:

For Version 1.0 volume groups prior to the September 2009 Update, the volume group can be left in an inconsistent configuration showing an additional pseudomirror copy for the extents being moved. You can remove the extra mirror copy using the lvrreduce command with the -m option on each of the affected logical volumes; there is no need to specify a disk.

For version 1.0 with September 2009 Update or Version 2.0 volume groups: On abnormal termination, the contiguous LV in question might have its extents laid out in a non-contiguous manner. If so, it is recommended that you run lvdisplay -v to check if the allocation policy for the LV is broken. If it is broken, then run the steps below to make contiguous the extents of this LV:

1. Change the allocation policy of this LV to default (see *lvchange(1M)*).
 2. Move the extents such that all the physical extents of the logical volume are contiguous (see *pvmmove(1M)*).
 3. Change the allocation policy back to contiguous (see *lvchange(1M)*.)
-

Moving Data for Disk Space Balancing: Auto Re-balancing

Definition

In this section, “balancing” refers to maintaining the percentage of free and used space on each involved physical volume equal to the total percentage of free and used space on all the physical volumes that are selected in the balance operation; calculations are based upon the number of physical extents. Auto re-balancing refers to the automatic movement of extents based on the optimal number of extents calculated automatically for each logical volume on each physical volume involved in the re-balance operation. Beginning with the September 2009 Update of HP-UX, you can achieve automatic re-balancing on Version 2.x volume groups using the -a option of the pvmmove command.

Purpose

Balancing achieves better space utilization for existing storage, especially after physical volumes have been recently added or deleted. Additionally, balancing can give better I/O performance for the target logical volumes.

Usage

There are three different ways to use the -a option:

- `pvmmove -a vg_name`
all logical volumes within the volume group `vg_name` are auto re-balanced
- `pvmmove -a lv_path [pv_path | pvg_name]`
only the logical volumes specified by `lv_path` are auto re-balanced. They are balanced across all the physical volumes within the volume group or the optional targets specified by the physical volume `pv_path` or physical volume group `pvg_name`
- `pvmmove -a -f pv_path`
all extents are moved from the physical volume `pv_path` to the remaining physical volumes within the volume group; then, all logical volumes are re-balanced across those same remaining physical volumes. You can use this when you want to remove or replace the physical volume specified by `pv_path`.

Preview Mode

Before performing a pvmmove for balancing, you can run a preview of the operation using the -p option. In addition, the -s option can be used to generate a preview summary report of the requested re-balance without actually performing any move. The -s option is only valid with the -p and -a options. See the *pvmmove(1M)* manpage for more details on the pvmmove features.

Example

Consider a volume group having the configuration below:

- Three physical volumes each of size 1245 extents:
/dev/disk/disk10, /dev/disk/disk11, and /dev/disk/disk12.
- Three logical volumes all residing on same disk, /dev/disk/disk10:
/dev/vg_01/lvol1 (Size = 200 extents),
/dev/vg_01/lvol2 (Size = 300 extents),
/dev/vg_01/lvol3 (Size = 700 extents).

The pvmove below generates a summary report to preview the re-balance of volume group vg_01:

```
# pvmove -p -s -a vg_01
--- Summary Report of re-balance operation ---
The optimal percentage of free space per PV = 68%
The optimal percentage of used space per PV = 32%
--- Before re-balance operation ---
The average percentage of free space on each PV =~ (68 +/- 32)%
The average percentage of used space on each PV =~ (32 +/- 32)%
--- After re-balance operation ---
The average percentage of free space on each PV =~ (68 +/- 0)%
The average percentage of used space on each PV =~ (32 +/- 0)%
```

The output above is interpreted as follows:

- The average percentage free space on each PV, before re-balance operation, is approximately in the range of (68 - 32 =) 36% to (68 + 32 =) 100% and after re-balance operation would be approximately (68 +/- 0 =) 68%.
- The average percentage used space on each PV, before re-balance operation, is approximately in the range of (32 - 32 =) 0% to (32 + 32 =) 64% and after re-balance operation would be approximately (32 +/- 0 =) 32%.

The pvmove command below actually moves the data for balancing:

```
# pvmove -a vg_01
Automatic re-balance operation successfully completed.
Volume Group configuration for /dev/vg_01 has been saved in
/etc/lvmconf/vg_01.conf
```

NOTE: In the automatic re-balance mode, the pvmove command tries to achieve an optimal rebalance, but it does not guarantee an optimal rebalance; and there are scenarios where the user can perform a more optimal rebalance manually than the one provided by the pvmove auto rebalance operation.

The auto re-balance operation does not move/re-balance boot, root, swap, or dump physical volumes, so resulting system is still bootable.

The auto re-balance will fail if the operation involves moving pre-allocated extents belonging to a snapshot. For information about pre-allocated extents of snapshots, see "[Creating and Administering Snapshot Logical Volumes](#)" (page 103).

Creating a Spare Disk

NOTE: Disk sparing requires the optional product HP MirrorDisk/UX.

Version 2.x volume groups do not support disk sparing.

To configure a spare physical volume into a volume group for which you want protection against disk failure, follow these steps before a disk failure actually occurs:

1. Use the `pvcreate` command to initialize the disk as an LVM disk.

NOTE: Do not use the `-B` option because spare physical volumes cannot contain boot information.

```
# pvcreate /dev/rdisk/disk3
```

2. Ensure the volume group has been activated, as follows:

```
# vgchange -a y /dev/vg01
```

3. Use the `vgextend` command with `-z y` to designate one or more physical volumes as spare physical volumes within the volume group. For example:

```
# vgextend -z y /dev/vg01 /dev/disk/disk3
```

Alternately, you can change a physical volume with no extents currently allocated within it into a spare physical volume using the `pvchange` command with the `-z y` option. For example:

```
# pvchange -z y /dev/disk/disk3
```

For more information on disk sparing, see [“Increasing Disk Redundancy Through Disk Sparing” \(page 27\)](#).

Reinstating a Spare Disk

NOTE: Disk sparing requires the optional product HP MirrorDisk/UX.

Version 2.x volume groups do not support disk sparing.

After a failed disk has been repaired or a decision has been made to replace it, follow these steps to reinstate it and return the spare disk to its former standby status:

1. Physically connect the new or repaired disk.

2. Restore the LVM configuration to the reconnected disk using `vgcfgrestore` as follows:

```
# vgcfgrestore -n /dev/vg01 /dev/rdisk/disk1
```

3. Ensure the volume group has been activated as follows:

```
# vgchange -a y /dev/vg01
```

4. Be sure that allocation of extents is now allowed on the replaced disk as follows:

```
# pvchange -x y /dev/disk/disk1
```

5. Use the `pvmove` command to move the data from the spare to the replaced physical volume. For example:

```
# pvmove /dev/disk/disk3 /dev/disk/disk1
```

The data from the spare disk is now back on the original disk or its replacement, and the spare disk is returned to its role as a standby empty disk.

For more information on disk sparing, see [“Increasing Disk Redundancy Through Disk Sparing” \(page 27\)](#).

Modifying Physical Volume Characteristics

The `vgmodify` command enables you to modify a volume group to adapt to changes in physical volumes. In particular, you can adjust the volume group to recognize size changes in physical volumes, and (for Version 1.0 volume groups only) you can change a physical volume type between bootable and nonbootable.

Beginning with the March 2009 Update, the `vgmodify` command can be run online (volume group activated and applications running) for Version 1.0 and Version 2.x volume groups. This allows physical volume sizes to be adjusted online. This is known as Dynamic LUN Expansion (DLE) and Dynamic LUN Contraction (DLC).

Handling Size Increases

From the LUN Side:

Disk arrays typically allow a LUN to be resized. If the volume group is activated during size increase, this is known as Dynamic LUN Expansion (DLE). If you increase the size of a LUN, follow these steps to incorporate the additional space into the volume group:

1. Increase the LUN size using the instructions for the array.
2. Verify that the LUN is expanded by running the `diskinfo` command.

From the LVM Side (for Version 1.0 Volume Groups):

Do the following for Version 1.0 volume groups.

1. Run `vgmodify` to detect any physical volume size changes. It also reports whether all of the space can be made available to the volume group.
2. If `vgmodify` reports that the maximum number of physical extents per physical volume (`max_pe`) is too small to accommodate the new size, use `vgmodify` with the `-t` and `-n` options to determine a new value for `max_pe`, as described in “[Modifying Volume Group Parameters](#)” (page 58).
3. Review the values by running `vgmodify` with the new settings and the `-r` option.
4. Deactivate the volume group, if needed.

TIP: Beginning with the September 2008 Update, you can use `vgmodify` to recognize and accommodate size increases *without* deactivating the volume group. For more information, see `vgmodify(1M)`.

5. Commit any new value of `max_pe` and update the physical volume information by running `vgmodify` without the `-r` option.
6. If needed, activate the volume group. To verify that the increased space is available, run the `vgdisplay` and `pvdisplay` commands.

From the LVM Side (for Version 2.x Volume Groups):

Do the following for Version 2.x volume groups.

NOTE: Beginning with the September 2009 Update, the `vgmodify` also supports Version 2.x volume groups, but the volume groups must be in active mode to run `vgmodify`. For more information, see `vgmodify(1M)`.

If the LUN for a physical volume was dynamically expanded using an array utility, the `vgmodify -E` option can be used to make available the new space for user data. After the physical volume is reconfigured, you can expand logical volumes with `lvextend`, or you can create new logical volumes with `lvcreate`.

The extra space added to a physical volume through an array utility can also be used for additional LVM configuration data when increasing the maximum VG size through the `vgmodify -S` option (see “[Modifying Volume Group Parameters](#)” (page 58)). The `vgmodify -E` and `-S` options can also be used together to perform DLE and to increase maximum VG size in a single step.

To handle dynamic LUN expansion for a PV in a Version 2.x volume group, follow these steps:

1. Use `vgdisplay -v` to determine which PV in the volume group require expansion and decide how much each LUN needs to be expanded. This decision can be based on which logical volumes need to grow and which PVs they can reside on. This decision can also be based on whether or not the DLE is being performed as part of increasing the maximum VG size. See the “[Modifying Volume Group Parameters](#)” (page 58) section for a description for DLE in combination with increasing VG size limit.
2. Use the `vgmodify -r` (review) option to verify which physical volumes require reconfiguration for DLE and to check if the reconfiguration will yield the desired result. If no PVs are specified

on the command line, `vgmodify` checks *all* PVs in the volume group. Optionally, you can specify the PVs you want to check in the command line after the VG name. .

3. Perform the actual DLE reconfiguration (run `vgmodify` without the `-r` review option). If no PVs are specified, `vgmodify` attempts reconfiguration on *all* PVs in the volume group (if it detects them as having been expanded). Optionally, you can list the specific PVs you want to reconfigure in the command line after the VG name.
4. Run `vgdisplay -v` to verify the result of the DLE reconfiguration: to check the total number of extents for the VG, the total number of free extents for the VG, the number of extents for each PV, and the number of free extents for each PV.

Example: Version 1.0 Volume Group

For example, to increase the size of the physical volume `/dev/rdisk/disk6` from 4 GB to 100000000 KB, follow these steps:

1. Increase the LUN size using the instructions for the disk array.
2. Run `vgmodify` with the `-v` and `-r` options to check whether any disks have changed in size and whether all of the space on the physical volumes can be used.

```
# vgmodify -v -r vg32
```

Current Volume Group settings:

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

`/dev/rdisk/disk6` Warning: Max_PE_per_PV for the volume group
(1016) too small for this PV (3051).

Using only 1016 PEs from this physical volume.

"`/dev/rdisk/disk6`" size changed from 4194304 to 100000000kb

An update to the Volume Group IS required

New Volume Group settings:

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

Review complete. Volume group not modified

The expanded physical volume requires 3051 physical extents to use all its space, but the current `max_pe` value limits this to 1016.

3. To determine the optimal values for `max_pv` and `max_pe`, run `vgmodify -t`, with and without the `-n` as follows:

```
# vgmodify -t vg32
```

Current Volume Group settings:

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176
VGRA space (Kbytes) without PE renumbering	896
VGRA space (Kbytes) PE renumbering lower	32768

Volume Group optimized settings (no PEs renumbered):

max_pv(-p)	max_pe(-e)	Disk size (Mb)
2	53756	1720193
3	35836	1146753
4	26876	860033
...		

28	3836	122753
30	3580	114561
32	3324	106369
35	3068	98177
38	2812	89985
...		
255	252	8065

The table shows that without renumbering physical extents, a `max_pv` of 35 or lower permits a `max_pe` sufficient to accommodate the increased physical volume size.

```
# vgmodify -v -t -n vg32
Volume Group configuration for /dev/vg32 has been saved in
/etc/lvmconf/vg32.conf
```

Current Volume Group settings:

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

VGRA space (Kbytes) on all Physical Volumes:

PV	current	-n
/dev/rdisk/disk6	896	32768
/dev/rdisk/disk5	896	32768
Summary	896	32768

Volume Group optimized settings (PEs renumbered lower):

max_pv (-p)	max_pe (-e)	Disk size (Mb)
61	65535	2097152
62	65532	2097056
63	64252	2056096
...		
251	16124	516000
252	16048	513568
255	15868	507808

The table shows that if physical extents are renumbered, all values of `max_pv` permit a `max_pe` large enough to accommodate the increased physical volume size.

For this example, select a `max_pv` of 10, which permits a `max_pe` value of 10748.

- Preview the changes by using the `-r` option to `vgmodify` as follows:

```
# vgmodify -p 10 -e 10748 -r vg32
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1016
PE Size (Mbytes)	32
VGRA Size (Kbytes)	176

The current and new Volume Group parameters differ.

`"/dev/rdisk/disk6"` size changed from 4194304 to 100000000kb
An update to the Volume Group IS required

New Volume Group settings:

Max LV	255
Max PV	10
Max PE per PV	10748
PE Size (Mbytes)	32
VGRA Size (Kbytes)	896

Review complete. Volume group not modified

- Deactivate the volume group as follows:

```
# vgchange -a n vg32
Volume group "vg32" has been successfully changed.
```

6. Commit the new values as follows:

```
# vgmodify -p 10 -e 10748 vg32
Current Volume Group settings:
      Max LV        255
      Max PV         16
      Max PE per PV 1016
      PE Size (Mbytes) 32
      VGRA Size (Kbytes) 176
The current and new Volume Group parameters differ.
"/dev/rdisk/disk6" size changed from 4194304 to 1000000000kb
An update to the Volume Group is required

New Volume Group settings:
      Max LV        255
      Max PV         10
      Max PE per PV 10748
      PE Size (Mbytes) 32
      VGRA Size (Kbytes) 896
New Volume Group configuration for "vg32" has been saved in
"/etc/lvmconf/vg32.conf"
Old Volume Group configuration for "vg32" has been saved in
"/etc/lvmconf/vg32.conf.old"
Starting the modification by writing to all Physical Volumes
Applying the configuration to all Physical Volumes from
"/etc/lvmconf/vg32.conf"
Completed the modification process.
New Volume Group configuration for "vg32" has been saved in
"/etc/lvmconf/vg32.conf.old"
Volume group "vg32" has been successfully changed.
```

7. Activate the volume group and verify the changes by entering the following commands:

```
# vgchange -a y vg32
Activated volume group
Volume group "vg32" has been successfully changed.
```

```
# vgdisplay vg32
--- Volume groups ---
VG Name          /dev/vg32
VG Write Access  read/write
VG Status        available
Max LV           255
Cur LV           0
Open LV          0
Max PV           10
Cur PV           2
Act PV           2
Max PE per PV   10748
VGDA             4
PE Size (Mbytes) 32
Total PE         3119
Alloc PE         0
Free PE          3119
Total PVG        0
Total Spare PVs  0
Total Spare PVs in use 0
VG Version       1.0
```

Example: Version 2.x Volume Group

In this example, the `vgmodify -E` option is used to increase the number of extents available for user data for two LUNs whose size has been expanded by an array utility from 100 GB to 200 GB.

1. Run online vgmodify in review mode to verify that the physical volumes require reconfiguration for the DLE and to preview the number of new extents to be added:

```
# vgmodify -r -a -E vg1
Physical volume "/dev/disk/disk46" requires reconfiguration for expansion.
Current number of extents: 12790
Number of extents after reconfiguration: 25590
Physical volume "/dev/disk/disk46" was not changed.
Physical volume "/dev/disk/disk47" requires reconfiguration for expansion.
Current number of extents: 12790
Number of extents after reconfiguration: 25590
Physical volume "/dev/disk/disk47" was not changed.
```

2. Run online vgmodify in change mode to perform the reconfiguration and add new extents to the physical volumes:

```
# vgmodify -a -E vg1
Reconfiguration to expand physical volume "/dev/disk/disk46" succeeded.
Previous number of extents: 12790
Number of extents after reconfiguration: 25590
Physical volume "/dev/disk/disk46" was changed.
Volume Group configuration for /dev/vg1 has been saved in
/etc/lvmconf/vg1.conf

Reconfiguration to expand physical volume "/dev/disk/disk47" succeeded.
Previous number of extents: 12790
Number of extents after reconfiguration: 25590
Physical volume "/dev/disk/disk47" was changed.

Volume Group configuration for /dev/vg1 has been saved in
/etc/lvmconf/vg1.conf
```

3. Verify that the physical volumes were reconfigured and that there are new extents available with the `vgdisplay -v` command:

```
# vgdisplay -v vg1

--- Volume groups ---
VG Name          /dev/vg1
...
PE Size (Mbytes)    8
Total PE          51180
Alloc PE          25580
Free PE           25600
...
VG Version        2.1
VG Max Size       32t
VG Max Extents   4194304
...

--- Logical volumes ---
LV Name          /dev/vg1/lvol1
LV Status         available/syncd
LV Size (Mbytes) 102320
Current LE        12790
Allocated PE      25580
Used PV           2

--- Physical volumes ---
PV Name          /dev/disk/disk46
PV Status         available
Total PE          25590
Free PE           12800
...
PV Name          /dev/disk/disk47
PV Status         available
Total PE          25590
Free PE           12800
```

For more information and examples, see the *Using the vgmodify Command to Perform LVM Volume Group Dynamic LUN Expansion (DLE) and Contraction (DLC)* white paper.

Handling Size Decreases

- ⚠ **CAUTION:** A similar procedure can also be used when the size of a physical volume is decreased. However, there are limitations:
- Sequence:
The sequence must be reversed to avoid data corruption.
For an increase in size, the sequence is:
 1. Increase the LUN size from the array side.
 2. Then, increase the volume group size from the LVM side.For a decrease in size, the sequence is:
 1. Decrease the volume group size from the LVM side.
 2. Then, decrease the LUN size from the array side.
 - Volume Group Status
Prior to the March 2009 Update, `vgmodify` for handling LUN contraction was only supported on Version 1.0 volume groups, and the volume group had to be deactivated before attempting any reduction. If you reduce the size of the LUN while the volume group was activated, LVM would mark the physical volume as unavailable.
Beginning with the March 2009 Update, `vgmodify` for handling LUN contraction can now be performed on Version 1.0 or Version 2.x volume groups and the volume group can be activated (this is known as Dynamic LUN Contraction, or DLC). For an active volume group, you can use the `-C` and `-a` options of `vgmodify` to recognize and accommodate size decreases *without* deactivating the volume group. For more information, see `vgmodify(1M)`.
 - Physical Extent Quantity
If the physical volume has any allocated physical extents beyond the target size, `vgmodify` prints an error message and exits without changing the physical volume. In this case, you must be prepared to restore the LUN to its original size (ensuring the same disk space is allocated).

Supported LUN Changes per Release

This table shows what is supported with regards to whether the volume group needs to be deactivated for LUN size changes:

Table 8 Supported LUN Changes Per Release

	LUN Size Increase	LUN Size Decrease
Prior to September 2008 Update	Volume group must be deactivated; only Version 1.0 volume groups supported.	Not supported
Beginning with the September 2008 Update	Volume group can be either deactivated or active (use <code>vgmodify-E-a</code>); only Version 1.0 volume group supported.	Not supported
Beginning with the March 2009 Update	Volume group can be either deactivated or active (use <code>vgmodify-E-a</code>); Version 1.0 and Version 2.x volume groups supported.	Volume group can be either deactivated or active (use <code>vgmodify-C-a</code>); Version 1.0 and 2.x volume groups supported

Changing Physical Volume Boot Types

NOTE: Version 2.0 and 2.1 volume groups do not support bootable physical volumes.

When a physical volume is initialized for LVM use, it can be made bootable or nonbootable. Bootable physical volumes require additional space in their LVM metadata for boot utilities and

information. If a physical volume was accidentally initialized as bootable, you can convert the disk to a nonbootable disk and reclaim LVM metadata space.

CAUTION: The boot volume group requires at least one bootable physical volume. Do not convert all of the physical volumes in the boot volume group to nonbootable, or your system will not boot.

To change a disk type from bootable to nonbootable, follow these steps:

1. Use `vgcfgrestore` to determine if the volume group contains any bootable disks.
2. Run `vgmodify` twice, once with the `-B n` and once without it. Compare the available values for `max_pe` and `max_pv`.
3. Choose new values for `max_pe` and `max_pv`. Review the values by running `vgmodify` with the new settings and the `-r` option.
4. Deactivate the volume group.
5. Commit the changes by running `vgmodify` without the `-r` option.
6. Activate the volume group. Run the `vgcfgrestore` or `pvdisplay` commands to verify that the disk type has changed.

For example, to convert any bootable disks in volume group `vg`, follow these steps:

1. Check if any physical volumes in `vg01` are bootable as follows:

```
# vgcfgrestore -l -v -n vg01
Volume Group Configuration information in "/etc/lvmconf/vg01.conf"
VG Name /dev/vg01
---- Physical volumes : 1 ----
PV      Type      Size (kb) Start (kb) PVkey
c2t1d0 Bootable  35566480    2912      0

max_pv 16 max_pe 1085 max_lv 255
```

2. To determine which values of `max_pe` and `max_pv` are available, run the following command:

```
# vgmodify -t -B n vg01 /dev/rdsk/c2t1d0
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1085
PE Size (Mbytes)	32
VGRA Size (Kbytes)	208
VGRA space (Kbytes) without PE renumbering	2784
VGRA space (Kbytes) PE renumbering lower	32768

```
Volume Group optimized settings (no PEs renumbered) :
max_pv(-p) max_pe(-e) Disk size (Mb)
5          65535      2097122
6          56828      1818498
...
255        1276       40834
```

Compare the values if the disk is made non-bootable, and if it is not. Enter the following command:

```
# vgmodify -t vg01
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1085
PE Size (Mbytes)	32
VGRA Size (Kbytes)	208
VGRA space (Kbytes) without PE renumbering	768
VGRA space (Kbytes) PE renumbering lower	768

```
Volume Group optimized settings (no PEs renumbered) :
max_pv(-p) max_pe(-e) Disk size (Mb)
```

```

1           65535      2097120
2           45820      1466240
...
255          252       8064

```

If you change the disk type, the VGRA space available increases from 768 KB to 2784KB (if physical extents are not renumbered) or 32768 KB (if physical extents are renumbered).

Changing the disk type also permits a larger range of `max_pv` and `max_pe`. For example, if `max_pv` is 255, the bootable disk can only accommodate a disk size of 8064 MB, but after conversion to nonbootable, it can accommodate a disk size of 40834 MB.

- For this example, select a `max_pv` value of 6, which permits a `max_pe` value of 56828. Preview the changes by entering the following command:

```
# vgmodify -r -p 6 -e 56828 -B n vg01 /dev/rdsk/c2t1d0
Current Volume Group settings:
```

Max LV	255
Max PV	16
Max PE per PV	1085
PE Size (Mbytes)	32
VGRA Size (Kbytes)	208

The current and new Volume Group parameters differ.

An update to the Volume Group IS required

New Volume Group settings:

Current Volume Group settings:

Max LV	255
Max PV	6
Max PE per PV	56828
PE Size (Mbytes)	32
VGRA Size (Kbytes)	2784

Review complete. Volume group not modified

- Deactivate the volume group as follows:

```
# vgchange -a n vg01
```

Volume group "vg01" has been successfully changed.

- Commit the new values as follows:

```
# vgmodify -p 6 -e 56828 -B n vg01 /dev/rdsk/c2t1d0
```

Current Volume Group settings:

Max LV	255
Max PV	16
Max PE per PV	1085
PE Size (Mbytes)	32
VGRA Size (Kbytes)	208

The current and new Volume Group parameters differ.

An update to the Volume Group IS required

New Volume Group settings:

Current Volume Group settings:

Max LV	255
Max PV	6
Max PE per PV	56828
PE Size (Mbytes)	32
VGRA Size (Kbytes)	2784

New Volume Group configuration for "vg01" has been saved in
`"/etc/lvmconf/vg01.conf"`

Old Volume Group configuration for "vg01" has been saved in
`"/etc/lvmconf/vg01.conf.old"`

Starting the modification by writing to all Physical Volumes
Applying the configuration to all Physical Volumes from

`"./etc/lvmconf/vg01.conf"`

Completed the modification process.

New Volume Group configuration for "vg01" has been saved in

```

"/etc/lvmconf/vg01.conf.old"
Volume group "vg01" has been successfully changed.

6. Activate the volume group and verify the changes as follows:
# vgchange -a y vg01
Activated volume group
Volume group "vg01" has been successfully changed.

# vgcfgbackup vg01
Volume Group configuration for /dev/vg01 has been saved in
/etc/lvmconf/vg01.conf

# vgcfgrestore -l -v -n vg01
Volume Group Configuration information in "/etc/lvmconf/vg01.conf"
VG Name /dev/vg01

---- Physical volumes : 1 ----
PV      Type      Size (kb) Start (kb) PVkey
c2t1d0 Non-Boot 35566480    2912      0

max_pv 6 max_pe 56828 max_lv 255

```

Disabling a Path to a Physical Volume

- ① **IMPORTANT:** This procedure only disables LVM's use of the link. The `pvchange` command does not prevent diagnostics or an application from accessing the physical volume.

By default, the mass storage stack uses all the paths available to access a physical volume, independently of the paths configured in LVM. Disabling a path in LVM does not prevent the native multipathing from using that path. Use the `scsimgr` command to disable I/O along a path or to disable the native multipathing.

You can temporarily disable LVM's use of one or all of the physical paths to a physical volume using the `pvchange` command. Disabling a path, also known as **detaching the link**, causes LVM to close that path to the device and stop using it. This can be useful if you want to guarantee that a link is idle, such as when you are running diagnostics on an I/O card, replacing an I/O card, or replacing the disk containing the physical volume.

Detaching a link to a physical volume is a temporary operation, not a permanent one. If you want to permanently remove a link or physical volume from the volume group, use `vgreduce` instead, as described in “[Removing a Disk from a Volume Group](#)” (page 51).

To detach a link to a physical volume, use the `-a` option to `pvchange`. For example, to disable the link through the device `/dev/disk/disk33`, enter the following command:

```
# pvchange -a n /dev/disk/disk33
```

If you are using LVM's alternate links for multipathed disks, each link uses a different legacy device files. In that situation, to detach all links to a physical volume, use `N` as the argument to the `-a` option:

```
# pvchange -a N /dev/dsk/c5t0d0
```

Detaching one or more links to a physical volume does not necessarily cause LVM to stop using that physical volume entirely. If the detached link is the primary path to the device, LVM begins using any available alternate link to it. LVM stops using the physical volume only when all the links to it are detached.

If all the links to a device are detached, the associated physical volume becomes unavailable to the volume group. The links remain associated with the volume group but LVM does not send any I/O requests to the physical volume until it is reattached. This means that the data on that physical volume becomes temporarily unavailable; consequently, you must make sure that any availability requirements for that data can be satisfied by mirroring before you make the device unavailable by detaching it.

Detaching a link does not disable sparing. That is, if all links to a physical volume are detached and a suitable spare physical volume is available in the volume group, LVM uses it to reconstruct the detached disk. For more information on sparing, see “[Increasing Disk Redundancy Through Disk Sparing](#)” (page 27).

You can view the LVM status of all links to a physical volume using `vgdisplay` with the `-v` option.

Restoring a detached link to a physical volume, or **reattaching** it, makes that link available to the volume group. LVM can begin using that link as necessary to access the disk.

To reattach a specific path to a physical volume, use the `pvchange` command with the `-a` option. For example, enter the following command:

```
# pvchange -a y /dev/dsk/c5t0d0
```

Because detaching a link to a physical volume is temporary, all detached links in a volume group are reattached when the volume group is activated, either at boot time or with an explicit `vgchange` command, such as the following:

```
# vgchange -a y /dev/vg02
```

Creating an Alternate Boot Disk

NOTE: Version 2.0 and 2.1 volume groups do not support bootable physical volumes. You cannot create an alternate boot disk in a Version 2.0 or 2.1 volume group.

With non-LVM disks, a single root disk contains all the attributes needed for boot, system files, primary swap, and dump. Using LVM, a single root disk is replaced by a pool of disks, a **root volume group**, which contains all of the same elements but allowing a root logical volume, a boot logical volume, a swap logical volume, and one or more dump logical volumes. Each of these logical volumes must be contiguous, that is, contained on a single disk, and they must have bad block relocation disabled. (Other noncontiguous logical volumes can be used for user data.) For more information on the swap and dump devices and their configuration, see *HP-UX System Administrator's Guide: Configuration Management*.

The root logical volume contains the operating system software and the root file system (/). The boot logical volume contains the boot file system (/stand). You can combine the root and boot logical volumes into a single logical volume or keep them separate. Whether you use a single combined root-boot logical volume, or separate root and boot logical volumes, the logical volume used to boot the system must be the first logical volume on its physical volume. It must begin at physical extent 0000 to boot the system in maintenance mode.

If you newly install your HP-UX system and choose the LVM configuration, a root volume group is automatically configured (/dev/vg00), as are separate root (/dev/vg00/lvo13) and boot (/dev/vg00/lvo11) logical volumes. If you currently have a combined root and boot logical volume and you want to reconfigure to separate them after creating the boot logical volume, use the `lvlnboot` command with the `-b` option to define the boot logical volume to the system, taking effect the next time the system is booted.

The swap logical volume is the system’s primary swap area and is typically used for dump. The swap logical volume is often on the same physical disk as the root logical volume. However, you can configure it (and the dump logical volumes) on a different physical disk than the root logical volume.

If you create your root volume group with multiple disks, use the `lvextend` command to place the boot, root, and primary swap logical volumes on the boot disk.

TIP: You can use `pvmove` to move the data from an existing logical volume to another disk if necessary to make room for the root logical volume. For more information, see “[Moving Data to a Different Physical Volume](#)” (page 73).

To create a new root volume group with an alternate boot disk, follow these steps:

1. Create a bootable physical volume.
 - a. On an HP Integrity server, partition the disk using the `idisk` command and a partition description file, then run `insf` as described in "[Mirroring the Boot Disk on HP Integrity Servers](#)" (page 92).
 - b. Run `pvcreate` with the `-B` option. On an HP Integrity server, use the device file denoting the HP-UX partition:

```
# pvcreate -B /dev/rdisk/disk6_p2
```

On an HP 9000 server, use the device file for the entire disk:

```
# pvcreate -B /dev/rdisk/disk6
```
2. Create a directory for the volume group. For example:

```
# mkdir /dev/vgroot
```
3. Create a device file named `group` in the previously described directory as follows:

```
# mknod /dev/vgroot/group c 64 0xnn0000
```
4. Create the root volume group, specifying each physical volume to be included, as follows:

```
# vgcreate /dev/vgroot /dev/disk/disk6
```
5. Place boot utilities in the boot area as follows:

```
# mkboot /dev/rdisk/disk6
```
6. Add an autoboot file to the disk boot area as follows:

```
# mkboot -a "hpx" /dev/rdisk/disk6
```
7. Create the boot logical volume. To create a boot logical volume named `bootlv` with size 512 MB, enter the following commands:

```
# lvcreate -C y -r n -n bootlv /dev/vgroot
# lvextend -L 512 /dev/vgroot/bootlv /dev/disk/disk6
```

NOTE: The `-r n` option for bad block reallocation (in this and next two steps) is used for Version 1.0 volume groups only. It is ignored by Version 2.x volume groups.

8. Create the primary swap logical volume. For example, to create a primary swap logical volume named `swalv` with size 2 GB on the same disks as the boot logical volume, enter the following commands:

```
# lvcreate -C y -r n -n swalv /dev/vgroot
# lvextend -L 2048 /dev/vgroot/swalv /dev/disk/disk6
```
9. Create the root logical volume. For example, to create a root logical volume named `rootlv` with size 1 GB, enter the following commands:

```
# lvcreate -C y -r n -n rootlv /dev/vgroot
# lvextend -L 1024 /dev/vgroot/rootlv /dev/disk/disk6
```
10. Specify that `bootlv` is the boot logical volume as follows:

```
# lvnboot -b /dev/vgroot/bootlv
```
11. Specify that `rootlv` is the root logical volume as follows:

```
# lvnboot -r /dev/vgroot/rootlv
```
12. Specify that `swalv` is the primary swap logical volume as follows:

```
# lvnboot -s /dev/vgroot/swalv
```
13. Specify that `swalv` is also to be used for dump as follows:

```
# lvnboot -d /dev/vgroot/swalv
```
14. Verify the configuration as follows:

```
# lvnboot -v /dev/vgroot
```

Boot Definitions for Volume Group /dev/vgroot:
Physical Volumes belonging in Root Volume Group:

```
/dev/disk/disk6 -- Boot Disk
Boot: bootlv    on:      /dev/disk/disk6
Root: rootlv   on:      /dev/disk/disk6
Swap: swaplv   on:      /dev/disk/disk6
Dump: swaplv   on:      /dev/disk/disk6, 0
```

15. Once the boot and root logical volumes are created, create file systems for them. For example:

```
# mkfs -F hfs /dev/vgroot/rbootlv
# mkfs -F vxfs /dev/vgroot/rrootlv
```

NOTE: On HP Integrity servers, the boot file system can be VxFS. Enter the following command:

```
# mkfs -F vxfs /dev/vgroot/rbootlv
```

Mirroring the Boot Disk

NOTE: Mirroring requires the optional product, HP MirrorDisk/UX.

Version 2.0 and 2.1 volume groups do not support boot disks, so you cannot mirror the boot disk in a Version 2.0 or 2.1 volume group.

After you create mirror copies of the root, boot, and primary swap logical volumes, if any of the underlying physical volumes fail, the system can use the mirror copy on the other disk and continue. When the failed disk comes back online, it is automatically recovered, provided the system has not been rebooted.

If the system reboots before the disk is back online, reactivate the volume group to update the LVM data structures that track the disks within the volume group. You can use vgchange -a y even though the volume group is already active.

For example, you can reactivate volume group vg00 by entering the following command:

```
# vgchange -a y /dev/vg00
```

As a result, LVM scans and activates all available disks in the volume group vg00, including the disk that came online after the system rebooted.

The procedure for creating a mirror of the boot disk is different for HP 9000 and HP Integrity servers. HP Integrity servers use partitioned boot disks.

NOTE: These examples include creating a mirror copy of the primary swap logical volume. The primary swap mirror does not need to be on a specific disk or at a specific location, but it must be allocated on contiguous disk space. The recommended mirror policy for primary swap is to have the mirror write cache and the mirror consistency recovery mechanisms disabled.

When primary swap is mirrored and your primary swap device also serves as a dump area, be sure that mirror write cache and mirror consistency recovery are set to off at boot time to avoid losing your dump. To reset these options, reboot your system in maintenance mode and use the lvchange command with the -M n and -c n options.

Mirroring the Boot Disk on HP 9000 Servers

To set up a mirrored root configuration, you must add a disk to the root volume group, mirror all the root logical volumes onto it, and make it bootable. For this example, the disk to be added is at path 0/1/1/0.0x1.0x0 and has device special files named /dev/rdisk/disk4 and /dev/disk/disk4. Follow these steps:

1. Make sure the device files are in place. For example:

```
# insf -e -H 0/1/1/0.0x1.0x0
```

The following device files now exist for this disk:

```
/dev/[r]disk/disk4
```

2. Create a bootable physical volume as follows:

```
# pvcreate -B /dev/rdisk/disk4
```
 3. Add the physical volume to your existing root volume group as follows:

```
# vgextend /dev/vg00 /dev/disk/disk4
```
 4. Place boot utilities in the boot area as follows:

```
# mkboot /dev/rdisk/disk4
```
 5. Add an autoboot file to the disk boot area as follows:

```
# mkboot -a "hpx" /dev/rdisk/disk4
```
-

NOTE: If you expect to boot from this disk only when you lose quorum, you can use the alternate string `hpx -lq` to disable quorum checking. However, HP recommends configuring your root volume group to minimize quorum losses, by using at least three physical volumes and no single points of failure, as described in “[Planning for Recovery](#)” (page 34).

6. The logical volumes on the mirror boot disk must be extended in the same order that they are configured on the original boot disk. Determine the list of logical volumes in the root volume group and their order. For example:

```
# pvdisplay -v /dev/disk/disk0 | grep 'current.*0000 $'
 00000 current /dev/vg00/lvol1 00000
 00038 current /dev/vg00/lvol2 00000
 00550 current /dev/vg00/lvol3 00000
 00583 current /dev/vg00/lvol4 00000
 00608 current /dev/vg00/lvol5 00000
 00611 current /dev/vg00/lvol6 00000
 00923 current /dev/vg00/lvol7 00000
 01252 current /dev/vg00/lvol8 00000
```

7. Mirror each logical volume in `vg00` (the root volume group) onto the specified physical volume. For example:

```
# lvextend -m 1 /dev/vg00/lvol1 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
# lvextend -m 1 /dev/vg00/lvol2 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
# lvextend -m 1 /dev/vg00/lvol3 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
# lvextend -m 1 /dev/vg00/lvol4 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
# lvextend -m 1 /dev/vg00/lvol5 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
# lvextend -m 1 /dev/vg00/lvol6 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
# lvextend -m 1 /dev/vg00/lvol7 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
# lvextend -m 1 /dev/vg00/lvol8 /dev/disk/disk4
The newly allocated mirrors are now being synchronized.
This operation will take some time. Please wait ....
```

TIP: To shorten the time required to synchronize the mirror copies, use the `lvextend` and `lvsync` command options introduced in the September 2007 release of HP-UX 11i Version 3. These options enable you to resynchronize logical volumes in parallel rather than serially. For example:

```
# lvextend -s -m 1 /dev/vg00/lvol1 /dev/disk/disk4
# lvextend -s -m 1 /dev/vg00/lvol2 /dev/disk/disk4
# lvextend -s -m 1 /dev/vg00/lvol3 /dev/disk/disk4
# lvextend -s -m 1 /dev/vg00/lvol4 /dev/disk/disk4
# lvextend -s -m 1 /dev/vg00/lvol5 /dev/disk/disk4
# lvextend -s -m 1 /dev/vg00/lvol6 /dev/disk/disk4
# lvextend -s -m 1 /dev/vg00/lvol7 /dev/disk/disk4
# lvextend -s -m 1 /dev/vg00/lvol8 /dev/disk/disk4
# lvsync -T /dev/vg00/lvol*
```

8. Update the root volume group information as follows:

```
# lvnboot -R /dev/vg00
```

9. Verify that the mirrored disk is displayed as a boot disk and that the boot, root, and swap logical volumes appear to be on both disks as follows:

```
# lvnboot -v
```

10. Specify the mirror disk as the alternate boot path in nonvolatile memory as follows:

```
# setboot -a 0/1/1/0.0x1.0x0
```

11. Add a line to `/stand/bootconf` for the new boot disk using `vi` or another text editor as follows:

```
# vi /stand/bootconf
1 /dev/disk/disk4
```

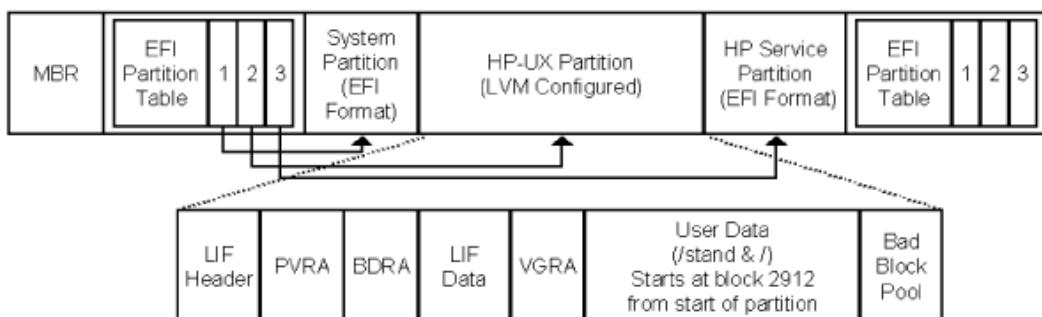
Where the literal "1" (lower case L) represents LVM.

Mirroring the Boot Disk on HP Integrity Servers

The procedure to mirror the root disk on Integrity servers is similar to the procedure for HP 9000 servers. The difference is that Integrity server boot disks are partitioned; you must set up the partitions, copy utilities to the EFI partition, and use the HP-UX partition device files for LVM commands.

Figure 5 shows the disk layout of a boot disk. The disk contains a Master Boot Record (MBR) and EFI partition tables that point to each of the partitions. The `idisk` command creates the partitions (see `idisk(1M)`).

Figure 5 Sample LVM Disk Layout on an HP Integrity Server



For this example, the disk to be added is at hardware path 0/1/1/0.0x1.0x0, with device special files named /dev/disk/disk2 and /dev/rdisk/disk2. Follow these steps:

1. Partition the disk using the idisk command and a partition description file.

- a. Create a partition description file. For example:

```
# vi /tmp/idf
```

In this example, the partition description file contains the following information:

```
3
EFI 500MB
HPUX 100%
HPSP 400MB
```

The values in the example represent a boot disk with three partitions: an EFI partition, an HP-UX partition, and an HPSP. Boot disks of earlier HP Integrity servers might have an EFI partition of only 100 MB and might not contain the HPSP partition.

- b. Partition the disk using idisk and your partition description file as follows:

```
# idisk -f /tmp/idf -w /dev/rdisk/disk2
```

- c. To verify that your partitions are correctly laid out, enter the following command:

```
# idisk /dev/rdisk/disk2
```

2. Create the device files for all the partitions. For example:

```
# insf -e -H 0/1/1/0.0x1.0x0
```

The following device files now exist for this disk:

```
/dev/[r]disk/disk2 (this refers to the entire disk)
/dev/[r]disk/disk2_p1 (this refers to the efi partition)
/dev/[r]disk/disk2_p2 (this will be the hp-ux partition)
/dev/[r]disk/disk2_p3 (this refers to the service partition)
```

3. Create a bootable physical volume using the device file denoting the HP-UX partition. For example:

```
# pvcreate -B /dev/rdisk/disk2_p2
```

4. Add the physical volume to your existing root volume group as follows:

```
# vgextend vg00 /dev/disk/disk2_p2
```

5. Place boot utilities in the boot area. Copy EFI utilities to the EFI partition, and use the device special file for the entire disk as follows:

```
# mkboot -e -l /dev/rdisk/disk2
```

6. Add an autoboot file to the disk boot area as follows:

```
# mkboot -a "boot vmunix" /dev/rdisk/disk2
```

NOTE: If you expect to boot from this disk only when you lose quorum, you can use the alternate string boot vmunix -lq to disable quorum checking. However, HP recommends configuring your root volume group to minimize quorum losses, by using at least three physical volumes and no single points of failure, as described in ["Planning for Recovery" \(page 34\)](#).

7. The logical volumes on the mirror boot disk must be extended in the same order that they are configured on the original boot disk. Determine the list of logical volumes in the root volume group and their order. For example:

```
# pvdisplay -v /dev/disk/disk0_p2 | grep 'current.*0000 $'
00000 current /dev/vg00/lvol1 00000
00010 current /dev/vg00/lvol2 00000
00138 current /dev/vg00/lvol3 00000
00151 current /dev/vg00/lvol4 00000
00158 current /dev/vg00/lvol5 00000
00159 current /dev/vg00/lvol6 00000
```

```
00271 current /dev/vg00/lvol7 00000  
00408 current /dev/vg00/lvol8 00000
```

8. Mirror each logical volume in vg00 (the root volume group) onto the specified physical volume.
For example:

```
# lvextend -m 1 /dev/vg00/lvol1 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....  
# lvextend -m 1 /dev/vg00/lvol2 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....  
# lvextend -m 1 /dev/vg00/lvol3 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....  
# lvextend -m 1 /dev/vg00/lvol4 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....  
# lvextend -m 1 /dev/vg00/lvol5 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....  
# lvextend -m 1 /dev/vg00/lvol6 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....  
# lvextend -m 1 /dev/vg00/lvol7 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....  
# lvextend -m 1 /dev/vg00/lvol8 /dev/disk/disk2_p2  
The newly allocated mirrors are now being synchronized.  
This operation will take some time. Please wait ....
```

NOTE: If lvextend fails with following message:

"m": illegal option

HP MirrorDisk/UX is not installed.



TIP: To shorten the time required to synchronize the mirror copies, use the lvextend and lvsync command options introduced in the September 2007 release of HP-UX 11i Version 3. These options enable you to resynchronize logical volumes in parallel rather than serially.
For example:

```
# lvextend -s -m 1 /dev/vg00/lvol1 /dev/disk/disk2_p2  
# lvextend -s -m 1 /dev/vg00/lvol2 /dev/disk/disk2_p2  
# lvextend -s -m 1 /dev/vg00/lvol3 /dev/disk/disk2_p2  
# lvextend -s -m 1 /dev/vg00/lvol4 /dev/disk/disk2_p2  
# lvextend -s -m 1 /dev/vg00/lvol5 /dev/disk/disk2_p2  
# lvextend -s -m 1 /dev/vg00/lvol6 /dev/disk/disk2_p2  
# lvextend -s -m 1 /dev/vg00/lvol7 /dev/disk/disk2_p2  
# lvextend -s -m 1 /dev/vg00/lvol8 /dev/disk/disk2_p2  
# lvsync -T /dev/vg00/lvol*
```

9. Update the root volume group information as follows:

```
# lvnboot -R /dev/vg00
```

10. Verify that the mirrored disk is displayed as a boot disk and that the boot, root, and swap logical volumes appear to be on both disks as follows:

```
# lvnboot -v
```

11. Specify the mirror disk as the alternate boot path in nonvolatile memory as follows:

```
# setboot -a 0/1/1/0.0x1.0x0
```

12. Add a line to /stand/bootconf for the new boot disk using vi or another text editor as follows:

```
# vi /stand/bootconf
l /dev/disk/disk2_p2
```

Where the literal "l" (lower case L) represents LVM.

Migrating a Volume Group to New Disks: vgmove

Beginning with September 2009 Update, LVM provides a new vgmove command to migrate data in a volume group from an old set of disks to a new set of disks.

The vgmove operation moves data in a volume group from old storage to new storage while adhering to these conditions:

- During migration, the volume group is online with no downtime to customer applications.
- The volume group device special file (DSF) is not changed. From an application perspective, the volume group remains unchanged.
- The extent allocation policies used by the logical volumes stay honored.
- The migration succeeds even if the number of destination physical volumes is different from the number of original physical volumes.
- High availability and recoverability are not compromised at any time during the data migration.

To migrate a volume group, you provide a file containing the mapping of source (old) to destination (new) disks, one line per source disk, in the following format:

```
source_pv_1 destination_pv_1_1 destination_pv_1_2 ....
source_pv_2 destination_pv_2_1 destination_pv_2_2 ....
...
source_pv_n destination_pv_n_1 destination_pv_n_2 ....
```

This diskmap file can be manually created or generated with the -i option in vgmove, by simply providing a list of destination disks to which to migrate the volume group. Once the diskmap file is created, you can review and edit it as needed, then run vgmove again to perform the actual migration.

The vgmove syntax is:

```
/usr/sbin/vgmove [-A autobackup] [-p] -f diskmapfile vg_name
```

```
/usr/sbin/vgmove [-A autobackup] [-p] -i diskfile -f diskmapfile vg_name
```

where:

- | | |
|-----------------------|---|
| -A <i>autobackup</i> | Sets automatic backup during this operation. |
| -p | Provides a preview of the operation without actually migrating the data. |
| -f <i>diskmapfile</i> | Specifies the file containing the mapping of source (old) to destination (new) disks |
| -i <i>diskfile</i> | Provides a list of destination disks to generate a <i>diskmapfile</i> that will contain the mapping of source to destination disks. When this option is used, no actual data is migrated. |
| <i>vg_name</i> | Specifies the name of the volume group whose data is to be migrated. |

Example of Migrating a Volume Group to New Disks

In this example, the volume group to be migrated is called /dev/vg00 and currently resides in old disk /dev/disk/disk1. The new disks to migrate the data to are /dev/disk/disk10 and /dev/disk/disk11.

- Instead of manually creating a diskmap file, mapping the old source to new destination disks, the `-i` option is used to generate a diskmap file for the migration. The user provides a list of destination disks, called `newdiskfile` in this example.

```
# cat newdiskfile  
/dev/disk/disk10  
/dev/disk/disk11  
  
# vgmove -i newdiskfile -f diskmap.txt /dev/vg00
```

- The resulting `diskmap.txt` file contains the mapping of old source disks to new destination disks:

```
# cat diskmap.txt  
/dev/disk/disk1 /dev/disk/disk10 /dev/disk/disk11
```

- Then another `vgmove` command is issued to actually perform the data migration:

```
# vgmove -f diskmap.txt /dev/vg00
```

Upon successful completion of data migration by `vgmove`, the original source disk will be automatically removed from the volume group using `vgreduce`. If a destination disk is not already part of the volume group, it will be added using `vgextend`.

After successful migration, the destination disks are added to the LVM configuration files; namely, `/etc/lvmtab` or `/etc/lvmtab_p`. The source disks along with their alternate links are removed from the LVM configuration files.

The volume group must be activated before running the `vgmove` command. If the `vgmove` command is interrupted before it completes, the volume group is in the same state it was at the beginning of the `vgmove` command. The migration can be continued by running the `vgmove` command with the same options and disk mapping file.

NOTE: For volume groups that support bootable physical volumes (Version 1.0 and Version 2.2 or higher), `vgmove` supports moving boot PVs. If the diskmap file contains a mix of bootable and non-bootable source PVs, the bootable source PV must be specified first.

Migrating a Logical Volume to New Disks: `lvmmove`

Beginning with the HP-UX 11i v3 March 2010 Update, LVM provides a new `lvmmove` command to migrate a logical volume within a volume group. The intent is to balance a logical volume either within the existing set of physical volumes or within a new set of physical volumes in the volume group.

The volume group must be activated before running the `lvmmove` command. If the `lvmmove` command is interrupted before it completes, the volume group is left in a consistent state, and can still be activated. The migration can be continued by running the `lvmmove` command with the same options and mapping file. See the `lvmmove(1M)` manpage for more information about the `lvmmove` command and its options.

NOTE: For volume groups that support bootable PVs (Version 1.0 and Version 2.2 or higher), `lvmmove` will not move root, boot, swap, or dump PVs.

The `lvmmove` migration will fail for a space efficient snapshot that has pre-allocated extents. For more information about snapshots, see “[Creating and Administering Snapshot Logical Volumes](#)” ([page 103](#))

Administering File System Logical Volumes

This section describes special actions you must take when working with file systems inside logical volumes. It addresses the following topics:

- [“Creating a File System” \(page 97\)](#)
- [“Extending a File System” \(page 98\)](#)
- [“Reducing the Size of a File System” \(page 99\)](#)
- [“Backing Up a VxFS Snapshot File System” \(page 101\)](#)

TIP: When dealing with file systems, you can use HP SMH or a sequence of HP-UX commands. For most tasks, using HP SMH is quicker and simpler. You do not have to explicitly perform each of the following distinct tasks; rather, proceed from the HP SMH Disk and File Systems area. HP SMH performs all the necessary steps for you.

Creating a File System

When creating either an HFS or VxFS file system in a logical volume, you can use HP SMH or a sequence of HP-UX commands. If you choose to use HP-UX commands directly, the following list describes the subtasks for creating a file system.

1. If you create a new file system of a type other than HFS, you might need to reconfigure the new type into the kernel. Normally, VxFS has been configured into the kernel as part of the default configuration. See *HP-UX System Administrator’s Guide: Configuration Management* for information on how to add a file system type.
2. Estimate the size required for the logical volume. To estimate the size requirements for a logical volume containing a file system, see “[Setting Up Logical Volumes for File Systems](#)” ([page 21](#)).
3. Determine if there is sufficient disk space available in the volume group. Use the `vgdisplay` command to calculate this information, as described in “[Creating a Logical Volume](#)” ([page 52](#)). If there is not enough space within the volume group, you can add a disk to the volume group, as described in “[Adding a Disk to a Volume Group](#)” ([page 51](#)).
4. Create the logical volume. Use `lvcreate`, as described in “[Creating a Logical Volume](#)” ([page 52](#)).

5. Create the file system using the character device file. For example:

```
# newfs -F fstype /dev/vg02/lvol1
```

If you do not use the `-F fstype` option, then `newfs` creates a file system based on the content of your `/etc/fstab` file. If there is no entry for the file system in `/etc/fstab`, then the file system type is determined from the file `/etc/default/fs`. For information on additional options, see `newfs(1M)`.

When creating a VxFS file system, file names will be long automatically.

For HFS, use the `-s` or `-l` option to specify a file system with short or long file names, respectively. By default, the length of file system names are consistent with those of the root file system. Short file names are 14 characters maximum. Long file names allow up to 255 characters. HP recommends using long file names for flexibility; files created on other systems that use long file names can be moved to your system without being renamed.

6. After you have created a file system, mount it for users to access it, and add it to `/etc/fstab` so that it is automatically mounted at boot time.

Extending a File System

Extending a file system inside a logical volume is a two-step task: extending the logical volume, then extending the file system. The first step is described in “[Extending a Logical Volume](#)” (page 53). The second step, extending the file system itself, depends on the following factors:

- What type of file system is involved? If it is HFS or VxFS? HFS requires the file system to be unmounted to be extended.

Check the type of file system. For example:

```
# /usr/sbin/fstyp /dev/vg01/lvol2
vxfs
```

- If the file system is VxFS, do you have the base VxFS product or the OnlineJFS product? If you have only the base VxFS product, you must unmount the file system before extending it.

To see if the OnlineJFS product is installed, enter the following command:

```
# swlist -l product | grep -i OnlineJFS
OnlineJFS B.11.31 Online features of the VxFS File System
```

- Can you unmount the file system? To unmount system directories such as `/var` and `/usr`, you must be in single-user mode.

- Is the file system the root file system (/)? If so, there are two complications:

- The logical volume containing the root file system is created with the contiguous allocation policy, so it might not be possible to extend it in place.
- The root file system cannot ever be unmounted, even if you change to single-user state.

If you are using VxFS as your root file system and have the OnlineJFS product, you can extend the original root file system without unmounting, provided contiguous disk space is available.

Otherwise, to extend the current root file system, you must create and mount another root disk which enables you to work with the unmounted original root disk, extending it if contiguous disk space is still available. If the original disk does not have contiguous disk space available, instead of expanding the original root disk, you can create a new root file system on another larger disk.



CAUTION: This procedure can fail on a VxFS file system already at 100% capacity (Error 28). You must remove some files before you attempt this operation.

Once you have the answers to these questions, follow these steps:

1. If the file system must be unmounted, unmount it.
 - a. Be sure no one has files open in any file system mounted to this logical volume and that it is no user's current working directory. For example:

```
# fuser -cu /work/project5
```

If the logical volume is in use, confirm that the underlying applications no longer need it. If necessary, stop the applications.

NOTE: If the file system is exported using NFS to other systems, verify that no one is using those other systems, then unmount it on those systems.

 - b. If you cannot stop the applications using the logical volume, or it is a system directory such as /var or /usr, change to single-user state as follows:

```
# /sbin/shutdown
```
 - c. Unmount the file system as follows:

```
# /sbin/umount /dev/vg01/lvol1
```
 2. Extend the logical volume. For example:

```
# /sbin/lvextend -L 332 /dev/vg01/lvol1
```

This increases the size of this volume to 332 MB.
 3. Extend the file system size to the logical volume size. If the file system is unmounted, use the `extendfs` command as follows:

```
# /sbin/extendfs /dev/vg01/r1vol1
```

If you did not have to unmount the file system, use the `fsadm` command instead. The new size is specified in terms of the block size of the file system. In this example, the block size of the file system /work/project5 is 1 KB. To extend the file system to 332 MB, the number of blocks is 339968 (332 times 1024). For example:

```
# fsadm -b 339968 /work/project5
```
 4. If you unmounted the file system, mount it again.
 - a. If you had to change to single-user state, reboot the system.

```
# /sbin/reboot -r
```

You can skip any additional steps to mount the file system and export it, since the boot process mounts and exports any file systems.
 - b. Remount the file system as follows:

```
# /sbin/mount /dev/vg01/r1vol1 /mount_point
```

NOTE: If the file system will continue to be used by NFS clients, export it on the server (`exportfs -a`) and remount it on the clients (`mount -a`).

 5. Verify that the file system reflects the expansion by entering `bdf`, `df`, or `fsadm -E`.

Reducing the Size of a File System

You might want to shrink a file system that has been allocated more disk space than it needs, allowing that disk space to be freed for other uses.

Reducing the size of a file system is more complicated than extending it. Because of file system block allocation strategies, data can be scattered throughout the logical volume. Reducing the logical volume reclaims space at the end of the logical volume, requiring file system drivers to coalesce and rearrange data blocks ahead of time. Most types of file system are unable to do such coalescence, so you must back up the data in the file system, reduce the logical volume, create a new file system in the smaller logical volume, and restore the data from your backup.

The only current file system type able to do online coalescence and size reduction is OnlineJFS, and it can fail in some cases.

Reducing a File System Created with OnlineJFS

Using the `fsadm` command shrinks the file system, provided the blocks it attempts to deallocate are not currently in use; otherwise, it fails. If sufficient free space is currently unavailable, file system defragmentation of both directories and extents might consolidate free space toward the end of the file system, allowing the contraction process to succeed when retried.

For example, suppose your VxFS file system is currently 6 GB. However, you decide you only need 2 GB with an additional 1 GB for reserve space. Thus, you want to resize the file system to 3 GB. Use `fsadm` with the `-b` option to specify the new size of the file system in sectors, then reduce the size of the logical volume to match. Assuming the file system sector size is 1K, use the following commands:

```
# fsadm -b 3145728 /home  
# lvreduce -L 3072 /dev/vg01/lvol5
```

Reducing a File System Created with HFS or VxFS

To reduce the size of a file system created with HFS or VxFS, follow these steps:

1. Be sure no one has files open in any file system on the logical volume and that the logical volume is no user's current working directory. For example::

```
# fuser -cu /dev/vg01/lvol5
```

NOTE: If the file system is exported using NFS to other systems, verify that no one is using those other systems, then unmount the file system on those systems before unmounting it on the server.

2. Back up the data in the logical volume.

For example, to back up `/work/project5` to the system default tape device, enter the following command:

```
# tar cv /work/project5
```

3. Remove the data in the file system the logical volume is mounted to, as follows:

```
# rm -r /work/project5
```

Because `/work/project5` is a mount point, `rm -r` does not remove the directory itself.

4. Unmount the file system to which the logical volume is mounted as follows:

```
# umount /work/project5
```

5. Reduce the size of the logical volume as follows:

```
# lvreduce -L 500 /dev/vg01/lvol5
```

This command reduces the logical volume `/dev/vg01/lvol5` to 500 MB.

6. Create a new file system in the reduced logical volume, using the character device file. For example:

```
# newfs -f fstype /dev/vg01/r1vol5
```

7. Mount the logical volume as follows:

```
# mount /dev/vg01/lvol5 /work/project5
```

8. Recover the data from the backup. For example:

```
# tar xv
```

This recovers all the contents of a tape in the system default drive.

9. If `/work/project5` continues to be used by NFS clients, re-export it on the server (`exportfs -a`) and remount it on the clients (`mount -a`).

Backing Up a VxFS Snapshot File System

NOTE: Creating and backing up a VxFS snapshot file system requires that you have the optional HP OnlineJFS product installed on your system. For more information, see *HP-UX System Administrator's Guide: Configuration Management*.

VxFS enables you to perform backups without taking the file system offline by making a snapshot of the file system, a read-only image of the file system at a moment in time. The primary file system remains online and continues to change. After you create the snapshot, you can back it up with any backup utility except the `dump` command.

To create and back up a VxFS snapshot file system, follow these steps:

1. Determine how large the snapshot file system must be, and create a logical volume to contain it. Use `bdf` to assess the primary file system size and consider the following:
 - Block size of the file system (1,024 bytes per block by default)
 - How much the data in this file system is likely to change (HP recommends 15 to 20% of total file system size)

For example, to determine how large to make a snapshot of `lvol4`, mounted on the `/home` directory, examine its `bdf` output:

```
# bdf /home
filesystem          kbytes    used    avail %used mounted on
/dev/vg00/lvol4      40960   38121     2400   94% /home
```

Allowing for 20% change to this 40 MB file system, create a logical volume of eight blocks (8 MB).

2. Use `lvcreate` to create a logical volume to contain the snapshot file system. For example:

```
# lvcreate -L 8 -n snap /dev/vg02
```

This creates an 8 MB logical volume called `/dev/vg02/snap`, which can contain a snapshot file system of `lvol4`.

3. Create a directory for the mount point of the snapshot file system. For example:

```
# mkdir /tmp/house
```

4. Create and mount the snapshot file system. For example:

```
# mount -f vxfs -o snapof=/dev/vg00/lvol4 /dev/vg02/snap /tmp/house
```

In this example, a snapshot is taken of logical volume `/dev/vg00/lvol4`, contained in logical volume `/dev/vg02/snap`, and mounted on `/tmp/house`.

5. Back up the snapshot file system with any backup utility except `dump`.

For example, use `tar` to archive the snapshot file system `/tmp/house`, ensuring that the files on the tape have relative path names, as follows:

```
# cd tmp; tar cf /dev/rtape/tape0BEST house
```

Alternatively, the following `vxdump` command backs up a snapshot file system `/tmp/house`, which has extent attributes:

```
# vxdump -0 -f /dev/rtape/tape0BEST /tmp/house
```

Administering Swap Logical Volumes

NOTE: Version 2.0 and 2.1 volume groups do not support swap logical volumes.

When you enable a swap area within a logical volume, HP-UX determines how large the area is, and it uses no more space than that. If your disk has enough remaining contiguous space, you can increase the size of your primary swap area by using the `lvextend` command (or HP SMH) to extend the logical volume, then reboot the system. This procedure allows HP-UX to use the extra space.

If you plan device swap areas in addition to primary swap, you get the best performance when the device swap areas are on different physical volumes. This configuration allows for the interleaving of I/O to the physical volumes when swapping occurs.

To create interleaved swap, create multiple logical volumes for swap, with each logical volume on a separate disk. You must use HP-UX commands to help you obtain this configuration. HP SMH does not allow you to create a logical volume on a specific disk. See “[Extending a Logical Volume to a Specific Disk](#)” (page 54).

You can configure your swap space as described in *HP-UX System Administrator's Guide: Overview*.

NOTE: You must reboot the system for the system to recognize changes to the swap configuration.

Creating a Swap Logical Volume

To create a swap logical volume, use the `lvcreate` command. You must set a contiguous allocation policy using the `-C y` option. For example:

```
# lvcreate -C y -n swap_lv1 /dev/vgnn
```

For more information, see `lvcreate(1M)`.

If you create a logical volume to use as primary swap, use the `lvolnboot` command with the `-s` option to update the swap information used by LVM at boot. For example:

```
# lvolnboot -s /dev/vgnn/swap_lv1
```

Extending a Swap Device

If you are using a logical volume for swap, you must increase the logical volume size *before* increasing the swap size. You can extend the logical volume using `lvextend` or HP SMH.

Swap logical volumes must be contiguous, so extending the logical volume succeeds only if there are physical extents available at the end of the existing logical volume. If contiguous disk space is not available, create a new contiguous logical volume for primary swap within the root volume group. You do not need to designate a specific disk. For example:

```
# lvcreate -C y -L 48 -n pswap /dev/vgroot
```

After creating a logical volume to use as primary swap, use `lvolnboot` to update the boot information:

```
# lvolnboot -s /dev/vgroot/pswap
```

Reducing the Size of a Swap Device

If you are using a logical volume for swap, you must reduce the swap size *before* reducing the size of the logical volume. You can reduce the size of the logical volume using `lvreduce` or HP SMH.

Administering Dump Logical Volumes

NOTE: Version 2.0 and 2.1 volume groups do not support dump logical volumes.

This section describes LVM information about using logical volumes as dump devices. For information on configuring and managing dump devices, see *HP-UX System Administrator's Guide: Overview*.

Creating a Dump Logical Volume

To create a dump logical volume, use the `lvcreate` command. You must set a contiguous allocation policy using the `-C y` option, and disable bad block relocation using the `-r n` option. (Note: the `-r` option is only required for Version 1.0 volume groups. It is not supported on Version 2.x and will be ignored.) For example:

```
# lvcreate -C y -r n -n dump_lv1 /dev/vgnn
```

For more information, see `lvcreate(1M)`.

After creating a logical volume to be used as a dump device, use the `lvolnboot` command with the `-d` option to update the dump information used by LVM. If you created a logical volume `/dev/vg00/lvol2` for use as a dump area, update the boot information by entering the following:

```
# lvolnboot -d /dev/vg00/lvol2
```

Removing a Dump Logical Volume

To discontinue the use of a currently configured logical volume as a dump device, use the `lvrmboot` command with the `-d` option. For example:

```
# lvrmboot -d /dev/vg00/lvol2
```

You can then use the logical volume for other purposes. To remove the logical volume completely, see "["Removing a Logical Volume" \(page 57\)](#)".

Creating and Administering Snapshot Logical Volumes

Beginning with the HP-UX 11i v3 March 2010 Update, LVM provides the ability to create *snapshot* logical volumes. A snapshot represents a point-in-time image of a logical volume. LVM snapshots let you:

- Use snapshots to back up data on the logical volume without splitting the logical volume. This reduces the space requirement.
- Create snapshots faster than manually copying over multiple point-in-time copies of the logical volume.
- Create multiple snapshots of a logical volume. This enables you to have images of the logical volume for multiple point in time without allocating space equivalent to the entire size of the logical volume per copy.

Up to 255 snapshots can be created off a single logical volume. The logical volume (referred to as the *original* logical volume) and all of its snapshots together form a *snapshot tree* and maintains a successor/predecessor relationship with each other.

When created, a snapshot shares all its data with that of the original logical volume. The snapshot gets a copy of its own data only when a write occurs (copy before write) onto itself or its successor. This process is called *data unsharing*. The smallest unit for data unsharing is the volume group's unshare unit, specified by the `vgcreate -U` option. See `vgcreate(1M)` for details.

Snapshots are by default read-only, but the user can choose to create writable snapshots. See `lvcreate(1M)` and `lvchange(1M)` for more details.

For more information about snapshot logical volumes and snapshot trees, see the `lvm(7)` manpage and the *LVM Snapshot Logical Volumes* white paper.

Types of Snapshots

Snapshots can be of two types: *fully-allocated* and *space-efficient*.

- When a fully allocated snapshot is created, the number of extents required for the snapshot is allocated immediately, just like for a normal logical volume. However, the data contained in the original logical volume is not copied over to these extents. The copying of data occurs through the data unsharing process.
- When a space-efficient snapshot is created, the user is expected to specify the number of extents that LVM needs to set aside to be used for unsharing in the future. These extents are referred to as pre-allocated extents.

When an unshare operation on a snapshot can not find the required extents in the pre-allocated extent pool, the snapshot will be marked as over-committed and inoperative. To avoid this, after the snapshot creation, the number of pre-allocated extents may be increased using the `lvextend` command with the `-l` or `-L` option. Refer to the `lvcreate(1M)` and `lvextend(1M)` manpages for details..

Beginning with the HP-UX 11i v3 September 2010 Update, LVM can be enabled to automatically increase the pre-allocated extents of a snapshot when the threshold is reached. With this update, space efficient snapshots' threshold value can also be configured by the user. Administrators can use the `lvcreate` or `lvchange` command to enable automatic increase of pre-allocated extents (via the `-e` option) and specify the threshold value (via the `-P` option). The threshold value and whether automatic pre-allocation is enabled are displayed in the `lvdisplay` command output. When automatic increase of pre-allocated extents is enabled, additional messages are printed to `syslog` when the threshold is reached. See these manpages for full details: `lvcreate(1M)`, `lvchange(1M)`, `lvdisplay(1M)`, and `lvm(7)`.

Creating a Snapshot Logical Volume

LVM snapshot logical volumes are created with the `-S` option of the `lvcreate` command. For example:

```
# lvcreate -S s -L 10 /dev/vg03/lvol1
```

where:

`-S snap_type` Specifies the creation of a snapshot logical volume with the `s` value for `snap_type`.

`-L space` Allocates space for the space-efficient snapshot logical volume, in megabytes. Instead of the `-L` option, you can use the `-l le_number` option to allocate space in terms of number of pre-allocated extents. The `-L` and `-l` options are mutually exclusive. If neither the `-L` nor `-l` option is used, a fully allocated snapshot is created.

Refer to the `lvcreate(1M)` manpage for more details and full syntax for creating snapshot logical volumes. Refer to the “[Naming Conventions for LVM](#)” (page 13) section for naming conventions used for snapshot logical volumes.

NOTE: Once snapshots are created, their size and the number of mirror copies cannot be changed using the `lvextend` or `lvreduce` commands.

Removing Snapshot Logical Volumes

Any snapshot on the snapshot tree can be deleted using the `lvremove` command. A new `lvremove` option, `-F`, has been introduced which allows the user to delete a snapshot and all its predecessors using a single `lvremove` command.

For example, for a LV having 4 snapshots, where `/dev/vg01/lvol1_S5` is the latest one, all of the snapshots can be deleted using the following command:

```
# lvremove -F /dev/vg01/lvol1_S5
```

Refer to the `lvremove(1M)` manpage for more details and full syntax for deleting logical volumes on the snapshot tree.

Displaying Snapshot Information

The `vgdisplay`, `lvdisplay`, and `pvdisplay` commands will now display additional information when snapshots are involved.

A summary of the additional fields displayed by these commands is listed here. See the respective manpages and the *LVM Snapshot Logical Volumes* white paper for more detailed information.

The `vgdisplay` command now additionally displays the following data:

- The number of snapshot logical volumes in the volume group.
- The unshare unit size for the volume group.
- The number of pre-allocated physical extents in the volume group.

- The current and the maximum snapshot capacity for the volume group.
- The listing of snapshot logical volumes and some snapshot related attributes for each.

The `lvdisplay` command now additionally displays the following data for a logical volume which has snapshots associated with it:

- The number of snapshots associated with the logical volume
- A listing of the names of the snapshot logical volumes associated with the logical volume.

For a snapshot logical volume, the following new details are displayed:

- The LV Status will now also display the type of snapshot.
- The state of the snapshot (operative, inoperative, over-committed, etc).
- The total number and the current number of pre-allocated logical extents associated with the snapshot volume.
- The number of logical extents that are currently allocated or unshared for the snapshot.
- The current number of pre-allocated physical extents for the snapshots.
- The name of the logical volume of which this logical volume is a snapshot.
- The creation timestamp for the snapshot.
- The state of whether automatic increase of pre-allocated extent is enabled or not.
- The threshold value of a space efficient snapshot logical volume.
- A listing of the distribution of the unshared logical and physical extents for the snapshot.
- A listing of the pre-allocated extents for the snapshot.
- A listing of the logical extents of the snapshot that have been unshared, and mapping of these logical extents to physical extents.

The `pvdisplay` command now additionally displays the following data:

- The number of pre-allocated physical extents on it.
- A listing of the distribution of unshared physical extents per logical volume that lie on this physical volume.
- The extent mapping for the physical volume will indicate pre-allocated extents with a status of "pre-allocated".

NOTE: The value of "Pre-allocated LE" should be the sum of "Current pre-allocated LE" and "Unshared LE." But, in some instances, this might not be shown as the case while an operation that changes the logical volume size or while unsharing of extents is in progress. The correct information will be displayed once the operation is complete.

For more information about the LVM snapshot feature and limitations when snapshots are involved, see the `lvm(7)` manpage and the *LVM Snapshot Logical Volumes* white paper.

Managing the `lvmpud` Daemon

The `lvmpud` daemon is used to handle online shared volume group reconfiguration (introduced in the HP-UX 11i v3 September 2009 Update) and automatic increase of pre-allocated extents for space efficient snapshots (introduced in the 11i v3 September 2010 Update). This daemon must be running to enable these two features.

A startup script (`/sbin/init.d/lvm`) is delivered as of the HP-UX 11i v3 September 2009 Update to start and stop this daemon. A configuration file (`/etc/rc/config.d/lvmconf`) is provided to configure automatic startup. To have `lvmpud` started automatically at each boot, include the following line in the `/etc/rc.config.d/lvmconf` file:

```
START_LVMPUD=1
```

To stop the daemon manually, use the following command:

```
#/sbin/init.d/lvm stop
```

To start the daemon manually again, use the following command:

```
#/sbin/init.d/lvm start
```

To start the `lvmpud` daemon even if `START_LVMPUD` is not set to 1, enter:

```
#lvmpud
```

See the `lvmpud(1M)` manpage for full details of this daemon.

Hardware Issues

This section describes hardware-specific issues dealing with LVM.

Integrating Cloned LUNs

Certain disk arrays can create clones of their LUNs. For example, the HP XP product enables you to split off a set of LUNs, called Business Copies (BCs), which are copies of existing LUNs.

Cloned disks have the same information in their LVM headers as the original disks, which violates LVM's requirement that each disk have a unique identifier. To make the cloned disks usable with LVM, use the `vgchgid` command to change their volume group identifier (VGID).

All of the physical volumes to be changed must belong to the same volume group. Therefore, if you are changing multiple physical volumes, specify all of them in a single invocation of `vgchgid`. Otherwise, they are assigned different VGIDs.

For example, you have a volume group containing four physical volumes and need to create a BC for each physical volume. If you run `vgchgid` on only two BCs, `vgchgid` modifies the VGID on those two BCs. If you then run `vgchgid` again with all four BCs, `vgchgid` reports that they belong to different volume groups. To correct this, you can either run `vgchgid` on the two unmodified BCs and then use the four BCs in two separate volume groups, or you can merge back the two modified BCs and split them off again before finally running `vgchgid` with all four BCs.

After running `vgchgid` on a set of physical volumes, use `vgimport` to import them into a new volume group. For example:

1. Make BC copies and create new device files using the instructions for the array.
2. Change the VGID on the cloned disks as follows:

```
# vgchgid /dev/rdisk/disk49 /dev/rdisk/disk50
```

3. If you are using an HP-UX release before March 2008, create the volume group group file using the procedure in “[Creating the Volume Group Device File](#)” (page 44).

4. Import the physical volumes as follows:

```
# vgimport /dev/vg04 /dev/rdisk/disk49 /dev/rdisk/disk50
```

5. Back up the volume group configuration information as follows:

```
# vgcfgbackup /dev/vg04
```

6. Activate the volume group as follows:

```
# vgchange -a y /dev/vg04
```

4 Troubleshooting LVM

This chapter provides conceptual troubleshooting information as well as detailed procedures to help you plan for LVM problems, troubleshoot LVM, and recover from LVM failures. It contains the following information:

- “Troubleshooting Overview” (page 107)
- “I/O Errors” (page 109)
- “Volume Group Activation Failures” (page 111)
- “Root Volume Group Scanning” (page 114)
- “LVM Boot Failures” (page 115)
- “Problems After Reducing the Size of a Logical Volume” (page 115)
- “Disk Troubleshooting and Recovery Procedures” (page 116)
- “Reporting Problems” (page 139)

Troubleshooting Overview

This section describes LVM troubleshooting areas and tools available for troubleshooting LVM problems.

Information Collection

You can collect information about your LVM configuration using the `vgdisplay`, `lvdisplay`, `pvdisplay`, and `lvolnboot` commands. As noted in “[Planning for Recovery](#)” (page 34), periodically collect the outputs from the commands listed in [Table 9](#).

Table 9 LVM Information to Collect and Maintain

Command	Scope	Purpose
<code>ioscan -f</code>		Prints I/O configuration
<code>lvolnboot -v</code>		Prints information on root, boot, swap, and dump logical volumes.
<code>vgcfgrestore -l</code>	All volume groups	Prints volume group configuration from the backup file.
<code>vgdisplay -v</code>	All volume groups	Prints volume group information, including status of logical volumes and physical volumes.
<code>lvdisplay -v</code>	All logical volumes	Prints logical volume information, including mapping and status of logical extents.
<code>pvdisplay -v</code>	All physical volumes	Prints physical volume information, including status of physical extents.
<code>ioscan -m lun</code>		Print I/O configuration listing the hardware path to the disk, LUN instance, LUN hardware path and lunpath hardware path to the disk.

In addition, use the `lvmadm` command for two purposes:

- To determine which volume group versions are supported by your release of HP-UX 11i Version 3. For example, if your release supports Version 2.1 volume groups, `lvmadm` displays the following:

```
# lvmadm -t -v 2.1
--- LVM Limits ---
VG Version          2.1
Max VG Size (Tbytes) 2048
Max LV Size (Tbytes) 256
```

Max PV Size (Tbytes)	16
Max VGs	2048
Max LVs	2047
Max PVs	2048
Max Mirrors	5
Max Stripes	511
Max Stripe Size (Kbytes)	262144
Max LXs per LV	33554432
Max PXs per PV	16777216
Max Extent Size (Mbytes)	256

If your release does *not* support Version 2.1 volume groups, it displays the following:

```
# lvmadm -t -V 2.1
Error: 2.1 is an invalid volume group version.
```

- To display the contents of the /etc/lvmtab and /etc/lvmtab_p files in a human-readable fashion. For example, the following command displays the contents of the LVM configuration files for all Version 1.0 volume groups on your system:

```
# lvmadm -l -V 1.0
--- Version 1.0 volume groups ---
VG Name /dev/vg00
PV Name /dev/disk/disk34_p2
```

In addition, there are some tools available from your HP support representative:

- dump_lvmtab: prints the contents of the /etc/lvmtab file in human-readable fashion.
- vfcfgdisplay: prints the contents of an LVM volume group configuration backup file (as created by vfcfgbackup), such as the volume group information, the logical volume information, physical volume information and logical extent distribution.

Consistency Checks

Most LVM commands perform consistency checking. You can inspect your LVM configuration with the vgdisplay, lvdisplay, and pvdisplay commands, and look for inconsistencies.

In addition, the pvck command performs explicit consistency checking on a physical volume. This command detects bad checksums caused by a forward system migration after a backward system migration. Run pvck only on deactivated volume groups. For more information, see *pvck(1M)*.

NOTE: The pvck command does not support Version 2.x volume groups.

Maintenance Mode Boot

LVM **maintenance mode boot** is a special way to boot your system that bypasses the normal LVM structures. Use it only for problems that prevent the system from otherwise booting. It is similar to single-user state in that many of the processes that normally start do not start, and many of the system checks are not performed. LVM maintenance mode is intended to enable you to boot your system long enough to repair damage to the system LVM data structures typically using vfcfgrestore, which then enables you to boot your system normally.

Normally, the boot loader uses the LABEL file in the LIF volume to determine the location of the boot file system and the kernel /stand/vmunix. The LABEL file also contains the starting block and size of the root file system.

Under a maintenance mode boot, the boot loader attempts to find the boot file system at the start of the boot disk's user data area rather than using information from the LIF volume. To obtain the root file system's starting block and size, the boot loader reads the file /stand/rootconf. Since LVM is not enabled, the root file system must be allocated contiguously.

A maintenance mode boot differs from a standard boot as follows:

- The system is booted in single-user mode.
- No volume groups are activated.
- Primary swap and dump are not available.
- Only the root file system and boot file system are available.
- If the root file system is mirrored, only one copy is used. Changes to the root file system are not propagated to the mirror copies, but those mirror copies are marked stale and will be synchronized when the system boots normally.

To boot in maintenance mode on a system with a root disk configured with LVM, use the `-1m` option to the boot loader. On an HP 9000 server, enter the following command:

```
ISL> hpx -1m
```

On an HP Integrity server, enter the following command:

```
HPUX> boot -1m
```

△ CAUTION: When you boot your system in maintenance mode, do not activate the root volume group and do not change to multiuser mode (for example, by specifying `/sbin/init 2`). Doing so can corrupt the root file system.

When you have repaired or restored the LVM configuration information, reboot the system using the following command:

```
# /sbin/reboot
```

For more information about LVM maintenance mode boots and troubleshooting problems with LVM structures, see *Disk and File Management Tasks on HP-UX*, published by Prentice Hall PTR, 1997.

I/O Errors

When a device driver returns an error to LVM on an I/O request, LVM classifies the error as either **recoverable** or **nonrecoverable**.

Recoverable Errors

When LVM encounters a recoverable (correctable) error, it internally retries the failed operation assuming that the error will correct itself or that you can take steps to correct it. Examples of recoverable errors are the following:

- Device power failure
- A disk that goes missing *after the volume group is activated*
- A loose disk cable (which looks like a missing disk)

In these cases, LVM logs an error message to the console, but it does not return an error to the application accessing the logical volume.

If you have a current copy of the data on a separate, functioning mirror, then LVM directs the I/O to a mirror copy, the same as for a nonrecoverable error. Applications accessing the logical volume do not detect any error. (To preserve data synchronization between its mirrors, LVM retries recoverable write requests to a problematic disk, even if a current copy exists elsewhere. However, this process is managed by a daemon internal to LVM and has no impact on user access to the logical volume.)

However, if the device in question holds the only copy of the data, LVM retries the I/O request until it succeeds—that is, until the device responds or the system is rebooted. Any application performing I/O to the logical volume might block, waiting for the device to recover. In this case, your application or file system might appear to be stalled and might be unresponsive.

Temporarily Unavailable Device

By default, LVM retries I/O requests with recoverable errors until they succeed or the system is rebooted. Therefore, if an application or file system stalls, your troubleshooting must include checking the console log for problems with your disk drives and taking action to restore the failing devices to service.

Permanently Unavailable Device

If retrying the I/O request never succeeds (for example, the disk was physically removed), your application or file system might block indefinitely. If your application is not responding, you might need to reboot your system.

As an alternative to rebooting, you can control how long LVM retries a recoverable error before treating it as nonrecoverable by setting a timeout on the logical volume. If the device fails to respond within that time, LVM returns an I/O error to the caller. This timeout value is subject to any underlying physical volume timeout and driver timeout, so LVM can return the I/O error seconds after the logical volume timeout expired.

The timeout value is normally zero, which is interpreted as an infinite timeout. Thus, no I/O request returns to the caller until it completes successfully.

View the timeout value for a logical volume using the `lvdisplay` command, as follows:

```
# lvdisplay /dev/vg00/lvol1 | grep Timeout  
IO Timeout (Seconds)           default
```

Set the timeout value using the `-t` option of the `lvchange` command. This sets the timeout value in seconds for a logical volume. For example, to set the timeout for `/dev/vg01/lvol1` to one minute, enter the following command:

```
# lvchange -t 60 /dev/vg01/lvol1
```

CAUTION: Setting a timeout on a logical volume increases the likelihood of transient errors being treated as nonrecoverable errors, so any application that reads or writes to the logical volume can experience I/O errors. If your application is not prepared to handle such errors, keep the default infinite logical volume timeout.

TIP: Set the logical volume timeout to an integral multiple of any timeout assigned to the underlying physical volumes. Otherwise, the actual duration of the I/O request can exceed the logical volume timeout. For details on how to change the I/O timeout value on a physical volume, see `pvchange(1M)`.

Nonrecoverable Errors

Nonrecoverable errors are considered fatal; there is no expectation that retrying the operation will work.

If you have a current copy of the data on a separate, functioning mirror, then LVM directs reads and writes to that mirror copy. The I/O operation for the application accessing the logical volume completes successfully.

However, if you have no other copies of the data, then LVM returns an error to the subsystem accessing the logical volume. Thus, any application directly accessing a logical volume must be prepared for I/O requests to fail. File systems such as VxFS and most database applications are designed to recover from error situations; for example, if VxFS encounters an I/O error, it might disable access to a file system or a subset of the files in it.

LVM considers the following two situations nonrecoverable.

Media Errors

If an I/O request fails because of a media error, LVM typically prints a message to the console log file (`/var/adm/syslog/syslog.log`) when the error occurs. In the event of a media error, you must replace the disk (see “[Disk Troubleshooting and Recovery Procedures](#)” (page 116)).

If your disk hardware supports automatic bad block relocation (usually known as hardware sparing), enable it, because it minimizes media errors seen by LVM.

NOTE: LVM does not perform software relocation of bad blocks. It recognizes and honors software relocation entries created by previous releases of LVM but does not create new ones. Enabling or disabling bad block relocation using `lvchange` has no effect.

Missing Device When the Volume Group Was Activated

If the device associated with the I/O was not present *when the volume group was activated*, LVM prints an error message to the user's terminal at activation time. You must either locate the disk and restore it to service, or replace it, then activate the volume group again.

Volume Group Activation Failures

Normally, volume groups are automatically activated during system startup. Unless you intentionally deactivate a volume group using `vgchange`, you do not need to manually activate volume groups. In all cases, LVM requires that a quorum of disks in a volume group be available.

A **quorum** is the required number of physical volumes that must be available in a volume group to activate that volume group or for it to remain activated. To activate a volume group, *more than half* its disks that were available during the last activation must be online and in service. For the volume group to remain fully operational, *at least half* the disks must remain present and available.

During run time, when a volume group is already active, if a disk fails or is taken offline, the quorum might become lost. This condition occurs if less than half of the physical volumes defined for the volume group now remain fully operational. For example, if two disks belong to a volume group, the loss of one does not cause a loss of quorum as is the case when activating the volume group. To lose quorum, both disks must become unavailable. In that case, your volume group remains active, but a message prints to the console, indicating that the volume group has lost quorum. Until the quorum is restored (at least one of the LVM disks in the volume group in the previous example is again available), LVM does not allow you to complete most commands that affect the volume group configuration. Further, some of the I/O accesses to the logical volumes for that volume group might hang because the underlying disks are not accessible. Also, until quorum is restored, the MWC will not be updated because LVM cannot guarantee the consistency (integrity) of the LVM information.

Use the `vgchange -q n` option to override the system's quorum check when the volume group is activated. This option has no effect on the runtime quorum check. Overriding quorum can result in a volume group with an inaccurate configuration (for example, missing recently creating logical volumes). This configuration change might not be reversible.

There are ways to override quorum requirements at volume group activation time or boot time. Even when allowed by LVM, HP recommends that you do not make changes to the LVM configuration for active volume groups that do not have a quorum of disks present. To correct quorum issues, HP recommends returning the unavailable disks to service.

Quorum Problems with a Nonroot Volume Group

If you attempt to activate a nonroot volume group when not enough disks are present to establish a quorum, `vgchange` displays error messages similar to the following:

```
# vgchange -a y /dev/vg01
vgchange: Warning: Couldn't attach to the volume group
                  physical volume "/dev/dsk/c1t0d2":
```

The path of the physical volume refers to a device that does not exist,

or is not configured into the kernel.
vgchange: Couldn't activate volume group "/dev/vg01":
Either no physical volumes are attached or no valid VGAs were found on
the physical volumes.

If a nonroot volume group does not activate because of a failure to meet quorum, follow these steps:

1. Check the power and data connections (including Fibre Channel zoning and security) of all the disks that are part of the volume group that you cannot activate. Return all disks (or at least enough to make a quorum) to service. Then use the vgchange command to activate the volume group again.
2. If there is no other way to make a quorum available, use the -q option with the vgchange command to override the quorum requirement.

```
# vgchange -a y -q n /dev/vg01
```

The volume group activates without a quorum. You might get messages about not being able to access certain logical volumes because part or all of a logical volume might be located on one of the disks that is not present.

Whenever you override a quorum requirement, you run the risk of using data that is not current. Be sure to check the data on the logical volumes in the activated volume group and the size and locations of the logical volumes to ensure that they are up to date.

Return the disabled disks to the volume group as soon as possible. When you return a disk to service that was not online when you originally activated the volume group, use the vgchange command as follows:

```
# vgchange -a y /dev/vg01
```

Quorum Problems with a Root Volume Group

Your root volume group can also report a quorum problem. If not enough disks are present in the root volume group to constitute a quorum, a message indicating that not enough physical volumes are present appears during the boot sequence. This error might occur if you have physically removed a disk from your system because you no longer intended to use it, but did not remove the physical volume from the volume group using vgreduce. Do not remove an LVM disk from a system without first removing it from its volume group. However, you can try to recover by booting the system with the quorum override option -lq.

On an HP 9000 server, enter the following command:

```
ISL> hpxx -lq
```

On an HP Integrity server, enter the following command:

```
HPUX> boot -lq
```

Version 2.x Volume Group Activation Failures

Version 2.x volume groups can fail to activate because of insufficient quorum. For example, vgchange could display error messages similar to the following:

```
# vgchange -a y /dev/vgtest
vgchange: Warning: Couldn't attach to the volume group physical
volume "/dev/disk/disk1":
I/O error
vgchange: I/O error
vgchange: Couldn't activate volume group "/dev/vgtest":
Quorum not present, or some physical volume(s) are missing.
```

If a Version 2.x volume group does not activate because of a failure to meet quorum, follow the steps described in “[Quorum Problems with a Nonroot Volume Group](#)” (page 111).

A Version 2.x volume group can also fail activation if the necessary commands or kernel drivers are not present. For example, vgchange could display the following error message:

```
# vgchange -a y /dev/vgtest  
vgchange: Error: The "lvmp" driver is not loaded.
```

Here is another possible error message:

```
# vgchange -a y /dev/vgtest  
vgchange: Warning: Couldn't attach to the volume group physical  
volume "/dev/disk/disk1":  
Illegal byte sequence  
vgchange: Couldn't activate volume group "/dev/vgtest":  
Quorum not present, or some physical volume(s) are missing.
```

A third possible error message is the following:

```
# vgchange -a y /dev/vgtest  
vgchange: Warning: Couldn't attach to the volume group  
physical volume "/dev/dsk/c1t0d0":  
Cross-device link  
vgchange: Warning: couldn't query physical volume "/dev/dsk/c1t0d0":  
The specified path does not correspond to physical volume  
attached to this volume group  
vgchange: Couldn't query the list of physical volumes.  
vgchange: Couldn't activate volume group "/dev/vgtest":  
Quorum not present, or some physical volume(s) are missing.
```

To troubleshoot these activation failures, run the lvmadm command as follows:

```
# lvmadm -t -V 2.0
```

If the necessary commands and drivers are present, lvmadm displays the following:

```
--- LVM Limits ---  
VG Version          2.0  
Max VG Size (Tbytes) 2048  
Max LV Size (Tbytes) 256  
Max PV Size (Tbytes) 16  
Max VGs             512  
Max LVs             511  
Max PVs             511  
Max Mirrors         6  
Max Stripes         511  
Max Stripe Size (Kbytes) 262144  
Max LXs per LV      33554432  
Max PXs per PV       16777216  
Max Extent Size (Mbytes) 256
```

If lvmadm displays no output, your operating system release does not support Version 2.x volumes. You must update your system to the March 2008 release of HP-UX 11i Version 3 or a newer release.

If the kernel driver to support Version 2.x volume groups is not loaded, lvmadm displays the following error:

```
# lvmadm -t -V 2.0  
lvmadm: Error: The "lvmp" driver is not loaded.
```

Load the lvmp module using the kmodule command as follows:

```
# kmodule lvmp=best  
==> Update the automatic 'backup' configuration first? n  
* Future operations will ask whether to update the backup.  
* The requested changes have been applied to the currently  
running configuration.  
Module           State   Cause      Notes  
lvmp  (before)    unused     loadable, unloadable  
      (now)        loaded    best  
      (next boot) loaded    explicit
```

You do not need to reboot. After you load the lvmp module, lvmadm succeeds:

```
# lvmadm -t -V 2.0  
--- LVM Limits ---  
VG Version          2.0
```

Max VG Size (Tbytes)	2048
Max LV Size (Tbytes)	256
Max PV Size (Tbytes)	16
Max VGs	512
Max LVs	511
Max PVs	511
Max Mirrors	6
Max Stripes	511
Max Stripe Size (Kbytes)	262144
Max LXs per LV	33554432
Max PXs per PV	16777216
Max Extent Size (Mbytes)	256

TIP: If your system has no Version 2.x volume groups, you can free up system resources associated with `lvmp` by unloading it from the kernel. Run the `kcmodule` command as follows:

```
# kcmodule lvmp=unused
    ==> Update the automatic 'backup' configuration first? n
        * Future operations will ask whether to update the backup.
        * The requested changes have been applied to the currently
            running configuration.

Module          State   Cause      Notes
lvmp    (before)  loaded  explicit loadable, unloadable
                (now)    unused
```

If you later want to create Version 2.x volume groups, load the `lvmp` driver as described previously.

Root Volume Group Scanning

If the LVM subsystem detects that vital information is corrupted on the boot disk, it scans all the attached devices to try to find the physical volumes that are part of the root volume group. You then see the following messages on the system console and in `/var/adm/syslog/syslog.log`:

```
LVM : Failure in attaching PV (dev=0x10000nn) to the root volume group.
The physical volume does not belong to the root volume group
LVM : Failure in attaching PV (dev=0x10000nn) to the root volume group.
The physical volume does not belong to the root volume group
LVM : Activation of root volume group failed
Quorum not present, or some physical volume(s) are missing
LVM: Scanning for Root VG PVs (VGID 0xnnnnnnnnn 0xnnnnnnnnn)
```

If this root volume group scanning succeeds, messages similar to the following appear:

```
LVM: Rootvgscan detected 10 PV(s). Will attempt root VG activation
using the following PV(s):
    0x100005f 0x1000060 0x1000061 0x1000062 0x1000063 0x1000064
    0x1000065 0x1000067 0x1000068 0x100006e
LVM: WARNING: Root VG activation required a scan. The PV information in
the on-disk BDRA may be out-of-date from the system's current IO
configuration. To update the on-disk BDRA, first update /etc/lvmtab
using vgscan(1M), then update the on-disk BDRA using lvlnboot(1M) .
For example, if the root VG name is /dev/vg00:
    1. vgscan -k -f /dev/vg00
    2. lvlnboot -R /dev/vg00
```

LVM: Root VG activated

If this root volume group scan fails to find all physical volumes, the following message appears:

```
LVM: WARNING: Rootvgscan did not find any PV(s) matching root VGID.
Will attempt root VG activation using the boot device (0x10000nn) .
```

Or the following message:

```
LVM: WARNING: BDRA lists the number of PV(s) for the root VG as nn,
but rootvgscan found only nn. Proceeding with root VG activation.
```

LVM Boot Failures

There are several reasons why an LVM configuration cannot boot. In addition to the problems associated with boots from non-LVM disks, the following problems can cause an LVM-based system not to boot.

Insufficient Quorum

In this scenario, not enough disks are present in the root volume group to meet the [quorum](#) requirements. At boot time, a message indicating that not enough physical volumes are available appears:

```
panic: LVM: Configuration failure
```

To activate the root volume group and successfully boot the system, the number of available LVM disks must be more than half the number of LVM disks that were attached when the volume group was last active. Thus, if during the last activation there were two disks attached in the root volume group, the “more than half” requirement means that both must be available. For information on how to deal with quorum failures, see [“Volume Group Activation Failures” \(page 111\)](#).

Corrupted LVM Data Structures on Disk

The LVM bootable disks contain vital boot information in the BDRA. This information can become corrupted, not current, or just no longer present. Because of the importance of maintaining up-to-date information within the BDRA, use the `lvrmboot` or `lvlnboot` commands whenever you make a change that affects the location of the root, boot, primary swap, or dump logical volumes.

To correct this problem, boot the system in maintenance mode as described in [“Maintenance Mode Boot” \(page 108\)](#), then repair the damage to the system LVM data structures. Use `vgcfgrestore` on the boot disk.

Corrupted LVM Configuration File

Another problem with activation of a volume group is a missing or corrupted `/etc/lvmtab` or `/etc/lvmtab_p` file. After booting in maintenance mode, you can use the `vgscan` command to re-create the `/etc/lvmtab` and `/etc/lvmtab_p` files. For more information, see `vgscan(1M)`.

Problems After Reducing the Size of a Logical Volume

When a file system is first created within a logical volume, it is made as large as the logical volume permits.

If you extend the logical volume without extending its file system, you can subsequently safely reduce the logical volume size, as long as it remains as big as its file system. (Use the `bdf` command to determine the size of your file system.) After you expand the file system, you can no longer safely reduce the size of the associated logical volume.

If you reduce the size of a logical volume containing a file system to a size smaller than that of a file system within it using the `lvreduce` command, you corrupt the file system. If you subsequently attempt to mount the corrupt file system, you might crash the system. If this occurs, follow these steps:

1. Reboot your system in single-user mode.
2. If you already have a good current backup of the data in the now corrupt file system, skip this step. If you do not have backup data and if that data is critical, try to recover whatever part of the data that might remain intact by attempting to back up the files on the file system.

Before you attempt any current backup, consider the following:

- When your backup program accesses the corrupt part of the file system, your system will crash again. You must reboot your system again to continue with the next step.
- There is no guarantee that all (or any) of your data on that file system will be intact or recoverable. This step is an attempt to save as much as possible. That is, any data

successfully backed up in this step will be recoverable, but some or all of your data might not be successfully backed up because of file corruption.

3. Immediately unmount the corrupted file system if it is mounted.
4. Use the logical volume for swap space or raw data storage, or use HP SMH or the newfs command to create a new file system in the logical volume. This new file system now matches the current reduced size of the logical volume.
5. If you have created a new file system on the logical volume, do one of the following:
 - If you have a good prior backup (not the backup from step 2), restore its contents. Because the new file system in the smaller logical volume is smaller than the original file system, you might not have enough space to restore all your original files.
 - If you do not have a good prior backup, attempt to restore as many files as possible from any backup you made in step 2.
 - Use the new file system for creating and storing a new set of files (not for trying to restore the original files).

Disk Troubleshooting and Recovery Procedures

This section provides step-by-step procedures to handle disk failure planning, troubleshooting, and recovery, as follows:

1. ["Step 1: Preparing for Disk Recovery" \(page 116\)](#)
2. ["Step 2: Recognizing a Failing Disk" \(page 118\)](#)
3. ["Step 3: Confirming Disk Failure" \(page 120\)](#)
4. ["Step 4: Determining Action for Disk Removal or Replacement" \(page 123\)](#)
5. ["Step 5: Removing a Bad Disk" \(page 126\)](#)
6. ["Step 6: Replacing a Bad Disk \(Persistent DSFs\)" \(page 129\)](#)
7. ["Step 7: Replacing a Bad Disk \(Legacy DSFs\)" \(page 138\)](#)

Step 1: Preparing for Disk Recovery

Recovering from system failure is a critical part of system administration. Consider the following guidelines before you experience a disk failure.

Defining a Recovery Strategy

As you create logical volumes, choose one of the following recovery strategies. Each choice strikes a balance between cost, data availability, and speed of data recovery.

- **Mirroring:** If you mirror a logical volume on a separate disk, the mirror copy is online and available while recovering from a disk failure. With hot-swappable disks, users will have no indication that a disk was lost.
- **Restoring from backup:** If you choose not to mirror, make sure you have a consistent backup plan for any important logical volumes. The tradeoff is that you will need fewer disks, but you will lose time while you restore data from backup media, and you will lose any data changed since your last backup.
- **Initializing from scratch:** If you do not mirror or back up a logical volume, be aware that you will lose data if the underlying hard disk fails. This can be acceptable in some cases, such as a temporary or scratch volume.

Using LVM Online Disk Replacement (LVM OLR)

LVM online disk replacement (LVM OLR) simplifies the replacement of disks under LVM. You can temporarily disable LVM use of a disk in an active volume group without deactivating the volume group or removing the logical volumes on the disk.

The LVM OLR feature uses a new option (-a) in the `pvchange` command. The -a option disables or re-enables a specified path to an LVM disk, as used to halt LVM access to the disk under “[Step 6: Replacing a Bad Disk \(Persistent DSFs\)](#)” ([page 129](#)) or “[Step 7: Replacing a Bad Disk \(Legacy DSFs\)](#)” ([page 138](#)). For more information, see the `pvchange(1M)` manpage.

Mirroring Critical Information, Especially the Root Volume Group

By making mirror copies of the root, boot, and primary swap logical volumes on another disk, you can use the copies to keep your system in operation if any of these logical volumes fail.

For details on mirroring the boot disk (which contains the root, boot, and primary swap logical volumes), see “[Mirroring the Boot Disk](#)” ([page 90](#)).

There are three corollaries to the mirroring recommendation:

1. Use the strict allocation policy for all mirrored logical volumes. Strict allocation forces mirrors to occupy different disks. Without strict allocation, you can have multiple mirror copies on the same disk; if that disk fails, you will lose all your copies. To control the allocation policy, use the -s option with the `lvcreate` and `lvchange` commands. By default, strict allocation is enabled.
2. To improve the availability of your system, keep mirror copies of logical volumes on separate I/O busses if possible. With multiple mirror copies on the same bus, the bus controller becomes a single point of failure—if the controller fails. If you create physical volume groups and set the allocation policy to *PVG-strict*, LVM helps you avoid inadvertently creating multiple mirror copies on a single bus. For more information about physical volume groups, see `lvm_pv(4)`.
3. Consider using one or more free disks within each volume group as spares. If you configure a disk as a spare, then a disk failure causes LVM to reconfigure the volume group so that the spare disk takes place of the failed one. That is, all the logical volumes that were mirrored on the failed disk are automatically mirrored and resynchronized on the spare, while the logical volume remains available to users. You can then schedule the replacement of the failed disk at a time of minimal inconvenience to you and your users. Sparing is particularly useful for maintaining data redundancy when your disks are not hot-swappable, since the replacement process may have to wait until your next scheduled maintenance interval. For information about disk sparing, see “[Increasing Disk Redundancy Through Disk Sparing](#)” ([page 27](#))

NOTE: Version 2.x volume groups do not support disk sparing.

Step 2: Recognizing a Failing Disk

This section explains how to look for signs that one of your disks is having problems, and how to determine which disk it is.

I/O Errors in the System Log

Often an error message in the system log file, `/var/adm/syslog/syslog.log`, is your first indication of a disk problem. You might see the following error

```
Asynchronous write failed on LUN (dev=0x3000015)
IO details : blkno : 2345, sector no : 23
```

See the system log errors fully described in ["Matching Error Messages to Physical Disks and Volume Groups" \(page 163\)](#), where it shows how you map this type of error message to a specific disk.

Disk Failure Notification Messages from Diagnostics

If you have Event Monitoring Service (EMS) hardware monitors installed on your system, and you enabled the disk monitor `disk_em`, a failing disk can trigger an event to the (EMS). Depending on how you configured EMS, you might get an email message, a message in `/var/adm/syslog/syslog.log`, or messages in another log file. EMS error messages identify a hardware problem, what caused it, and what must be done to correct it. The following example is part of an error message:

```
Event Time.....: Tue Oct 26 14:06:00 2004
Severity.....: CRITICAL
Monitor.....: disk_em
Event #.....: 18
System.....: myhost

Summary:
Disk at hardware path 0/2/1/0.2.0 : Drive is not responding.

Description of Error:
The hardware did not respond to the request by the driver.
The I/O request was not completed.

Probable Cause / Recommended Action:
The I/O request that the monitor made to this device failed because the
device timed-out. Check cables, power supply, ensure the drive is powered ON,
and if needed contact your HP support representative to check the drive.
```

For more information on EMS, see the *Diagnostics section* on the <http://docs.hp.com> website

LVM Command Errors

Sometimes LVM commands, such as `vgdisplay`, return an error suggesting that a disk has problems. For example:

```
#vgdisplay -v | more
...
--- Physical volumes ---
PV Name          /dev/dsk/c0t3d0
PV Status        unavailable
Total PE         1023
Free PE          173
...
```

The physical volume status of `unavailable` indicates that LVM is having problems with the disk. You can get the same status information from `pvdisplay`.

The next two examples are warnings from `vgdisplay` and `vgchange` indicating that LVM has no contact with a disk:

```
#vgdisplay -v vg
vgdisplay: Warning: couldn't query physical volume "/dev/dsk/c0t3d0":
The specified path does not correspond to physical volume attached to
```

```
this volume group
vgdisplay: Warning: couldn't query all of the physical volumes.
#vgchange -a y /dev/vg01
vgchange: Warning: Couldn't attach to the volume group physical volume
"/dev/dsk/c0t3d0":
A component of the path of the physical volume does not exist.
Volume group "/dev/vg01" has been successfully changed.
```

Another sign of disk problem is seeing stale extents in the `lvdisplay` command output. If there are stale extents on a logical volume even after running the `vgsync` or `lvsync` commands, you might have an issue with an I/O path or one of the disks used by the logical volume, but not necessarily the disk showing stale extents. For example:

```
# lvdisplay -v /dev/vg01/lvol3 | more
...
LV Status           available/stale
...
--- Logical extents ---
LE   PV1          PE1  Status 1  PV2          PE2  Status 2
0000 /dev/dsk/c0t3d0  0000 current  /dev/dsk/c1t3d0  0100 current
0001 /dev/dsk/c0t3d0  0001 current  /dev/dsk/c1t3d0  0101 current
0002 /dev/dsk/c0t3d0  0002 current  /dev/dsk/c1t3d0  0102 stale
0003 /dev/dsk/c0t3d0  0003 current  /dev/dsk/c1t3d0  0103 stale
...
```

All LVM error messages indicates which device file is associated with the problematic disk. This is useful for the next step, confirming disk failure.

See “[Messages For All LVM Commands](#)” (page 164) for more information on messages from LVM commands, their causes and recommended actions.

Step 3: Confirming Disk Failure

Once you suspect a disk has failed or is failing, make certain that the suspect disk is indeed failing. Replacing or removing the incorrect disk makes the recovery process take longer. It can even cause data loss. For example, in a mirrored configuration, if you were to replace the wrong disk—the one holding the current good copy rather than the failing disk—the mirrored data on the good disk is lost.

It is also possible that the suspect disk is not failing. What seems to be a disk failure might be a hardware path failure; that is, the I/O card or cable might have failed. If a disk has multiple hardware paths, also known as pvlinks, one path can fail while an alternate path continues to work. For such disks, try the following steps on all paths to the disk.

If you have isolated a suspect disk, you can use hardware diagnostic tools, like Support Tools Manager, to get detailed information about it. Use these tools as your first approach to confirm disk failure. They are documented on <http://docs.hp.com> in the *Diagnostics* area. If you do not have diagnostic tools available, follow these steps to confirm that a disk has failed or is failing:

1. Use the `ioscan` command to check the S/W state of the disk. Only disks in state `CLAIMED` are currently accessible by the system. Disks in other states such as `NO_HW` or disks that are completely missing from the `ioscan` output are suspicious. If the disk is marked as `CLAIMED`, its controller is responding. For example:

```
# ioscan -fCdsk
Class I H/W Path     Driver  S/W State H/W Type   Description
=====
disk  0 8/4.5.0      sdisk   CLAIMED  DEVICE    SEAGATE ST34572WC
disk  1 8/4.8.0      sdisk   UNCLAIMED UNKNOWN  SEAGATE ST34572WC
disk  2 8/16/5.2.0   sdisk   CLAIMED  DEVICE    TOSHIBA CD-ROM XM-5401TA
```

In this example, the disk at hardware path 8/4.8.0 is not accessible.

If the disk has multiple hardware paths, be sure to check all the paths.

2. You can use the `pvdisplay` command to check whether the disk is attached or not. A physical volume is considered to be attached, if the `pvdisplay` command is able to report a valid status (unavailable/available) for it. Otherwise, the disk is unattached. In that case, the disk was defective or inaccessible at the time the volume group was activated. For example, if `/dev/dsk/c0t5d0` is a path to a physical volume that is attached to LVM, enter:

```
# pvdisplay /dev/dsk/c0t5d0 | grep "PV Status"
PV Status           available
```

If `/dev/dsk/c1t2d3` is a path to a physical volume that is detached from LVM access using a `pvchange -a n` or `pvchange -a N` command, enter:

```
# pvdisplay /dev/dsk/c1t2d3 | grep "PV Status"
PV Status           unavailable
```

If the disk responds to the `ioscan` command, test it with the `diskinfo` command. The reported size must be nonzero; otherwise, the device is not ready. For example:

```
# diskinfo /dev/rdsk/c0t5d0
SCSI describe of /dev/rdsk/c0t5d0:
        vendor: SEAGATE
        product id: ST34572WC
        type: direct access
        size: 0 Kbytes
        bytes per sector: 512
```

In this example the size is 0, so the disk is malfunctioning.

3. If both `ioscan` and `diskinfo` succeed, the disk might still be failing. As a final test, try to read from the disk using the `dd` command. Depending on the size of the disk, a comprehensive read can be time-consuming, so you might want to read only a portion of the disk. If the disk is functioning properly, no I/O errors are reported. The following example shows a successful read of the first 64 megabytes of the disk: When you enter the following command, look for the solid blinking green LED on the disk:

```
# dd if=/dev/rdsk/c0t5d0 of=/dev/null bs=1024k count=64
64+0 records in
64+0 records out
```

NOTE: If the dd command hangs or takes a long time, Ctrl+C stops the read on the disk. To run dd on the background, add & at the end of the command.

The following command shows an unsuccessful read of the whole disk:

```
# dd if=/dev/rdsk/c1t3d0 of=/dev/null bs=1024k
dd read error: I/O error
0+0 records in 0+0 records out
```

4. If the physical volume is attached but cannot be refreshed via an lvsync, it is likely there is a media problem at a specific location. Reading only the extents associated with the LE can help isolate the problem. Remember the stale extent might not have the problem. The lvsync command starts refreshing extents at LE zero and stops if it encounters an error. Therefore, find the first LE in any logical volume that is stale and test this one. For example:

- a. Find the first stale LE:

```
# lvdisplay -v /dev/vg01/lvol3 | more
LV Status available/stale
.
.
.
--- Logical extents ---
LE PV1 PE1 Status 1 PV2 PE2 Status 2
0000 /dev/dsk/c0t3d0 0000 current /dev/dsk/c1t3d0 0100 current
0001 /dev/dsk/c0t3d0 0001 current /dev/dsk/c1t3d0 0101 current
0002 /dev/dsk/c0t3d0 0002 current /dev/dsk/c1t3d0 0102 stale
0003 /dev/dsk/c0t3d0 0003 current /dev/dsk/c1t3d0 0103 stale
```

In this case, LE number 2 is stale.

- b. Get the extent size for the VG:

```
# vgdisplay /dev/vg01 | grep -I "PE Size"
PE size (Mbytes) 32
```

- c. Find the start of PE zero on each disk:

For a version 1.0 VG, enter:

```
xd -j 0x2048 -t uI -N 4 /dev/dsk/c0t3d0
```

For a version 2.x VG, enter:

```
xd -j 0x21a4 -t uI -N 4 /dev/dsk/c0t3d0
```

In this example, this is a version 1.0 VG.

```
# xd -j 0x2048 -t uI -N 4 /dev/dsk/c0t3d0
0000000 1024
0000004
# xd -j 0x2048 -t uI -N 4 /dev/dsk/c1t3d0
0000000 1024
0000004
```

- d. Calculate the location of the physical extent for each PV. Multiply the PE number by the PE size and then by 1024 to convert to Kb:

2 * 32 * 1024 = 65536

Add the offset to PE zero:

65536 + 1024 = 66560

- e. Enter the following dd commands:

```
# dd bs=1k skip=66560 count=32768 if=/dev/rdsk/c0t3d0 of=/dev/null
# dd bs=1k skip=66560 count=32768 if=/dev/rdsk/c1t3d0 of=/dev/null
```

Note the value calculated is used in the `skip` argument. The count is obtained by multiplying the PE size by 1024.

Step 4: Determining Action for Disk Removal or Replacement

Once you know which disk is failing, you can decide how to deal with it. You can choose to remove the disk if your system does not need it, or you can choose to replace it. Before deciding on your course of action, you must gather some information to help guide you through the recovery process.

- **Is the disk hot-swappable?**

You can remove or add an inactive hot-swappable hard disk drive module to a system while power is still on and the SCSI bus is still active. That is, you can replace or remove a hot-swappable disk from a system without turning off the power to the entire system.

See your system hardware manuals for information about which disks in your system are hot-swappable. Specifications for other hard disks are available in their installation manuals..

If your disk is not hot-swappable, you must schedule system down time to replace the disk.

- **Is the disk the root disk or part of the root volume group?**

If the root disk is failing, the replacement process includes steps to set up the boot area; in addition, you might have to boot from its mirror if the primary root disk has failed. If a failing root disk is not mirrored, you must reinstall to the replacement disk or recover it from an Ignite-UX backup.

To determine whether the disk is in the root volume group, use the `lvolnboot` command with the `-v` option. It lists the disks in the root volume group, and any special volumes configured on them. For example:

```
# lvolnboot -v
Boot Definitions for Volume Group /dev/vg00:
Physical Volumes belonging in Root Volume Group:
  /dev/disk/disk47_p2 -- Boot Disk
Boot: lvol1      on:      /dev/disk/disk47_p2
Root: lvol3      on:      /dev/disk/disk47_p2
Swap: lvol2      on:      /dev/disk/disk47_p2
Dump: lvol2      on:      /dev/disk/disk47_p2, 0
```

- **What logical volumes are on the disk, and are they mirrored?**

If you plan to replace the disk, you might need to restore data from backups. However, you must only recover data for a subset of the logical volumes in the volume group. Only the logical volumes that actually have physical extents on the disk are affected. In addition, if a logical volume is mirrored, there is probably a current copy of the data on the mirror, so it does not need to be recovered from backup.

You can find the list of logical volumes using the disk with the `pvdisplay` command. With the `-v` option, `pvdisplay` shows a listing of all the physical extents on a physical volume and to what logical volume they belong. This list is long; pipe it to `more` or send it to a file. For example:

```
# pvdisplay -v /dev/disk/disk3 | more
...
--- Distribution of physical volume ---
  LV Name           LE of LV  PE for LV
  /dev/vg00/lvol5    50       50
  /dev/vg00/lvol6   245      245
...
```

In this example, logical volumes `/dev/vg00/lvol5` and `/dev/vg00/lvol6` have physical extents on this disk, so you must restore `lvol5` and `lvol6` only.

If `pvdisplay` fails, you have the following options. You can refer to any configuration documentation you created in advance.

Alternatively, you can run `lvdisplay -v` on all the logical volumes in the volume groups and see if any extents are mapped to an unavailable physical volume. The `lvdisplay`

command shows '????' for the physical volume if it is unavailable. The issue with this approach is that it does not show precisely how many disks are unavailable. To ensure that multiple simultaneous disk failures have not occurred, run `vgdisplay` to check the difference between the number of *active* and number of *current* physical volumes. For example, a difference of one means only one disk is failing.

A third option for determining which logical volumes are on the disk is to use the `vgcfgdisplay` command, available from your HP support representative.

For each of the logical volumes affected, use `lvdisplay` to determine if the number of mirror copies is greater than zero. This verifies that the logical volume is mirrored. For example:

```
# lvdisplay /dev/vg00/lvol1
--- Logical volumes ---
LV Name          /dev/vg00/lvol1
VG Name          /dev/vg00
LV Permission    read/write
LV Status        available/syncd
Mirror copies   1
Consistency Recovery MWC
Schedule         parallel
LV Size (Mbytes) 300
Current LE      75
Allocated PE    150
Stripes          0
Stripe Size (Kbytes) 0
Bad block        off
Allocation       strict/contiguous
IO Timeout (Seconds) default
```

The number of mirror copies is not zero; therefore, the logical volume is mirrored.

Use `lvdisplay` again to determine which logical extents are mapped onto the suspect disk, and whether there is a current copy of that data *on another disk*. With the `-v` option, `lvdisplay` shows every logical extent, its mapping to any physical extents, and the status of those physical extents (stale or current).

This listing can be quite long, so use `grep` to confine the listing to the disk that is being replaced. For example:

```
# lvdisplay -v /dev/vg00/lvol1 | grep -e /dev/disk/disk3 -e '????'
 00000 /dev/disk/disk3  00000 current  /dev/disk/disk6  00000 current
 00001 /dev/disk/disk3  00001 current  /dev/disk/disk6  00001 current
 00002 /dev/disk/disk3  00002 current  /dev/disk/disk6  00002 current
 00003 /dev/disk/disk3  00003 current  /dev/disk/disk6  00003 current
 00004 /dev/disk/disk3  00004 current  /dev/disk/disk6  00004 current
 00005 /dev/disk/disk3  00005 current  /dev/disk/disk6  00005 current
...
```

In this example, all `lvol1` physical extents on `/dev/disk/disk3` have a current copy elsewhere on the system, specifically on `/dev/disk/disk6`. If `/dev/disk/disk3` was unavailable when the volume group was activated, its column contains a '????' instead of the disk name.

NOTE: There might be an instance where you see that only the failed physical volume holds the current copy of a given extent (and all other mirror copies of the logical volume hold the stale data for that given extent), and LVM does not permit you to remove that physical volume from the volume group. In this case, use the `lvunstale` command (available from your HP support representative) to mark one of the mirror copies as “nonstale” for that given extent.

HP recommends you use the `lvunstale` tool with caution.

Based on the gathered information, choose the appropriate disk removal or disk replacement procedure, detailed on the following pages. The possible scenarios are:

- [“Step 5: Removing a Bad Disk” \(page 126\)](#)
- [“Step 6: Replacing a Bad Disk \(Persistent DSFs\)” \(page 129\)](#)
- [“Step 7: Replacing a Bad Disk \(Legacy DSFs\)” \(page 138\)](#)

Step 5: Removing a Bad Disk

You can elect to remove the failing disk from the system instead of replacing it if you are certain that another valid copy of the data exists or the data can be moved to another disk.

Removing a Mirror Copy from a Disk

If you have a mirror copy of the data already, you can stop LVM from using the copy on the failing disk by reducing the number of mirrors. To remove the mirror copy from a specific disk, use `lvreduce`, and specify the disk from which to remove the mirror copy.

For example (if you have a single mirror copy):

```
# lvreduce -m 0 -A n /dev/vgname/lvname bad_disk_path
```

Or, if you have two mirror copies:

```
# lvreduce -m 1 -A n /dev/vgname/lvname bad_disk_path
```

The `-A n` option is used to prevent the `lvreduce` command from performing an automatic `vgcfgbackup` operation, which might hang while accessing a defective disk.

If you have only a single mirror copy and want to maintain redundancy, as soon as possible, create a second mirror of the data on a different, functional disk, subject to the mirroring guidelines, described in “[Step 1: Preparing for Disk Recovery](#)” (page 116), before you run `lvreduce`.

Removing Mirror Copy from Ghost Disk

One might encounter a situation where you have to remove from the volume group a failed physical volume or a physical volume that is not actually connected to the system but is still recorded in the LVM configuration file. Such a physical volume is sometimes called a *ghost disk* or *phantom disk*. You can get a ghost disk if the disk has failed before volume group activation, possibly because the system was rebooted after the failure.

A ghost disk is usually indicated by `vgdisplay` reporting more current physical volumes than active ones. Additionally, LVM commands might complain about the missing physical volumes as follows:

```
# vgdisplay vg01
vgdisplay: Warning: couldn't query physical volume "/dev/dsk/c5t5d5":
The specified path does not correspond to physical volume attached
to this volume group
vgdisplay: Couldn't query the list of physical volumes.
--- Volume groups ---
VG Name                  /dev/vg01
VG Write Access          read/write
VG Status                available
Max LV                  255
Cur LV                  3
Open LV                 3
Max PV                  16
Cur PV                  2  (#No. of PVs belonging to vg01)
Act PV                  1  (#No. of PVs recorded in the kernel)
Max PE per PV            4350
VGDA                     2
PE Size (Mbytes)         8
Total PE                 4341
Alloc PE                 4340
Free PE                  1
Total PVG                0
Total Spare PVs           0
Total Spare PVs in use   0
```

In these situations where the disk was not available at boot time, or the disk has failed before volume group activation (`pvdisplay` failed), the `lvreduce` command fails with an error that it could not query the physical volume. You can still remove the mirror copy, but you must specify the physical volume **key** rather than the name.

The physical volume key of a disk indicates its order in the volume group. The first physical volume has the key 0, the second has the key 1, and so on. This need not be the order of appearance in /etc/lvmtab file although it is usually the case, at least when a volume group is initially created.

You can use the physical volume key to address a physical volume that is not attached to the volume group. This usually happens if it was not accessible during activation, for example, because of a hardware or configuration problem. You can obtain the key using lvdisplay with the -k option as follows:

```
# lvdisplay -v -k /dev/vg00/lvol1
...
--- Logical extents ---
  LE   PV1          PE1   Status 1 PV2          PE2   Status 2
  00000  0           00000 stale      1           00000 current
  00001  0           00001 stale      1           00001 current
  00002  0           00002 stale      1           00002 current
  00003  0           00003 stale      1           00003 current
  00004  0           00004 stale      1           00004 current
  00005  0           00005 stale      1           00005 current
...
```

Compare this output with the output of lvdisplay without -k, that was used to check the mirror status. The column that contained the failing disk (or '????') now holds the key. For this example, the key is 0. Use this key with lvreduce. For example, if you have a single mirror copy:

```
# lvreduce -m 0 -A n -k /dev/vgname/lvname key
```

Or, if you have two mirror copies:

```
# lvreduce -m 1 -A n -k /dev/vgname/lvname key
```

Moving the Physical Extents to Another Disk

If the disk is partially available and you can still read from it, you can move the data onto another disk by moving the physical extents onto another disk. The pvmove command moves logical volumes or certain extents of a logical volume from one physical volume to another. It is typically used to free up a disk; that is, to move all data from that physical volume so it can be removed from the volume group. In its simplest invocation, you specify the disk to free up, and LVM moves all the physical extents on that disk to any other disks in the volume group, subject to any mirroring allocation policies. For example:

```
# pvmove pvname
```

NOTE: The pvmove command will fail if the logical volume is striped.

The pvmove command provides many features to move data from a disk. See “[Moving Data to a Different Physical Volume](#)” (page 73) for more information and pvmove examples. The pvmove(1M) manpage describes all the command’s features and options.

Removing the Disk from the Volume Group

After the disk no longer holds any logical extents, you can use the vgreduce command to remove the physical volume from the volume group so it is not inadvertently used again. Check for alternate links before removing the disk, since you must remove all the paths to a multipathed disk. Use the pvdisplay command as follows:

```
# pvdisplay /dev/dsk/c0t5d0
--- Physical volumes ---
PV Name          /dev/dsk/c0t5d0
PV Name          /dev/dsk/c1t6d0  Alternate Link
VG Name          /dev/vg01
PV Status        available
Allocatable      yes
VGDA            2
Cur LV           0
PE Size (Mbytes) 4
```

Total PE	1023
Free PE	1023
Allocated PE	0
Stale PE	0
IO Timeout (Seconds)	default
Autoswitch	On

In this example, there are two entries for PV Name. Use the `vgreduce` command to reduce each path as follows:

```
# vgreduce vgname /dev/dsk/c0t5d0
# vgreduce vgname /dev/dsk/c1t6d0
```

If the disk is unavailable, the `vgreduce` command fails. You can still forcibly reduce it, but you must then rebuild the `lvmtab`, which has two side effects. First, any deactivated volume groups are left out of the `lvmtab`, so you must manually `vgimport` them later. Second, if any multipathed disks have their link order reset, and if you arranged your pvlinks to implement load-balancing, you might have to arrange them again.

Starting with the HP-UX 11i v3 release, there is a new feature introduced in the mass storage subsystem that also supports multiple paths to a device and allows access to the multiple paths simultaneously. If the new multi-path behavior is enabled on the system, and the imported volume groups were configured with only persistent device special files, there is no need to arrange them again.

If needed, use the following steps to rebuild the LVM configuration files (`/etc/lvmtab` or `/etc/lvmtab_p`):

```
# vgreduce -f vgname pvname
# vgscan -f vgname
```

In cases where the physical volume is not readable (for example, when the physical volume is unattached either because the disk failed before volume group activation or because the system has been rebooted after the disk failure), running the `vgreduce` command with the `-f` option on those physical volumes removes them from the volume group, provided no logical volumes have extents mapped on that disk. Otherwise, if the unattached physical volume is not free, `- vgreduce -f` reports an extent map to identify the associated logical volumes. You must free all physical extents using `lvreduce` or `lvremove` before you can remove the physical volume with the `vgreduce` command.

This completes the procedure for removing the disk from your LVM configuration. If the disk hardware allows it, you can remove it physically from the system. Otherwise, physically remove it at the next scheduled system reboot.

Step 6: Replacing a Bad Disk (Persistent DSFs)

If instead of removing the disk, you need to replace the faulty disk, this section provides a step-by-step guide to replacing a faulty LVM disk, for systems configured with persistent DSFs. For systems using legacy DSFs, refer to the next step “[Step 7: Replacing a Bad Disk \(Legacy DSFs\)](#)” ([page 138](#)).

If you have any questions about the recovery process, contact your local HP Customer Response Center for assistance.

Because disks are physical devices, their hardware can fail, necessitating their replacement. After a failing disk is replaced with a new one (retaining the hardware address of the original disk to avoid confusion), the data must be restored to that disk from a backup.

Since the disk was under LVM control, it can have physical extents for several logical volumes on it. The layout of those logical volumes must first be restored and the data for each of those logical volumes restored from backup (if the data is not on a mirror copy).

There are basically four scenarios for replacing a bad disk. Determine your particular case and follow the steps provided for your scenario:

- [“Replacing a Mirrored Nonboot Disk” \(page 129\)](#)
- [“Replacing an Unmirrored Nonboot Disk” \(page 131\)](#)
- [“Replacing a Mirrored Boot Disk” \(page 134\)](#)
- [“Replacing an Unmirrored Boot Disk” \(page 137\)](#)

Replacing a Mirrored Nonboot Disk

Use this procedure if *all* the physical extents on the disk have copies on another disk, and the disk is *not* a boot disk. If the disk contains any unmirrored logical volumes or any mirrored logical volumes without an available and current mirror copy, see [“Replacing an Unmirrored Nonboot Disk” \(page 131\)](#).

For this example, the disk to be replaced is at lunpath hardware path 0/1/1/1.0x3.0x0, with device special files named /dev/disk/disk14 and /dev/rdisk/disk14. Follow these steps:

1. Save the hardware paths to the disk.

Run the `ioscan` command and note the hardware paths of the failed disk.

```
# ioscan -m lun /dev/disk/disk14
Class I Lun H/W Path      Driver S/W State H/W Type Health Description
=====
disk 14 64000/0xfa00/0x0  esdisk CLAIMED DEVICE offline HP MSA Vol
          0/1/1/1.0x3.0x0
          /dev/disk/disk14      /dev/rdisk/disk14
```

In this example, the LUN instance number is 14, the LUN hardware path is 64000/0xfa00/0x0, and the lunpath hardware path is 0/1/1/1.0x3.0x0.

When the failed disk is replaced, a new LUN instance and LUN hardware path are created. To identify the disk after it is replaced, you must use the lunpath hardware path (0/1/1/1.0x3.0x0).

2. Halt LVM access to the disk.

If the disk is not hot-swappable, power off the system to replace it. By shutting down the system, you halt LVM access to the disk, so you can skip this step.

If the disk is hot-swappable, detach it using the `-a` option of the `pvchange` command:

```
# pvchange -a N /dev/disk/disk14
```

3. Replace the disk.

For the hardware details on how to replace the disk, see the hardware administrator's guide for the system or disk array.

If the disk is hot-swappable, replace it.

If the disk is not hot-swappable, shut down the system, turn off the power, and replace the disk. Reboot the system.

- Notify the mass storage subsystem that the disk has been replaced.

If the system was not rebooted to replace the failed disk, then run `scsimgr` before using the new disk as a replacement for the old disk. For example:

```
# scsimgr replace_wwid -D /dev/rdisk/disk14
```

This command allows the storage subsystem to replace the old disk's LUN World-Wide-Identifier (WWID) with the new disk's LUN WWID. The storage subsystem creates a new LUN instance and new device special files for the replacement disk.

- Determine the new LUN instance number for the replacement disk.

For example:

```
# ioscan -m lun
Class I Lun H/W Path      Driver  S/W State H/W Type Health  Description
=====
disk 14 64000/0xfa00/0x0  esdisk  NO_HW    DEVICE   offline HP MSA Vol
                           /dev/disk/disk14        /dev/rdisk/disk14
...
disk 28 64000/0xfa00/0x1c esdisk  CLAIMED   DEVICE   online  HP MSA Vol
0/1/1/1.0x3.0x0           /dev/disk/disk28        /dev/rdisk/disk28
```

In this example, LUN instance 28 was created for the new disk, with LUN hardware path 64000/0xfa00/0x1c, device special files /dev/disk/disk28 and /dev/rdisk/disk28, at the same lunpath hardware path as the old disk, 0/1/1/1.0x3.0x0. The old LUN instance 14 for the old disk now has no lunpath associated with it.

NOTE: If the system was rebooted to replace the failed disk, then `ioscan -m lun` does not display the old disk.

- Assign the old instance number to the replacement disk.

For example:

```
# io_redirect_dsf -d /dev/disk/disk14 -n /dev/disk/disk28
```

This assigns the old LUN instance number (14) to the replacement disk. In addition, the device special files for the new disk are renamed to be consistent with the old LUN instance number. The following `ioscan -m lun` output shows the result:

```
# ioscan -m lun /dev/disk/disk14
Class I Lun H/W Path      Driver  S/W State H/W Type Health  Description
=====
disk 14 64000/0xfa00/0x1c esdisk  CLAIMED   DEVICE   online  HP MSA Vol
0/1/1/1.0x3.0x0           /dev/disk/disk14        /dev/rdisk/disk14
```

The LUN representation of the old disk with LUN hardware path 64000/0xfa00/0x0 was removed. The LUN representation of the new disk with LUN hardware path 64000/0xfa00/0x1c was reassigned from LUN instance 28 to LUN instance 14 and its device special files were renamed as /dev/disk/disk14 and /dev/rdisk/disk14.

- Restore LVM configuration information to the new disk.

For example:

```
# vgcfgrestore -n /dev/vgnn /dev/rdisk/disk14
```

8. Restore LVM access to the disk.

If you did *not* reboot the system in [Step 2](#), “Halt LVM access to the disk,” reattach the disk as follows:

```
# pvchange -a y /dev/disk/disk14
```

If you did reboot the system, reattach the disk by reactivating the volume group as follows:

```
# vgchange -a y /dev/vg01
```

NOTE: The `vgchange` command with the `-a y` option can be run on a volume group that is deactivated or already activated. It attaches all paths for all disks in the volume group and resumes automatically recovering any disks in the volume group that had been offline or any disks in the volume group that were replaced. Therefore, run `vgchange` only after all work has been completed on all disks and paths in the volume group, and it is necessary to attach them all.

Because all the data on the replaced disk was mirrored, you do not need to do anything else; LVM automatically synchronizes the data on the disk with the other mirror copies of the data.

Replacing an Unmirrored Nonboot Disk

Use this procedure if *any* of the physical extents on the disk do not have mirror copies elsewhere, and your disk is not a boot disk.

In this example, the disk to be replaced is at lunpath hardware path 0/1/1/1.0x3.0x0, with device special files named `/dev/disk/disk14` and `/dev/rdisk/disk14`. Follow these steps:

1. Save the hardware paths to the disk.

Run the `ioscan` command and note the hardware paths of the failed disk.

```
# ioscan -m lun /dev/disk/disk14
Class I Lun H/W Path      Driver S/W State H/W Type Health Description
=====
disk 14 64000/0xfa00/0x0  esdisk CLAIMED DEVICE offline HP MSA Vol
0/1/1/1.0x3.0x0
          /dev/disk/disk14      /dev/rdisk/disk14
```

In this example, the LUN instance number is 14, the LUN hardware path is 64000/0xfa00/0x0, and the lunpath hardware path is 0/1/1/1.0x3.0x0.

When the failed disk is replaced, a new LUN instance and LUN hardware path are created. To identify the disk after it is replaced, you must use the lunpath hardware path (0/1/1/1.0x3.0x0).

2. Halt LVM access to the disk.

If the disk is not hot-swappable, power off the system to replace it. By shutting down the system, you halt LVM access to the disk, so you can skip this step.

If the disk is hot-swappable, disable user and LVM access to all unmirrored logical volumes.

First, disable *user* access to all unmirrored logical volumes. Halt any applications and unmount any file systems using these logical volumes. This prevents the applications or file systems from writing inconsistent data over the newly restored replacement disk.

For each unmirrored logical volume using the disk:

- Use the `fuser` command to make sure no one is accessing the logical volume, either as a raw device or as a file system. If users have files open in the file system or it is their current working directory, `fuser` reports their process IDs.

For example, if the logical volume was `/dev/vg01/lvol1`, enter the following command:

```
# fuser -cu dev/vg01/lvol1
/dev/vg01/lvol1:    27815c(root)    27184c(root)
```

- b. If fuser reports process IDs using the logical volume, use the ps command to map the list of process IDs to processes, and then determine whether you can halt those processes.

For example, look up processes 27815 and 27184 as follows:

```
# ps -fp27815 -p27184
  UID      PID  PPID   C   STIME  TTY          TIME COMMAND
root    27815 27184   0 09:04:05 pts/0        0:00 vi test.c
root    27184 27182   0 08:26:24 pts/0        0:00 -sh
```

- c. If so, use fuser with the -k option to kill all processes accessing the logical volume.

The example processes are noncritical, so kill them as follows:

```
# fuser -ku dev/vg01/lvol1
/dev/vg01/lvol1:    27815c(root)    27184c(root)
```

- d. If the logical volume is being used as a file system, unmount it as follows:

```
# umount /dev/vg01/lvol1
```

NOTE: If you cannot stop the applications using the logical volume, or you cannot unmount the file system, you must shut down the system.

After disabling user access to the unmirrored logical volumes, disable LVM access to the disk:

```
# pvchange -a N /dev/disk/disk14
```

3. Replace the disk.

For the hardware details on how to replace the disk, see the hardware administrator's guide for the system or disk array.

If the disk is hot-swappable, replace it.

If the disk is not hot-swappable, shut down the system, turn off the power, and replace the disk. Reboot the system.

4. Notify the mass storage subsystem that the disk has been replaced.

If the system was not rebooted to replace the failed disk, then run scsimgr before using the new disk as a replacement for the old disk. For example:

```
# scsimgr replace_wwid -D /dev/rdisk/disk14
```

This command allows the storage subsystem to replace the old disk's LUN World-Wide-Identifier (WWID) with the new disk's LUN WWID. The storage subsystem creates a new LUN instance and new device special files for the replacement disk.

5. Determine the new LUN instance number for the replacement disk.

For example:

```
# ioscan -m lun
Class I Lun H/W Path      Driver S/W State H/W Type Health Description
=====
disk 14 64000/0xfa00/0x0 esdisk NO_HW DEVICE offline HP MSA Vol
                  /dev/disk/disk14      /dev/rdisk/disk14
...
disk 28 64000/0xfa00/0x1c esdisk CLAIMED DEVICE online HP MSA Vol
  0/1/1/1.0x3.0x0
                  /dev/disk/disk28      /dev/rdisk/disk28
```

In this example, LUN instance 28 was created for the new disk, with LUN hardware path 64000/0xfa00/0x1c, device special files /dev/disk/disk28 and /dev/rdisk/disk28, at the same lunpath hardware path as the old disk, 0/1/1/1.0x3.0x0. The old LUN instance 14 for the old disk now has no lunpath associated with it.

NOTE: If the system was rebooted to replace the failed disk, then `ioscan -m lun` does not display the old disk.

6. Assign the old instance number to the replacement disk.

For example:

```
# io_redirect_dsf -d /dev/disk/disk14 -n /dev/disk/disk28
```

This assigns the old LUN instance number (14) to the replacement disk. In addition, the device special files for the new disk are renamed to be consistent with the old LUN instance number. The following `ioscan -m lun` output shows the result:

```
# ioscan -m lun /dev/disk/disk14
Class I Lun H/W Path      Driver S/W State H/W Type Health Description
=====
disk 14 64000/0xfa00/0x1c esdisk CLAIMED DEVICE online HP MSA Vol
0/1/1/1.0x3.0x0
          /dev/disk/disk14      /dev/rdisk/disk14
```

The LUN representation of the old disk with LUN hardware path 64000/0xfa00/0x0 was removed. The LUN representation of the new disk with LUN hardware path 64000/0xfa00/0x1c was reassigned from LUN instance 28 to LUN instance 14 and its device special files were renamed as `/dev/disk/disk14` and `/dev/rdisk/disk14`.

7. Restore LVM configuration information to the new disk.

For example:

```
# vgcfgrestore -n /dev/vgnn /dev/rdisk/disk14
```

8. Restore LVM access to the disk.

If you did *not* reboot the system in [Step 2](#), “Halt LVM access to the disk,” reattach the disk as follows:

```
# pvchange -a y /dev/disk/disk14
```

If you did reboot the system, reattach the disk by reactivating the volume group as follows:

```
# vgchange -a y /dev/vggn
```

NOTE: The `vgchange` command with the `-a y` option can be run on a volume group that is deactivated or already activated. It attaches all paths for all disks in the volume group and resumes automatically recovering any disks in the volume group that had been offline or any disks in the volume group that were replaced.

Therefore, run `vgchange` only after all work has been completed on all disks and paths in the volume group, and it is necessary to attach them all.

9. Recover any lost data.

LVM recovers all the mirrored logical volumes on the disk, and starts that recovery when the volume group is activated.

For all the unmirrored logical volumes that you identified in [Step 2](#), "Halt LVM access to the disk," restore the data from backup and reenable user access as follows:

- For raw volumes, restore the full raw volume using the utility that was used to create your backup. Then restart the application.
- For file systems, you must re-create the file systems first. For example:

```
# newfs -F fstype /dev/vgnn/r1volnn
```

Use the logical volume's character device file for the `newfs` command. For file systems that had nondefault configurations, see `newfs(1M)` for the correct options.

After creating the file system, mount it under the mount point that it previously occupied. Then restore the data for that file system from your full backups.



TIP: To make the file system restoration step easier, record how they were originally created. You can change other file system parameters, such as those used to tune file system performance. The file system must be at least as large as before the disk failure.

Replacing a Mirrored Boot Disk

There are two additional operations you must perform when replacing a mirrored boot disk:

1. You must initialize boot information on the replacement disk.
2. If the replacement requires rebooting the system, and the primary boot disk is being replaced, you must boot from the alternate boot disk.

In this example, the disk to be replaced is at lunpath hardware path 0/1/1/1.0x3.0x0, with device special files named `/dev/disk/disk14` and `/dev/rdisk/disk14`. The system is an HP Integrity server, so the physical volume names must specify the HP-UX partition on the boot disk (`/dev/disk/disk14_p2` and `/dev/disk/disk14_p2`).

1. Save the hardware paths to the disk.

Run the `ioscan` command and note the hardware paths of the failed disk as follows:

```
# ioscan -m lun /dev/disk/disk14
Class I Lun H/W Path      Driver  S/W State H/W Type Health  Description
=====
disk 14 64000/0xfa00/0x0  esdisk  CLAIMED   DEVICE    offline HP MSA Vol
          0/1/1/1.0x3.0x0
                  /dev/disk/disk14      /dev/rdisk/disk14
                  /dev/disk/disk14_p1    /dev/rdisk/disk14_p1
                  /dev/disk/disk14_p2    /dev/rdisk/disk14_p2
                  /dev/disk/disk14_p3    /dev/rdisk/disk14_p3
```

In this example, the LUN instance number is 14, the LUN hardware path is 64000/0xfa00/0x0, and the lunpath hardware path is 0/1/1/1.0x3.0x0.

When the failed disk is replaced, a new LUN instance and LUN hardware path are created. To identify the disk after it is replaced, you must use the lunpath hardware path (0/1/1/1.0x3.0x0).

2. Halt LVM access to the disk.

If the disk is not hot-swappable, power off the system to replace it. By shutting down the system, you halt LVM access to the disk, so you can skip this step.

If the disk is hot-swappable, detach the device using the `-a` option of the `pvchange` command:

```
# pvchange -a n /dev/disk/disk14_p2
```

NOTE: On an HP 9000 server, the boot disk is not partitioned so the physical volume refers to the entire disk, not the HP-UX partition. Use the following command:

```
# pvchange -a N /dev/disk/disk14
```

3. Replace the disk.

For the hardware details on how to replace the disk, see the hardware administrator's guide for the system or disk array.

If the disk is hot-swappable, replace it.

If the disk is not hot-swappable, shut down the system, turn off the power, and replace the disk. Reboot the system. Two problems can occur:

- If you replaced the disk that you normally boot from, the replacement disk does not contain the information needed by the boot loader. In this case, interrupt the boot process and boot from the mirror boot disk, which is configured as the alternate boot path.
- If there are only two disks in the root volume group, the system probably fails its quorum check as described in ["Volume Group Activation Failures" \(page 111\)](#). It can panic early in the boot process with the message:

```
panic: LVM: Configuration failure
```

In this situation, you must override quorum to boot successfully. Do this by interrupting the boot process and adding the `-lq` option to the boot command.

For information on the boot process and how to select boot options, see *HP-UX System Administrator's Guide: Configuration Management*.

4. Notify the mass storage subsystem that the disk has been replaced.

If the system was not rebooted to replace the failed disk, then run `scsimgr` before using the new disk as a replacement for the old disk. For example:

```
# scsimgr replace_wwid -D /dev/rdisk/disk14
```

This command allows the storage subsystem to replace the old disk's LUN World-Wide-Identifier (WWID) with the new disk's LUN WWID. The storage subsystem creates a new LUN instance and new device special files for the replacement disk.

5. Determine the new LUN instance number for the replacement disk.

For example:

```
# ioscan -m lun
Class I Lun H/W Path      Driver S/W State H/W Type Health Description
=====
disk 14 64000/0xfa00/0x0 esdisk NO_HW   DEVICE offline HP MSA Vol
        /dev/disk/disk14    /dev/rdisk/disk14
        /dev/disk/disk14_p1  /dev/rdisk/disk14_p1
        /dev/disk/disk14_p2  /dev/rdisk/disk14_p2
        /dev/disk/disk14_p3  /dev/rdisk/disk14_p3
...
disk 28 64000/0xfa00/0x1c esdisk CLAIMED DEVICE online  HP MSA Vol
        0/1/1/1.0x3.0x0
        /dev/disk/disk28     /dev/rdisk/disk28
```

In this example, LUN instance 28 was created for the new disk, with LUN hardware path 64000/0xfa00/0x1c, device special files /dev/disk/disk28 and /dev/rdisk/disk28, at the same lunpath hardware path as the old disk, 0/1/1/1.0x3.0x0. The old LUN instance 14 for the old disk now has no lunpath associated with it.

NOTE: If the system was rebooted to replace the failed disk, then `ioscan -m lun` does not display the old disk.

6. (HP Integrity servers only) Partition the replacement disk.

Partition the disk using the `idisk` command and a partition description file, and create the partition device files using `insf`, as described in “[Mirroring the Boot Disk on HP Integrity Servers](#)” (page 92).

7. Assign the old instance number to the replacement disk.

For example:

```
# io_redirect_dsf -d /dev/disk/disk14 -n /dev/disk/disk28
```

This assigns the old LUN instance number (14) to the replacement disk. In addition, the device special files for the new disk are renamed to be consistent with the old LUN instance number. The following `ioscan -m lun` output shows the result:

```
# ioscan -m lun /dev/disk/disk14
Class I Lun H/W Path      Driver S/W State H/W Type Health Description
=====
disk 14 64000/0xfa00/0x1c esdisk CLAIMED DEVICE online HP MSA Vol
  0/1/1.0x3.0x0
    /dev/disk/disk14      /dev/rdisk/disk14
    /dev/disk/disk14_p1   /dev/rdisk/disk14_p1
    /dev/disk/disk14_p2   /dev/rdisk/disk14_p2
    /dev/disk/disk14_p3   /dev/rdisk/disk14_p3
```

The LUN representation of the old disk with LUN hardware path `64000/0xfa00/0x0` was removed. The LUN representation of the new disk with LUN hardware path `64000/0xfa00/0x1c` was reassigned from LUN instance 28 to LUN instance 14 and its device special files were renamed as `/dev/disk/disk14` and `/dev/rdisk/disk14`.

8. Restore LVM configuration information to the new disk.

For example:

```
# vgcfgrestore -n /dev/vg00 /dev/rdisk/disk14_p2
```

NOTE: On an HP 9000 server, the boot disk is not partitioned, so the physical volume refers to the entire disk, not the HP-UX partition. Use the following command:

```
# vgcfgrestore -n /dev/vg00 /dev/rdisk/disk14
```

9. Restore LVM access to the disk.

If you did *not* reboot the system in [Step 2](#), “Halt LVM access to the disk,” reattach the disk as follows:

```
# pvchange -a y /dev/disk/disk14_p2
```

On an HP 9000 server, use this command:

```
# pvchange -a y /dev/disk/disk14
```

If you did reboot the system, reattach the disk by reactivating the volume group as follows:

```
# vgchange -a y /dev/vg00
```

NOTE: The vgchange command with the -a y option can be run on a volume group that is deactivated or already activated. It attaches all paths for all disks in the volume group and resumes automatically recovering any disks in the volume group that had been offline or any disks in the volume group that were replaced. Therefore, run vgchange only after all work has been completed on all disks and paths in the volume group, and it is necessary to attach them all.

10. Initialize boot information on the disk.

For an HP Integrity server, set up the boot area and update the autoboot file in the disk's EFI partition as described in step 5 and step 6 of "[Mirroring the Boot Disk on HP Integrity Servers](#)" ([page 92](#)).

For an HP 9000 server, set up the boot area and update the autoboot file as described in step 4 and step 5 of "[Mirroring the Boot Disk on HP 9000 Servers](#)" ([page 90](#)).

Replacing an Unmirrored Boot Disk

With the failure of an unmirrored boot disk, you have lost the only copy of information that is required to boot the system. You must reinstall onto the replacement disk, or recover it from an Ignite-UX backup.

Step 7: Replacing a Bad Disk (Legacy DSFs)

Follow these steps to replace a bad disk if your system is configured with only legacy DSFs.

NOTE: LVM recommends the use of persistent device special files, because they support a greater variety of load balancing options. For replacing a disk with persistent device special files, see “[Step 6: Replacing a Bad Disk \(Persistent DSFs\)](#)” (page 129)

To replace a bad disk, follow these steps.

1. Halt LVM Access to the Disk.

If the disk is not hot-swappable, power off the system to replace it. By shutting down the system, you halt LVM access to the disk, so you can skip this step.

If the disk is hot-swappable, disable user and LVM access to all unmirrored logical volumes.

- First, disable user access to all unmirrored logical volumes or any mirrored logical volumes without an available and current mirror copy. Halt any application and umount any file systems using these logical volumes. This prevents the applications or file systems from writing inconsistent data over the newly restored replacement disk.
- After disabling user access to the unmirrored logical volumes, disable LVM access to the disk, as follows:

```
# pvchange -a N pvname
```

The following recommendation are intended to maximize system uptime and access to the volume group, but you can use a stronger approach if your data and system availability requirement allows.

- If `pvdisplay` shows “PV status” as available, halt LVM access to the disk by removing it from the volume group.
- If `pvdisplay` show “PV status” as unavailable, or if `pvdisplay` fails to print the status, use `ioscan` to determine if the disk can be accessed at all. If `ioscan` reports the disk status as NO_HW on all its hardware paths, you can remove the disk. If `ioscan` shows any other status, halt LVM access to the disk by deactivating the volume group.

2. Replace the Faulty Disk

If the disk is hot-swappable, you can replace it without powering down the system. Otherwise, power down the system before replacing the disk. For the hardware details on how to replace the disk, see the hardware administrator’s guide for the system or disk array. If you powered down the system, reboot it normally. The only exception is if you replaced a disk in the root volume group.

- If you replaced the disk that you normally boot from, the replacement disk does not contain the information needed by the boot loader. If your root disk is mirrored, boot from it by using the alternate boot path. If the root disk was not mirrored, you must reinstall or recover your system.
- If there are only two disks in the root volume group, the system probably fails its quorum check as described in “[Volume Group Activation Failures](#)” (page 111). It can panic early in the boot process with the message:

```
panic: LVM: Configuration failure
```

In this situation, you must override quorum to boot successfully. Do this by interrupting the boot process and adding the `-lq` option to the boot command.

For information on the boot process and how to select boot options, see *HP-UX System Administrator’s Guide: Configuration Management*.

3. Initialize the Disk for LVM

This step copies LVM configuration information onto the disk, and marks it as owned by LVM so it can subsequently be attached to the volume group.

If you replaced a mirror of the root disk on an Integrity server, run the `idisk` and `insf` commands as described in “[Mirroring the Boot Disk on HP Integrity Servers](#)” (page 92). For PA-RISC servers or non-root disks, this step is unnecessary.

For any replaced disk, restore LVM configuration information to the disk using the `vgcfgrestore` command as follows:

```
# vgcfgrestore -n vgname pvname
```

NOTE: If the `vgcfgrestore` command fails to write the original LVM header back to the new disk because a valid LVM configuration backup file (`/etc/lvmconf/vgXX.conf[.old]`) is missing or corrupted, you must remove the physical volume that is being restored from the volume group (with the `vgreduce` command) to get a clean configuration.

In these situations the `vgcfgrestore` command might fail to restore the LVM header, issuing this message:

Mismatch between the backup file and the running kernel

If you are sure that your backup is valid, you can override this check by using the `-R` option. To remove a physical volume from a volume group, you must first free it by removing all of the logical extents. If the logical volumes on such a disk are not mirrored, the data is lost anyway. If it is mirrored, you must reduce the mirror before removing the physical volume.

4. Re-Enable LVM Access to the Disk.

This step re-attaches the disk, as follows

```
# pvchange -a y pvname
```

5. Restore Lost Data to the Disk

This final step can be a straightforward re-synchronization for mirrored configurations, or a recovery of data from backup media.

- If a mirror of the root disk was replaced, initialize its boot information as follows:
 - For an Integrity server, follow steps 5, 6, and 8 in “[Mirroring the Boot Disk on HP Integrity Servers](#)” (page 92)
 - For a PA-RISC server, follow steps 4, 5, and 7 in “[Mirroring the Boot Disk on HP 9000 Servers](#)” (page 90)
- If all the data on the replaced disk was mirrored, you do not have to do anything; LVM automatically synchronizes the data on the disk with the other mirror copies of the data.
- If the disk contained any unmirrored logical volumes (or mirrored logical volumes that did not have a current copy on the system), restore the data from backup, mount the file systems, and restart any applications you halted in Step 1.

Reporting Problems

If you are unable to solve a problem with LVM, follow these steps:

1. Read the *HP-UX Logical Volume Manager and MirrorDisk/UX Release Notes* to see if the problem is known. If it is, follow the solution offered to solve the problem.
2. Determine if the product is still under warranty or if your company purchased support services for the product. Your operations manager can supply you with the necessary information.
3. Access <http://www.itrc.hp.com> and search the technical knowledge databases to determine if the problem you are experiencing has been reported already. The type of documentation and resources you have access to depend on your level of entitlement.

NOTE: The ITRC resource forums at <http://www.itrc.hp.com> offer peer-to-peer support to solve problems and are free to users after registration.

If this is a new problem or if you need additional help, log your problem with the HP Response Center, either online through the support case manager at <http://www.itrc.hp.com>, or by calling HP Support. If your warranty has expired or if you do not have a valid support contract for your product, you can still obtain support services for a fee, based on the amount of time and material required to solve your problem.

4. If you are asked to supply any information pertaining to the problem, gather the requested information and submit it.

5 Support and Other Resources

New and Changed Information in This Edition

The eighth edition of *HP-UX System Administrator's Guide: Logical Volume Management* addresses the following new topics:

- Added information about converting cDSFs back to their corresponding persistent DSFs, in [Table 5 \(page 39\)](#).
- Provided new information on LVM I/O timeout parameters, see ["Configuring LVM I/O Timeout Parameters" \(page 33\)](#) and ["LVM I/O Timeout Parameters" \(page 161\)](#).
- Added error message information for vgmodify in ["Warning and Error Messages" \(page 163\)](#).
- Added error message and recommendation for failure in deletion of snapshot logical volume device, see ["Warning and Error Messages" \(page 163\)](#).

About this Series

The *HP-UX System Administrator's Guide* documents the core set of tasks (and associated concepts) necessary to administer systems running HP-UX 11i Version 3.

The *HP-UX System Administrator's Guide* is a set of documents comprised of the following volumes:

Overview	Provides a high-level view of HP-UX 11i, its components, and how they relate to each other.
Configuration Management	Describes many of the tasks you must perform to configure and customize system settings and the behavior of subsystems.
Logical Volume Management	Documents how to configure physical volumes, volume groups, and logical volumes using the HP Logical Volume Manager (LVM).
Security Management	Documents the data and system security features of HP-UX 11i.
Routine Management Tasks	Documents many of the ongoing tasks you must perform to keep your system running smoothly.

Typographic Conventions

This document uses the following typographical conventions:

%, \$, or #	A percent sign represents the C shell system prompt. A dollar sign represents the system prompt for the Bourne, Korn, and POSIX shells. A number sign represents the superuser prompt.
<code>audit(5)</code>	A manpage. The manpage name is <i>audit</i> , and it is located in Section 5.
Command	A command name or qualified command phrase.
Computer output	Text displayed by the computer.
Ctrl+x	A key sequence. A sequence such as Ctrl+x indicates that you must hold down the key labeled Ctrl while you press another key or mouse button.
Document Title	The title of a document. On the web and on the Instant Information media, it may be a hot link to the document itself.
ENVIRONMENT VARIABLE	The name of an environment variable, for example, PATH.
ERROR NAME	The name of an error, usually returned in the <code>errno</code> variable.

Key	The name of a keyboard key. Return and Enter both refer to the same key.
Term	The defined use of an important word or phrase.
User input	Commands and other text that you type.
<i>Variable or Replaceable</i>	The name of a placeholder in a command, function, or other syntax display that you replace with an actual value.
- <i>chars</i>	One or more grouped command options, such as -ikx. The <i>chars</i> are usually a string of literal characters that each represent a specific option. For example, the entry -ikx is equivalent to the individual options -i, -k, and -x. The plus character (+) is sometimes used as an option prefix.
- <i>word</i>	A single command option, such as -help. The <i>word</i> is a literal keyword. The difference from - <i>chars</i> is usually obvious and is clarified in an Options description. The plus character (+) and the double hyphen (--) are sometimes used as option prefixes.
[]	The bracket metacharacters enclose optional content. If the contents are a list separated by , you choose one of the items.
{ }	The brace metacharacters enclose required content. If the contents are a list separated by , you must choose one of the items.
...	The preceding element can be repeated an arbitrary number of times. Ellipsis is sometimes used to indicate omitted items in a range.
:	Indicates the continuation of a code example.
	Separates items in a list of choices.
WARNING	A warning calls attention to important information that if not understood or followed will result in personal injury or nonrecoverable system problems.
CAUTION	A caution calls attention to important information that if not understood or followed will result in data loss, data corruption, or damage to hardware or software.
IMPORTANT	This alert provides essential information to explain a concept or to complete a task.
NOTE	A note contains additional information to emphasize or supplement important points of the main text.

Examples and Shells

This document describes practices used by the system administrator. Since the `root` user (superuser) is required to use the POSIX shell `/sbin/sh`, all command examples use that shell. The POSIX shell is defined in *sh-posix(1)*. For information on other shells, see the *Shells User's Guide* and *sh(1)*.

Related Information

HP-UX technical documentation can be found on HP's documentation website at <http://www.hp.com/go/hpxx-core-docs>.

In particular, LVM documentation is provided on this web page:<http://www.hp.com/go/hpxx-LVM-VxVM-docs>. See the *HP-UX Logical Volume Manager and Mirror Disk/UX Release Notes* for information about new features and defect fixes on each release. In addition, the following white papers are available:

- *LVM Migration from Legacy to Agile Naming Model*
- *LVM Online Disk Replacement (LVM OLR)*
- *LVM Volume Group Quiesce/Resume*
- *SLVM Online Reconfiguration*
- *Using the vgmodify Command to Perform LVM Volume Group Dynamic LUN Expansion (DLE) and Contraction (DLC)*
- *LVM Snapshot Logical Volumes*

For detailed description of each LVM command and their options, see the LVM manpages that are installed with the LVM product, accessible via the `man` command. The manpage are also published on <http://www.hp.com/go/hpxx-clickable-manpages>

Finding HP-UX Information

The following table outlines where to find system administration information for HP-UX. This table does not include information for specific products.

Table 10 Finding HP-UX Information

If you need to	Refer To	Located at .
Find out: <ul style="list-style-type: none">• What has changed in HP-UX releases• The contents of the Operating Environments• Firmware requirements and supported systems for a specific release	The HP-UX 11i Release Notes specific to your version of HP-UX.	<ul style="list-style-type: none">• HP Instant Information media• HP-UX Technical Documentation website at:http://www.hp.com/go/hpxx-core-docs
Install or update HP-UX	<ul style="list-style-type: none">• <i>Read Before Installing or Updating to HP-UX</i>• <i>HP-UX 11i Installation and Update Guide</i>	<ul style="list-style-type: none">• Media Kit (supplied with the Operating Environment)• HP Instant Information media• HP-UX Technical Documentation website at:http://www.hp.com/go/hpxx-core-docs
Administer an HP-UX system	Releases beginning with HP-UX 11i Version 3: <ul style="list-style-type: none">• <i>HP-UX System Administrator's Guide</i> (a multivolume set) Other sources of system administration information: <ul style="list-style-type: none">• <i>nPartition Administrator's Guide</i>• <i>Planning Superdome Configuration White Paper</i>• <i>HP Superdome 2 Partitioning Administrator Guide</i>	<ul style="list-style-type: none">• HP Instant Information media• HP-UX Technical Documentation website at:http://www.hp.com/go/hpxx-core-docs

HP-UX 11i Release Names and Operating System Version Identifiers

With HP-UX 11i, HP delivers a highly available, secure, and manageable operating system that meets the demands of end-to-end Internet-critical computing. HP-UX 11i supports enterprise, mission-critical, and technical computing environments. HP-UX 11i is available on both HP 9000 systems and HP Integrity systems.

Each HP-UX 11i release has an associated release name and release identifier. The `uname` command with the `-r` option returns the release identifier. [Table 11](#) shows the releases available for HP-UX 11i.

Table 11 HP-UX 11i Releases

OS Version Identifier	Release Name	Supported Processor Architecture
B.11.11	HP-UX 11i Version 1	HP 9000
B.11.23	HP-UX 11i Version 2	Integrity
B.11.23.0409	HP-UX 11i Version 2 September 2004 Update	HP 9000 and Integrity
B.11.31	HP-UX 11i Version 3	HP 9000 and Integrity

For information on supported systems and processor architecture for various versions of HP-UX 11i, see the HP-UX 11i system release notes specific to your version of HP-UX.

Determining Your System Version

The `uname`, `model`, and `swlist` commands can help you determine information about your system, including its hardware type, machine model, operating system version, and operating environment update status. See `uname(1)`, `model(1)`, and `swlist(1M)`.

For OS naming conventions, see "[HP-UX 11i Release Names and Operating System Version Identifiers](#)" (page 144).

Table 12 OS Version, System Architecture, and Machine Model

Topic	Command	Sample Output
OS Version	\$ <code>uname -r</code>	B.11.31 ¹
Architecture	\$ <code>uname -m</code>	ia64 ² 9000/800 ²
Machine Model	\$ <code>model</code> ³	ia64 hp server rx5670 9000/800/S16K-A
Operating Environment	\$ <code>swlist HPUX*OE*</code>	# HPUX11i-OE-MC B.11.31 HP-UX Mission Critical Operating Environment ¹
OS Version.Update	\$ <code>swlist HPUX*OE*</code>	# HPUX11i-TCOE B.11.23.0409 HP-UX Technical Computing OE Component ¹

¹ HP-UX 11i OS version identifiers have the form B.11.23 or B.11.23.0409, where B.11.23 is the OS version and 0409 is the year-month of the operating environment (OE) update.

² ia64 = Integrity. All others = HP 9000.

³ The `getconf MACHINE_MODEL` command gives the same output. See `getconf(1)`.

Publication History

The document printing date and part number indicate the document's current edition. The printing date changes when a new edition is printed. Minor changes can be made at reprint without changing the printing date. The document part number changes when extensive changes are made.

Document updates can be issued between editions to correct errors or document product changes. To ensure that you receive the updated or new editions, subscribe to the appropriate product support service. See your HP sales representative for details. You can find the latest version of this document online at:

<http://www.hp.com/go/hpux-LVM-VxVM-docs> .

Manufacturing Part Number	Supported Operating Systems	Supported Versions	Edition Number	Publication Date
5991-6481	HP-UX	HP-UX 11i Version 3	1	February 2007
5992-3385	HP-UX	HP-UX 11i Version 3	2	March 2008
5992-4589	HP-UX	HP-UX 11i Version 3	3	September 2008
5992-6576	HP-UX	HP-UX 11i Version 3	4	September 2009
B3921-90000	HP-UX	HP-UX 11i Version 3	5	March 2010
B3921-90014	HP-UX	HP-UX 11i Version 3	6	September 2010
B3921-90038	HP-UX	HP-UX 11i Version 3	7	March 2011
B3921-90053	HP-UX	HP-UX 11i Version 3	8	September 2011

NOTE: The volumes in the *HP-UX System Administrator's Guide* can be updated independently. Therefore, the latest versions of the volumes in the set can vary with time and with respect to each other. The latest versions of each volume are available at <http://www.hp.com/go/hpux-core-docs>.

HP Insight Remote Support Software

HP strongly recommends that you install HP Insight Remote Support software to complete the installation or upgrade of your product and to enable enhanced delivery of your HP Warranty, HP Care Pack Service or HP contractual support agreement. HP Insight Remote Support supplements your monitoring, 24x7 to ensure maximum system availability by providing intelligent event diagnosis, and automatic, secure submission of hardware event notifications to HP, which will initiate a fast and accurate resolution, based on your product's service level. Notifications may be sent to your authorized HP Channel Partner for on-site service, if configured and available in your country. The software is available in two variants:

- **HP Insight Remote Support Standard:** This software supports server and storage devices and is optimized for environments with 1-50 servers. Ideal for customers who can benefit from proactive notification, but do not need proactive service delivery and integration with a management platform.
- **HP Insight Remote Support Advanced:** This software provides comprehensive remote monitoring and proactive service support for nearly all HP servers, storage, network, and SAN environments, plus selected non-HP servers that have a support obligation with HP. It is integrated with HP Systems Insight Manager. A dedicated server is recommended to host both HP Systems Insight Manager and HP Insight Remote Support Advanced.

Details for both versions are available at:

<http://www.hp.com/go/insightremotesupport>

To download the software, go to Software Depot:

<http://www.software.hp.com>

Select **Insight Remote Support** from the menu on the right.

HP Encourages Your Comments

HP encourages your comments concerning this document. We are committed to providing documentation that meets your needs. Send any errors found, suggestions for improvement, or compliments to:

<http://www.hp.com/bizsupport/feedback/ww/webfeedback.html>

Include the document title, manufacturing part number, and any comment, error found, or suggestion for improvement you have concerning this document.

A LVM Specifications and Limitations

This appendix discusses LVM product specifications.

NOTE: Do not infer that a system configured to these limits is automatically usable.

Table 13 Volume Group Version Maximums

	Version 1.0 Volume Groups	Version 2.0 Volume Groups	Version 2.1 Volume Groups	Version 2.2 Volume Groups
Maximum data on a single HP-UX system	128 PB	1024 PB	1024 PB	1024 PB
Maximum number of volume groups on a system	256	512 ¹	2048 ¹	2048 ¹
Maximum number of physical volumes in a volume group	255	511	2048	2048
Maximum number of logical volumes in a volume group	255	511	2047	2047 ²
Maximum size of a physical volume	2 TB	16 TB	16 TB	16 TB
Maximum size of a volume group	510 TB	2048 TB	2048 TB	2048 TB
Maximum size of a logical volume	16 TB	256 TB	256 TB	256 TB
Maximum size of a physical extent	256 MB	256 MB	256 MB	256 MB
Maximum size of a stripe	32 MB	256 MB	256 MB	256 MB
Maximum number of stripes	255	511	511	511
Maximum number of logical extents per logical volume	65535	33554432	33554432	33554432
Maximum number of physical extents per physical volume	65535	16777216	16777216	16777216
Maximum number of mirror copies (MirrorDisk/UX product required)	2	5	5	5
Maximum number of snapshots per LV	Not supported	Not supported	Not Supported	255
Maximum size for the unit of unsharing between logical volume and its snapshot	Not supported	Not supported	Not Supported	4M
Maximum snapshot capacity per VG	Not supported	Not supported	Not Supported	Maximum VG capacity

- The limit of 2048 volume groups is shared among Version 2.x volume groups. Volume groups of Versions 2.x can be created with volume group numbers ranging from 0-2047. However, the maximum number of Version 2.0 volume groups that can be created is 512.
- For volume group Version 2.2 or higher, the total number of logical volumes includes normal logical volumes as well as snapshot logical volumes.

Table 14 Version 1.0 Volume Group Limits

Parameter	Command to Set/Change Parameter	Minimum Value	Default Value	Maximum Value
Number of volume groups on a system	n/a	0	n/a	256
Number of physical volumes in a volume group	<code>vgcreate -p max_pv</code> <code>vgmodify -p max_pv</code>	1	16	255
Number of logical volumes in a volume group	<code>vgcreate -l max_lv</code> <code>vgmodify -l max_lv</code>	1	255	255
Size of a physical volume	<code>pvcreate -s pv_size</code>	1 PE	LUN Capacity	2 TB
Size of a logical volume	<code>lvcreate -L lv_size</code> <code>lvextend -L lv_size</code>	0	0	16 TB
Size of a physical extent	<code>vgcreate -s pe_size</code>	1 MB	4 MB	256 MB
Size of a stripe	<code>lvcreate -l stripe_size</code>	4 KB	8 KB	32 MB
Number of stripes	<code>lvcreate -i stripes</code>	2	n/a	Number of PVs in VG
Number of logical extents per logical volume	<code>lvcreate -l max_le</code> <code>lvextend -l max_le</code>	0	0	65535
Number of physical extents per physical volume	<code>vgcreate -e max_pe</code> <code>vgmodify -e max_pe</code>	1	1016 ¹	65535
Number of mirror copies (MirrorDisk/UX product required)	<code>lvcreate -m copies</code> <code>lvextend -m copies</code>	0	0	2

¹ If more than 1016 extents are necessary to access all of the first physical volume in the volume group, `vgcreate` increases the default value to the size of the first physical volume divided by the physical extent size.

Table 15 Version 2.x Volume Group Limits

Parameter	Command to Set/Change Parameter	Minimum Value	Default Value	Maximum Value
Number of volume groups on a system	n/a	0	n/a	2048 ¹
Number of physical volumes in a volume group	n/a	511	511 (2.0) 2048 (2.1, 2.2)	511 (2.0) 2048 (2.1, 2.2)
Number of logical volumes in a volume group	n/a	511	511 (2.0) 2047 (2.1, 2.2)	511 (2.0) 2047 (2.1, 2.2) ²
Size of a volume group	<code>vgcreate -S max_vgsizE</code>	1 MB ³	n/a	2 PB
Size of a physical volume	<code>pvcREATE -s pv_size</code>	1 PE	IUN Capacity	16 TB
Size of a logical volume	<code>lvCREATE -L lv_size</code> <code>lvEXTEND -L lv_size</code>	0	0	256 TB
Size of a physical extent	<code>vgCREATE -s pe_size</code>	1 MB	n/a	256 MB
Size of a stripe	<code>lvCREATE -l stripe_size</code>	4 KB	n/a	256 MB
Number of stripes	<code>lvCREATE -i stripes</code>	2	n/a	511
Number of logical extents per logical volume	<code>lvCREATE -l max_le</code> <code>lvEXTEND -l max_le</code>	0	0	33554432
Number of physical extents per physical volume	n/a	1	IUN Capacity ÷ PE Size	16777216
Number of snapshots per logical volume (for volume group Version 2.2 and higher only)	n/a	0	0	255
Size of unshare unit for snapshots (for volume group Version 2.2 and higher only)	<code>vgCREATE -U unshare_unit</code>	512KB	1024KB	4096KB
Maximum snapshot capacity (for volume group Version 2.2 and higher only)	n/a	<i>max_vgsizE</i>	<i>max_vgsizE</i>	<i>max_vgsizE</i>
Number of mirror copies (MirrorDisk/UX product required)	<code>lvCREATE -m copies</code> <code>lvEXTEND -m copies</code>	0	0	5

- 1 The limit of 2048 volume groups is shared among Version 2.x volume groups. Volume groups of Versions 2.x can be created with volume group numbers ranging from 0-2047. However, the maximum number of Version 2.0 volume groups that can be created is 512.
- 2 For volume group Version 2.2 or higher, the total number of logical volume includes normal logical volumes as well as snapshot logical volumes.
- 3 If the total size of the specified physical volumes is larger than *max_vgsizE*, `vgCREATE` adjusts the minimum volume group size to the total size.

Determining LVM's Maximum Limits on a System

The March 2008 update to HP-UX 11i v3 (11.31) introduced a new command that enables the system administrator to determine the maximum LVM limits supported on the target system for a given volume group version. The `lvmadm` command displays the implemented limits for Version 1.0 and Version 2.x volume groups. It is impossible to create a volume group that exceeds these limits.

Example

```
# lvmadm -t
--- LVM Limits ---
VG Version          1.0
Max VG Size (Tbytes) 510
Max LV Size (Tbytes) 16
Max PV Size (Tbytes) 2
Max VGs             256
Max LVs              255
Max PVs              255
Max Mirrors          2
Max Stripes          255
Max Stripe Size (Kbytes) 32768
Max LXs per LV      65535
Max PXs per PV       65535
Max Extent Size (Mbytes) 256

VG Version          2.0
Max VG Size (Tbytes) 2048
Max LV Size (Tbytes) 256
Max PV Size (Tbytes) 16
Max VGs             512
Max LVs              511
Max PVs              511
Max Mirrors          5
Max Stripes          511
Max Stripe Size (Kbytes) 262144
Max LXs per LV      33554432
Max PXs per PV       16777216
Max Extent Size (Mbytes) 256

VG Version          2.1
Max VG Size (Tbytes) 2048
Max LV Size (Tbytes) 256
Max PV Size (Tbytes) 16
Max VGs             2048
Max LVs              2047
Max PVs              2048
Max Mirrors          5
Max Stripes          511
Max Stripe Size (Kbytes) 262144
Max LXs per LV      33554432
Max PXs per PV       16777216
Max Extent Size (Mbytes) 256

VG Version          2.2
Max VG Size (Tbytes) 2048
Max LV Size (Tbytes) 256
Max PV Size (Tbytes) 16
Max VGs             2048
Max LVs              2047
Max PVs              2048
Max Mirrors          5
Max Stripes          511
Max Stripe Size (Kbytes) 262144
Max LXs per LV      33554432
```

Max PXs per PV	16777216
Max Extent Size (Mbytes)	256
Min Unshare unit (Kbytes)	512
Max Unshare unit (Kbytes)	4096
Max Snapshots per LV	255

B LVM Command Summary

This appendix contains a summary of the LVM commands and descriptions of their use.

Table 16 LVM Command Summary

Command	Description and Example
extendfs	Extends a file system: <code># extendfs /dev/vg00/lvol3</code>
lvmadm	Displays the limits associated with a volume group version: <code># lvmadm -t -v 2.0</code>
lvchange	Changes the characteristics of a logical volume: <code># lvchange -t 60 /dev/vg00/lvol3</code>
lvcreate	Creates a logical volume in a volume group: <code># lvcreate -L 100 /dev/vg00</code>
lvdisplay	Displays information about logical volumes: <code># lvdisplay -v /dev/vg00/lvol1</code>
lvextend -m	Adds a mirror to a logical volume: <code># lvextend -m 1 /dev/vg00/lvol3</code>
lvextend -L	Increases the size of a logical volume: <code># lvextend -L 120 /dev/vg00/lvol3</code>
lvlnboot	Prepares a logical volume to become a root, swap, or dump area: <code># lvlnboot -d /dev/vg00/lvol2</code>
lvmerge	Merges split volumes into one logical volume: <code># lvmerge /dev/vg00/lvol4b /dev/vg00/lvol4</code>
lvmove	Migrate a logical volume to new disks: <code># lvmove -f newdisks.txt /dev/vg00/lvol2</code>
lvmpud	Daemon to handle online shared LVM reconfiguration, and pre-allocation of extents for space-efficient snapshots: <code># lvmpud</code>
lvreduce -L	Decreases the size of a logical volume: <code># lvreduce -L 100 /dev/vg00/lvol3</code>
lvreduce -m	Decreases the number of mirror copies of a logical volume: <code># lvreduce -m 0 /dev/vg00/lvol3</code>
lvremove	Removes logical volumes from a volume group: <code># lvremove /dev/vg00/lvol6</code>
lvrmboot	Removes a logical volume link to root, swap, or dump: <code># lvrmboot -d /dev/vg00/lvol2</code>
lvsplit	Splits a mirrored logical volume into two logical volumes: <code># lvsplit /dev/vg00/lvol4</code>
lvsync	Synchronizes stale logical volume mirrors: <code># lvsync /dev/vg00/lvol1</code>

Table 16 LVM Command Summary (continued)

Command	Description and Example
pvchange	Changes the characteristics of a physical volume: <code># pvchange -a n /dev/disk/disk2</code>
pvck	Performs a consistency check on a physical volume: <code># pvck /dev/disk/disk47_p2</code>
pvcreate	Creates a physical volume to be used as part of a volume group: <code># pvcreate /dev/rdisk/disk2</code>
pvdisplay	Displays information about a physical volume: <code># pvdisplay -v /dev/disk/disk2</code>
pvmove	Moves extents from one physical volume to another: <code># pvmove /dev/disk/disk2 /dev/disk/disk3</code>
pvremove	Removes LVM data structures from a physical volume: <code># pvremove /dev/rdisk/disk2</code>
vgcfgbackup	Saves the LVM configuration for a volume group: <code># vgcfgbackup vg00</code>
vgcfgrestore	Restores the LVM configuration: <code># vgcfgrestore -n /dev/vg00 /dev/rdisk/disk2</code>
vgchange	Turns a volume group off or on: <code># vgchange -a y /dev/vg00</code>
vgchgid	Changes the volume group ID of a physical volume: <code># vgchgid /dev/rdisk/disk3</code>
vgcdsf	Converts persistent device special files in a volume group to cluster device special files: <code># vgcdsf /dev/vg10</code>
vgcreate	Creates a volume group: <code># vgcreate /dev/vg01 /dev/disk/disk2 /dev/disk/disk3</code>
vgdisplay	Displays information about a volume group: <code># vgdisplay -v /dev/vg00</code>
vgdsf	Migrates volume groups from legacy to persistent device files: <code># vgdsf</code>
vgextend	Extends a volume group by adding a physical volume: <code># vgextend /dev/vg00 /dev/disk/disk2</code>
vgexport	Removes a volume group from the system: <code># vgexport /dev/vg01</code>
vgimport	Adds an existing volume group to the system: <code># vgimport -v /dev/vg04</code>
vgmodify	Modifies the configuration parameters of a volume group: <code># vgmodify -v -t -n -r vg32</code>
vgmove	Migrates a volume group to different disks: <code># vgmove -f diskmap.txt vg10</code>

Table 16 LVM Command Summary (continued)

Command	Description and Example
vgscan	Scans the system disks for volume groups: <code># vgscan -v</code>
vgreduce	Reduces a volume group by removing one or more physical volumes from it: <code># vgreduce /dev/vg00 /dev/disk/disk2</code>
vgremove	Removes the definition of a volume group from the system and the disks: <code># vgremove /dev/vg00 /dev/disk/disk2</code>
vgsync	Synchronizes all mirrored logical volumes in the volume group: <code># vgsync vg00</code>
vgversion	Migrates a volume group to a different version: <code># vgversion -v 2.1 vg14</code>

C Volume Group Provisioning Tips

This appendix contains recommendations for parameters to use when creating your volume groups.

Choosing an Optimal Extent Size for a Version 1.0 Volume Group

When creating a Version 1.0 volume group, the `vgcreate` command may fail and display a message that the extent size is too small or that the VGRA is too big. In this situation, you must choose a larger extent size and run `vgcreate` again.

Increasing the extent size increases the data area marked stale when a write to a mirrored logical volume fails. That can increase the time required to resynchronize stale data. Also, more space may be allocated to each logical volume because the space is allocated in units of extent size. Therefore, the optimal extent size is the smallest value that can be used to successfully create the volume group with the desired configuration parameters.

The minimum extent size for a volume group is calculated using the maximum number of logical volumes (MAXLVs) and physical volumes (MAXPVs) in the volume group and the maximum number of physical extents (MAXPXs) per each physical volume.

For a volume group with bootable physical volumes, the metadata must fit within 768 KB. Therefore, running `vgcreate` with a set of values for MAXLVs, MAXPVs and MAXPXs that succeed on a volume group without bootable physical volumes may fail on a volume group with bootable physical volumes. In this situation, if you must add a bootable physical volume to a volume group, recreate the volume group by giving lesser values for these arguments. By far the biggest factor in the size of the metadata is the values for MAXPVs and MAXPXs. Alternatively, convert the bootable physical volume to a normal physical volume by running `pvcreate` on that physical volume without the `-B` option and then adding it to the volume group. For a physical volume that is already part of a volume group, you can use `vgmodify` to change a physical volume from a bootable to a normal physical volume.

Sample Shell Script

The following shell script creates and compiles a small program that displays the minimum extent size for a given volume group:

```
#!/usr/bin/sh
cat << EOF > vgrasize.c
#include <stdio.h>

#define BS 1024 /* Device block Size */
#define roundup(val, rnd) (((val + rnd - 1) / rnd) * rnd)

main(int argc, char *argv[])
{
    int i, length, lvs, pvs, pxs;
    if (argc != 4) {
        /* Usage example:
         * Maximum LVs in the VG = 255
         * Maximum PVs in the VG = 16
         * Maximum extents per PV = 2500
         *
         * $ vgrasize 255 16 2500
         */
        printf("USAGE: %s MAXLVs MAXPVs MAXPXs\n", argv[0]);
        exit(1);
    }
    lvs = atoi(argv[1]);
    pvs = atoi(argv[2]);
    pxs = atoi(argv[3]);
    length = 16 + 2 * roundup(2 +
                                (roundup(36 + ((3 * roundup(pvs, 32)) / 8) +
                                         roundup(pxs, 8) / 8) * pvs, BS) +
```

```

        roundup(16 * lvs, BS) +
        roundup(16 + 4 * pxs, BS) * pvs) / BS, 8);

if (length > 768) {
    printf("Warning: A bootable PV cannot be added to a VG \n"
    "created with the specified argument values. \n"
    "The metadata size %d Kbytes, must be less \n"
    "than 768 Kbytes.\n"
    "If the intention is not to have a boot disk in this \n"
    "VG then do not use '-B' option during pvcreate(1M) \n"
    "for the PVs to be part of this VG. \n", length);
}

length = roundup(length, 1024) / 1024;

if (length > 256 ) {
    printf("Cannot configure a VG with the maximum values"
    " for LVs, PVs and PXs\n");
    exit(1);
}

for (i = 1; i < length ; i = i << 1) { }

printf("\nMinimum extent size for this configuration = %d MB\n", i);

exit(0);
}
EOF
make vgrasize

```

Choosing an Optimal Extent Size for a Version 2.x Volume Group

Like Version 1.0 Volume Groups, there is a relationship between extent size and maximum volume group size. There is also a limitation on the number of extents an individual volume group can contain. To determine the proper size, the `vgcreate` command has a new `-E` option. This option displays the maximum volume group size based on the given physical extent size or minimum physical extent size based on the maximum volume group size.

Example

After you know the VG size you want to provision for, use `vgcreate` with the `-E` option to determine the minimum extent size required to achieve it.

What is the minimum extent size to provision a 2.0 VG for 1 PB?

```
# vgcreate -V 2.0 -E -s 1p
Max_VG_size=1p:extent_size=32m
```

What is the maximum 2.0 volume group size with an extent size of 16MB?

```
#vgcreate -V 2.0 -E -s 16
Max_VG_size=512t:extent_size=16m
```

What is the minimum extent size to provision a 2.1 VG for 2 PB?

```
# vgcreate -V 2.1 -E -s 2p
Max_VG_size=2p:extent_size=64m
```

What is the maximum 2.1 volume group size with an extent size of 32MB?

```
# vgcreate -V 2.1 -E -s 32
Max_VG_size=1p:extent_size=32m
```

D Striped and Mirrored Logical Volumes

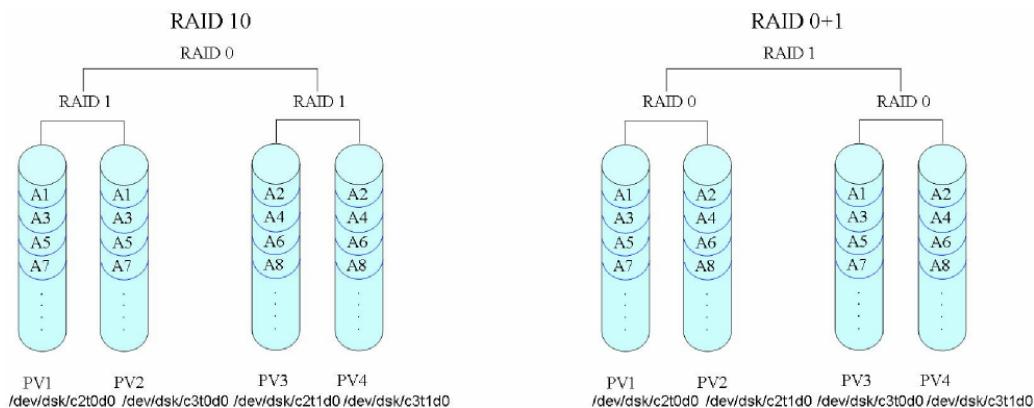
This appendix provides more details on striped and mirrored logical volumes. It describes the difference between standard hardware-based RAID and LVM implementation of RAID.

Summary of Hardware Raid Configuration

RAID 0, commonly referred to as *striping*, refers to the segmentation of logical sequences of data across disks. RAID 1, commonly referred to as *mirroring*, refers to creating exact copies of logical sequences of data. When implemented in a device hardware, RAID 10 (or RAID 1+0) and RAID 01 (or RAID 0+1) are nested RAID levels. The difference between RAID 0+1 and RAID 1+0 is the location of each RAID system: RAID 0+1 is a mirror of stripes whereas RAID 1+0 is a stripe of mirrors. [Figure 6](#) shows the RAID 10 and RAID 01 configurations (A1, A2...Ax are stripe chunks of a logical volume).

With a hardware-based RAID 10 configuration, I/O operations are striped first then each strip is mirrored. With hardware-based RAID 01, I/Os are mirrored first then striped. RAID 10 and RAID 01 can have the same physical disk layout.

Figure 6 Hardware Configuration: Raid 10 and Raid 01



The advantages of hardware-based RAID 10 over RAID 01:

- When one disk fails and is replaced, only the amount of data on this disk needs to be copied or re-synchronized.
- RAID 10 is more tolerant to multiple disk failures before data becomes unavailable.

The advantages of hardware-based RAID 01 over RAID 10:

- Simpler to configure striped volumes and then extend mirroring.
- Able to split the mirror copy and have two usable volume sets

LVM Implementation of RAID Levels in HP-UX

LVM implementation of RAID management is different from the hardware based solutions because it does not nest the RAID levels, but processes them simultaneously. Typically with hardware solutions, you create a LUN with a RAID level and the RAID functions are stacked. LVM provides more flexibility on how logical volumes are created amongst a set of disks as compared to hardware solutions.

LVM allocates the physical extents for striped and mirrored logical volumes in sets of stripe width multiplied by the number of copies of the data. For example, if the logical volume is 1-way mirrored and striped across two disks, extents are allocated to the logical volume, four at a time. LVM enforces that the physical extents of a single set are from different physical volumes. Within this

set, the logical extents are stripped and mirrored to obtain the data layout displayed in [Figure 6 \(page 157\)](#).

Striping and mirroring in LVM combines the advantages of the hardware implementation of RAID 1+0 and RAID 0+1, and provides the following benefits:

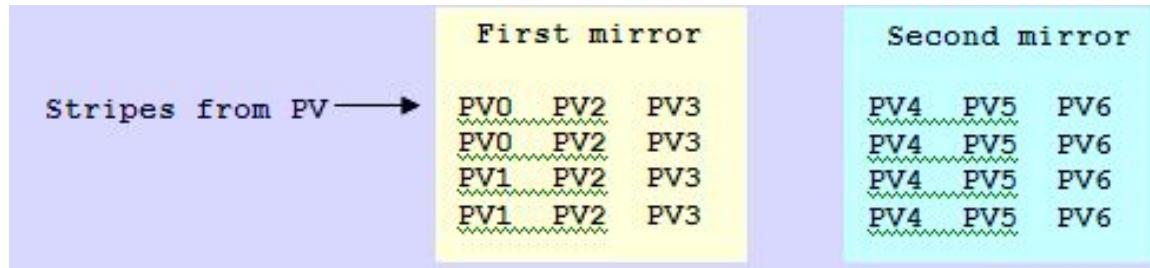
- Better write performance. Write operations take place in parallel and each physical write operation is directed to a different physical volume.
- Excellent performance for read. Even in the case where several disks are out of service, the read of a stripe can be done in parallel on different physical volumes with one I/O operation per physical volume.
- High availability of data. With multiple copies of the user data residing on different physical volumes, LVM avoids single point of failure ensuring high availability.

LVM Striped and Mirrored Logical Volume Configuration

You can create a new striped and mirrored logical volume from scratch, or extend (add mirror to) an existing striped logical volume. With the use of physical volume groups (PVGs), mirroring and striping can be directed to specific physical volumes. Without the constraint created by the nesting of levels in hardware based RAID solutions, LVM is combining the striping and mirroring processing, which allows LVM to provide the best of both RAID 10 and RAID 01.

To create a pure RAID 10, ensure that the physical volumes have enough extents in terms of size and stripe width to accommodate the logical volume. Otherwise, LVM optimizes the extent allocation to enable you to create a striped and mirrored logical volume as long as there are enough free extents in the volume group. The logical volume layout then uses a slight variation of RAID 10 and all the benefits of RAID 10 are retained. For example, [Figure 7](#) shows a volume group with seven physical volumes and a logical volume striped across three disks with one mirror copy.

Figure 7 Example of LVM's Implementation of Raid 10 Variation



If the first physical volume used in the creation has enough extents (size of $LV/3$) to create RAID 10, only six physical volumes are used, resulting in a strict RAID 1 0 configuration. But if one of the physical volumes, for example PV0, does not have enough free extents, LVM uses another physical volume, PV1, to complete the operation. PV0 was short by two extents so two extents were allocated from PV1. LVM created the logical volume using seven physical volumes resulting in a slight variation of RAID 10.

You can extend your existing striped logical volumes to a striped and mirrored configuration easily, using the `lvextend` command to add mirrors. You cannot convert existing mirrored logical volumes to a striped and mirrored configuration. You can create any new logical volumes with striped and mirrored configuration as shown in ["Examples" \(page 158\)](#).

Examples

Use any of the following procedures to create a striped and mirrored logical volume:

- To create a logical volume of size 90MB striped across two physical volumes with one mirror copy and stripe size of 64 KB, enter:

```
#lvccreate -L 90 -i 2 -I 64 -m 1 /dev/vgtest
```

NOTE: Striping with mirroring always uses strict allocation policies where copies of data do not exist on the same physical disk. This results in a configuration similar to the RAID 01 as illustrated in [Figure 7 \(page 158\)](#).

- To create a logical volume of size 90MB striped across two physical volumes with one mirror copy and stripe size of 64 KB and to create the mirror copies on specific disks (configuration equivalent to RAID 01 as illustrated in [Figure 7 \(page 158\)](#), enter):

```
#lvcreate -L 90 -i 2 -I 64 /dev/vgtest
#lvedit -m 1 /dev/vgtest/lvol1 /dev/disk/disk3 /dev/disk/disk4
Contents of /etc/lvmpvg
VG /dev/vgtest
PVG PVG0
/dev/disk/disk1
/dev/disk/disk2
PVG PVG1
/dev/disk/disk3
/dev/disk/disk4
```

- Adding a mirror to existing striped logical volumes.

The existing logical volume has the striped property as shown by the `lvdisplay` output:

```
# lvdisplay /dev/vgtest/lvol1
--- Logical volumes ---
LV Name          /dev/vgtest/lvol1
VG Name          /dev/vgtest
LV Permission    read/write
LV Status        available/syncd
Mirror copies  0
Consistency Recovery
Schedule      striped
LV Size (Mbytes) 1024
Current LE       256
Allocated PE     256
Stripes        2
Stripe Size (Kbytes) 64
Bad block        on
Allocation       strict
IO Timeout (Seconds) default
```

To get a striped and mirrored configuration, extend the logical volume as follows:

```
# lvedit -m 1 /dev/vgtest/lvol1
```

Note that the volume group must have enough physical volumes and extents to accommodate the mirror copy.

Now the `lvdisplay` output shows the logical volume is striped and mirrored:

```
# lvdisplay /dev/vgtest/lvol1
--- Logical volumes ---
LV Name          /dev/vgtest/lvol1
VG Name          /dev/vgtest
LV Permission    read/write
LV Status        available/syncd
Mirror copies  1
Consistency Recovery
Schedule      striped
LV Size (Mbytes) 1024
Current LE       256
Allocated PE     256
Stripes        2
Stripe Size (Kbytes) 64
Bad block        on
Allocation       strict
IO Timeout (Seconds) default
```

Compatibility Note

Releases prior to HP-UX 11i v3 only support striped or mirrored logical volumes and do not support combination of striped and mirrored logical volumes. If a logical volume using simultaneous mirroring and striping is created on HP-UX 11i v3, attempts to import or activate its associated volume group fails on a previous HP-UX release.

To import the volume group with striped and mirrored logical volume to releases prior to HP-UX 11i v3, you must remove the incompatible logical volumes or reduce them to no mirrors.

NOTE: Striped and mirrored logical volume is supported on HP-UX 11i v2 with the PHCO_36744, PHKL_36745 and PHCO_36746 patches. With these patches or superseding patches installed on a HP-UX 11i v2 system, you can successfully import and activate a volume group with striped and mirrored logical volume created on HP-UX 11i v3.

E LVM I/O Timeout Parameters

When LVM receives an I/O to a logical volume, it converts this **logical I/O** to **physical I/Os** to one or more physical volumes from which the logical volume is allocated. There are two LVM timeout values which affect this operation:

- Logical volume timeout (LV timeout).
- Physical volume timeout (PV timeout).

Logical Volume Timeout (LV timeout)

LV timeout controls how long LVM retries a logical I/O after a recoverable physical I/O error.

LVM timeout can be configured for a specific logical volume using `lvchange`. See the `lvchange(1M)` manpage for details. Also, note the following:

- The default LV timeout is zero which means infinite timeout to retry physical I/Os with recoverable errors.
- If the LV timeout is changed, it should be kept greater than or equal to the PV timeout.
- Set higher LV timeout for greater I/O resiliency. Set lower LV timeout for faster I/O failure detection (which, in turn, facilitates faster failover).

Physical Volume Timeout (PV timeout)

PV timeout is the time budget set by LVM for each physical I/O originating from LVM. The mass storage stack underlying LVM manages completion of the physical I/O subject to this timeout, resulting in success or failure for the request.



WARNING! While mass storage stack strives to ensure the physical I/O completion by this PV timeout budget, there could be situations where the I/O completion exceeds this timeout. It could be due to failures in recovery mechanism of underlying transport layers.

PV timeout can be configured for a specific physical volume using `pvchange`. See the `pvchange(1M)` manpage for details. Also, note the following:

- With 11i v3, native multi-pathing is enabled by default in mass storage stack. When it is enabled:
 - For clarity, it is recommended that only agile DSFs and cluster DSFs should be used in LVM when configuring or adding physical volumes to the volume group. Native multi-pathing will still be in effect even if legacy DSFs are specified and this may cause confusion
 - PV timeout is the time budget given to mass storage stack for its I/O retries on all of the PV links to this physical volume. However upon the expiration of this PV timeout for a particular physical I/O, the mass storage stack will return this I/O back to LVM even if not all the PV links have been tried within the duration specified by the given PV timeout value. For example, if PV timeout is 30 seconds and if the physical volume has 3 links, the mass storage stack can utilize at the most 30 seconds to retry any I/O on as many links as possible.
- The PV timeout can be from 2 seconds to a maximum of 300 seconds. The default timeout, shown by a value of "0", is 30 seconds. However, to ensure reliable handling of device failure, the driver stack requires a PV timeout of no less than 15 seconds. Using values lower than 15 seconds can lead to false failure detection and possibly longer delays as LVM is immediately recovering from a failure. Ideal PV timeout value for a physical volume depends on several factors including number of links to the physical volume, quality/speed/recovery-time of the transport links, physical volume model/type, and so on.

Timeout Differences: 11i v2 and 11i v3

Since native multi-pathing is included in 11i v3 mass storage stack and it is enabled by default, the LVM timeout concepts may vary between 11i v2 and 11i v3 in certain cases.

- Meaning of PV timeout.

In 11i v2, LVM utilizes the configured PV timeout fully for a particular PV link to which it is set. If there is any I/O failure, LVM will retry the I/O on a next available PV link to the same physical volume with the new PV timeout budget.

But in 11i v3, when native multi-pathing is enabled, the mass storage stack does the I/O retries on all/subset of the available PV links within the single configured PV timeout period. In spite of this difference, the default PV timeout is retained as 30 seconds in 11i v3 to guarantee backward compatibility.

Care should be taken in selecting an appropriate value for the PV timeout in 11i v3, so that retries may be attempted on all available PV links within the timeout period.

- Timeout considerations for mirrored LV environment.

With 11i v3, in certain situations, LVM I/Os to a mirrored Logical volume may take longer time to complete (comparing to 11i v2) than the configured LVM LV timeout. It happens when one or more PV links goes offline while write I/Os to a mirrored logical volume are in-progress. The situation gets worse when DasProvider, SAM or lvnboot runs while PV link goes offline. However this situation remains only for some period of time (usually up to transient_secs but could be higher in some cases), until the associated LUNs transitions out of transient state either to online or offline state, after which it is recovered automatically. This could be due to the higher timeout value set in the SCSI tunable transient_secs and path_fail_secs. If the physical volume has few links underneath, reducing these parameters may result in faster I/O completion.

transient_secs and path_fail_secs are mass storage stack tunable parameters that can be configured through the scsimgr and scsimgr_esdisk commands. Regardless of these parameter settings, the mass storage stack never holds onto any physical I/O originating from LVM for more than the PV timeout budget. These two parameter settings never replace the PV timeout but control the threshold of a LUN and Lunpath before failing I/Os.

⚠ **CAUTION:** Lowering the parameters transient_secs and path_fail_secs may result in less I/O resiliency.

For more details on these two tunable parameters, see the *scsimgr(1M)* and *scsimgr_esdisk(7)* manpages. Also refer to the *Path Recovery Policies and Managing Loss of Accessibility to a LUN* sections of the *HP-UX 11i v3 Native Multi-Pathing for Mass Storage* whitepaper on <http://www.hp.com/go/hpux-core-docs>.

NOTE: When you use disks with a single link, such as most internal disks, with 11i v3 you can set SCSI parameters transient_secs and path_fail_secs to the PV timeout value to get similar 11i v2 timeout behavior in 11i v3 for the mirrored LV I/Os.

F Warning and Error Messages

This appendix lists some of the warning and error messages reported by LVM. For each message, the cause is described and an action is recommended.

Matching Error Messages to Physical Disks and Volume Groups

Often an error message contains the device number for a device, rather than the device file name. For example, you might see the following message in /var/adm/syslog/syslog.log:

```
Asynchronous write failed on LUN (dev=0x3000015)
IO details : blkno : 2345, sector no : 23
```

To map this error message to a specific disk, search under the /dev directory for a device file with a device number that matches the printed value. More specifically, search for a file whose minor number matches the lower six digits of the number following dev:. The device number in this example is 3000015; its lower six digits are 000015, so search for that value using the following command:

```
# ll /dev/*dsk | grep 000015
brw-r----- 1 bin          sys            3 0x000015 May 26 20:01 disk43
crw-r----- 1 bin          sys            23 0x000015 May 26 20:01 disk43
```

Use the pvdisplay command to determine which volume group contains this physical volume as follows:

```
# pvdisplay /dev/disk/disk43 | grep "VG Name"
VG Name                  /dev/vgtest
```

If the pvdisplay command fails, search for the physical volume in the LVM configuration files /etc/lvmtab and /etc/lvmtab_p as follows:

```
# lvmadm -l -v 1.0
--- Version 1.0 volume groups ---
VG Name /dev/vg00
PV Name /dev/disk/disk36_p2

VG Name /dev/vgtest
PV Name /dev/disk/disk43

VG Name /dev/vg03
PV Name /dev/dsk/c2t2d0
```

If your version of lvmadm does not recognize the -l option, use the strings command as follows:

```
# strings /etc/lvmtab | more
/dev/vg00
/dev/disk/disk36_p2
/dev/vgtest
/dev/disk/disk43
/dev/vg03
/dev/dsk/c2t2d0
```

Based on the output of these commands, the error message refers to physical volume /dev/disk/disk43, which belongs to volume group vgtest.

Similarly, some LVM error messages in /var/adm/syslog/syslog.log contain the device number for a volume group. For example:

```
LVM: VG 128 0x002000: Lost quorum.
```

The major number 128 indicates that this is a Version 2.x volume group. A Version 1.0 volume group has a major number of 64. To map this error message to a volume group, search under the /dev directory for a volume group device file with a device number that matches the major and minor numbers. In this example, the major number is 128 and the minor number is 0x002000, so search for those values using the following command:

```
# ll /dev/*group | grep 128.0x002000
crw-r----- 1 root    sys   128 0x002000 Jan  7 08:27 /dev/vgtest2/group
```

The example error message refers to the Version 2.x volume group vgtest2.

Messages For All LVM Commands

Message Text

```
vvcfgbackup: /etc/lvmtab is out of date with the running kernel:  
Kernel indicates # disks for "/dev/vgname"; /etc/lvmtab has # disks.  
Cannot proceed with backup.
```

Cause

The number of current physical volumes (Cur PV) and active physical volumes (Act PV) are not the same. Cur PV and Act PV must always agree for the volume group. This error also indicates that /etc/lvmtab or /etc/lvmtab_p, which is used to match physical volumes to a volume group, is out of date with the LVM data structures in memory and on disk.

Recommended Action

Try to locate any missing disks. For each of the disk in the volume group, use `ioscan` and `diskinfo` to confirm that the disk is functioning properly.

lvchange(1M)

Message Text

```
"m": Illegal option.
```

Cause

The system does not have HP MirrorDisk/UX installed.

Recommended Action

Install HP MirrorDisk/UX.

lvextend(1M)

Message Text

```
lvextend: Not enough free physical extents available.  
Logical volume "/dev/vgname/lvname" could not be extended.  
Failure possibly caused by strict allocation policy
```

Cause

There is not enough space in the volume group to extend the logical volume to the requested size. This is typically caused by one of the following situations:

1. There are not enough free physical extents in the volume group. Run `vgdisplay` to confirm the number of available physical extents, and multiply that number by the extent size to determine the free space in the volume group. For example:

```
# vgdisplay vg00  
--- Volume groups ---  
VG Name                  /dev/vg00  
VG Write Access           read/write  
VG Status                available  
Max LV                  255  
Cur LV                  10  
Open LV                  10  
Max PV                  16  
Cur PV                  1  
Act PV                  1
```

Max PE per PV	4350
VGDA	2
PE Size (Mbytes)	4
Total PE	4340
Alloc PE	3740
Free PE	600
Total PVG	0
Total Spare PVs	0
Total Spare PVs in use	0
VG Version	1.0
VG Max Size	1082g
VG Max Extents	69248

In this example, the total free space is 600 physical extents of 4 MB, or 2400 MB.

2. The logical volume is mirrored with a strict allocation policy, and there are not enough extents on a separate disk to comply with the allocation policy. To confirm this, run `lvdisplay` to determine which disks the logical volume occupies, and then check whether there is sufficient space on the other disks in the volume group.
3. In a SAN environment, one of the disks was dynamically increased in size. LVM did not detect the asynchronous change in size.

Recommended Action

1. Choose a smaller size for the logical volume, or add more disk space to the volume group.
2. Choose a smaller size for the logical volume, or add more disk space to the volume group. Alternatively, free up space on an available disk using `pvmove`.
3. Use the `vgmodify` command to detect the disk size change and incorporate the new space into the volume group.

Message Text

"m": Illegal option.

Cause

The system does not have HP MirrorDisk/UX installed.

Recommended Action

Install HP MirrorDisk/UX.

lvlnboot(1M)

Message Text

```
lvlnboot: Unable to configure swap logical volume.
Swap logical volume size beyond the IODC max address.
```

Cause

The boot disk firmware cannot access the entire range of the swap logical volume. This happens with older host bus adapters when primary swap is configured past 4 GB on the disk.

Recommended Action

Upgrade the system firmware or use a newer host bus adapter that supports block addressing. If neither of these actions succeeds, reduce the size of the primary swap logical volume so that it does not exceed 4 GB.

pvchange(1M)

Message Text

Unable to detach the path or physical volume via the pathname provided.
Either use pvchange(1M) -a N to detach the PV using an attached path
or detach each path to the PV individually using pvchange(1M) -a n

Cause

The specified path is not part of any volume group, because the path has not been successfully attached to the otherwise active volume group it belongs to.

Recommended Action

Check the specified path name to make sure it is correct. If the error occurred while detaching a physical volume, specify a different path. If it is not clear whether any path was attached before, individually detach each path to the physical volume using pvchange with the -a n option.

Message Text

Warning: Detaching a physical volume reduces the availability of data within the logical volumes residing on that disk.

Prior to detaching a physical volume or the last available path to it, verify that there are alternate copies of the data available on other disks in the volume group. If necessary, use pvchange(1M) to reverse this operation.

Cause

This warning is advisory. It is generated whenever a path or physical volume is detached.

Recommended Action

None.

vgcfgbackup(1M)

Message Text

Invalid LVMREC on Physical Volume *pvname*

Cause

The LVM header on the disk is incorrect. This can happen when an existing LVM disk is overwritten with a command like dd or pvcreate. If the disk is shared between two systems, one system might not be aware that the disk is already in a volume group. The corruption can also be caused by running vgchgid incorrectly when using BC split volumes.

Recommended Action

Restore a known good configuration to the disk using vgcfgrestore. Be sure to use a valid copy dated before the first occurrence of the problem. For example:

```
# vgcfgrestore -n vgname pvname
```

vgcfgrestore(1M)

Message Text

Cannot restore Physical Volume *pvname*
Detach the PV or deactivate the VG, before restoring the PV.

Cause

The `vgcfgrestore` command was used to initialize a disk that already belongs to an active volume group.

Recommended Action

Detach the physical volume or deactivate the volume group before attempting to restore the physical volume. If the disk may be corrupted, detach the disk and mark it using `vgcfgrestore`, then attach it again without replacing the disk. This causes LVM to reinitialize the disk and synchronize any mirrored user data mapped there.

`vgchange(1M)`

Message Text

```
vgchange: WARNING: The "lvmp" driver is not loaded.
```

Cause

You are activating a Version 2.x volume group, and the kernel driver for Version 2.x volume groups is not loaded.

Recommended Action

Load the `lvmp` driver, as follows:

```
# kcmodule lvmp=best
    ==> Update the automatic 'backup' configuration first? n
        * Future operations will ask whether to update the backup.
        * The requested changes have been applied to the currently
            running configuration.
Module          State   Cause      Notes
lvmp (before)   unused      loadable, unloadable
           (now)     loaded best
           (next boot) loaded explicit
```

Message Text

```
vgchange: Warning: Couldn't attach to the volume group
physical volume "pvname":
Illegal byte sequence
vgchange: Couldn't activate volume group "vgname":
Quorum not present, or some physical volume(s) are missing.
```

Cause

You are activating a Version 2.x volume group, and your operating system release does not support Version 2.x volumes.

Recommended Action

Update your system to the March 2008 release of HP-UX 11i Version 3 or a newer release.

Message Text

```
Warning: couldn't query physical volume "pvname":
The specified path does not correspond to physical volume
attached to this volume group
Couldn't query the list of physical volumes.
```

Cause

This error has the following possible causes:

1. The disk was missing when the volume group was activated, but was later restored. This typically occurs when a system is rebooted or the volume group is activated with a disk missing, uncabled, or powered down.
2. The disk LVM header was overwritten with the wrong volume group information. If the disk is shared between two systems, one system might not be aware that the disk was already in a volume group. To confirm, check the volume group information using the `dump_lvmtab` command (available from your HP support representative) and look for inconsistencies. For example:

```
# dump_lvmtab -s | more
SYSTEM : 0x35c8cf58
TIME   : 0x3f9acc69 : Sat Oct 25 15:18:01 2003
FILE   : /etc/lvmtab
HEADER : version:0x03e8 vgnum:7
VG[00]  VGID:35c8cf58 3dd13164 (@0x00040c) pvnum:2 state:0 /dev/vg00
        (00) VGID:35c8cf58 3dd13164 PVID:35c8cf58 3dd13164 /dev/dsk/c0t6d0
        (01) VGID:35c8cf58 3dd13164 PVID:35c8cf58 3dda4694 /dev/dsk/c4t6d0
VG[01]  VGID:065f303f 3e63f01a (@0x001032) pvnum:92 state:0 /dev/vg01
        (00) !VGID:35c8cf58 3f8df316 PVID:065f303f 3e63effa /dev/dsk/c40t0d0
        (01) !VGID:35c8cf58 3f8df316 PVID:065f303f 3e63effe /dev/dsk/c40t0d4
        (02) !VGID:35c8cf58 3f8df316 PVID:065f303f 3e63f003 /dev/dsk/c40t1d0
...

```

In this example, the VGIDs for the disks in /dev/vg01 are not consistent; inconsistencies are marked !VGID.

Recommended Action

1. Use `ioscan` and `diskinfo` to confirm that the disk is functioning properly. Reactivate the volume group using the following command:


```
# vgchange -a y vgname
```
2. There are several methods of recovery from this error. If you are not familiar with the commands outlined in the following procedures, contact your HP support representative for assistance.
 - a. Restore a known good configuration to the disks using `vgcfgrestore`. Be sure to use a valid copy dated before the first occurrence of the problem. For example:


```
# vgcfgrestore -n vgname pvname
```
 - b. Recreate the volume group and its logical volumes, restoring the data from the most current backup. See "[Creating a Volume Group](#)" (page 44) and "[Creating a Logical Volume](#)" (page 52).
 - c. Export and reimport the volume group, as described in "[Exporting a Volume Group](#)" (page 57) and "[Importing a Volume Group](#)" (page 58). For example:


```
# vgexport -m vgname.map -v -f vgname.file /dev/vgname
# vgimport -m vgname.map -v -f vgname.file /dev/vgname
```

Message Text

```
vgchange: Couldn't set the unique id for volume group "/dev/vgname"
```

Cause

There are multiple LVM group files with the same minor number.

Recommended Action

List the LVM group files. If there are any duplicate minor numbers, export one of the affected volume groups, optionally create a new group file with a unique minor number, and reimport the volume group. If you are not familiar with this procedure, contact your HP support representative for assistance.

```
# ll /dev/*/*group
# vgexport -m vgname.map -v -f vgname.file /dev/vgname
```

```
# mkdir /dev/vgname
# mknod /dev/vgname/group c 64 unique_minor_number
# vgimport -m vgname.map -v -f vgname.file /dev/vgname
```

vgcreate(1M)

Message Text

vgcreate: "/dev/vgname/group": not a character device.

Cause

The volume group device file does not exist, and this version of the `vgcreate` command does not automatically create it.

Recommended Action

Create the directory for the volume group and create a group file, as described in “[Creating the Volume Group Device File](#)” (page 44).

Message Text

```
vgcreate: Volume group "/dev/vgname" could not be created:
Error: The physical volume "pvname" contains BDRA.
It cannot be added into volume group "/dev/vgname" because
this version does not support bootable disks.
```

Cause

The physical volume `pvname` is a bootable disk, and `vgname` is a Version 2.0 or Version 2.1 volume group. Version 2.0 and 2.1 volume groups do not support bootable physical volumes.

Recommended Action

Use the `pvcreate` command without the `-B` option to reinitialize the disk, as described in “[Initializing a Disk for LVM Use](#)” (page 43). For example:

```
# pvcreate -f pvname
```

Use the `vgmodify` command to convert the disk to a nonbootable disk, as described in “[Changing Physical Volume Boot Types](#)” (page 84). Alternately, use the `pvcreate` command without the `-B` option to reinitialize the disk, as described in “[Initializing a Disk for LVM Use](#)” (page 43). For example:

```
# pvcreate -f pvname
```

Message Text

```
vgcreate: Volume group "/dev/vgname" could not be created:
VGRA for the disk is too big for the specified parameters.
Increase the extent size or decrease max_PVs/max_LVs and try again.
```

Cause

The volume group reserved area (VGRA) at the front of each LVM disk cannot hold all the information about the disks in this volume group. This error typically occurs if you use disks larger than 100 GB.

Recommended Action

Adjust the volume group creation parameters. Use the `-s` option of the `vgcreate` command to select an extent size larger than 4 MB, or use the `-p` option to select a smaller number of physical volumes. For more information on these options, see [vgcreate\(1M\)](#). For recommendations on extent sizes, see [Appendix C](#) (page 155).

vgdisplay(1M)

Message Text

```
vgdisplay: Couldn't query volume group "/dev/vgname".  
Possible error in the Volume Group minor number;  
Please check and make sure the group minor number is unique.  
vgdisplay: Cannot display volume group "/dev/vgname".
```

Cause

This error has the following possible causes:

1. There are multiple LVM group files with the same minor number.
2. Serviceguard was previously installed on the system and the /dev/slvmvg device file still exists.

Recommended Action

1. List the LVM group files. If there are any duplicate minor numbers, export one of the affected volume groups, optionally create a new group file with a unique minor number, and reimport the volume group. If you are not familiar with this procedure, contact your HP support representative for assistance.

```
# ll /dev/*group  
# vgexport -m vgname.map -v -f vgname.file /dev/vgname  
# mkdir /dev/vgname  
# mknod /dev/vgname/group c 64 unique_minor_number  
# vgimport -m vgname.map -v -f vgname.file /dev/vgname
```

2. Remove the /dev/slvmvg device file and re-create the /etc/lvmtab and /etc/lvmtab_p files using the following commands:

```
# rm /dev/slvmvg  
# mv /etc/lvmtab /etc/lvmtab.old  
# mv /etc/lvmtab_p /etc/lvmtab_p.old  
# vgscan -v
```

Message Text

```
vgdisplay: /etc/lvmtab: No such file or directory  
vgdisplay: No volume group name could be read from "/etc/lvmtab".  
vgdisplay: /etc/lvmtab: No such file or directory  
vgdisplay: No volume group name could be read from "/etc/lvmtab".
```

Cause

One of the LVM configuration files, either /etc/lvmtab or /etc/lvmtab_p, is missing.

Recommended Action

Create the /etc/lvmtab and /etc/lvmtab_p files using the following command:

```
# vgscan -v
```

For additional information on the vgscan command and its option, see [vgscan\(1M\)](#).

Message Text

```
Warning: couldn't query physical volume "pvname":  
The specified path does not correspond to physical volume  
attached to this volume group  
Couldn't query the list of physical volumes.
```

Cause

The possible causes of this error are described under the "[vgchange\(1M\)](#)" (page 167) error messages.

Recommended Action

See the recommended actions under the “[vgchange\(1M\)](#)” (page 167) error messages.

vgextend(1M)

Message Text

```
vgextend: Not enough physical extents per physical volume.  
Need: #, Have: #.
```

Cause

The disk size exceeds the volume group maximum disk size. This limitation is defined when the volume group is created, as a product of the extent size specified with the `-s` option of `vgcreate` and the maximum number of physical extents per disk specified with the `-e` option. Typically, the disk is successfully added to the volume group, but not all of the disk is accessible.

Recommended Action

Use the `vgmodify` command to adjust the maximum number of physical extents per disk. Alternately, you can re-create the volume group with new values for the `-s` and `-e` options.

vgimport(1M)

Message Text

```
vgimport: "/dev/vgname/group": not a character device.
```

Cause

The volume group device files do not exist, and this version of the `vgimport` command does not automatically create them.

Recommended Action

Create the directory for the volume group, and create a group file, as described in “[Creating the Volume Group Device File](#)” (page 44).

Message Text

```
Verification of unique LVM disk id on each disk in the volume group  
/dev/vgname failed.
```

Cause

There are two possible causes for this message:

1. The `vgimport` command used the `-s` option and two or more disks on the system have the same LVM identifier; this can happen when disks are created with BC copy or cloned with `dd`.
2. LVM was unable to read the disk header; this can happen when you create new logical units on a SAN array.

Recommended Action

1. Use `vgimport` without the `-s` option. Alternately, use `vgchgid` to change the LVM identifiers on copied or cloned disks.
2. Retry the `vgimport` command.

vgmodify(1M)

Message Text

Error: Cannot reduce *max_pv* below *n+1* when the volume group is activated because the highest *pvkey* in use is *n*.

Cause

The command is trying to reduce *max_pv* below the highest *pvkey* in use. This is disallowed when a version 1.0 volume group is activated since it requires the compacting of *pvkeys*.

Recommended Action

Try executing the `vgmodify` operation with *n+1* PVs. The permissible *max_pv* values can be obtained using the `vgmodify -t` option (used with and without the `-a` option).

If *max_pv* needs to be reduced below *n+1*, deactivate the volume group before running the `vgmodify` operation.

vgversion(1M)

Message Text

Error: The Physical Volume *pvname* has bad blocks relocated on it.

Error: Volume Group version 2.x does not support bad block relocation.

Cause

Version 2.x volume groups cannot handle bad blocks on physical volumes. If any physical volume of a Version 1.0 volume group has relocated bad blocks, `vgversion` fails the migration.

Recommended Action

This physical volume should be replaced by a new disk. For more information on how to replace a disk, see "[Step 6: Replacing a Bad Disk \(Persistent DSFs\)](#)" (page 129) or "[Step 7: Replacing a Bad Disk \(Legacy DSFs\)](#)" (page 138).

Message Text

Error: Could not find any free minor number for volume group version 2.x

Cause

`vgversion` could not find a free minor number to create the group file for the volume group because all the minor numbers in the supported range for Version 2.x are used by existing Version 2.x volume groups. This can occur when the maximum supported number of Version 2.x volume groups is already created or when some group files exist, but do not correspond to any existing volume group.

Recommended Action

Determine the number of Version 2.x volume groups configured on the system using `lvmadm -l`. If this is less than the maximum supported number of Version 2.x volume groups, identify the group files that do not correspond to any of these volume groups. Delete one of these group files and re-run the `vgversion` command. If you are not familiar with the procedure of finding an unused group file, contact your HP support representative for assistance.

Message Text

Error: The Physical Volume *pvname* has the *pvkey #*, which is not supported by volume group version 2.x

Cause

The *pvkey* of a physical volume can range between 0 to a value equal to one less than the maximum supported number of physical volumes for a volume group version. If the *pvkey* of the physical volume is not in this range for the target Version 2.x volume group, *vgversion* fails the migration.

Recommended Action

1. Use *lvmadm* to determine the maximum supported number of physical volumes for the target volume group version.
2. Use *vgcfgrestore -l -v* to check if there is any unused *pvkey* in the range supported by the target volume group version.
3. If a valid *pvkey* is found in this range, add a disk to the volume group and verify that this physical volume has a *pvkey* in the supported range. Use *pvmove* to move all extents from the physical volume *pvname* to this new physical volume, so that the physical volume *pvname* can be removed from the volume group.
4. If the volume group already contains more than the maximum supported number of physical volumes for the target volume group version, then the migration cannot be achieved.

If you are not familiar with any of the steps mentioned above, contact your HP support representative for assistance.

Message Text

Error: The Logical Volume /dev/vgname/lvname has the lv number #, which is not supported by volume group version 2.x.

Cause

The logical volume number of a logical volume can range between 1 to the maximum supported number of logical volumes for a volume group version. If the logical volume number of the logical volume is not in this range for the target Version 2.x volume group, *vgversion* fails the migration.

Recommended Action

1. Use *lvmadm* to determine the maximum supported number of logical volumes for the target volume group version.
2. For each logical volume device file in the /dev/vgname directory, determine the logical volume number. (Section “[Device Number Format](#)” ([page 15](#)) describes the format of the logical volume device file.) Based on the volume group version, decode the bits associated with the logical volume number and convert this value to decimal notation.
3. Check if there is any unused logical volume number in the range supported by the target volume group version. If a valid number is found in this range, create a logical volume with the same properties as that of /dev/vgname/lvname. Copy the entire data from this logical volume to the one created newly, so that the logical volume /dev/vgname/lvname can be removed from the volume group.
4. If the volume group already contains more than the maximum supported number of logical volumes for the target volume group version, then the migration will fail.

If you are not familiar with any of the steps mentioned above, contact your HP support representative for assistance.

Log Files and Trace Files: /var/adm/syslog/syslog.log

LVM does not have a dedicated log file or trace file. It logs errors and warnings to /var/adm/syslog/syslog.log. Below are LVM messages that might appear in this file.

Message Text

LVM: Begin: Contiguous LV (VG *mmm* 0x00n000, LV Number: *p*) movement:
LVM: End: Contiguous LV (VG *mmm* 0x00n000, LV Number: *p*) movement:

Cause

This message is advisory. It is generated whenever the extents of a contiguous logical volume belonging to version 2.x volume group is moved using pvmove. This message is generated beginning with the September 2009 Update.

Recommended Action

None, if both the Begin and End message appear for a particular contiguous LV.

If the End message for a contiguous LV is missing in the syslog file, it probably means that the pvmove operation could have been aborted while moving the extents of the LV in question. In this case, the contiguous LV in question could have its extents laid out in a non-contiguous manner. If so, it is recommended that the user run lvdisplay -v to check if the allocation policy for the LV is broken. If it is broken, then perform the steps below to make the extents of this LV contiguous:

1. Change the allocation policy of this LV to default (see *lvchange(1M)*).
2. Move the extents such that all the physical extents of the logical volume are contiguous (see *pvmove(1M)*).
3. Change the allocation policy back to contiguous (see *lvchange(1M)*).

Message Text

LVM: VG *mm* 0x*nn*n0000:
Data in one or more logical volumes on PV *nn* 0x0*nn*n000
was lost when the disk was replaced.
This occurred because the disk contained the only copy of the data.
Prior to using these logical volumes, restore the data from backup.

Cause

LVM cannot synchronize the data on a replaced disk automatically, for example, when LVM discovers an unmirrored logical volume residing on a disk that was just replaced.

Recommended Action

Restore the contents of the logical volume from backup.

Message Text

LVM: VG *mm* 0x*nn*n0000: PVLink *nn* 0x0*nn*n000 Detached.

Cause:

This message is advisory. It is generated whenever a disk path is detached.

Recommended Action

None.

Message Text

LVM: vg[*nn*] pv[*nn*] No valid MCR, resyncing all mirrored MWC LVs on the PV

Cause

This message appears when you import a volume group from a previous release of HP-UX. The format of the MWC changed at HP-UX 11i Version 3, so if the volume group contains mirrored logical volumes using MWC, LVM converts the MWC at import time. It also performs a complete resynchronization of all mirrored logical volumes, which can take substantial time.

Recommended Action

None.

Message Text

```
LVM: vg 64 0xnnnnnnn: Unable to register  
for event notification on device 0xnnnnnnnn (1)
```

Cause

This message can be displayed on the first system boot after upgrading to HP-UX 11i Version 3. It is a transient message caused by updates to the I/O configuration. Later in the boot process, LVM registers for event notification again, and succeeds.

Recommended Action

None.

Message Text

```
LVM: WARNING: Snapshot LV (VG mmm 0x00n000, LV Number: p) threshold value reached.  
Please increase the number of pre-allocated extents for this snapshot LV
```

Cause

These are the possible causes:

1. This message is generated when the number of extents remaining in a space-efficient snapshot logical volume's pre-allocated pool falls beyond a certain internally computed threshold. The pre-allocated extents of the space-efficient snapshot are consumed as a result of writes either on its successor logical volume on the snapshot tree or on the snapshot itself. This message is displayed so that the user is informed that further unsharing of data between the snapshot and its successors might end up depleting the extents in the pre-allocated pool and lead to the snapshot and its predecessors being marked as inoperative.
2. Beginning with the HP-UX 11i v3 September 2010 Update, you can enable LVM to automatically increase the number of pre-allocated extents when the threshold is reached. If this feature is enabled, this message will be logged every time LVM tries to increase pre-allocated extents.

Recommended Action

1. Increase the number of pre-allocated extents for the space-efficient snapshot using the `lvextend` command to prevent the snapshot from becoming inoperative.
2. If there is no error reported following this warning message, this message can be ignored. If there is an error, please follow the relevant recommended action for the error. If this error is seen more frequently, you can either increase the number of pre-allocated extends and/or threshold value, as appropriate. Please see the *LVM Snapshot Logical Volumes* white paper for more details.

Message Text

```
vmunix: LVM:ERROR: The task to increase the pre-allocated extents could not be posted  
for this snapshot LV (VG 128 0x000000, LVM Number 3).  
Please check if lvmpud is running.
```

Cause

If the automatic increase of pre-allocated extents is enabled for a space efficient snapshot, the numbers of pre-allocated extents are automatically increased when the threshold value is reached. The `lvmpud` daemon must be running for this to succeed. When `lvmpud` is not running, the above message is logged.

Recommended Action

Start the lvmpud daemon. Refer to the “[Managing the lvmpud Daemon](#)” (page 105) section.

Message Text

```
vmunix: LVM: ERROR: The task posted for increasing the pre-allocated extents failed  
for the snapshot LV (VG 128 0x004000, LV Number: 48).
```

Cause

The above message is logged when LVM fails to increase the pre-allocated extents automatically. This message should not be seen frequently. If seen frequently, it must be a LVM software or configuration issue.

Recommended Action

Increase the number of pre-allocated extents for the space-efficient snapshot using the lvextend command to prevent the snapshot from becoming inoperative.

Message Text

```
lvmpud->lvextend[]: LogicalExtentsNumber is not bigger than current setting.  
LVM: ERROR: The task posted for increasing the pre-allocated extents failed  
for the snapshot LV (VG 128 0x004000, LV Number: 48).
```

Cause

This is a known issue (QXCR1001039721) occurring only in the HP-UX 11i v3 September 2010 Update. This message can be ignored.

Recommended Action

None.

Message Text

```
LVM: Deletion of snapshot LV device [major: nnn, minor:0xnnnnnn] failed since there are  
pending I/Os on the successor LV device [major: nnn, minor:0xnnnnnn] waiting for one  
or more offline PVs to come back online. Please make the PVs available and retry.
```

Or,

```
LVM: Deletion of snapshot LV device [major: nnn, minor:0xnnnnnn] failed since there are  
pending I/Os on the predecessor LV device [major: nnn, minor:0xnnnnnn] waiting for one  
or more offline PVs to come back online. Please make the PVs available and retry.
```

Cause

While deleting a snapshot logical volume, if LVM detects pending I/Os due to unavailability of underlying physical volumes on either the successor or the predecessor logical volumes, LVM will abort the delete operation.

Recommended Action

Make sure that the disk devices being used by the entire snapshot tree (the original logical volume and all of its snapshots) are available and healthy before retrying the delete operation.

Glossary

Agile Addressing	The ability to address a LUN with the same device special file regardless of the physical location of the LUN or the number of paths leading to it. In other words, the device special file for a LUN remains the same even if the LUN is moved from one Host Bus Adaptor (HBA) to another, moved from one switch/hub port to another, presented via a different target port to the host, or configured with multiple hardware paths. Also referred to as persistent binding .
Agile View	The representation of LUNs using the new agile addressing and persistent DSFs, introduced in HP-UX 11i v3.
Allocation Policy	The LVM allocation policy governing how disk space is distributed to logical volumes and how extents are laid out on an LVM disk. LVM allocates disk space in terms of strict vs. non-strict and contiguous vs. noncontiguous. Strict allocation requires that mirror copies reside on different LVM disks. Contiguous allocation requires that no gaps exist between physical extents on a single disk.
Device Special File (DSF)	HP-UX applications access peripheral devices such as tape drives, disk drives, and printers via special files in the /dev directory called Device Special Files (DSFs). Device Special Files allow applications to access devices with minimum knowledge of the system's underlying hardware. Beginning with HP-UX 11i v3, HP-UX supports two types of Device Special Files: legacy and persistent .
Disk Spanning	The allocation of a logical volume across multiple disks, allowing the volume size to exceed the size of a single disk.
Hardware Path	As opposed to the <i>physical path-specific lunpath</i> , the LUN hardware path presents a <i>virtualized</i> LUN hardware path. This hardware path represents the device or LUN itself rather than a single physical path to the device or LUN. For more information, see the <i>intro(7)</i> manpage.
I/O Channel Separation	A configuration of disks useful for segregating highly I/O-intensive areas. For example, you might have a database on one channel and file systems on another. When mirroring logical volumes using HP MirrorDisk/UX, you can spread the mirrored copies over different I/O channels to increase system and data availability.
Legacy Device Special Files (DSFs)	Legacy DSFs are used in releases prior to HP-UX 11i v3, and also supported on HP-UX 11i v3. A legacy DSF is a DSF with the hardware path information such as SCSI bus, target, and LUN embedded in the file's minor number and file name, such as /dev/dsk/c2t3d4. Because legacy DSFs are based on physical paths, a single multi-pathed LUN often yields multiple legacy DSFs.
Legacy View	The representation of legacy hardware paths and legacy Device Special Files, as in releases prior to HP-UX 11i v3.
Logical Extents	Fixed-size addressable areas of space on a logical volume. The basic allocation unit for a logical volume, a logical extent is mapped to a physical extent; thus, if the physical extent size is 4 MB, the logical extent size will also be 4 MB. The size of a logical volume is determined by the number of logical extents configured.
Logical Volume	A virtual storage device of flexible size that can hold a file system, raw data, dump area, or swap. Because its data are distributed logically (rather than physically), a single logical volume can be mapped to one LVM disk or span multiple disks. A logical volume appears to the administrator as though it was a single disk.
Logical Volume Manager	An operating system software module that implements virtual (logical) disks to extend, mirror, and improve the performance of physical disk access.
LUN	LUN represents a virtual disk unit. Most servers store data on external disk arrays. Storage in an array is subdivided into Logical Units (LUNs). The array assigns each LUN a globally unique WW Identifier (WWID), which is a 64-bit number typically displayed in hexadecimal form. Array administrators also assign each LUN an easier to remember LUN ID number. HP-UX sees each LUN as a disk device.
Lunpath	A <i>physical</i> hardware path leading to a SCSI logical unit. A SCSI LUN can have more than one lunpath if the LUN can be accessed via different physical hardware paths. Compare this with the single <i>virtualized</i> hardware path to the LUN, called the hardware path . For more information, see the <i>intro(7)</i> manpage.

Mirroring	Simultaneous replication of data, ensuring a greater degree of data availability. LVM can map identical logical volumes to multiple LVM disks, thus providing the means to recover easily from the loss of one copy (or multiple copies in the case of multi-way mirroring) of data. Mirroring can provide faster access to data for applications using more data reads than writes. Mirroring requires the MirrorDisk/UX product.
Persistent Device Special Files (DSFs)	A Device Special File conforming to the naming model introduced in HP-UX 11i v3 to support agile addressing. The Device Sepcial File name contains an instance number, such as <code>/dev/disk/disk#</code> , and the minor number has no hardware path information. Persistent DSFs provide a persistent, path-independent representation of a device bound to the device's LUN hardware path and World Wid Identifier (WWID). The persistent DSFs represent the LUN itself (rather than a specific path to a LUN such as legacy DSFs), so a single/persistent DSF is created for a LUN regardless of the number of underlying paths to the LUN.
Physical Extents	Fixed-size addressable areas of space on an LVM disk. They are the basic allocation units for a physical volume. Physical extents map to areas on logical volumes called logical extents.
Physical Volume	A disk that has been initialized by LVM for inclusion in a volume group; also called an LVM disk. As with standard disks, an LVM disk (physical volume) is accessed via a raw device file (for example, <code>/dev/rdisk/disk3</code>). Use the HP SMH or the <code>pvcreate</code> command to initialize a disk as a physical volume.
Physical Volume Group	A subset of physical volumes within a volume group, each with a separate I/O channel or interface adapter to achieve higher availability of mirrored data.
Quorum	The requirement that a certain number of LVM disks be present in order to change or activate a volume group. To activate a volume group, quorum requires the number of available LVM disks to be <i>more than</i> half the number of configured LVM disks that were present when the volume group was last active. To make a configuration change, the quorum requirement is <i>at least</i> half. If there is no quorum, LVM prevents the operation. Quorum is checked both during configuration changes (for example, when creating a logical volume) and at state changes (for example, if a disk fails). Quorum ensures the consistency and integrity of the volume groups. The <code>vgchange</code> command with the <code>-q n</code> option can be used to override quorum check, but this should be used with caution.
Snapshot Logical Volume	A point-in-time image of a logical volume, used to create another copy of the logical volume without taking up as much of the physical space as the size of the logical volume.
Striping	Disk striping distributes logically contiguous data blocks (for example, chunks of the same file) across multiple disks, which speeds up I/O throughput for large files when they are read and written sequentially (but not necessarily when access is random).
Synchronization	The process of updating stale (non-current) copies of mirrored logical extents by copying data from a fresh (current) copy of the logical volume. Synchronization keeps mirrored logical volumes consistent by ensuring that all copies contain the same data.
Unit of Information	<p>The units of information used are as follows:</p> <ul style="list-style-type: none"> • KB is a kilobyte unit of information equal to 2^{10} or 1024 bytes. • MB is a megabyte unit of information equal to 2^{20} or 1,048,576 bytes. • GB is a gigabyte unit of information equal to 2^{30} or 1,073,741,824 bytes. • TB is a terabyte unit of information equal to 2^{40} or 1,099,511,627,776 bytes. • PB is a petabyte unit of information equal to 2^{50} or 1,125,899,906,842,624 bytes.
Unshare Unit	The smallest unit at which data is unshared between a logical volume and its snapshot when a write occurs onto the snapshot or its successor.
Volume Group	A collection of one or more LVM disks from which disk space may be allocated to individual logical volumes. A disk can belong to only one volume group. A volume group is accessed through the group file (for example, <code>/dev/vg01/group</code>) in that volume group's directory. Use HP SMH or the <code>vgcreate</code> command to create a volume group.

Index

Symbols

/etc/default/fs, 98
/etc/fstab, 35, 57, 67, 98
/etc/lvmconf/ directory, 18, 35, 70
/etc/lvmpvg, 33
/etc/lvmtab, 11, 21, 39, 57, 71, 72, 108, 115, 163, 164
/stand/bootconf, 92, 95
/stand/rootconf, 108
/stand/system, 23
/var/adm/syslog/syslog.log, 111, 114, 163, 173

A

adding a mirror to a logical volume, 55
adding a multipathed disk, 28
adding a physical volume to a volume group, 51
agile view, 12
allocation policy, 24
 contiguous and noncontiguous, 25
 strict and non-strict, 24
alternate boot disk
 creating, 88
alternate links see multipathing

B

backups
 mirrored logical volumes, 68
 volume group configuration, 69
 VxFS snapshot file system, 101
bad block relocation, 88, 102, 103, 111
BDRA
 area on disk, 16
 corrupted, 115
 requirement for boot disks, 89
 updating with lvnboot, 36
block device file, 13
Boot Data Reserved Area see BDRA
boot logical volume, 89
 information in BDRA, 16
 lvlnboot, 89
 mirroring, 90
 requirements, 88

C

character device file, 13
cluster-wide device special files, 13
contiguous allocation
 and logical volume size, 22
 defined, 11, 25
 for dump logical volume, 102
 for swap logical volume, 102
converting a physical volume from bootable to
 nonbootable, 84, 169
creating a dump logical volume, 102
creating a file system logical volume, 97

creating a logical volume, 52
creating a mirror of the boot disk, 90
creating a mirrored logical volume, 53
creating a physical volume, 43
creating a spare disk, 76
creating a striped logical volume, 52
creating a swap logical volume, 102
creating a volume group, 12, 44
creating an alternate boot disk, 88

D

database partitions
 stripe size for, 32
device file
 agile view, 12
 block, 13
 character, 13
 creating, 44, 66, 72, 89, 90, 93
 format, 15
 legacy, 12, 13, 28, 87
 legacy view, 12
 logical volume, 14, 15, 52, 56
 persistent, 12, 13
 physical volume, 13, 15, 73, 90, 93, 163
 volume group, 15, 44, 57, 66, 72, 89
disabling a path to a physical volume, 87
disk failure, 129
disk sparing see sparing
disk striping see striping
disks, 9
 see also physical volumes
 moving, 71, 72
displaying LVM information, 40
du command, 21
dual cabling (dual controllers) see multipathing
dump logical volume, 23
 creating, 102
 guidelines, 23
 removing, 103
 requirements, 102

E

error handling, 109
 media errors, 111
 non-recoverable errors, 110
 recoverable errors, 109
exporting a volume group, 57
extends command, 22, 99, 152
extending a file system logical volume, 98
extending a logical volume, 53
extending a logical volume to a specific disk, 54
extending a swap logical volume, 102

F

file system logical volume, 21
 and /etc/default/fs, 98

backing up via mirroring, 68
boot file system see boot logical volume
creating, 97
determining who is using, 57, 68, 99, 100, 131
extending, 98
guidelines, 22
in /etc/fstab, 98
initial size, 21
OnlineJFS, 98
overhead, 21
performance considerations, 22
reducing, 99, 115
 HFS or VxFS, 100
 OnlineJFS, 100
resizing, 22
root file system see root logical volume
short or long file names, 98
stripe size for HFS, 32
stripe size for VxFS, 32
unresponsive, 109
finding logical volumes using a disk, 123
fsadm command, 99, 100
fully allocated snapshot, 103
fuser command, 55, 57, 68, 99, 100, 131
 and NFS, 99, 100

G

GB, defined, 178
group device file, 44

H

hot-swappable disks, 123
HP SMH, 9, 38
 managing LVM, 38
 naming conventions, 13
 running, 38

I

idisk command, 89, 92, 93, 136, 139
importing a volume group, 58
insf command, 89, 90, 93, 136, 139
Insight Remote Support, 145
interleaved swapping, 102
io_redirect_dsf command, 130, 133, 136
ioscan command, 35, 107, 130, 133, 136, 164
 to determine device files, 43, 73
 to determine hardware paths, 129, 131, 134
 to determine instance numbers, 130, 132, 135

K

KB, defined, 178

L

legacy view, 12
LIF volume
 area on disk, 17
 maintenance mode boot, 108
 requirement for boot disks, 44, 89
logical extents

and rounding logical volume size, 21
defined, 9
mapping to physical extents, 10
Logical Interface Format see LIF volume
logical volumes, 23, 101
 see also boot logical volume
 see also dump logical volume
 see also root logical volume
 see also swap logical volume
commands for, 40
configuration information, 71
creating, 52
creating on a specific disk, 54
defined, 9
determining who is using, 55, 57, 68, 99, 100, 131
device file, 14, 15, 52, 56
displaying information, 42
extending, 53
extending to a specific disk, 54
for swap, 22
naming convention, 14
performance issues, 21
reducing, 55
removing, 57
renaming, 56
size, 20
lv timeout, 161
lvchange command, 40, 53, 152
 bad block relocation, 111
 errors, 164
 setting allocation policy, 24
 setting scheduling policy, 25
 setting synchronization policy, 25, 90
 setting timeout, 110
lvcreate command, 40, 52, 152
 for dump logical volume, 102
 for swap logical volume, 102
 setting allocation policy, 24
 setting scheduling policy, 25
 setting synchronization policy, 25
 striped logical volumes, 52
 without a size, 22
lvdisplay command, 10, 26, 35, 40, 42, 107, 108, 152
 displaying extent status, 123
 displaying mirror status, 123
 displaying timeout value, 110
 scheduling policy, 29
lvextend command, 40
 adding a mirror, 56, 91, 94, 152
 errors, 94, 164
 extending a file system, 99
 extending a logical volume, 53, 152
 extending swap, 102
 extending to a specific disk, 54, 88
lvinboot command, 12, 17, 35, 40, 152
 displaying boot information, 16, 89, 92, 94, 107, 123
 errors, 165
 for boot logical volume, 88, 89
 for dump logical volume, 89, 103

for root logical volume, 89
for swap logical volume, 89, 102
updating boot information, 36, 92, 94, 115
lvmadm command, 12, 39, 107, 108, 112, 150, 152, 163
lvmchk command, 39
lvmerge command, 40, 68, 152
 synchronization, 26
lvmmove command, 40, 152
lvmpud, 105
lvmpud command, 39, 152
lvreduce command, 40
 and **pvmove** failure, 75
 reducing a file system, 100, 115
 reducing a logical volume, 55, 152
 reducing a swap device, 102
 removing a mirror, 56, 152
 removing a mirror from a specific disk, 56
lvremove command, 40, 57, 152
 splitting a volume group, 67
lvrmboot command, 12, 16, 17, 40, 103, 115, 152
lvsplit command, 40, 68, 152
lvsync command, 26, 40, 152

M

maintenance mode boot, 90, 115
 booting, 108
 requirement for boot logical volume, 88
major number, 16, 44
mapping logical to physical extents, 10
mass storage stack, 28
 tunable parameters, 162
MB, defined, 178
merging a split mirror logical volume, 68
minfree, 21
minor number, 16, 44
Mirror Consistency Recovery, 26
mirror copy
 removing from disk, 126
Mirror Write Cache, 25
mirrored logical volumes, 24
 adding a mirror, 55, 152
 allocation policy, 24
 backing up, 68
 benefits, 24
 boot logical volume, 90
 creating a spare disk, 76
 creating mirrored copies, 53
 defined, 8, 178
 determining if a logical volume is mirrored, 123
error messages, 164, 165, 174
handling non-recoverable errors, 110
handling recoverable errors, 109
logical to physical extent mapping, 10
maintenance mode, 109
merging, 68, 152
Mirror Write Cache, 29
mirroring the boot disk, 90
mirroring the boot disk,

HP 9000 servers, 90
HP Integrity servers, 92
modifying mirror copies, 53
physical volume groups, 30, 33
primary swap logical volume, 90
quiescing a logical volume, 66
reinstating a spare disk, 77
removing a mirror, 56, 152
replacing a boot disk, 134
replacing a non-boot disk, 129
root logical volume, 90
scheduling policy, 25, 29
separating I/O channels, 33
sparing, 27
splitting, 68, 152
stale data, 26
strict allocation policy, 165
striping, 33
 synchronization policy, 25
 synchronizing, 26, 152, 154
mkboot command, 17, 89, 91, 93, 137
mknod command, 44, 66, 72, 89, 153
modifying a mirrored logical volume, 53
modifying a volume group's parameters, 58
moving data, 73, 75
moving disks, 71, 73
moving physical volumes, 71, 72
multipathing, 28
 benefits, 28
 defined, 8
 disabling a path, 87
 importing volume groups, 58
 moving multipathed disks, 72
 native, 13, 28
 persistent device file, 13
 removing multipathed disks, 52
 setting up, 28
 switching between links, 28
 using LVM, 28

N

naming conventions, 13
newfs command, 32, 98, 100, 134
NFS
 and fuser, 99, 100
non-strict allocation policy, 24
noncontiguous allocation, 11
 defined, 25

P

parallel scheduling policy, 25
path_fail_secs parameter, 162
PB, defined, 178
physical extents, 10
 beyond the size of a physical volume, 84
 defined, 9
 finding which logical volume is using, 123
 moving to another disk, 52, 59, 63
performance considerations, 29

policies for allocating, 24
 policies for writing, 25
 size, 9, 17
 synchronizing, 26
 physical volume groups, 30, 33
 naming convention, 15
Physical Volume Reserved Area see PVRA
 physical volumes
 adding, 51
 auto re-balancing, 75
 commands for, 39
 converting from bootable to nonbootable, 84, 169
 creating, 43
 defined, 9
 device file, 13, 15, 163
 disabling a path, 87
 disk layout, 16
 displaying information, 41
 moving, 71, 72
 moving data between, 73
 naming convention, 13
 removing, 51
 resizing, 77
 pre-allocated extents, 103
 primary swap logical volume, 23
 as a dump area, 23
 mirroring, 90
 PV timeout, 161
 pvchange command, 39, 153
 disabling a path, 87, 129, 132, 134, 138
 errors, 166
 multipathing, 28
 restoring volume group configuration, 70
 setting timeout, 110
 sparing, 77
 pvck command, 12, 39, 108, 153
 pvcreate command, 9, 12, 39, 43, 153
 for boot disks, 44, 89, 91, 93
 sparing, 77
 pvdisplay command, 10, 35, 39, 41, 107, 108, 153
 finding logical volumes using a disk, 123
 sparing, 27
 pvlinks see multipathing
 pvmove command, 39, 73, 153
 abnormal termination, 75
 auto re-balancing, 75
 moving physical extent 0, 59, 60, 63, 65
 sparing, 77
PVRA
 area on disk, 17
 pvremove command, 39, 153

Q
 quiescing a volume group, 65
 quorum, 34
 defined, 111, 178
 enabling, 68
 error messages, 111
 overriding, 67, 111, 112, 178

overriding at boot, 91, 93, 112, 135, 138
 requirements for booting, 115

R
 raw data logical volume
 stripe size for, 32
 reducing a logical volume, 55
 reducing the size of a file system logical volume, 99
 reinstating a spare disk, 77
 remote support, 145
 removing a dump logical volume, 103
 removing a logical volume, 57
 removing a mirror copy from a disk, 126
 removing a mirror from a logical volume, 56
 from a specific disk, 56
 removing a physical volume from a volume group, 51
 removing a volume group, 68
 renaming a logical volume, 56
 renaming a volume group, 66
 replacing a failed disk, 129
 mirrored, 129, 134
 unmirrored, 131
 resizing a physical volume, 77
 resizing a swap logical volume, 102
 restoring data
 volume group configuration, 70
 resuming a quiesced volume group, 65
root logical volume, 89
 creating, 88
 information in BDRA, 16
 lvolboot, 89
 mirroring, 90
 requirements, 88, 98
root volume group
 and dump, 23

S
 scheduling policy, 25
 parallel, 25
 sequential, 25
scsimgr command
 during disk replacement, 130, 132, 135
 multipathing, 28, 87
secondary swap, 23
 configuration, 102
sequential scheduling policy, 25
snapshot logical volumes
 creating, 103
 defined, 178
 displaying, 104
 fully-allocated, 103
 removing, 104
 space-efficient, 103
 types, 103
space-efficient snapshot, 103
spanning
 defined, 8
 performance considerations, 30
 sparing, 12, 27

creating a spare disk, 76
 defined, 8, 27
 detaching links, 87
 reinstating a spare disk, 77
 requirements, 27
 splitting a mirrored logical volume, 68
 splitting a volume group, 67
 stale data, 26
 strict allocation policy, 24
 stripe size, 32
 striping, 31
 and mirroring, 33
 benefits, 31
 creating a striped logical volume, 52
 defined, 8
 interleaved disks, 31
 performance considerations, 31
 selecting stripe size, 32
 setting up, 31
 swap logical volume, 22, 23, 101
 see also primary swap logical volume
 creating, 89, 102
 extending, 102
 guidelines, 22
 information in BDRA, 16
 interleaving, 22
 IODC errors, 165
 lVlnboot, 89, 102
 maintenance mode, 109
 mirroring policy, 26, 90
 performance considerations, 22
 printing information, 16
 requirements, 11, 88, 102
 resizing, 102
 secondary swap, 23
 stripe size for, 32
 synchronization policy, 25
 Mirror Consistency Recovery, 26
 Mirror Write Cache, 25
 none, 26
 synchronizing a mirror, 26
 automatically, 26
 manually, 26
 System Management Homepage *see* HP SMH

T

TB, defined, 178
 timeout
 LV timeout, 33, 161
 PV timeout, 33, 161
 transient_secs parameter, 162
 troubleshooting tools, 107

U

unshare unit
 defined, 178

V

Version 1.0 volume groups, 11, 16, 44, 147
 Version 2.0 volume groups, 12, 147
 Version 2.1 volume groups, 12, 147
 Version 2.x volume groups, 12, 16, 22, 23, 27, 40, 44, 45, 76, 77, 84, 88, 90, 101, 102, 108, 112, 167, 169
 vgcdf command, 39, 153
 vcfgbackup command, 35, 39, 66, 69, 153
 backup location, 70
 errors, 164, 166
 vcfgrestore command, 39, 70, 107, 153
 backup location, 70
 checking if physical volume is bootable, 85
 errors, 166
 recovering corrupted LVM data, 115
 restating a spare disk, 77
 vgchange command, 39, 90, 153
 activating a volume group, 58, 61, 66, 71, 72, 73, 81, 87, 112
 activation failures, 111
 attaching all detached links, 88, 131, 133, 137
 automatic synchronization, 26
 deactivating a volume group, 57, 61, 66, 67, 71, 72, 80, 86
 disabling quorum checks, 67, 111
 enabling quorum checks, 68, 112
 errors, 167
 quiescing a volume group, 66
 resuming a quiesced volume group, 66
 vgchgid command, 39, 153
 integrating cloned LUNs, 106
 splitting a volume group, 67
 vgcreate command, 39, 44, 45, 153
 adding a multipathed disk, 28
 errors, 169
 for alternate boot disk, 89
 vgdisplay command, 35, 39, 41, 107, 108, 153
 displaying free space, 52, 53, 164
 errors, 170
 sparing, 27
 while quiesced, 65
 vgdsf command, 39, 153
 vgexport command, 39, 57, 71, 153
 map file, 57
 moving disks, 71, 72
 renaming a volume group, 66
 splitting a volume group, 67
 vgextend command, 39, 51, 153
 errors, 171
 sparing, 77
 with multipathed disks, 28, 58
 VGID, 17, 168
 changing, 106
 splitting a volume group, 67
 with vgexport, 57
 with vgimport, 58
 vgimport command, 39, 58, 71, 153
 errors, 171
 moving disks, 72, 73
 renaming a volume group, 66

- splitting a volume group, 67
- with multipathed disks, 58

vgmodify command, 12, 39, 58, 153

- changing physical volume type, 84, 169
- collecting information, 59, 62
- errors, 172
- modifying volume group parameters, 59, 171
- resizing physical volumes, 77, 165

vgmove command, 95, 153

VGRA

- and `vgmodify`, 59, 63
- area on disk, 17
- size dependency on extent size, 17, 169

vgreduce command, 39, 51, 154

- with multipathed disks, 58

vgremove command, 39, 68, 154

vgscan command, 39, 154

- moving disks, 72
- recreating `/etc/lvmtab`, 115

vgsync command, 26, 39, 154

vgversion command, 40, 46, 154

- errors, 172

volume group configuration

- backing up, 69
- location, 70
- restoring, 70

volume group identifier see `VGID`

volume group migration to new disks, 95

Volume Group Reserved Area see `VGRA`

volume group versions, 46

- Version 1.0, 11, 16, 44, 147
- Version 2.0, 12, 147
- Version 2.1, 12, 147
- Version 2.x, 12, 16, 22, 23, 27, 40, 44, 45, 76, 77, 84, 88, 90, 101, 102, 108, 112, 167, 169

volume groups

- activation failures, 111, 114
- commands for, 40
- creating, 12, 44
- defined, 18
- device file, 15, 44, 57, 66, 72, 89
- displaying information, 41
- exporting, 57
- importing, 58
- modifying parameters, 58
- moving, 71, 73
- naming convention, 14
- performance considerations, 30
- quiescing and resuming, 65
- removing, 68
- renaming, 66
- space available within, 52
- splitting, 67

VxFS snapshot file system, 101