

Ignite-UX Custom Configuration Files



Table of Contents

Abstract	6
Introduction	6
Typographic Conventions	7
HP-UX 11i release names and release identifiers	8
Configuration files and INDEX files	9
The INDEX file	9
The CINDEX file	10
The per-client configuration file	10
The global config.local file	11
The recovery config.local file	11
Order of precedence of configuration files	11
Testing the order of precedence	12
What is in a configuration (cfg) clause?	13
The make_net_recovery configuration files	14
The make_tape_recovery configuration files	14
Files created by make_config	15
Using the manage_index command	15
A variety of uses	15
Adding a configuration file to a clause or "release"	16
Adding scripts to the INDEX file	17
Removing cfg clauses from an INDEX file	18
Setting the default cfg clause in an INDEX file	19
Listing the names of cfg clauses in an INDEX file	20
Listing the name of the default cfg clause in an INDEX file	21

Renaming a cfg clause in an INDEX file	21
Creating a new cfg clause from an existing clause	22
Removing a configuration file from a cfg clause	22
Removing a script from an INDEX file	23
List the names of all configuration files in a cfg clause	24
Display the description of a cfg clause.....	25
 Using the make_bundles command	25
Why do you need to use make_bundles ?	25
Choosing which form of make_bundles to use.....	25
The make_bundles first form	26
The make_bundles second form.....	31
The make_bundles third form.....	33
 Using the instl_dbg command	35
Introduction.....	35
Requirements.....	35
Using the itool command	35
Combining instl_dbg and itool	36
Running instl_dbg	36
Other instl_dbg options	41
The hw.info and host.info files	41
Creating both files	41
 Miscellaneous configuration tips.....	43
Analyzing the HP-UX default B.11.11 cfg clause	43
The release-specific configuration file	44
Analyzing the HP-UX default B.11.31 cfg clause	64
The release-specific configuration file	65
 Special variables.....	85
<code>_hp_locale</code>	85
<code>_hp_cfg_detail_level</code>	86
<code>_hp_pri_swap</code>	87
<code>_hp_min_swap</code>	87
<code>_hp_disk_layout</code>	88
<code>_hp_default_cur_lan_dev</code>	90
<code>_hp_default_final_lan_dev</code>	90
<code>_hp_keyboard</code>	91
<code>_hp_root_disk</code>	92
<code>_hp_boot_dev_path</code>	92
<code>_hp_primary_path</code>	93
<code>_hp_primary_partition_size</code>	93
<code>_hp_efi_partition_size</code>	93
<code>_hp_service_partition_size</code>	94
<code>_hp_root_grp_disks</code>	94
<code>_hp_root_grp_stripped</code>	94
<code>_hp_addnl_fs_free_pct</code>	95
<code>_hp_ignore_sw_impact</code>	95
<code>_hp_custom_sys</code>	97
<code>_hp_lanadmin_args</code>	98
<code>_hp_nfs_mount_opts</code>	99
<code>_hp_nfs_mount_retries</code>	99
<code>_hp_tftp_cmds</code>	100

<code>_hp_hide_other_disks</code>	100
<code>_hp_saved_detail_level</code>	100
<code>_hp_os_bitness</code>	101
<code>_hp_force_autoboot</code>	101
<code>_hp_ikernel_os_release</code>	102
<code>_hp_current_client_release</code>	102
<code>_HP_CLONING</code>	102
<code>_hp_console_verbosity</code>	103
<code>_hp_patch_save_files</code>	103
<code>_hp_umask</code>	104
<code>_hp_ht_enable</code>	104
<code>_hp_debug_level</code>	105
 Configuration for software to be installed	105
Application software depots	105
Core operating system depot configuration	110
Impacts statements	116
Overstated SD impacts	118
Categories and other Ignite-UX software attributes	118
Defining a custom software configuration	120
Looking at a network recovery <code>sw_source</code> and <code>sw_sel</code>	120
Using a <code>sw_sel</code> to run commands instead of installing software	122
Using a <code>sw_sel</code> to apply kernel parameters	123
Forcing software (<code>sw_sel</code>) clauses to be installed	127
Automating dependencies in software	127
Installing patches	129
 Configuration for volume and disk groups	129
Overview	130
 Configuration examples	131
Example One (custom disk layout)	131
Example Two (selection of disk layout based on hardware)	136
Example Three (be careful what you ask for, you just might get it)	138
Example Four (installation file system configuration)	139
Part A (custom configuration in installation file system)	139
Part B (Installation file system custom network config)	140
Example Five (handling user input)	140
Part A (restricting the values that a variable can take using an enum)	141
Part B (looking at the same thing when we don't have an enum)	142
Example Six (handling command line arguments)	142
Example Seven (regular expression matching)	143
 Configuration parameters in the installation file system	148
Networking	148
Problems that can be solved with <code>_hp_lanadmin_args</code>	150
Control	151
Environment variables	152
 Managing configurations with <code>unifdef</code>	153
 Coping with <code>auto_adm</code> and boot changes in HP-UX B.11.23	155
Looking at <code>auto_adm</code>	156
ISL and CONF format data	156
CONF data	156
ISL data	158
Usage examples	158

Creating new files	158
Adding new menu entries to a file.....	159
Using an "append" file.....	160
Updating a menu entry in a file	160
Deleting a menu entry from a file	161
Changing the default menu choice	162
Changing the timeout	162
Updating the prompt message	163
Into and out of an LIF file.....	163
Installation configurations using Software Distributor depots	163
Getting started	164
Creating the core operating system depot.....	164
Operating Environments containing multiple CDs or DVDs.....	166
Other methods for creating the Core OS depot	167
Using the Ignite-UX GUI	167
The pitfalls of using multiple media.....	173
Creating the configuration file to describe the depot	173
Creating a minimalist <code>cfg</code> clause for installation	175
Comments on the SD based <code>cfg</code> clause	176
Extending the generic SD <code>cfg</code> clause with applications.....	176
Creating the application depot	176
Creating a configuration for the application depot	176
An example problem	177
Follow on consequences	178
Resolving applications that are not in SD format.....	179
Package an application in SD format	179
Define an application in a non-SD format.....	180
Using <code>noncore.cfg</code> to define applications	183
SD and archive bitness comparison	186
Adding the non-SD application configuration file to the <code>INDEX</code> file	187
Installing patches from a depot.....	187
Setting up the depot	187
Packaging the patches into a bundle	187
Generating a configuration	188
Customizing configuration	189
Installing with SD wrap-up	191
Performance considerations for SD-UX based installs	191
Tuning SD-UX for concurrent access	192
Issues to consider related to the HP-UX revision	192
Issues Independent of HP-UX revision	193
Memory and the buffer cache.....	193
Network Bandwidth	194
Installation configurations using golden images.....	195
Golden image configuration file explanation.....	195
Instances that may require modifying <code>os_arch_post_1</code>	201
Creating a golden image.....	202
Final words about golden image installations	202
Understanding what <code>is_net_info_temporary</code> does.....	203
Understanding how VxFS file system versions are set.....	204
How do I.....	205
How do I remove the warning message that occurs when compiling a kernel on PA-RISC systems?	205

How do I recognize if a disk exists or not from within a configuration file?	206
How do I create the CD equivalent of a tape created by <code>make_boot_tape</code> ?	207
How do I enable the X server and CDE during a golden image install?	211
Summary.....	212
For more information	212

Abstract

Ignite-UX for HP-UX addresses the need for system administrators to perform operating system installations, deployment, and recovery, often on a large scale. It provides the means for creating and reusing standard operating system configurations. Additionally, Ignite-UX delivers the ability to archive operating system configurations, and to use these archives to replicate systems, with the added benefit of speeding up the process. Ignite-UX also permits various customizations, and is capable of both interactive and unattended operating modes.

Introduction

The *Ignite-UX Administration Guide* explains how to use the Ignite-UX product. However, it does not cover in detail the configuration files used by Ignite-UX (and custom configuration files that you can write).

This document supplements the *Ignite-UX Administration Guide* to help make creating and writing custom configuration files easier. It is assumed readers understand how to install a system using Ignite-UX. The information in this document addresses managing, writing, and modifying configuration files.

In developing this document, Ignite-UX versions from B.5.1.x to C.7.1.x were used.

Note:

Care was taken to prevent unintended line wraps from occurring in the examples in this document; however, due to formatting limitations, some might exist. This may affect the usefulness of examples if you use them without verifying the syntax using `instl_adm` with the `-T` option.

Typographic Conventions

The following typographical conventions are used throughout this document.

audit(5)

An HP-UX manpage. "audit" is the name and "5" is the section in the *HP-UX Reference*. On the Web and on the Instant Information media, it may be a hot link to the manpage itself. From the HP-UX command line, enter "man audit" or "man 5 audit" to view the manpage. For more information, refer to *man(1)*.

Book Title

The title of a book. On the Web and on the Instant Information media, it may be a hot link to the book itself.

Emphasis

Text that is emphasized.

Emphasis

Text that is strongly emphasized.

ComputerOut

Text displayed by the computer.

Command

A command name or qualified command phrase.

Computer

Computer font indicates literal items displayed by the computer. For example:

file not found

Filename

Text that shows a filename and/or filepath.

UserInput

Commands and other text that you type.

Variable

The name of a variable that you may replace in a command or function or information in a display that represents several possible values.

[]

The contents are optional in formats and command descriptions.

{ }

The contents are required in formats and command descriptions. If the contents are in a list separated by |, you must choose one of the items

...

The preceding element may be repeated an arbitrary number of times.

|

Separates items in a list of choices.

<MAC>

When used as part of a file name, this placeholder variable refers to the Media Access Control (MAC) address of the Ignite-UX client system under /var/opt/ignite/clients. This may also be referred to as Link-Level Address (LLA).

<DATE-TIME>

This indicates a directory created by Ignite-UX that has a fixed format containing the date and time it was created, for example "2004-01-12,15:23".

HP-UX 11i release names and release identifiers

Each HP-UX 11i release has an associated release name and release identifier.

Table 1 shows the releases available for HP-UX 11i.

Table 1

Release Name	Release Identifier	Supported Processor Architecture
HP-UX 11i v1	B.11.11	PA-RISC
HP-UX 11i v2	B.11.23	Intel® Itanium PA-RISC ¹
HP-UX 11i v3	B.11.31	Intel® Itanium PA-RISC ¹

The *uname(1)* command with the *-r* option returns the release identifier.

You can also determine the update release date and the Operating Environment by entering the following:

```
# swlist | grep HPUX11i
```

The resulting output lists the current release identifier, update release date, and Operating Environment. For example:

```
HPUX11i-OE-MC      B.11.23.0606      HP-UX Mission Critical Operating Environment
Component
```

The preceding revision string represents the following:

```
B.11.23 = HP-UX 11i v2
0606 = June 2006 Update Release (in the format YYMM)
```

¹ HP9000 systems are supported on HP-UX 11i v2 starting with the September 2004 release.

Configuration files and INDEX files

This section discusses the contents of configuration files and INDEX files to help you understand the purpose of the files, not to document what is occurring in the files.

The INDEX file

The file /var/opt/ignite/INDEX is used during cold-installs. The following example only has one cfg clause; usually there are more. Each cfg clause defines a series of configuration files, along with a description, that together form a cohesive configuration that is used to install a system.

```
$ pwd  
/var/opt/ignite  
$ cat INDEX  
# /var/opt/ignite/INDEX  
# This file is used to define the Ignite-UX configurations  
# and to define which config files are associated with each  
# configuration. See the ignite(5), instl_adm(4), and  
# manage_index(1M) man pages for details.  
#  
# NOTE: The manage_index command is used to maintain this file.  
#       Comments, logic expressions and formatting changes are not  
#       preserved by manage_index.  
#  
# WARNING: User comments (lines beginning with '#'), and any user  
#           formatting in the body of this file are not preserved  
#           when the version of Ignite-UX is updated.  
#  
cfg "HP-UX B.11.11 Default" {  
    description "This selection supplies the default system configuration  
that HP supplies for the B.11.11 release."  
    "/opt/ignite/data/Rel_B.11.11/config"  
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"  
    "/var/opt/ignite/config.local"  
}
```

If the cfg clause in the INDEX file were to be the default cfg clause (in a non-interactive installation this would cause the cfg clause to be selected for installation by default) it would look like the following:

```
cfg "HP-UX B.11.11 Default" {  
    description "This selection supplies the default system configuration  
that HP supplies for the B.11.11 release."  
    "/opt/ignite/data/Rel_B.11.11/config"  
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"  
    "/var/opt/ignite/config.local"  
}=TRUE
```

Of course, the INDEX file is never to be edited manually; instead, the manage_index command should be used to maintain this file. The actual command needed to set a cfg clause to be the default clause is as follows:

```
# manage_index -e -c "HP-UX B.11.11 Default"
```

The manage_index command operates on the /var/opt/ignite/INDEX file by default, so the INDEX file does not need to be specified on the command line.

The CINDEX file

The CINDEX file is very different from the INDEX file. The CINDEX file is created and managed by the make_net_recovery command. The format of the CINDEX file is the same as the INDEX file. For example:

```
# cat /var/opt/ignite/clients/<MAC>/CINDEX
# This file is used to define the Ignite-UX configurations
# and to define which config files are associated with each
# configuration. See the ignite(5), instl_adm(4), and
# manage_index(1M) man pages for details.
#
# NOTE: The manage_index command is used to maintain this file.
#       Comments, logic expressions and formatting changes are not
#       preserved by manage_index.
#
# WARNING: User comments (lines beginning with '#'), and any user
#           formatting in the body of this file are not preserved
#           when the version of Ignite-UX is updated.
#
cfg "2003-10-08,12:41 Recovery Archive" {
    description "Recovery Archive"
    "recovery/2003-10-08,12:41/system_cfg"
    "recovery/2003-10-08,12:41/control_cfg"
    "recovery/2003-10-08,12:41/archive_cfg"
}
cfg "2003-10-08,12:45 Recovery Archive" {
    description "Recovery Archive"
    "recovery/2003-10-08,12:45/system_cfg"
    "recovery/2003-10-08,12:45/control_cfg"
    "recovery/2003-10-08,12:45/archive_cfg"
}=TRUE
```

This file contains the cfg clauses needed to recover a system from a network recovery archive. The make_net_recovery command automatically manages the CINDEX file so older information is removed² and the latest recovery archive is set as the default archive that is recovered.

The per-client configuration file

The configuration file that can exist in a per-client directory (var/opt/ignite/clients/<MAC>) overrides all other files that an Ignite-UX client may use. Any configuration values set in this file override all other files with one exception that is discussed elsewhere (see "Order of precedence of configuration files" for more information). During the installation process, Ignite-UX saves to this file the configuration information that was used, which facilitates repeated installations of the same system. If this file exists, it typically selects a cfg clause from the CINDEX or INDEX files. For example:

```
# cat config
cfg "2003-10-08,12:45 Recovery Archive"=TRUE
_hp_cfg_detail_level="ipvs"
#
# Variable assignments
```

² For more information, refer to *make_net_recovery(1M)* and review the description of the -n option.

```

#
init _hp_root_disk="8/16/5.6.0"
init _hp_root_grp_disks=1
init _hp_root_grp_striped="NO"
...

```

In this case, a recovery archive has been selected. (This also indicates that a recovery archive may have been recovered onto this system at some time in the past.)

The global **config.local** file

The file `/var/opt/ignite/config.local` is included by default into all new `cfg` clauses in the `/var/opt/ignite/INDEX` file. For this reason, configuration data stored in this file is used on all new installations.

The recovery **config.local** file

For `make_net_recovery`, if the file

`/var/opt/ignite/clients/<MAC>/recovery/config.local` exists on the Ignite-UX server, the contents are automatically included into the configuration used by the client during a network recovery.

For `make_tape_recovery`, if the file `/var/opt/ignite/recovery/config.local` exists, the contents are automatically included into the configuration used by the client even though it is not explicitly referenced. The content of this file is placed into the LIF on the tape.

Note:

If you use the `-s` option with `make_tape_recovery`, the same `config.local` file used by `make_net_recovery` is used instead of `/var/opt/ignite/recovery/config.local`, as previously described.

Order of precedence of configuration files

Configuration files have an order of precedence when a client is reading configurations from an Ignite-UX server. The precedence is as follows:

1. The per-client configuration file (`/var/opt/ignite/clients/<MAC>/config`) — Any variable set in the per-client config file takes precedence over any variable set anywhere else. Using the per-client config file in the preceding example, the `cfg` clause `2003-10-08,12:45 Recovery Archive` in the `CINDEX` file is selected no matter what the `INDEX` or `CINDEX` have set to TRUE. (This, of course, you can change this using the Ignite-UX Graphical User Interface (GUI).)
2. The per-client `CINDEX` file (`/var/opt/ignite/clients/<MAC>/CINDEX`) — This file contains recovery archive configurations. Any `cfg` clause set to TRUE in this file automatically selects that `cfg` clause if there is no per-client config file.
3. The global `INDEX` file (`/var/opt/ignite/INDEX`) — This is the last file to have precedence. The default clause from the `INDEX` file is the one that is set to TRUE. (None of the `cfg` clauses in the `INDEX` file has been set to TRUE.)

Configuration also has an order of precedence when you are booting a system. The precedence is as follows:

1. The first 8 KB of the installation file system has the highest precedence since it is the first configuration read by Ignite-UX.
2. During media installations, if there is a LIF file called CUSTOM on any device, this file is loaded and parsed after the configuration from the installation file system has been processed. Parsing the configuration after loading all other configurations means that it has the highest order of precedence (when booting from a CD or tape)³.

Testing the order of precedence

It is possible to test the order of precedence using the `instl_dbg` command:

```
# pwd
/var/opt/ignite/clients/<MAC>
# cat CINDEX
...
cfg "2003-10-08,12:41 Recovery Archive" {
    description "Recovery Archive"
    "recovery/2003-10-08,12:41/system_cfg"
    "recovery/2003-10-08,12:41/control_cfg"
    "recovery/2003-10-08,12:41/archive_cfg"
}
cfg "2003-10-08,12:45 Recovery Archive" {
    description "Recovery Archive"
    "recovery/2003-10-08,12:45/system_cfg"
    "recovery/2003-10-08,12:45/control_cfg"
    "recovery/2003-10-08,12:45/archive_cfg"
}=TRUE
# manage_index -e -c "2003-10-08,12:41 Recovery Archive" \
> -i /var/opt/ignite/clients/box1/CINDEX
# cat CINDEX
...
cfg "2003-10-08,12:41 Recovery Archive" {
    description "Recovery Archive"
    "recovery/2003-10-08,12:41/system_cfg"
    "recovery/2003-10-08,12:41/control_cfg"
    "recovery/2003-10-08,12:41/archive_cfg"
}=TRUE
cfg "2003-10-08,12:45 Recovery Archive" {
    description "Recovery Archive"
    "recovery/2003-10-08,12:45/system_cfg"
    "recovery/2003-10-08,12:45/control_cfg"
    "recovery/2003-10-08,12:45/archive_cfg"
}
```

So far, the CINDEX for this client has been changed so that the older recovery archive is the default archive. This is needed in order to show how precedence works. In the config file, the following `cfg` clause is selected:

```
# cat config
cfg "2003-10-08,12:45 Recovery Archive"=TRUE
_hp_cfg_detail_level="ipvs"
...
```

³ This feature is being considered for obsolescence. This part of the installation process is very time consuming for systems with a very large number of disks.

You can now use the `instl_dbg` command to print out the fully parsed configuration (The `-D` option points it to a per-client directory containing the files it needs. The `-f` option designates where to write the parsed configuration). This will occur with and without the `config` file being present in the per-client directory:

```
# instl_dbg -D . -f /tmp/client.with.config
# mv config config.save
# instl_dbg -D . -f /tmp/client.with.no.config
```

With the per-client config file present, the `cfg` clause 2003-10-08,12:45 Recovery Archive is selected:

```
# head /tmp/client.with.config
cfg "2003-10-08,12:45 Recovery Archive"=TRUE
server="10.0.0.3"
is_net_info_temporary=FALSE
init _hp_keyboard="PS2_DIN_US_English"
system_name="host"
ip_addr[]="10.0.0.1"
netmask[]="0xffffffff00"
route_gateway[0]="10.0.0.2"
route_destination[0]="default"
_hp_cfg_detail_level="ibnpvcdsrlLtfh"
...
...
```

Without the per-client config file present, the `cfg` clause is selected by the `CINDEX` file (2003-10-08,12:41 Recovery Archive) instead:

```
# head /tmp/client.with.no.config
cfg "2003-10-08,12:41 Recovery Archive"=TRUE
server="10.0.0.3"
is_net_info_temporary=FALSE
init _hp_keyboard="PS2_DIN_US_English"
system_name="host"
ip_addr[]="10.0.0.1"
netmask[]="0xffffffff00"
route_gateway[0]="10.0.0.2"
route_destination[0]="default"
_hp_cfg_detail_level="ibnpvcdsrlLtfh"
...
...
```

The `instl_dbg` command is covered in more detail in a later section.

What is in a configuration (`cfg`) clause?

The question now is, what is in a `cfg` clause? The following example shows a `cfg` clause taken from the `/var/opt/ignite/INDEX` file:

```
cfg "HP-UX B.11.11 Default" {
    description "This selection supplies the default system configuration
that HP supplies for the B.11.11 release."
    "/opt/ignite/data/Rel_B.11.11/config"
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"
    "/var/opt/ignite/config.local"
}=TRUE
```

The `cfg` clause needs the following information:

1. A name (for example, HP-UX B.11.11 Default)
2. A description (see the `description` keyword)
3. A list of one or more configuration files referenced by this `cfg` clause
4. One or more configuration files that define software that may be installed⁴

When a `cfg` clause is selected (it has been set to `TRUE` and nothing of higher precedence changes this), its configuration files are processed in order and evaluated. This turns the set of configuration files into something that a client can use to install itself.

The `make_net_recovery` configuration files

A `make_net_recovery` session always creates the same configuration files, and you can see what they are by looking at a `cfg` clause in a `CINDEX` file:

```
cfg "2003-10-08,12:45 Recovery Archive" {  
    description "Recovery Archive"  
    "recovery/2003-10-08,12:45/system_cfg"  
    "recovery/2003-10-08,12:45/control_cfg"  
    "recovery/2003-10-08,12:45/archive_cfg"  
}
```

The three configuration files are as follows:

1. `system_cfg` –
This file is produced by the `save_config` command. It contains the file system layout, networking information, and hardware instance numbers for the system.
2. `control_cfg` –
This file contains definitions of Ignite-UX variables and the commands needed to import volume groups back into the final system.
3. `archive_cfg` –
This file contains the software definition of the archive that is used for recovery (including `impacts` keywords, etc.).

The `CINDEX` file provides a higher level of precedence than the `INDEX` file, so any `cfg` clause selected in the `CINDEX` file is used in preference to anything in the `INDEX` file in a non-interactive installation for a system. In an interactive installation, the recovery archive is automatically selected for installation so you can manually select another configuration for installation on the **Basic** tab in the Ignite-UX GUI.

The `make_tape_recovery` configuration files

The `make_tape_recovery` command still creates the same three configuration files that `make_net_recovery` does. For example:

```
# pwd
```

⁴ This `cfg` clause is a default clause set up by Ignite-UX. This clause, as shown, does not have any configuration files that define software to be installed. Therefore, it is incomplete and cannot be used to install a system.

```
/var/opt/ignite/recovery/2004-01-12,15:23
# ll *cfg
-rw-r--r-- 1 root      sys          3032 Jan 12 15:24 archive_cfg
-rw-r--r-- 1 root      sys          963 Jan 12 15:24 control_cfg
-rw-r--r-- 1 root      sys         7223 Jan 12 15:24 system_cfg
```

The main difference is that the three configuration files end up in a single LIF file on the tape when the recovery tape is created. The order of precedence here is unimportant since the tape is the only device used and all of the configuration files are concatenated together into one LIF file (CONFIG).

Files created by `make_config`

The `make_config` command is used to create configuration files that provide enough information to allow the installation of software bundles from Software Distributor (SD) depots. The `make_config` command is covered in a later section.

Using the `manage_index` command

A variety of uses

The `manage_index` command has many different options. This section describes every form the command can take, as well as its uses.

You should never manually maintain INDEX files. Instead, you should always use `manage_index` to maintain them. The `manage_index` command does not maintain any formatting or comments that may have been added to index files by an Ignite-UX administrator.

The following are all the forms defined in `manage_index(1M)`:

```
manage_index -a -f config_filename [-c cfg_clause_name| -r release]
              [-p] [-v] [-i index_filename]

manage_index -a -s script_file_name [-p] [-v] [-i index_filename]

manage_index -d -c cfg_clause_name | -r release [-p] [-v]
              [-i index_filename]

manage_index -e -c cfg_clause_name [-p] [-v] [-i index_filename]

manage_index -l -c cfg_clause_name | -r release [-o] [-v]
              [-i index_filename]

manage_index -l -o [-v] [-i index_filename]

manage_index -m old_clause_name -c new_clause_name [-p] [-v]
              [-i index_filename]

manage_index -n existing_clause_name -c new_clause_name [-p] [-v]
              [-i index_filename]

manage_index -t -f config_file_name [-c cfg_clause_name| -r release]
              [-p] [-v] [-i index_filename]
```

```
manage_index -t -s script_file_name [-p] [-v] [-i index_filename]  
manage_index -w -c cfg_clause_name [-v] [-i index_filename]  
manage_index -x -c cfg_clause_name [-v] [-i index_filename]
```

Note:

The examples presented are for illustration purposes only; they show the intended purpose of the command and use fictional configuration files, index files, and scripts.

Although the **-p** (preview, make no changes) and **-v** (verbose) options are available to most or all forms of `manage_index`, they are not discussed with any of the examples.

Adding a configuration file to a clause or "release"

When you need to add a configuration file to a clause or release you need to use the **-a** option. The **-i** option must be given a fully qualified path if it is used (it cannot be relative).

```
manage_index -a -f config_filename [-c cfg_clause_name] [-r release]  
[-p] [-v] [-i index_filename]
```

Neither the **-c** nor **-r** options are required. If you use the **-c** option, you can explicitly name the configuration clause to which to add the configuration file. If you use the **-r** option to give an HP-UX release identifier (for example, `B.11.23`), the configuration file will be added to each configuration clause that satisfies one of the following conditions:

- One of the configuration files in a `cfg` clause defines the `release` keyword with a value that matches the release identifier given with the **-r** option.
- One of the configuration files in a `cfg` clause starts with the path `/var/opt/ignite/data/Rel_`, and the value of the release identifier immediately after it matches the release identifier given with the **-r** option.
- One of the configuration files in a `cfg` clause starts with the path `/opt/ignite/data/Rel_`, and the value of the release identifier immediately after it matches the release identifier given with the **-r** option.

Without the **-c** or **-r** option, the value of the release identifier is determined in the same way that clauses are matched with the **-r** option. The following, in order, is performed to try to determine the release identifier that will be used:

- The configuration file defines the `release` keyword, and the value of the `release` keyword provides the value for the release identifier.
- The configuration file starts with the path `/var/opt/ignite/data/Rel_`, and the value of the release identifier from this path provides the value for the release identifier.
- The configuration file starts with the path `/opt/ignite/data/Rel_`, and the value of the release identifier from this path provides the value for the release identifier.

Once a release identifier has been determined for the configuration file, the `manage_index` command behaves as though it was given that release identifier with the `-r` option.⁵

In the following example, if you subsequently tried to add a configuration file based upon a release, it would not work because none of the configuration files contains an operating system release in their path names. However, you can add a file based upon a release if you first add the expected information in the path.

```
$ cat INDEX
cfg "testing" {
    description "testing clause"
}
$ touch config_a config_b
$ print "release=B.11.11" > config_c
$ manage_index -a -f /var/tmp/config_c -c "testing" \
> -i /var/tmp/INDEX
$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
}
```

The following example adds the configuration file `/opt/ignite/data/Rel_B.11.11/config` to the `cfg` clause that then enables you to add the configuration file `/var/tmp/config_a`.

```
manage_index -a -f /opt/ignite/data/Rel_B.11.11/config -c "testing" \
> -i /var/tmp/INDEX
$ manage_index -a -f /var/tmp/config_a -r B.11.11 \
> -i /var/tmp/INDEX
$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
```

Adding scripts to the **INDEX** file

As part of the syntax for configuration index files, there is a keyword called `scripts`, which enables you to define user-selectable scripts that can be run during installation and recovery:

```
manage_index -a -s script_file_name [-p] [-v] [-i index_filename]
```

The `manage_index` command with the `-a` and `-s` options allows you to add scripts into an `INDEX` file. For example:

```
$ manage_index -a -s /var/tmp/script_a \
> -i /var/tmp/INDEX
$ cat INDEX
...
```

⁵ This information is applicable to all forms of the `manage_index` command that need to search for or match release identifiers.

```

scripts {
    "/var/tmp/script_a"
}

```

On the **Advanced** tab in the Ignite-UX GUI, the sample script appears so it can be selected and run by a user. The `scripts` clause in the `INDEX` file is global in scope; scripts defined here are available for selection with any `cfg` clause in the `INDEX` file. You should be careful not to define scripts that make assumptions about the release they will be run on or are not portable across all HP-UX releases.

Removing `cfg` clauses from an `INDEX` file

The `-d` option is used to remove `cfg` clauses from an `INDEX` file using `manage_index`:

```

manage_index -d -c cfg_clause_name | -r release [-p] [-v]
[-i index_filename]

```

The following example removes the `cfg` clause "testing two" from the `INDEX` file completely:

```

$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
$ manage_index -d -c "testing two" -i /var/tmp/test/INDEX

```

You can see that the `cfg` clause "testing two" has been removed from the `INDEX` file:

```

$ cat INDEX...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"

```

The `-r` option operates on all `cfg` clauses applicable to a release identifier. If you reference `cfg` clauses by release, you need to be careful. In the following example, both `cfg` clauses have a release identifier that matches `B.11.11`, so both are removed from the `INDEX` file.

```

$ cat INDEX.save
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}

```

```

cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
$ manage_index -d -r B.11.11 -i /var/tmp/test/INDEX
$ cat INDEX
...

```

Setting the default `cfg` clause in an `INDEX` file

Using the following command, you can easily set which of the `cfg` clauses in an `INDEX` file is the default `cfg` clause:

```
manage_index -e -c cfg_clause_name [-p] [-v] [-i index_filename]
```

The following example changes the default `cfg` clause between two different `cfg` clauses:

```

$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
$ manage_index -e -c "testing" -i /var/tmp/INDEX
$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}=TRUE
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
$ manage_index -e -c "testing two" -i /var/tmp/INDEX
$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"

```

```

        "/opt/ignite/data/Rel_B.11.11/config"
        "/var/tmp/config_a"
}=TRUE

```

Listing the names of `cfg` clauses in an `INDEX` file

The following syntax for the `manage_index` command lists the names of all of the `cfg` clauses in an `INDEX` file:

```
manage_index -l -c cfg clause name | -r release ] [-o] [-v]
[-i index filename]
```

For example:

```
$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}=TRUE
$ manage_index -l -i /var/tmp/INDEX
testing
testing two
```

Being able to list the `cfg` clauses also makes it easier to script Ignite-UX tasks. An example of this could be a sanity-checking script that ensures all of the configuration files referenced from within `cfg` clauses exist and are readable.⁶ Such a script might look like:

```
#!/usr/bin/sh
/opt/ignite/bin/manage_index -l | while read CFG
do
{
eval /opt/ignite/bin/manage_index -w -c \"${CFG}\\" | while read FILE
do
if [[ ! -f $FILE ]]
then
print -u2 "Error: file $FILE does not exist"
fi
if [[ ! -r $FILE ]]
then
print -u2 "Error: file $FILE is not readable"
fi
done
}
done
```

⁶ Note that this script is a trivial example. The `inst1_adm` command with the `-T` option would perform these tasks as well as sanity-check the syntax within the configuration files.

Listing the name of the default `cfg` clause in an `INDEX` file

This form of the `manage_index` command enables you to see which `cfg` clause is the current default `cfg` clause:

```
manage_index -l -o [-v] [-i index_filename]
```

For example:

```
$ cat INDEX
...
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}=TRUE
$ manage_index -l -o -i /var/tmp/INDEX
Default config clause is "testing two"
    in index file: /var/tmp/INDEX
```

Renaming a `cfg` clause in an `INDEX` file

The following form of `manage_index` enables you to rename a `cfg` clause. (This may be useful when you need to swap two clauses.)

```
manage_index -m old_clause_name -c new_clause_name [-p] [-v]
[-i index_filename]
```

The following example swaps two `cfg` clauses:

```
$ manage_index -l -i /var/tmp/INDEX
testing
testing two
$ manage_index -m "testing two" -c "t two" -i /var/tmp/INDEX
$ manage_index -l -i /var/tmp/INDEX
testing
t two
$ manage_index -m "testing" -c "testing two" -i /var/tmp/INDEX
$ manage_index -l -i /var/tmp/INDEX
testing two
t two
$ manage_index -m "t two" -c "testing" -i /var/tmp/INDEX
$ manage_index -l -i /var/tmp/INDEX
testing two
testing
```

Creating a new `cfg` clause from an existing clause

Using the following form of `manage_index`, you can easily to create a new `cfg` clause from an existing clause:

```
manage_index -n existing_clause_name -c new_clause_name [-p] [-v]
[-i index_filename]
```

For example:

```
$ manage_index -n "testing" -c "old testing two" -i /var/tmp/INDEX
$ manage_index -l -i /var/tmp/INDEX
testing two
testing
old testing two
```

Removing a configuration file from a `cfg` clause

Earlier you learned how to add a configuration file to a clause. Now you can do the opposite, which is to remove configuration files from `cfg` clauses:

```
manage_index -t -f config_file_name [-c cfg_clause_name] -r release
[-p] [-v] [-i index_filename]
```

In the following example, a configuration file is removed from a `cfg` clause and then, a file is removed from all of the `cfg` clauses for one release:

```
$ cat INDEX
...
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}=TRUE
cfg "old testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
$ manage_index -t -f "/opt/ignite/data/Rel_B.11.11/config" -c "testing" \
> -i /var/tmp/INDEX
$ cat INDEX
...
cfg "testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
```

```

cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/var/tmp/config_a"
}=TRUE
cfg "old testing two" {
    description "testing clause"
    "/var/tmp/config_c"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}

```

In the previous example, you removed a specific configuration file from the `cfg` clause "testing". The following example removes `config_c` from all B.11.11 clauses:

```

$ manage_index -t -f /var/tmp/config_c -r B.11.11 \
> -i /var/tmp/INDEX
$ cat INDEX
...
cfg "testing two" {
    description "testing clause"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/var/tmp/config_a"
}=TRUE
cfg "old testing two" {
    description "testing clause"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}

```

The `cfg` clause "testing" is not modified in this example. Earlier in the section Adding a configuration file to a clause or "release" on page 6, we discussed how `manage_index` associates a release identifier with a `cfg` clause. The configuration files associated with the `cfg` clause "testing" no longer meets any of the conditions that allows `manage_index` to determine a release identifier. Since no release identifier could be determined, the configuration file `/var/tmp/config_c` was not removed from the `cfg` clause.

Important:

A `cfg` clause must have a `release` keyword in one of the configuration files associated with it.

Removing a script from an **INDEX** file

Earlier you learned how to add a script into the configuration so that it appears on the **Advanced** tab in the Ignite-UX GUI. You use the following form of the command to remove the script:

```
manage_index -t -s script_file_name [-p] [-v] [-i index_filename]
```

The following example adds a script and then removes it:

```

$ manage_index -a -s /var/tmp/script_a -i /var/tmp/INDEX
$ cat INDEX
...
cfg "testing two" {
    description "testing clause"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/var/tmp/config_a"
}=TRUE
cfg "old testing two" {
    description "testing clause"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
scripts {
    "/var/tmp/script_a"
}
$ manage_index -t -s /var/tmp/script_a -c "testing" \
> -i /var/tmp/INDEX
$ cat INDEX
cfg "testing two" {
    description "testing clause"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}
cfg "testing" {
    description "testing clause"
    "/var/tmp/config_c"
    "/var/tmp/config_a"
}=TRUE
cfg "old testing two" {
    description "testing clause"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/var/tmp/config_a"
}

```

Files (or in this case a script) must exist before they can be removed (or added) with `manage_index`. This applies to configuration files as well as scripts.

```

$ manage_index -a -s /var/tmp/script_a -i /var/tmp/INDEX
ERROR: Cannot read script file: /var/tmp/script_a
ERROR: No read access to file specified by -f

```

List the names of all configuration files in a `cfg` clause

The following command enables you to list all the configuration files named in a `cfg` clause:

```
manage_index -w -c cfg_clause_name [-v] [-i index_filename]
```

For example:

```

$ manage_index -w -c "testing" -i /var/tmp/INDEX
/var/tmp/config_c
/var/tmp/config_a

```

In conjunction with the earlier form of `manage_index` used to list out the names of `cfg` clauses in an `INDEX` file, the preceding command can be used to ensure that all of the configuration files referenced in an `INDEX` file are present and, as far as can be determined, correct.

Display the description of a `cfg` clause

This last form of `manage_index` is used to display the description of the `cfg` clause:

```
manage_index -x -c cfg_clause_name [-v] [-i index_filename]
```

For example:

```
$ manage_index -x -c "testing" -i /var/tmp/INDEX
testing clause
```

Using the `make_bundles` command

Why do you need to use `make_bundles`?

The `make_bundles` command is necessary when using Ignite-UX because commands like `make_config` only work with software that is contained within a bundle. Therefore, if you create a configuration file for a depot and it contains software that is packaged only as an SD product, it is not listed in a configuration file created by `make_config`.

Choosing which form of `make_bundles` to use

The `make_bundles` command has the following three forms:

```
/opt/ignite/bin/make_bundles {-b|-B} [-i] [-n name] [-t title] [-c
category] [-o psf] [-r revision] depot_path
/opt/ignite/bin/make_bundles [-b] [-p|-f] [-i] [-c category] [-o psf]
depot_path
/opt/ignite/bin/make_bundles {-p|-f|-B} [-i] [-c category] [-o psf]
[-l file| product/fileset...] depot_path
```

The `make_bundles` command is primarily controlled by the `-B` or `-b` options.⁷ The `-b` option causes `make_bundles` to operate on all product filesets⁸ that are not contained within a bundle.⁹ The `-B` option causes all product filesets in the depot (or those given on the command line) to be operated on regardless of whether they are currently in a bundle or not. The different forms of the `make_bundles` command, which are covered in this section, create either one or many bundles. Each form has advantages over the other forms depending on what you want to achieve.

⁷ In some forms of `make_bundles`, the use of `-b` or `-B` is optional. However, in the most useful forms of the `make_bundles` command, you will almost always use one or the other option.

⁸ You cannot have a standalone fileset because it must be contained within a product.

⁹ If you perform a `swlist` operation on a depot, the default output shows all bundles in the depot followed by all products not contained within a bundle. The `-b` option operates only on products not contained within a bundle.

You can use the `-b` option when you want to operate on all of the unbundled product filesets in a depot and the `-B` if you want to operate on all product filesets within a depot.

You can use the `-p` and `-f` options to modify what the `-b` option does (although you cannot use them with the `-B` option). The `-p` option creates one bundle for every product in the depot, making `-B` unnecessary. When used with the `-b` option, the command only processes products that are not already contained within a bundle. The `-f` option has a similar function to `-p` but it operates at the fileset level; that is, when you use `-f` it causes a bundle to be created for all filesets in the depot (except when `-b` is used as well; then only filesets in products not already contained within a bundle are processed).

The form of `make_bundles` that you use depends upon what you want to do with it. The following sections demonstrate what can be done with three forms of the command.

The `make_bundles` first form

The first form of `make_bundles` uses `-b` or `-B` and not `-p` or `-f`, so you can operate either on all unbundled product filesets or on all product filesets. This form, however, can only create one bundle wrapper.

```
/opt/ignite/bin/make_bundles {-b|-B} [-i] [-n name] [-t title] [-c  
    category] [-o psf] [-r revision] depot_path
```

With the addition of the `-i` option, you can set the `is_reference` SD bundle attribute to TRUE. When set to TRUE, the `is_reference` causes the bundle wrapper to be installed whenever a product or fileset within a bundle is installed. When set to FALSE, the bundle wrapper is installed only when selected. HP-UX patch bundles always have the `is_reference` set to TRUE because they can then be installed with either `patch_match_target` or `autoselect_patches` set to TRUE. In contrast, if `is_reference` is set to FALSE, the bundle wrapper for the the HP-UX patch bundle would not be installed with the patches since SD selects patches to be installed, not bundles (leaving the patches unbundled after installation). This allows the bundle wrapper to be automatically installed when patches have been automatically selected for installation by SD. (This can happen if there are software bundles in the same depot as the patches that are installed, and the patches are for filesets in the bundles.)

Important:

You should always use the `-i` option when defining bundles that contain patches so that the bundle definition is installed when any patch from the bundle is installed.

With the `-n` option, you can give the bundle a name, which must be 16 characters or fewer. With the `-t` option, you can give a longer descriptive title for the bundle, which must be 256 characters or fewer and can only contain one line. With the `-c` option you can assign a category to the bundle; the default, if not given, is to make the bundle uncategorized. See the section "Categories and other Ignite-UX software attributes" on page 118 for more information on software categories.

With the `-o` option, `make_bundles` writes out a Product Specification File (PSF) instead of modifying the target depot. This option is useful if you need to make customizations to the PSF before applying the changes to the depot. The PSF contains the `swpackage` command required to

package the bundle in the depot. The example `swpackage` command contains the path given to PSF "as is". You may need to change the path to the PSF file in the example if you change your current working directory or move the PSF file.

With the `-r` option, you can specify a revision for a bundle. This allows you to create bundle wrappers with revision numbers, enabling you to better manage changes to software.

The examples in this section operate on the following depot:

```
# swlist -d @ /var/opt/ignite/depots/Rel_B.11.23/patches
# Initializing...
# Contacting target "test"...
#
# Target: test:/var/opt/ignite/depots/Rel_B.11.23/patches
#
#
# No Bundle(s) on test:/var/opt/ignite/depots/Rel_B.11.23/patches
# Product(s):
#
PHCO_31622      1.0          Cumulative changes to frecover(1M)
PHCO_31634      1.0          Cumulative changes to pax(1)
PHCO_33431      1.0          tar(1) catalogs ; directory timestamp
PHCO_36056      1.0          find(1) cumulative patch
```

Because there are no existing bundles, using the `-b` option or the `-B` options achieves the same result.

The following example bundles all the patches:

```
# make_bundles -b -i -n "PB_Sept_2007" \
> -t "backup commands patches for Sept 2007" \
> -r 1.0 /var/opt/ignite/depots/Rel_B.11.23/patches
Generating list of unbundled filesets...

===== 08/30/07 16:25:12 EST BEGIN swpackage SESSION

* Session started for user "root@test".

* Source:           test:/var/tmp/psf.28120
* Target:           test:/var/opt/ignite/depots/Rel_B.11.23/patches
* Software selections:
  *

* Beginning Selection Phase.
* Reading the Product Specification File (PSF)
  "/var/tmp/psf.28120".
* Reading the bundle "PB_Sept_2007" at line 11.

* Selection Phase succeeded.

* Beginning Analysis Phase.
* Analysis Phase succeeded.

* Beginning Package Phase.

* Packaging the bundle
```

```

    "PB_Sept_2007,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
* Package Phase succeeded.

===== 08/30/07 16:25:13 EST  END swpackage SESSION

With swlist, you can verify that all of the patches are now contained within the bundle:

# swlist -d @ /var/opt/ignite/depots/Rel_B.11.23/patches
# Initializing...
# Contacting target "test"...
#
# Target: test:/var/opt/ignite/depots/Rel_B.11.23/patches
#
#
# Bundle(s):
#
PB_Sept_2007          1.0          backup commands patches for Sept 2007

```

To continue the example, assume you missed a patch from the bundle that you created and you still need to add it in. In this case, you need to use the **-B** option in a similar command to what was used to create the bundle initially. Therefore, after you copy in the latest ls patch PHCO_33976, you have the following in the depot:

```

# swlist -d @ /var/opt/ignite/depots/Rel_B.11.23/patches
# Initializing...
# Contacting target "test"...
#
# Target: test:/var/opt/ignite/depots/Rel_B.11.23/patches
#
#
# Bundle(s):
#
PB_Sept_2007          1.0          backup commands patches for Sept
2007
#
# Product(s) not contained in a Bundle:
#
PHCO_33976            1.0          ls(1) cumulative patch

```

You can then run `make_bundles` again¹⁰ but this time we can't use the **-b** option because we want all of the patches in the depot to be placed into the bundle. To do that we need to use the **-B** option instead, as **-b** only operates on product filesets not currently in a bundle.

```

make_bundles -B -i -n "PB_Sept_2007" \
> -t "backup commands patches for Sept 2007" \
> -r 1.1 /var/opt/ignite/depots/Rel_B.11.23/patches

===== 08/30/07 16:40:40 EST  BEGIN swpackage SESSION

* Session started for user "root@test".

```

¹⁰ The reason for specifying **-r 1.1** is to identify a collection of patches by a unique bundle specification composed of the bundle name and revision. Revision 1.0 of the bundle `PB_July_2004` does not include `PHCO_30150`; revision 1.1 does include `PHCO_30150`. If the option **-r 1.0** were specified, the original bundle would be replaced.

```

* Source:          test:/var/tmp/psf.28291
* Target:         test:/var/opt/ignite/depots/Rel_B.11.23/patches
* Software selections:
*
* Beginning Selection Phase.
* Reading the Product Specification File (PSF)
  "/var/tmp/psf.28291".
* Reading the bundle "PB_Sept_2007" at line 11.
*
* Selection Phase succeeded.

* Beginning Analysis Phase.
* Analysis Phase succeeded.

* Beginning Package Phase.

* Packaging the bundle
  "PB_Sept_2007,r=1.1,a=HP-UX_B.11.23_IA/PA,v=HP".
* Package Phase succeeded.
=====
  08/30/07 16:40:41 EST END swpackage SESSION

```

Now the depot lists two revisions of the bundle:

```

# swlist -d @ /var/opt/ignite/depots/Rel_B.11.23/patches
# Initializing...
# Contacting target "test"...
#
# Target:  test:/var/opt/ignite/depots/Rel_B.11.23/patches
#
#
# Bundle(s):
#
PB_Sept_2007      1.0          backup commands patches for Sept 2007
PB_Sept_2007      1.1          backup commands patches for Sept 2007

```

By looking at the filesets for revision 1.1 or PB_Sept_2007, you can also verify that the new bundle was created with the correct contents:

```

swlist -d -l fileset PB_Sept_2007,r=1.1 @
/var/opt/ignite/depots/Rel_B.11.23/patches
# Initializing...
# Contacting target "test"...
#
# Target:  test:/var/opt/ignite/depots/Rel_B.11.23/patches
#
#
# PB_Sept_2007          1.1          backup commands
patches for Sept 2007
# PB_Sept_2007.PHCO_31622          1.0          Cumulative
changes to frecover(1M)
  PB_Sept_2007.PHCO_31622.SYS-ADMIN          1.0          OS-Core.SYS-
ADMIN
...
  PB_Sept_2007.PHCO_31634.UX2-CORE          1.0          OS-Core.UX2-CORE
  PB_Sept_2007.PHCO_31634.UX2-CORE          1.0          OS-Core.UX2-CORE

```

In the preceding example, note that most of the output is not shown. If you wanted the ls patch PHCO_33976 to have its own wrapper bundle, you could have instead used -b and changed the bundle name, revision, and so forth and created a new bundle to hold the patch.

The following example demonstrates what happens when you use the -o option to create a PSF. It also shows how you can then use the PSF with swpackage to implement the changes yourself.

Before creating the PSF however we need to remove the existing bundle wrappers since they will prevent the PSF from being created with the -b option:

```
# swmodify -u -d PB_Sept_2007,r=\* @ /var/opt/ignite/depots/Rel_B.11.23/patches
```

Now we can create the PSF file.

```
# make_bundles -b -i -n "PB_Sept_2007" \
> -t "backup commands patches for Sept 2007" \
> -o ./PB-Sept_2007.psf -r 1.0 \
> /var/opt/ignite/depots/Rel_B.11.23/patches
Generating list of unbundled filesets...
# ll ./PB-Sept_2007.psf
-rw-r--r-- 1 root      sys          1011 Aug 30 16:53 ./PB-Sept_2007.psf
```

The PSF contains all of the information needed to create the bundle and the swpackage command needed to create the bundle.¹¹

```
# cat PB-Sept_2007.psf
#####
# swpackage product specification file generated by make_bundles
# on Thu Aug 30 16:53:26 EST 2007. For depot:
/var/opt/ignite/depots/Rel_B.11.2
3/patches
#
# To run swpackage to apply the bundle definitions to the
# depot, use the command:
#
#   swpackage -s ./PB-Sept_2007.psf -xlayout_version=1.0 -
xreinstall_files=true
-d /var/opt/ignite/depots/Rel_B.11.23/patches
#
#####
bundle
tag PB_Sept_2007
title backup commands patches for Sept 2007
os_name HP-UX
is_reference true
revision 1.0
architecture HP-UX_B.11.23_IA/PA
vendor_tag "HP"
contents PHCO_31622,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP,fr=1.0
contents PHCO_31634,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP,fr=1.0
contents PHCO_33431,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP,fr=1.0
contents PHCO_33976,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP,fr=1.0
contents PHCO_36056,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP,fr=1.0
end
```

You can then package the bundle manually:

¹¹ Getting the make_bundles command to produce a PSF also means that the file can be kept in a source management system so changes can be tracked and monitored. Because the PSF can be checked out, the same configuration can be validated and recreated at will.

```

# swpackage -s ./PB-Sept_2007.psf -xlayout_version=1.0 \
> -xreinstall_files=true -d /var/opt/ignite/depots/Rel_B.11.23/patches
=====
  08/30/07 16:55:31 EST  BEGIN swpackage SESSION

* Session started for user "root@test ".

* Source:          test:./PB-Sept_2007.psf
* Target:          test:/var/opt/ignite/depots/Rel_B.11.23/patches
* Software selections:
  *

* Beginning Selection Phase.
* Reading the Product Specification File (PSF)
  "./PB-Sept_2007.psf".
* Reading the bundle "PB_Sept_2007" at line 11.

* Selection Phase succeeded.

* Beginning Analysis Phase.
* Analysis Phase succeeded.

* Beginning Package Phase.

* Packaging the bundle
  "PB_Sept_2007,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP" .
* Package Phase succeeded.

=====
  08/30/07 16:55:31 EST  END swpackage SESSION

```

This new revision 1.0 bundle is the same as the previous 1.1 bundle since we did not remove the ls patch from the depot.

Running `make_bundles` with the `-B` option to repackage the bundle has the same effect:

```
# make_bundles -B -i -n "PB_Sept_2007" \
> -t "backup commands patches for Sept 2007" \
> -r 1.0 /var/opt/ignite/depots/Rel_B.11.23/patches
```

The `make_bundles` second form

The second form of `make_bundles` uses the `-b` option. The `-b` option causes `make_bundles` to operate on all products and filesets that are not contained within a bundle.

```
/opt/ignite/bin/make_bundles [-b] [-p|-f] [-i] [-c category] [-o psf]
depot_path
```

The `-p` and `-f` options specify that multiple bundles are to be created.

- If `-p` is specified and `-b` is not, a bundle is created for each product in the depot.
- If both `-p` and `-b` are specified, a bundle is created for every product in the depot which does not already belong to a bundle.
- If `-f` is specified and `-b` is not, a bundle is created for each fileset in the depot.
- If both `-f` and `-b` are specified, a bundle is created for each fileset in the depot which does not already belong to a bundle.

The `-i`, `-c`, and `-o` options are discussed in the "The `make_bundles` first form" section.

The following examples show what happens when you use the `-p` and `-f` options with the `make_bundles` command. These examples do not have a depot with any bundles, so the `-b` option has no effect.

The first example uses the `-p` option to show bundles being produced at the product level:

```
===== 08/31/07 14:24:04 EST BEGIN swpackage SESSION

    * Session started for user "root@test".

    * Source:          test:/var/tmp/psf.29896
    * Target:          test:/var/opt/ignite/depots/Rel_B.11.23/patches
    * Software selections:
        *

    * Beginning Selection Phase.
    * Reading the Product Specification File (PSF)
      "/var/tmp/psf.29896".
    * Reading the bundle "b_PHCO_31622" at line 11.
    * Reading the bundle "b_PHCO_31634" at line 20.
    * Reading the bundle "b_PHCO_33431" at line 29.
    * Reading the bundle "b_PHCO_33976" at line 38.
    * Reading the bundle "b_PHCO_36056" at line 47.

    * Selection Phase succeeded.

    * Beginning Analysis Phase.
    * Analysis Phase succeeded.

    * Beginning Package Phase.

    * Packaging the bundle
      "b_PHCO_31622,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    * Packaging the bundle
      "b_PHCO_31634,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    * Packaging the bundle
      "b_PHCO_33431,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    * Packaging the bundle
      "b_PHCO_33976,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    * Packaging the bundle
      "b_PHCO_36056,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    * Package Phase succeeded.

===== 08/31/07 14:24:05 EST END swpackage SESSION
```

The preceding example covers packaging of all the patches in the depot that are not already part of the bundle. The command would have to be modified to replace `-b` with `-B` to package every patch¹² into a bundle.

Next, the same command is used but with the `-f` option instead of `-p`. Since the `-b` option only operates on unbundled filesets, the previously existing bundle wrappers were removed before the command was run. All of the bundle names are produced at the fileset level:

¹² Actually every product would be placed into an individual bundle if the `-B` option was used instead of `-b`, that is it would not just apply to patches.

```

# make_bundles -b -f -i /var/opt/ignite/depots/Rel_B.11.23/patches
Generating list of unbundled filesets...
Generating swpackage PSF...

===== 08/31/07 14:30:41 EST BEGIN swpackage SESSION

    * Session started for user "root@test".

    * Source:          test:/var/tmp/psf.29995
    * Target:          test:/var/opt/ignite/depots/Rel_B.11.23/patches
    * Software selections:
        *

    * Beginning Selection Phase.
    * Reading the Product Specification File (PSF)
      "/var/tmp/psf.29995".
    * Reading the bundle "b_PHCO_31622_SYS" at line 11.
    * Reading the bundle "b1_PHCO_31622_SY" at line 20.
    ...
    * Reading the bundle "b1_PHCO_36056_CM" at line 299.
    * Reading the bundle "b_PHCO_36056_CMI" at line 308.

    * Selection Phase succeeded.

    * Beginning Analysis Phase.
    * Analysis Phase succeeded.

    * Beginning Package Phase.

    * Packaging the bundle
      "b10_PHCO_33431_U,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    * Packaging the bundle
      "b11_PHCO_33431_U,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    ...
    * Packaging the bundle
      "b_PHCO_36056_CMD,r=1.0,a=HP-UX_B.11.23_IA,v=HP".
    * Packaging the bundle
      "b_PHCO_36056_CMI,r=1.0,a=HP-UX_B.11.23_IA/PA,v=HP".
    * Package Phase succeeded.

===== 08/31/07 14:30:44 EST END swpackage SESSION

```

Note:

The preceding example merely demonstrates how the `-f` option works - patches need to be installed fully at the product level, not at the fileset level. You should not package and attempt to install patches in this way.

The `make_bundles` third form

The third form of `make_bundles` uses `-p`, `-f`, or `-B` (all optional) with either a list of products and filesets or a list contained in a file given with the `-l` option to process. This form of `make_bundles` gives you improved control over exactly what is packaged into a bundle regardless of what fileset in which it may already be bundled.

```
/opt/ignite/bin/make_bundles {-p|-f|-B } [-i] [-c category] [-o psf]  
[-l file| product/fileset...] depot path
```

Note:

With this form of the command you cannot give the bundles you are creating a name. The most useful form of this command involves giving the -o option to create a PSF that can be used for subsequent packaging.

In this example, you create some of the same bundles that were created in previous examples. However, this time you use an explicit list of products to include (or patches in this case):

```
# make_bundles -B -i -o ./PB_Sept_2007_1.0.psf -r 1.0 PHCO_31622 \  
> PHCO_31634 PHCO_33431 PHCO_36056 \  
> /var/opt/ignite/depots/Rel_B.11.23/patches  
# ll ./PB_Sept_2007_1.0.psf  
-rw-r--r-- 1 root sys 951 Aug 31 16:46 ./PB_Sept_2007_1.0.psf
```

Next, you create another PSF but this time with the extra pax patch PHCO_33967:

```
# make_bundles -B -i -o ./PB_Sept_2007_1.1.psf -r 1.1 PHCO_31622 \  
> PHCO_31634 PHCO_33431 PHCO_36056 PHCO_33976 \  
> /var/opt/ignite/depots/Rel_B.11.23/patches  
# ll ./PB_Sept_2007_1.1.psf  
-rw-r--r-- 1 root sys 1015 Aug 31 16:47 ./PB_Sept_2007_1.1.psf
```

You can now apply the commands stored in these two PSF files against the depot. When you do, you end up with the same results as with the make_bundles first form:

```
# swlist -d @ /var/opt/ignite/depots/Rel_B.11.23/patches  
# Initializing...  
# Contacting target "test"...  
#  
# Target: test:/var/opt/ignite/depots/Rel_B.11.23/patches  
#  
#  
# Bundle(s):  
#  
patches 1.1 /var/opt/ignite/depots/Rel_B.11.23/patches  
patches 1.0 /var/opt/ignite/depots/Rel_B.11.23/patches
```

With the third form, you cannot specify an explicit name. You must edit the PSF to change the name of the bundle.

Tip:

If you use change control it is better to generate a PSF and place it under change control than to use make_bundles to generate a temporary PSF and apply the changes to the depot. By placing the PSF in your change process, you know exactly how a bundle wrapper in a depot was created and what was in it. This allows you to, at some later time, recreate bundles exactly as they were before.

Using the `instl_dbg` command

Introduction

The `instl_dbg` command is often overlooked but offers valuable functionality in testing the effect of a change in a configuration file. This section introduces some of the useful functionality of `instl_dbg` and how it can be combined with the `itool` command to rapidly prototype and test configuration changes.

Requirements

The `instl_dbg` command requires some information that can only be created during an installation session either by adding a new client to an Ignite-UX server, by running `make_net_recovery`, or by running `make_tape_recovery` with the `-s` option. The files that are most needed are `hw.info`¹³, `host.info`, and `config.sys`.

```
/opt/ignite/bin/instl_dbg -D client_directory [-f file] [-v] [-a {l|r}]  
[-cdls|-A] [-V var[=value]] [-S swsel[=TRUE|!=FALSE]]  
[-U use_model[=TRUE|!=FALSE]] [-?]
```

The one command-line option required by `instl_dbg` is the `-D` option. This would usually point to a per-client directory on an Ignite-UX server. It must contain all of the information that `instl_dbg` needs to run. If you add a new client to an Ignite-UX server, the information you need is created automatically. However, this does not help when you have a new client that has never had an operating system installed.

Using the `itool` command

The `itool` command provides the Wizard and Advanced interfaces during recovery and installation. It is not well known that the `itool` command can be run from the command line on an Ignite-UX server to manipulate the configuration for a client. The `itool` command accepts the following options¹⁴:

```
# itool -?  
itool: illegal option -- ?  
usage: itool [ arguments ]  
      -d <directory path>  
      -i <client name>  
      -m 'demo', 'push', 'pull', 'quick', or 'wiz'  
      -w show welcome screen in push mode  
(The default is -m pull)
```

The only method you really need to worry about when using `itool` is as follows:

```
# itool -d <client dir> -m pull
```

¹³ As of Ignite-UX C.7.0 the `io.info` file is also required.

¹⁴ HP does not recommend the use of `itool` outside of its normal use (running on a client system during installation or recovery or via the Ignite user interface). `ignite(1)`. This information is presented here so `itool` can be run in conjunction with `instl_dbg` for rapid prototyping of configurations.

This runs the `itool` user interface as though it were running on the client system.¹⁵ You must run `itool` as root.

Important:

Do not change time-related information in the Ignite-UX GUI unless the system you are on is considered non-production; altering the time information changes the system clock.

Configuration changes are saved into the configuration file in the per-client directory.

When exiting `itool`, there is no need for concern with respect to the list of options it presents; no action should be taken for any of them. The different options set a different return code from `itool`. The return code is how `itool` communicates the next action to the program that calls it when it is running during a normal installation or recovery.

Combining `instl_dbg` and `itool`

If a configuration file (a file named `config`) exists in the client directory, its configuration information is used (for example, which `cfg` clause) instead of the default. This is the reason the `itool` program was discussed in the previous section.

When you use `itool` to change configurations and then run `instl_dbg` to see what is happening in the configurations, it makes it possible to test configurations and changes without having to boot a system and attempt to install the system.

Running `instl_dbg` in this way enables you to see the configuration the way Ignite-UX parses it on that client.

Running `instl_dbg`

You can use `instl_dbg` to test some aspects of your configuration; for instance, with the `-c` option you can perform system configuration sanity checks. In the following example, it is assumed that you have changed working directories to a per-client directory on an Ignite-UX server (`/var/opt/ignite/clients/<MAC>`).

```
# instl_dbg -D . -c
=====
System Configuration Checks =====
WARNING: The disk at: 0/8/0/0.2.0 (SEAGATE_ST39173WC) appears to contain a
file
        system and boot area. Continuing the installation will destroy any
existing data on this disk.
```

The `-c` option allows you to see configuration check messages that would normally only be seen during an installation or recovery session. This allows you to preview what you will see during an installation or recovery session.

¹⁵ When an installation is controlled from the server, `itool` is actually run on the server. However, it is run with the `-m` push option since configuration changes are pushed to the client system.

With the `-d` option, you can get `instl_dbg` to print out the expected size of the volumes it creates and what the impacts are for each file system. The expected output of percent (%) used is also shown with the file system type. If the disk does not contain enough space, the amount `shy` (KB) column indicates the additional space needed.

```
# instl_dbg -D . -d
===== Disk and Volume allocation =====
  Mount      Size    Usage   Disk /      impact      amount
  Directory   (Mb)          Group     (MB / %)      shy (KB)
-----
/stand        300    HFS      vg00       21    7%       0
(swap_dump)  1280   SWAP_DUMP vg00        0    0%       0
/            200    VxFS     vg00       92   46%       0
/tmp          200    VxFS     vg00        0    0%       0
/home         20     VxFS     vg00        0    0%       0
/opt          976    VxFS     vg00      723   74%       0
/usr          1296   VxFS     vg00     1035   79%       0
/var          1780   VxFS     vg00      229   12%       0
-----
Size unallocated from group: vg00: 2624Mb
  * Expected physical allocation of volumes on disks:
  *   0/8/0/0.2.0 (/dev/dsk/c0t2d0) (vg00)
  *     3997 -  311197 KB  /stand
  *     311197 - 1621917 KB  primary
  *     1621917 - 1826717 KB  /
  *     1826717 - 2031517 KB  /tmp
  *     2031517 - 2051997 KB  /home
  *     2051997 - 3051421 KB  /opt
  *     3051421 - 4378525 KB  /usr
  *     4378525 - 6201245 KB  /var
  *     2686976 KB unallocated
```

The `-l` option to `instl_dbg` performs some lint type checks on the configuration. Variables that are not assigned with the `init` keyword appear as follows:

```
# instl_dbg -D . -l
...
***** Variables assigned without init keyword *****
* _my_variable
```

Note:

Every variable used in an Ignite-UX configuration file should be initialized using the `init` keyword. If you do not want the variable to be changed, use the `visible_if` keyword to prevent it from being displayed by the Ignite-UX GUI. (See `instl_adm(1M)` for more information about the `init` and `visible_if` keywords.)

```
# instl_dbg -D . -l
=====
  Lint-type Checks =====
*****
  Variables assigned without init keyword *****
All variables assigned with the init keyword.
```

With the `-s` option, the value of all variables, `sw_sel` objects, and usage models are shown.

The information is quite detailed. From the information for `sw_sel` objects, you can tell if `sw_sel` has been selected, if it has been selected because of a dependency, or if it has been unselected because of an exrequisite.

The variable information provides the name of the variable, the value that it holds, the type of variable, and whether the variable is visible (based on the setting of `visible_if`). If the variable has an "effects" relationship, the name of the variable it has a relationship with is shown, and if the variable has a list of potential values, the list of values is shown.

For use models, the name is shown and whether it is selectable, selected, or visible.

```
# instl_dbg -D . -s

===== Software Selection Attributes: =====
sw_sel      name: "100BaseT-01"
            description: "HP-PB 100BaseT;Supptd HW=A3495A;SW=J2759BA"
            selected: 0
            dep_selected: 0
            ex_unselected: 0
            is_defined: 1

sw_sel      name: "100BaseT-00"
            description: ""
            selected: 0
            dep_selected: 0
            ex_unselected: 0
            is_defined: 1
...
sw_sel      name: "perl"
            description: "Perl Programming Language"
            selected: 0
            dep_selected: 1
            ex_unselected: 0
            is_defined: 1

===== Variable Attributes: =====
Variable   name: "_hp_boot_dev_path"
            value: "0/8/0/0.2.0"
            type: string
            visible: 0

Variable   name: "_hp_default_cur_lan_dev"
            value: "lan5"
            type: string
            visible: 0
...
Variable   name: "_hp_disk_layout"
            value: "Logical Volume Manager (LVM) with VxFS"
            type: string
            visible: 0
            val_list[0]: "Whole disk (not LVM) with HFS"
            val_list[1]: "Logical Volume Manager (LVM) with HFS"
            val_list[2]: "Logical Volume Manager (LVM) with VxFS"
            val_list[3]: "VERITAS Volume Manager (VxVM) with VxFS"
            Effects var: _hp_pri_swap
            Effects var: _hp_group_name

Variable   name: "_hp_root_disk"
            value: "0/8/0/0.2.0"
```

```

        type: string
        visible: 0
Effects var: _hp_pri_swap
Effects var: _hp_min_swap
Effects var: _hp_disk_layout
Effects var: _hp_root_grp_disks
...
===== Use-Model Attributes: =====
Use model name: "Create /export volume"
    selected: 0
    selectable: 1
    visible: 1

Use model name: "Create separate volumes (/usr, /var, ...)"
    selected: 1
    selectable: 1
    visible: 1

```

Using the `-f` option, the `instl_dbg` command also provides a completely parsed version of the configuration. This output builds on top of the configuration files parsed on the system. In this case, the configuration was modified using the Ignite-UX GUI before running `instl_dbg`:

```

# instl_dbg -D . -f ./parsed.output
# cat parsed.output
cfg "HP-UX B.11.11 Default"=TRUE
server="10.0.0.60"
is_net_info_temporary=FALSE
init _hp_keyboard="Not_Applicable"
system_name="spectre"
netmask["lan0"]="255.255.255.0"
ip_addr[]="10.0.0.162"
netmask[]="255.255.255.0"
route_gateway[0]="10.0.0.1"
route_destination[0]="default"
_hp_cfg_detail_level="ibnpscrlLtfh"
#
# Variable assignments
#
init _hp_boot_dev_path="0/8/0/0.2.0"
init _hp_default_cur_lan_dev="lan5"
init _hp_primary_path="0/8/0/0.2.0"
init _hp_current_client_release="B.11.11"
init _hp_HFS_blksize=65536
init _hp_HFS_fragsize=8192
init _hp_VxFS_blksize=8192
init _hp_FS_stripe_size=64K
init _hp_disk_layout="Logical Volume Manager (LVM) with VxFS"
...
init _hp_os_bitness="64"
init _hp_patch_save_files="YES"
init _hp_custom_sys="Current System Parameters"
init _hp_locale="SET_NULL_LOCALE"
init "Create /export volume"=FALSE
init "Create separate volumes (/usr, /var, ...)"=TRUE
sd_command_line=" -x os_release=B.11.11 -x os_name=HP-UX:64 "
#
# Software Selections
#
init sw_sel "100BaseT-01"=FALSE
init sw_sel "ATM-00"=FALSE
init sw_sel "ATM-01"=FALSE
init sw_sel "B5725AA"=TRUE

```

```

...
init sw_sel "b_PHSS_28764"=FALSE
init sw_sel "b_PartitionManag"=FALSE
# sw_sel "perl" will be loaded due to other software.
(cfg "HP-UX B.11.11 Default") {
...
    sw_source "core" {
        description="HP-UX coreSoftware"
        source_type="NET"
        source_format=SD
        sd_server="10.0.0.60"
        sd_depot_dir="/var/opt/ignite/depots/Rel_B.11.11/core"
        sd_command_line=" -xpatch_filter=.* -xpatch_save_files=true
"
        sd_use_ui=FALSE
        load_order=0
    }
    sw_source "Kernel Configuration" {
        source_format=CMD
        load_order=11
    }
    sw_source "site commands" {
        source_format=CMD
        load_order=10
    }
}
RUN_UI=TRUE
RUN_SD=TRUE
SD_USE_UI=FALSE
CONTROL_FROM_SERVER=FALSE
HALT_WHEN_DONE=FALSE
USE_EXPERT_UI=TRUE
CLEAN_ALL_DISKS=FALSE
HIDE_BOOT_DISK=FALSE
ERROR_IF_BAD_SW=FALSE
RECOVERY_MODE=FALSE
DISABLE_DHCP=TRUE
ALLOW_DISK_REMAP=FALSE
release="B.11.11"
#
# System/Networking Parameters
#
_hp_custom_sys+= {"Current System Parameters"}
init _hp_custom_sys="Current System Parameters"
_hp_custom_sys visible_if TRUE
_hp_custom_sys help_text "System/Networking Info"
(_hp_custom_sys=="Current System Parameters") {
    final system_name="testa"
    final ip_addr["lan5"]="10.0.0.162"
    final netmask["lan5"]="255.255.255.0"
    final ip_addr[]="10.0.0.47"
    final netmask[]={}=255.255.255.0"
    final route_gateway[0]="10.0.0.1"
    final route_destination[0]="default"
    TIMEZONE="EST-10EDT"
    ROOT_PASSWORD="/aIc61gXgbz2A"
    is_net_info_temporary=FALSE
    run_setparms=FALSE
} # end "Current System Parameters"

#
# Disk and File systems
#

```

```

_hp_disk_layout+= { "Modified LVM Layout" }
init _hp_disk_layout="Modified LVM Layout"
(_hp_disk_layout=="Modified LVM Layout") {
    # Disk/File system Layout:
    volume_group "vg00" {
        usage=LVM
        physical_volume disk[_hp_root_disk] {
            } # end pv_options
        max_physical_extents=2500
        logical_volume {
            mount_point="/stand"
            disk[_hp_root_disk]
            usage=HFS
            size=307200K
            FILE_LENGTH=LONG
            FRAGSIZE=8192
            BLKSIZE=65536
            bad_block_relocate=FALSE
            contiguous_allocation=TRUE
        } # end logical_volume
    ...
    logical_volume {
        mount_point="/var"
        usage=VxFS
        size=73728K | remaining | 1822720K
        BLKSIZE=8192
        largefiles=TRUE
    } # end logical_volume
} # end volume_group "vg00"
} # end "Modified LVM Layout"

```

Other **instl_dbg** options

The **-vvv** option (very verbose) lists all the selections and values of all variables. You can select use models (for example, "Logical Volume Manager (LVM) with VxFS") and set values for variables to test the impact that a change has on the configuration.

Refer to *instl_dbg(1M)* for more information on other options available.

The **hw.info** and **host.info** files

Both the **hw.info** and **host.info** files can be created in either of two ways:

- When you add a new client to the Ignite-UX server for recovery, the `add_new_client` script creates these two files (indirectly).
- When a new client boots from the Ignite-UX server and creates a directory for itself, information about the client hardware is written to the **hw.info** file.

As of Ignite-UX version C.7.0, the **io.info** file is also created. This is used by Ignite-UX instead of the **hw.info** file.

Creating both files

You can easily create both **host.info** and **hw.info**¹⁶ files on a system by using the same method as the `add_new_client` script. The following example, run on the client system as `root`, places the **host.info** and **hw.info** files in the current working directory:

¹⁶ As of Ignite-UX version C.7.0, you will require the **io.info** file.

```

# export INST_CLIENT_DIR=./
# export PATH=$PATH:/opt/ignite/lbin
# rescan_hw_host
    * Scanning system for IO devices...
    * Querying disk device: 0/0/1/1.15.0 ...
    * Querying disk device: 0/0/2/1.15.0 ...

# ll
total 4
-rw-r--r--  1 root      sys          214 Apr 27 15:31 host.info
-rw-r--r--  1 root      sys          688 Apr 27 15:31 hw.info
# cat host.info
MEMORY=1310720K
HARDWARE_MODEL="9000/800"
MODEL="9000/800/A400-6X"
can_run_32bit=TRUE
can_run_64bit=TRUE
is_numa=FALSE
is_ia64=FALSE
is_hppa=TRUE
_hp_boot_dev_path="0/0/1/1.15.0"
_hp_boot_dev_path visible_if FALSE
# cat hw.info
cdrom: 0/0/1/0.1.0 2 sdisk 188 31 1000 HP_DVD-ROM_304 0 /dev/rdsk/c0t1d0
/dev/dsk/c0t1d0 -1 -1 0 1 0
disk: 0/0/1/1.15.0 0 sdisk 188 31 1f000 SEAGATE_ST336704LC 35566480
/dev/rdsk/c1t15d0 /dev/dsk/c1t15d0 -1 -1 5 1 10
disk: 0/0/2/1.15.0 1 sdisk 188 31 3f000 SEAGATE_ST336704LC 35566480
/dev/rdsk/c3t15d0 /dev/dsk/c3t15d0 -1 -1 4 1 10
lan: 0/0/0/0 0 lan0 btlan 000F201D7D1F HP_PCI_10/100Base-TX_core0
ext_bus: 0/0/1/0 0 c720 n/a SCSI_C896_Ultra_Wide_Single-Ended
ext_bus: 0/0/1/1 1 c720 n/a SCSI_C896_Ultra_Wide_Single-Ended
ext_bus: 0/0/2/0 2 c720 n/a SCSI_C87x_Fast_Wide_Single-Ended
ext_bus: 0/0/2/1 3 c720 n/a SCSI_C87x_Ultra_Wide_Single-Ended
processor: 160 0 processor n/a Processor

```

Note:

The lines are not folded in the actual hw.info file as they appear in the preceding example.

You can test configuration changes with `inst1_dbg` by copying an existing per-client directory to a new directory. However, it is not a good idea to copy the directory to another directory under `/var/opt/ignite/clients` in case the <MAC> directory you choose to copy it to conflicts with a potential future client; instead, copy it somewhere else.

You can either copy the `hw.info` and `host.info` files from the client system or freely change the existing `hw.info` and `host.info` files for testing. For example, the disk size is the field after the description, but before the raw device file name (both disk devices are set to 35566480 in the preceding example).

If you had to test a configuration with 72 GB disks, you could manually change those sizes. While you might consider doing this for testing purposes, you should *never* change the `hw.info` or `host.info` files in a per-client directory on the Ignite-UX server that is used by a "real" client system, as doing so may cause many potentially fatal problems. You should not attempt to

manually change the `io.info` file, this file is considerably more complex than the `hw.info` file and manual changes are likely to leave the file unusable.

Miscellaneous configuration tips

- In Ignite-UX configuration files, `TRUE` and `FALSE` should be stated only in upper- or lower-case and never in mixed case. Ignite-UX does not understand `TRUE` or `FALSE` when they are in mixed case. This applies to other keywords as well.
- Be careful when deciding to use the `=` or `+=` operators. The `+=` operator enables you to add to something, while the `=` operator removes any current value for a variable and replaces it with the new value being assigned. See the "`_hp_disk_layout`" section on page 88 for information about how `=` and `+=` can be used to replace existing disk layouts or add more layout choices. The `error` keyword is a good example of the need to use `+=` at times. If you use `=`, you remove any previous error messages that may have been encountered. If you use `+=`, all of the error messages, both previous and new, can be shown to the user.
- When you are using the Ignite-UX GUI to add new logical volumes, you might want to consider choosing the logical volume that is most similar to what you want to create and use that as a starting point. Change the attributes and instead of selecting **Modify**, then select **Add** to create your new logical volume.
- When setting a network configuration, review `inst1_adm(4)` to verify whether the modifier `final` (placed before a keyword, such as `dns_domain`) is required to retain the configuration when the system is installed. If it is required and you do not provide this modifier, the configuration is only kept for the duration of the installation and is discarded at the final system reboot.
- If you do not want to preserve instance numbers when cloning a system, you should remove all of the `hw_instance_num` lines from the `system_cfg` file. Unlike network recovery archives, you cannot remove `hw_instance_num` lines from a recovery tape after creation. Instead, if you use a recovery tape for cloning and do not want to preserve instance numbers, you must preview using `make_tape_recovery` with the `-p` option, edit the associated configuration file (`/var/opt/ignite/recovery/latest/system_cfg`), and then resume creation of the recovery tape using `make_tape_recovery` with the `-r` option. For a network recovery archive, you would find the `system_cfg` file in the following location:
`/var/opt/ignite/clients/<MAC>/recovery/<DATE-TIME>/system_cfg`
- If you need to create swap spaces larger than 32GB with HP-UX 11i v2 or later releases of HP-UX, you should be running Ignite-UX version C.7.0 or later - previous versions of Ignite-UX do not allow you to do this.

Analyzing the HP-UX default B.11.11 cfg clause

The following is from a default `/var/opt/ignite/INDEX` showing the configuration files referenced by the HP-UX B.11.11 cfg clause:

```
cfg "HP-UX B.11.11 Default" {
    description "This selection supplies the default system configuration
that HP supplies for the B.11.11 release."
    "/opt/ignite/data/Rel_B.11.11/config"
```

```

        "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"
        "/var/opt/ignite/config.local"
}

```

This section describes these configuration files and explains what is happening and why certain things have been done in certain ways.

The release-specific configuration file

The config file defines all of the file system and volume management configurations for the release (and a few other things).

```

#####
# config:
#
# (c) Copyright Hewlett-Packard Company 1994
#
# This file contains the default system definitions used during a
# cold-install of HP-UX. This file should not be modified. If
# modifications are needed, the file "config.local" should be modified
# instead.
# See the instl_adm(4) man page for details.
#
# Changes to this file will not be preserved if the fileset is re-installed.
#
# If you are editing this file via the advanced options during a
# cold-install, then notice that the "config.local" contents are appended
# to the end of this file.
#
# @(#) config $Revision: 10.82 $

```

The release keyword is probably the single-most important keyword in an Ignite-UX configuration file. This keyword *must* be set or else any installation of HP-UX is likely to fail or not work properly. Ignite-UX has no idea what version of HP-UX it is installing unless it is set.

As of Ignite-UX C.6.0.x, the release keyword is checked against the value of the _hp_ikernel_os_release variable (the client creates the value of this variable in the host.info file). If the HP-UX release in the release keyword is not supported by the kernel version given in _hp_ikernel_os_release the installation is not allowed to proceed and the following error appears:

```

ERROR: The version of HP-UX you have chosen to install on the system
(B.11.23) is not supported by the version of the Ignite-UX install
kernel that the system booted (B.11.11). You will need to reboot the
target system from an install kernel matching the desired release
from the menu at the console. If using the bootsys command, use the
'-R' option to specify the install kernel version.

```

Table 2 shows the different HP-UX releases, beginning with Ignite-UX C.7.3.x, which can be installed with the different versions of an installation kernel¹⁷.

Table 2

¹⁷ This assumes that you have installed the necessary filesets to support those HP-UX releases.

HP-UX Release ¹⁸	PA-RISC	Itanium®-based
B.11.11	B.11.11	N/A
B.11.23	B.11.23	B.11.23
B.11.31	B.11.31	B.11.31

Therefore, if you boot the B.11.11 installation kernel you can install HP-UX B.11.11 but you cannot install B.11.23. The reverse applies - if you boot the B.11.23 PA-RISC installation kernel, you cannot install B.11.11 with it.

When you create custom configuration files and do not plan to include the default configuration files you *must* set this variable.

```
release="B.11.11"

#####
#####
```

Some time ago, changes were made to Ignite-UX to increase the block and fragment sizes for VxFS and HFS file systems. This can lead to differences between systems installed in the last few years and much older systems that had smaller default block and fragment sizes (for example, the file system attributes are different so the block and fragment sizes from older systems are probably smaller).

The use of <variable> visible_if false ensures that the variables do not appear in the **Additional** dialog box from the **Basic** tab.

```
init _hp_HFS_blksize = 65536      # HFS block size.
init _hp_HFS_fragsize = 8192      # HFS frag size.
init _hp_VxFS_blksize = 8192      # VxFS block size.
init _hp_FS_stripe_size = 64Kb   # File system stripe size.

_hp_HFS_blksize visible_if false
_hp_HFS_fragsize visible_if false
_hp_VxFS_blksize visible_if false
_hp_FS_stripe_size visible_if false
#####
#####
```

The following example shows the start of a definition for the _hp_disk_layout variable, which is set to one of four possible values.

Around the _hp_disk_layout is a test for is_hppa (is PA-RISC). This only exists in this file since the file is generated from one common source file using unifdef (refer to *unifdef(1)* for more information) for all HP-UX releases. Obviously, this test is not actually needed for HP-UX B.11.11 since it is only supported on PA-RISC systems. The entire release-specific configuration files are based on one original file and they are processed with unifdef to produce the release-specific configuration file.

You actually get to see the value of the variable _hp_disk_layout when you use the Ignite-UX GUI. The values currently assigned to _hp_disk_layout are shown on the **File System:** list on

¹⁸ HP-UX release of the install kernel

the **Basic** tab. The form used in the next example to initialize a list of possible values allows the Ignite-UX GUI to modify the list later.

Further in the configuration file, when giving `_hp_disk_layout` a final value you only assign its value using `init`. When using the form "`init _hp_disk_layout=..."` the Ignite-UX GUI is still allowed to modify the value. If you instead define "`_hp_disk_layout=..."`" without the `init`, the Ignite-UX GUI cannot change this value. If the Ignite-UX GUI were not able to change the value of `_hp_disk_layout`, no changes made in the Ignite-UX GUI to file system attributes are applied to the system¹⁹.

This is because the Ignite-UX GUI must change the value of `_hp_disk_layout` to "Modified LVM Layout", "Modified VxVM Layout", or "Modified Whole-Disk Layout" when changes have been made. If the Ignite-UX GUI is not allowed to change the variable, you cannot keep any file system customizations. This concept becomes very important when creating custom configuration files.

```
#####
# Any _hp_disk_layout variable setting besides those below will disable
# the default disk layout described in this file. This can be
# done if a custom disk layout is defined elsewhere.
#
# For PA architecture only
is_hppa {
    _hp_disk_layout = {
        "Whole disk (not LVM) with HFS",
        "Logical Volume Manager (LVM) with HFS",
        "Logical Volume Manager (LVM) with VxFS",
        "VERITAS Volume Manager (VxVM) with VxFS"
    }
}
```

You can now set the variable `_hp_disk_root` to the value of `_hp_primary_path`, assuming that the device that `_hp_primary_path` points to exists and `_hp_root_disk` does not already have a value. You can use an Extended Regular Expression (ERE)²⁰ to see if `_hp_root_disk` is set to any value²¹.

```
#
# Set the default root disk to the current primary path. The
# _hp_primary_path will be "" if it was set to a non-existent dev.
# also don't reset the _hp_root_disk if it was initialized in
# a different config file and non-null (see 15654FSDdt). Use
# the ~ operator instead of == to detect if _hp_root_disk is not
# yet set (See FSDdt22058).
#
({_hp_primary_path != "" & !(_hp_root_disk ~ ".*")})
{
    init _hp_root_disk=_hp_primary_path
}
```

¹⁹The `itool` command Ignite-UX GUI accepts changes to file system attributes and does not caution you, but no file system attribute changes made in the Ignite-UX GUI are applied to the system.

²⁰For more information regarding ERE, refer to `regexp(5)`.

²¹The ERE `".*"` means zero or more of any character so it matches a zero length string.

The next section of the configuration file presents a critical concept that it is important to understand – the idea of the effects relationship. An effects relationship creates a relationship between two variables so that when one variable changes the other variable is re-evaluated – even if its final value is not based upon that variable.

In the next example, you can see that the value of `_hp_pri_swap` has been set to the size of the root disk and to the value of `_hp_disk_layout` (after that primary swap is initialized to a value more useful for swap).

This forces Ignite-UX to reevaluate the value of `_hp_pri_swap` whenever the size of the root disk changes (a smaller disk might force Ignite-UX to define a smaller swap space) or when the disk layout changes.

```
#  
# The next two statements are used only to establish an effects  
# relationship between the _hp_pri_swap variable and _hp_root_disk &  
# _hp_disk_layout variables. The value is overwritten below.  
# This makes the UI change the swap anytime either of these two  
# values change.  
init _hp_pri_swap = disk[_hp_root_disk].size  
init _hp_pri_swap = _hp_disk_layout  
  
# default (recommended) swap size is 2 X memory  
# Round up to the nearest 64MB so that we don't get a warning  
# about it being adjusted in the UI. Really this just needs  
# to be rounded to be a multiple of the physical-extent-size, but  
# that can vary depending on the size of disk (4,8,16,32,64...)  
# 64MB should work on the majority of the disks available.  
init _hp_pri_swap = (((MEMORY * 2) + 65535KB)/64MB)*64MB
```

Now you have other calculations for swap for different sized disks. The following example limits the size of swap for small disks. If you use `init` when setting both `_hp_pri_swap` and `_hp_sec_swap`, then the configuration only uses methods for setting the variables that allows the Ignite-UX GUI to modify them.

```
# Put an upper bounds of 1024Mb to the default swap size for disks  
# less than 5Gb. And 4Gb swap for other disks  
_hp_pri_swap > 1024Mb {  
    disk[_hp_root_disk].size < 5120Mb {  
        init _hp_pri_swap = 1024Mb  
    } else {  
        _hp_pri_swap > 4096Mb { init _hp_pri_swap = 4096Mb }  
    }  
}  
  
# Use a 128Mb as the default minimum amount of swap configured on any  
# system. The real swap space will be reduced down to _hp_min_swap if  
# there is not enough file system space.  
(_hp_pri_swap < 128Mb)  
{ init _hp_pri_swap = 128Mb }  
  
# Initialize the swap range minimum to what the default is  
_hp_min_swap = _hp_pri_swap  
  
# If the system is limited on resources, then reduce the minimum so  
# that the OS has a better chance of fitting. The swap size will  
# still be set to the recommended value if there is enough disk space.  
(_hp_min_swap > 512Mb & disk[_hp_root_disk].size < 3000Mb)
```

```

        { _hp_min_swap = 512Mb }
(_hp_min_swap > 192Mb & disk[_hp_root_disk].size < 1800Mb)
        { _hp_min_swap = 192Mb }
(_hp_min_swap > 96Mb & disk[_hp_root_disk].size < 700Mb)
        { _hp_min_swap = 96Mb }
(_hp_min_swap > 68Mb & disk[_hp_root_disk].size < 600Mb)
        { _hp_min_swap = 68Mb }

```

Now `_hp_sec_swap` has been initialized with a set of values that can be chosen from; however, they are not a fixed set of values. The variable `_hp_sec_swap` can be set using the **Additional** button because of the way the variable was defined (`<variable>={ value1, value2, ... }`). The variable `_hp_sec_swap` has its value set from a selection list or by typing it into the corresponding field directly. This can happen because `_hp_sec_swap` was not defined as an "enum"; if it had been defined as an enum only a value from the list of values it contains could be used.

Although, in this configuration, `hp_pri_swap` cannot be modified using the **Additional** button. This value appears directly in the Ignite-UX GUI on the **Basic** tab using the **Root Swap (MB)** button and it allows you to select from a list of possible values. The list of values shown in the selection list for `_hp_pri_swap` comes from the list defined in the configuration. If you create a custom disk configuration file, you should supply a value or a list of values for `_hp_pri_swap` (or somehow otherwise calculate a default value for it).

```

_hp_pri_swap = { 128Mb, 256Mb, 512Mb, 1024Mb,
                 1536Mb, 2048Mb, 3072Mb, 4096Mb,
                 8192Mb, 12288Mb, 16384Mb }

_hp_sec_swap = { 0Mb, 128Mb, 256Mb, 512Mb, 1024Mb,
                  1536Mb, 2048Mb, 3072Mb, 4096Mb,
                  8192Mb, 12288Mb, 16384Mb }
init _hp_sec_swap = 0Mb
_hp_sec_swap help_text "Secondary Swap space (KB)"

```

The next example gives `_hp_disk_layout` a value of `disk[_hp_root_disk].model`. If the value of that expression ever changes, it causes the value of `_hp_disk_layout` to be re-evaluated.

Doing this makes good sense in the context of this configuration file because if you change the root disk, you may need to re-evaluate what the disk configuration is. Tests later on may override some defaults for `_hp_disk_layout` in the following example:

Note:

The models being tested on are not likely ever to run HP-UX B.11.11.
This is old data being carried along into later releases.

```

# This next statement is used only to establish an effects relationship
# between the _hp_disk_layout variable and _hp_root_disk. The value is
# overwritten below.
init _hp_disk_layout = disk[_hp_root_disk].model

```

The following section defines some limitations on non-LVM configurations due to firmware limitations in older HP9000 systems - it is somewhat unlikely any of these systems are running HP-UX 11.11 (some of them supported HP-UX revisions as early as 7.x).

```

#
# Define the defaults based on the system architecture 700 vs 800
#
is_hppa {
    HARDWARE_MODEL ~ "9000/7.*" {
        init _hp_disk_layout = "Logical Volume Manager (LVM) with VxFS"
    } else {
        # For a certain class of S800's, non-LVM takes >30 minutes to boot due
        # to sequential access firmware. So make the default LVM on them no
        # matter what size of disk it is.
        disk[_hp_root_disk].size >= 600Mb |
            HARDWARE_MODEL ~ "9000/825.*" |
            HARDWARE_MODEL ~ "9000/834.*" |
            HARDWARE_MODEL ~ "9000/835.*" |
            HARDWARE_MODEL ~ "9000/635.*" |
            HARDWARE_MODEL ~ "9000/845.*" |
            HARDWARE_MODEL ~ "9000/645.*" |
            HARDWARE_MODEL ~ "9000/822.*" |
            HARDWARE_MODEL ~ "9000/832.*" |
            HARDWARE_MODEL ~ "9000/842.*" |
            HARDWARE_MODEL ~ "9000/852.*"
    {
        init _hp_disk_layout = "Logical Volume Manager (LVM) with VxFS"
    } else {
        init _hp_disk_layout = "Whole disk (not LVM) with HFS"
    }
}
#
# For disks over 2Gb+300k the default must be LVM.
# Set the default here, and sanity_checks will ensure it.
#
disk[_hp_root_disk].size > 2097452K {
    init _hp_disk_layout = "Logical Volume Manager (LVM) with VxFS"
}
} else {
    init _hp_disk_layout = "VERITAS Volume Manager (VxVM) with VxFS"
}

```

The enumeration `_hp_root_grp_striped` in the following example is a perfect example of how to ask a yes/no or Boolean question using the **Additional** button on the **Basic** tab in the Ignite-UX GUI. The first step is to define the variable as an enum, then give it two values, an initial value (with the `init` keyword so the user is allowed to change it using the Ignite-UX GUI), and some help text.

Without defining `_hp_root_grp_striped` as an enum, the user would be able to select **YES** or **NO** on the **Additional** button or also enter any value. In this case, it does not matter because the configuration only ever tests on the value **YES** (later) to decide on an action, but if you could enter a lower-case "yes" the tests would fail; hence the enumeration type is used.

```

#
# Boolean describing whether or not the root disk group
# will be striped
#
enum _hp_root_grp_striped
_hp_root_grp_striped = { "YES", "NO" }
init _hp_root_grp_striped = "NO"
_hp_root_grp_striped help_text "Stripe root VG disks?"

```

The expression in the next example is defining a use model. Use models are not variables; they are true / false expressions that are used to control how things are done²². Like variables, use models can be viewed using the **Additional** button of the **Basic** tab in the Ignite-UX GUI. When you are in the **Additional** dialog box, use models look visibly different from variables. Later in the configuration file, is the impact statement that this use model would have when it is set to true.

```
INIT23 "Create /export volume" = false
```

Now you have another effects relationship, this time between a new variable `_hp_root_grp_disks` and `_hp_root_disk`. The variable `_hp_root_grp_disks` is meant to be an integer so when doing the effects relationship you add zero (0) to `disk[_hp_root_disk].size` to ensure that the final value is an integer. Since variables are not explicitly typed, you want to imply an integer type.

The variable `_hp_root_grp_disks` is also an enum; it can take one of the values from 1 up to the value of variable `num_disks`,²⁴ and it has a default initial value of 1.

```
#  
# Number of disks in the root volume group  
# The fist line below is just to define an "effects" relationship  
# between the root disk and _hp_root_grp_disks which dependent  
# upon each other in the logic to follow.  
#  
init _hp_root_grp_disks = (disk[_hp_root_disk].size+0) # +0 convert to int  
enum _hp_root_grp_disks  
_hp_root_grp_disks = { 1..num_disks }  
init _hp_root_grp_disks = 1  
_hp_root_grp_disks help_text "# of disks in root VG"
```

In the next example, you test on the size of the root disk and the number of disks attached to the system. If the size of the root disk is <600MB and there are only two disks, the second disk is automatically included into the root volume group (by setting `_hp_root_grp_disks` to be 2).

When the root disk is <600MB, the default value for `_hp_disk_layout` is "Whole disk (not LVM) with HFS". If there are only two disks in the system, the second disk is included and the default for `_hp_disk_layout` is changed to "Logical Volume Manager (LVM) with VxFS".

```
# If there are only two disks and the "_hp_root_disk" disk is small, default  
to  
# LVM configuration with both disks in root volume group.  
disk[_hp_root_disk].size < 600Mb & num_disks == 2  
{  
    init _hp_disk_layout = "Logical Volume Manager (LVM) with VxFS"  
    init _hp_root_grp_disks = 2  
}
```

²² Unfortunately, variables cannot have effects relationships with use models or software selections.

²³ The "init" keyword is case insensitive, so it can appear is `init` or `INIT`. It cannot appear in mixed case such as "Init". If you do this, Ignite-UX will not recognize it as a keyword.

²⁴ A built-in variable that is a count of the number of disks discovered on the system. Refer, refer to `inst1_adm(4)`.

Next is another use model defined called "Create separate volumes (/usr, /var, ...)". It is used later for `_hp_disk_layouts` other than "Whole disk (not LVM) with HFS" only. By default, this is TRUE (and the Ignite-UX GUI does show separate volumes for /usr, /var, ... by default).

```
INIT "Create separate volumes (/usr, /var, ...)" = TRUE
```

The first disk configuration "Whole disk (not LVM) with HFS" is protected by a test of the variable `_hp_disk_config` to see if it is equal to "Whole disk (not LVM) with HFS".

After that, the `partitioned_disk` keyword starts the definition of a whole disk layout for a system. The `partitioned_disk` disk definition must contain the following:

- A `physical_volume` definition
- An `fs_partition` definition
- A `swap_partition` definition

If you have no `fs_partition` definition that at least defines the `root` file system, the following error appears after pressing **Go!** in the Ignite-UX GUI:

```
ERROR: There is no root volume (mount point = /) defined in the configuration.
```

If you have no `swap_partition` defined, you see the following error after selecting **Go!** in the Ignite-UX GUI:

```
ERROR: There is no swap volume defined in the root volume group (or on the root disk).
```

For the physical volume statement, you need only supply the name of the disk that is used as the root disk (the `disk[]` keyword takes a hardware path or index number and returns the name of the disk). This discussion here assumes that you are defining the boot disk as a whole disk²⁵.

However, if you were writing custom configurations you probably would not worry about whole disk layouts on any system running B.11.11 (you need a 2GB or less disk drive to do this on PA-RISC because of potential firmware limitations and HP-UX B.11.11 recommends at least a 4GB root disk).

```
_hp_disk_layout == "Whole disk (not LVM) with HFS"
{
    partitioned_disk
    {
        physical_volume disk[_hp_root_disk]
```

The file system partition on a partitioned disk must be HFS. It follows from the fact that the PA-RISC boot loader can only read from some types of HFS file system (large files-enabled HFS file systems cannot be booted from on PA-RISC systems). Therefore, the usage statement is set to be HFS. You set the block size and the fragment size²⁶ to the HFS defaults defined earlier.

²⁵ You can have multiple whole disks defined in a configuration. Each whole disk is defined by different `partitioned_disk` definitions and has a different `mount_point`.

²⁶ These concepts are only meaningful in a HFS file system. The block size is the largest block of data a file can have allocated to it (files that can occupy a full free block will do so and the smallest amount of space that can be allocated by a file is the fragment size. In extent-based file systems these concepts are usually meaningless. For example, the smallest allocation unit is 1KB no matter what the fragment size is set to for the file system.

Now size is set to the remaining, which takes whatever space is left over once "other" things (in this case swap at the end of the disk) have been allocated. The mount point naturally is "/" (since it is the root file system). Lastly, if this is a very small disk change minfree to be 5% (you are installing B.11.11 and there is no way you can fit B.11.11 onto a 300MB disk anyway).

```
fs_partition {
    usage = HFS
    blksize = _hp_HFS_blksize
    fragsize = _hp_HFS_fragsize
    size = remaining
    mount_point = "/"
    disk[_hp_root_disk].size < 300Mb {
        # For really small disks, tune down minfree
        # in order to gain some disk space.
        minfree = 5
    }
}
```

Now you have our primary (and presumably only) swap partition definition, you can see that it is swap from the usage. The interesting thing is the size specification. This sets the size to be at least _hp_min_swap but also includes any remaining space up to _hp_pri_swap in size. The swap space in this case attempts to get all of the space that it can (up to _hp_pri_swap) but does not go lower than _hp_min_swap while ensuring that the file system partition gets enough to meet its impacts²⁷ statements requirements.

```
swap_partition {
    usage = SWAP
    mount_point = "primary"
    size = _hp_min_swap | remaining | _hp_pri_swap
}
}
```

Now you have all the volume manager layouts. You may have noticed at the beginning of this section that LVM and VxVM both use the same ways of defining volume definitions; only the usage statement at the volume group level determines if an LVM volume group or VxVM disk group will be created (see below).

```
_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with HFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with VxFS"
{
```

Here you are defining the variable _hp_group_name. However, you must first make sure that it is not visible on the **Additional** button (since you do not want anyone to be able to change it). If the value of _hp_disk_layout is set to be VxVM ("VERITAS Volume Manager (VxVM) with VxFS") you set _hp_group_name to be rootdg; otherwise it is LVM you are using and vg00 is the name you are going to use²⁸.

²⁷ Impact statements and software are discussed later.

²⁸ VxVM 3.5 and earlier requires the root volume group be named rootdg. This is enforced by Ignite-UX.

The `_hp_group_name` does not actually do anything except hold the initial name of the volume group or disk group that this configuration is defining. If this variable could be changed using the **Additional** button, it would not do anything since it is only a temporary variable to hold the name²⁹.

```
_hp_group_name visible_if FALSE

(_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFS") {
    _hp_group_name = "rootdg"
} else {
    _hp_group_name = "vg00"
}
```

The temporary variable is necessary to give the volume group a name. You do this by assigning the name to a temporary variable and then starting your volume group definition with that name.

A volume group definition must have the following things in it:

- volume group attributes
- one or more `physical_volume` definitions (including "group"³⁰ definitions that also include more physical volume definitions)
- one or more (usually more) `logical_volume` definitions.

```
volume_group _hp_group_name
{
```

The `physical_volume` definition places the first disks (whose number is determined by `_hp_root_grp_disks`) returned by the `disk[]` construct into the root volume group³¹. When Ignite-UX encounters the first volume group that uses `*=<index>` it attempts to return the preset root disk as the first disk. This causes things to look as though the disk in `_hp_root_disk` (initialized from `_hp_primary_path`) was always placed into the list of physical volumes put into the root volume group and disk group.

```
physical_volume disk[*=_hp_root_grp_disks]
```

Next is a test to choose the volume manager. Currently the only choice that enables VxVM is "VERITAS Volume Manager (VxVM) with VxFS". All of the other choices use LVM.

```
(_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFS") {
    usage = VxVM
} else {
    usage = LVM
}
```

All the volume group attributes such as `max_physical_extents` are optional and only affect LVM volume groups. If you do not give them, they take on the defaults of the `vgcreate` command.

²⁹ To change the name of a volume group you use the Ignite-UX GUI to select the **File system** tab then the **Additional Tasks** button, and click **Group Parameters**. From the **Group Parameters** dialog box, select the name of the group to change from the selection list, change the group name, press modify, and then click **Ok**.

³⁰ Defining physical volume groups is discussed later when discussing custom configuration.

³¹ This is evaluated once. After that, changes made via the Ignite-UX GUI affect which physical volumes is used in the root volume group.

```

# In order to accommodate adding 9GB disks now or in the
# future, set the default max_physical_extents large
# enough to handle it (based on 4Mb extent size).
# Note that this will be increased by IUX automatically
# for disks larger than 9GB.
max_physical_extents = 2500

```

Very large disks are the most common disks on new systems with ≥ 140 GB disks considered large, 36/72GB disks average, and 18GB disks small. These tests adjust the physical extent size for the LVM volume group being defined and ensures that the complete disk can be used with the disks. It is only based upon the size of the root disk, not the size of every disk that might be included. This might influence the amount of usable space in the volume group. Additionally, Ignite-UX examines all the disks being used and selects the LVM parameters that work for that entire disk group. The setting of physical extent size here is just a first approximation; Ignite-UX can change these values at a later stage.

In a custom configuration, it might be better to scale both `max_physical_extents` and `physical_extent_size` based upon the disk size.

```

#
# For very large disks, the root group VGRA meta-data area
# will outgrow its bounds and trigger an IUX sanity check
# error. So for root disks large enough, set the physical
# extent size. Note that this only handles the root disk, but
# any disk in the root VG could trigger the error...
#
(disk[_hp_root_disk].size > 21504MB) # >21GB
{
    physical_extent_size=8
}
(disk[_hp_root_disk].size > 44032MB) # >43GB
{
    physical_extent_size=16
}
(disk[_hp_root_disk].size > 84992MB) # >83GB
{
    physical_extent_size=32
}
(disk[_hp_root_disk].size > 179200MB) # >175GB
{
    physical_extent_size=64
}

```

Next are the definitions of the logical volumes in the configuration. The first one is, of course, the root file system (note the `mount_point` definition). Depending on the value of the disk layout, the file system type is either HFS or VxFS.

Further, you have a test for a use model "Create separate volumes (`/usr`, , ...)". If there are separate file systems, then the default size for the root file system is 140MB if the root disk is less than 8192MB, or 200MB if it is greater than or equal to 8192MB.

If you are not expected to create separate logical volumes for file systems like `/usr` and `/var` then the sizes are the remaining space up to 1200MB if the size of the disk is less than 8192MB or the remaining space up to 1200MB plus the size of `_hp_pri_swap`³².

The root file system cannot have bad block relocation on and must be contiguous; therefore `contiguous_allocation` is set to true, and `bad_block_relocation` is set to false.

```
logical_volume {
    mount_point = "/"
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
    "Create separate volumes (/usr, /var, ...)" {
        disk[_hp_root_disk].size < 8192Mb {
            size = 140Mb
        } else {
            size = 200Mb
        }
    } else {
        disk[_hp_root_disk].size < 8192Mb {
            size = remaining | 1200Mb
        } else {
            size = remaining | 1200Mb + _hp_pri_swap
        }
    }
    contiguous_allocation = true
    bad_block_relocate = false
}
```

Next, you define primary swap by setting the usage to be swap and dump. The `mount_point` is set to "primary". The `vgdisk[0]` value enables you to specify that this logical volume be placed onto the first disk in the root volume group. Refer to *instl_adm(4)* for more information on the `vgdisk` keyword.

Next are the usual settings for primary swap of contiguous allocation and disabled bad block relocation. The size must be at least `_hp_min_swap` and up to `_hp_pri_swap` depending on the size of the disk.

```
logical_volume {
    usage = SWAP_DUMP
    mount_point = "primary"
    # Map it to the root disk to ensure it is considered primary
    # without this the secondary swap gets sorted first and considered
    # primary.
    vgdisk[0]
    contiguous_allocation = true
    bad_block_relocate = false
    size = _hp_min_swap | remaining | _hp_pri_swap
}
```

³² HP-UX B.11.11 may not actually fit into the amount of space defined here; the impacts statements for software defined may also force the sizes of file systems to be adjusted.

You looked at `_hp_sec_swap` earlier in the configuration file to see how it was defined. The only place you can give it a value is by using the **Additional** button on the **Basic** tab in the Ignite-UX GUI. You can manually define secondary swap using the **File system** tab, but you must manually select the disk to avoid it being on the same disk as primary swap. The configuration enables you to have secondary swap/dump. You only have to care about its size, not placement.

So if the variable `_hp_sec_swap` is `>0MB`, swap/dump space is created. If there is more than one disk in the root volume group, this swap/dump space is created on the second disk. If you provide equal sized primary and secondary swap, it creates an interleaved swap setup.

Also, you set the options required for this swap space to also be used as a dump space (contiguous allocation and bad block relocation off), and lastly you set the size to be `_hp_sec_swap`.

Secondary swap, if not used for dump, does not require contiguous allocation and bad_block relocation to be `_relocate` set to true.

```
_hp_sec_swap > 0Mb {
    logical_volume {
        usage = SWAP_DUMP
        mount_point = "secondary"
        (num_vgdisks > 1)
        {
            # This maps the secondary swap space to
            # the second disk when one is defined.
            vgdisk[1]
        }
        contiguous_allocation = true    # allows use as dump
        bad_block_relocate      = false # allows use as dump
        size = _hp_sec_swap
    }
}
```

Next `/stand` is defined. Of course, it must be HFS since the HP9000 secondary load cannot read a VxFS file system. Only HP Integrity servers (starting with HP-UX B.11.23) can have a VxFS file system in `/stand`. The block and fragment sizes are initialized from constants defined earlier. If you have a disk less than 2050MB, `/stand` defaults to 84MB. A disk greater than or equal to 2050MB but less than 3075MB is 112MB in size; otherwise you get a 300MB `/stand` file system by default.

```
logical_volume {
    mount_point = "/stand"
    # /stand must be HFS!
    usage = HFS
    blksize = _hp_HFS_blksize
    fragsize = _hp_HFS_fragsize
    # The /stand volume needs to be able to hold at least 5 kernels,
    # for various debugging reasons. The typical i1i kernel is
    # about 18Mb. The debug IA kernel is 42Mb. The smallest /stand
    # area needs to be at least twice the largest possible kernel.
    # Unless a user tries to put on more than the "normal" amount of
    # software on a system with a large root disk, the default /stand
    # directory will be 300Mb. To see the use of the size option,
    # see the man page for "instl_adm".
    disk[_hp_root_disk].size < 3075Mb {
        disk[_hp_root_disk].size < 2050Mb {
            # Most 2G drives are actually bigger than exactly 2G
            size = 84Mb
        }
    }
}
```

```

        } else {
            size = 112Mb
        }
    } else {
        size = 300Mb
    }
contiguous_allocation = true
bad_block_relocate    = false
}

```

300MB is the default for large disk configurations because additional Dynamically Loadable Kernel Modules (DLKMs) require more space in /stand.

Now at this stage you might be wondering why you had to make sure that you set up those logical volumes with certain values for contiguous allocation and bad block relocation. Table 3 contains a brief summary of the reasons.

Table 3

Name	Contiguous Allocation	Bad Block Relocation ³³	Reason
/	TRUE	FALSE	The root file system must be usable before the volume manager is active to enable this. It cannot allow the volume manager to relocate blocks and it must be contiguous on a disk.
Primary swap	TRUE	FALSE	Primary swap is activated before the volume manager is active, so it requires the same options as the root file system.
Dump space	TRUE	FALSE	A dump is written when no volume manager is active; because of this, it must be contiguous, and cannot have bad blocks relocated by the volume manager.
Secondary swap	N/A	N/A	Secondary swap has no requirements unless it is also a dump space; in which case, it has the dump space requirements.
/stand	TRUE	FALSE	The /stand volume is read by the boot loader. The boot loader knows nothing about volume managers and currently only understands HFS ³⁴ file systems, so the file system must be contiguous and cannot have bad block relocation enabled.

Note:

On HP-UX B.11.11, it is assumed that all IODC (I/O Dependent Code) in firmware can perform operations with up to 4GB offsets³⁵. For some I/O controllers this limits the maximum offset for dump. Newer systems have block-based I/O implemented in IODC. This allows the offset of an I/O operation to be given in blocks not bytes. This change dramatically increases the amount allowed for offset of dump spaces on most new I/O interfaces. Itanium®-based systems do not have a 4GB

³³ This is bad block relocation performed by a volume manager not automatic relocation performed by the disk drives.

³⁴ Itanium®-based systems may have /stand as a VxFS file system. The HFS restriction is applicable only to PA-RISC systems.

³⁵ Various combinations of 2GB and 4GB offsets (depending on the I/O interface) were allowed in older releases of Ignite-UX (before B.3.0 and in all of the A.x.x versions).

offset limit for dump since they all support block-based dump. In these cases the size is limited to the size of the disk itself.

If the option to create separate volumes is true (the default), then the following very long block of definitions is examined:

```
"Create separate volumes (/usr, /var, ...)" {
```

Next is a definition for `/usr`. The file system has large files enabled, and if the disk layout is "Logical Volume Manager (LVM) with HFS" then a HFS file system is created with the default block and fragment sizes set up earlier. Otherwise, it is a VxFS file system with the default VxFS block size (note that the fragment size is irrelevant³⁶ for VxFS).

```
logical_volume {
    mount_point = "/usr"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
}
```

Now if the size of the root disk is less than 1800MB, the size of the volume is at least 332MB, after all other volumes have had their sizes allocated. The size of `/usr` can be increased to have up to 20% free space³⁷ (if available). The impact of disk space requirements for software being installed that affects `/usr` may force the size to be higher.

```
disk[_hp_root_disk].size < 1800Mb {
    size = 332Mb | remaining | 20% free
} else {
```

If the root disk size is >3072 MB, the size of `/usr` is 500MB. Once all other volumes have had their space allocated, the size can be increased to ensure that there is 20% free space (noting the issues discussed previously can impact the size requirements). If the root disk size is between 1800MB and 3072MB, then the minimum size is 400MB.

```
disk[_hp_root_disk].size > 3072MB {
    size = 500Mb | remaining | 20% free
} else {
    size = 400Mb | remaining | 20% free
}
```

³⁶ VxFS is an extent-based file system and allocates available space to a file. Fragments can be as low as 1KB but are typically larger. There is no way to control fragment size in a VxFS file system.

³⁷ This amount does not include the amount of space that can be added by the `_hp_addnl_fs_free_pct` configuration item.

If the configuration item `_hp_root_grp_striped` is set to "YES" then you set up striping. The `"stripes = *"` set the number of stripes equal to the number of disks in the volume group. The stripe size `_hp_FS_stripe_size` was set up earlier as 64KB.

```

    _hp_root_grp_striped == "YES"  {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}

```

The definition of `/opt` is very similar to the definition of `/usr`. For this volume and the others defined by setting "Create separate volumes (`/usr`, `/var`, ...)" to true, you only look at the differences rather than the similarities.

The only difference between `/usr` and `/opt` is the size definition: rather than giving a percentage free, there is a specific size requirement for free space. For example, in the first size definition, when the size of the root disk is <1800MB, `/opt` is a minimum of 100MB. After all space has been allocated, you take from the remaining up to an extra 100MB and place it into `/opt`³⁸.

```

logical_volume {
    mount_point = "/opt"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
    disk[_hp_root_disk].size < 1800Mb {
        size = 100Mb | remaining | 100Mb free
    } else {
        disk[_hp_root_disk].size > 3072MB {
            size = 100Mb | remaining | 252Mb free
        } else {
            size = 100Mb | remaining | 152Mb free
        }
    }
    _hp_root_grp_striped == "YES"  {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}

```

If the size of the root disk is greater than or equal to 8192MB, the final size can be 500MB plus the size of `_hp_pri_swap`. The value is based on primary swap being dump and some space is required in `/var/adm/crash` to save crash dumps. This assumption is invalid if you define a separate volume for `/var/adm/crash` using the Ignite-UX GUI.³⁹

³⁸ Again, the impact of software being installed can increase the size of this volume, and the `_hp_addnl_fs_free_pct` configuration item can increase the final size of the volume (if free space is still left unallocated).

³⁹ The size requirement is up to 500MB + `_hp_pri_swap`. Depending on free disk space, there may not be enough to increase the size of a file system this large. During an interactive installation, the size of `/var` can be manually tuned down to a more appropriate value.

```

logical_volume {
    mount_point = "/var"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        # /var needs lots of inodes due to SD's IPD
        nbpi = 2048
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
    disk[_hp_root_disk].size < 8192Mb {
        size = 72Mb | remaining | 500Mb
    } else {
        size = 72Mb | remaining | 500Mb + _hp_pri_swap
    }
    _hp_root_grp_striped == "YES" {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}

```

The definition of `/tmp` is a simpler case than the other volumes. In this case, the size is a fixed value. It is unlikely that any impacts caused by installing software would affect `/tmp`.

```

logical_volume {
    mount_point = "/tmp"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        nbpi = 2048
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
    disk[_hp_root_disk].size < 8192Mb {
        size = 64Mb
    } else {
        size = 200Mb
    }
    _hp_root_grp_striped == "YES" {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}
} # "Create separate volumes (/usr, /var, ...)"

```

Although it is rare to see an `/export` volume created, you can enable it by using the **Additional** button on the **Basic** tab in the Ignite-UX GUI.

```

"Create /export volume" {
    logical_volume {
        mount_point = "/export"
        largefiles = true
        _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {

```

```

        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
    size = remaining | 100Mb
    _hp_root_grp_striped == "YES" {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}
}

```

The `/home` volume always gets defined and has a rather small size with no minimum, going up to 20MB.

```

logical_volume {
    mount_point = "/home"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFSS
        blksize = _hp_VxFSS_blksize
    }
    size = remaining | 20Mb
    _hp_root_grp_striped == "YES" {
        stripes = *
            stripe_size = _hp_FS_stripe_size
        }
    }
} # root group
} # Volume Manager disk layout.

```

The following special variables are used by Ignite-UX to control aspects of the installation or recovery. The "Special variables" section describes all of the special variables in detail.

```
# Set the variables that are recognized by the UI as invisible  
# so that they don't show on the "Additional" screen:  
_hp_locale           visible_if false  
_hp_cfg_detail_level visible_if false  
_hp_pri_swap         visible_if false  
_hp_min_swap         visible_if false  
_hp_disk_layout      visible_if false  
_hp_default_cur_lan_dev  visible_if false  
_hp_default_final_lan_dev visible_if false  
_hp_keyboard          visible_if false  
_hp_root_disk         visible_if false  
_hp_boot_dev_path     visible_if false
```

If `_hp_disk_layout` is one of the volume manager values then you allow some things to become visible in the **Additional** button from the **Basic** tab in the Ignite-UX GUI. Some of these control some of the configuration that you have already looked at. They are only applicable if you are using a volume manager, so the `else` construct makes sure that none of them can be seen on the **Additional** button.

```

_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with HFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with VxFS"
{
    "Create /export volume"                      visible_if true
    "Create separate volumes (/usr, /var, ...)"  visible_if true
    _hp_sec_swap      visible_if true
    _hp_root_grp_striped visible_if (_hp_root_grp_disks > 1)
    _hp_root_grp_disks  visible_if true
}
else
{
    "Create /export volume"                      visible_if false
    "Create separate volumes (/usr, /var, ...)"  visible_if false
    _hp_sec_swap      visible_if false
    _hp_root_grp_striped visible_if false
    _hp_root_grp_disks  visible_if false
}

```

The following `sw_sel` definition adds some impacts and enables the addition of the impacts to the file systems to take into account file system usage caused by the installation. When the `visible_if` variable is set to `false` the `sw_sel` definition does not appear in the Ignite-UX GUI; this variable is preset to `true` so the impacts always affect the installation.

```

#
# In order to more accurately specify the software impact to
# specific volumes that normally do not have files loaded
# directly, but get created by scripts, we specify some
# typical sizes here so the disk space is accounted for.
#
init sw_sel "impact by scripts" {
    sw_source = "core"
    visible_if = false
    impacts = "/var"   60Mb
    impacts = "/stand" 20Mb
} = TRUE

```

The following configuration statements add a selection to the **Additional** button that enables the user to control the use of autoboot. Ignite-UX normally runs no matter what the `autoboot` flag is set to (refer to `setboot(1M)`). It forces it on if it was off, then later turns it back off for the final reboot when the system would normally come back. This configuration enables the user to manually boot the system when `autoboot` is disabled.

Ignite-UX normally forces the system to perform the first reboot automatically when rebooting from the installation kernel to the newly created or recovered kernel. This reboot is forced no matter what the `autoboot` flag is set to (refer to `setboot(1M)`). Ignite-UX forces it on if it was off; prior to the final reboot Ignite-UX returns the `autoboot` flag to off. If `_hp_force_autoboot` is set to NO and the system's `autoboot` flag is not set to on, the system does not perform a reboot.

Important:

If you disable `autoboot` and then install a system, the system never executes a final reboot without manual intervention. HP does not

support booting into any other mode than the default multi-user run level at this point. The installation or recovery session has not yet completed and some steps must still be performed by Ignite-UX. Booting into anything other than the default run level does not allow the installation or recovery to complete. This is the intended behavior of Ignite-UX and not a flaw. In general (except when investigating Ignite-UX issues on advice from HP Support), you should never change the value of the _hp_force_autoboot variable.

```
#  
# Variable to give user control over the use of setboot. When the variable  
# is YES, autoboot is set before the first reboot and restored after the  
# first reboot. When the variable is NO, autoboot is never changed.  
#  
enum _hp_force_autoboot  
_hp_force_autoboot = { "YES", "NO" }  
init _hp_force_autoboot = "YES"  
_hp_force_autoboot help_text "Force Ignite-UX autoboot?"
```

The following example enables you to set, using the **Additional** button, whether the final system uses DHCP to obtain an IP address and hostname once it is installed or recovered. Ignite-UX uses DHCP to get initial networking information (you must disable DHCP [disable_dhcp=true] in the installation file system to stop a system booting over the network from using DHCP to obtain an IP address and hostname).

```
#  
# Variable to give user control over DHCP.  
# If YES, then DHCP is turned off on the system.  
#  
enum _disable_dhcp  
_disable_dhcp = { "YES", "NO" }  
init _disable_dhcp = "NO"  
_disable_dhcp help_text "Disable DHCP?"  
_disable_dhcp == "YES" {  
    DISABLE_DHCP=TRUE  
}
```

In the next example, you can configure the value of _hp_addnl_fs_free_pct. The variable is defined as a list of values that can be selected from the **Additional** button. If an enum had been placed before the _hp_addnl_fs_free_pct you would only be able to select a value from 0 to 10. Instead, you can select a value from 0 to 10 or enter another value when changing the variable using the **Additional** button.

```
#  
# JAGae82704 fix. Variable to give user control over the amount of  
# free space left over in a logical volume. It defaults to 10. The  
# range provided if you click on the button is from 0-10, but a larger  
# value can be entered directly in the field if desired.  
#  
_hp_addnl_fs_free_pct = { 0,1,2,3,4,5,6,7,8,9,10 }  
init _hp_addnl_fs_free_pct = 10  
_hp_addnl_fs_free_pct help_text "Additional free space %"
```

The last part of configuration accomplishes several things:

1. Sets the os_release attribute for the swinstall command line.

2. Makes sure that `_hp_os_bitness` is set and set correctly (it should be set by the configuration file that defines the core operating system depot or archive `sw_sel` clause).
3. Makes sure that the system can run the bitness of HP-UX.
4. Sets the operating system Name attribute for the `swinstall` command line.

```
#####
# Ensure that the swinstall command line used contains the correct pose_as
# attributes for the type of system being installed.
#
sd_command_line += " -x os_release=" +40 ${release}

(_hp_os_bitness == "64")
{
    (!can_run_64bit)
    {
        ERROR += "This system model: \\" +${model}+ "\"" is not supported for
running 64bit HP-UX, you must select the 32bit selection"
    }
    sd_command_line += " -x os_name=HP-UX:64 "
}
else
{
    (_hp_os_bitness == "32")
    {
        (!can_run_32bit)
        {
            ERROR += "This system model: \\" +${model}+ "\"" is not supported for
running 32bit HP-UX, you must select the 64bit selection"
        }
        sd_command_line += " -x os_name=HP-UX:32 "
    }
    else
    {
        ERROR += "The _hp_os_bitness variable is not set to \"32\" or \"64\".
This variable must be set in the config file that supplies the core archive or
depot. If using an archive, make sure you start with the core11.cfg example
config file. When using a depot, make_config will set this automatically."
    }
}
_hp_os_bitness visible_if false
```

You should exercise care when setting the value of `sd_command_line` in any other configuration file because the `os_name` and `os_release` variables are required to install software correctly during an initial installation. Using `=` rather than `+=` to add an option to `sd_command_line` removes any current settings from `sd_command_line`.

Analyzing the HP-UX default B.11.31 cfg clause

The following is from a default `/var/opt/ignite/INDEX` showing the configuration files referenced by the default HP-UX B.11.31 cfg clause:

⁴⁰ Concatenates two strings together using the `+` operator. For more information, refer to `instl_adm(4)`.

```

cfg "HP-UX B.11.31 Default" {
    description "This selection supplies the default system configuration
that HP supplies for the B.11.31 release."
    "/opt/ignite/data/Rel_B.11.31/config"
    "/opt/ignite/data/Rel_B.11.31/hw_patches_cfg"
    "/var/opt/ignite/config.local"
}

```

Of course, by itself, this can't be used to install a system as it doesn't have a configuration file referencing any software to be installed.

This section describes the per release configuration file `/opt/ignite/data/Rel_B.11.31/config` and explains what is happening and why certain things have been done in certain ways.

Please note the following text from the Ignite-UX C.7.0 release notes:

- Default file system sizes, primary swap size and other default system configuration settings have been changed in HP-UX 11i v3 as compared to prior HP-UX releases. This was important to better adjust to Operating Environment software and computer system hardware. Further tuning of these configuration defaults may be done in a future Ignite-UX release. As before, it may be necessary or appropriate to create custom configurations specific for your systems. Such custom configurations or recovery configurations will not be affected by future default configuration tuning.

The configuration presented in this section may not match that provided with Ignite-UX revisions greater than C.7.0, please be careful if you rely on the information provided here and ensure it is consistent with the release of Ignite-UX you are using.

The release-specific configuration file

The `config` file defines all of the file system and volume management configurations for the release (and a few other things). This description is based on the release-specific config file from Ignite-UX version C.7.0.

```

#####
# config:
#
# This file contains the default system definitions used during a
# cold-install of HP-UX. This file should not be modified. If
# modifications are needed, the file "config.local" should be modified
# instead.
# See the instl_adm(4) man page for details.
#
# Changes to this file will not be preserved if the fileset is re-installed.
#
# If you are editing this file via the advanced options during a
# cold-install, then notice that the "config.local" contents are appended
# to the end of this file.
#
# @(#) config $Revision: 10.11 $

# IGNITE_UX_COPYRIGHT 1996 2006 #
# "(C) Copyright 1996-2006 Hewlett-Packard Development Company, L.P."#
# IGNITE_UX_COPYRIGHT #

```

The `release` keyword is probably the single-most important keyword in an Ignite-UX configuration file. This keyword *must* be set or else any installation of HP-UX is likely to fail or not work properly. Ignite-UX has no idea what version of HP-UX it is installing unless it is set.

As of Ignite-UX C.6.0.x, the `release` keyword is checked against the value of the `_hp_ikernel_os_release` variable (the client creates the value of this variable in the `host.info` file). If the HP-UX release in the `release` keyword is not supported by the kernel version given in `_hp_ikernel_os_release`, the installation is not allowed to proceed and the following error appears:

```
ERROR: The version of HP-UX you have chosen to install on the system  
       (B.11.31) is not supported by the version of the Ignite-UX install  
       kernel that the system booted (B.11.23). You will need to reboot the  
       target system from an install kernel matching the desired release  
       from the menu at the console. If using the bootsys command, use the  
       '-R' option to specify the install kernel version.
```

Therefore, if you boot the B.11.31 installation kernel, you can't install HP-UX B.11.11 or B.11.23, you can only install HP-UX B.11.31.

When you create custom configuration files and do not plan to include the default configuration files you *must* set this variable.

```
release="B.11.31"
```

In this section we have some HP-UX 11i v3 restrictions that Ignite-UX enforces via configuration.

HP-UX 11i v3 does not support any HP9000 workstations, so the first test checks to see if the system being installed is a PA-RISC system, and then checks to see if the hardware model starts with 9000/7 followed by one or more of any character. If this is true, an error will eventually be shown to the user saying "HP-UX B.11.31 does not support 9000/7xx workstation systems".

The second thing that is enforced relates to memory. HP-UX 11i v3 requires a minimum of 1GB of memory in order to be installed. If the system contains less than 950MB of memory, you will eventually see an error saying "HP-UX B.11.31 requires at least 1GB of system memory" and you will not be allowed to continue the install.

```
#####
# Release supported configuration checking for unsupported configs
# that are easy to recognize.
#
is_hppa {
    HARDWARE_MODEL ~ "9000/7.*" {
        error += "HP-UX " + ${release} + " does not support 9000/7xx
workstation systems."
    }
}

( memory < 950MB ) {
    error += "HP-UX " + ${release} + " requires at least 1GB of system memory."
}
```

Sometime ago, changes were made to Ignite-UX to increase the block and fragment sizes for VxFS and HFS file systems. This can lead to differences between systems installed in the last few years

and much older systems that had smaller default block and fragment sizes. (For example, the file system attributes are different so the block and fragment sizes from older systems are probably smaller.)

The use of <variable> `visible_if false` ensures that the variables do not appear in the **Additional** dialog box from the **Basic** tab.

```
#####
init _hp_HFS_blksize = 65536      # HFS block size.
init _hp_HFS_fragsize = 8192       # HFS frag size.
init _hp_VxFS_blksize = 8192       # VxFs block size.
init _hp_FS_stripe_size = 64Kb    # File system stripe size.

_hp_HFS_blksize visible_if false
_hp_HFS_fragsize visible_if false
_hp_VxFS_blksize visible_if false
_hp_FS_stripe_size visible_if false
#####
```

This section defining the EFI and HPSP partitions for HP Integrity servers is the same as it was for HP-UX B.11.23.

The test around the section ensures that partition sizes are only defined for HP Integrity systems. The variable `_hp_efi_partition_size` is given a set of values it can have. Note that it is not defined as an enum, so although there is a list of values to choose from, you can still set the size of the EFI partition directly if you want a very specific size. The `_hp_efi_partition` variable is then given an initial (default) size, and then given some help text to make it easier to understand what is does on via the additional screen.

Note that on the additional screen sizes appear in KB not MB, so although the size is defined in MB you will see the size in KB when in itool.

For the service partition size we have some configuration that looks almost identical to the EFI partition size we've already discussed.

```
#####
# Set the boot disk's EFI boot partition and HP Service Partition (HPSP)
# sizes. The EFI boot partition may actually hold the HP-UX kernel at
# some point in the future. The HPSP hold off-line diagnostics and dump
# information related to machine state (clearly not a full memory dump).
# A zero size (no HPSP) should probably be kept as an option.
#
is_ia64 {
    _hp_efi_partition_size = {
        500Mb,
        1000Mb,
        2000Mb,
        4000Mb
    }
    init _hp_efi_partition_size = 500Mb
    _hp_efi_partition_size help_text "EFI Boot Partition (KB)"
    _hp_service_partition_size = {
        0Mb,
        400Mb,
        1000Mb,
        2000Mb,
        4000Mb
    }
}
```

```

    init _hp_service_partition_size = 400Mb
    _hp_service_partition_size help_text "HP Service Partition (KB)"
} else {
    _hp_efi_partition_size visible_if false
    _hp_service_partition_size visible_if false
}

```

The following shows the start of a definition for the `_hp_disk_layout` variable, which is set to one of several possible values depending on if this is a HP9000 or HP Integrity system.

Around the `_hp_disk_layout` is a test for `is_ia64` and `is_hppa` (is HP Integrity or PA-RISC/HP9000 system). Disk layouts specific to those different types of systems are defined. The reason for converting some of the string for backwards compatibility is that older custom configuration files may refer to the names previously used for a disk layout. This would break those older configuration files, so the new names that appear in itool for the disk layouts are mapped to the older names for compatibility with older configuration files.

You actually get to see the value of the variable `_hp_disk_layout` when you use the Ignite-UX GUI. The values currently assigned to `_hp_disk_layout` are shown on the **File System**: list on the **Basic** tab. The form used in the next example to initialize a list of possible values allows the Ignite-UX GUI to modify the list later.

Further in the configuration file, when giving `_hp_disk_layout` a final value, you only assign its value using `init`. When using the form "`init _hp_disk_layout="..."`", the Ignite-UX GUI is still allowed to modify the value. If you instead define "`_hp_disk_layout="..."`" without the `init`, the Ignite-UX GUI cannot change this value. If the Ignite-UX GUI were not able to change the value of `_hp_disk_layout`, no changes made in the Ignite-UX GUI to file system attributes are applied to the system⁴¹.

This is because the Ignite-UX GUI must change the value of `_hp_disk_layout` to "Modified LVM Layout", "Modified VxVM Layout", or "Modified Whole-Disk Layout" when changes have been made on the **File System** tab in the Ignite-UX GUI. If the Ignite-UX GUI is not allowed to change the variable, you cannot keep any file system customizations. This concept becomes very important when creating custom configuration files.

```

#####
# Any _hp_disk_layout variable setting besides those below will disable
# the default disk layout described in this file. This can be
# done if a custom disk layout is defined elsewhere.
#
is_ia64 {
    _hp_disk_layout = {
        "VERITAS Volume Manager (VxVM) with VxFs",
        "Whole disk with VxFs",
        "Logical Volume Manager (LVM) with VxFs"
    }
}

# For PA architecture only
is_hppa {
    _hp_disk_layout = {
        "Whole disk (not LVM) with HFS",

```

⁴¹The itool command Ignite-UX GUI accepts changes to file system attributes and does not caution you, but no file system attribute changes made in the Ignite-UX GUI are applied to the system.

```

        "Logical Volume Manager (LVM) with HFS",
        "Logical Volume Manager (LVM) with VxFS",
        "VERITAS Volume Manager (VxVM) with VxFS"
    }
}

# convert strings for backwards compatibility
({_hp_disk_layout == "Whole disk (not LVM) with VxFS")
{
    _hp_disk_layout = "Whole disk with VxFS"
}

({_hp_disk_layout == "VxVM with VxFS")
{
    _hp_disk_layout = "VERITAS Volume Manager (VxVM) with VxFS"
}

```

Here the variable `_hp_disk_root` is set to the value of `_hp_primary_path`, assuming that the device `_hp_primary_path` points to exists and `_hp_root_disk` does not already have a value. An Extended Regular Expression (ERE)⁴² is used to see if `_hp_root_disk` is set to any value⁴³.

```

#####
# Set the default root disk to the current primary path.
# The _hp_primary_path will be "" if it was set to a
# non-existent dev. Do not change the root disk variable
# if it's already set. The '~' operator is used instead
# of '==' to deal with the variable not yet being defined.
#
({_hp_primary_path != "" & !(_hp_root_disk ~ ".*"))
{
    init _hp_root_disk = _hp_primary_path
}

```

The definition of swap has changed significantly in this release-specific configuration file in order to deal better with two issues – large memory configurations and large disk luns. Systems containing very large memory configurations and very large disk luns (or JBODs) should be commonplace with HP-UX B.11.31, so the definition of primary swap takes both of these factors into consideration.

```

#####
# Default (recommended) swap size for very small memory systems is
# two times memory. Swap size needs to be reduced if available disk
# space is limited. As system memory space increases it may or may
# not be necessary to increase swap size. Some systems may run using
# only system memory and little additional swap space. For some system
# application loads additional swap space may be needed. By default
# little additional swap space is added for large memory systems.
#
# Note that the swap size should be a multiple of 64MB to avoid a UI
# adjustment warning.

```

⁴² For more information regarding EREs, refer to `regexp(5)`.

⁴³ The ERE ".*" means zero or more of any character so it matches a zero length string.

Although we cover effects relationships in other places, it is worthwhile revisiting it again. Normally, once a value is worked out for a variable it is not reevaluated again. Effects relationships set relationships between values, so when one changes the other is forced to be reevaluated. The effects relationships below cause `_hp_pri_swap` to be reevaluated whenever the size of the root disk changes, or when the disk layout is changed.

```
# The next two statements are used only to establish an effects
# relationship between the _hp_pri_swap variable and _hp_root_disk &
# _hp_disk_layout variables. The value is overwritten below.
# This makes the UI change the swap anytime either of these two
# values change.
init _hp_pri_swap = disk[_hp_root_disk].size
init _hp_pri_swap = _hp_disk_layout
```

Now we have a sanity check based on the size of the root disk. If it is less than 9GB, and primary swap is more than 1GB, and it is the only disk in the root volume or disk group, an error will be given that will force you to change the size of primary swap to be 1GB or less. Note that Ignite-UX won't set primary swap to more than 1GB under these circumstances - this sanity check is here to prevent it from being set manually.

```
# Recognize swap size being set too large for the root volume config.
( disk[_hp_root_disk].size < 9000MB & _hp_pri_swap > 1GB &
_hp_root_grp_disks == 1 ) {
    error+="Primary swap size is too large for root disk. Reduce primary swap
size to 1GB, add disks to root volume / disk group, or use a larger disk."
}
```

Here we set the size of primary swap. Primary swap is sized this way because on lower memory configurations you are more likely to allocate or use device swap, so more swap is required compared to the size of memory. Ignite limits the size of primary swap to 8GB.

```
# Start with 2x or cap on swap based on memory size
( memory < 4GB ) {
    init _hp_pri_swap = (((memory * 2) + 65535KB)/64MB)*64MB
}
( memory >= 4GB ) {
    init _hp_pri_swap = 8GB
}
```

In this section we adjust the size of primary swap for a small root disk. If you only have one disk in the root volume or disk group, and it is less than 9GB in size, primary swap will be set to 1GB. The second part of the configuration checks if the size of the root disk is between 9G and 18GB. Then, if the size of memory is greater than or equal to 4GB, primary swap will be set to 4GB. Otherwise it will be set to 2GB.

```
# Adjust swap size to deal with small disks.
( disk[_hp_root_disk].size < 9000MB & _hp_root_grp_disks == 1 ) {
    init _hp_pri_swap = 1GB
}
( disk[_hp_root_disk].size >= 9000MB & disk[_hp_root_disk].size < 18000MB ) {
    ( memory >= 4GB ) {
        init _hp_pri_swap = 4GB
    } else {
        init _hp_pri_swap = 2GB
    }
}
```

```
    }
}
```

Here we set up an effects relationship between `_hp_min_swap` and `_hp_pri_swap`. That is followed by some final adjustments to the size of primary swap based on the size of the root disk and the amount of memory in the system. Depending on the size of the root disk and the amount of memory in the system, Ignite-UX will attempt to place a sensible lower bound (`_hp_min_swap`) on the size of primary swap.

```
# The next statement is used only to establish an effects relationship
# between the _hp_pri_swap variable and _hp_min_swap variable. The
# value is overwritten below. This ensure that changes will be handled
# correctly. Depending how _hp_min_swap is calculated it may not change
# due to UI changes.
init _hp_min_swap = _hp_pri_swap

# Adjust swap size to deal with a lot of system memory
( memory < 16GB) {
    _hp_min_swap = 1GB
}
( memory >= 16GB & memory < 64GB) {
    _hp_min_swap = 2GB
}
( memory >= 64GB) {
    _hp_min_swap = 4GB
}
```

Now `_hp_sec_swap` and `_hp_pri_swap` are initialized with a set of values that can be chosen from; however, they are not a fixed set of values. The variable `_hp_sec_swap` can be set using the **Additional** button because of the way the variable was defined (`<variable>={ value1, value2, ... }`). The variable `_hp_sec_swap` has its value set from a selection list or by typing it into the corresponding field directly. This can happen because `_hp_sec_swap` was not defined as an "enum"; if it had been defined as an enum only a value from the list of values it contains could be used.

Although, in this configuration, `hp_pri_swap` cannot be modified using the **Additional** button. This value appears directly in the Ignite-UX GUI on the **Basic** tab using the **Root Swap (MB)** button, and it allows you to select from a list of possible values. The list of values shown in the selection list for `_hp_pri_swap` comes from the list defined in the configuration. If you create a custom disk configuration file, you should supply a value or a list of values for `_hp_pri_swap` (or somehow otherwise calculate a default value for it).

Note that large values can be used for primary and secondary swap. From Ignite-UX version C.7.0 onwards, HP-UX 11i v2 and later releases of HP-UX are no longer limited to 32GB swap spaces.

```
# Note that 0 (zero) is not a valid size for primary swap
_hp_pri_swap = {      1GB,     2GB,     4GB,
                     8GB,    16GB,    24GB,    32GB,
                    48GB,    64GB,   98GB, 128GB,
                   192GB, 256GB }

_hp_sec_swap = {     0GB,     1GB,     2GB,     4GB,
                     8GB,    16GB,    24GB,    32GB,
                    48GB,    64GB,   98GB, 128GB,
                   192GB, 256GB }

init _hp_sec_swap = 0Mb
```

```
_hp_sec_swap help_text "Secondary Swap space (KB)"
```

From here we are going to be looking at the disk layouts. Some parts of the disk layout have changed significantly for HP-UX 11i v3.

```
#####
# Determine disk layout for volumes and file systems.
#
```

You should have noticed by now that for HP-UX 11i v3, all of the tests related to legacy HP9000 systems have been removed. None of these systems can run HP-UX 11i v3 (they couldn't run HP-UX 11i v2 either) so all of the tests are gone.

The next example gives `_hp_disk_layout` a value of disk `[_hp_root_disk].model`. If the value of that expression changes, it causes the value of `_hp_disk_layout` to be re-evaluated.

Doing this makes good sense in the context of this configuration file, because if you change the root disk, you may need to re-evaluate what the disk configuration is. After that, we set the default disk layout to be LVM with VxFS.

```
# This next statement is used only to establish an effects relationship
# between the _hp_disk_layout variable and _hp_root_disk. The value is
# overwritten below.
init _hp_disk_layout = disk[_hp_root_disk].model

# Default disk layout
init _hp_disk_layout = "Logical Volume Manager (LVM) with VxFS"
```

The enumeration `_hp_root_grp_striped` in the following example is a perfect example of how to ask a yes/no or Boolean question using the **Additional** button on the **Basic** tab in the Ignite-UX GUI. The first step is to define the variable as an enum, then give it two values, an initial value (with the `init` keyword so the user is allowed to change it using the Ignite-UX GUI), and some help text.

Without defining `_hp_root_grp_striped` as an enum, the user would be able to select **YES** or **NO** on the **Additional** button, or also enter any value. In this case, it does not matter because the configuration only tests on the value **YES** (later) to decide on an action. If you could enter a lower-case "yes", the tests would fail; hence the enumeration type is used.

```
# Boolean describing whether or not the root disk group
# will be striped
#
enum _hp_root_grp_striped
_hp_root_grp_striped = { "YES", "NO" }
init _hp_root_grp_striped = "NO"
_hp_root_grp_striped help_text "Stripe root VG disks?"
```

The expression in the next example defines a use model. Use models are not variables; they are true / false expressions that are used to control how things are done⁴⁴. Like variables, use models can be viewed using the **Additional** button of the **Basic** tab in the Ignite-UX GUI. In the

⁴⁴ Unfortunately, variables cannot have effects relationships with use models or software selections.

Additional dialog box, use models look visibly different from variables. Later in the configuration file is the impact statement this use model would have when it is set to true.

```
INIT "Create /export volume" = false
```

Now you have another effects relationship, this time between a new variable `_hp_root_grp_disks` and `_hp_root_disk`. The variable `_hp_root_grp_disks` is meant to be an integer, so when doing the effects relationship you add zero (0) to `disk[_hp_root_disk].size` to ensure that the final value is an integer. Since variables are not explicitly typed, you want to imply an integer type.

The variable `_hp_root_grp_disks` is also an enum; it can take one of the values from 1 up to the value of variable `num_disks`,⁴⁵ and it has a default initial value of 1.

That is followed by the usage model "Create separate volumes (/usr, /var, ...)". This is used to decide if there will be only one large root file system or if separate file systems will be created (it would be unusual to set this to **FALSE**).

```
# Number of disks in the root volume group
# The fist line below is just to define an "effects" relationship
# between the root disk and _hp_root_grp_disks which dependent
# upon each other in the logic to follow.
#
init _hp_root_grp_disks = (disk[_hp_root_disk].size+0) # +0 convert to int
enum _hp_root_grp_disks
_hp_root_grp_disks = { 1..num_disks }
init _hp_root_grp_disks = 1
_hp_root_grp_disks help_text "# of disks in root VG"

INIT "Create separate volumes (/usr, /var, ...)" = TRUE
```

The first disk configuration "Whole disk (not LVM) with HFS" is protected by a test of the variable `_hp_disk_config` to see if it is equal to "Whole disk (not LVM) with HFS".

After that, the `partitioned_disk` keyword starts the definition of a whole disk layout for a system. The `partitioned_disk` disk definition must contain the following:

- A `physical_volume` definition
- An `fs_partition` definition
- A `swap_partition` definition

If you have no `fs_partition` definition that at least defines the root file system, the following error appears after pressing **Go!** in the Ignite-UX GUI:

```
ERROR: There is no root volume (mount point = /) defined in the configuration.
```

If you have no `swap_partition` defined, you see the following error after selecting **Go!** in the Ignite-UX GUI:

```
ERROR: There is no swap volume defined in the root volume group (or on the
root disk).
```

⁴⁵ A built-in variable that is a count of the number of disks discovered on the system. Refer, refer to `inst1_adm(4)`.

For the physical volume statement, you need only supply the name of the disk used as the root disk. (The disk[] keyword takes a hardware path or index number, and returns the name of the disk.) This discussion assumes you are defining the boot disk as a whole disk⁴⁶.

```
_hp_disk_layout == "Whole disk with VxFS" |
_hp_disk_layout == "Whole disk (not LVM) with HFS"
{
    partitioned_disk
{
```

The file system partition on a partitioned disk must be HFS for HP9000 systems. It follows from the fact that the PA-RISC boot loader can only read from some types of HFS file systems (large files-enabled HFS file systems cannot be booted from on PA-RISC systems). Therefore, the usage statement is set to HFS. You set the block size and the fragment size⁴⁷ to the HFS defaults defined earlier.

Size is set to "remaining", which takes whatever space is left over once "other" things (in this case swap at the end of the disk) have been allocated. The mount point naturally is "/" (since it is the root file system). The size of the file system is all the remaining space on the disk (after swap has been allocated).

When the disk layouts were defined earlier, "Whole disk with VxFS" was only available for HP Integrity servers.

```
physical_volume disk[_hp_root_disk]
fs_partition {
    _hp_disk_layout == "Whole disk with VxFS"
    {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    } else {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    }
    size = remaining
    mount_point = "/"
}
```

Now you have your primary swap partition definition - you can see that it is swap from the usage. The interesting thing is the size specification. This sets the size to be at least `_hp_min_swap`, but also includes any remaining space up to `_hp_pri_swap` in size. The swap space in this case attempts to get all the space it can (up to `_hp_pri_swap`), but does not go lower than `_hp_min_swap`, while ensuring the file system partition gets enough to meet its impacts⁴⁸ statements requirements.

⁴⁶ You can have multiple whole disks defined in a configuration. Each whole disk is defined by different `partitioned_disk` definitions and has a different `mount_point`.

⁴⁷ These concepts are only meaningful in a HFS file system. The block size is the largest block of data a file can have allocated to it (files that can occupy a full free block will do so and the smallest amount of space that can be allocated by a file is the fragment size. In extent-based file systems these concepts are usually meaningless. For example, the smallest allocation unit is 1KB no matter what the fragment size is set to for the file system.

⁴⁸ Impact statements and software are discussed later.

```

        swap_partition {
            usage = SWAP
            mount_point = "primary"
            size = _hp_min_swap | remaining | _hp_pri_swap
        }
    }
}

```

Here are the volume manager disk layouts. The definition of logical volumes and volumes are performed identically with Ignite-UX. The usage definition (LVM or VxVM), not the volume group definition, determines if LVM or VxVM will be used.

```

_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with HFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with VxFS"
{

```

Here you are defining the variable `_hp_group_name`. However, you must first make sure it is not visible on the **Additional** button (since you do not want anyone to be able to change it). If the value of `_hp_disk_layout` is set to be VxVM ("VERITAS Volume Manager (VxVM) with VxFS"), set `_hp_group_name` to be `rootdg`; otherwise it is LVM you are using, and `vg00` is the name you are going to use⁴⁹.

The `_hp_group_name` does not actually do anything except hold the initial name of the volume group or disk group this configuration is defining. If this variable could be changed using the **Additional** button, it would not do anything since it is only a temporary variable to hold the name⁵⁰.

```

_hp_group_name visible_if FALSE

(_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFS") {
    _hp_group_name = "rootdg"
} else {
    _hp_group_name = "vg00"
}

```

Here is the definition of a root volume or disk group. The construct in the `disk[]` keyword `*=_hp_root_grp_disks` means the first `_hp_root_grp_disks` disks. Since the value of `_hp_root_disk` is always `disk[1]`, we know `_hp_root_disk` will be one of the disks defined in this volume group, and there may be `_hp_root_grp_disks-1` additional disks.

Depending if VxVM was chosen as the volume manager, we then choose either LVM or VxVM as the volume manager via the `usage` keyword. That is followed by some housekeeping configuration to maintain good values for the number of extents in a volume group and the extent size (only applies to LVM).

The temporary variable is necessary to give the volume group a name. You do this by assigning the name to a temporary variable, and then starting your volume group definition with that name.

⁴⁹ VxVM 3.5 and earlier requires the root volume group be named `rootdg`. This is enforced by Ignite-UX.

⁵⁰ To change the name of a volume group you use the Ignite-UX GUI to select the **File system** tab then the **Additional Tasks** button, and click **Group Parameters**. From the **Group Parameters** dialog box, select the name of the group to change from the selection list, change the group name, press modify, and then click **Ok**.

A volume group definition must have the following:

- volume group attributes
- one or more physical_volume definitions (including "group"⁵¹ definitions that also include more physical volume definitions)
- one or more (usually more) logical_volume definitions.

```
volume_group _hp_group_name
{
```

The physical_volume definition places the first disks (whose number is determined by _hp_root_grp_disks) returned by the disk[] construct into the root volume group⁵². When Ignite-UX encounters the first volume group that uses *=<index>, it attempts to return the preset root disk as the first disk. This causes things to look as though the disk in _hp_root_disk (initialized from _hp_primary_path) was always placed into the list of physical volumes put into the root volume group and disk group.

```
physical_volume disk[*=_hp_root_grp_disks]
```

Next is a test to choose the volume manager. Currently the only choice that enables VxVM is "VERITAS Volume Manager (VxVM) with VxFs". All of the other choices use LVM.

```
(_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFs") {
    usage = VxVM
} else {
    usage = LVM
}
```

All the volume group attributes, such as max_physical_extents, are optional and only affect LVM volume groups. If you do not assign them to a value, they take on the defaults of the vgcreate command.

```
# In order to accommodate adding 9GB disks now or in the
# future, set the default max_physical_extents large
# enough to handle it (based on 4Mb extent size).
# Note that this will be increased by IUX automatically
# for disks larger than 9GB.
max_physical_extents = 2500
```

Very large disks are the most common disks on new systems, with \geq 140GB disks becoming normal, 36/72GB disks less common, and 18GB or smaller disks rare. The following tests adjust the physical extent size for the LVM volume group being defined, and ensures the complete disk can be used with the disks. It is only based upon the size of the root disk, not the size of every disk that might be included. This might influence the amount of usable space in the volume group. Additionally, Ignite-UX examines all the disks being used and selects the LVM parameters that work for that entire disk group. The setting of physical extent size here is just a first approximation; Ignite-UX can change these values at a later stage.

⁵¹ Defining physical volume groups is discussed later when discussing custom configuration.

⁵² This is evaluated once. After that, changes made via the Ignite-UX GUI affect which physical volumes is used in the root volume group.

In a custom configuration, it might be better to scale both `max_physical_extents` and `physical_extent_size` based upon the disk size.

```

#
# As disk size increases the root group VGRA meta-data area
# will outgrow it's bounds and trigger an IUX sanity check
# error. So for root disks large enough, set the physical
# extent size. Note that this only handles the root disk, but
# any disk in the root VG could trigger the error...
# For consistency with previous releases these disk sizes
# are specified in MB.
#
#
(disk[_hp_root_disk].size > 21504MB) # (about 23GBDISK)
{
    physical_extent_size=8
}
(disk[_hp_root_disk].size > 44032MB) # (about 46GBDISK)
{
    physical_extent_size=16
}
(disk[_hp_root_disk].size > 84992MB) # (about 89GBDISK)
{
    physical_extent_size=32
}
(disk[_hp_root_disk].size > 179200MB) # (about 188GBDISK)
{
    physical_extent_size=64
}

```

We start with our first definition of a logical volume. (We don't define names for the logical volumes here, so they take the default names applicable to the relevant volume manager.) The logical volume has a mount point of "/" (so it is the root file system), and only one disk layout is defined as HFS (other disk layouts are VxFS file systems). Because this is the root file system, we must set contiguous allocation to true and bad block relocation to false. Note the size change: the root file system has a default size of 1GB now. Before Ignite-UX version C.7.0, HP-UX 11i v2 only required a minimum of 200MB.

```

logical_volume {
    mount_point = "/"
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
    contiguous_allocation = true
    bad_block_relocate = false

    "Create separate volumes (/usr, /var, ...)" {
        size = 1GB
    } else {
        size = remaining | 1GB + _hp_pri_swap
    }
}

```

Next, primary swap is defined. The usage has been set to swap and dump. The `mount_point` is set to "primary". The `vgdisk[0]` value enables you to specify that this logical volume be placed onto the first disk in the root volume group. Refer to `instl_adm(4)` for more information on the `vgdisk` keyword.

Next are the usual settings for primary swap of contiguous allocation and disabled bad block relocation. The size must be at least `_hp_min_swap` and up to `_hp_pri_swap` depending on the size of the disk.

```
logical_volume {
    mount_point = "primary"
    usage = SWAP_DUMP
    # Map it to the root disk to ensure it is considered primary
    # without this the secondary swap gets sorted first and considered
    # primary.
    vgdisk[0]
    contiguous_allocation = true
    bad_block_relocate = false

    size = _hp_min_swap | remaining | _hp_pri_swap
}
```

You looked at `_hp_sec_swap` earlier in the configuration file to see how it was defined. The only place you can give it a value is by using the **Additional** button on the **Basic** tab in the Ignite-UX GUI. You can define secondary swap using the **File system** tab, but you must select the disk to avoid having it on the same disk as primary swap. The configuration enables you to have secondary swap/dump. You only have to care about its size, not placement.

So if the variable `_hp_sec_swap` is >0MB, swap/dump space is created. If there is more than one disk in the root volume group, this swap/dump space is created on the second disk. If you provide equal-sized primary and secondary swap, it creates an interleaved swap setup.

Set this swap space to be also used as a dump space (contiguous allocation and bad block relocation off), and lastly set the size to be `_hp_sec_swap`. Secondary swap, if not used for dump, does not require contiguous allocation, and `bad_block_relocate` can be set to true.

```
_hp_sec_swap > 0Mb {
    logical_volume {
        mount_point = "secondary"
        usage = SWAP_DUMP
        (num_vgdisks > 1)
        {
            # This maps the secondary swap space to
            # the second disk when one is defined.
            vgdisk[1]
        }
        contiguous_allocation = true # allows use as dump
        bad_block_relocate = false # allows use as dump

        size = _hp_sec_swap
    }
}
```

There is one slight difference here compared to the other logical volume definitions we've seen so far. For HP9000 systems of the disk layout LVM with HFS, `/stand` must be an HFS file system. The HP9000 boot loader can't read a VxFS file system, so it must be an HFS file system (otherwise the

system won't boot). The size defined here is fixed at 2000MB. Note that impacts exist later that prevent you from decreasing the size below approximately 1GB. The reason for the much larger size is that kernels are considerably larger on HP-UX 11i v3 than previous HP-UX releases.

```

logical_volume {
    mount_point = "/stand"
    (is_hppa |
     _hp_disk_layout == "Logical Volume Manager (LVM) with HFS") {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }
    contiguous_allocation = true
    bad_block_relocate = false

    # The /stand volume should be at least large enough to hold
    # 5 kernels to support multiple configs and debugging.
    # The /stand volume absolutely needs to be large enough to
    # hold 2 kernels.
    #
    # Note that for PA-RISC systems /stand must be no larger than
    # 2000MB to avoid firmware addressing issues. The end of the
    # /stand volume must be < 2GB into the disk. Some space for
    # initial boot LIF content and volume overhead must be
    # considered in sizing /stand for PA-RISC systems.
    size = 2000MB
}

```

The other logical volumes are controlled by the use model "Create separate volumes (/usr, /var, ...)" . When it is set to true (and it is by default as we saw earlier) other file systems are created.

```
"Create separate volumes (/usr, /var, ...)" {
```

Below is the definition of /usr. Compared to previous releases of HP-UX the definition is simplified, since the size no longer changes depending on the size of the root disk. We have large files enabled, and only for the disk layout "Logical Volume Manager (LVM) with HFS" is the file system type set to HFS.

The size is a minimum of 512MB, and can have a maximum size of the impacts with a minimum of 20% free. It can be resized by Ignite-UX to take into account the remaining space if there is not enough to meet the maximum size.

The configuration dealing with _hp_root_grp_striped is last, and enables the logical volume to be striped over all the disks in the root volume or disk group (assuming there is more than one since it won't do much if only one disk is present). This configuration is present in other logical volume definitions and won't be mentioned again.

```

logical_volume {
    mount_point = "/usr"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS

```

```

        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }

    size = 512MB | remaining | 25% free

    _hp_root_grp_striped == "YES" {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}

```

The definition of /opt is identical to /usr. Please refer to the discussion of /usr for more information.

```

logical_volume {
    mount_point = "/opt"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }

    size = 512MB | remaining | 25% free

    _hp_root_grp_striped == "YES" {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}

```

The definition of /var is fairly similar to the logical volume definitions we have already seen. One main difference is that for the disk layout “Logical Volume Manager (LVM) with HFS”, the inode density for HFS⁵³ file systems has been changed because of a need for more inodes in /var. (The reason given is for the SD IPD, but if lots of files are created in /var/tmp or other places in /var, you would also need a higher inode density.)

The file system is sized so it has a minimum of 512MB and a maximum size of 512MB plus _hp_pri_swap. This provides enough space to save a crash dump. (The calculation is a best guess since you could have multiple dump spaces, but most crash dumps are only partial dumps. This assumption is not valid if you define a separate file system for /var/adm/crash.)

```

logical_volume {
    mount_point = "/var"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        # /var needs lots of inodes due to SD's IPD
        nbspi = 2048
        usage = HFS

```

⁵³ Not required for VxFS file systems since they automatically allocate new inodes as needed.

```

        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxF
        blksize = _hp_VxF_blksize
    }

    size = 512Mb | remaining | 512Mb + _hp_pri_swap

    _hp_root_grp_striped == "YES" {
        stripes =
        stripe_size = _hp_FS_stripe_size
    }
}

```

In previous releases of HP-UX, the /tmp file system size was based upon the size of the root disk (64MB if less than approximately 8GB, or 200MB if greater). The size of the file system has been increased to be simply 512MB. The new default size is much more usable than the much smaller default size for previous releases of HP-UX.

```

logical_volume {
    mount_point = "/tmp"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        nbpi = 2048
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxF
        blksize = _hp_VxF_blksize
    }

    size = 512Mb

    _hp_root_grp_striped == "YES" {
        stripes =
        stripe_size = _hp_FS_stripe_size
    }
}
} # "Create separate volumes (/usr, /var, ...)"

```

If the use mode “Create /export volume” is set to **TRUE**, then the /export file system is created. The size allocated for the file system is up to 100MB, and can be reduced by Ignite-UX depending on any impacts statements affecting the file system and available space. This file system is rarely created.

```

"Create /export volume" {
    logical_volume {
        mount_point = "/export"
        largefiles = true
        _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
            usage = HFS
            blksize = _hp_HFS_blksize
            fragsize = _hp_HFS_fragsize
        } else {
            usage = VxF
            blksize = _hp_VxF_blksize
        }
    }
}

```

```

        size = remaining | 100Mb

        _hp_root_grp_striped == "YES"  {
            stripes = *
            stripe_size = _hp_FS_stripe_size
        }
    }
}

```

The /home directory has also changed in size from previous releases of HP-UX, making it much more usable in a default state. The new size is up to 100MB (depending on free space) - previously it was up to 20MB in size.

```

logical_volume {
    mount_point = "/home"
    largefiles = true
    _hp_disk_layout == "Logical Volume Manager (LVM) with HFS" {
        usage = HFS
        blksize = _hp_HFS_blksize
        fragsize = _hp_HFS_fragsize
    } else {
        usage = VxFS
        blksize = _hp_VxFS_blksize
    }

    size = remaining | 100Mb

    _hp_root_grp_striped == "YES"  {
        stripes = *
        stripe_size = _hp_FS_stripe_size
    }
}
} # root group
} # Volume Manager disk layout.

```

Little explanation is required here. These variables are used in the UI (itool). Since they're directly manipulated by the UI, they shouldn't be visible via the additional button, so they're made invisible. Following that, there are some variables that need to be visible via the additional button depending on the disk layout chosen.

```

# Set the variables that are recognized by the UI as invisible
# so that they don't show on the "Additional" screen:
_hp_locale           visible_if false
_hp_cfg_detail_level visible_if false
_hp_pri_swap          visible_if false
_hp_min_swap          visible_if false
_hp_disk_layout       visible_if false
_hp_default_cur_lan_dev visible_if false
_hp_default_final_lan_dev visible_if false
_hp_keyboard          visible_if false
_hp_root_disk         visible_if false
_hp_boot_dev_path     visible_if false

_hp_disk_layout == "VERITAS Volume Manager (VxVM) with VxFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with HFS" |
_hp_disk_layout == "Logical Volume Manager (LVM) with VxFS"
{
    "Create /export volume"                      visible_if true
    "Create separate volumes (/usr, /var, ...)"  visible_if true

```

```

        _hp_sec_swap      visible_if true
        _hp_root_grp_striped visible_if (_hp_root_grp_disks > 1)
        _hp_root_grp_disks  visible_if true
    }
else
{
    "Create /export volume"           visible_if false
    "Create separate volumes (/usr, /var, ...)" visible_if false
    _hp_sec_swap      visible_if false
    _hp_root_grp_striped visible_if false
    _hp_root_grp_disks  visible_if false
}

```

The following impacts have changed from previous releases. Note the large impacts on /stand. This prevents you from decreasing the size of /stand below about 1GB in size.

```

#
# During install Ignite-UX will build an HP-UX kernel. The size
# of this kernel is not included in software impacts since it is
# not directly installed. It is also important that the size of
# /stand be large enough to hold at least 2 or 3 kernels. This
# impacts specification forces /stand to have at least enough
# space for 2 or 3 kernels.
#
init sw_sel "impact for min space in /stand" {
    sw_source = "core"
    visible_if = false
    impacts = "/stand" 800Mb
} = TRUE

#
# SD control scripts run during install and most system software
# relies on /var having available free space to operate. Temporary
# files placed in /var are not part of the installed content and
# that space is often not included in product packages. This /var
# impact provides an overall /var impact for these system software
# packages.
#
init sw_sel "impact by scripts on /var" {
    sw_source = "core"
    visible_if = false
    impacts = "/var" 300Mb
} = TRUE

```

The following configuration statements add a selection to the **Additional** button that enables the user to control the use of autoboot. Ignite-UX normally runs no matter what the autoboot flag is set to (refer to *setboot(1M)*). It forces it on if it was off, then later turns it back off for the final reboot when the system would normally come back. This configuration enables the user to manually boot the system when autoboot is disabled.

Ignite-UX normally forces the system to perform the first reboot automatically when rebooting from the installation kernel to the newly created or recovered kernel. This reboot is forced no matter what the autoboot flag is set to (refer to *setboot(1M)*). Ignite-UX forces it on if it was off; prior to the final reboot Ignite-UX returns the autoboot flag to off. If *_hp_force_autoboot* is set to NO and the system's autoboot flag is not set to on, the system does not perform a reboot.

Important:

If you disable autoboot and then install a system, the system never executes a final reboot without manual intervention. HP does not support booting into any other mode than the default multi-user run level at this point. The installation or recovery session has not yet completed, and some steps must still be performed by Ignite-UX. Booting into anything other than the default run level does not allow the installation or recovery to complete. This is the intended behavior of Ignite-UX and not a flaw. In general (except when investigating Ignite-UX issues on advice from HP Support), you should never change the value of the _hp_force_autoboot variable.

```
#  
# Variable to give user control over the use of setboot. When the variable  
# is YES, autoboot is set before the first reboot and restored after the  
# first reboot. When the variable is NO, autoboot is never changed.  
#  
enum _hp_force_autoboot  
_hp_force_autoboot = { "YES", "NO" }  
init _hp_force_autoboot = "YES"  
_hp_force_autoboot help_text "Force Ignite-UX autoboot?"
```

This variable is visible via the **Additional** button on the **Basic** tab in itool, and allows you to control whether DHCP will be active on the final system or not. In this context it does not control DHCP usage by Ignite-UX during a network boot. To affect a network boot so Ignite-UX will not attempt to use DHCP to gain an IP address, this variable must be defined in the install file system.

```
#  
# Variable to give user control over DHCP.  
# If YES, then DHCP is turned off on the system.  
#  
enum _disable_dhcp  
_disable_dhcp = { "YES", "NO" }  
  
init _disable_dhcp = "NO"  
_disable_dhcp help_text "Disable DHCP?"  
_disable_dhcp == "YES" {  
    DISABLE_DHCP=TRUE  
}
```

This configuration (like the above configuration for DHCP) exists in previous releases of Ignite-UX for earlier revisions of HP-UX. The _hp_addnl_fs_free_pct allows you to define an additional percentage free space that will be added to all file systems being created by Ignite-UX. The default for this has always been 10%, but it can be changed to 0%. The additional free space allowed by this variable is flexible. If Ignite-UX cannot fit the existing volumes into a disk once the extra free space has been added, the amount added by this variable can be reduced until the volumes fit into the disk space.

You can change the value of this variable via the **Additional** button on the **Basic** tab in the Ignite-UX GUI..

```
#  
# JAGae82704 fix. Variable to give user control over the amount of  
# free space left over in a logical volume. It defaults to 10. The
```

```

# range provided if you click on the button is from 0-10, but a larger
# value can be entered directly in the field if desired.
#
_hp_addnl_fs_free_pct = { 0,1,2,3,4,5,6,7,8,9,10 }
init _hp_addnl_fs_free_pct = 10
_hp_addnl_fs_free_pct help_text "Additional free space %"

```

The configuration that used to exist before this variable for previous releases is no longer present. It is no longer required with HP-UX 11i v3. Most of the configuration involved makes sure that 32bit-only systems did not attempt to install 64bit HP-UX (and vice versa) - HP-UX 11i v3 only supports 64bit systems.

```
_hp_os_bitness visible_if false
```

This variable allows you to control Hyper-Threading support on those HP Integrity servers that support it. By default it is turned on if the system can support it. Note the introduction of a new configuration language feature (similar to `is_hppa` and `is_ia64`) called `is_ht_capable`. If you need to tune the kernel or install software based on whether a system supports Hyper-Threading, you can use this new feature to do this.

```

#
# Add definition of _hp_ht_enable, to provide hook to enable
# hyperthreading on CPUs that support this for 11.31 and later. Provide an
# install-time method for configuring the HT firmware enabled/disabled
# state.
#
is_ht_capable {
    enum _hp_ht_enable
    _hp_ht_enable = {"true", "false"}
    init _hp_ht_enable="true"
    _hp_ht_enable help_text "Enable firmware hyperthreading"
} else {
    _hp_ht_enable visible_if false
}

```

There have been significant changes to the per-release configuration file for HP-UX 11i v3, however these changes have both reduced the size and complexity of the configuration file.

Special variables

Ignite-UX uses the following special variables to control aspects of installation or recovery, or to modify the state of the final system. The descriptions of these variables come from the `inst1_adm(4)` manpage, with further explanation provided where possible. If the variable appears directly in or impacts the Ignite-UX GUI, this information is also explained.

`_hp_locale`

The `geocustoms` user interface (UI) sets this string variable to the language locale that the user has chosen. If set to the string `ASK_AT_FIRST_BOOT` then the `geocustoms(1M)` application is invoked when the installation is complete, allowing the user to pick the desired language. (The `geocustoms` application is not available prior to HP-UX 10.30). Setting the value to `SET_NULL_LOCALE` leaves the system at default with no `LANG` variable set, which causes the commands to use the internal messages that provide better performance.

If `_hp_locale` is not set, it is set internally to the first value of the locale keyword for a selected `sw_sel`. If multiple selected `sw_sele`s contain locale keywords, the `_hp_locale` variable is set to the first value of the locale of the first selected `sw_sel`.

The `geocustoms`⁵⁴ application is used to manage the default locale on a system. If a locale is not important, do not set one. If you set `_hp_locale`, set it to a locale that exists on the final installed system, it should be visible in the output of `locale -a`.

You can run the `geocustoms` command at any time to manage the languages installed on your system (for example, you can remove some or change the default locale). You cannot use the `geocustoms` command to install non-existent locales back again.

The `geocustoms` program modifies `/etc/rc.config.d/LANG` to set a locale. The value of `SET_NULL_LOCALE` for `_hp_locale` leaves `/etc/rc.config.d/LANG` with no locale set.

The `geocustoms` program is easy to use because it provides a full-screen menu-driven interface for managing locale support. It does have command line options if you prefer not to use the user interface.

The standard configuration files set the variable `_hp_locale visible_if false` to prevent this variable from being visible or changeable when using the **Additional** button in the Ignite-UX GUI. The locale is chosen using the **Languages** button on the **Basic** tab in the Ignite-UX GUI.

`_hp_cfg_detail_level`

This internal string variable is fundamental to the configuration file management done by Ignite-UX. It is used primarily in client-specific configuration files. It contains a list of option characters that represent which aspects of the configuration file are modified by the user interface. This represents the areas of information that the configuration file, written by the UI, contains.

The Ignite-UX process uses this information in case it needs to rewrite the configuration file. The configuration file is rewritten both by the user interface and by the client installation process. In these cases, the process uses `_hp_cfg_detail_level` to determine how much information is represented by the configuration file and so it can write out just the minimal amount of information and let the rest be supplied by the other configuration files in the `INDEX` file.

The option characters recognized in the `_hp_cfg_detail_level` string that is found in a client-specific config file are as follows:

- i `INDEX cfg` clause selection (file contains the line telling which `INDEX` file entry should be used).
- v Variable and use model settings.
- s Software selection settings.
- S Modified software selection definitions.
- r Modified software source definitions (depot information).
- f Modified file system information.
- p System identity information.
- T The `post_config_script` selections settings.
- h Hardware control information (`hw_instance_num` statements).
- l Other control information (global `mod_kernel` statements for example).

⁵⁴ The `geocustoms` command is deprecated and is planned for obsolescence in a future release of HP-UX.

For example, the line `_hp_cfg_detail_level="ivsp"` in a client configuration file would indicate that the file contains information about which `cfg INDEX` selection is to be used, as well as the variable settings, software selection settings, and system parameters.

The user interface rewrites the `config` and `config.full` files in the per-client directories on the Ignite-UX server for network installation and recovery, or in the `/var/opt/ignite/local` directory for local media installation and recovery.

For network recoveries and installations, if you manually modify a file that Ignite-UX can modify, you must be careful that the `_hp_cfg_detail_level` contains the appropriate level of detail since it tells Ignite-UX what information is represented in the file that is what should be written back out into the file when it is rewritten by anything. If you add information that is not listed in the variable `_hp_cfg_detail_level`, it can be read, but then, the unlisted information is not written back out again even if other information in the file is modified and the file needs to be rewritten.

The standard configuration files use `_hp_cfg_detail_level visible_if false` to prevent this variable from being visible or changeable when using the **Additional** button on the **Basic** tab in the Ignite-UX GUI.

`_hp_pri_swap`

The `_hp_pri_swap` is the integer variable that is set by the Ignite-UX GUI to indicate how much swap the user desires to have allocated on the root disk/volume-group. The default LVM primary swap volume is defined as a size range with `_hp_pri_swap` being the desired size and `_hp_min_swap` being the minimum size.

On the **Basic** tab in the Ignite-UX GUI, if you change the value of the **Root Swap (MB)...** field you have changed the value of `_hp_pri_swap`.

Important:

If you navigate to the File system tab and select the primary swap volume to change its size, it does not reflect the change into the value on the Basic tab. That value remains the size as the primary swap. However, if you return to the Basic tab and change the field labeled Root Swap (MB)... at all, it changes the primary swap size on the File system tab, so exercise caution.

The standard configuration files use the variable `_hp_pri_swap visible_if false` to prevent this variable from being visible or changeable from the **Additional** button in the Ignite-UX GUI.

`_hp_min_swap`

The `_hp_min_swap` is the integer variable that is used as the minimum size the primary swap volume can be reduced to if there is not enough disk space for all other volumes. This value defaults to an amount that should allow the system to boot and run HP-UX.

On the **File system** tab in the Ignite-UX GUI, if you select primary swap in the list of file systems you see an immediate change to the fields below the section list; the field labeled **Min:** is the value of `_hp_min_swap`.

The standard configuration files use `_hp_min_swap visible_if false` to prevent this variable from being visible or changeable when using the **Additional** button in the Ignite-UX GUI.

_hp_disk_layout

The `_hp_disk_layout` string variable is set by the Ignite-UX GUI to indicate which disk layout (LVM, whole-disk, etc) that you have selected. As configurations are saved, the list of values that it can have is increased to contain any modified layouts.

There are ways you should and should not define this variable. The following is the correct way to define the `_hp_disk_layout` variable; this assumes that `_hp_root_disk` has already been set.

```
_hp_disk_layout= { "N4000 with 9Gb disk", "N4000 with 18Gb disk",
" N4000 with 36Gb disk", "N4000 with 72Gb disk" }
(model ~ "9000/N4000")
{
    (disk[_hp_root_disk].size<10000Mb)
    {
        init _hp_disk_layout=" N4000 with 9Gb disk"
    }
    else
    {
        (disk[_hp_root_disk].size<20000Mb)
        {
            init _hp_disk_layout=" N4000 with 18Gb disk"
        }
    }
    else
    {
        (disk[_hp_root_disk].size<400000Mb)
        {
            init _hp_disk_layout=" N4000 with 36Gb disk"
        }
    }
    else
    {
        init _hp_disk_layout=" N4000 with 72Gb disk"
    }
}
}
```

The preceding configuration correctly defines the settings for `_hp_disk_layout` to enable the Ignite-UX GUI to modify the disk configuration. If you want to set `_hp_disk_layout` so you cannot make any changes using the Ignite-UX GUI, follow the next example. In this configuration, you cannot change the disk layout after it has been selected by these tests:

```
_hp_disk_layout= { "N4000 with 9Gb disk", "N4000 with 18Gb disk",
" N4000 with 36Gb disk", "N4000 with 72Gb disk" }
(model ~ "9000/N4000")
{
    (disk[_hp_root_disk].size<10000Mb)
    {
        _hp_disk_layout="N4000 with 9Gb disk"
    }
    else
    {
        (disk[_hp_root_disk].size<20000Mb)
        {
            _hp_disk_layout="N4000 with 18Gb disk"
        }
    }
}
```

```
        (disk[_hp_root_disk].size<400000Mb)
    {
        _hp_disk_layout="N4000 with 36Gb disk"
    }
}
else
{
    _hp_disk_layout="N4000 with 72Gb disk"
}
}
}
```

The difference between the two sets of configuration is that the `init` keyword is missing from the second set of configuration statements when `_hp_disk_layout` is set to be a specific value. The `instl_adm(4)` manpage states:

init variable=value
Preceding the assignment with the init keyword means that the variable is to be initialized to the given value, but the user interface is allowed to alter the value later.

variable=value
When the init keyword is not used, then the variable cannot be changed by the user interface. This type of assignment is not recommended for "visible" variables.

Therefore, the `init` keyword must precede the value that `_hp_disk_layout` is set to; otherwise the Ignite-UX GUI cannot change it. When you modify anything using the Ignite-UX GUI, it wants to add a new value to the list that `_hp_disk_layout` can have and set it to the value of `_hp_disk_layout`. Unless `init` is used to give `_hp_disk_layout` its initial value, the Ignite-UX GUI cannot change it and all modifications to the disk layout are lost.

If you have an Ignite-UX server and a client that has been added or installed recently, you can test this as follows:

- ## 1. Enter:

```
cd /var/opt/ignite/clients  
itool -m pull -d <client>
```

2. Change the type of disk layout by selecting the **File system** tab and choosing the correct layout from the list.
 3. Quit the Ignite-UX GUI. You can choose any option on the exit screen; none of the options reboots or halts the system.
 4. Edit `/var/opt/ignite/clients/<MAC>/config` and remove `init` from the front of the `_hp_disk_layout` variable.
 5. Save the file.
 6. Rerun the Ignite-UX GUI using the `itool` command in Step 1. If you change the file system layout, it reverts to the original value.

Important:

Only perform this test with one of the standard unmodified layouts. Do not perform this test with a modified disk layout. Removing init prevents you from modifying anything to do with the file system.

The values in `_hp_disk_layout` are only indirectly visible. The list of values in the `_hp_disk_layout` variable is shown on the **Basic** tab in the Ignite-UX GUI in the **File System**: list. The use of `init` indirectly affects how `_hp_disk_layout` works.

The standard configuration files use the variable `_hp_disk_layout visible_if false` to prevent this variable from being visible or changeable using the **Additional** button in the Ignite-UX GUI.

As defined, the preceding configuration replaces any other values assigned to `_hp_disk_layout`. If you want to retain any currently defined disk layouts, you need to start the preceding examples with the following:

```
_hp_disk_layout+= { "N4000 with 9Gb disk", "N4000 with 18Gb disk",
" N4000 with 36Gb disk", "N4000 with 72Gb disk" }
```

The `+=` operator adds to current values rather than replacing any existing values.

`_hp_default_cur_lan_dev`

The `_hp_default_cur_lan_dev` string variable is set to the LAN device that is enabled during the Ignite-UX process. It is used when a LAN device is omitted from keywords that can accept a LAN interface specifier. This defaults to the interface that you picked during the install; in the non-interactive case, it defaults to the LAN device the system booted from or to `lan0` if not a network boot.

Therefore, a configuration that looks like this:

```
( lan[].driver ~ "btlan" ) {
...
}
```

With this variable, the preceding configuration is really evaluated as:

```
( lan[_hp_default_cur_lan_dev].driver ~ "btlan" ) {
...
}
```

The standard configuration files use `_hp_default_cur_lan_dev visible_if false` to prevent this variable from being visible or changeable using the **Additional** button in the Ignite-UX GUI.

`_hp_default_final_lan_dev`

The `_hp_default_final_lan_dev` string variable is similar to `_hp_default_cur_lan_dev`, but is used when network information is specified by using the `final` keyword. It defaults to `_hp_default_cur_lan_dev` if not set. The IP address associated

with the LAN device specified by this variable is used for the entry in the /etc/hosts file for the system's hostname.

That means that if you have the following partial configuration:

```
final ip_addr[ ]="15.30.129.47"
final netmask[ ]="255.255.255.0"
```

The IP address information is applied to the LAN interface defined by the variable _hp_default_final_lan_dev, and this IP address is the one associated with the name of the host in the /etc/hosts file. If the value of the variable _hp_default_final_lan_dev is not set explicitly, it defaults to the value of the variable _hp_default_cur_lan_dev.

The standard configuration files use _hp_default_final_lan_dev visible_if false to prevent this variable from being visible or changeable using the **Additional** button in the Ignite-UX GUI.

hp_keyboard

The _hp_keyboard string variable is set by the Ignite-UX GUI to the keyboard language and mapping desired. The kbdlang keyword that was used for this in past releases is equivalent to this variable. Setting this variable in the INSTALLFS file (using instl_adm) prevents the system from prompting for a keyboard type during the install. This information is stored in /etc/kbdlang on the final system.

If kbdlang is not set, it causes Ignite-UX to run the following command on workstations and servers that have graphics consoles:

```
# /sbin/itemap -i -L -w /etc/kbdlang
```

Not all servers support graphics consoles. If you run the itemap command on a system that does not have the required hardware installed, the following message appears when run with the -v option:

```
# itemap -i -L -v -w /etc/kbdlang
No framebuffer device.
```

The -L option causes itemap to load the appropriate keymap for non-PS2 keyboards into the Internal Terminal Emulator (ITE)⁵⁵ and if a PS2 keyboard is found, the command will interactively prompt for the type of keyboard. The list is created by showing all of the keyboards found in the file /etc/X11/XHPKeymaps that start with PS2_DIN. A list of possible values that this variable can take is generated with the following command:

```
# keymap_ed -l | awk '$5 ~ "^PS2_DIN" { print $5 }'
```

The X11 patches that modify /etc/X11/XHPKeymaps may add or remove keyboards from this list (although not very often) so the list is dependent on X graphics patch revisions. If you need to inspect the XHPkeymaps files from archives or other sources, the keymap_ed command, with the

⁵⁵ Internal Terminal Emulator is an emulation of an HP style terminal in the kernel when a graphics display is the console device.

`-k` option, enables you to specify an alternate file to read. Refer to *keymap_ed(1)* for more information.

The standard configuration files use `_hp_keyboard visible_if false` to prevent this variable from being visible or changeable using the **Additional** button in the Ignite-UX GUI. It can be set for clients with required hardware from the **Keyboard** button on the **Basic** tab. If the hardware is not present, the only choice presented is "Not Applicable". If "Not Applicable" is chosen or left as a default, then the keyboard language/mapping is prompted for when the system boots the first time if the system has a graphics console and keyboard attached.

`_hp_root_disk`

The `_hp_root_disk` string variable is set by the Ignite-UX GUI to contain the hardware path of the disk that the user has chosen to be the root disk. The value is initialized to the value of `_hp_primary_path` in `/opt/ignite/data/Rel_*/config`, and can be overridden in other config files. If no config file were to set the initial value, then it would default to the disk with the hardware path with the highest SCSI priority.

The standard installation configuration attempts to set the variable `_hp_root_disk` to the value of `_hp_primary_path` if `_hp_primary_path` is set. The variable `_hp_primary_path` holds the hardware path to the disk that the system is set up to boot from by default (set by `setboot` or selected at the BCH⁵⁶ prompt). If the primary path setup for the system points to a nonexistent device, `_hp_root_disk` is not be set by this configuration. Instead, Ignite-UX chooses what it determines to be the "best" device for the variable instead.

```
(_hp_primary_path != "" & !(_hp_root_disk ~ ".*"))
{
    init _hp_root_disk=_hp_primary_path
}
```

During an interactive installation or recovery, you can modify this value using the Ignite-UX GUI using the **Root Disk...** button on the **Basic** tab to get a selection of choices. It can also be changed using the **File system** tab in **Add/Remove Disks...** dialog box by changing the disks that are assigned to root volume or disk group. However, in the **Add/Remove Disks...** dialog box Ignite-UX chooses which of the disks (if more than one) is assigned to `_hp_root_disk` (initially anyway, after that they can be modified).

The standard configuration files use `_hp_root_disk visible_if false` to prevent this variable from being from visible or changeable using the **Additional** button in the Ignite-UX GUI.

`_hp_boot_dev_path`

The variable `_hp_boot_dev_path` is not defined by *inst1_adm(4)*. Do not rely on this description and the behavior defined here because they may change without notice.

If `bootsys` was used on a disk that was not destroyed during the installation, then the disk has an `AUTO` file left over from `bootsys`. Typically, you want to be able to boot from that disk after the installation is done. The disk that you booted from originally is pointed to by `_hp_boot_dev_path`.

⁵⁶ Boot Console Handler: the firmware boot prompt on PA-RISC systems.

The `bootsys` command stores the original `AUTO` file line as a commented-out string in the `AUTO` file itself (using '#' as the comment char). Ignite-UX uses the value in `_hp_boot_dev_path` to reset the `AUTO` file in the `LIF` of that device if that disk is not used in the install.

On Itanium®-based systems the `AUTO` file is stored in the Extensible Firmware Interface (EFI) partition. Ignite-UX treats that `AUTO` file exactly the same way as the PA-RISC `AUTO` file stored in a boot `LIF`.

Do not have this variable set by any custom configuration file; its use is intended only for internal use within Ignite-UX.

The standard configuration files use `_hp_boot_dev_path visible_if false` to prevent this variable being visible or changeable using the **Additional** button in the Ignite-UX GUI.

`_hp_primary_path`

The `_hp_primary_path` string variable is set during the initial startup of Ignite-UX on the client. The variable is set to the system's primary boot path (as set in stable storage and read with `setboot`). If the system's primary boot path is incorrectly set to a non-existent device, then `_hp_primary_path` is set to the empty string ("").

Use this variable for reference only, or to set other variables. Do not have a configuration file assign it a value. Ignite-UX sets the value.

`_hp_primary_partition_size`

The `_hp_primary_partition_size` variable controls the size of the HP-UX partition on an EFI boot disk on HP Itanium®-based systems. Normally, the value of this variable is calculated as the remainder of the disk being installed to after the other partition sizes have been subtracted. For more information regarding the EFI and HPSP partitions, see "`_hp_efi_partition_size`" and "`_hp_service_partition_size`".

Typically, there is no reason to set the `_hp_primary_partition_size` variable in a custom configuration. If you do set this variable and the size is too large, it is automatically reduced; if set too low, part of the boot disk remains unused.

`_hp_efi_partition_size`

The `_hp_efi_partition_size` integer variable is set by the Ignite-UX GUI to indicate the size of the HP-UX Extensible Firmware Interface (EFI) boot partition on Itanium®-based systems. This disk partition holds loader software and other utilities used to boot the HP-UX operating system. This value defaults to a size that allows some space for additional boot partition content in future HP-UX releases. This variable is adjusted to the minimum EFI boot partition size if set to zero or a value less than the minimum. This variable is supported starting with HP-UX 11.23 and is only supported on Itanium®-based systems.

This variable sets the size of the EFI partition located on bootable disks on Itanium®-based systems. The minimum size is 150MB (default size 500MB) currently.

This variable can affect recoveries. For example, if there is no free space left on the boot disk and an Ignite-UX version is released that increases the size of the EFI partition, the recovery or installation requires manual intervention if the EFI partition size minimum is increased because of lack of disk space.

This variable can affect new installations or reinstallations. If you have sized the boot disk of a system to a certain size and have not included the size of the EFI Partition (and Service Partition, see "`_hp_service_partition_size`") you may not have enough disk space to install the system in the way that you want.

`_hp_service_partition_size`

The `_hp_service_partition_size` integer variable is set by the Ignite-UX GUI and indicates the size of the HP Service Partition on Itanium®-based systems. This disk partition holds diagnostics software, saved system and hardware state data, and other utilities used to verify and update hardware functionality. This value defaults to a size expected by the diagnostics software. A zero value indicates that no HP Service Partition is to be created. This variable is adjusted to the minimum HP Service Partition size if set to a non-zero value less than the minimum. This variable is supported starting with HP-UX 11.23 and only supported on Itanium®-based systems.

This variable defines the size of the HP Service Partition located on bootable disks on Itanium®-based systems. The minimum size is 400MB (default size 400MB) at this writing. If the size of the HP Service Partition is set to zero, the partition is not created.

This variable can affect recoveries. For example, if there is no free space left on the boot disk and an Ignite-UX version is released that increases the size of the HP Service Partition, the recovery or installation requires manual intervention if the HP Service Partition size minimum is increased because of lack of disk space

This variable can have an impact on new installations or reinstallations. If you have sized the boot disk of a system to a certain size and have not included the size of the HP Service Partition (and EFI Partition, see "`_hp_efi_partition_size`"), you may not have enough disk space to install the system in the way that you want.

`_hp_root_grp_disks`

The `_hp_root_grp_disks` integer variable indicates how many disks to put into the root volume group.

The `itool` wizard prompts you for the number of disks in the root volume/disk group and whether they should be striped. However, in the Ignite-UX GUI you must use the **Additional** button on the **Basic** tab to display the **# of disks in root VG**.

The default configuration described earlier (for B.11.11) sets `_hp_root_grp_disks` to an enum from one to the value of the internal variable `num_disks`. During an installation, Ignite-UX adds disks automatically to the root volume/disk group to make up this number. Because the disks are assigned automatically, you must ensure that the disks do not contain data that you do not want to lose. For example, in a SAN environment those disks may belong to another system.

`_hp_root_grp_striped`

The `_hp_root_grp_striped` string variable uses the possible values YES and NO to indicate if LVM data striping should be used on all disks in the root volume group if multiple disks are in the root group.

This variable is defined within the default installation configuration so you can easily specify whether the volumes within the root volume/disk group should be striped across all disks in the root volume/disk group.

The `i tool` wizard prompts you to specify whether the disks in the root volume/disk groups should be striped (as well as to specify the number of disks in the root volume/disk group on the same screen). However, in the Ignite-UX GUI you must use the **Additional** button on the **Basic** tab to display the item **Stripe root VG disks?**.

Even though `inst1_adm(4)` discusses LVM, striping also applies to VxVM (VxVM licensing may restrict its use). Remember that the root, boot, and primary swap+dump cannot be striped, so this only applies to other volumes in the root volume/disk group.

`_hp_addnl_fs_free_pct`

The `_hp_addnl_fs_free_pct` integer variable is used to control the amount of additional free space allocated to volumes beyond the space required to load the software. If this variable is not set, it defaults to 10 (10%).

The `_hp_addnl_fs_free_pct` variable can be changed using `Additional free space %` during an installation because of the following configuration.

```
_hp_addnl_fs_free_pct = { 0,1,2,3,4,5,6,7,8,9,10 }
init _hp_addnl_fs_free_pct = 10
_hp_addnl_fs_free_pct help_text "Additional free space %"
```

This variable is not present for recoveries.

The `inst1_adm(4)` manpage states the following about how `_hp_addnl_fs_free_pct` affects the size calculations for volumes:

In addition to this minimum size, an additional percentage is added to the volumes to ensure sufficient room after the system is installed. This percentage is by default 10%, but can be controlled by the configuration file variable `_hp_addnl_fs_free_pct`. If the disk capacity is insufficient to accommodate this additional free space, this additional percentage value is continually reduced by 1/2 until a value is found that allows all the volumes to fit. If the volumes still do not fit with 0% additional space, then the installation is not allowed to proceed.

Therefore, the variable does nothing when there is no free disk space; however, if there is free disk space, volumes on the system may be expanded in an unexpected manner.

`_hp_ignore_sw_impact`

The `_hp_ignore_sw_impact` integer variable can be set to one if you want to disable all effects that the impacts statements declared in the software selections have on the volume size calculations. This may be helpful if you want to ensure that Ignite-UX does not automatically modify the file-system volume sizes.

Be careful when setting this variable to 1. Ignite-UX does not change the file system sizes based upon the impacts statements associated with software products that are being installed. The installer must correctly set file system sizes so the system does not fill file systems during installation.

So, when do you use this variable? To answer that you need to look at how impacts are created. Consider the following commands⁵⁷:

⁵⁷ You do not call `gen_impacts` in this way, because the format that it accepts is not documented. This example is only to illustrate how impacts are generated.

```
$ print "f tmp/a 1024" | /opt/ignite/lbin/gen_impacts
    impacts = "tmp" 8Kb
$ print "f tmp/a 10240" | /opt/ignite/lbin/gen_impacts
    impacts = "tmp" 16Kb
```

The first command specifies a file size `tmp/a` of 1KB, but the impact is 8KB. The second command specifies a file size of `tmp/a` as 10KB, and the impact is 16KB.

Most systems these days use VxFS instead of HFS for file systems. VxFS file systems are extent-based file systems and can usually allocate (space allowing) blocks as small as 1KB to a file.

When you have the impacts for a HFS file system applied to a VxFS file system, the size requirements can be significantly overstated because of the assumptions that `gen_impacts`⁵⁸ makes.

The `gen_impacts` command assumes that all files are located on HFS file systems and have an 8KB fragment size and a 64KB block size. The `gen_impacts` command gets the approximate size correct when the file is large enough to use indirect block to track the space used by the file.

These overstated impacts are not considered wrong, just conservative. Ignite-UX cannot predict when impacts are being generated or what kind of file system on which the software is installed. The correct approach is to take the conservative approach.

Lastly, consider the following example of what effect different fragment sizes have on the impacts generated for a depot containing a copy of the B.11.11 Version Mission Critical (MC) Operating Environment (OE) with the default fragment size (8KB), a 2KB fragment size, and a 1KB fragment size:

```
# /usr/sbin/swlist -l file -d -a type -a name -a size @/depot/11lvl_mc/core |
awk ' $3 == "f" || $3 == "d" { printf("%s %s %s\n", $3, $2, $4) } ' |
/opt/ignite/lbin/gen_impacts -l 2 -f 1024
    impacts = "/" 1138Kb
    impacts = "/usr/lib" 599090Kb
    impacts = "/sbin/lib" 561Kb
    impacts = "/usr/conf" 107843Kb
    impacts = "/sbin/init.d" 664Kb
...
# /usr/sbin/swlist -l file -d -a type -a name -a size @/ depot/11lvl_mc/core |
awk ' $3 == "f" || $3 == "d" { printf("%s %s %s\n", $3, $2, $4) } ' |
/opt/ignite/lbin/gen_impacts -l 2 -f 2048
    impacts = "/" 1138Kb
    impacts = "/usr/lib" 603124Kb
    impacts = "/sbin/lib" 578Kb
    impacts = "/usr/conf" 108359Kb
    impacts = "/sbin/init.d" 734Kb
...
# /usr/sbin/swlist -l file -d -a type -a name -a size @/depot/11lvl_mc/core |
awk ' $3 == "f" || $3 == "d" { printf("%s %s %s\n", $3, $2, $4) } ' |
/opt/ignite/lbin/gen_impacts -l 2 -f 8192
    impacts = "/" 1138Kb
    impacts = "/usr/lib" 632140Kb
    impacts = "/sbin/lib" 684Kb
    impacts = "/usr/conf" 111703Kb
```

⁵⁸ The `archive_impact` and `make_arch_config` commands both call `gen_impacts` to generate impacts statements.

```

impacts = "/sbin/init.d" 1256Kb
...

```

The size differences in the first four impacts are summarized in Table 4 (the percentage difference is between the 2KB and 8KB fragments compared to the 1KB fragment impacts):

Table 4

Impacts	Size (1KB)	Size(2KB)	% difference	Size (8KB)	%difference
/	1138	1138	0.0%	1138	0.0%
/usr/lib	599090	603124	0.7%	632140	5.5%
/sbin/lib	561	578	3.0%	684	21.9%
/usr/conf	107843	108359	0.5%	111703	3.6%
/sbin/init.d	664	743	11.9%	1256	89.2%

The percentage increases are somewhat variable but when more space is used for the larger impacts, the larger the fragment size. The impacts are made much larger when a directory structure contains many very small files.

The reason why "/" is the same size is that `gen_impacts` allocates 1KB to an impact level whenever it sees a directory mentioned. The "/" directory in this case appears 1138 times in the `swlist` output (so multiple mentions of a directory in the output from `swlist` can also bloat impacts).

Caution:

If you change or somehow recalculate impacts in any way, you must be very careful. Impacts that do not fully take into account the software impacts to a system may lead to configurations that cannot install a system (because of file system full problems).

Another issue that may cause overstated impacts is hard links. Before Ignite-UX version C.6.4 Ignite-UX would assume that every file that was to be installed from an archive (this includes recovery archives) would occupy the same amount of space as its size ignoring the link count. The space contributed by hard linked files will not be counted multiple times (once for each link) from Ignite-UX version C.6.4 onwards. If you have upgraded an Ignite server from a version prior to C.6.4 you might consider regenerating impacts for golden images.

`_hp_custom_sys`

The `_hp_custom_sys` string variable is used in conditional statements surrounding network and system identity information when the Ignite-UX GUI writes out a configuration file. This variable can be used in conditionals in configuration files to define multiple sets of network parameters that can be selected from the **Additional** screen. You can see how this variable is used by clicking the **Save As** button during an installation after modifying the parameters under the **System** tab and looking at the resulting file in the `/var/opt/ignite/saved_cfgs` directory.

Never use the `_hp_custom_sys` variable in custom configurations. Ignite-UX uses this variable in the configuration files it can produce so setting its value outside of those files may interfere with its use by Ignite-UX.

_hp_lanadmin_args

If your network requires that the default MTU or speed be set using the lanadmin command to operate correctly, you can specify the arguments to the lanadmin command using the _hp_lanadmin_args variable. This variable setting must be set in the `INSTALLFS` file using the `instl_adm` command. Any change made with this variable only affects the installation and is not permanently applied to the system. For example, to set the MTU size to 1500, the line would be:

```
init _hp_lanadmin_args="-M 1500"
```

To set the speed for a 100-Base-T interface to full duplex (the default is half-duplex) you could use the setting:

```
init _hp_lanadmin_args="-S 1"
```

Setting the MTU value with lanadmin only works for the NIO/HPPB FDDI interface.

Additional lanadmin libraries have been added over time to allow many interfaces to be controlled with the `-X` option to lanadmin.

This variable is only useful when placed into the installation file system so the changes are available at boot time. For example:

```
( lan[ ].driver ~ "btlan" ) {  
    _hp_lanadmin_args="-X 100FD"  
}
```

You might choose to place the preceding configuration lines into the installation file system if you have systems (using LAN interfaces controlled by the `btlan` driver) connected to switches that have been set to operate only at 100 Full Duplex.

The reason you would need to do this is that there are no startup scripts when installing a system. The `btlan` driver attempts to auto-negotiate the speed and duplex settings with the device to which it is connected. If the switch is set to 100 Full Duplex the auto-negotiation fails and the system runs at 100 Half Duplex.

When the system and the switch attempt to communicate using mismatched duplexes, the values result in problems and performance suffers. This leads to either very long installation times or a failed installation.

For additional information, see "

Problems that can be solved with `_hp_lanadmin_args`.

`_hp_nfs_mount_opts`

The `_hp_nfs_mount_opts` string variable is set in the `INSTALLFS` file and is used to supply additional options to the NFS mounts that are performed during the installation. This is intended for use when the default options are not appropriate for your network. (Refer to `mount_nfs(1M)` for valid options). For example, to set the read and write NFS buffer size to 1K:

```
init _hp_nfs_mount_opts="-orsize=1024 -owsize=1024"
```

This option must be given in the installation file system for it to take effect, although a read buffer size of 4-8KB performs better for recovery over a local network. Otherwise, if you were to read the value of this variable from the Ignite-UX server, the NFS file system would already be mounted and the options would be too late to be applied.

Before Ignite-UX version C.6.0.x the `inst1_adm(4)` manpage is incorrect; the NFS mount command accepts one and only one `-o` option. All of the options to the `-o` parameter must be given in a comma separated list. The following example does work:

```
init _hp_nfs_mount_opts="-orsize=1024,wszie=1024"
```

Do not use NFS mount options documented in `mount_nfs(1M)` that interfere with the correct operation of Ignite-UX such as the following:

`-oro` — If you mount the file system read-only it is very hard for Ignite-UX to update files on it. It is also incompatible with the `-orw` option provided by Ignite-UX (some file systems are mounted read-only by Ignite-UX so the `-orw` option should not be given either).

`-obg` — Never place the mount attempts into the background as once the mount command returns successfully, Ignite-UX assumes that the file system has mounted and then continues. Ignite-UX provides a retry mechanism using the `_hp_nfs_mount_retries` variable.

`-overs=2` — Unless absolutely necessary do not use this option as NFS v2 does not allow you to use large files (golden image and recovery archives more than 2GB in size fails to install/recover).

Other NFS options may be used as needed unless they interfere with the operation of Ignite-UX.

`_hp_nfs_mount_retries`

The `_hp_nfs_mount_retries` integer variable is used to modify the default number of times that the NFS mounts are retried before failing. If this variable is not set, the mounts are retried 4 times before giving up. If you need to change this default, this variable should be specified in the `INSTALLFS` file. For example, to set the number of retries to 8:

```
init _hp_nfs_mount_retries=8
```

This is the number of retry attempts that Ignite-UX performs when attempting to mount an NFS file system. If Ignite-UX attempts to retry an NFS mount, messages similar to the following appear on the console and in the client's `install.log` file:

NOTE: Retrying: "/usr/sbin/mount -orw

```
10.0.0.1:/var/opt/ignite/clients /var/opt/ignite/clients"
NOTE:    Retrying: "/usr/sbin/mount -orw
10.0.0.1:/var/opt/ignite/clients /var/opt/ignite/clients"
NOTE:    Retrying: "/usr/sbin/mount -orw
10.0.0.1:/var/opt/ignite/clients /var/opt/ignite/clients"
...
...
```

The messages continue until either the amount of retries given in `_hp_nfs_mount_retries` is exceeded or the NFS file system mounts successfully.

`_hp_tftp_cmds`

The `_hp_tftp_cmds` string variable can be specified in the `INSTALLFS` file to supply additional instructions to the `tftp` commands that are used to transfer data during an installation. The commands supplied with this variable are passed as input to the `tftp` command along with the usual commands supplied by Ignite-UX. The most likely use of this would be to modify the retransmission-timeout (`rexmt`) and overall timeout (`timeout`) values. The default values that Ignite-UX uses are `rexmt=2 timeout=25`. (Refer to `tftp(1)` for more details). The string assigned to this variable should contain one `tftp` command statement per line. For example:

```
init _hp_tftp_cmds="rexmt 5
timeout 40"
```

Setting this variable can increase the number of retransmissions and the timeout. You may want to increase these values if the system being installed is a long way (network-wise) from the Ignite-UX server or if you have a fairly unreliable or overloaded network.

`_hp_hide_other_disks`

The `_hp_hide_other_disks` string variable can be set to one or more space-separated hardware paths of disks that should be "hidden" from being configured or otherwise modified during the install. It allows more disks to be hidden than what is provided by the `hide_boot_disk` keyword.

This variable is not especially useful for installation purposes, as you need to know all of the disk hardware paths that need to be hidden, which is difficult at installation time unless you know exactly what is in the system beforehand and what it looks like.

This is more useful during recoveries as Ignite-UX currently does them, as you can ensure that some disks cannot be accessed or modified during a recovery session.

`_hp_saved_detail_level`

The `_hp_save_detail_level` internal string variable is used in configuration files that have been created by a **Save-As** operation in the Ignite-UX GUI. Like `_hp_cfg_detail_level`, it contains a list of option characters that represent which aspects of the configuration file have been modified. The format is identical to the `_hp_cfg_detail_level` variable.

You should not use this variable in custom configurations. If you are using `itool` or other Ignite-UX internal commands to create configuration files that place this variable into a configuration file, you should remove any reference to this variable.

_hp_os_bitness

The `_hp_os_bitness` string variable is used to package software for running on a 32-bit operating system, 64-bit operating system, or both. The `_hp_os_bitness` variable is set to either 32 or 64 when an operating system is chosen. This is normally done in a configuration file by keying off the `sw_sel` for a 32- or 64-bit operating system. The `sw_sel` statements for certain applications rely on the setting of this variable to tell which version of the application to install.

This variable is set by the core operating system `sw_sel` that was selected, for example:

```
sw_sel "HPUXBase64" {  
    description = "HP-UX 64-bit Base OS"  
    ...  
}  
(sw_sel "HPUXBase64") {  
    _hp_os_bitness = "64"  
}
```

When the software `HPUXBase64` is selected (so the test "`(sw_sel "HPUXBase64")`" returns `true`) the variable `_hp_os_bitness` is set to 64. The `make_config` command automatically does this for you. However, if you manually write configuration files containing definitions of a core operating system depot or archive, you must set this variable to the correct bitness.

_hp_force_autoboot

The `_hp_force_autoboot` string variable is used to modify the behavior of Ignite-UX with respect to stable store's `autoboot` flag. The Ignite-UX configuration process has two parts separated by a reboot. By default (`_hp_force_autoboot="YES"`), Ignite-UX guarantees that `autoboot` is set during the installation process and then reset to its previous state (if necessary) at the end of the installation. If `_hp_force_autoboot="NO"`, Ignite-UX does not modify the `autoboot` flag in stable storage. This may mean that you do not have to do a manual boot from the primary path between the two parts of the Ignite-UX installation.

Ignite-UX normally forces the system to perform the first reboot automatically when rebooting from the installation kernel to the newly created or recovered kernel. This reboot is forced no matter what the `autoboot` flag is set to (refer to `setboot(1M)`). Ignite-UX forces it on if it was off; prior to the final reboot Ignite-UX returns the `autoboot` flag to off. If `_hp_force_autoboot` is set to NO and the system's `autoboot` flag is set to off, the system cannot complete the recovery and boot the final recovered or installed system, which leaves the system at the BCH, prompt.

Important:

If you disable `autoboot` and then install a system, the system never executes a final reboot without manual intervention. HP does not support booting into any other mode than the default multi-user run level at this point. The installation or recovery session has not yet completed and some steps must still be performed by Ignite-UX. Booting into anything other than the default run level does not allow the installation or recovery to complete. This is the intended behavior of Ignite-UX and not a flaw. In general (except when investigating Ignite-UX issues on advice from HP Support), you should never change the value of the `_hp_force_autoboot` variable.

`_hp_ikernel_os_release`

The `_hp_ikernel_os_release` variable is not defined in the Ignite-UX documentation and it should never be set explicitly by any custom written Ignite-UX configuration files.

When a client boots, part of starting Ignite-UX for installation or recovery is to create a file called "host.info". As of Ignite-UX version C.6.0.x, this variable is added to the host.info file and gives the HP-UX revision that the installation kernel is running. Ignite-UX uses this information to prevent you from selecting HP-UX releases that cannot be installed by the installation kernel.

For example, with Ignite-UX version C.6.0.x for PA-RISC systems, the B.11.11 installation kernel is used to install B.11.11 client systems but cannot be used to install B.11.23 systems. Therefore, if `_hp_ikernel_os_release` is set to B.11.11 Ignite-UX will not allow you to install HP-UX B.11.23.

`_hp_current_client_release`

The `_hp_current_client_release` variable determines what release of HP-UX the client is currently running and is used by the `ignite` and `make_net_recovery` commands. This variable should not be set in any custom-written configuration file. It is not used during installation or recovery.

`_HP_CLONING`

The `_HP_CLONING` variable is only created by `make_net_recovery` and `make_tape_recovery` as part of the configuration for a recovery archive. There seems to be some confusion about exactly what this variable does. It is defined as follows in the `control_cfg` file:

```
enum _HP_CLONING
    _HP_CLONING help_text "Cloning to different HW?"
    _HP_CLONING = { "TRUE", "FALSE" }
    (MODEL == "9000/800/SD16000")
    { init _HP_CLONING = "FALSE" }
    else
    { init _HP_CLONING = "TRUE" }
```

This defines `_HP_CLONING` as an `enum` and means that it can only have a value from one of the list of possible values that are defined for it (on the third line of the definition). The second line defines the help text that appears when using the **Additional** button on the **Basic** tab. If the model string matches the pattern, it defaults to FALSE; otherwise, it defaults to TRUE.

If you are using a recovery archive containing this configuration to clone one system to another with the same model string, you should manually change this value to TRUE using the **Additional** button on the **Basic** tab. This is to force a new kernel to be created. Although the model of system may be the same, it may have different I/O cards present. If the driver for those I/O cards is not present in the kernel being recovered, the final system will not have the drivers and those I/O cards will not be claimed.

This variable should never be defined in a custom configuration. It should **only** be defined in configurations created by `make_net_recovery` and `make_tape_recovery`.

The code that does this is in one of the scripts provided with Ignite-UX that is run during a recovery (`os_arch_post_1`):

```

# If cloning to a different model of machine, _HP_CLONING will
# be true (it can also be set via the Additional UI screen). In
# this case remove the kernel which will allow IUX to rebuild it,
# plus remove device files which will get rebuilt by insf or
# might be inappropriate.
if [[ "$_HP_CLONING" = "TRUE" ]] ; then
    remove_devs
    rm -f /stand/vmunix
fi

```

This variable does not control anything else.

`_hp_console_verbosity`

The `_hp_console_verbosity` variable allows you to reduce the amount of information that is printed to the screen when Ignite-UX executes a recovery or an installation. It has a default value of five that sets Ignite-UX to print everything; however, it only has to be three or more to print everything.

When set to two or less, it greatly reduces the amount of messages printed on the console. The messages not printed on the console are printed to the `install.log` file. This variable maybe useful when creating media for use by non-technical personnel who may not understand the messages being printed or may become concerned with some of the warning messages that Ignite-UX can print that may not be indicative of a real problem.

This variable is set in the installation file system. The following example reduces the number of messages written to the console during a recovery or installation by setting the variable to zero:

```
init _hp_console_verbosity=0
```

Warning:

If you reduce the amount of messages printed and a failure occurs, this may cause difficulties when attempting to determine the root cause of failures that occur early in an installation or recovery when the `install.log` file has not yet been moved to permanent storage.

`_hp_patch_save_files`

The `_hp_patch_save_files` variable is not actually a special variable and by itself, it means nothing to Ignite-UX. It is a construct used within the `hw_patches_cfg` file in the per release data directory (`/opt/ignite/data/Rel_B.#.#.#`). This construct is an enum that is accessible from the **Additional** button on the **Basic** tab in the Ignite-UX GUI. The following is an example of the `_hp_patch_save_files` enum definition:

```

#
# _hp_patch_save_files is used to determine if patches loaded from the
# core are committed, or not committed and can later be backed out.
#
enum _hp_patch_save_files
_hp_patch_save_files = { "YES", "NO" }
init _hp_patch_save_files = "YES"

```

```
_hp_patch_save_files help_text "Save patched files?"  
However, the description in the previous example is wrong because the patches are not committed.  
This variable construct only controls whether the swinstall option -x  
patch_save_files=true or false command is added to the swinstall command line.  
Later in the same file, the following construct is found:
```

```
(_hp_patch_save_files == "NO")  
{  
    sd_command_line += " -xpatch_save_files=false "  
}  
else  
{  
    # This is actually the default...  
    sd_command_line += " -xpatch_save_files=true "  
}
```

Again, this variable construct is specific to one configuration file. It is not a special variable that controls all software installations during an Ignite-UX installation session.

hp_umask

The `_hp_umask` variable allows you to set the umask to a custom value during a recovery or installation session. If not set, the umask value used by Ignite-UX defaults to 022 (removes write permission from “other” and “group”).

When a recovery or installation session first starts, the umask value is set to the value 022; this is not configurable. The umask defined by `_hp_umask` is only set once the configuration user interface has finished and phase one of the recovery or installation session starts.

After the first reboot and the recovery or installation continues, the umask value from `_hp_umask` is used to set the umask again.

Warning:

HP does not test Ignite-UX with any umask setting other than 022; setting the umask to a more restrictive value may cause installation failures. HP does not consider installation failures caused by setting `_hp_umask` to a value other than the default to be a defect in Ignite-UX.

hp_ht_enable

This variable is used by Ignite-UX to control the state of Hyper-Threading on systems containing processors that support it. You do not need to define this variable; on HP-UX 11.31 the following definition is included in the default cold installation per-release config file (`/opt/ignite/data/Rel_B.11.31/config`):

```
#  
# Add definition of _hp_ht_enable, to provide hook to enable  
# hyperthreading on CPUs that support this for 11.31 and later. Provide an  
# install-time method for configuring the HT firmware enabled/disabled  
# state.  
  
is_ht_capable {  
    enum _hp_ht_enable  
    _hp_ht_enable = {"true", "false"}
```

```

    init _hp_ht_enable="true"
        _hp_ht_enable help_text "Enable firmware hyperthreading"
} else {
    _hp_ht_enable visible_if false
}

```

This allows you to set the state of Hyper-Threading during a cold install using SD depots. Note that setting this variable to "true" on systems that do not support Hyper-Threading has no effect.

`_hp_debug_level`

The `_hp_debug_level` variable allows you to set the debug level without having to interact with the boot loader. This variable only takes effect when placed into the install file system (via `instl_adm`), if used in any other configuration file it will be ignored. The variable should be set to a decimal via between 1 and 5, for example:

```

init _hp_debug_level=5
_hp_debug_level visible_if false

```

You should never set this variable to a string - if you do it will be ignored and have no effect. This variable is known to exist from Ignite-UX version C.7.0 onwards.

Configuration for software to be installed

This section discusses the configuration files required for software, beginning with applications.

Application software depots

The following example system has multiple depots, one for the HP-UX 11i v1 (B.11.11) core operating system depot and the other two for different versions of vPars (HP product number T1335AC):

```

# swlist -l depot
# Initializing...
# Target "box2" has the following depot(s):
/var/opt/ignite/depots/Rel_B.11.11/core
/var/opt/ignite/depots/T1335AC_A.02.02.00
/var/opt/ignite/depots/recovery_cmds
/var/opt/ignite/depots/T1335AC_A.02.03.02

```

You must create some configuration files for the depots with the `make_config` command:

```

# make_config -l 2 -s /var/opt/ignite/depots/T1335AC_A.02.03.02 \
-c /var/opt/ignite/data/Rel_B.11.11/vpars_A.02.03.02.cfg
NOTE: make_config can sometimes take a long time to complete. Please be
patient!
# make_config -l 2 -s /var/opt/ignite/depots/T1335AC_A.02.02.00 \
-c /var/opt/ignite/data/Rel_B.11.11/vpars_A.02.02.00.cfg
NOTE: make_config can sometimes take a long time to complete. Please be
patient!

```

The configuration file produced for version A.02.02.00 of vPars contains the `sw_source` statement for the depot. The depot is available over the network. It is an SD depot, and is located on the server 10.0.0.1 at `/var/opt/ignite/depots/T1335AC_A.02.02.00`:

```

# cat vpars_A.02.02.00.cfg
#####
## Software Sources
#####

sw_source "/var/opt/ignite/depots/T1335AC_A.02.02.00" {
    source_type = "NET"
    sd_server = "10.0.0.1"
    sd_depot_dir = "/var/opt/ignite/depots/T1335AC_A.02.02.00"
    source_format = SD
}

```

Next is patch software. The patch is only applicable to 64-bit systems (so it becomes invisible on 32-bit systems):

```

#####
## Other Software
#####

sw_sel "b_PHSS_28764" {
    (_hp_os_bitness == "64") {
        description = "(PHSS_28764) vPar (A.02.02) monitor cumulative patch"
        sw_source = "/var/opt/ignite/depots/T1335AC_A.02.02.00"
        sw_category = "Uncategorized"
        sd_software_list = "b_PHSS_28764,r=1.0,a=HP-UX_B.11.11_64,v=HP"
        impacts = "/stand" 2288Kb
        visible_if = TRUE
    }
    else {
        visible_if = FALSE
    }
}

```

Although not shown here, `make_bundles` was run to create the bundle `b_PHSS_28764`.

Next is the vPars product:

```

sw_sel "T1335AC" {
    (_hp_os_bitness == "64") {
        description = "HP-UX Virtual Partitions"
        sw_source = "/var/opt/ignite/depots/T1335AC_A.02.02.00"
        sw_category = "OrderedApps"
        sd_software_list = "T1335AC,r=A.02.02.00,a=HP-UX_B.11.11_64,v=HP"
        impacts = "/usr" 2Kb
        impacts = "/etc/rc.config.d" 24Kb
        impacts = "/stand" 1136Kb
        impacts = "/usr/share" 1640Kb
        impacts = "/usr/sbin" 3712Kb
        impacts = "/usr/lib" 2747Kb
        impacts = "/usr/lbin" 1376Kb
        impacts = "/usr/contrib" 15304Kb
        impacts = "/sbin" 1595Kb
        impacts = "/usr/conf" 9578Kb
        impacts = "/usr/include" 640Kb
        impacts = "/usr/newconfig" 5664Kb
        impacts = "/usr/ccs" 10248Kb
        impacts = "/usr/bin" 208Kb
        impacts = "/sbin/init.d" 40Kb
        impacts = "/" 2Kb
    }
}

```

```

        visible_if = TRUE
    }
    else {
        visible_if = FALSE
    }
}

```

The configuration file is not yet 100% complete and requires a manual change that causes problems. Specifically, the patch defined by the `sw_sel`, `b_PHSS_28764` needs to be installed at the same time as T1335AC. Therefore, you need to make a small change to the `sw_sel` clause for T1335AC:

```

sw_sel "T1335AC" {
    (_hp_os_bitness == "64") {
        description = "HP-UX Virtual Partitions"
        sw_source = "/var/opt/ignite/depots/T1335AC_A.02.02.00"
        sw_category = "OrderedApps"
        sd_software_list = "T1335AC,r=A.02.02.00,a=HP-UX_B.11.11_64,v=HP"
        impacts = "/usr" 2Kb
        impacts = "/etc/rc.config.d" 24Kb
        impacts = "/stand" 1136Kb
        impacts = "/usr/share" 1640Kb
        impacts = "/usr/sbin" 3712Kb
        impacts = "/usr/lib" 2747Kb
        impacts = "/usr/lbin" 1376Kb
        impacts = "/usr/contrib" 15304Kb
        impacts = "/sbin" 1595Kb
        impacts = "/usr/conf" 9578Kb
        impacts = "/usr/include" 640Kb
        impacts = "/usr/newconfig" 5664Kb
        impacts = "/usr/ccs" 10248Kb
        impacts = "/usr/bin" 208Kb
        impacts = "/sbin/init.d" 40Kb
        impacts = "/" 2Kb
        visible_if = TRUE
        corequisite="b_PHSS_28764" ←
    }
    else {
        visible_if = FALSE
    }
}

```

Because you want the patch to be installed automatically with vPars, you need to manually add the corequisite configuration line. When located in the same depot, the default `swinstall` behavior of `autoselect_patches=true` automatically selects the patch with the installation of the relevant vPars bundle. The configuration file changes used here are strictly needed only when the patch is not in the same depot.

Now you can look at the vPars software configuration file for the other version of vPars (version A.02.03.02). The `sw_source` definition is to the previous version except the depot is in a different location, so the `sw_source` has a different name as well.

```

# cat vpars_A.02.03.02.cfg
#####
## Software Sources
#####

```

```

sw_source "/var/opt/ignite/depots/T1335AC_A.02.03.02" {
    source_type = "NET"
    sd_server = "10.0.0.1"
    sd_depot_dir = "/var/opt/ignite/depots/T1335AC_A.02.03.02"
    source_format = SD
}

```

Now vPars version A.02.03.02 contains many software products, as depicted in the following:

```

#####
## Other Software
#####

sw_sel "VPARAMGR" {
    description = "Virtual Partition Manager - HP-UX"
    sw_source = "/var/opt/ignite/depots/T1335AC_A.02.03.02"
    sw_category = "OrderedApps"
    sd_software_list = "VPARAMGR,r=B.11.11.01.02,a=HP-UX_B.11.11_32/64,v=HP"
    impacts = "/var" 1Kb
    impacts = "/opt/webadmin" 1544Kb
    impacts = "/opt/vparamgr" 204Kb
    impacts = "/opt" 1Kb
    impacts = "/usr/share" 8Kb
}

sw_sel "VPARDOC" {
    description = "VPAR Documentation Bundle"
    sw_source = "/var/opt/ignite/depots/T1335AC_A.02.03.02"
    sw_category = "HPUXAdditions"
    sd_software_list = "VPARDOC,r=B.11.11.59,a=HP-UX_B.11.11_32/64,v=HP"
    impacts = "/tmp" 8Kb
    impacts = "/cdrom/vParsWINSTALL" 49664Kb
    impacts = "/cdrom/DOCS" 1192Kb
}

sw_sel "T1335AC" {
    (_hp_os_bitness == "64") {
        description = "HP-UX Virtual Partitions"
        sw_source = "/var/opt/ignite/depots/T1335AC_A.02.03.02"
        sw_category = "OrderedApps"
        sd_software_list = "T1335AC,r=A.02.03.02,a=HP-UX_B.11.11_64,v=HP"
        impacts = "/usr" 2Kb
        impacts = "/etc/rc.config.d" 24Kb
        impacts = "/stand" 1176Kb
        impacts = "/usr/share" 1640Kb
        impacts = "/usr/sbin" 3712Kb
        impacts = "/usr/lib" 2747Kb
        impacts = "/usr/lbin" 1376Kb
        impacts = "/usr/contrib" 15304Kb
        impacts = "/sbin" 1603Kb
        impacts = "/usr/conf" 9194Kb
        impacts = "/usr/include" 640Kb
        impacts = "/usr/newconfig" 5664Kb
        impacts = "/usr/ccs" 10248Kb
        impacts = "/usr/bin" 208Kb
        impacts = "/sbin/init.d" 40Kb
        impacts = "/" 2Kb
        visible_if = TRUE
    }
    else {
        visible_if = FALSE
    }
}

```

```

        }
    }

sw_sel "KRMonitor" {
    description = "EMS Kernel Resource Monitor"
    sw_source = "/var/opt/ignite/depots/T1335AC_A.02.03.02"
    sw_category = "HPUXAdditions"
    sd_software_list = "KRMonitor,r=B.11.11.04,a=HP-UX_B.11.11_32/64,v=HP"
    ...
}

```

Using both of these configuration files together (or the configuration file /var/opt/ignite/data/Rel_B.11.11/vpars_A.02.03.02.cfg with anything else) causes conflicts unless manual changes are made.

When make_config finds two or more bundles of the same name in one depot, it places a combination of the revision, architecture, and vendor into the sw_sel clause name (enough of one or more of them to make the sw_sel name unique). When make_config finds that a bundle name is unique in the depot, it does not place the revision, architecture, or vendor in the name of the sw_sel clause.

If you included both of the preceding configurations, the duplicate names (in this case "T1335AC") would cause issues. The information in the first definition of the sw_sel is overruled with the information from the second definition. The reason for this is that sw_sel clauses are stored by name, T1335AC, without regard to the configuration file from which they were derived.

To prevent this issue, you need to distinguish between the products in the first configuration file and the second configuration file. Change vPars A.02.02 to the following:

```

sw_sel "T1335AC,r=A.02.02.00" {
    (_hp_os_bitness == "64") {
    ...
        exrequisite="T1335AC,r=A.02.03.02"
        visible_if = TRUE
    }
    else {
        visible_if = FALSE
    }
}

```

This puts the revision number into the name of the sw_sel. When you have multiple versions of a product available to an installation session (but potentially in different depots), it is a good idea to manually include the version number in the configuration. In the second configuration file the same product needs to become the following:

```

sw_sel "T1335AC,r=A.02.03.02" {
    (_hp_os_bitness == "64") {
    ...
        exrequisite="T1335AC,r=A.02.02.00"
        visible_if = TRUE
    }
    else {
        visible_if = FALSE
    }
}

```

The `exrequisite` statements are added to make sure that you cannot select both products at the same time. In this case, to change from one revision of the vPars bundle to the other, simply select it. The originally selected version is automatically unselected because of the `exrequisite` statement.

Warning:

You should be careful when creating application depots containing software from different sources, such as the application media and other media. If you receive media from HP that has restrictions on what can be installed, it may affect the installation of applications from that depot if you do not honor those restrictions. For example, installing more than one bundle from the Serviceguard Storage Management Suite can cause an installation session to be aborted due to a `checkinstall` script associated with a fileset in one of the bundles. If other bundles are being installed from the same application depot, it may have negative consequences for the system being installed if the installation session is aborted by a `checkinstall` script midway through an installation session.

In this case, keeping the Serviceguard Storage Management Suite in a different depot will prevent it from affecting other software to be installed because of a `checkinstall` script aborting an installation session. Since the software would be in a different depot it would only abort the software installation session from that depot.

When you encounter this issue with Serviceguard Storage Management Suite, the `checkinstall` script aborting the software installation session will print an error containing the following message: "The bundles in HP Storage Management Suite are mutually exclusive. Cannot select more than one bundle for the installation.".

You can resolve this problem by adding `exquisite` information into the configuration file describing the bundles in the Serviceguard Storage Management Suite, or by not selecting more than one bundle from the Serviceguard Storage Management Suite for installation.

Core operating system depot configuration

The following example shows a configuration file for a core operating system depot:

```
# cat core.cfg
#####
## Software Sources
#####

sw_source "core" {
    description = "HP-UX Core Software"
    source_format = SD
    sd_server = "10.0.0.1"
    sd_depot_dir = "/var/opt/ignite/depots/Rel_B.11.11/core"
    source_type = "NET"
    load_order = 0
}
```

The `sw_source` definition for the core operating system depot contains the information needed to define an SD depot. The `load_order` of 0 (zero) is important for a core operating system depot. The load order determines the order in which Ignite-UX loads the software. Other software products can be loaded at load order 0 except when the core operating system is an archive rather than a depot. When loading a core operating system archive, you can have one and only one `sw_sel` installed at `load_order` 0.

If you place any other software in the core operating system depot, `swinstall` resolves dependencies from the source depot that it is installing from before attempting to satisfy dependencies from the target. If you were to use "`-x reinstall=true`" on a `swinstall` command to install a product that had dependencies on the core operating system, you could also reinstall parts of the core operating system along with the product. The easiest way to prevent this from happening is to avoid placing additional software products in the core operating system depot or to never load software from that depot outside of Ignite-UX (you may know to never use `-x reinstall=true` but you may want to use that depot in the future with that option).

Next are the `sw_sel` definitions of the 32-bit and 64-bit Base HP-UX bundles. Effectively the same definitions are generated for a core operating system depot. Never include the definitions of multiple core operating system depots within one clause in the `/var/opt/ignite/INDEX` file because they start to override each other and cause issues with the install.

In the preceding example, the impacts statements for the bundles have been removed to make the clauses easier to read. Now, look at the 64-bit `sw_sel` clause. The name of the `sw_sel` clause is `HPUXBase64`. Its description is set to `HP-UX 64bit Base OS`. It has a `sw_source` set to `core` that was defined earlier. Always try to place the `sw_source` definition and the software products that reference it in the same configuration file to prevent dependencies between configuration files.

Next is the software category. For more information on categories and other software attributes (including special handling for some attributes), see the section on "Categories and other Ignite-UX software attributes" because some categories are handled specially. Some categories are not shown to the user as being available for selection.

The `sd_software_list` next defines the software that this `sw_sel` references. In this case, it is the bundle `HPUXBase64` with the attributes `revision ("r=") B.11.11`, `architecture ("a=") HP-UX_B.11.11`, and `vendor ("v=") HP`.

Next `exerequisite = sw_category` makes sure that no other software with the same software category can be selected at the same time as this bundle. This automatically unselects any other `sw_sel` with the same software category set. This is to prevent both the 64- and 32-bit bundles from being selected simultaneously. That is the end of the software definition.

The following example tests to see if `sw_sel` `HPUXBase64` is selected (the surrounding parenthesis make this a test) and the variable `_hp_os_bitness` is set to 64, if the 64-bit Base operating system has been selected.

```
#####
## HPUX Base OS
#####

sw_sel "HPUXBase64" {
    description = "HP-UX 64-bit Base OS"
    sw_source = "core"
```

```

    sw_category = "HPUXBaseOS"
    sd_software_list = "HPUXBase64,r=B.11.11,a=HP-UX_B.11.11_64,v=HP"
...
    exrequisite = sw_category
}
(sw_sel "HPUXBase64") {
    _hp_os_bitness = "64"
}

```

The following example is the definition of the 32-bit Base HP-UX bundle. It is very similar to the 64-bit example. It also has the `exrequisite = sw_category` attribute that prevents the selection of the 32-bit and 64-bit Base HP-UX bundle at the same time.

```

sw_sel "HPUXBase32" {
    description = "HP-UX 32-bit Base OS"
    sw_source = "core"
    sw_category = "HPUXBaseOS"
    sd_software_list = "HPUXBase32,r=B.11.11,a=HP-UX_B.11.11_32,v=HP"
...
    exrequisite = sw_category
}
(sw_sel "HPUXBase32") {
    _hp_os_bitness = "32"
}

```

Next are the operating environment (OE) definitions for the Base HP-UX 11i v1 OEs. These definitions determine whether the system can run a 64-bit operating system or not. Ignoring the contents of the `sw_sel` clauses, the `sw_sel "OE90BaseOS64"` is automatically selected if the system can run 64-bit HP-UX. The `visible_if` prevents the Operating Environment (OE) from being seen if the system cannot run a 64-bit version of HP-UX).

```

init sw_sel "OE90BaseOS64" {
    description = "HP-UX 11i Base OS-64bit"
    sw_source = "core"
    sw_category = "HPUXEnvironments"
    corequisite = "HPUXBase64"
    visible_if = can_run_64bit
} = (can_run_64bit)

```

The 32-bit OE is only visible if the system can run a 32-bit version of HP-UX; some older systems can run either. The 32-bit OE is automatically selected if the system cannot run 64-bit HP-UX, or if the system can run 32-bit HP-UX and the 64-bit Base OE is not currently selected (so you must end up with one or the other selected).

```

init sw_sel "OE91BaseOS32" {
    description = "HP-UX 11i Base OS-32bit"
    sw_source = "core"
    sw_category = "HPUXEnvironments"
    corequisite = "HPUXBase32"
    visible_if = can_run_32bit
} = (! can_run_64bit) | (can_run_32bit) & !(sw_sel "OE90BaseOS64")

```

Now you have some code selecting the model it runs on and choosing a language (the `sw_sel` clause for "English" is not shown in this section but it does exist in the configuration file).

The items of most interest here are the regular expressions shown since they need improvement. The next section discusses what the regular expressions do.

The regular expression "9000/8.*" means match for the pattern 9000/8 followed by zero or more characters. The ". ." means any character, and the following "*" changes that to mean zero (0) or more of that character. The regular expression "9000/8" is more precise since it just matches any string containing "9000/8"; the ". *" is superfluous. Since the expression is not anchored,⁵⁹ nothing else is needed in the regular expression.

The same thing applies to the "ia64 .* server .*", it could be more simply written as "ia64.*server". Since it is an unanchored expression, you do not have to worry about what comes before or after the string. The changed expression matches ia64 followed by zero or more of any character followed by server. Again, because it is unanchored you do not care about what comes after the server part of the expression so there is no need for an explicit ". *" at the end of the expression.

For PA-RISC servers and Itanium®-based servers, when the Base HP-UX OE is selected, the sw_sel "English" is automatically selected.

```
HARDWARE_MODEL ~ "9000/8.*" | MODEL ~ "ia64 .* server .*" {  
    (sw_sel "OE91BaseOS32") | (sw_sel "OE90BaseOS64") {  
        init sw_sel "English" = TRUE  
    }  
} # end of hardware-specific section  
else {
```

The same thing discussed previously about regular expressions also applies to the workstation regular expressions that follow.

If you have a workstation, initialize the sw_sel clause, "Global", to be TRUE.

```
HARDWARE_MODEL ~ "9000/7.*" | MODEL ~ "ia64 .* workstation .*" {  
    (sw_sel "OE91BaseOS32") | (sw_sel "OE90BaseOS64") {  
        init sw_sel "Global" = TRUE  
    }  
} # end of hardware-specific section
```

Otherwise the sw_sel "English" is set to TRUE.

```
else {  
    (sw_sel "OE91BaseOS32") | (sw_sel "OE90BaseOS64") {  
        init sw_sel "English" = TRUE  
    }  
}
```

⁵⁹ Regular Expression anchoring involves a "^" to indicate the start of the line to match against, and "\$" to indicate the end of the line. For example 'print "baaaa" |grep "^aaaa"' does not match anything because the text being passed through grep starts with a "b". The "^aaaa" says to only match lines starting with "aaaa" at the very start of the line, so the expression is anchored to the start of the line. The use of "\$" similarly anchors the expression to the end of the line.

The following example shows the OE definitions for all of the OEs on the B.11.11 OE media. The definitions are similar to the other `sw_sel` definitions already discussed. The impacts statements have been removed, but the impacts are different for 32-bit and 64-bit installations. Subsequent OE definitions and the language clauses that follow them have been removed from this discussion of the core configuration file.

```
#####
## Operating Environments
#####

sw_sel "HPUX11i-OE-MC" {
    description = "HP-UX Mission Critical Operating Environment Component"
    sw_source = "core"
    sw_category = "OpEnvironments"
    sd_software_list= "HPUX11i-OE-MC,r=B.11.11.0212,a=HP-
UX_B.11.11_32/64,v=HP"
        (_hp_os_bitness == "32") {
...
}
        (_hp_os_bitness == "64") {
...
}
}

sw_sel "HPUX11i-OE-Ent" {
    description = "HP-UX Enterprise Operating Environment Component"
    sw_source = "core"
    sw_category = "OpEnvironments"
    sd_software_list = "HPUX11i-OE-Ent,r=B.11.11.0212,a=HP-
UX_B.11.11_32/64,v=HP"
        (_hp_os_bitness == "32") {
...
}
        (_hp_os_bitness == "64") {
...
}
}

sw_sel "HPUX11i-OE" {
    description = "HP-UX 11i Operating Environment Component"
    sw_source = "core"
    sw_category = "OpEnvironments"
    sd_software_list = "HPUX11i-OE,r=B.11.11.0212,a=HP-UX_B.11.11_32/64,v=HP"
        (_hp_os_bitness == "32") {
...
}
        (_hp_os_bitness == "64") {
...
}
}
```

Next is a discussion on the definition of some other software that can be installed.

```
#####
## Other Software
#####
```

The following definition automatically selects Perl for installation from the OE media. The `sw_sel` configuration is initialized to TRUE. However, you can deselect the software from installation on the **Software** tab in the Ignite-UX GUI.

```
init sw_sel "perl" {
    description = "Perl Programming Language"
    sw_source = "core"
    sw_category = "OrderedApps"
    sd_software_list = "perl,r=D.5.8.0.B,a=HP-UX_B.11.11_32/64,v=HP"
    impacts = "/opt" 92854Kb
} = TRUE
```

The following software is the Online Diagnostics bundle. This is an example of always-installed software. It is automatically loaded with any `HPUXBaseOS` bundle. You cannot deselect this software bundle from being installed.

```
sw_sel "OnlineDiag" {
    description = "HPUX 11.11 Support Tools Bundle, Jun 2004"
    sw_source = "core"
    sw_category = "HPUXAdditions"
    sd_software_list="OnlineDiag,r=B.11.11.14.15,a=HP-UX_B.11.11_32/64,v=HP"
    impacts = "/usr" 66792Kb
    impacts = "/opt" 943Kb
    impacts = "/var" 12395Kb
    impacts = "/sbin" 80Kb
    impacts = "/etc" 5342Kb
    impacts = "/dev" 1Kb
    load_with_any ~ "HPUXBaseOS" . "*"
}
```

Next is a list of available keyboard layouts. If the system has some hardware controlled by the `ps2` driver, then the keyboards controlled by the `has_ps2` test are added to the list of available keyboards. Not Applicable is always a valid choice even if the system does have a keyboard attached.

The follow-on test, `has_usb`, adds more keyboard choices if the system being installed has USB hardware attached to the system. The default is set at the very end to have the keyboard set to "Not Applicable". HP workstations and servers do not have `ps2` and USB keyboards so you can have one or the other but not both types of keyboards connected to a workstation or server.

The available keyboards are shown in the Ignite-UX GUI on the **Basic** tab using the **Keyboards...** button. When you select that button, a list of keyboards available for selection appears and is generated from this configuration (most servers only show Not Applicable).

```
#####
## Keyboards
#####

_hp_keyboard = { "Not_Applicable" }

has_ps2 {
    _hp_keyboard += {
        "PS2_DIN_Arabic",
        "PS2_DIN_Belgian",
        "PS2_DIN_Belgian_Euro",
        "PS2_DIN_Bulgarian",
```

```

    "PS2_DIN_Canada_TBITS-3",
    "PS2_DIN_Canadian_French",
    "PS2_DIN_Czech",
    "PS2_DIN_Czech_Euro",
    "PS2_DIN_Danish",
    "PS2_DIN_Danish_Euro",
    "PS2_DIN_Euro_Spanish",
    "PS2_DIN_Euro_Spanish_Euro",
    "PS2_DIN_French",
    "PS2_DIN_French_Euro",
    ...
    "PS2_DIN_T_Chinese",
    "PS2_DIN_Turkish",
    "PS2_DIN_UK_English",
    "PS2_DIN_UK_English_Euro",
    "PS2_DIN_US_English",
    "PS2_DIN_US_English_Euro"
}
}

has_usb {
    _hp_keyboard += {
        "USB_PS2_DIN_Belgian",
        "USB_PS2_DIN_Belgian_Euro",
        "USB_PS2_DIN_Danish",
        "USB_PS2_DIN_Danish_Euro",
        "USB_PS2_DIN_Euro_Spanish",
        "USB_PS2_DIN_Euro_Spanish_Euro",
        ...
        "USB_PS2_DIN_Swiss_French2_Euro",
        "USB_PS2_DIN_Swiss_German2",
        "USB_PS2_DIN_Swiss_German2_Euro",
        "USB_PS2_DIN_T_Chinese",
        "USB_PS2_DIN_UK_English",
        "USB_PS2_DIN_UK_English_Euro",
        "USB_PS2_DIN_US_English",
        "USB_PS2_DIN_US_English_Euro"
    }
}
init _hp_keyboard = "Not_Applicable"

```

Impacts statements

Previous sections have explained impacts statements and the circumstances in which they may not accurately represent space usage. This section describes how to structure your impacts, and what level to use with your impacts - such decisions may not be straightforward.

HP-UX installation media is discussed first. The impacts statements on the installation media are all at level one. Consider what limitations this may cause, that is, what problems the following set of (partial) impacts could cause:

```

impacts = "/usr" 687920Kb
impacts = "/opt" 95300Kb
impacts = "/var" 23395Kb

```

Most people encounter no problems. The previous statements assume that you did not create any extra file systems apart from the default file systems that Ignite-UX would create. The impacts do not cause a problem for most people, because Ignite-UX only creates file systems at the top level by default, for example: /home, /var, /opt, and /usr.

What happens if you want to create the following extra file systems (for an application)?

```
/opt/application  
/opt/application/static  
/var/opt/application  
/var/opt/application/data  
/var/opt/application/tmp
```

Ignite-UX has no information about what data might be stored in these file systems because it has no impacts statements that affect them. Suppose that the impacts are only one level deep and everything is being tracked against /var and /opt for the application:

```
impacts = "/opt"    1048576Kb  
impacts = "/var"   3673088Kb
```

If you create those extra file systems, you might easily misjudge their size. Because the impacts statements have tracked all of the space against /var and /opt, Ignite-UX cannot know how much space is needed under those file systems so it does not alert you. If your situation is like this, it would make sense to set up your own Ignite-UX server with the correct impacts keywords rather than attempting to use installation media.

Assume that the file systems needed for this application are those in Table 5:

Table 5

Mount point	File System Size (Mb)	Space Used (Mb)
/opt/application/	1024Mb	768Mb
/opt/application/static	512Mb	384Mb
/var/opt/application	2048Mb	1512Mb
/var/opt/application/data	16384Mb	2011Mb
/var/opt/application/tmp	16384Mb	64Mb

Impacts statements for the sw_sel clause associated with this application look like the following:

```
impacts = "/opt/application"      786432Kb  
impacts = "/opt/application/static" 393216Kb  
impacts = "/var/opt/application"    1548288Kb  
impacts = "/var/opt/application/data" 2059264Kb  
impacts = "/var/opt/application/tmp" 65536Kb
```

The impacts statements place the required impacts at the places where the application (in the environment in which it is installed) has mount points. With this information, Ignite-UX is aware of the impacts associated with each file system.

Creating impacts in this manner (for our optional application with only five impacts statements at up to four levels deep) is a manual job for two reasons:

First, Ignite-UX normally creates impacts at every directory up to the level that you specify. If you have an archive and you use the archive_impact command to generate the impacts with the option -l 4, the archive_impact command lists the space used in every directory up to the fourth level deep. With a large directory structure, this can generate thousands of impacts

statements. When Ignite-UX processes a large amount of impacts statements it can slow down to a significant degree.

Combining those large numbers of impacts statements down into a few that match planned mount points is a manual task (and, unfortunately, if the software ever changes you need to do it all again).

Impact Hints

- Keep impact levels to the minimum required to minimize their affect Ignite-UX performance.
- If you need detailed impacts, you can always summarize impacts into the most appropriate form (impacts based upon expected mount points) without having to have thousands of impacts statements.

Second, applications may require extra space. In the preceding application, what would you do if the application only installed a few MB into the file system /var/opt/application/data but then when configured needed the rest of the space?

You have to plan to have the space there when the application is installed. Sometimes you must manually change impacts statements to increase the impacts on a directory or directory structure because of space requirements on a file system that are not reflected in the size of the software to be installed.

An example of this might be database table spaces that may be required when an application creates a database at initial installation. If this occurs, manually changing impacts statements is the only way to account for the extra space required. Manually increasing the impacts in this case prevents anyone from decreasing the file system size to the point that failures occur.

Overstated SD impacts

There are conditions that can generate overstated impacts for SD based installs:

1. Overlapping bundles

You should be careful with bundles that contain overlapping contents. Ignite-UX only tracks impacts at the bundle level and cannot detect if bundles contain overlapping contents. If you select overlapping bundles (for example all of the Ignite-UX bundles) disk impacts may become overstated. This is, however, the way Ignite-UX was designed to work and is not considered a defect.

2. Impacts from patches

Most patches do not deliver new content. Apart from the patched files that may be saved, their impact on most file systems is only incremental. Prior to Ignite-UX version C.6.3 there was no way to tell Ignite-UX to ignore some of the impacts generated by patches. As of Ignite-UX version C.6.3 you can use the -d option with `make_config` to discount the space impacts by the percentage given (1-75 are the allowed values.) This will not discount the space that is assumed for files saved by the patch in /var/adm/sw/save, and you must test to ensure that the patches you are installing do not deliver large amounts of new files (as this may fill a file system during installation.)

Categories and other Ignite-UX software attributes

The `inst1_adm(4)` manpage states the following about `sw_category`:

```
sw_category cat-tag-string { description = string }
```

The `sw_category` definition provides a grouping mechanism for `sw_sel`'s to reference. The user interface uses the `sw_category` to help the user browse the software more easily. There are six values of `cat-tag-string` that the user interface treats special. Software that are in these groups are represented in special locations and by special methods in the user interface. These values are: "HPUXBaseOS", "HPUXEnvironments", "Languages", "LanguagesUI", "OpEnvironments", and "UserLicenses". The only attribute associated with a `sw_category` is a description.

The following is what is special about these categories:

- OpEnvironments— OEs defined in a configuration are available for selection in the Ignite-UX GUI on the **Basic** tab in the **Environments:** list. They are not selectable using the **Software** tab.
- HPUXBaseOS— used for the HPUXBase32 and HPUXBase64 bundles that are included in HP-UX B.11.11. In HP-UX B.11.23, there is a single HPUXBaseOS bundle that is 64-bit only (the tag and category have the same name). These bundles also control the setting of `_hp_os_bitness` in a core configuration file.
- Languages— In HP-UX B.11.11 and later, these were pseudo-bundles that when selected cause LanguagesUI category bundles to get installed.
- LanguagesUI— These bundles are not shown at all using the user interface. Instead, they are automatically selected based upon the bundles with the `sw_category` tag set to Languages. Take, for example, Swedish from the B.11.11 core configuration file:

```
init sw_sel "Swedish"{
    description="Swedish CDE Environment"
    sw_source="core"
    sw_category="Languages"
} = FALSE

init sw_sel "CDE-Swedish"{
    description="Swedish CDE Environment"
    sw_source="core"
    sw_category="LanguagesUI"
    sd_software_list="CDE-Swedish,r=B.11.11,a=HP-
UX_B.11.11_32/64,v=HP"
    load_with_all="Swedish"
    impacts="/usr" 70992K
    impacts="/opt" 9098K
    impacts="/sbin" 54K
    impacts="/" 55K
    visible_if=FALSE
    locale={"sv_SE.iso88591:Swedish","sv_SE.utf8:Swedish",
    "SET_NULL_LOCALE:English","ASK_AT_FIRST_BOOT","C:English"}
} = FALSE
```

When Swedish is selected, the bundle CDE-Swedish is automatically selected for installation because it has the configuration `load_with_any` set to "Swedish".

- UserLicenses— This category is no longer very relevant with HP-UX B.11.11 and later releases of HP-UX. An unlimited user license is provided as standard.

Defining a custom software configuration

The section of `inst1_adm(4)` that defines software configuration is quite long but very useful - only sections of it are discussed here. If you require more information, review the section titled "Software Source and Selections".

Do not attempt to write configuration files for SD depots. Instead use the `make_config` command because it correctly defines all configuration data needed. You can then customize the configuration file produced by `make_config`.

The general structure for specifying software is as follows:

```
sw_source src-tag-string
{
source-attributes...
}
sw_category cat-tag-string
{
    description = string
}
sw_sel sel-tag-string
{
    sw_source src-tag-string
    sw_category cat-tag-string
    selection-attributes...
}
init sw_sel sel-tag-string=Boolean
```

Looking at a network recovery sw_source and sw_sel

To see how software is defined, look at the way that `make_net_recovery` builds a definition of an archive:

```
sw_source "core archive"{
    description = "Recovery Archive"

    load_order = 0
    source_format = archive
    change_media = FALSE

    post_load_script = "/opt/ignite/data/scripts/os_arch_post_l"
    post_config_script = "/opt/ignite/data/scripts/os_arch_post_c"

    # if nfs_source is used, be sure to export the source.
    (source_type == "NET") {
        nfs_source = "10.0.0.1:/var/opt/ignite/recovery/archives/systema"
    }

    # ftp_source and remsh_source are alternate ways to copy
    # the archive. An example of the ftp syntax:
    # ftp_source = "anonymous@15.1.54.123:iux"
    # remsh_source = "user@15.1.54.123"
}
```

In the above example, there is a `sw_source` definition called `core archive` that has a description of Recovery Archive. The archive is loaded at load order 0 (this is a requirement

for a core operating system archive if it is a recovery archive or a golden image; you can only have one archive at load order 0). The attribute `change_media` is set to `FALSE` because you do not want to be prompted to change the media. It does not make sense to change media when the source is over the network using NFS.

You then test to see if the `source_type` is "NET" and if it defines the `nfs_source` to point to the archive.

In the comments regarding other network access methods, the `source_type` variable defaults to the type of media the system was booted from (although this automatically becomes "NET" when an Ignite-UX server is being used, regardless of whether the system was booted from disk or tape).

Next, categories are defined because you are going to have to provide one `sw_sel` in the `Languages` category and another in the `HPUXEnvironments` category:

```
#####
## Software Categories
#####

sw_category "Languages" {
    description = "Languages"
}

sw_category "HPUXEnvironments" {
    description = "HP-UX Operating Environments"
}
```

You can now define a `sw_sel`, called "golden image1", to hold the definition of the Recovery Archive. The `init` command is used to initialize the `sw_sel` and at the end of the definition it is set to `TRUE`. This allows the Ignite-UX GUI to change this value (although changing the value in this case does not make sense).

You then define the `sw_source`; the `sw_source` appears earlier in this configuration file. The `sw_category` for the `sw_sel` is set to `HPUXEnvironments`. This so that the description Recovery Archive appears in the **Environments**: selection list in the Ignite-UX GUI so it can be easily selected. The `archive_type` is then defined, in this case it is a gzipped `tar` file (refer to `inst1_adm(4)` for other valid types).

Next, the archive path relative to the NFS mount point in the `sw_source` is given. The recovery archive is only visible if the system can run 64-bit HP-UX. If the `sw_sel` "golden image1" is selected, set `_hp_os_bitness` to 64, and then set the variable so it cannot be changed using the **Additional** button on the **Basic** tab of the Ignite-UX GUI.

```
#####
## Operating Environments
##
#####

#####
## 64-bit OS archives
#####

init sw_sel "golden image1" {
    description = "Recovery Archive"

    sw_source = "core archive"
    sw_category = "HPUXEnvironments"
```

```

archive_type = gzip tar
# For NFS, the path to the archive is relative to the mount point
# specified in the sw_source:

(source_type == "NET") {
    nfs_source = "10.0.0.1:/var/opt/ignite/images"
}

...
visible_if  = can_run_64bit
} = TRUE

(sw_sel "golden image1") {
    _hp_os_bitness = "64"
}
_hp_os_bitness visible_if false

```

Next, a language is set up so there is a choice of no locale English, or C locale and English. Your choice depends on the languages that would be available depend on the system at the time the recovery archive was created:

```

sw_source "no select" {
    source_format = cmd
}

init sw_sel "English" {
    description = "English Language Environment"
    sw_source = "no select"
    sw_category = "Languages"
    locale = { "SET_NULL_LOCALE:English", "C:English" }
}
} = TRUE

```

Using a sw_sel to run commands instead of installing software

The following configuration acts in a similar way to the language selection earlier. The source_format of cmd in the sw_source definitions does not cause any software to be loaded. The load_order statement forces the command to run after everything else (in case someone does anything at the default load order of 5). The source_format of cmd gives you a great deal of flexibility in running scripts.

You have to define any sw_category that you want to use. Before you can use a software category, you must define a sw_category. This is not true for the sw_category "Uncategorized"; this sw_category is applied to any software that does not explicitly define a sw_category. If all software is given, the **Uncategorized** sw_category will not appear on the **Software** tab in the Ignite-UX GUI. You should not explicitly assign a sw_category **Uncategorized** to any software. In the following example, the sw_category of "SiteSpecific" is defined.

The sw_sel, however, defines a post_load_cmd that runs:

```

sw_category "SiteSpecific" {
    description = "Local Site Commands"
}

sw_source "site commands" {
    source_format = cmd
    load_order = 10

```

```

}

init sw_sel "Run site commands" {
    description = "Run site commands after OS loaded"
    sw_source = "site commands"
    sw_category = "SiteSpecific"
    post_load_script="/var/opt/ignite/scripts/run_after_load"
} = TRUE

```

The difference between this method and placing a `post_load_cmd` into the core operating system load `sw_sel` is that you have the opportunity to prevent the commands from running the Ignite-UX GUI by deselecting the software (even though in this case there really is no "software").

However, be aware that any script you run must be either accessible using `tftp`, if this is a network installation, or located in the `SCRIPTS` file, if the installation is happening from media (in the LIF produced by `make_medialif`).

If you want to run a script that has already been installed with other software, use a `cmd` hook instead of a script hook:

```

init sw_sel "Run site commands" {
    description = "Run site commands after OS loaded"
    sw_source = "site commands"
    sw_category = "SiteSpecific"
    post_load_cmd="/opt/site/bin/run_after_load"
} = TRUE

```

You must ensure that the load order associated with the `sw_source` prevents the "cmd" `sw_source` from running until after any software that it depends on has been loaded (if it has dependencies).

Using a `sw_sel` to apply kernel parameters

You can use a `sw_sel` clause to control the application of kernel parameters to a system, as well. The Ignite-UX configuration items that can be used to change kernel parameters (from the `instl_adm(4)` manpage) are as follows:

```

mod_kernel = cplx-string
mod_kernel += cplx-string

```

This keyword can be used to add drivers or tunable parameters to the system's kernel that is built during the Ignite-UX process. The format of `cplx-string` may be either "driver" or "tunable value".

The largest of any tunable parameter that exists in either the `/stand/system` file or that is specified will be used. Beginning with the 11.23 release, if a tunable is not found in `/stand/system`, then it will also be compared with the default value as reported by `kctune`. Ignite-UX does not compare the values of formulas to discrete numerical values, or two formulas, or hexadecimal values, in order to determine which is larger. It will issue a note message stating that it will assume the last `mod_kernel` keyword parsed is larger (regardless of whether it is a formula or discrete numerical value) and will apply it. There is no bounds checking done on tunable parameters.

Decimal values should be in the range from 0 to 2147483647 (2³¹-1). The shell that is used to evaluate these values uses

signed 32-bit arithmetic for decimal values, and it has problems when values exceed this limit. If larger values are needed, convert them to hexadecimal since the shell does not evaluate them.

Note that any allowable syntax supported by the config command can be used in mod_kernel keywords. This includes formulas. For example, it is possible to use:

```
mod_kernel += "ninode (20+8*MAXUSERS+NPTY+" + ${"%d" _inc} + ")"
```

assuming _inc has been initialized to some value in the configuration file.

The = operator will override any prior global mod_kernel assignments. The += operator will add to any prior settings. Notice that mod_kernel statements may also be associated with a sw_sel definition. The = operator does not have any effect on mod_kernel assignments made in a sw_sel.

Beginning with the 11.23 release, the format for cplx-string arguments was enhanced. There are two new formats specifically that are understood (in addition to the older formats). They are:

```
mod_kernel += "tunable name value"  
mod_kernel += "module name [state]"
```

The former is the same as how tunables were handled before only the keyword tunable is put at the beginning. The latter is how kernel modules are added to the system. The optional state argument has one of four values: static, auto, loaded and best. No checking is performed on this value. See kcmodule(1M) for more information.

Actions involving mod_kernel are done before those for both set_kernel and rm_kernel.

```
set_kernel = cplx-string  
set_kernel += cplx-string
```

This keyword is the same for drivers as the mod_kernel keyword in that it will add the driver to the kernel. For tunables it differs from mod_kernel in that it will set the tunable to the arbitrary value as defined by cplx-string. No comparisons or checks are performed with prior or default values. Actions involving set_kernel are done after those for mod_kernel but before those for rm_kernel.

```
rm_kernel = cplx-string  
rm_kernel += cplx-string
```

This keyword will remove the driver or tunable from the /stand/system file. For drivers, this implies the driver will be removed from the kernel. For tunables, this implies that the tunable will revert back to its default value. Actions involving rm_kernel are done after those for both mod_kernel and set_kernel.

The following example shows that a new category is created, called KernelConfig, which defines sets of kernel parameters that can be applied to a system. Ignite-UX does not do any parameter validation. That is followed by the "no-op" sw_source clause:

```
sw_category "KernelConfig" {
```

```

        description = "Kernel Configuration sets"
    }

sw_source "Kernel Configuration" {
    source_format = cmd
    load_order = 11
}

```

The first `sw_sel` defines a set of kernel parameters that can be applied to a system when Java is installed. The only new items here are the `exrequisite` statement and `set_kernel`. The `set_kernel` statements enable you to easily set the formula as the value to be assigned to a kernel tunable since Ignite-UX does not attempt to check the value being assigned to the kernel tunable. It is simply set to that value.

The `exrequisite` ensures that if this `sw_sel` is selected, none of the `exquisites` are also selected (they are unselected when this `sw_sel` is selected).

```

sw_sel "Kernel Config for Java" {
    description = "Tunable settings required for Java 1.4.2"
    sw_source = "site commands"
    sw_category = "KernelConfig"
    exrequisite = "Kernel Config for Servers"
    exrequisite += "Kernel Config for Workstations"
    set_kernel = "maxusers 400"
    set_kernel += "max_thread_proc maxusers*3"
    set_kernel += "maxfiles 2048"
    set_kernel += "maxfiles_lim 2048"
    set_kernel += "ncallout 2*(((nproc*7)/4)+16)*2"
    set_kernel += "nkthread 2*max_thread_proc"
    set_kernel += "nfile (2*nproc)+1000"
    set_kernel += "nproc ((maxusers*5)+64)"
}

```

The next two `sw_sels` reinforce the use of `exquisites`. Because all the `sw_sel` clauses explicitly state that if one is selected the other two cannot be, only one of them can be selected at any time.

```

sw_sel "Kernel Config for Servers" {
    description = "Tunable settings required for Servers"
    sw_source = "site commands"
    sw_category = "KernelConfig"
    exrequisite = "Kernel Config for Java"
    exrequisite += "Kernel Config for Workstations"
    set_kernel = "maxdsiz 0x4000000"
    set_kernel += "maxusers 128"
}
sw_sel "Kernel Config for Workstations" {
    description = "Tunable settings required for Workstations"
    sw_source = "site commands"
    sw_category = "KernelConfig"
    exrequisite = "Kernel Config for Java"
    exrequisite += "Kernel Config for Servers"
    set_kernel = "maxdsiz 0x4000000"
    set_kernel += "maxusers 96"
}

```

In the following example, you have a configuration that will automatically select which `sw_sel` containing the kernel configuration changes to apply. If the `sw_sel "B9789AA, r=1.3.1.09.08"` is selected, then the `sw_sel "Kernel Config for Java"` will be

automatically selected; otherwise, if this is a server selected, then "Kernel Config for Servers" will be selected. If neither of these are selected, the "Kernel Config for Workstations" will be selected.

```
( sw_sel "B9789AA,r=1.3.1.09.08" ) {
    init sw_sel "Kernel Config for Java" = TRUE
} else {
    ( HARDWARE_MODEL ~ "9000/8" | MODEL ~ "ia64.*server" ) {
        init sw_sel "Kernel Config for Servers" = TRUE
    } else {
        init sw_sel "Kernel Config for Workstations" = TRUE
    }
}
```

It is important to understand that, starting with HP-UX B.11.23, the format of the cplx-string arguments to mod_kernel, set_kernel, and rm_kernel have changed and are incompatible with the previous format. The following information is from the *Ignite-UX Reference*:

"Beginning with the B.11.23 release, the format for cplx-string arguments is enhanced. There are additional formats:

```
mod_kernel += "tunable name value"
mod_kernel += "module name [state]"
```

The first performs the same as it did previously with the addition of the keyword tunable. The second describes how kernel modules are added to the client. The optional state argument has one of four values: static, auto, loaded, and best. No syntax checking is performed for this value. See *kcmodule(1M)* for more information."

So to tune maxdsiz at HP-UX B.11.11, you would need to set it as follows:

```
set_kernel = "maxdsiz 0x4000000"
```

At HP-UX B.11.23, you should use:

```
set_kernel = "tunable maxdsiz 0x4000000"
```

Under a lot of circumstances configuration files are not shared between HP-UX revisions, so this does not create an immediate problem. However, if you do share a configuration file that tunes the kernel between HP-UX releases, you should probably do it as follows. This example assumes that you wish different tunables for HP-UX B.11.11, B.11.23, and B.11.31:

```
...
( release == "B.11.11" ) {
    set_kernel = "maxdsiz 0x4000000"
    set_kernel += "maxusers 64"
} else {
    ( release == "B.11.23" ) {
        set_kernel = "maxdsiz 0x4000000"
    } else {
        ( release == "B.11.31" ) {
            set_kernel = "tunable maxdsiz 0x4000000"
        } else {
            error += "release " + ${release} + " not supported by this
configuration."
```

```
        }
    }
...
}
```

Forcing software (sw_sel) clauses to be installed

Be careful when marking software automatically for installation using a configuration file. If you mark it in the following way (set it directly equal to TRUE), the software is marked as required in the Ignite-UX GUI and you cannot manually deselect the software:

```
sw_sel "Kernel Config" {
    description = "Tunable settings"
    sw_source = "site commands"
    sw_category = "KernelConfig"
    set_kernel = "maxdsiz 0x4000000"
    set_kernel += "maxusers 96"
} =TRUE
```

Compare the previous software clause to the following. Because it starts with init you can change its state using the Ignite-UX GUI during installation, although by default it is marked for installation.

```
init sw_sel "Kernel Config" {
    description = "Tunable settings"
    sw_source = "site commands"
    sw_category = "KernelConfig"
    set_kernel = "maxdsiz 0x4000000"
    set_kernel += "maxusers 96"
} =TRUE
```

You can force software installation in other ways, such as with the keywords load_with_any and load_with_all. However, you should not combine usage of the init keyword with these methods.

Keep this in mind when defining software because you do not want to force software onto a system unintentionally.

Automating dependencies in software

This section explains the ways of enforcing dependencies between software with Ignite-UX.

Corequisites are used to enforce dependencies between software that you want to load at the same time:

```
corequisite [+]= tag-string
```

Indicates that the sw_sel referred to by tag-string should be loaded along with this sw_sel. Multiple corequisites may be listed using one corequisite statement per, and by using the += operator.

In the following example, there are the sw_sel "Product XYZ" and "Product XYZ - 4GL". The 4GL product requires the other product to be installed as well, so introduced here is a corequisite for the 4GL to install automatically "Product XYZ" when it is selected for installation.

```

sw_sel "Product XYZ" {
    description = "Product XYZ base runtime"
    sw_source = "Applications for 11i Version 1"
    sw_category = "SiteApps"
    sd_software_list = "ProductXYZ,r=4.0,v=XYZCorp"
}

sw_sel "Product XYZ - 4GL" {
    description = "4GL option for Product XYZ"
    sw_source = "Applications for 11i Version 1"
    sw_category = "SiteApps"
    corequisite = "Product XYZ"
    sd_software_list = "ProductXYZ4GL,r=4.0,v=XYZCorp"
}

```

See "Using a `sw_sel` to apply kernel parameters" for `exrequisite` examples.

`exrequisite [+]= tag-string`
 Defines an exclusive relationship between the current `sw_sel` and the one referenced by tag-string. This will prevent the referenced `sw_sel` from being selected any time that this `sw_sel` is selected (and vice-versa). The `+ =` operator may be used to define multiple `exrequisites`.

Instead of automatically selecting something for installation based upon the name of a `sw_sel` clause, `exrequisite` instead works on the name of a `sw_category`⁶⁰. This form could be used to set up two sets of incompatible software in different software categories and to use the name of the `sw_category` as an `exrequisite` to software in the other category.

`exrequisite = sw_category`
 Specifies that this `sw_sel` is to be exclusive with all other software in the same category. This attribute can be used in addition to other `exrequisites` and does not override them even though the `+ =` operator is not used in this case.

The `load_with_any` configuration item gives you flexibility in automatically selecting software to be installed, based upon what other software may be selected for installation. By using the `=` sign with `load_with_any`, you specify that when any of the listed `sw_sel` clauses are selected, this `sw_sel` is selected also.

`load_with_any = tag-string [[| tag-string]...]`
 Specifies that when any of the one or more `sw_sels` listed are selected, that this `sw_sel` should be selected for loading as well. Multiple tags may be listed separated by the `|` character.

The following form enables regular expression matching (not EREs, only those allowed by `fnmatch(3)`) to select the software or category that software should be loaded with.

`load_with_any ~ tag-regexp [[| tag-regexp]...]`
`load_with_any ~ category.tag-regexp ...`

⁶⁰ Be careful when creating `sw_category` names that are the same as `sw_sel` names. Using `exrequisites` can cause Ignite-UX to produce unexpected results if you use the same name for an existing `sw_sel` and `sw_category`.

Similar to above, except that when the ~ operator is used instead of =, then the sw_sel tags listed are treated as [fnmatch\(3C\)](#) regular expressions that can be used to match any selection fitting the given pattern. (Note that this operator uses [fnmatch\(3C\)](#) pattern matching expressions. Other uses of the ~ operator in the configuration file, uses extended regular expressions. This difference is for compatibility with some existing data.) The [category](#) string can be specified (with a ".." separator) in order to limit the matches to those selections in the specified software category.

The following example uses expression matching:

```
sw_sel "perl" {
    description = "Perl Programming Language"
    sw_source = "core"
    sw_category = "OrderedApps"
    sd_software_list = "perl,r=B.5.6.1.C,a=HP-UX_B.11.11_32/64,v=HP"
    impacts = "/opt" 49480Kb
    load_with_any ~ "HPUXBaseOS" . "*"
}
```

The preceding load_with_any automatically selects perl to be loaded when any sw_sel with a category of HPUXBaseOS is selected.

The load_with_all configuration item enables you to specify a list of sw_sel clauses that must all be selected automatically to select a sw_sel for installation:

```
load_with_all = tag-string [[ & tag-string ]...]
Specifies that when all the sw_sels listed are selected,
then this sw_sel should be selected as well. Multiple tags
may be listed separated by the & character.

If multiple load_with_* keywords are specified, the list of
tag-strings will be added to any already listed for the
respective keyword.
```

The load_with_* configuration items automatically deselect sw_sel clauses as appropriate when you are changing what sw_sel clauses are selected for installation (that is, enforce dependencies). There is one exception to this rule: if you manually select a sw_sel for installation and it has any load_with_* configuration items that are not satisfied, it cannot be automatically deselected because it was manually selected.

Installing patches

Patches are best installed using software configuration rather than command hooks like "post_load_cmd" because a new kernel is built after all software is installed. If you install patches using command hooks, a new kernel is not built, which may cause some unusual behavior in HP-UX where commands depend on kernel functionality that is not in the kernel.

Configuration for volume and disk groups

In this section, the default configuration file for an HP-UX release is assumed always to be included before any others. Because internal variables defined by that configuration file must be set for Ignite-UX to work correctly.

Overview

The general format for Disk Group, Volume Group, and whole disk layout is as follows:

```
partitioned_disk
{
    physical_volume disk[hw_path|index]
    {
        physical_volume_attributes...
    }
    fs_partition
    {
        file-system-attributes...
    }
    swap_partition
    {
        swap-attributes...
    }
}
volume_group group_name
{
    volume_group_attributes...
    physical_volume group "pvname"
    {
        physical_volume disk[hw_path|index]
        {
            physical_volume_attributes...
        }
        ...
    }
    ...
    physical_volume disk[hw_path|index]
    {
        physical_volume_attributes...
    }
    ...
    logical_volume
    {
        file-system/swap-attributes...
    }
    logical_volume "lvname"
    {
        file-system/swap-attributes...
    }
    ...
}
```

Note that partitioned disks are not used very often in more modern systems, as the size of the boot disk cannot be more than approximately 2GB in size on PA-RISC systems. There is no such restriction on Itanium®-based systems.

Before looking at anything in detail, the attributes that can be set in each of the preceding examples are introduced in the following configuration examples. For more details on the individual attributes, refer to *instl_adm(4)*.

Configuration examples

The following section presents several different disk configuration examples. They may prove useful when developing your own configurations.

This section does not discuss the extra space that may be allocated to file systems because of impacts configuration statements or the effect that variables such as `_hp_addnl_pct_free` or `_hp_efi_partition_size` may have.

Example One (custom disk layout)

This example defines a configuration that is dependent on whether or not a particular software product is installed:

```
#  
# This config file builds on top of the default config file for  
# an OS release since it reuses some of the same variables.  
# It should be included into a cfg clause in the index file after  
# the default config file for a release. The following software  
# should be defined in any clauses that includes this file:  
#   "Product XYZ"  
#     "Product XYZ - 4GL"  
#
```

Remember that if you change the disk configuration, the disk layout name is changed to "Modified LVM Layout". Consequently, things defined in the configuration may no longer be evaluated anymore. For example, you can set the name of the volume group in the configuration in this example using the **Additional** button on the **Basic** tab in the Ignite-UX GUI (see `_my_volume_name` later in this section). However, once you modify some part of the disk layout, the name of the layout becomes "Modified LVM Layout" and changing the value of `_my_volume_name` no longer has any effect.

The name is defined by using the selection list produced by the **Additional** button on the **Basic** tab in the Ignite-UX GUI. For the selection, "Name of the root volume group" you can select `vg00` or `vgroot`. However, you can also enter a new name into the field and that will be used (since it was not defined as an enum you cannot enforce a choice from the available names).

The name can be changed on the **Additional** button. However, it only takes effect when the disk layout selected is "Custom configuration for Product XYZ".

```
#  
# Allow the user to change the name of the root volume group  
#  
_my_volume_name = { "vg00" , "vgroot" }  
init _my_volume_name = "vg00"  
_my_volume_name help_text "Name of root volume group"  
_my_volume_name visible_if TRUE
```

Here a new disk layout name is being added to the configuration without setting it as one of the choices. This configuration file builds on top of the existing default operating system release configuration file. It is not a replacement because many variables need to be defined that are not worth the effort to set.

```

#
# Add a new disk layout type
#
_hp_disk_layout += { "Custom configuration for Product XYZ" }

```

The test on `sw_sel` Product XYZ to the variable set `_hp_disk_layout` to Custom configuration for Product XYZ is only evaluated once. If Product XYZ is not selected for installation by the time this configuration file is parsed, you must manually change the disk layout to Custom configuration for Product XYZ during installation. That is, after selecting the software products that the disk configuration depends on, if they are to be installed, to make sure that the optional logical volume is created.

To make the disk configuration modifiable, remember to use `init` before the `_hp_disk_layout`.⁶¹

```

#
# Automatically select the disk layout if product XYZ is being
# installed.
#
(sw_sel "Product XYZ") {
    init _hp_disk_layout = "Custom configuration for Product XYZ"62
}

```

Next is the definition of the disk layout Custom configuration for Product XYZ:

```
( _hp_disk_layout == "Custom configuration for Product XYZ" ) {
```

First, make sure that all disks are greater than 36GB. An error will be issued if there are smaller disks because this disk layout is only intended for disks 36GB or higher.

Note:

This prevents the user from proceeding with an installation (or force a non-interactive installation to become interactive) if the size of the root disk is less than what is being tested for.⁶³

Looking for errors where possible can help ensure that a configuration does not allow someone to do something that cannot work.

```
(disk[_hp_root_disk].size < 34000Mb) {
    error+="This configuration cannot be used with disks of less than
36Gb"
```

⁶¹ Without "init" in front of the `_hp_disk_layout` variable, the Ignite-UX GUI cannot make changes to the disk configuration since all changes (in this case) are saved into the disk layout called "Modified LVM Layout". If the Ignite-UX GUI cannot change `_hp_disk_layout`, then the name of the disk layout cannot be changed and no file system changes asked for on the file system tab are honored.

⁶² You cannot have an effects relationship between a variable and a software selection (or a use model). If this software product is selected and the disk layout is set to "Custom configuration for Product XYZ", unselecting the software does not cause the disk layout to be reevaluated.

⁶³ Keep in mind that GB for disks is not $36 * 1024 * 1024 * 1024$ (38654705664) bytes but approximately $36 * 1000 * 1000 * 1000$ (36000000000) bytes, a difference of over 2GB. This test, to verify whether the size is less than 34000MB (35651584000 bytes), is approximately correct.

```
}
```

You can now start the definition of the volume group that has its name set by the variable `_my_volume_group`; the volume group is explicitly defined to be the LVM:

```
volume_group _my_volume_name
{
    #
    # Only allow LVM
    #
    usage = LVM
```

You now define some volume group attributes, specifically the maximum number of physical extents (Ignite-UX changes this to suit the disks as required). In addition, you set up the PE size here as well. The default configuration file sets the variable `up _hp_root_disk` for you.

```
#
# Setup for defaults allowing of 36 and 72Gb disks as boot disk
#
max_physical_extents = 2500
(disk[_hp_root_disk].size < 40000MB) {
    physical_extent_size=8
} else {
    physical_extent_size=16
}
#
# The root disk was found by Ignite-UX for us
#
physical_volume disk[_hp_root_disk]
```

Now you define the root file system. In all of these file systems, you are not giving them explicit names; this is left up to Ignite-UX. Here, you only set the minimum number of fields for the root file system attributes:

```
#
# Define the root file system
#
logical_volume
{
    mount_point = "/"
    usage=VxFS
    size=400Mb
    contiguous_allocation = true
    bad_block_relocate = false
}
```

Next is primary swap and dump. This configuration does not define a secondary swap. The size is interesting; allocate at least `_hp_min_swap` as the size. The variable `_hp_min_swap` again is set by the default configuration file. The maximum size is `_hp_pri_swap`; however, depending on the available disk space, the actual amount of space allocated is somewhere between `_hp_min_swap` and `_hp_pri_swap`.

```
#
# Define the swap/dump
#
logical_volume
```

```

{
    mount_point = "primary"
    usage=SWAP_DUMP
    contiguous_allocation = true
    bad_block_relocate = false
    size = _hp_min_swap | remaining | _hp_pri_swap
}

```

Currently, for PA-RISC systems, the boot file system *must* be HFS. For Itanium®-based systems, the boot file system can be either HFS or VxFS. So here, you define a fixed-size boot file system (/stand) to be HFS. HFS is the most compatible, although you could add tests to see if the system is an HP9000 or HP Integrity server, and then change the file system type based upon that information.

```

#
# Define /stand
#
logical_volume
{
    mount_point = "/stand"
    usage=HFS
    blksize = _hp_HFS_blksize
    fragsize = _hp_HFS_fragsize
    contiguous_allocation = true
    bad_block_relocate = false
    size = 300Mb
}

```

Next, you define /tmp; the size is predicated on the current selection state of Product XYZ – 4GL at the time the configuration is parsed. If the software is selected when the configuration is parsed, the size is different from when it is not selected. If the software changes state after this disk layout is chosen, the size does not change because of it.⁶⁴

Important:

Size is only defined based upon software selection when the configuration is initially parsed.

If the product Product XYZ – 4GL is selected, the size is a minimum 100MB and up to 400MB depending on the remaining disk space. If the product Product XYZ – 4GL is not selected, the size is a minimum 50MB and up to 200MB depending on the remaining disk space.

```

#
# Define /tmp
#
logical_volume
{
    mount_point = "/tmp"
    usage=VxFS
    (sw_sel "Product XYZ - 4GL" ) {
        size = 100Mb | remaining | 400Mb
    } else {

```

⁶⁴ You cannot have effects relationships between variables and software selections and use models.

```

        size = 50Mb | remaining | 200Mb
    }
}

```

The next file system is `/usr`. The minimum size is 500MB. Depending on available disk space, the volume may be expanded to require a minimum of 20 percent free space after the file system impacts statements associated with software being installed into `/usr` are taken into account.

Minimum free percentages of disk space may be a good compromise for file systems in systems where you know how much space is needed and you only require some additional overhead space to be left available (this is true for file systems such as `/usr` and `/opt`).

```

#
# Define /usr
#
logical_volume
{
    mount_point = "/usr"
    usage=VxFS
    size = 500Mb | remaining | 20% free
}

```

The `/var` file system has a different definition for its size. It is whatever space can be allocated up to 1000MB plus the size of `_hp_pri_swap`. Therefore, the final size for this file system depends on the free disk space once minimum sizes have been allocated to file systems.

```

#
# Define /var
#
logical_volume
{
    mount_point = "/var"
    usage=VxFS
    size = remaining | 1000Mb + _hp_pri_swap
}

```

The `/var/opt/pxyz` file system is only defined to be created if the software product Product XYZ - 4GL is selected for installation when this disk layout is selected. This is similar to the previous size tests for `/tmp`. This file system is not defined if the product Product XYZ - 4GL is selected for installation after the disk layout has been changed to Custom configuration for Product XYZ. Similarly, the file system is not removed from the disk layout if the product Product XYZ - 4GL is unselected unless the disk layout is changed and then changed back to Custom configuration for Product XYZ. Of course, it could be manually removed.

```

#
# Define /var/opt/pxyz
#
(sw_sel "Product XYZ - 4GL" ) {
    logical_volume
    {
        mount_point = "/var/opt/pxyz"
        usage=VxFS
        size = remaining | 2000Mb
    }
}

```

The last volumes are defined as follows:

```
#  
# Define /var/tmp  
#  
logical_volume  
{  
    mount_point = "/var/tmp"  
    usage=VxFS  
    size = remaining | 1000Mb  
}  
#  
# Define /opt  
#  
logical_volume  
{  
    mount_point = "/opt"  
    usage=VxFS  
    (sw_sel "Product XYZ - 4GL" ) {  
        size = remaining | 1500Mb  
    } else {  
        size = remaining | 20% free  
    }  
}  
#  
# Define /home  
#  
logical_volume  
{  
    mount_point = "/home"  
    usage=VxFS  
    size = 500Mb  
}
```

The default configuration files make setting up a disk configuration much more complex than for a relatively simple disk configuration.

To change the layout of this root volume group to be VxVM, you must set the name of the disk group to `rootdg`; instead of giving the user a choice of names, change `usage = LVM` to `usage = VXVM`.

Example Two (selection of disk layout based on hardware)

The configuration example in this section shows a fragment of a disk layout definition that selects the disk layout based upon the model string and the size of the disk you are installing to (you do not need to worry about the disk layout itself).

You must first add the new disk layout names into `_hp_disk_layout`. Remember to add them as choices because you want the user to be able to select them using Ignite-UX GUI:

```
#  
# Add new disk layout types  
#  
_hp_disk_layout+= { "Custom configuration rp8400 w/18Gb disks",  
                    "Custom configuration rp8400 w/36Gb disks",  
                    "Custom configuration rp8400 w/72Gb disks",  
                    "Custom configuration rp8400 w/>100Gb disks" }
```

The following test makes sure that the root disk size is at least the minimum required size for installing an rp8400 system⁶⁵. If the root disk is too small, an error message is presented to users to prevent them from being able eventually to select **Go!** in the Ignite-UX GUI and install the system. The only recourse is to change the disk layout or increase the size of the root disk by selecting a new one.

```
(hardware_model == "9000/800/SD16K-A" & disk[_hp_root_disk].size < 16000Mb )
{
    error += "This configuration does not support installing disks of
<18Gb please choose another disk to install to that is at least 18Gb in size"
}
```

The following test determines whether the disk layout should be changed to Custom configuration rp8400 w/18GB disks. A similar test is done for all of the other disk layouts that you can set.

```
(hardware_model == "9000/800/SD16K-A" & disk[_hp_root_disk].size >= 16000Mb &
disk[_hp_root_disk].size < 30000Mb ) {
    init _hp_disk_layout = "Custom configuration rp8400 w/18Gb disks"
}

(hardware_model == "9000/800/SD16K-A" & disk[_hp_root_disk].size >= 30000Mb &
disk[_hp_root_disk].size < 60000Mb ) {
    init _hp_disk_layout = "Custom configuration rp8400 w/36Gb disks"
}

(hardware_model == "9000/800/SD16K-A" & disk[_hp_root_disk].size >= 60000Mb &
disk[_hp_root_disk].size < 100000Mb ) {
    init _hp_disk_layout = "Custom configuration rp8400 w/72Gb disks"
}

(hardware_model == "9000/800/SD16K-A" & disk[_hp_root_disk].size >= 100000Mb )
{
    init _hp_disk_layout = "Custom configuration rp8400 w/>100Gb disks"
}
```

Next, you start the Custom configuration rp8400 w/18GB disks disk layout definition. The tests that you perform inside the definition prevent anyone from misusing the defined layout for the wrong system type and disk size. These tests reinforce the previous default selection code.

```
(_hp_disk_layout == "Custom configuration rp8400 w/18Gb disks") {
    volume_group "vg00"
    {
        usage = LVM

        (hardware_model != "9000/800/SD16K-A") {
            error += "This disk layout is for use with a rp8400"
        }
        (disk[_hp_root_disk].size < 16000Mb |
        disk[_hp_root_disk].size > 30000Mb ) {
            error += "Disk size not between 16000Mb and 30000Mb"
        }
    ...
}
```

⁶⁵ The rp8400 is a marketing name for a system. The model string of the system is actually 9000/800/SD16K-A.

You leave the custom configuration at that point since you are only looking for some other tests for automating selection and enforcing it with tests and the `error` keyword. Instead of using `error`, using `warning` might alert the user to choose the selection when the tests fail (for example, to use it on an rp7410 instead.)

Example Three (be careful what you ask for, you just might get it)

This example disk layout asks for something that just can't be done. It's a cautionary example to show the errors displayed if you attempt to implement an impossible configuration. This is not a complete configuration, just part of one.

```
#  
# Test configuration  
#  
release="B.11.23"  
...  
init _hp_pri_swap = disk[_hp_root_disk].size  
...  
_hp_disk_layout = { "My Customised LVM" }  
init _hp_disk_layout = disk[_my_second_disk_path].model  
init _hp_disk_layout = "My Customised LVM"  
_hp_group_name visible_if FALSE  
...  
(_hp_disk_layout == "My Customised LVM") {  
  
    volume_group "vg00" {  
        physical_volume disk[_hp_root_disk]  
        physical_volume disk[_my_second_disk_path]  
        usage = LVM  
    ...  
        logical_volume {  
            usage = SWAP_DUMP  
            mount_point = "primary"  
        # Map it to the root disk to ensure it is considered primary  
        # without this the secondary swap gets sorted first and considered  
        # primary.  
            vgdisk[0]  
            contiguous_allocation = true  
            bad_block_relocate = false  
            size = _hp_pri_swap  
        }  
    ...  
}
```

In the definition above, the primary swap size is set to `_hp_pri_swap`, which was initialized to the size of the disk with “`init _hp_pri_swap = disk[_hp_root_disk].size`”

Because we've explicitly set the size of primary swap, it cannot be changed. This causes a problem given the small amount of overhead for LVM disk structures, and because the root and boot file systems must be on the disk `_hp_root_disk` as well as primary swap. They simply cannot fit, and this will give errors.

During installation those errors would look like:

```
ERROR: There is not enough disk space to allocate for volume: "primary"  
(short by 28Mb). For more information run "instl_dbg -D <clientdir>  
-d".  
ERROR: There is not enough disk space to allocate for volume: "/" (short by
```

```
1000Mb). For more information run "instl_dbg -D <clientdir> -d".  
ERROR: The root volume (/) size plus the boot volume (/stand) size plus  
primary swap size plus dump volume(s) size (1024000Kb + 24576Kb +  
17780736Kb + 0Kb) is larger than the root disk (0/0/2/1.15.0)capacity  
(17776640Kb) by 1052672Kb. These volumes must fit contiguously on the  
root disk. For more information run "instl_dbg -D <clientdir> -d".
```

In this case there were two disks in the root volume group. Ignite-UX was able to move all of the logical volumes except root, boot, and primary swap to the second disk. But because the size value for primary swap was so large, Ignite-UX was unable to fit root, boot, and primary swap onto the disk _hp_root_disk, and the installation failed.

Assuming the value for _hp_min_swap (if you've included the per-release config file⁶⁶ into the cfg clause) is not too large, the size would be better given as:

```
size = _hp_min_swap | remaining | _hp_pri_swap
```

This would allow the size of primary swap to be at least _hp_min_swap, but can take all the remaining space on the disk up to the value of _hp_pri_swap (which we know is larger than the size of the disk.)

Doing this would allow the disk configuration to successfully create a system.

Example Four (installation file system configuration)

The example in this section is intended for use in an installation file system, not in configuration files. The best use of this configuration is in an installation file system contained in LIF and used to boot a system using tape|CD|DVD so it does not use DHCP to gain an IP address. This section does not discuss other configuration items you might set like _hp_lanadmin_args.

This example is intended for use with HP9000 systems. For information on how to achieve these same results with HP Integrity systems (or to create media that can be used on both hardware architectures) see "Building an Installation DVD" in the *Ignite-UX Administration Guide*, Edition 18 or later. Additionally, this section contains information about the placement of a LIF inside an El-Torito boot image for use with Itanium®-based systems.

Part A (custom configuration in installation file system)

This part of the example deals with simply providing an IP address and network information that can be applied to an LIF (and hence to the installation file system within the LIF.) This method requires you to build a different LIF for every system and place it onto media separately.

The following configuration is needed to contact the Ignite-UX server:

```
server="10.0.0.2"  
netmask[ ]="255.255.255.0"  
route_gateway[0]="10.0.1.1"  
route_destination[0]="default"  
disable_dhcp=true
```

This line gives the system its IP address:

⁶⁶ Since this example is about 11.23 the per-release config file would be /opt/ignite/data/Rel_B.11.23/config.

```
ip_addr[ ]="10.0.1.45"
```

Doing this might not be very useful on multi-homed hosts because the LAN interface at the lowest hardware address is given the IP address. To give the IP address to a specific LAN interface you must do the following:

```
ip_addr[0/0/0/0]="10.0.1.45"
```

This is applied to an LIF (created with `make_medialif -m`) with the `instl_adm` command. The LIF can then be placed on media.

Note:

Unless you are using unallocated IP addresses, a boot tape or CD|DVD created in this way is specific to a system. You do not want to boot another system with it while the IP address is in use on a live system.

Part B (Installation file system custom network config)

This part of the example deals with providing a more generic (for example, can be used on multiple systems and has information unique to each) solution for IP address and network information that can be applied to an LIF (and therefore to the installation file system within the LIF).

The following section would be generic to all systems located on one LAN segment; they can all access the same gateway:

```
server="10.0.0.2"
netmask[ ]="255.255.255.0"
route_gateway[0]="10.0.1.1"
route_destination[0]="default"
disable_dhcp=true
```

For the system-specific parts, each of the definitions of the IP address is protected by a test on the MAC address:

```
( lla[0/0/0/0] == "080009654321" ) {
    ip_addr[0/0/0/0]="10.0.1.45"
}
( lla[2/0/2] == "080009654123" ) {
    ip_addr[2/0/2]="10.0.1.46"
}
( lla[1/0/0/0] == "080009654213" ) {
    ip_addr[1/0/0/0]="10.0.1.47"
}
...
...
```

This enables the tape or CD|DVD to be used on all systems within one subnet.

Example Five (handling user input)

In this example we're going to look at creating a variable that behaves in different ways. In the first method we're going to create an `enum` and force someone to only choose values that are allowed.

In the second method we will look at a way to define a variable that allows someone to pick from a list of values or allows them to enter any value.

Part A (restricting the values that a variable can take using an enum)

The following configuration is what we want to end up with. We will look at what it contains line by line and then look at how the variable can be used:

```
enum _my_location_codes
_my_location_codes = { "Not Valid", "AAA000", "AAA001", "AAA002", "AAA003" }
init _my_location_codes = "Not Valid"
_my_location_codes visible_if TRUE
_my_location_codes help_text "Choose a location code"
```

The first line defines the variable `_my_location_codes` to be an enum. That means that it can only hold a value from the enumerated values assigned to it. It is not possible to set this variable to any other value.

The second line assigns a list of values to the variable. A variable doesn't have to be an enum to give it a list of potential values like this, more on that in the next part of this example.

The third line gives the variable an initial value. Since this variable is an enum it should only be assigned a value from the list of allowed values. Note the use of `init` before the variable - the `init` keyword is only required if you want the user to be allowed to change this value via the additional button on the **basic** tab in itool. Without the `init` keyword the user would not be allowed to change the value.

The last line sets the variable so it can be seen via the **additional** button on the **basic** tab in itool. If you did not want to allow the user to see this variable you must use `FALSE` instead of `TRUE`.

You should be aware that enums behave differently than other variables on the **additional** button on the **basic** tab in itool. You cannot directly enter a value into an enum field, instead press **Enter** where the help text is. When you do this you will be presented a list of values to choose from. That is how to change the value for an enum.

The end result of this is to present to the user a field that they can change but can only select one of a predefined set of values. The question you may be asking is why you would do this. In this case the value of the `_my_location_codes` variable could control networking information. If there were one location code for each subnet you have HP-UX systems on, you could choose the appropriate location code to supply information such as routing and other networking information common to that subnet. This would require additional configurations that tested the value of `_my_location_codes` and set up networking information accordingly.

You may have noticed that `_my_location_codes` contained "Not Valid" as one of its possible values. There is a good reason for this - in this case we want to force the person doing the installation to set this value. If there is an initial value that you know is not valid, you can create some configuration like:

```
( _my_location_codes == "Not Valid" ) {
    error+="You have not selected a location code."
} else {
    note+="The location code selected was " + ${_my_location_codes}
```

```
}
```

Since we can test on the initial value of `_my_location_codes`, we can show an error and prevent any installation from continuing if a location code has not been set. If one has been chosen we can print out information about the location code.

Part B (looking at the same thing when we don't have an enum)

The following configuration is what we want to end up with in this case. We will look at what it contains line by line and then look at how the variable can be used:

```
init _my_location_codes = "Not Valid"
_my_location_codes = { "AAA000", "AAA001", "AAA002", "AAA003" }
_my_location_codes visible_if TRUE
_my_location_codes help_text "Choose a location code"
```

We can tell that we're looking at something different from the previous part of this example since the `enum` keyword isn't being used. Since we haven't used `enum`, the variable is not restricted to the list of predefined values.

That means the person performing the installation can choose a value from a list (similar to `enum`) or they can enter a value directly into the field using the **additional** button on the **basic** tab in `itool`.

This approach offers more flexibility than an `enum`. One example use would be the name of a user in a configuration variable - there might be some commonly set values, but you don't want to restrict input to just that list. In such a case you might provide the common user names to choose from but allow the person performing the install to enter any user name.

The same configuration used above can still be used to validate the variable. We can ensure that the value has been set to something other than the initial value of "Not Valid":

```
( _my_location_codes == "Not Valid" ) {
    error+="You have not selected a location code."
} else {
    note+="The location code selected was " + ${_my_location_codes}
}
```

Again, note the use of the keyword `init` before `_my_location_codes`. The `init` keyword is required when the configuration assigns a value to a variable the user will be able to change. If it is not present, the user interface (the **additional** button on the **basic** tab in `itool`) will not accept changes to the value of the variable (you can change it but the changes will not take effect.)

An example of a variable defined in this way is the configuration variable `_hp_addnl_fs_free_pct` (covered elsewhere in this document.)

Example Six (handling command line arguments)

This example examines building variables that are based upon other variables. For this example we are going to construct a set handlers for command line arguments. We have three arguments to cater for, those arguments are:

- `-a` which is fixed and doesn't change
- `-b` which only takes a value between 1 and 10

- -c which needs to be able to take any specified value (a user name)

Since -a is fixed we don't do anything about it. We only need to do something about the variable parts.

For option -b we need to create it as an enum because it only takes one of a specific set of values. We don't need to worry about special configurations to ensure that the user sets a specific value, as in this case we have chosen a suitable default of "5".

```
enum _my_command_b_option
_my_command_b_option = { "1", "2", "3", "4", "5", "6", "7", "8", "9", "10" }
init _my_command_b_option = "5"
_my_command_b_option visible_if TRUE
_my_command_b_option help_text "command -b value"
```

For -c it's a bit easier since we want to allow any value (since you can't validate a user name):

```
init _my_command_c_option = "user1"
_my_command_c_option visible_if TRUE
_my_command_c_option help_text "command -c value"
```

It's important to remember the init in front of the variable. If you left it out the user would not be able to change the value via the **additional** button on the **basic** tab in itool.

We can put together a command line as follows:

```
_my_command_line = "-a -b " + ${_my_command_b_option} + " -c " +
${_my_command_c_option}
_my_command_line visible_if FALSE
```

The variable `_my_command_line` could be used in the following way:

```
post_config_cmd+=" /opt/myapp/bin/command " + ${_my_command_line}
```

or be placed into the environment for when a script is ran later:

```
env_vars+="COMMAND_OPTS=" + ${_my_command_line}
```

Example Seven (regular expression matching)

In the previous example we mentioned that it wasn't possible to validate a user name:

```
init _my_command_c_option = "user1"
_my_command_c_option visible_if TRUE
_my_command_c_option help_text "command -c value"
```

That is true, but it is still possible to perform sanity checking on the value. We can do that with the help of regular expressions (information on regular expressions is available in the *regexp(5)* manual page – Ignite-UX uses extended regular expressions.) We will look at some examples of what is possible.

This example is not a reference on how to write regular expressions (complex regular expression writing can be very difficult) but we will look at some common operators and how you might use them. The information provided here is only a very basic introduction to extended regular expressions.

An important thing to remember is that `~` is the operator that controls regular expression matching within Ignite-UX configuration files. When writing regular expressions there are several operators you might find useful⁶⁷:

- `^` is an anchor to the start of what is being matched against. This might allow you to say something like `^u` which would match anything that started with a `u`. Without the `^`, a regular expression that just contained `u` would match any occurrence of `u` in what was being matched against.
- `$` is an anchor to the end of what is being matched against. This might allow you to say something like `12$` which would match anything that ended with `12`.
- `|` in a regular expression acts as an `or` operator. This allows you to give several regular expressions together where either expression is allowed to match. For example `^user|^sysadm` would match anything that started with `user` or `sysadm`.
- `.` allows you to say any character. For example, `u.s` would match anything that contained (no anchor is in the expression) the letter `u` followed by any character followed by the letter `s`.
- `[]` forms a bracket expression in a regular expression. This construct can be used to give multiple values that can be used to match a single character. For example `[ahj]` means anything that contains the letters `a`, `h` or `j`.
- You can give a character class in the regular expression, for example `[:alpha:]` can be used to match any letter (upper or lowercase). The manual page `regexp(5)` contains more information on character classes (including the names of other character classes.) Note that you must use this within a bracket expression⁶⁸ so it will end up looking like `[[[:alpha:]]]`.
- `*` can be used to perform wildcard matching, the regular expression `u*r` matches anything containing a `u` followed by zero or more characters followed by an `r`. For example, this would match `ur`, `usr`, and `user` (and many other things.)
- Do not use the construct `.*` (zero or more of any character) at the start or end of an unanchored regular expression, it is pointless. For example `.*user.*` means zero or more of any character followed by `user` followed by zero or more of any character. The reason why this is pointless is that `user` by itself will match any part of the string it is being compared against. If compared against the string `auser` it will match without any additional `.*` in the regular expression.
- `+` Allows you to match one or more occurrences of a pattern within an expression.

Note that the meaning of any of the regular expression operators mentioned above can change depending on their position. An example of this would be `|` inside a bracket expression: inside a bracket expression it means the character `|` not the `or` operator. Please refer to the `regexp(5)` manual page for additional information.

Now we can look at the use of regular expressions for adding some sanity checks against the user name from Example 5:

⁶⁷ Some of these explanations have been simplified; you should consult the `regexp(5)` manual page for more complete explanations.

⁶⁸ Outside of a bracket expression `[:alpha:]` in a regular expression would mean match the characters `:alpha:` but `[[[:alpha:]]]` means match any upper or lower case letter.

```
( _my_command_c_option ! ~ "^\user" ) {
    error+="The user name must start with user"
} else {
    note+="The user name chosen was " + ${_my_command_c_option}
}
```

Note the use of the “!~” operator in the test: this reverses the comparison, so an error is given if the value in `_my_command_c_option` does *not* match the regular expression. The example is presented this way since you may not want to always print the note - the `else` and the note can be skipped.

The test above compares the value in `_my_command_c_option` against “`^\user`”. This expression is anchored to the start of what we are comparing against, so matches must start with “`user`”. It is important to understand that anything can be after the “`user`”. Unlike some other regular expressions, such as shell regular expressions, extended regular expressions do not require a wildcard operator after the “`user`” to get a match.

```
( _my_command_c_option ! ~ "^\[:alpha:]\[:digit:]*" ) {
    error+="The user name must start with letter followed by only digits"
} else {
    note+="The user name chosen was " + ${_my_command_c_option}
}
```

In the above example you might think that the regular expression matches something that starts with a letter followed by one or more numbers. It doesn’t actually mean that. It really matches something that starts with a letter followed by zero or more numbers. The above regular expression would match just a single letter.

If you want to only allow user names in the form of a single letter (upper or lower case) followed by exactly six numbers, you would be better-off with a regular expression like:

```
^\[:alpha:]\[:digit:]\[:digit:]\[:digit:]\[:digit:]\[:digit:]\[:digit:]\[:digit:]\$
```

The expression is anchored at the start and the end, so we know what we will be matching against will match exactly or not at all. There are no wildcards or other constructs that can cause problems for us. The above regular expression will only match something that starts with a single letter followed exactly six digits.

If your requirements aren’t that strict, you can use a slightly different expression:

```
"^\[:alpha:]\[:digit:]+\[:digit:]\$"
```

This matches a user name starting with a letter, followed by one or more numbers, and ending with a number. This will match the following: “`u12`”, “`u123`”, and “`u1234`”. The only requirements for a match are that the user name starts with a letter and contains at least two digits⁶⁹ after it (with nothing else after that.)

⁶⁹ The regular expression “[[:digit:]]+[:digit:]” requires that there be at least a starting and ending number there can be 0 or more numbers in between.

You can test regular expressions against expected data outside of Ignite-UX using `grep` with the `-E`⁷⁰ option:

```
# cat data
u1
u12
u123
u1234
U1a
U123456
# cat data | grep -E "^[[:alpha:]][[:digit:]]+[:digit:]]$"
u12
u123
u1234
U123456
```

This allows you to have some level of confidence that your regular expression will do what you want it to.

Another example:

```
( _my_command_c_option ! ~ "^[[:alpha:]][[:digit:]]+[:digit:]]$|root" ) {
    error+="The user name must start with letter followed by only digits"
} else {
    note+="The user name chosen was " + ${_my_command_c_option}
}
```

The regular expression above adds “`|root`” to one of the previous expressions we’ve looked at, which matches the user `root` as well as the other users allowed by the expression. Remember that the “`|`” operator means “or” and allows you to give multiple regular expressions to match against.

Testing this with `grep` you can see that this does work:

```
# cat data
u1
root
u12
u123
u1234
U1a
U123456
# cat data | grep -E "^[[:alpha:]][[:digit:]]+[:digit:]]$|root"
root
u12
u123
u1234
U123456
```

⁷⁰ This option is required as it enables extended regular expressions in `grep` so we can test in exactly the same way that Ignite-UX will perform its regular expression matching. Without the `-E` option `grep` will only use basic regular expression matching. As an alternative to `grep -E` you can just use the `egrep` command.

Although we've concentrated on using character classes for matching, there are cases where this may not be appropriate. In this example we will restrict the letter the user name starts with: "u", "a", or "s".

```
( _my_command_c_option ! ~ "^[uas][[:digit:]]+[:digit:]$" ) {  
    error+="The user name must start with u, a, or s followed by only digits"  
} else {  
    note+="The user name chosen was " + ${_my_command_c_option}  
}
```

So instead of allowing any upper or lower case letter to start the user name, we now need one of: "u", "a", or "s".

It is important to understand what bracket expressions match. A bracket expression is used to match a single character. Look at the following test:

```
$ cat data  
ab000123  
cd000243  
ad999233  
ef000033  
gh888444  
$ grep -E "^[acg][bdh][[:digit:]]+[:digit:]$" data  
ab000123  
cd000243  
ad999233  
gh888444  
$ grep -E "^(ab|cd|gh)[[:digit:]]+[:digit:]$" data  
ab000123  
cd000243  
gh888444
```

The bracket expressions [acg][bdh] mean: match one character, "a", "c", or "g", followed by one character, "b", "d", or "h". If you wish to match something containing "ab", "cd", or "gh" then a bracket expression is not what should be used.

In that case you can use parentheses with the "or" operator to group expressions together, so any of the expressions are used for matching. In the following example we have (ab|cd|gh), which matches "ab", "cd", or "gh". The parentheses are important. If you didn't have them, the following extended regular expression:

```
"^ab|cd|gh[[:digit:]]+[:digit:]$"
```

would match:

1. Something that starts with "ab"
2. Something that contains "cd"
3. Something that ends with the pattern "gh" followed by one or more numbers, and ends in a number (since it is anchored to the end of what is being matched against).

So the following expression tests for a user name that starts with "ab", "cd", or "gh" followed by one or more numbers and ends in a number:

```

(_my_command_c_option ! ~ "^(ab|cd|gh)[[:digit:]]+[:digit:]]$") {
    error+="The user must start with ab, cd, or gh followed by only digits"
} else {
    note+="The user chosen was " + ${_my_command_c_option}
}

```

Try to restrict yourself to simple regular expressions. If complex regular expressions are needed, test them extensively against large amounts of potential values to ensure they behave the way you expect.

Configuration parameters in the installation file system

This section discusses configuration parameters that are most suited to being placed in an installation file system using `inst1_adm`. The most likely reason you would place them in the installation file system would be that they need to be seen by Ignite-UX to control the way Ignite-UX works before it attempts to mount the client's directory from an Ignite-UX server (in a network installation).

For more information regarding the configuration parameters discussed in this section, see "Configuration examples" and "Example Three (be careful what you ask for, you just might get it)."

Networking

The information in this section only applies to the usage of Ignite-UX after the installation kernel has started running AND after the installation file system has been loaded (since the configuration is stored in the first 8KB of the installation file system).

The following networking configuration parameters are covered elsewhere in this document. For more information, see the following:

```

_hp_lanadmin_args (see "_hp_lanadmin_args")
_hp_nfs_mount_opts (see "_hp_nfs_mount_opts")
_hp_nfs_mount_retries (see "_hp_nfs_mount_retries")
_hp_tftp_cmds (see "_hp_tftp_cmds")

```

The following networking options are best set within an installation file system since Ignite-UX must be aware of them either before it contacts the Ignite-UX server or when it is attempting to start networking.

The DHCP configuration parameters are mentioned here because Ignite-UX requests an IP address using DHCP when it starts and before it first attempts to make contact with the Ignite-UX server.

- `dhcp_class_id`—this configuration item can be used "as is" (for example, applies to every LAN interface) or can be given optional LAN information so the DHCP class id is only expected to be used with one particular LAN interface, for example,

```
dhcp_class_id[0/0/0]="Ignite_clients".
```

This item is extremely useful if you have no control over the other DHCP servers on the network and do not want to accept temporary IP addresses from those servers. For example, the following configuration splits Itanium®-based and PA-RISC clients into two different DHCP classes:

```

is_hppa {
    dhcp_class_id = "Ignite_PA_clients"
} else {
    dhcp_class_id = "Ignite_IA_clients"
}

```

The `else` condition implicitly means Itanium®-based systems since Ignite-UX currently supports only PA-RISC and Itanium®-based systems. This also assumes that you would want to share the same installation file system configuration. This condition effectively limits what the client responds to⁷¹.

- `dhcp_server`—this keyword can be used to limit the responses that are accepted by a particular server.

This is not as flexible as the `dhcp_class_id` keyword. If the DHCP server is down, no response is accepted. Compare that to a setup where you have two DHCP servers, each one willing to serve out different IP addresses to the one class id. If one Ignite-UX server goes down, the other DHCP server can still provide an IP address (although not as many are available).

The `dhcp_server` configuration item can apply to all LAN interfaces (for example, `dhcp_server = "10.0.0.53"`) or to one particular LAN interface (for example, `dhcp_server[0/0/0] = "10.0.0.54"`).

- `dhcp_misc_opts`—the `dhcp_misc_opts` configuration item can apply to all LAN interfaces (for example, `dhcp_misc_opts = "-l 6"`) or to one particular LAN interface (for example, `dhcp_misc_opts[0/0/0] = "-l 6"`).

Care should be taken when using the `dhcp_misc_opts` configuration item. The options provided are for /usr/lbin/dhcpclient. Since this is a back-end command, its options are not officially documented and this option should only be used when requested by HP.

The only option you are likely to be asked to add is the option to print debugging output:
`dhcp_misc_opts = "-l 6".`

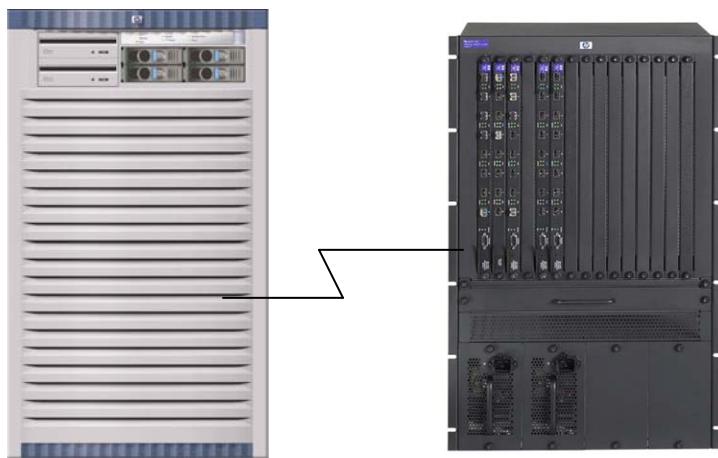
- `disable_dhcp`—when set to TRUE in the installation file system, this item prevents Ignite-UX from attempting to get an IP address using DHCP. You can use this item when you want to manually give Ignite-UX its information (or if you have defined the IP address information in the installation file system).

⁷¹ In a recovery situation, if you want to use DHCP to gain an IP address, Ignite-UX only adds the specific client IP address that actually created the recovery archive into the /etc(exports file. If you can accept DHCP responses from anyone, the potential list of IP addresses that you may need to configure can get quite large, so many manual changes maybe required to the /etc(exports file on the Ignite-UX server. This assumes that archives are written to the Ignite-UX server.

Problems that can be solved with `_hp_lanadmin_args`

When a kernel first boots (this is true even of a normal kernel that is not used for installation) all of the 100bt and Gigabit LAN interfaces attempt to autonegotiate the LAN speed and duplex settings.

In this example, an rp8400 is connected to an HP Procurve 9315m routing switch. When the switch point to which port the system is connected is set to autonegotiate, the switch, and the



system both autonegotiate the speed and duplex. When the switch is set to a speed and duplex value, autonegotiation fails between the system and the switch. This causes the system to revert to 100HD (half-duplex) and the switch continues using its hard-set speed (100FD or full-duplex, for example). Duplex mismatches that can occur in these circumstances can cause extreme throughput problems

and Ignite-UX can appear hung or take a significantly longer time to complete tasks.

During a normal system startup, the system does not experience these problems because the system startup scripts have ideally been modified to set the speed and duplex values correctly for the LAN interface, assuming it is not going to autonegotiate. When using Ignite-UX, startup scripts cannot be used to set the speed and duplex values for your LAN interfaces.

To match the speed and duplex values, you must use the `_hp_lanadmin_args` variable to set these values correctly in the installation file system. This variable should *only* be set using `instl_adm`. For example, you could add the following configuration:

```
( lan[].driver == "btlan" )
{
    _hp_lanadmin_args="-X 100FD"
}
```

This sets all LAN interfaces controlled by the LAN driver `btlan` to 100 full duplex in the installation file system by obtaining the current settings from the installation file system using the following command:

```
# instl_adm -d > /tmp/installfs.config
```

Next, you update the file to contain the extra configuration. If a similar configuration exists already, you may need to consider it when you add a new configuration. The updated configuration must be placed back into the installation file system using the following command:

```
# instl_adm -f /tmp/installfs.config
```

The `instl_adm` command attempts to keep the installation file systems synchronized. If you want only to apply a configuration to specific file systems, do not use the `-f` option; only use the `-F`

option to give explicit installation file system names. However, you should try to keep the various installation file systems consistent.

This issue creates difficulties when you mix system types that are running at different duplexes and speeds, rather than setting the systems to autonegotiate both. To avoid this situation, you should develop and follow a consistent approach to speed and duplex settings.

If you have a gigabit switch set to only operate at 1000FD (gigabit speed with full duplex) and the gigabit interface on system is attempting to autonegotiate the speed and duplex, a 1000FD link will be achieved. This is because the only duplex supported at gigabit speeds is full duplex. However, this is only true of 1000Base-T interfaces so 1000Base-SX interfaces must autonegotiate or they will fail to form a link.

Note that if you are using 1000Base-T interfaces and the switch is autonegotiating link speed and duplex or it is hard set to use 1000FD, you should not need to use `_hp_lanadmin_args` to force any speed/duplex values.

Control

The `_hp_keyboard` (see the "`_hp_keyboard`" section) control configuration item although not strictly an installation file system variable, it is a good idea to set this to a keyboard type if all of the systems that have keyboards are the same type. This prevents HP-UX from asking for a keyboard type when you have a graphics console.

The following control options are best set within an installation file system because Ignite-UX must be aware of them before it contacts either the Ignite-UX server or when it is attempting to start networking:

- `run_ui`—this configuration item controls whether the installation process should be interactive or not. If set to `FALSE`, the Ignite-UX GUI does not run and it is assumed that the configuration files provide the complete configuration required to install the system. If this is not `true` and there is not enough information, the installation changes to become interactive.

Be careful when setting this item to `FALSE` as it is possible (if the system is booted from the network) to reinstall the system without any way to prevent it (except using the console, powering off, or resetting the system). This is regardless of how you set `control_from_server` (see next configuration item).

- `control_from_server`—this configuration item is used to determine if the Ignite-UX server that Ignite-UX has connected to should control the session or not. If set to `TRUE`, you must run `ignite` (the command) on the server and then manage the client remotely (you can break out of server control on the client console).
- `use_expert_ui`—this configuration item is very useful if you are an advanced user of Ignite-UX and want to run the Ignite-UX GUI most of the time instead of using "Wizard" mode. Using the Ignite-GUI, provides a much finer level of control over the system to be installed than Wizard mode does. To make the Ignite-UX GUI the default, add the following line to the installation file system:

```
use_expert_ui = true
```

Environment variables

With the `env_vars` configuration item, you can place environment variables into the environment of the programs that are run by Ignite-UX. Setting an environment variable is done in the following way:

```
env_vars += "TZ=JST-9"
```

The preceding example sets the time zone for the installation to Japanese Standard Time.

Although this time zone does not define daylight savings, if you set a time zone that does have daylight savings you should be very careful. Take, for example, the time zone Australian Eastern time zone (EST-10EDT). Since the installation file system does not have a `tztab` file, when the time zone US EASTERN (EST5EDT) is in daylight savings it assumes EST-10EDT is also in daylight savings. In reality these time zones never actually overlap because one is in the northern hemisphere and the other in the southern.

The following environment variables are special and affect the installation or recovery process:

- `INST_NET_RESPONSE_TIMEOUT`—this variable provides safety for systems. The `inst1_adm(4)` manpage states the following information about the variable:

When booting an install client from the network, this sets the amount of time that the system will wait for a user response before it reboots in assumption that the system booted from the install server accidentally. Setting `seconds` to 0 (zero) will disable the timeout and the system will not prompt for a response.

If this value is set to >0 the following message appears:

```
Please press <return/enter> (within %d seconds) to continue loading the  
network-install utility:
```

If that times out, it is followed by the following message, so users know that the system is rebooting:

```
Unless you select <return/enter> within 10 seconds, the system will  
reboot:
```

If neither `Enter` nor `Return` is selected within 10 seconds, the system reboots.

- `INST_ALLOW_WARNINGS`—this variable again provides for the safety of systems (or in this case, disks):

Setting this environment variable is useful for non-interactive install sessions when warnings about disks containing data cause the installation to switch to interactive mode. Setting `seconds` to 1 will cause all warnings to be ignored and the installation will proceed. Setting `seconds` to greater than 1 will allow the user that many seconds to read the warning and stop the installation by pressing `<return>`.

You should be very careful about setting this variable away from the default of 0. A non-default value can be dangerous when you are doing a clean installation or cloning a system (especially

to one with the same model string⁷²). The following happens, depending on what value you set it to:

0 – A yes/no question is asked: Do you wish to cancel the non-interactive installation in order to respond to the warnings above?. There is no timeout associated with this question; Ignite-UX waits indefinitely for a response.

1 – This message appears: Continuing despite above warnings because `INST_ALLOW_WARNINGS=1`. The installation/recovery proceeds regardless.

>1 – This message appears: Press <Return/Enter> within <num> seconds to cancel batch-mode installation: . You have that number of seconds to press **Return** or **Enter** to cancel the non-interactive session to enter the user interface.

- `INST_ENABLE_NETWORK`–this environment variable is useful when you have created media from which to boot a system, but you need the installation process to start networking (for example, one of the scripts used needs to contact another host). For example, setting this variable to 1 does this:

```
env_vars += "INST_ENABLE_NETWORK=1"
```

Unless `use_dhcp` has been set to `FALSE`, the system attempts to contact a DHCP server at that point.

- `LOADFILE_RETRY_COUNT`–this has the following description in the *inst1_adm(4)* manpage:

This can be used change the default number of times that the internal loadfile command will retry a failed attempt to retrieve data from the server or media. Usually this retry mechanism is used to overcome tftp transfer problems. The default value is 5.

As an alternative, you can use the `_hp_tftp_cmds` to change the re-transmission and timeout values when tftp is actually used.

Managing configurations with `unifdef`

You may find when writing configuration files that many of them tend to end up looking the same. There are at least two programs that you can use in conjunction with configuration files to make similar files more maintainable.

The program you can use at is `unifdef` (The second program is `m4`, but how to use `m4` is not discussed.)

The following example looks at the configuration files discussed in "Configuration examples – Part B". Suppose you have one very long configuration file that you want to save into an LIF, but over time you have finally exceeded the 8KB limit on configuration data so you have a configuration file that looks like the following:

⁷² When you set `INST_ALLOW_WARNINGS` to a non-zero value and you are performing a non-interactive recovery (especially cloning to the same hardware) or installation, it is possible that the warnings can be extremely important. For example, if you are cloning between systems of the same model but the boot disks are at different hardware paths (especially on a Fibre Channel attached array) and you do not interrupt the boot process in time, the installation/recovery session may have overwritten production data or data belonging to a different system.

```

server="10.0.0.2"
netmask[ ]="255.255.255.0"
route_gateway[0]="10.0.1.1"
route_destination[0]="default"
disable_dhcp=true

( lla[0/0/0/0] == "080009654321" ) {
    ip_addr[0/0/0/0]="10.0.1.45"
}
( lla[2/0/2] == "080009654123" ) {
    ip_addr[2/0/2]="10.0.1.46"
}
( lla[1/0/0/0] == "080009654213" ) {
    ip_addr[1/0/0/0]="10.0.1.47"
}
...
( lla[0/0/0/0] == "08000A654321" ) {
    ip_addr[0/0/0/0]="10.0.2.45"
}
( lla[2/0/2] == "08000A654123" ) {
    ip_addr[2/0/2]="10.0.2.46"
}
( lla[1/0/0/0] == "08000A654213" ) {
    ip_addr[1/0/0/0]="10.0.2.47"
}

```

You will want to use `unifdef` on this file so that you can maintain it in one file but also be able to create separate files to apply to LIFs (because you want to be able to set IP addresses automatically on the systems you boot). First, though, you need to insert some other things into the file.

The `unifdef` command recognizes the following C preprocessor directives:

```

#define
#ifndef
#if
#else
#endif

```

With these directives, you can change the configuration file as follows:

```

server="10.0.0.2"
netmask[ ]="255.255.255.0"
route_gateway[0]="10.0.1.1"
route_destination[0]="default"
disable_dhcp=true

#ifdef LAN_SEGMENT_1
( lla[0/0/0/0] == "080009654321" ) {
    ip_addr[0/0/0/0]="10.0.1.45"
}
( lla[2/0/2] == "080009654123" ) {
    ip_addr[2/0/2]="10.0.1.46"
}
( lla[1/0/0/0] == "080009654213" ) {
    ip_addr[1/0/0/0]="10.0.1.47"
}
...
#endif

```

```

#define LAN_SEGMENT_2
( lla[0/0/0/0] == "08000A654321" ) {
    ip_addr[0/0/0/0]="10.0.2.45"
}
( lla[2/0/2] == "08000A654123" ) {
    ip_addr[2/0/2]="10.0.2.46"
}
( lla[1/0/0/0] == "08000A654213" ) {
    ip_addr[1/0/0/0]="10.0.2.47"
}
...
#endif

```

You can now use unifdef to create a configuration file for LAN segment one:

```

$ unifdef -t73 -DLAN_SEGMENT_1 -ULAN_SEGMENT_2 lan.cfg
server="10.0.0.2"
netmask[]="255.255.255.0"
route_gateway[0]="10.0.1.1"
route_destination[0]="default"
disable_dhcp=true

( lla[0/0/0/0] == "080009654321" ) {
    ip_addr[0/0/0/0]="10.0.1.45"
}
( lla[2/0/2] == "080009654123" ) {
    ip_addr[2/0/2]="10.0.1.46"
}
( lla[1/0/0/0] == "080009654213" ) {
    ip_addr[1/0/0/0]="10.0.1.47"
}
...

```

Note that unifdef is not a general C preprocessor; you must explicitly define or undefine variables before they are evaluated in the input file. If you require more functionality than unifdef provides, refer to *m4(1)*.

Coping with **auto_adm** and boot changes in HP-UX B.11.23

Because PA-RISC systems have two possible kernels that they may need to boot with the release of Ignite-UX C.6.0, there needs to be some way to interact with the boot loader to boot the desired kernel.

Two criteria must be met:

- The installation process must be capable of a non-interactive installation or recovery.
- The installation process must provide a mechanism for selecting a kernel version.

These changes were made to the initial system loader (ISL) to enable it to understand a new format of AUTO file (this ISL is provided as part of Ignite-UX in the LIF file

⁷³ The **-t** option treats the input as plain text rather than C style source code. Because this configuration does not look like C style source code, always use the **-t** option.

`/opt/ignite/boot/boot_lif`). This new format is capable of showing you a menu-like selection and asking you to choose a kernel to boot. Of course, to provide for non-interactive installation and recovery, there is a time-out value for a response and a default response.

With PA-RISC systems using Ignite-UX version C.6.0.x or later, an HP-UX B.11.11 kernel is used to install and recover HP-UX B.11.11 systems, and an HP-UX B.11.23 kernel is used to install and recover HP-UX B.11.23 systems. You cannot install or recover HP-UX B.11.11 systems using the HP-UX B.11.23 kernel, and the HP-UX B.11.11 kernel cannot install or recover an HP-UX B.11.23 system. There is only a 64-bit HP-UX B.11.23 PA-RISC kernel. Ignite-UX prevents you from performing an unsupported installation using an incorrect kernel during installation.

Note the following:

- Currently, `auto_adm` only affects PA-RISC systems. There is only one Itanium®-based kernel shipped with Ignite-UX; no method is required to select between it and another kernel.
- This command exists from Ignite-UX version C.6.0.x onwards.
- With Ignite-UX version C.6.0.x, the new `auto_adm` formats in the `AUTO` file provide an enhancement to the current form of the `AUTO` file; they are not a replacement.

Looking at `auto_adm`

As stated in the `auto_adm(1M)` manpage, the following forms are allowed:

```
/opt/ignite/bin/auto_adm -n -T timeout -M message -l label -c command -b  
    device -i image [-o outfile] [-O output_format] [-?]

/opt/ignite/bin/auto_adm -a -l label -c command -b device -i image  
    [-f infile] [-o outfile] [-O output_format] [-?]

/opt/ignite/bin/auto_adm -A appendfile -f infile [-o outfile]  
    [-O output_format] [-?]

/opt/ignite/bin/auto_adm -L label record_modifier [-f infile]  
    [-o outfile] [-O output_format] [-?]

/opt/ignite/bin/auto_adm -d -L label [-f infile] [-o outfile]  
    [-O output_format] [-?]

/opt/ignite/bin/auto_adm -D -L label [-f infile] [-o outfile]  
    [-O output_format] [-?]

/opt/ignite/bin/auto_adm -T timeout [-f infile] [-o outfile]  
    [-O output_format] [-?]

/opt/ignite/bin/auto_adm -M message [-f infile] [-o outfile]  
    [-O output_format] [-?]
```

ISL and CONF format data

CONF data

The `auto_adm` manpage does not correctly document all of the allowed combinations of options. For instance, the preceding forms of the command do not allow the following example:

```

# auto_adm -f /opt/ignite/boot/boot_lif
timeout = 0
default = 1
message = Choose a boot action

label = target OS is B.11.11
bootcmd = boot
boot = (:0)
image = /boot/Rel_B.11.11/INSTALL

label = target OS is B.11.23 PA
bootcmd = boot
boot = (:0)
image = /boot/Rel_B.11.23/WINSTALL

label = target OS is B.11.31 PA
bootcmd = boot
boot = (:0)
image = /boot/Rel_B.11.31/WINSTALL

```

This reads the AUTO file from the LIF and prints it in the auto_conf (CONF) format.

In the CONF format, the timeout line specifies the amount of time (in seconds) to wait before taking the default action. The default line specifies the default action to take, which, in this case, boots a kernel for HP-UX 11.11. The message line specifies the message that is displayed after booting and showing the menu options to select a target operating system to boot.

The *auto_adm(4)* manpage lists the following definitions for label, bootcmd, boot, and image:

label=string

A string label that is displayed to the user on the console describing what action is performed if this boot image is loaded. For example, "Install 64-bit HP-UX B.11.11". The string need not be enclosed in quotes.

bootcmd=cmd

The SSL command to be executed if this image is selected. Refer to *hpxx(1M)* for more information.

boot=hxpath

The hardware path relative to the boot LIF containing the desired boot image. Typically (0;). See *hpxx(1M)* for more information.

image=path

The path of the image to be booted relative to the hardware path specified in the boot directive.

The bootcmd field should match the following form of the secondary loader (see *hpxx(1M)*):

```

hpxx [-F] [-lm] [-vm] [-lq] [-a[C|R|S|D] devicefile] [-fnumber]
[-istring] [boot] [devicefile]

```

If, for example, you wanted to enable debug-level logging from boot logging automatically, you need to change the bootcmd string to the following:

```
bootcmd = -i3
```

The `auto_adm` command is intended for the management of the `AUTO` file in an LIF that is used by Ignite-UX. The `bootcmd` field only accepts one argument.

ISL data

The following example presents the same data as the `CONF` data although it is translated into ISL format:

```
# auto_adm -f ./output -o ./isl -O ISL
# cat isl
hpx KernelPrompt "Choose a boot action" 0 2
reset
"target OS is B.11.11" boot (;0)/boot/Rel_B.11.11/INSTALL
"target OS is B.11.23 PA" boot (;0)/boot/Rel_B.11.23/WINSTALL
"target OS is B.11.31 PA" boot (;0)/boot/Rel_B.11.31/WINSTALL
"Exit" reboot

# auto_adm -f ./lif -D -L "target OS is B.11.11" -O ISL
hpx KernelPrompt "Choose a boot action" 0 1
reset
"target OS is B.11.11" boot (;0)/boot/Rel_B.11.11/INSTALL
"target OS is B.11.23 PA" boot (;0)/boot/Rel_B.11.23/WINSTALL
"target OS is B.11.31 PA" boot (;0)/boot/Rel_B.11.31/WINSTALL
"Exit" reboot
```

The `auto_adm` command operates on `CONF` or `ISL` format data equally well. The default destination, if `-o` is not given, is to send the result to `stdout`.

Usage examples

Creating new files

This is the form of the `auto_adm` command used to create new files:

```
/opt/ignite/bin/auto_adm -n -T timeout -M message -l label -c command -b
device -i image [-o outfile] [-O output format] [-?]
```

The `-n` means that a new file is being created and the other options allow you to provide enough data to give the higher-level menu information and the first option on the menu. The following example creates the file:

```
# auto_adm -n -T 10 -M "Please select something to boot" \
> -l "HP-UX 11.11 32bit74" -c boot -b "(;0)" \
> -i boot/Rel_B.11.11/INSTALL -o boot.adm -O CONF
$ cat boot.adm
timeout = 10
default = 1
message = Please select something to boot

label = HP-UX 11.11 32bit75
```

⁷⁴ This format is not strictly 32-bit, as it correctly boots 32-bit and 64-bit capable systems (but not V class systems).

⁷⁵ Even though this menu choice says 32-bit, the PA-RISC primary loader (ISL) changes the `INSTALL` kernel to `WINSTALL` if the system can only run a 64-bit kernel. So, even though the one line says 32-bit, it actually supports both 32-bit and 64-bit systems.

```

bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.11/INSTALL

```

To create the same thing but on ISL format, you can use the following format of the command:

```

# auto_adm -n -T 10 -M "Please select something to boot" \
> -l "HP-UX 11.11 32bit" -c boot -b "(;0)" \
> -i boot/Rel_B.11.11/INSTALL -o boot.isl -O ISL
$ cat boot.isl
hpux KernelPrompt "Please select something to boot" 10 1
reset
"HP-UX 11.11 32bit" boot (;0)boot/Rel_B.11.11/INSTALL
"Exit" reboot

```

Adding new menu entries to a file

The following the form of the auto_adm command is used to add new menu selections into the file:

```

/opt/ignite/bin/auto_adm -a -l label -c command -b device -i image
[-f infile] [-o outfile] [-O output format] [-?]

```

First, add some entries into the auto_conf file in ISL format using auto_adm:

```

# auto_adm -a -l "HP-UX 11.11 32bit w/debug logging" -c "-i3" \
> -b "(;0)" -i boot/Rel_B.11.11/INSTALL -f boot.isl -o boot.isl2 -O ISL
# auto_adm -a -l "HP-UX 11.11 64bit" -c "boot" \
> -b "(;0)" -i boot/Rel_B.11.11/INSTALL -f boot.isl2 -o boot.isl -O ISL
$ cat boot.isl
hpux KernelPrompt "Please select something to boot" 10 1
reset
"HP-UX 11.11 64bit" boot (;0)boot/Rel_B.11.11/INSTALL
"HP-UX 11.11 32bit w/debug logging" -i3 (;0)boot/Rel_B.11.11/INSTALL
"HP-UX 11.11 32bit" boot (;0)boot/Rel_B.11.11/INSTALL
"Exit" reboot

```

You can make some similar changes in CONF format:

```

$ auto_adm -a -l "HP-UX 11.11 64bit" -c "boot" \
> -b "(;0)" -i boot/Rel_B.11.11/WINSTALL -f boot.adm -o boot.adm2 -O CONF
$ auto_adm -a -l "HP-UX 11.23 64bit" -c "boot" \
> -b "(;0)" -i boot/Rel_B.11.23/WINSTALL -f boot.adm2 -o boot.adm -O CONF
$ cat boot.adm
timeout = 10
default = 1
message = Please select something to boot

label = HP-UX 11.23 64bit
bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.23/WINSTALL

label = HP-UX 11.11 64bit
bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.11/WINSTALL

label = HP-UX 11.11 32bit

```

```
bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.11/INSTALL
```

Using an "append" file

The `auto_adm` command also enables you to take records sitting in one file and append them to an input file, combining them into one output file. The following form does this:

```
/opt/ignite/bin/auto_adm -A appendfile -f infile [-o outfile]
[-O output format] [-?]
```

For example:

```
# cat boot.append
label = HP-UX 11.11 64bit
bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.11/WINSTALL

label = HP-UX 11.11 32bit
bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.11/INSTALL
# cat boot.isl
hpx KernelPrompt "Please select something to boot" 30 1
reset
"HP-UX 11.23 64bit w/debug logging" -i3 (;0)boot/Rel_B.11.23/WINSTALL
"HP-UX 11.23 64bit" boot (;0)boot/Rel_B.11.23/WINSTALL
"Exit" reboot
# auto_adm -A boot.append -f boot.isl -o boot.combined -O ISL
# cat boot.combined
hpx KernelPrompt "Please select something to boot" 30 1
reset
"HP-UX 11.23 64bit w/debug logging" -i3 (;0)boot/Rel_B.11.23/WINSTALL
"HP-UX 11.23 64bit" boot (;0)boot/Rel_B.11.23/WINSTALL
"HP-UX 11.11 64bit" boot (;0)boot/Rel_B.11.11/WINSTALL
"HP-UX 11.11 32bit" boot (;0)boot/Rel_B.11.11/INSTALL
"Exit" reboot
```

Note that the file given with the `-A` option must be in CONF format; if you use ISL format, the command generates errors.

Updating a menu entry in a file

The following form of `auto_adm` enables you to update a boot menu entry:

```
/opt/ignite/bin/auto_adm -L label record modifier [-f infile]
[-o outfile] [-O output format] [-?]
```

The record modifier consists of the following:

record modifier consists of at least one (and optionally any combination) of `-l label`, `-c command`, `-b device`, and/or `-i image`.

This enables you to update a boot menu entry. The following examples only look at CONF format files. First change the boot field for the label HP-UX 11.23 64bit:

```
# auto_adm -L "HP-UX 11.23 64bit" -c "-i3" \
> -f boot.adm -o boot.adm2 -O CONF
$ cat boot.adm2
timeout = 10
default = 1
message = Please select something to boot

label = HP-UX 11.23 64bit
bootcmd = boot
boot = -i3
image = boot/Rel_B.11.23/WINSTALL

label = HP-UX 11.11 64bit
bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.11/WINSTALL
```

Rename the label "HP-UX 11.23 64bit" to "HP-UX 11.23 64bit (PA-RISC)":

```
$ auto_adm -L "HP-UX 11.23 64bit" -l "HP-UX 11.23 64bit (PA-RISC)" \
> -f boot.adm2 -o boot.adm -O CONF
$ cat boot.adm
timeout = 10
default = 1
message = Please select something to boot

label = HP-UX 11.23 64bit (PA-RISC)
bootcmd = -i3
boot = (;0)
image = boot/Rel_B.11.23/WINSTALL

label = HP-UX 11.11 64bit
bootcmd = boot
boot = (;0)
image = boot/Rel_B.11.11/WINSTALL
```

Deleting a menu entry from a file

The boot menu entries are removed from a file using the following form of auto_adm:

```
/opt/ignite/bin/auto_adm -d -L label [-f infile] [-o outfile]
[-O output format] [-?]
```

To remove a boot menu entry, you simply need to provide the label name from the file:

```
$ auto_adm -L "HP-UX 11.23 64bit (PA-RISC)" -d -f boot.adm -o boot.adm2 -O
CONF
$ cat boot.adm2
timeout = 10
default = 1
message = Please select something to boot
label = HP-UX 11.11 64bit
bootcmd = boot
boot = (;0)
```

```
image = boot/Rel_B.11.11/WINSTALL
```

Changing the default menu choice

Using `auto_adm`, you can also change the default menu choice with the following form of the command:

```
/opt/ignite/bin/auto_adm -D -L label [-f infile] [-o outfile]  
[-O output format] [-?]
```

The following example shows the change in default menu choices from the first to the second option:

```
$ cat boot.adm  
hpx KernelPrompt "Please select something to boot" 10 1  
reset  
"HP-UX 11.23 64bit (PA-RISC)" -i3 (;0)boot/Rel_B.11.23/WINSTALL  
"HP-UX 11.11 64bit" boot (;0)boot/Rel_B.11.11/WINSTALL  
"Exit" reboot  
$ auto_adm -D -L "HP-UX 11.11 64bit" -f boot.adm -o boot.adm2 -O ISL  
$ cat boot.adm2  
hpx KernelPrompt "Please select something to boot" 10 2  
reset  
"HP-UX 11.23 64bit (PA-RISC)" -i3 (;0)boot/Rel_B.11.23/WINSTALL  
"HP-UX 11.11 64bit" boot (;0)boot/Rel_B.11.11/WINSTALL  
"Exit" reboot
```

Changing the timeout

The timeout value is used by the boot loader to determine how long to wait for user input before executing the default menu choice (see the next section for more information about how to set the default menu choice).

```
/opt/ignite/bin/auto_adm -T timeout [-f infile] [-o outfile]  
[-O output format] [-?]
```

In the following example, you can see the timeout change from 10 to 30 seconds (the second last field on the first line) for this file in ISL format:

```
$ cat boot.isl2  
hpx KernelPrompt "Please select something to boot" 10 1  
reset  
"HP-UX 11.11 32bit w/debug logging" -i3 (;0)boot/Rel_B.11.11/INSTALL  
"HP-UX 11.11 32bit" boot (;0)boot/Rel_B.11.11/INSTALL  
"Exit" reboot  
$ auto_adm -T 30 -f boot.isl2 -o boot.isl -O ISL  
$ cat boot.isl  
hpx KernelPrompt "Please select something to boot" 30 1  
reset  
"HP-UX 11.11 32bit w/debug logging" -i3 (;0)boot/Rel_B.11.11/INSTALL  
"HP-UX 11.11 32bit" boot (;0)boot/Rel_B.11.11/INSTALL  
"Exit" reboot
```

Updating the prompt message

The prompt message is what the boot loader prompts you with when you are asked to make a selection from the available options. The form of the command that enables you to change the prompt message is as follows:

```
/opt/ignite/bin/auto_adm -M message [-f infile] [-o outfile]  
[-O output format] [-?]
```

The following example uses this form to change the prompt message:

```
# auto_adm -M "Please select an option" -f boot.combined -o boot.final -O ISL  
# cat boot.final  
hpx KernelPrompt "Please select an option" 30 1  
reset  
"HP-UX 11.11 64bit" boot (;0)boot/Rel_B.11.11/WINSTALL  
"HP-UX 11.11 32bit" boot (;0)boot/Rel_B.11.11/INSTALL  
"Exit" reboot
```

Into and out of an LIF file

The `auto_adm` command automatically recognizes an LIF file when provided as the argument to `-f` or `-o` options and handles the format of the data correctly. In the following example, the data is extracted from the boot LIF provided with Ignite-UX and at the same time changes the timeout from 120 seconds to 30 seconds. You can compare the data with `lifcp`:

```
$ auto_adm -T 30 -f ./boot_lif -o boot.tmp -O ISL  
$ lifcp ./boot_lif:AUTO -  
hpx KernelPrompt "Choose Operating System to Install :" 120 1  
reset  
"target OS is B.11.11" boot (;0)/boot/Rel_B.11.11/INSTALL  
"target OS is B.11.23 PA" boot (;0)/boot/Rel_B.11.23/WINSTALL  
"Exit" reboot  
$ cat boot.tmp  
hpx KernelPrompt "Choose a boot action" 30 1  
reset  
"target OS is B.11.11" boot (;0)/boot/Rel_B.11.11/INSTALL  
"target OS is B.11.23 PA" boot (;0)/boot/Rel_B.11.23/WINSTALL  
"Exit" reboot
```

You can put the data back into an LIF by specifying the LIF as the output file (it is placed into the `AUTO` file within the LIF).

Installation configurations using Software Distributor depots

The following is an example process for configuring support on an Ignite-UX server to allow you to install systems using Software Distributor (SD) depots as the repository for the software.

It is important to note the following:

- How to configure `inst1_bootd` or DHCP servers to allow clients to boot over the network is not discussed. For more information, refer to the *Ignite-UX Administration Guide* and *inst1_bootd(1M)*.
- How to setup depots up remote to the Ignite-UX server is not discussed; the simplest location for depots is on the Ignite-UX server.
- How to use this process with versions earlier than the C.6.0.x version of Ignite-UX is not discussed. If you are using a version of Ignite-UX and some commands do not work as shown, you should consider updating your Ignite-UX.

You should review the entire section before following any of the instructions because consequences and issues related to running commands are discussed throughout.

Getting started

The first step is to choose the HP-UX 11i v1 Operating Environment (OE) media that you want to use. To make an informed decision you should review the *Successful System Cloning using Ignite-UX* White Paper. The issues that this paper discusses are relevant to installing a system (particularly a new system) because you must choose the OE that provides support for all of the hardware in the system.

Creating the core operating system depot

To create the core operating system depot the `make_depots` command is used. You will be copying the December 2003 HP-UX 11i Version 1 OE media into a core operating system depot in this example. You will use the defaults for the `make_depots` command; you only need to give the block device⁷⁶ of the DVD drive to define the location of the final depot and the release you will be copying as follows:

```
make_depots -r B.11.11 -s /dev/dsk/c0t1d0
```

If the command is successful, no output is displayed by `make_depots`.

Important:

To successfully copy the contents of the core operating system DVD⁷⁷ to a depot on the Ignite-UX server, you must have the patches that enable the Rock Ridge extensions natively on HP-UX installed on the Ignite-UX server. Also, to successfully copy the contents of the HP-UX 11i v2 May 2005 or later OE media on an 11i v1 Ignite-UX server, you must upgrade to Ignite-UX version C.6.2 and have the patches that enable the Rock Ridge extensions natively on HP-UX installed on the Ignite-UX server. On an 11i Version 2 Ignite-UX server, you need only to upgrade to Ignite-UX version C.6.2.

The system on which this example was developed (running HP-UX B.11.11) required these patches:

⁷⁶ The manpage for `make_depots` indicates that a DVD/CD-ROM drive must be given using the block device. A character device would be assumed to be a tape dive containing a serial core operating system depot on a tape. Also you can give a normal SD depot in the form <host>:<path>; a regular file name would be considered a tape depot on disk.

⁷⁷ HP-UX 11i v1 is delivered on both CD and DVD. HP-UX 11i v2 is only delivered on DVD.

PHCO_25841	1.0	Add Rock Ridge extension to mount_cdfs(1M)
PHKL_26269	1.0	Rock Ridge extension for ISO-9660
PHKL_28025	1.0	Rock Ridge extension for ISO-9660

There are no patches required for HP-UX B.11.23.

The `make_depots` command above assumes that there was no previous depot created at `/var/opt/ignite/depots/Rel_B.11.11/core`. If such a depot exists, you should use the `-d` option of `make_depots` to have the new depot created in a new location. If you need to maintain multiple core operating system (OE) depots for a release of HP-UX, you should consider placing them into depots that indicate what version of the media the depot contains for easy identification. For example, if you are creating depots for the December 2003 and June 2004 HP-UX 11i v1 Mission Critical OE media you could use the following commands:

```
make_depots -r B.11.11 -s /dev/dsk/c0t1d0 -d \
/var/opt/ignite/depots/Rel_B.11.11/core_1203_mc
make_depots -r B.11.11 -s /dev/dsk/c0t1d0 -d \
/var/opt/ignite/depots/Rel_B.11.11/core_0604_mc
```

Important:

If you need to install different OEs from your Ignite-UX server, you must ensure that the only the OEs that the client system is licensed to run are installed onto the client system.

Alternatively, you could use `swremove` on the depot⁷⁸ to remove all of the software it contains, or use the following commands to remove the contents before running the previous `make_depots` commands:

```
# swreg -u -l depot /var/opt/ignite/depots/Rel_B.11.11/core
# rm -rf /var/opt/ignite/depots/Rel_B.11.11/core
```

If everything is successful, you will have a depot at the following location containing the OE media for HP-UX B.11.11:

```
swlist -d -s /var/opt/ignite/depots/Rel_B.11.11/core

# Initializing...
# Contacting target "test"...
#
# Target: test:/var/opt/ignite/depots/Rel_B.11.11/core
#
#
# Bundle(s):
#
100BaseT-00          B.11.11.01      EISA 100BaseT;Supptd W=A4308B;SW=J2780BA
```

⁷⁸ Be extremely careful if you use `swremove` with a wildcard. Make sure you also use the `-d` option to tell `swremove` you will be operating on a depot, otherwise you may end up operating on the software installed on the local system.

```

100BaseT-01          B.11.11.01      HP-PB 100BaseT;Supptd HW=A3495A;SW=J2759BA
...
HPUXBase32          B.11.11          HP-UX 32-bit Base OS
HPUXBase64          B.11.11          HP-UX 64-bit Base OS
HPUXBaseAux         B.11.11.0312    HP-UX Base OS Auxiliary
HWEnable11i          B.11.11.0312.4  Hardware Enablement Patches for HP-UX 11i
v1, December 2003
...
hpuxwsWebmin        A.1.0.05.02    HP-UX Webmin-based Admin
hpuxwsXml           A.1.0.01.02    HP-UX XML Web Server Tools
perl                B.5.6.1.F       Perl Programming Language
scsiU320-00          B.11.11.00      PCI SCSI U320; Supptd HW=A7173A

```

The log file that `make_depot`s creates contains more detailed information. By default, the log file is in `/var/opt/ignite/logs/make_depot`.

Operating Environments containing multiple CDs or DVDs

HP-UX 11i v1 has been available on multiple CDs for a long time; as of March 2006 the 11i v2 OE media now comes on two DVDs. To copy those DVDs into the same depot, you do the same thing you would to copy 11i v1 from multiple CDs – run `make_depot`s multiple times.

Perform the following steps to copy 11i v2 from multiple DVDs:

1. Insert the first DVD from the March 2006 11.23 OE.
2. Run the following command:

```
make_depot -r B.11.23 -s /dev/dsk/c0t1d0 -d \
/var/opt/ignite/depots/Rel_B.11.23/core_Mar2006_mc
```

3. Remove the first DVD from the DVD drive.
4. Insert the second DVD from the March 2006 11.23 OE.
5. Run the following command:

```
make_depot -r B.11.23 -s /dev/dsk/c0t1d0 -d \
/var/opt/ignite/depots/Rel_B.11.23/core_Mar2006_mc
```

6. Remove the second DVD.

The same steps should be followed for 11i v1 OEs on CD: for each CD run `make_depot`s and change the media between each run of `make_depot`. An example of copying the September 2005 11i v1 OE media on CD to a depot follows:

1. Insert the first CD from the September 2005 11.11 OE.
2. Run the following command:

```
make_depot -r B.11.11 -s /dev/dsk/c0t1d0 -d \
/var/opt/ignite/depots/Rel_B.11.11/core_Sep2005_mc
```

3. Remove the first CD.
4. Insert the second CD from the September 2005 11.11 OE.

5. Run the following command:

```
make_depots -r B.11.11 -s /dev/dsk/c0t1d0 -d \  
/var/opt/ignite/depots/Rel_B.11.11/core_Sep2005_mc
```

6. Remove the second CD.
7. Insert the third CD from the September 2005 11.11 OE.
8. Run the following command:

```
make_depots -r B.11.11 -s /dev/dsk/c0t1d0 -d \  
/var/opt/ignite/depots/Rel_B.11.11/core_Sep2005_mc
```

9. Remove the third CD.
10. Insert the fourth CD from the September 2005 11.11 OE.
11. Run the following command:

```
make_depots -r B.11.11 -s /dev/dsk/c0t1d0 -d \  
/var/opt/ignite/depots/Rel_B.11.11/core_Sep2005_mc
```

12. Remove the fourth CD.

At this point you still need to run `make_config` on the depot you have created. See the section “Creating the configuration file to describe the depot” on page 173 for more information on how to do this.

Other methods for creating the Core OS depot

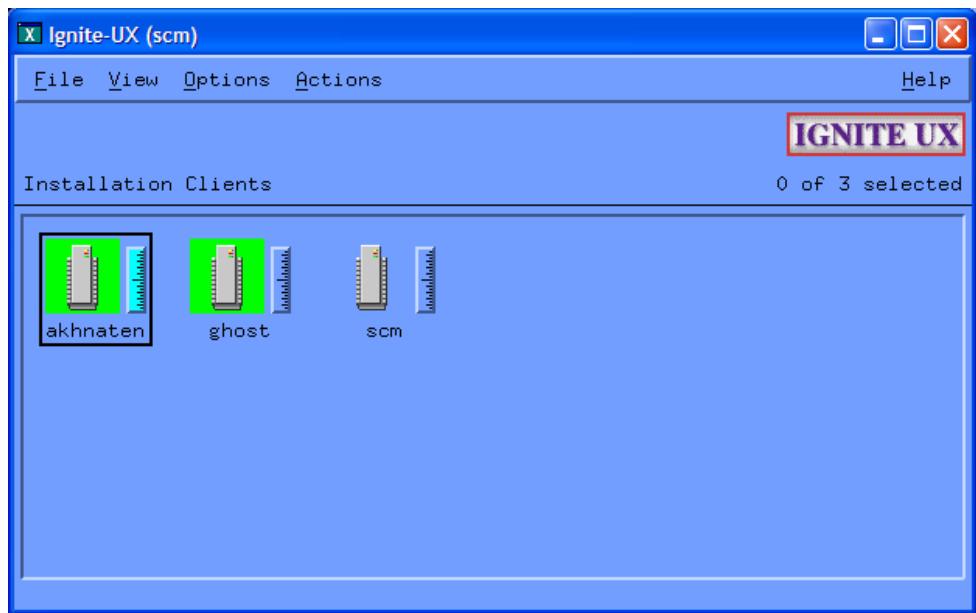
There are other methods for creating the Core OS depot⁷⁹. Some better than others and some having unique pitfalls.

Using the Ignite-UX GUI

You can use the Ignite-UX GUI on your Ignite-UX server to create a core operating system (OE) depot using the following process:

1. On your Ignite-UX server, enter `/opt/ignite/bin/ignite`.

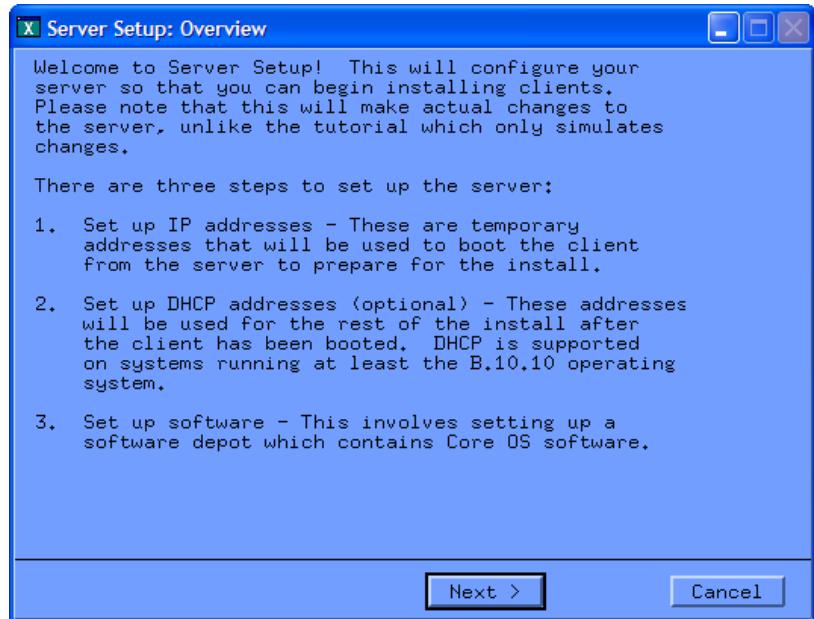
⁷⁹ Note that `make_depots` and the option in the Ignite-UX GUI to setup software depots is not just for core operating system (OE) media, you can copy other media using this method such as the applications CD/DVD as well.



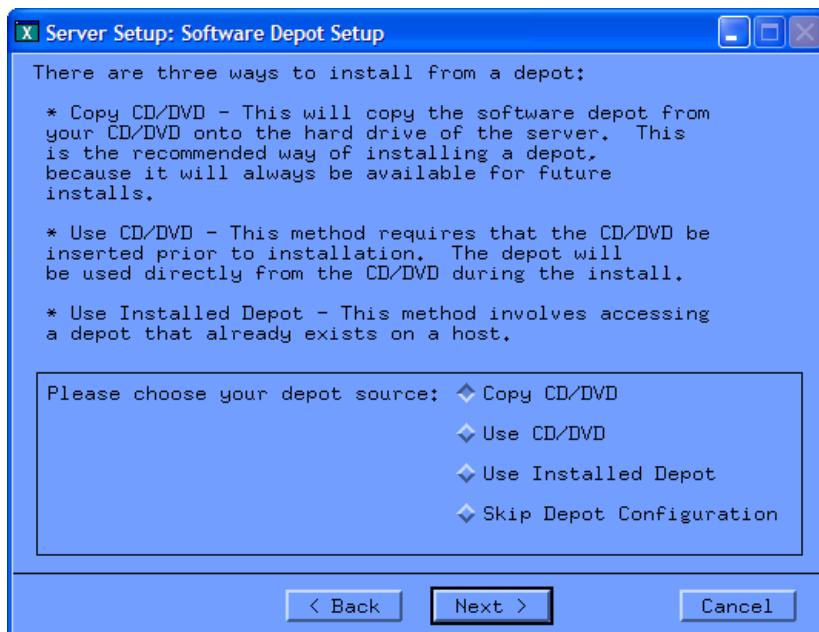
2. Select **Run Tutorial/Server Setup...** from the **Actions** menu.



3. Click the **Server Setup** button.

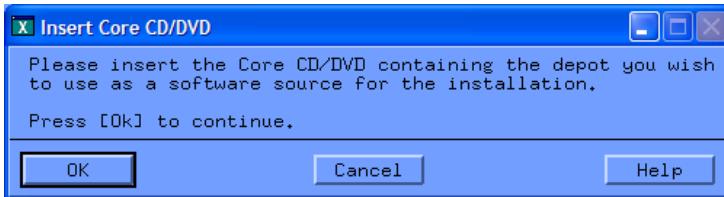


4. The third step, **Set up Software**, allows you to create a core operating system depot. Click the **Next** button in this and the **Next** or **OK** buttons on the following dialog boxes until the following dialog box appears.



5. Select **Copy CD/DVD** using the adjacent radio button then click **Next** to copy the core operating system (OE) media into a depot. Note the **Use CD/DVD** option, which allows you use a directly mounted DVD rather than have it copied to another depot for usage. If you had multiple DVD drives attached to your Ignite-UX server, you could mount all of the depots directly from their DVDs and save disk space.⁸⁰

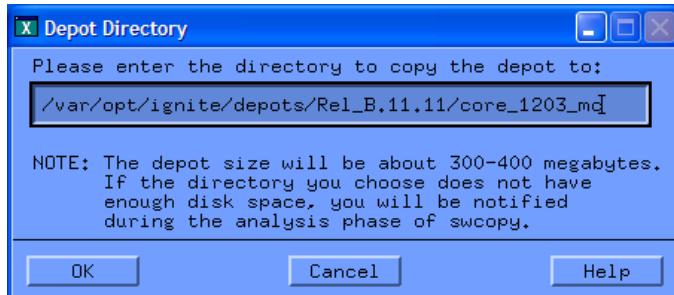
⁸⁰ If you manually mount a DVD, you can run `make_config` on the depot and then create a `cfg` clause referencing the configuration file. You *must* register the depot on the DVD using `swreg` to avoid errors. When remote systems attempt to access the mounted DVD as a depot



6. Insert the media as requested and click **OK**.



7. Select the correct CD or DVD device from the drop-down button then click **OK**.



8. Enter the location and name of the depot that the core operating system (OE) to which the media will be copied then click **OK**.

The size displayed in the dialog box may not be correct as it is a static message and the 3-400Mb is not based upon the size of the depot on the CD or DVD.



9. Review the information you have entered to be sure it is correct and then click **OK**. A status dialog box appears so that you can monitor the software copying.

they may be denied access to it. To determine whether a depot is registered or not, use `swlist -l depot`, which displays the path to the depot if it is registered.

x Server Setup: Logfile

server_setup Automatic Scrolling

```
# Ignite-UX Server Setup Logfile.

*****EXECUTING: /opt/ignite/bin/make_depots -s /iuxcdrom0 -d /var/opt/ignite/depots/Rel_B.11.11/core_1203_mc

===== 10/09/04 00:39:49 EST Beginning make_depots

NOTE: Checking SD tools version and layout_version of /iuxcdrom0
NOTE: Obtaining bundle information from /iuxcdrom0
NOTE: Processing bundles found in depot
NOTE: Beginning execution of swcopy commands
NOTE: Beginning swcopy from /iuxcdrom0 to
      /var/opt/ignite/depots/Rel_B.11.11/core_1203_mc
```

OK Help

You should not click **OK** in the status dialog box until the process of copying the CD/DVD is complete as an error results.

- Once complete click **OK** and a log file dialog box similar to the following appears:

x Server Setup: Logfile

server_setup Automatic Scrolling

```
/var/opt/ignite/depots/Rel_B.11.11/core_1203_mc

===== 10/09/04 01:17:27 EST Ending make_depots

SUCCESS: The command "/opt/ignite/bin/make_depots -s /iuxcdrom0 -d /var/opt/ignite/depots/Rel_B.11.11/core_1203_mc" completed successfully.

*****EXECUTING: /opt/ignite/bin/make_config -c /var/opt/ignite/config

NOTE: make_config can sometimes take a long time to complete. Please be patient!
SUCCESS: The command "/opt/ignite/bin/make_config -c /var/opt/ignite/config" completed successfully.

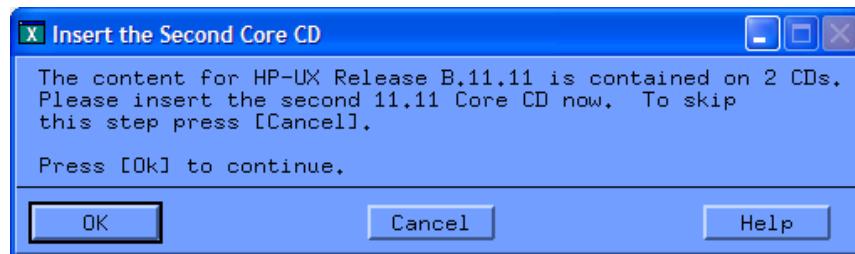
*****EXECUTING: /opt/ignite/bin/manage_index -a -f /var/opt/ignite/index

SUCCESS: The command "/opt/ignite/bin/manage_index -a -f /var/opt/ignite/index" completed successfully.

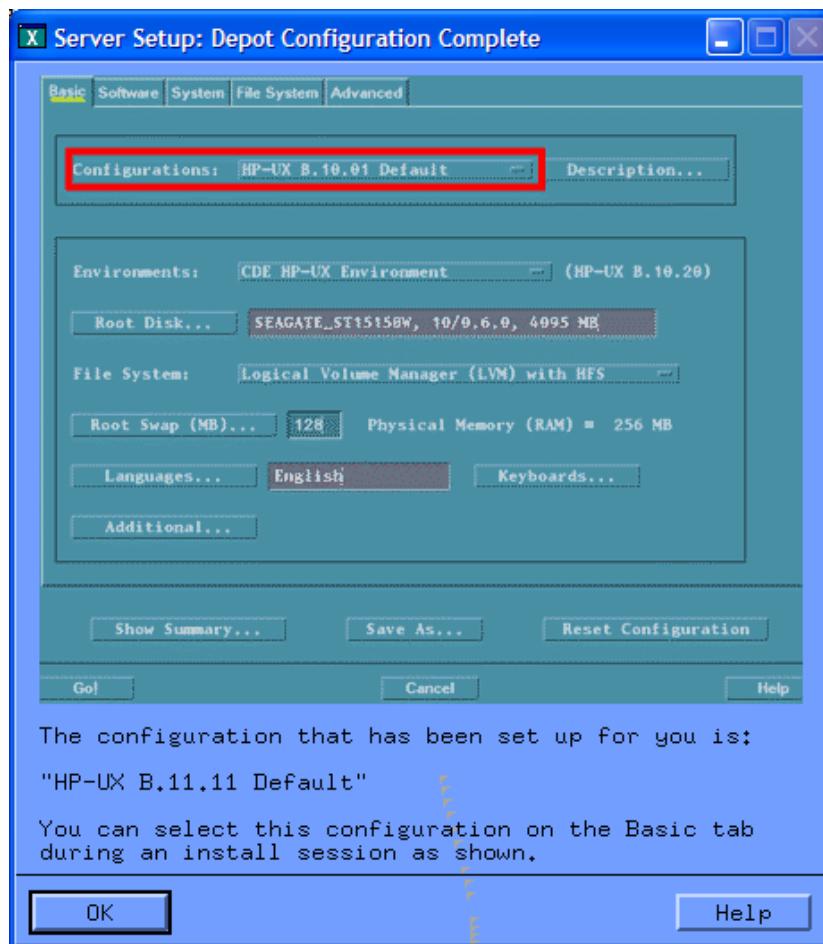
*****SETUP is complete! Press [OK] to continue.
```

OK Help

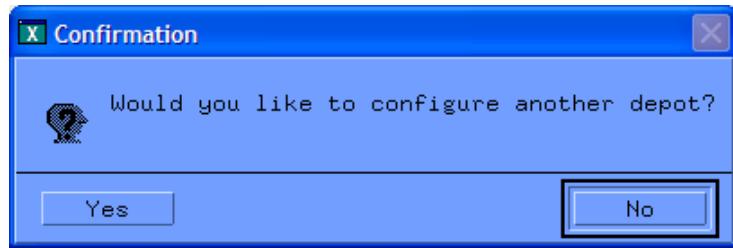
Once the process is completes successfully: the depot exists, a configuration file is created for it, and the configuration file is added to the appropriate default configuration clause /var/opt/ignite/INDEX. If you do not want the configuration file added to the cfg clauses in the INDEX file that it has been added to, you must remove it from each clause using manage_index. Allowing manage_index to update the /var/opt/ignite/INDEX file indiscriminately may have unintended side affects so you should review the INDEX file to ensure that it remains correct for your environment.



11. If the OE you are using contains multiple CDs or DVDs, change the CD/DVD and click **OK**. If you are using a single CD or DVD, or you have completed copying all the media from the OE, continue by clicking **Cancel** and then **OK** in the next dialog box. The following dialog box appears:



12. Click **OK**.



13. Click **Yes** to setup more depots or **No** to finish the process.



14. Click **Finish** to complete the process.

The pitfalls of using multiple media

If you are creating a depot, you should be very careful if the OE media contains more than one CD or DVD. If the OE contains multiple CDs or DVDs, it will be indicated on each CD or DVD. For example, the 11i v2 March 2006 Mission Critical DVDs will have "1 of 2" printed on the first DVD, and "2 of 2" printed on the second DVD.

The `make_depot` command is unable to determine the number of CDs or DVDs included in a media set, so it will not prompt for additional media. Only the first CD or DVD will be copied into the depot. It is safe to run the `make_depot` command for each of the CDs or DVDs in the OE to copy them all into the depot; do not copy them into different depots.

Creating the configuration file to describe the depot

The `make_config` command is used to create the configuration file that describes the core operating system (OE) depot you just created. The easiest way to do this since the depot is local and in `/var/opt/ignite/depots/Rel_B.11.11` is as follows:

```
make_config -s /var/opt/ignite/depots/Rel_B.11.11/core \
```

```
-c /var/opt/ignite/data/Rel_B.11.11/core_cfg
```

Warning:

HP does not support placing different versions of an OE in one depot.
Instead, you must keep different OEs in different depots.

This creates the configuration file that describes the depot in
`/var/opt/ignite/data/Rel_B.11.11` and it is named `core_cfg`⁸¹. However, it does not
update `/var/opt/ignite/INDEX` so that you can create your own custom `cfg` clause later. If
everything completes successfully, you should see the following message only:

NOTE: `make_config` can sometimes take a long time to complete. Please be patient!

An alternative to this is to use the following command:

```
make_config -r B.11.11
```

This command automatically creates configuration files for all of the depots under
`/var/opt/ignite/depots/Rel_B.11.11`. However, this form of the command will also
update the `/var/opt/ignite/INDEX` file (typically updating all `cfg` clauses affecting the given
HP-UX release with information about the configuration files). This may not be desirable.

For example, if you run the following command:

```
make_config -r B.11.11 -s core_1203_mc
```

Then the configuration file will be created and the `/var/opt/ignite/INDEX` file will be updated
(only showing the Default B.11.11 clause) to be:

```
# cat /var/opt/ignite/INDEX
...
cfg "HP-UX B.11.11 Default" {
    description "This selection supplies the default system configuration
that HP supplies for the B.11.11 release."
    "/opt/ignite/data/Rel_B.11.11/config"
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/core_1203_mc_cfg"
    "/var/opt/ignite/config.local"
}
```

If you then run the following command:

```
make_config -r B.11.11
```

⁸¹ This is why it was suggested earlier that if you have multiple core operating system (OE) depots, you should give them names that include the release date of the media and an indication (if required) of the OE type.

In addition to creating the configuration file, the /var/opt/ignite/INDEX file will be updated as in the following excerpt of the file:

```
cfg "HP-UX B.11.11 Default" {
    description "This selection supplies the default system configuration
that HP supplies for the B.11.11 release."
    "/opt/ignite/data/Rel_B.11.11/config"
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/core_1203_mc_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/core_cfg"
    "/var/opt/ignite/config.local"
}
```

You should avoid multiple core operating system (OE) definitions within the one cfg clause. The contents of the second configuration file defining a core operating system (OE) depot override the configuration in the first.

This command does not cause any additional changes to occur and make_config would only add the core_cfg configuration file to every B.11.11 clause in the INDEX file.

```
# make_config -r B.11.11 -s core
```

For information about the contents of core operating system (OE) configuration file, see the "Core operating system depot configuration" section.

Creating a minimalist cfg clause for installation

You can create a cfg clause in /var/opt/ignite/INDEX that will allow you to install a system.

```
manage_index -n "HP-UX B.11.11 Default" -c \
    "Custom B.11.11 installation using SD"
manage_index -a -f /var/opt/ignite/data/Rel_B.11.11/core_cfg \
    -c "Custom B.11.11 installation using SD"
```

You use manage_index to avoid using the variations of make_config that indiscriminately change cfg clauses in the INDEX file. In this case, you are copying the default B.11.11 clause (the changes from make_config in the previous section have been undone prior to doing this step).

When complete you have a new cfg clause in the INDEX file:

```
# cat /var/opt/ignite/INDEX
...
cfg "Custom B.11.11 installation using SD" {
    description "This selection supplies the default system configuration
that HP supplies for the B.11.11 release."
    "/opt/ignite/data/Rel_B.11.11/config"
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/core_cfg"
    "/var/opt/ignite/config.local"
}
```

Comments on the SD based `cfg` clause

The clause created does nothing except install a generic core operating system (OE) from an SD depot (with the patches included with the operating system). The important thing is that it will install HP-UX. To enable this all you have had to do is create a core operating system (OE) depot, run `make_config` once, and then `manage_index` twice.

In the following sections, you will add to this by adding new depots to allow additional applications to be installed and custom configuration to be executed.

You have saved a lot of time by using the default configuration files that Ignite-UX supplies for every HP-UX revision (the `config` and `hw_patches_cfg` files) as a basic configuration for file systems and some additional configuration for patching. The release-specific (in this case `/opt/ignite/data/Rel_B.11.11/config`) configuration file supplies some important variable definitions. Including this file prevents changes to Ignite-UX in the future from impacting you because you automatically receive changes to the release-specific `config` and `hw_patches_cfg` files as with each new release of Ignite-UX.

How to setup your environment to boot systems so that they can contact your Ignite-UX server to allow you to install systems is not discussed in this paper. For more information on completing the setup of your server, refer to the *Ignite-UX Administration Guide*.

Extending the generic SD `cfg` clause with applications

In this section, you will enhance the core operating system (OE) bundle that you previously created by making available some of the applications from the HP Application Media. In this example, you will be copying the December 2003 Application DVD. The application media is created and copied in much the same way as "Creating the core operating system depot".

Creating the application depot

First, you must copy the applications from the DVD into a depot so that you can use them. The following command begins the process:

```
# make_depots -d /var/opt/ignite/depots/Rel_B.11.11/apps_1203 \
-s /dev/dsk/c0t1d0
```

If the media contains codeword-protected applications, so you must supply the codewords directly to the `make_depots` command (for more information, see `make_depots(1M)`). This command can take a long time to complete.

Creating a configuration for the application depot

Since you created the depot, you can create the configuration file for the depot using the following command:

```
# make_config -s /var/opt/ignite/depots/Rel_B.11.11/apps_1203 \
-c /var/opt/ignite/data/Rel_B.11.11/apps_1203_cfg
```

Remember that you do not use the `-r` option so that `manage_index` is run by `make_config` to update every `cfg` clause that uses HP-UX B.11.11 to include the file. You should only add the `apps_1203_cfg` file to one `cfg` clause.

After you run the following command to add `apps_1203.cfg` to the `cfg` clause, you can install any applications that are in the depot:

```
# manage_index -a -f /var/opt/ignite/data/Rel_B.11.11/apps_1203.cfg \
-c "Custom B.11.11 installation using SD"
```

An example problem

It is often helpful to examine a real life problem. The previous process created a problem by adding the applications configuration file to the same `cfg` clause as the core operating system (OE) depot. This issue was touched on previously and this section describes it in further detail.

In `/var/opt/ignite/data/Rel_B.11.11/core_cfg` is the following definition of Perl:

```
init sw_sel "perl" {
    description = "Perl Programming Language"
    sw_source = "core"
    sw_category = "OrderedApps"
    sd_software_list = "perl,r=B.5.6.1.F,a=HP-UX_B.11.11_32/64,v=HP"
    impacts = "/opt" 68657Kb
} = TRUE
```

In `/var/opt/ignite/data/Rel_B.11.11/apps_1203.cfg` is this definition of Perl⁸²:

```
init sw_sel "perl" {
    description = "Perl Programming Language"
    sw_source = "/var/opt/ignite/depots/Rel_B.11.11/apps_1203"
    sw_category = "OrderedApps"
    sd_software_list = "perl,r=B.5.6.1.F,a=HP-UX_B.11.11_32/64,v=HP"
    impacts = "/opt" 68657Kb
} = TRUE
```

The two previous examples have the same `sw_sel` name so they are the same. When the second `perl` `sw_sel` clause is parsed, Ignite-UX assumes that it is an update to the existing definition of Perl already found because the clauses are named identically. The attributes of the `sw_sel`, such as the `sw_source`, are updated by the second definition, which replaces the first definition of Perl.

Perl is an always-loaded product; by default, it is marked to be loaded. There is the possibility that this could cause a problem if the application media and the core operating system (OE) media do not have the same release date. If you have newer core operating system (OE) media and a much older application depot being referenced, problems like the following could occur:

- An older version of Ignite-UX is referenced in the applications depot and it may not support the system being installed. For example, if you create a recovery tape and the installation kernel on the tape does not support the current system it may panic on first boot when you need to recover the system.
- The kernel drivers in the application depot are also located in the core operating system (OE) depot including SCSI, Fibre Channel, and Gigabit Ethernet drivers. If you are referencing older versions of the drivers and the system is newer, you may have hardware in the system that is not supported by the older drivers and this causes problems.

⁸² Perl is not the only product affected by the issue. There are other products that the core operating system (OE) media and the applications media have in common.

There are four solutions to this problem:

- Edit the `apps_1203.cfg` file and change the name of the `sw_sel` clauses so that they are unique and descriptive of the clause's content. This allows you to address the two software selection clauses uniquely. It also ensures that the intended software is selected from the core operating system (OE) depot, and that software from the applications media is not selected inadvertently. This is time consuming and if you want to update the configuration file using `make_config` you must reapply the changes.
- Remove all of the software bundles that exist in the core operating system (OE) depot and the applications depot. You would run `make_depots` on the applications DVD then use `swremove` on the applications depot. HP does not recommend this solution.
- Place the application media into the core operating system (OE) depot rather than putting it into a separate depot. This means that you must manage the depots as one. Note that you cannot update an OE in a depot. If you wish to use a new OE in an existing depot, you must remove the entire contents of the depot before using `make_depots` to copy in new software from OE and applications media. This also assumes that a different version of VxVM/VxFs is not available on the application media - you can't cold install a different version of VxVM/VxFs to the version available from the core OE media.
- Ensure that the Core OS (OE) media and applications media are always synchronized.

Follow on consequences

This can also lead to more serious problems when you have an application that depends on something that would normally be installed from the core OE depot, but because of the issues described above in "An example problem" it may be installed from somewhere else (the applications depot).

We will use the following `cfg` clause as the basis for our example:

```
cfg "HP-UX B.11.11" {
    description "HP-UX 11.11 MCOE with vPars"
    "/opt/ignite/data/Rel_B.11.11/config"
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/core_OE.cfg"
    "/var/opt/ignite/data/Rel_B.11.11/vPars_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/apps.cfg"
    "/var/opt/ignite/config.local"
}
```

Some vPars filesets in version A.03.03.06 have an SD dependency on the fileset `PartitionManager.PARMGR,r>=B.11.11.01.05`. That means that the fileset `PartitionManager.PARMGR` of that revision or higher must be installed on the system otherwise parts of the vPars software won't install⁸³.

If the bundle that defines the Partition Manager commands is present in the core OE and the applications depot, it will be installed from the applications depot not from the OE depot. That means if the vPars software is installed after the OE depot but before the applications depot, the vPars software will only partially install because some dependencies cannot be met.

⁸³ If available from the same depot as the vPars software, it would be automatically selected but may only end up partially installed (since only the dependencies would be selected). If you didn't keep the Partition Manager commands in sync with the application and core OE depots, you might end up with mixed revision filesets in the Partition Manager commands bundle.

As in the previous section, there are several ways to deal with this:

- Always place the applications depot into the core OE depot. This also assumes a different version of VxVM/VxFs is not available on the application media - you can't cold install a different version of VxVM/VxFs to the version available from the core OE media.
- Use the `load_order` keyword to change the load order of the depots so that the applications depot is always loaded after the core OE depot and before any other depots. The core OE depot is always loaded at load order 0, and by default other depots are loaded at load order 5. Change the applications depot so it has a load order of 1 to 4 to ensure it is loaded before the vPars software.

It is important to understand the impacts of the choices you make when deciding how your depots and configuration will be laid out.

Other software could be affected in similar ways to this vPars example.

Resolving applications that are not in SD format

There are two solutions when you find that you have an application that is not in SD format:

- Package an application in SD format
- Define an application in a non-SD format

The following topics describe both solutions.

Package an application in SD format

You could learn how to write an SD product specification file (psf) for use with `swpackage` to produce an SD depot of your application. Learning how to write a PSF can be difficult and time consuming. For information regarding the creation of PSF files, see *Software Distributor Administration Guide for HP-UX 11i*.

Alternatively, you can use the HP Software Package Builder (SPB) product makes it easy to package applications into SD format. SPB provides a window into the software package structure, showing attributes that can be set for each package element.

You can find SPB product information and download it from the HP Software Depot at:

<http://software.hp.com/portal/swdepot/displayProductInfo.do?productNumber=SwPkgBuilder>

Once you have the software in SD format, using either solution, you can:

1. copy it into a depot;
2. run `make_config` to generate a config file;
3. then run `manage_index` to add the configuration file to the desired `cfg` clause in `/var/opt/ignite/INDEX`.

Important:

SD depots must contain software in bundles after you have packaged the software. If it is not packaged at the bundle level, you will have to create a bundle for it with the `make_bundles` command.

Define an application in a non-SD format

You either have the option of providing applications in non-SD format in `tar` or `cpio` format archives that have been gzipped or compressed. This section explains how to write a configuration file that describes an application in non-SD format⁸⁴.

The archive that you will be installing the application with is assumed to contain no system specific configuration.

You should configure the application as necessary with the command and script hooks provided by Ignite-UX. The scripts that may be executed while installing the application should *never* query the user because not all installations are attended. For a non-interactive installation (the user does not have access to the console), a script that seeks input from the console would effectively hang the installation process.

Writing the configuration file for an application in non-SD format

In the following example, the application configuration file that this configuration will be placed in is `/var/opt/ignite/data/Rel_B.11.11/myapps_cfg`.

```
#  
# Definition of where my application is located  
#  
sw_source "my application archive" {  
    description = "my application"  
    source_format = archive  
    source_type="NET"  
    nfs_source = "10.2.72.150:/var/opt/ignite/My_Applications"  
}
```

Next, you define a `sw_source` for the application including its location on the NFS server (at IP address 10.2.72.150) in the `/var/opt/ignite/My_Applications` directory. Remember, this directory must be exported from the NFS server so that it can be mounted. If you choose to limit access to this directory, remember that you may want servers that obtain IP address via DHCP to be able to mount the directory when setting its permissions.

```
sw_category "SiteApps" {  
    description = "My site applications"  
}
```

The previous example defines a category that you can easily pick out in the list of applications in the **Software** tab of the Ignite-UX GUI.

Next, define a 64-bit version of the application and ensure that you can load both the 32-bit and 64-bit versions of the application by making them exrequisties of each other. The impacts statements are generated with the `archive_impact` command, and then manually tailored to account for extra space in `/opt`, `/var`, and `/tmp` that the application requires as in the following example.

```
sw_sel "MyApp,v=1.0,64bit" {
```

⁸⁴ This section is not relevant to core operating system (OE) golden images. Additional configuration steps are required to define a core operating system (OE) versus an application golden image.

```

description = "My Application version 1.0"
sw_source = "my application archive"
sw_category = "SiteApps"
archive_type = gzip tar
archive_path = "myapps64-1.0.tar.gz"
impacts = "/opt" 160154Kb
impacts = "/tmp" 500Kb
impacts = "/var" 450334Kb
exrequisite="MyApp,v=1.0,32bit"
post_load_cmd="/opt/myapp/bin/initialize"
}

```

You must ensure that only one version of the application is installed and that the user cannot change it. For example, allowing a user to install the 64-bit version with 32-bit HP-UX. The following example automatically selects the version based upon the HP-UX Base OS bundle has been selected.

```

sw_sel "MyApp,v=1.0,32bit" {
    description = "My Application version 1.0"
    sw_source = "my application archive"
    sw_category = "SiteApps"
    archive_type = gzip tar
    archive_path = "myapps32-1.0.tar.gz"
    impacts = "/opt" 165154Kb
    impacts = "/tmp" 500Kb
    impacts = "/var" 450334Kb
    exrequisite="MyApp,v=1.0,64bit"
    post_load_cmd="/opt/myapp/bin/initialize"
}

```

The software selection of MyApp was not created using init so the user cannot change the selection. It cannot be deselected nor can the version be changed.

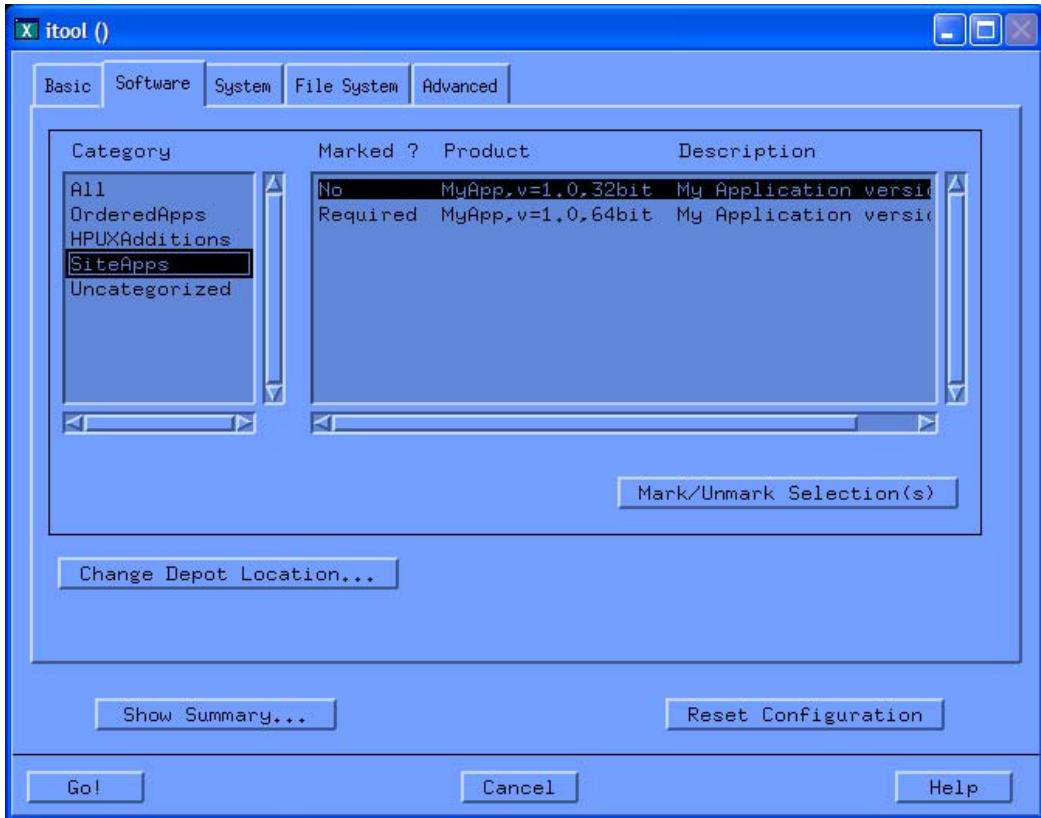
```

(sw_sel "HPUXBase64") {
    sw_sel "MyApp,v=1.0,64bit"=TRUE
}

(sw_sel "HPUXBase32") {
    sw_sel "MyApp,v=1.0,32bit"=TRUE
}

```

The new software selection you created, MyApp, appears on the **Software** tab of the Ignite-UX GUI:



The selection of the 64-bit product is marked as Required because the 64-bit OE has been selected.

If you placed the `init` before the `sw_sel`, the Ignite-UX GUI would allow you to change the software selection. The software would then be marked with Yes or No to indicate whether it will be installed. This is not recommended because it allows users to install the 64-bit version with a 32-bit installation; this should be avoided and is illustrated in the following example:

```
(sw_sel "HPUXBase64") {
    init sw_sel "MyApp,v=1.0,64bit"=TRUE
}

(sw_sel "HPUXBase32") {
    init sw_sel "MyApp,v=1.0,32bit"=TRUE
}
```

Since the examples presented are installing onto HP-UX B.11.11, you do not have to worry about issues resulting from code compiled for PA-RISC or Itanium-based systems. If this is a consideration as it is with HP-UX B.11.23 and later revisions of HP-UX, you must use the following test to determine the system type. For more information, see `instl_adm(4)`.

```
is_hppa
is_ia64
    Boolean values that indicate if a system is PA-RISC or
    Itanium®-based architecture, respectively.
```

Additional system hardware tests are explained in the following section.

Using **noncore.cfg** to define applications

Ignite-UX provides a template that you can use to define non-SD (archive based) applications in the file /opt/ignite/data/examples/noncore.cfg. The following is an excerpt from this file:

```
#####
## 
## Filename:      noncore.cfg
## 
## @(#) noncore.cfg $Revision: 10.10 $
## 
## 
## Description:
##
```

The template file provides a description of itself and how you can use it to define non-core operating system software installations.

```
##      If the software source is not an "SD" depot, then the
##      make_config command cannot be used to generate a config file
##      to represent it and therefore the config file would have to
##      be manually created. This example gives a starting template
##      to use for such a config file.
##
##      This config file would be used for a non-coreOS (application)
##      tar/cpio archive.
##
```

When you copy this file do not make the name **noncore.cfg**, rather rename it to a name that allows you to observe what application it is defining. The following excerpt describes this concept:

```
##      This config file should be copied to:
##          /var/opt/ignite/data/Rel_B.<release>/<filename>
##      then edit this new file to match your situation.
##
##      The values that you will most likely need to change in this
##      new file are:
##          nfs_source    - IP address and directory path to server
##                           directory containing the archives.
##          archive_path - path name of archive to be loaded from the
##                           nfs_source location.
##          sw_source     - tag name of the software source.
##          sw_sel        - tag name of the software select.
##          description   - description of software.
##          impacts       - amount of space needed, for a given selection,
##                           on each named file system.
##
##      Also, you may need to:
##          - add or delete any sw_sel's you do not need.
##          - run /opt/ignite/lbin/archive_impact on each
##              archive and replace the "impacts" statements with
##              those given by the archive_impact tool.
##
##
##      When done editing this config file, be sure to run:
##          /opt/ignite/bin/instl_adm -T -f <filename>
##      to check for syntax errors. Then either use manage_index, or
##      manually edit the /var/opt/ignite/INDEX file to add a reference
##      to the copy of the file you edited.
```

```

## Note:
## This example shows the use of an NFS mounted directory
## from the Ignite-UX server to the client in order to obtain
## access to the archives. You could also use the ftp_source or
## remsh_source keywords instead.
##
## In order to use the nfs_source keyword as shown, you must first
## export the directory given in the nfs_source keyword. This is
## accomplished by editing the Ignite-UX server file /etc(exports
## and adding the export information. When editing is complete,
## be sure to run the "/usr/sbin/exportfs -av" command (as root)
## to make the take effect.
##
## See also:
##   - instl_adm(4)      for syntax documentation
##   - instl_adm(1M)     for checking syntax (-T option)
##   - archive_impact(1M) for generating "impacts" statements
##   - ignite(5)          for general overview
##
#####

```

The `sw_source` stanzas define the location of the archives that you want to reference using `sw_sel` clauses. In the following excerpt, the source format is set to archive so Ignite-UX will search the archive definitions in the `sw_sel` clauses that use this `sw_source`.

```

#####
## Software Sources
#####

sw_source "Per-Discipline Packs" {
    description = "Software Apps for Individual Disciplines"
    source_format = archive
    source_type = "NET"

    # Change this to be your NFS server's IP and path:
    nfs_source = "14.12.99.113:/var/opt/ignite/archives"
}
```

The software categories, `sw_category` are defined next. Categories are a good way of grouping software together though if you create too many categories it may become more difficult to find software in the **Software** tab of the Ignite-UX GUI rather than easier. The following excerpt is an example of a `sw_category` definition:

```

#####
## Software Categories
#####

sw_category "Disciplines" {
    description = "Software for Individual Disciplines"
}
```

Next is a test of the hardware model string. The expression is set to `9000/*` meaning that if the string returned contains `9000` followed by zero or more slashes (`/`) then the first set of software defined in the following excerpt is for PA-RISC systems only. The software definitions that follow the `else` are for Itanium®-based systems.

```

## Note that the HARDWARE_MODEL clause is only needed if the sw_sels
## inside it can only be installed on a certain type of machine.
##
HARDWARE_MODEL ~ "9000/*" {

#####
## Software Selections
#####

```

The sw_sel clauses are self-explanatory. The important parts are:

- The src, which defines where the software is located.
- The archive type, which defines that it is a gzipped tar archive though it could be compressed or in cpio format.
- The archive name or path of the archive.
- The impacts statements that enable Ignite-UX to determine how much additional file system space is required for the application.

```

sw_sel "EE pack" {
    description = "Software Apps for Electrical Engineering PA clients"
    sw_source = "Per-Discipline Packs"
    archive_type = gzip tar
    archive_path = "ee_client_archive_PA.gz"
    sw_category = "Disciplines"
    impacts = "/var" 12568Kb
    impacts = "/usr" 23468Kb
    impacts = "/" 2Kb
}

```

The following sw_sel defines a different application, but shares all of the same attributes.

```

sw_sel "ME pack" {
    description = "Software Apps for Mechanical Engineering PA clients"
    sw_source = "Per-Discipline Packs"
    archive_type = gzip tar
    archive_path = "me_client_archive_PA.gz"
    sw_category = "Disciplines"
    impacts = "/var" 23678Kb
    impacts = "/usr" 12986Kb
    impacts = "/opt" 56892Kb
    impacts = "/" 64Kb
}

```

The following sw_sel is for Itanium®-based systems is very similar to the software product that definition for PA-RISC systems in the previous excerpts. The only differences are the archive_path names and the impacts statements. Notice that the software products have the same name since they cannot both be defined at the same time. Any software that depends on this software product can use the one name as a dependency (for example, in conjunction with corequisite or exrequisite).

```

} else { # IPF software selections:

#####
## Software Selections
#####

```

```

    sw_sel "EE pack" {
        description = "Software Apps for Electrical Engineering IPF clients"
        sw_source = "Per-Discipline Packs"
        archive_type = gzip tar
        archive_path = "ee_client_archive_IPF.gz"
        sw_category = "Disciplines"
        impacts = "/var" 18852Kb
        impacts = "/usr" 35202Kb
        impacts = "/" 3Kb
    }
}

```

The next example `sw_sel` defines a different application, but shares all of the same attributes.

```

    sw_sel "ME pack" {
        description = "Software Apps for Mechanical Engineering IPF clients"
        sw_source = "Per-Discipline Packs"
        archive_type = gzip tar
        archive_path = "me_client_archive_IPF.gz"
        sw_category = "Disciplines"
        impacts = "/var" 35517Kb
        impacts = "/usr" 19479Kb
        impacts = "/opt" 85338Kb
        impacts = "/" 96Kb
    }
}

```

SD and archive bitness comparison

You may have noticed an issue with the non-SD application example. Special effort was taken to ensure that a 64-bit version of the product was not installed on a 32-bit system. You could set SD to do this automatically. On PA-RISC systems, you would set the SD attribute `os_name`. For example using `swlist -l fileset -a name -a os_name` to look at some of the installed filesets results in the following:

```

# FCMassStorage                                HP-UX
FCMassStorage.FCMS-ENG-A-MAN                   HP-UX
FCMassStorage.FCMS-INIT                         HP-UX
FCMassStorage.FCMS-JPN-E-MAN                   HP-UX
FCMassStorage.FCMS-JPN-S-MAN                   HP-UX
FCMassStorage.FCMS-KRN                         HP-UX:*64
FCMassStorage.FCMS-RUN                         HP-UX
FDDI-KRN-COM                                  HP-UX
FDDI-KRN-COM.FDDI467-KRN                     HP-UX:*64
...

```

Notice that `os_name` is set to `HP-UX:*64` for some of the filesets in the previous example. In the file, `/opt/ignite/data/Rel_B.11.11/config` that is included into the configuration clause that you setup is the following configuration code:

```

sd_command_line += " -x os_release=" +${release}
(_hp_os_bitness == "64")
{
    (!can_run_64bit)
    {
        ERROR += "This system model: \\" +${model}+ \" is not supported for
running 64bit HP-UX, you must select the 32bit selection"
    }
    sd_command_line += " -x os_name=HP-UX:64 "
}

```

```
}
```

This code allows SD to select common and 64-bit only filesets (and not include 32-bit only filesets) into the operating system. Because there is nothing similar for archives, you must do this manually.

Adding the non-SD application configuration file to the INDEX file

Now you can easily add the application configuration file to our `cfg` clause in the configuration file using the following command:

```
# manage_index -a -f /var/opt/ignite/data/Rel_B.11.11/myapps_cfg \
-c "Custom B.11.11 installation using SD"
```

The software is now ready to install.

Installing patches from a depot

Installing patches from a depot is explained in the following three topics.

Setting up the depot

In this section, you will setup a depot containing the patches that you made sure were installed earlier to enable the system to read the DVD containing the core operating system (OE) software.

At this point, you have a depot containing the patches that is observed using `swlist`:

```
# swlist -l depot

# Initializing...
# Target "test" has the following depot(s):
...
/var/opt/ignite/depots/Rel_B.11.11/patches_iso
# swlist -d @ /var/opt/ignite/depots/Rel_B.11.11/patches_iso
# Initializing...
# Contacting target "test"...
#
# Target: test:@/var/opt/ignite/depots/Rel_B.11.11/patches_iso
#

#
# No Bundle(s) on test:@/var/opt/ignite/depots/Rel_B.11.11/patches_iso
# Product(s):
#
PHCO_25841      1.0          Add Rock Ridge extension to mount_cdfs(1M)
PHKL_26269      1.0          Rock Ridge extension for ISO-9660
PHKL_28025      1.0          Rock Ridge extension for ISO-9660
```

Unfortunately, Ignite-UX does not work with products and filesets so you must package the product patches into bundles.

Packaging the patches into a bundle

Previously the use of the `make_bundles` command to package products into bundles was discussed; see "Package an application in SD format". The following command quickly places the patches into a bundle:

```

# make_bundles -b -i -n "ISO_patches" \
-t "Patches to enable Rock Ridge natively" -r 1.0 \
/var/opt/ignite/depots/Rel_B.11.11/patches_iso

Generating list of unbundled filesets...

===== 10/11/04 23:37:34 EST BEGIN swpackage SESSION

* Session started for user "root@test.rc.aus.hp.com".
* Source:      test:/var/tmp/psf.24606
* Target:      test:/var/opt/ignite/depots/Rel_B.11.11/patches_iso
* Software selections:
*
* Beginning Selection Phase.
* Reading the Product Specification File (PSF)
  "/var/tmp/psf.24606".
* Reading the bundle "ISO_patches" at line 11.

* Selection Phase succeeded.

* Beginning Analysis Phase.
* Analysis Phase succeeded.

* Beginning Package Phase.
* Packaging the bundle
  "ISO_patches,r=1.0,a=HP-UX_B.11.11_32/64,v=HP".
* Package Phase succeeded.

===== 10/11/04 23:37:35 EST END swpackage SESSION

```

You can see that the patches have all been placed into a bundle with the following command:

```

# swlist -d @ /var/opt/ignite/depots/Rel_B.11.11/patches_iso

# Initializing...
# Contacting target "test"...
#
# Target:  test:/var/opt/ignite/depots/Rel_B.11.11/patches_iso
#
#
# Bundle(s):
#
ISO_patches    1.0          Patches to enable Rock Ridge natively

```

Generating a configuration

Next, you must create a configuration file to describe the SD depot. This is very easy and only requires the `make_config` command:

```
# make_config -s /var/opt/ignite/depots/Rel_B.11.11/patches_iso \
-c /var/opt/ignite/data/Rel_B.11.11/patches_iso_cfg
```

The generated configuration file looks like the following:

```

#####
## Software Sources
#####

sw_source "/var/opt/ignite/depots/Rel_B.11.11/patches_iso" {
    source_type = "NET"
    sd_server = "10.2.72.150"
    sd_depot_dir = "/var/opt/ignite/depots/Rel_B.11.11/patches_iso"
    source_format = SD
}

#####
## Other Software
#####

sw_sel "ISO_patches" {
    description = "Patches to enable Rock Ridge natively"
    sw_source = "/var/opt/ignite/depots/Rel_B.11.11/patches_iso"
    sw_category = "Uncategorized"
    sd_software_list = "ISO_patches,r=1.0,a=HP-UX_B.11.11_32/64,v=HP"
    (_hp_os_bitness == "32") {
        impacts = "/sbin" 224Kb
        impacts = "/usr" 184Kb
    }
    (_hp_os_bitness == "64") {
        impacts = "/sbin" 224Kb
        impacts = "/usr" 224Kb
    }
}

```

The last step is to add the configuration file to the `cfg` clause you are assembling in the `INDEX` file using the following command:

```
# manage_index -a -f /var/opt/ignite/data/Rel_B.11.11/patches_iso_cfg \
-c "Custom B.11.11 installation using SD"
```

Customizing configuration

In the current `INDEX` file, you have the following `cfg` clause:

```
cfg "Custom B.11.11 installation using SD" {
    description "This selection supplies the default system configuration
that HP supplies for the B.11.11 release."
    "/opt/ignite/data/Rel_B.11.11/config"
    "/opt/ignite/data/Rel_B.11.11/hw_patches_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/core_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/myapps_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/apps_1203_cfg"
    "/var/opt/ignite/data/Rel_B.11.11/patches_iso_cfg"
    "/var/opt/ignite/config.local"
}
```

You have a rich feature set to build on top of, the advantage to this is that the `config` and `hw_patches_cfg` files provide a good working base on which you can build a generic system or a customized system.

To begin customizing the system you must create a new configuration file then add it to the `cfg` clause with the following commands:

```
# touch /var/opt/ignite/data/Rel_B.11.11/custom_cfg  
# manage_index -a -f /var/opt/ignite/data/Rel_B.11.11/custom_cfg \  
-c "Custom B.11.11 installation using SD"
```

You have a file to which you can start adding configuration information to customize the existing configuration files. HP does not recommend that you customize configuration files generated by Ignite-UX commands. If the software in the depot changes, it is much easier to recreate them using Ignite-UX commands than it is to maintain them manually.

For example, rather than changing the configuration files that define a software product to set it to be installed you can add configuration information to this new file to install the software just for this `cfg` clause. This means the configuration file that describes the depot could be used in other `cfg` clauses as well to define the contents of a depot without having to take additional action. You could add the following configuration information:

```
init sw_sel "B3901BA"=TRUE  
init sw_sel "B5725AA"=TRUE  
sw_sel "ISO_patches"=TRUE
```

This would install the ANSI/C compiler, the complete version of Ignite-UX, and the patch bundle you created previously onto the system being installed. Users running the Ignite-UX GUI to interactively change the software to be installed could change whether the ANSI/C compiler or Ignite-UX is installed. The user has no choice in the installation of the ISO patches you packaged earlier since the software selection is set to `TRUE` without using `init`, which means that the Ignite-UX GUI cannot change the selection.

Next, you could define a custom disk layout that is a choice between the existing disk layouts defined in `/opt/ignite/data/Rel_B.11.11/config` or completely replacing them so you do not have a choice of any of the default disk layouts. To implement this choice you would add the following:

```
_hp_disk_layout = {  
    "MyApp custom disk layout"  
}
```

None of the default disk layout choices defined in `/opt/ignite/data/Rel_B.11.11/config` would be available (although not shown here you must define the disk layout that would be used).

Alternatively, you could use the following example that adds the new disk layout to list of available choices and then sets it to be the default choice (anyone installing a system could still use the other layouts defined in the `/opt/ignite/data/Rel_B.11.11/config`).

```
_hp_disk_layout += {  
    "MyApp custom disk layout"  
}  
init _hp_disk_layout="MyApp custom disk layout"
```

It is important to use a good foundation, building on top of the release default configuration files and the depot (software) configuration files enables you to easily create more complex configurations.

If you maintain the core operating system (OE) and the applications media in different depots, you may need to make customizations for the bundles you have created that define patches. If some products are installed from the applications depot instead of the core operating system (OE) depot and you have patches that depend on the filesets in the bundles that are taken from the applications depot, you may need to set the `load_order` on `sw_sel` clauses describing patches to be >10. This guarantees that the patches are loaded after all software from the applications depot is installed.

You can do this by directly modifying the configuration files created to describe the patches or by adding a `sw_sel` similar to the following to the customization file; not in the configuration file that actually defines the patches:

```
sw_sel "ISO_patches" {  
    load_order=15  
}
```

This takes advantage of something that described earlier. If you have a `sw_sel` mentioned twice, the information defined in the second clause overrides the information previously defined. In this case, no `load_order` was defined for the `sw_sel`, `ISO_patches` so the change adds the existing definition to change its load order since the default is 5.

A `load_order` of 10 ensures that the patches are loaded after the core operating system (OE), which is loaded at `load_order` 0, and the applications that default to `load_order` 5.

Installing with SD wrap-up

The previous sections illustrated the following concepts:

- How simple it can be to use SD to install systems; installation from an Ignite-UX server over the network was employed in the examples presented.
- How to setup your Ignite-UX server for use with SD installations.
- How to manage your configuration explaining some of the issues that you face when managing your configuration.
- Why it is a good idea to separate configuration that belongs to depots or archives from the default configuration supplied with Ignite-UX and custom configuration that you may layer over the default.
- Where possible you should not modify the configuration files created by Ignite-UX commands. If you change the software in a depot, you will probably want to recreate the configuration file that describes the depot; if you have manual changes you must review the file and reapply the same changes to the new configuration file. This is time consuming so keeping the configuration separate means you only have to be aware of changes and review the customizations layered on top of the other configuration files periodically.

Performance considerations for SD-UX based installs

The following section provides information for you to consider when you have an Ignite server that is intended to perform large numbers of concurrent cold installs using SD-UX (more than three at one time). In the following discussion, only the information regarding network bandwidth applies to all the methods Ignite-UX can use to perform install or recoveries (SD-UX-based installs, archive-based installs, and recovery sessions.)

Tuning SD-UX for concurrent access

Generally, SD will not be able to handle more than three installation sessions by remote systems without introducing either delays or timeouts. The following information also applies to swm on HP-UX 11i v3, since it uses swagentd to access SD-UX depots.

In a default setup, SD-UX is able to handle no more than six concurrent RPC connections. Any excess RPC connection attempts to swagentd will be queued, and if they take too long to service, may fail due to timeout or rejection.

A swinstall session may require up to two active RPC connections to swagentd. What happens when you want 10 systems to be able to install from your Ignite⁸⁵ server?

In this case there is no documented solution for increasing the number of RPC connections that SD-UX can handle concurrently, however the following undocumented SD-UX setting can be put into the /usr/lib/sw/sys.defaults file:

```
swagentd.rpc_max_calls=n
```

Where *n* is two times the number of expected concurrent sessions SD-UX needs to handle. For temporary changes (or testing) you can stop swagentd and restart it with the option *-x rpc_max_calls=n*. You could use the following command to do this:

```
# swagentd -r -x rpc_max_calls=n
```

When choosing a value for *n*, please take into account that local swlist commands do not make RPC connections to swagentd.

So to allow for 10 systems to install concurrently, and to allow a few RPC connections for the local system to perform swlist and other SD-UX commands, you would need 20 RPC connections for the installing systems (10x2). A few extra RPC sessions (4) are added to allow for local SD command usage (swlist, swcopy, etc), and some padding is added to ensure that if usage grows no problem will be immediately encountered (an extra 4). In this case you would need to set *rpc_max_calls* to be 28.

Issues to consider related to the HP-UX revision

There are some considerations for SD-UX related to the version of HP-UX used. On 11i v2 and later releases of HP-UX, SD-UX is kernel threaded; on 11i v1 SD-UX uses CMA or user space threads.

Increasing *rpc_max_calls* will increase the number of kernel threads in use in swagentd. This may have an impact on the following kernel tunables on HP-UX 11i v2:

1. nproc – There will be additional swagent processes running (one process per swagent ; one swagent per SD command that is issued to communicate to SD-UX on this system.) The increase in processes caused by swagent is unlikely to add enough additional processes on the system to force a change to nproc.

⁸⁵ The SD-UX server may be different from your Ignite server. This discussion assumes they are the same system.

2. nkthread – There will be a thread created in `swagentd` for each of the RPC connections that `swagentd` may need to handle. If you were to increase `rpc_max_calls` from the default of six to 28, there would be at least 22 more threads active in `swagentd`.
3. `max_thread_proc` – As more threads will be active in `swagentd`, you will need to ensure that `max_thread_proc` is set at a value that allows all of the additional threads to be created.

Because `swagentd` is kernel-threaded on HP-UX 11i v2, and there is a lot of work required of `swagentd`, it is able to utilize more than one CPU on a multi-processor system. With HP-UX 11i v1, `swagentd` is not able to use more than 100% of one CPU, as it is not kernel-threaded.

Issues Independent of HP-UX revision

When `swagent` processes are started in order to handle SD-UX requests, they will consume resources such as swap, CPU, and memory. The following issues should be considered when sizing a system:

1. The amount of memory SD-UX will consume on the system depends on how many concurrent sessions you have allowed for.
2. The number of concurrent SD-UX sessions will reduce the bandwidth available per system using SD-UX for installation via a network interface.
3. Depending on usage (if multiple `swagent` processes are attempting to read from disk) the buffer cache may not be large enough to effectively cache SD-UX depots, which would cause performance problems.

Further discussion of these issues follow:

Memory and the buffer cache

The amount of memory consumed by SD-UX depends on your planned concurrent usage. If you are planning for a worst case of 10 concurrent SD-UX-based installs, you should monitor the memory usage of one `swagent` process during an install. You can use the memory usage of that process to estimate the resources required for the worst case of 10 concurrent install sessions.

If you used `glance` to track the memory usage of a `swagent` process, and found that it used a maximum of 16Mb of real memory and 24Mb of virtual memory, you would need to ensure that you had 160Mb of real memory and 240Mb of virtual memory available at peak times to handle the SD-UX sessions⁸⁶.

For example, the following system would be able to cope with 10 `swagent` processes running since there is currently 531Mb of free memory:

# swapinfo -tm								
	Mb	Mb	Mb	PCT	START/	Mb		
TYPE	AVAIL	USED	FREE	USED	LIMIT	RESERVE	PRI	NAME
dev	1024	0	1024	0%	0	-	1	/dev/vg00/lvol2
dev	3000	0	3000	0%	0	-	0	/dev/vg00/lvol15
reserve	-	211	-211					
memory	924	393	531	43%				
total	4948	604	4344	12%	-	0	-	

⁸⁶ Real and virtual memory usage is likely to be different in your environment. You should see how much memory `swagent` uses in your environment, and use that value in any memory requirement calculations.

If you were planning for a maximum of 50 concurrent swagent processes, you might begin to run into problems since the system is likely to start to swap (800Mb of real memory and 1200Mb of virtual memory would be required).

Do not add real and virtual memory together when trying to determine if you need more swap or physical memory in a system; the amount of virtual memory consumed by swagent includes the real memory used.

When a lot of memory has been consumed by swagent processes, the buffer cache would be significantly reduced. This leads to performance bottlenecks due to accessing SD depots on disk. A large buffer cache (assuming that multiple swagent processes are accessing the same depot) will increase the likelihood that swagent processes will read data from the buffer cache rather than directly from disk. How large that buffer cache should be is beyond the scope of this discussion.

If you need to have a lot of concurrent SD sessions running on the Ignite server, carefully consider the trade-off between cost and performance. Large JBODs will hold a lot of SD depots, but the speed of your installs will be limited if you have a very small buffer cache and the data must be read from disk constantly by multiple swagent processes. If you have a JBOD that has a lot of read requests from swagent processes reading depots, you might consider relocating the data to a higher performance device such as an HP XP or EVA disk array, or you could add more physical memory to your system to allow it to be used as buffer cache.

Network Bandwidth

The following section is based on assumptions that may or may not be true in your environment. The main purpose of this section is to help you understand how multiple SD sessions using one LAN interface can affect how long it takes to install a system via the network.

The following is a comparison of 100BT and Gigabit Ethernet network interfaces. The examples assume throughput rates on the LAN interfaces, and that a number of systems would all start an install at the same time using the same SD depot.

If you have an 100BT interface, and assuming that SD gets about 60 megabits per second, 60 systems pulling a 3Gb depot SD gives you the following amount of time to install a system:

$$60 * 3 * 1024 / ((60/8) * 3600) = 184320 / 27000 = 6.82 \text{ hrs to complete}$$

If you can achieve 300 megabits through a Gigabit Ethernet interface, that dramatically drops the time required (assuming that the system is fast enough to push that much data from SD through):

$$60 * 3 * 1024 / ((300/8) * 3600) = 184320 / 27000 = 1.36 \text{ hrs to complete}$$

The calculation is:

$$\text{Number of concurrent systems} * \text{depot size (in Mb, hence } 3 * 1024)$$

That is divided by:

$$\text{Lan speed (in megabits)/8 (to get Megabytes per second)} * 3600$$

The result is the estimated number of hours to install a system. Please note that these calculations are only intended to illustrate that the time it takes to install a system is impacted by: 1) the size of the data to be installed by SD, 2) the number of systems running concurrently, and 3) the speed of the connection. The calculation is not intended to tell you exactly how long it takes to install an actual system.

An understanding of the network topology between the Ignite server and the systems being installed is important. If the Ignite server has a Gigabit Ethernet interface connected to a Gigabit Ethernet switch port, but is connected to the installing systems via a 100BT link between switches, the Ignite server will never achieve more than 100BT speeds in total to all of those clients.

Careful consideration should be given to your network topology as it can impact on network throughput significantly.

Installation configurations using golden images

It is important to look at how to create the `cfg` clause used for a golden image installation in the `INDEX` file.

Golden image configuration file explanation

Provided with Ignite-UX is a configuration file that you can use to define a golden image installation. This file is `/opt/ignite/data/examples/core11.cfg` and the following is an example of this file:

```
#####
## Filename:    core11.cfg  (template for B.11.* systems)
##
## @(#) core11.cfg $Revision: 10.10 $
##
## Description:
##   This is an sample configuration file that enables an OS archive
##   to be installed via Ignite-UX.
##
##   Before Ignite-UX can use an OS archive like that created by
##   make_sys_image(1M), the following must be done:
##
```

The instructions for how to use this config file are actually located in it as comments (see below).

```
##           - Make an OS archive of the desired system using
##             /opt/ignite/data/scripts/make_sys_image
##             (see make_sys_image(1M) for details).
##
##           - Determine if you want to use ftp, NFS, or remsh to access
##             the OS archive (stored on the Ignite-UX server).
##
##           - Make a copy this sample configuration file.
##
##           - Edit your copy of this configuration file and modify:
##             - The IP addresses of the Ignite-UX server.
##             - The paths to OS archives.
##             - The sw_sel keyword descriptions.
```

```

##           - remove any extra sw_sel clauses.
##           - run /opt/ignite/lbin/archive_impact on each OS archive
##             and replace the "impacts" statements with those given
##             by the archive_impact tool.
##
##           - Edit the /var/opt/ignite/INDEX file. If you already have a
##             stanza in the INDEX file that refers to an SD depot, then
##             you will probably want to make a new stanza ("cfg" clause)
##             that does not have references to the core SD cfg files, and
##             insert the path to your newly edited configuration file.
##             The end result should look something like this:
##

```

HP recommends that /opt/ignite/data/Rel_B.xx.yy/config is included first, similarly with SD installs there is information already defined that allows you to start with a very strong configuration that you can customize for your environment.

It is easiest to begin customizing by renaming core11.cfg to a new name and placing it in /opt/ignite/data/Rel_B.xx.yy/config. It is likely that you will have more than one golden image installation defined for any one HP-UX revision so customizing core11.cfg directly is not recommended.

```

##           cfg "HP-UX via an OS archive" {
##             description "11.11 system using archives."
##             "/opt/ignite/data/Rel_B.11.11/config"
##             "/opt/ignite/data/Rel_B.11.11/core11.cfg"
##             "/var/opt/ignite/config.local"
##           }
##
##           - If you have additional archives or applications in an SD
##             depot that you would like loaded along with this archive,
##             you may create a configuration file for them and add it to
##             the INDEX file entry as well.
##
##           See make_config(1M) for creating configuration files for an
##           SD depot. See the example file "noncore.cfg" for setting
##           up application archives.
##
##           - After making changes, always run:
##
##             # /opt/ignite/bin/instl_adm -T
##
##             to check the syntax.
##
#####
#####
## Software Sources
#####


```

In the sw_source clause for the core operating system archive, the source format is archive so sw_sel definitions that use this sw_source must be archives as well. Since this archive contains the core operating system, it must be loaded at load order zero (0)⁸⁷.

⁸⁷ Only one archive can be loaded at load order zero (0), and it must be the core operating system archive. For technical reasons if the core operating system archive cannot fit within a certain size (cannot be written to a CD) then you must attempt to split out the software that can be loaded after load order one. For example, /opt.

```

sw_source "core archive" {
    description = "HP-UX Core Operating System Archives"
    load_order = 0
    source_format = archive
    source_type="NET"

    # When using this configuration file with a tape or CD-ROM where the
    # archive is on the same piece of media, then un-comment the next
    # line to prevent a prompt to change media during install.
    # change_media=FALSE

```

The following instructions, the two scripts given as the `post_load_script` and the `post_config_script` are required and you should not remove them.

However, there may be instances in which you may need to change the contents of the two scripts (for more information, see "Instances that may require modifying `os_arch_post_1`"). If you need to modify these files, you must never change the versions under `/opt` because when you upgrade Ignite-UX new versions of the files are installed in that location. In other words, your modifications are overwritten and are lost. Instead, you should copy them to `/var/opt/ignite/scripts` and then modify them. Note that if you have multiple golden images that require different changes in `post_load|config_scripts`, you should rename the scripts and then reference the new names for each golden image installations.

Periodically after upgrading Ignite-UX you should check to see what differences exist between the new scripts in `/opt/ignite/data/scripts` and any scripts you may have created using these examples. Ignite-UX may deliver enhancements or defect fixes that you need in these scripts. If you want to take advantage of these updates, you must manually merge the changes.

```

# The following 2 scripts should be reviewed to see if the defined
# defaults are applicable. If you need to modify them, copy them
# to /var/opt/ignite/scripts first and change the paths here.
post_load_script = "/opt/ignite/data/scripts/os_arch_post_1"
post_config_script = "/opt/ignite/data/scripts/os_arch_post_c"

# If nfs_source is used, be sure to export the source.
nfs_source = "10.2.72.150:/var/opt/ignite/My_Golden_Images"

# ftp_source and remsh_source are alternate ways to copy the
# archive. For example:
#   ftp_source = "anonymous@10.2.72.150:myusername"
#   remsh_source = "myusername@10.2.72.150"
}

#####
## Software Categories
#####

```

Next, categories that define the software that can be installed using the Ignite-UX GUI are defined.

```

sw_category "Languages" {
    description = "Languages"
}

sw_category "HPUXEnvironments" {
    description = "HP-UX Operating Environments"
}

```

Note that there are no Operating Environment (OE) definitions. This is not possible because software packaged in SD format is required.

```
#####
## Operating Environments
#####

#####
## OS archives
#####

#####
```

Next is an example 32-bit operating system archive `sw_sel`. Note that the logic in this stanza dictates that if the system can run 32-bit HP-UX, it becomes the default golden image to be installed. If you want to change this test so that any system that could run 64-bit HP-UX, you would replace the logic that sets the `sw_sel` to TRUE or FALSE in the golden image - 32 bit OS clause with `!can_run_64bit` and in the golden image - 32 bit OS clause replace the final test with `can_run_64bit`.

The use of `exquisite` prevents users from choosing the 32-bit and the 64-bit core operating systems to load at the same time. Also, the use of `visible_if` prevents one or the other core operating system archive from appearing if the system can only run 32-bit or 64-bit HP-UX but not both. This is a good practice since you should not confuse users performing installations by offering software that cannot be installed onto the system as a choice.

If you are installing HP-UX B.11.23, you would remove the 32-bit `sw_sel` (since HP-UX B.11.23 for HP9000 systems does not provide a 32-bit kernel). If you needed to support a mixture of HP9000 and HP Integrity systems, consider placing it with the definition of an archive for HP Integrity systems.

```
init sw_sel "golden image - 32 bit OS" {
    description = "English HP-UX 11.11 CDE - 32 Bit OS"
    sw_source = "core archive"
    sw_category = "HPUXEnvironments"
    archive_type = gzip tar

    # For NFS, the path to the archive is relative to the mount point
    # specified in the "nfs_source" keyword within the sw_source stanza
    # above:
    archive_path = "hp_client1_B.11.11_32bitCDE.gz"

    # ftp and remsh sources can use a full path:
    # archive_path = "/pub/IUXarchives/hp_client1_B.11.11_32bitCDE.gz"

    # The data for the "impacts" statements are found by running:
    # /opt/ignite/lbin/archive_impact -t -g <OS_archive>
    impacts = "/"          27Kb
    impacts = "./dt"        35Kb
    impacts = "/etc"        1864Kb
    impacts = "/export"     1Kb
    impacts = "/home"       1Kb
    impacts = "/opt"         74096Kb
    impacts = "/sbin"        13449Kb
    impacts = "/stand"      1Kb
    impacts = "/tmp"         1Kb
    impacts = "/users"      40Kb
    impacts = "/usr"         225951Kb
```

```

    impacts = "/var"      5705Kb
    exrequisite += "golden image - 64 bit OS"
    visible_if  = can_run_32bit
} = (can_run_32bit)

```

Following is the example for a 64-bit golden archive. All of the information previously presented regarding the 32-bit golden image are applicable to a 64-bit golden archive as you can see in the example.

```

init sw_sel "golden image - 64 bit OS" {
    description = "English HP-UX 11.11 CDE - 64 Bit OS"
    sw_source = "core archive"
    sw_category = "HPUXEEnvironments"
    archive_type = gzip cpio
    archive_path = "hp_client1_B.11.11_64bitCDE.gz"
    impacts = "/"          421Kb
    impacts = "/sbin"     30086Kb
    impacts = "/opt"       78654Kb
    impacts = "/usr"       276420Kb
    impacts = "/var"       10059Kb
    exrequisite += "golden image - 32 bit OS"
    visible_if  = can_run_64bit
} = (!can_run_32bit)

```

You must set `_hp_os_bitness`; this variable is required by Ignite-UX to determine the "bitness" of the kernel being installed. Depending on which archive has been selected, this variable is set to the corresponding value.

Tip:

If you install HP-UX B.11.23, you must unconditionally set `_hp_os_bitness` to 64.

```

#####
## Set the variable _hp_os_bitness depending on which of the archive
## sw_sels was selected. This variable is used by configuration files
## generated via make_config to determine which version (32-bit or
## 64-bit) of a bundle will be loaded.
#####
(sw_sel "golden image - 32 bit OS") {
    _hp_os_bitness = "32"
}
(sw_sel "golden image - 64 bit OS") {
    _hp_os_bitness = "64"
}

```

Following is a definition for the English language. The `sw_source` associated with it is a dummy `sw_source`. In other words, the source format is `cmd` and does not have any scripts so it does nothing. Note that the `sw_sel` does set the locale that you can choose from in the Ignite-UX GUI.

```

#####
## Languages
## Specified so that the selector in the UI will not be empty.
## Cannot use the "core archive" sw_source because there are no

```

```

## archives associated with it.

## The "no select" sw_source is a no-op source since it is "cmd"
## format, and does not have any scripts specified. Every sw_sel needs
## a sw_source.
#####
sw_source "no select" {
    source_format = cmd
}

init sw_sel "English" {
    description = "English Language Environment"
    sw_source = "no select"
    sw_category = "Languages"
    locale = { "SET_NULL_LOCALE:English", "C:English" }
} = TRUE

```

If you wish to be able to select, for example, Japanese locales during installation, you would have to customize the configuration file to describe the locales available in the golden image. For example, you could replace the locale line above with the following configuration:

```

locale = { "ja_JP.SJIS:Japanese", "ja_JP.eucJP:Japanese",
"ja_JP.utf8:Japanese", "SET_NULL_LOCALE:English", "C:English" }

```

This would allow you to select any of the Japanese or English locales available on the system. Note that the "English" description can be misleading. The locale SET_NULL_LOCALE means the default system locale (POSIX/C) and C are explicitly the POSIX/C locale.

The locale line in the configuration file should not be wrapped; rather, it should be on one contiguous line. You should never add locales that do not exist in the golden image to the configuration.

The format of a entry in the locale list is:

```
"<locale>:<description>"
```

Where <locale> is the name of the locale shown in the output of the "locale -a" command on the system on which the golden image was created (for example ja_JP.SJIS). This is followed immediately by a colon and then by a description of the locale. The description cannot contain white space; if you need more than one word in the description, replace any white space by the underscore ("_") character. For example "ja_JP.SJIS:Japanese_SJIS" has the space between Japanese and SJIS replaced with an underscore. Remember that the locale list entry must be enclosed by double quotes, as seen above, when placed into a configuration file.

You should also be aware that even though a locale might not be listed, it only means that it cannot be selected via the Ignite-UX GUI. Any locale that was present when the golden image was created will still be available on any system installed using the golden image.

The next section defines default keyboards. You could add other keyboards depending on your environment or replace the defaults with those applicable to you.

```

#####
## Keyboards
## Newer machines use USB as the keyboard interface. This example
## shows a way to handle both older (PS2) and newer (USB) keyboards.
## There are many other keyboard types available, see itemap(1m).

```

```

#####
has_ps2 {
    _hp_keyboard = {
        "Not_Applicable",
        "PS2_DIN_US_English",
        "PS2_DIN_US_English_Euro"
    }
    init _hp_keyboard = "PS2_DIN_US_English"
}

has_usb {
    _hp_keyboard = {
        "Not_Applicable",
        "USB_PS2_DIN_US_English",
        "USB_PS2_DIN_US_English_Euro"
    }
    init _hp_keyboard = "USB_PS2_DIN_US_English"
}

```

Instances that may require modifying `os_arch_post_1`

There are circumstances in which you may want to modify `os_arch_post_1`. For example, the following excerpt is from the `os_arch_post_1` script:

```

#
# Networking files:
#   /etc/hosts
#   /etc/resolv.conf
#   /etc/rc.config.d/namesvrs
#       By default, these files will be constructed from the information
#       in the config file. The starting point for the hosts file is the
#       /usr/newconfig version, which just has a loopback entry. The other
files
#       files are built from scratch. To get the archive versions of these files,
#       uncomment only the save_file lines here and comment out the rm line.
#       Using save_file will restore the file as-is from the archive. Using
#       merge_file will allow Ignite-UX to make changes to the file based
#       on what the config file or changes made in the UI.
#
#save_file /etc/hosts
#merge_file /etc/hosts
rm -f /etc/resolv.conf
#save_file /etc/resolv.conf
#merge_file /etc/resolv.conf
#save_file /etc/rc.config.d/namesvrs
#merge_file /etc/rc.config.d/namesvrs:

```

By default, during an installation ⁸⁸ from an archive, the final versions of these files are taken from `/usr/newconfig`. If you have a complex configuration that needs to be preserved from the archive version of the file, you must uncomment the corresponding `merge_file` entry.

For example, if your `/etc/hosts` file contains extra entries that you need to preserve, you would uncomment the `merge_file /etc/hosts` line in your `os_arch_post_1` script. In the case of `/etc/hosts`, it may not be a good idea to use `save_file` as it means that any changes Ignite-UX made to `/etc/hosts` would be overwritten by the version saved from the archive.

⁸⁸ This does not apply to recovery situations. A `merge_file` is executed for the files in the `os_post_arch_1` script. You can find out what files maybe merged or saved automatically in a recovery situation by searching `os_arch_post_1` for "RECOVERY_MODE".

In summary you can use:

- `save_file` to keep the version of the file from the Operating System archive preventing Ignite-UX from making changes to the file;
- `merge_file` to keep the version of the file from the Operation System archive but allow Ignite-UX to make customizations to it based upon the configuration for the installation session.

Creating a golden image

The creation of a golden image is not discussed in this document. For more information, see `make_sys_image(1M)` and the *Ignite-UX Administration Guide*.

It is important to realize that a golden image is an archive created by `make_sys_image` at clean level 2 (`make_sys_image -l 2`) so it contains no host-specific information where possible. You should not call a `make_net_recovery` or `make_sys_image` archive created at any other clean level for a golden image. In other words, to be a golden image, the archive must be created at clean level 2 using `make_sys_image`.

Final words about golden image installations

If you look at the example `cfg` clause from the `/var/opt/ignite/INDEX` file described in the `core11.cfg` file, you can see similarities between it and the SD-based `cfg` clause:

```
cfg "HP-UX via an OS archive" {
    description "11.11 system using archives."
    "/opt/ignite/data/Rel_B.11.11/config"
    "/opt/ignite/data/Rel_B.11.11/core11.cfg"
    "/var/opt/ignite/config.local"
}
```

Both clauses include the release-specific configuration file `config` before any other configuration files are defined. Then the configuration file that defines the location of the core operating system (OE) is defined.

Because `config.local` must always be loaded last, you can see that after the entry for `core11.cfg` is the point at which you would begin adding additional configuration files to install other software like patches, additional applications, and customizations.

With virtually no changes, the contents, and methods explained in the "Installation configurations using Software Distributor depots" section can be used here. This is because the name of the core operating system `sw_sel` will change so some small changes are required if you have corequisites, exrequisites, or other dependencies on this name⁸⁹).

⁸⁹The SD-based installation examples include a dependency between an application and which of the `sw_sels`, `HPUXBase32` or `HPUXBase64`, was selected. This would change to be based upon the `sw_sel` name of the core operating system (OE) golden images.

Understanding what `is_net_info_temporary` does

Ignite-UX differentiates temporary and permanent network settings with the use of the `final` keyword. The `final` keyword is placed in front of networking attributes to indicate that they should apply to the final system. This behavior can change based upon the value of the `is_net_info_temporary` variable.

Information picked up from a DHCP/bootp server is placed into the configuration without the `final` keyword. Configuration files may also define network configuration without the `final` keyword. This configuration is treated in the same way as networking information picked up from a DHCP/bootp server.

When `is_net_info_temporary` is set to true, only networking information with the `final` keyword will be reflected in the network settings of the final system. When set to false (the default), networking information without the `final` keyword can override networking information specified with the `final` keyword.

To illustrate this, let's look at an example. The following partial configuration shows two DNS name servers defined for a system with the `final` keyword:

```
final dns_nameserver[0] = "10.1.1.10"
final dns_nameserver[1] = "10.1.1.30"
```

The entry from `/etc/bootptab` used to provide this system with network information when booting over the network looks like:

```
ignite-defaults:\
    ht=ethernet:\
    hn:\
    bf=/opt/ignite/boot/nbp.efi:\
    bs=48
ig0001:\
    tc=ignite-defaults:\
    ha=0073217D1429:\
    ip=10.1.1.119:\
    sm=255.255.255.0:\
    gw=10.1.1.1:\
    ds=10.1.1.10      *** DNS server
```

When booted using this `bootp` entry, Ignite-UX adds the following entry into the configuration automatically:

```
dns_nameserver[0] = "10.1.1.10"
```

With `is_net_info_temporary` set to false, the information Ignite-UX will write into `/etc/resolv.conf` for the DNS name server is:

```
nameserver 10.1.1.10
```

However with `is_net_info_temporary` set to true, the configuration marked with `final` is used instead:

```
nameserver 10.1.1.10
nameserver 10.1.1.30
```

This shows Ignite-UX using the networking configuration marked with the final keyword when `is_net_info_temporary` is set to true and using the networking configuration not marked with final when `is_net_info_temporary` set to false (in this case information picked up from a DHCP/bootp server).

If `is_net_info_temporary` is set to false, the following non-final networking information may be applied to the system unless the configuration already defines it with the final keyword:

- hostname
- DNS domain name
- NIS domain name
- NIS server
- NTP server

If `is_net_info_temporary` is set to false, the following networking information not marked final will be used instead of what is defined as final in the configuration:

- lan interface information (e.g. IP address, netmask, ...)
- DNS name servers
- DNS domain search list

If `is_net_info_temporary` is set to false, routing information not marked final will be applied to the system, but only if the configuration defines no final routing information.

This information is accurate as of Ignite-UX version C.6.6 (but it has not changed in a long time before that release of Ignite-UX).

If you encounter issues with any of the above information changing please investigate how you have set `is_net_info_temporary` (to true or false) and decide how it should be set to achieve the result you require.

Understanding how VxFS file system versions are set

Starting with Ignite-UX version C.6.3 with HP-UX 11.23, you are now able to interactively choose the VxFS file system version in `itool`. To do this you need to go to the **Filesystem** tab and from the **Additional Tasks** pull down choose **File System Parameters**. You can change the VxFS version at the bottom on the new screen that appears. You currently have a choice of: default, 4, 5, or 6. The value of `default` will be VxFS version 5 for the life of HP-UX 11.23 (the value of `default` for future revisions of HP-UX may differ but should be constant for the life of a revision of HP-UX.)

For HP-UX 11.11, this selection list is greyed-out and you cannot change it. The version of VxFS used on those versions of HP-UX will be the same as the VxFS layout for the 11.11 `mkfs` command.

With Ignite-UX C.6.3, all 11.23 cold installs will be done using VxFS layout version 5. The highest

layout version for VxFS 4.1 is 6. If your system will be running VxFS 3.5 you should not choose layout version 6, as VxFS 3.5 cannot read VxFS version 6 file systems.

With Ignite-UX version C.6.9 (released September 2006) on HP-UX 11.23 you have the choice of making a VxFS file system version 7. You must be careful when doing this as VxFS version 7 file systems require you to be running VxVM/VxFS 5.0. Although Ignite-UX will allow you to set the VxFS version to 7, it cannot detect if you are using VxVM/VxFS 5.0. A VxFS version 7 file system when you are running VxVM/VxFS 4.1 or below will not be mountable. If you choose to create a file system with VxFS version 7, you will see the following warning message during installation or recovery:

```
WARNING: The version keyword for the VxFS file system  
"/home" is set to "7". This requires loading the 5.0  
Base-VXFS bundle onto your system. You must ensure this is  
present otherwise your system will not boot successfully.
```

On Itanium-based servers running HP-UX 11.23 you must be very careful using vxupgrade to upgrade the VxFS version of a file system. The bootloader on Itanium-based servers does not understand VxFS version 6 file systems; upgrading the boot file system to VxFS version 6 will render the system unable to boot. Ignite-UX will refuse to make the boot volume (usually /stand) VxFS layout version 6 during a recovery - it will be forced to version 5 regardless of any configuration. If you accidentally upgrade the VxFS version of the boot file system to version 6, you will need to create a recovery archive and recover the system.

Except as noted above during a recovery, if you created a recovery archive with Ignite-UX version C.6.3 or above and you are running HP-UX 11.23 or later, Ignite-UX can preserve the VxFS file system version. To preserve the VxFS file system version use the `version` keyword. This keyword must be associated with only a VxFS file system, and accepts a number as an valid value. The `version` keyword does not accept `default` as a valid value. If you wish to use the default VxFS version for a file system do not use the `version` keyword.

How do I...

There are many times when questions arise regarding the use of Ignite-UX and the many utilities provided by the product. This section provides the answers to some of the more common questions that can occur.

How do I remove the warning message that occurs when compiling a kernel on PA-RISC systems?

When compiling a kernel in Ignite-UX on PA-RISC systems, it is likely that the following message from the linker, when it is linking the kernel, will appear:

```
/usr/ccs/bin/ld: (Warning) Linker features were used that may not be supported  
in future releases. The +vallcompatwarnings option can be used to display  
more details, and the ld(1) man page contains additional information. This  
warning can be suppressed with the +vnocompatwarnings option.
```

You can prevent this warning by adding the following to your custom configuration:

```

# Stop the kernel linker warning from printing # when compiling a new kernel
is_hppa {
    env_vars += "LD_OPTS+=vnocompatwarnings"
}

```

How do I recognize if a disk exists or not from within a configuration file?

It is easy to see if a disk exists or not within an Ignite-UX configuration file. For example:

```

( disk[0/1/0/8/0.15.0].size == 0 ) {
    error += "Disk at 0/1/0/8/0.15.0 does not exist"
}

```

The previous example configuration excerpt adds an error message if the disk does not exist. This allows you to create a configuration that will alert you when disks that should be there are not. However, this same `error` keyword will prevent an actual install from occurring; so use this only if it is an actual fatal problem.

If you want to be alerted conditionally if the disk does exist, the example changes to the following:

```

( disk[0/1/0/8/0.15.0].size != 0 ) {
    ...
}

```

You could use this logic to conditionally define a disk layout only if the disks that it refers to exist. This would allow you to create unique disk layouts for various systems that have disks at different hardware addresses.

The size of a disk that does not exist on the system should always evaluate to zero. However, some disk arrays may present zero size disk LUNs. If there is a possibility of encountering this issue then you should consider the following example as an alternative (note that a zero length LUN is not be useful in a disk layout definition):

```

( (disk[0/1/0/8/0.15.0].model ~ ".") ) {
    ...
}

```

The logic inside the inner parenthesis tests whether the model string is set or not. The `~` does an extended regular expression match and the `"."` means match (one of) any character. Using this logic ensures that there is a disk LUN present at that address because it contains a model string so the configuration within the braces is evaluated only if the disk does exist.

You can change the sense of the test by adding a logical not operator (!):

```

( ! (disk[0/1/0/8/0.15.0].model ~ ".") ) {
    ...
}

```

Outside the inner parenthesis the ! (logical not) changes the sense of the test so it becomes TRUE if no model string is reported for that hardware path. This test could be used to perform actions if the disk did not exist.

For a more complete test, you could combine the previous examples as follows:

```
( (disk[0/1/0/8/0.15.0].model ~ ".") & (disk[0/1/0/8/0.15.0].size != 0) ) {  
    init _my_var1=disk[0/1/0/8/0.15.0].model  
    init _my_var2=disk[0/1/0/8/0.15.0].size  
    note += "Disk at 0/1/0/8/0.15.0 has a model of " + ${_my_var1}  
    note += "Disk at 0/1/0/8/0.15.0 has a size of " + ${_my_var2} }
```

This test ensures that the disk at the given hardware path has a model string and a non-zero size. If either test fails, the configuration it protects will not be evaluated. On a test system, this configuration produced the following output:

```
NOTE: Disk at 0/1/0/8/0.15.0 has a model of SEAGATE_ST39173WC  
NOTE: Disk at 0/1/0/8/0.15.0 has a size of 8891556K
```

This result is because the disk did exist. The logic that you would use to see if a disk does *not* exist would be as follows:

```
( ! (disk[0/1/0/8/0.15.0].model ~ ".") & (disk[0/1/0/8/0.15.0].size ==  
0) ) {  
    ...  
}
```

How do I create the CD equivalent of a tape created by make_boot_tape?

This media can be used for several practical purposes:

1. A replacement for network booting so you do not have to setup boot helpers, DHCP, etc.
The version of Ignite-UX on the media must match the Ignite version on the Ignite server.
2. For use with tape recovery when you need to perform two-step or dual media recovery.
The version of Ignite-UX on the media must match the Ignite version on the tape.

Important:

HP does not provide support for the CD and DVD writing tools supplied with Ignite-UX. However, you can notify HP if you find defects. HP will try to report defects to the related open source communities and incorporate the appropriate fixes into future Ignite-UX releases. To provide feedback or report a defect, please open a call with the HP response center.

Please provide as much information as possible when you contact HP about a defect, as HP will not investigate problems related to the tools unless they relate to unmodified examples taken from HP documentation (allowing documentation issues to be addressed).

Important:

HP-UX does not currently support writing to USB DVD writers. This currently affects the built-in DVD writers on the rx3600 and rx6600 systems and is expected to be resolved on both HP-UX 11i v2 and 11i v3 with OE updates during 2007.

Although this question does not appear to be related to custom configuration, you can place custom configuration in the installation file systems contained on a custom installation CD that you create.

The following example instructions were developed using Ignite-UX version B.5.4. These instructions are applicable to Ignite-UX version C.6.0 and greater, however you will need to select an operating system revision using the `-r` option.

With Ignite-UX version B.5.4 this method creates a CD that can boot HP-UX 11i v1 systems and Itanium®-based systems running HP-UX 11i v2 (running OEs released before September 2004). With Ignite-UX C.6.0, you are queried for the HP-UX release so you can create a CD that supports Itanium®-based and PA-RISC systems running 11i v2; you must create a different CD to support PA-RISC systems running 11i v1.

You must install the correct Ignite-UX software to support the HP-UX revision for which you wish to create these CDs. For example, if you install the HP-UX 11i v1 bundles for Ignite-UX C.6.0, you will not be able to create a CD supporting 11i v2 since you have not installed the software supporting that release of HP-UX.

Create a LIF. If you are using Ignite-UX versions before C.6.0, enter:

```
# make_medialif -m -l ./my_lif
```

If you are using Ignite-UX versions C.6.0 and later, enter:

```
# make_medialif -m -r B.XX.YY -l ./my_lif
```

where, B.XX.YY is the release of HP-UX for which you want to create the LIF. See Column 1 in Table 2 for a list of HP-UX releases you can use and what it will allow you to install or recover.

To verify the lif contents, use the following commands:

```
# lifls -l ./my_lif
volume ISL10 data size 804676 directory size 3 04/11/15 13:57:02
filename type start size implement created
=====
ISL      -12800 16    306    0      04/11/15 13:57:03
AUTO     -12289 328   1       0      04/11/15 13:57:03
HPUX     -12928 336   848    0      04/11/15 13:57:03
FWWKAR4  BIN      1184   1       0      04/11/15 13:57:03
FWWKAR5  BIN      1192   1       0      04/11/15 13:57:03
FWWKAR6  BIN      1200   1       0      04/11/15 13:57:04
FWWKAR7  BIN      1208   1       0      04/11/15 13:57:04
FWWKAR8  BIN      1216   1       0      04/11/15 13:57:04
INSTALL   -12290 1224  68184   0      04/11/15 13:57:09
INSTALLFS -12290 69408 35840   0      04/11/15 13:57:12
VINSTALLFS -12290 69408 35840   0      04/11/15 13:57:12
```

WINSTALLFS	-12290	69408	35840	0	04/11/15 13:57:12
IINSTALL	-12290	105248	197630	0	04/11/15 13:57:26
VINSTALL	-12290	302880	73806	0	04/11/15 13:57:32
WINSTALL	-12290	376688	83837	0	04/11/15 13:57:38
IINSTALLFS	-12290	460528	116224	0	04/11/15 13:57:47
PAD	BIN	576752	256	0	04/11/15 13:57:48

If you have customizations to apply to the installation file systems, you should apply them now using `instl_adm`.

The `instl_adm` command , prior to Ignite-UX version C.6.0.x, will **not** update the Itanium®-based installation file system `IINSTALLFS` if any PA-RISC installation file systems are present (this issue is being addressed at this time; review the release notes for your version of Ignite-UX to determine if it has been resolved). If you wish to apply custom configuration to the Itanium®-based installation file system, you must perform the following commands:

```
# lifcp ./mylif:IINSTALLFS ./IINSTALLFS
# instl_adm -f output -F ./IINSTALLFS
# lifrm ./mylif:IINSTALLFS
# lifcp -r -T-12290 -K2 ./IINSTALLFS ./mylif:IINSTALLFS
```

This assumes that the configuration you wish to apply is located within the file called `output` above.

Remember that these steps are not required if the LIF does not contain any of the PA-RISC installation file systems (`INSTALLFS`, `WINSTALLFS`, or `VINSTALLFS`) or if you are using Ignite-UX C.6.0.x or greater. So with Ignite-UX C.6.0.x and greater, all of the installation file systems will be updated by `instl_adm`. If you wish to maintain different configurations in each file system, you must manually modify the Itanium-based installation file system, `IINSTALLFS` as previously described.

Tip:

If you only wish to support PA-RISC systems, you can write the LIF you just created directly onto a CD and use it to boot systems.

4. You must create a pseudo-root for our CD image; this is necessary to add support for Itanium®-based systems. Enter the following commands:

```
# mkdir top
# mv my_lif top
# cp /opt/ignite/boot/EFI_CD_image top/EFI_CD_image
# cd top
```

5. Create an ISO image that can be written to a CD. You may need to load additional software from your applications media to get the `mkisofs` command onto your system.

```
# /opt/ignite/lbin/mkisofs -no-emul-boot -b EFI_CD_image \
-eltorito-alt-boot -no-emul-boot -b my_lif -o ../cdfs.out $(pwd)
```

Messages similar to the following appear:

```

Size of boot image is 10240 sectors -> No emulation
Size of boot image is 288504 sectors -> No emulation
6.69% done, estimate finish Mon Nov 15 18:05:23 2004
13.39% done, estimate finish Mon Nov 15 18:05:31 2004
20.07% done, estimate finish Mon Nov 15 18:05:28 2004
26.77% done, estimate finish Mon Nov 15 18:05:27 2004
33.45% done, estimate finish Mon Nov 15 18:05:26 2004
40.15% done, estimate finish Mon Nov 15 18:05:26 2004
46.83% done, estimate finish Mon Nov 15 18:05:26 2004
53.54% done, estimate finish Mon Nov 15 18:05:25 2004
60.21% done, estimate finish Mon Nov 15 18:05:25 2004
66.92% done, estimate finish Mon Nov 15 18:05:25 2004
73.60% done, estimate finish Mon Nov 15 18:05:25 2004
80.30% done, estimate finish Mon Nov 15 18:05:25 2004
86.98% done, estimate finish Mon Nov 15 18:05:25 2004
93.68% done, estimate finish Mon Nov 15 18:05:25 2004
Total translation table size: 2048
Total rock ridge attributes bytes: 0
Total directory bytes: 0
Path table size(bytes): 10
Max brk space used 9000
74736 extents written (145 Mb)

```

6. Ensure that there are no `mkisofs` command errors. Older versions of `mkisofs` do not accept the `-no-emul-boot` option given twice, only accept the use of the `-b` option once, and would not accept the `-eltorito-alt-boot` option at all.
7. Place a LIF directory at the start of the ISO image with the `instl_combine` command so that the media will boot:

```
# /opt/ignite/lbin/instl_combine -C cdfs.out
```

HP Integrity systems do not require a LIF to be present to boot (PA-RISC systems do), but the kernel loader used to install Itanium®-based systems does read the installation kernel and file system from the LIF. Later in the installation process, other files will also be read from the LIF.

For more information regarding how to create custom installation media, refer to the *Ignite-UX Administration Guide*.

8. You can now write the ISO image onto a CD. The ISO image created in this example was written onto the CD using `Cdrecord` on an Itanium®-based system:

```
# cdrecord -data -v speed=4 dev=4,0,0 /var/tmp/cdfs.out
```

It is important to realize that some Windows-based CD writing software cannot create CDs that will boot on Itanium®-based systems. The CDs produced by some CD writing software may boot PA-RISC systems, but fail to boot Itanium®-based systems.

This CD, like a tape created by `make_boot_tape`, allows you to boot a system locally so that it can contact an Ignite-UX server. In addition, this media is capable of performing a dual-media recovery. For more information regarding dual-media recovery, refer to the *Ignite-UX Administration Guide*.

How do I enable the X server and CDE during a golden image install?

The information provided here is intended only for golden image installs and not for SD-UX based installations. You should *not* use these commands during an SD-UX based installation. It is also assumed that you are using a version of Ignite-UX recent enough to support the graphics card in question (for newer graphics cards), and that the golden image has the required products and patches loaded to support the graphics card and input device.

The following configuration can be placed into a custom configuration file to enable CDE:

```
( graphics[ ].model ~ ".*" ) {
    post_configure_cmd+="cp -p /usr/dt/config/Xservers
/etc/dt/config/Xservers"
    post_configure_cmd+="/usr/dt/bin/dtconfig -e"
} else {
    post_configure_cmd+="/usr/dt/bin/dtconfig -d"
}
```

When a golden image is created on a system without a graphics card, CDE and the X server are disabled from starting. When the golden image is then loaded onto a system with a graphics card, CDE and the X server still won't start by default. The reason for this is that CDE disabled the X server from starting on the system without a graphics card where the golden image was created, and CDE then never checks again to see if any other system where the golden image is loaded has a graphics card.

CDE disables the X server by commenting out a line in `/etc/dt/config/Xservers`. Therefore, replace this file with `/usr/dt/config/Xservers` to make CDE check to see if the local X server can be started on its next attempt at starting up. During a golden image installation, this allows you to automatically enable the X server so CDE will start it after installation.

Please note the following limitation, though. The above configuration assumes that the commands above are all that are needed to enable the X server. The contents of the configuration files in `/etc/X11` may not be compatible with the newly installed system if they have been customized. You may need to replace files in `/etc/X11` to successfully get the X server started and working as desired. How to configure files in `/etc/X11` is beyond the scope of this document, however. For information on how to configure the X server files in `/etc/X11`, please review the SD-UX configuration scripts associated with the X server filesets located under `/var/adm/sw/products`.

Summary

A firm understanding of how to design, maintain, test and troubleshoot configuration files is essential in reducing business resources. With careful consideration and planning, you can create custom configuration files that improve mass deployments and system standardization.

For more information

The following relevant documents are available online at the HP Technical Documentation Web site at <http://www.docs.hp.com/>:

Ignite-UX Administration Guide

Successful System Cloning Using Ignite-UX

Managing HP-UX Software With SD-UX

Software Distributor Administration Guide for HP-UX 11i

HP-UX 11.11 Installation and Update Guide

Release Notes for HP-UX 11.11

HP-UX 11i v[1|2|3] Installation and Update Guide

HP-UX 11i v[1|2|3] Release Notes

Managing Systems and Workgroups: A Guide for HP-UX System Administrators

Additionally, there are a number of white papers pertinent to Ignite-UX that are located at the HP Technical Documentation Web site.

Product information regarding Ignite-UX is available at

<http://www.docs.hp.com/en/IUX/>

© Copyright 2004-2007, 2009 Hewlett-Packard Development Company, L.P.

Confidential computer software. Valid license from HP required for possession, use or copying. Consistent with FAR 12.211 and 12.212, Commercial Computer Software, Computer Software Documentation, and Technical Data for Commercial Items are licensed to the U.S. Government under vendor's standard commercial license.

The information contained herein is subject to change without notice. The only warranties for HP products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. HP shall not be liable for technical or editorial errors or omissions contained herein.

Intel Itanium® Logo, Intel, Intel Inside and Itanium are trademarks or registered trademarks of Intel Corporation in the US and other countries and are used under license.

Intel® Itanium® Processor Family is a trademark of Intel Corporation in the US and other countries and is used under license

MPN 5992-1999, 4/2009